

# **CIS 450**

## **Computer Architecture and Organization**

### **Lecture 15: Code Optimization II**

**Mitch Neilsen**  
[neilsen@ksu.edu](mailto:neilsen@ksu.edu)  
**219D Nichols Hall**

# Topics

## Machine-Independent Optimizations

- Optimization Blockers

## Machine Dependent Optimizations

- Understanding Processor Operations
- Branches and Branch Prediction

# Getting High Performance

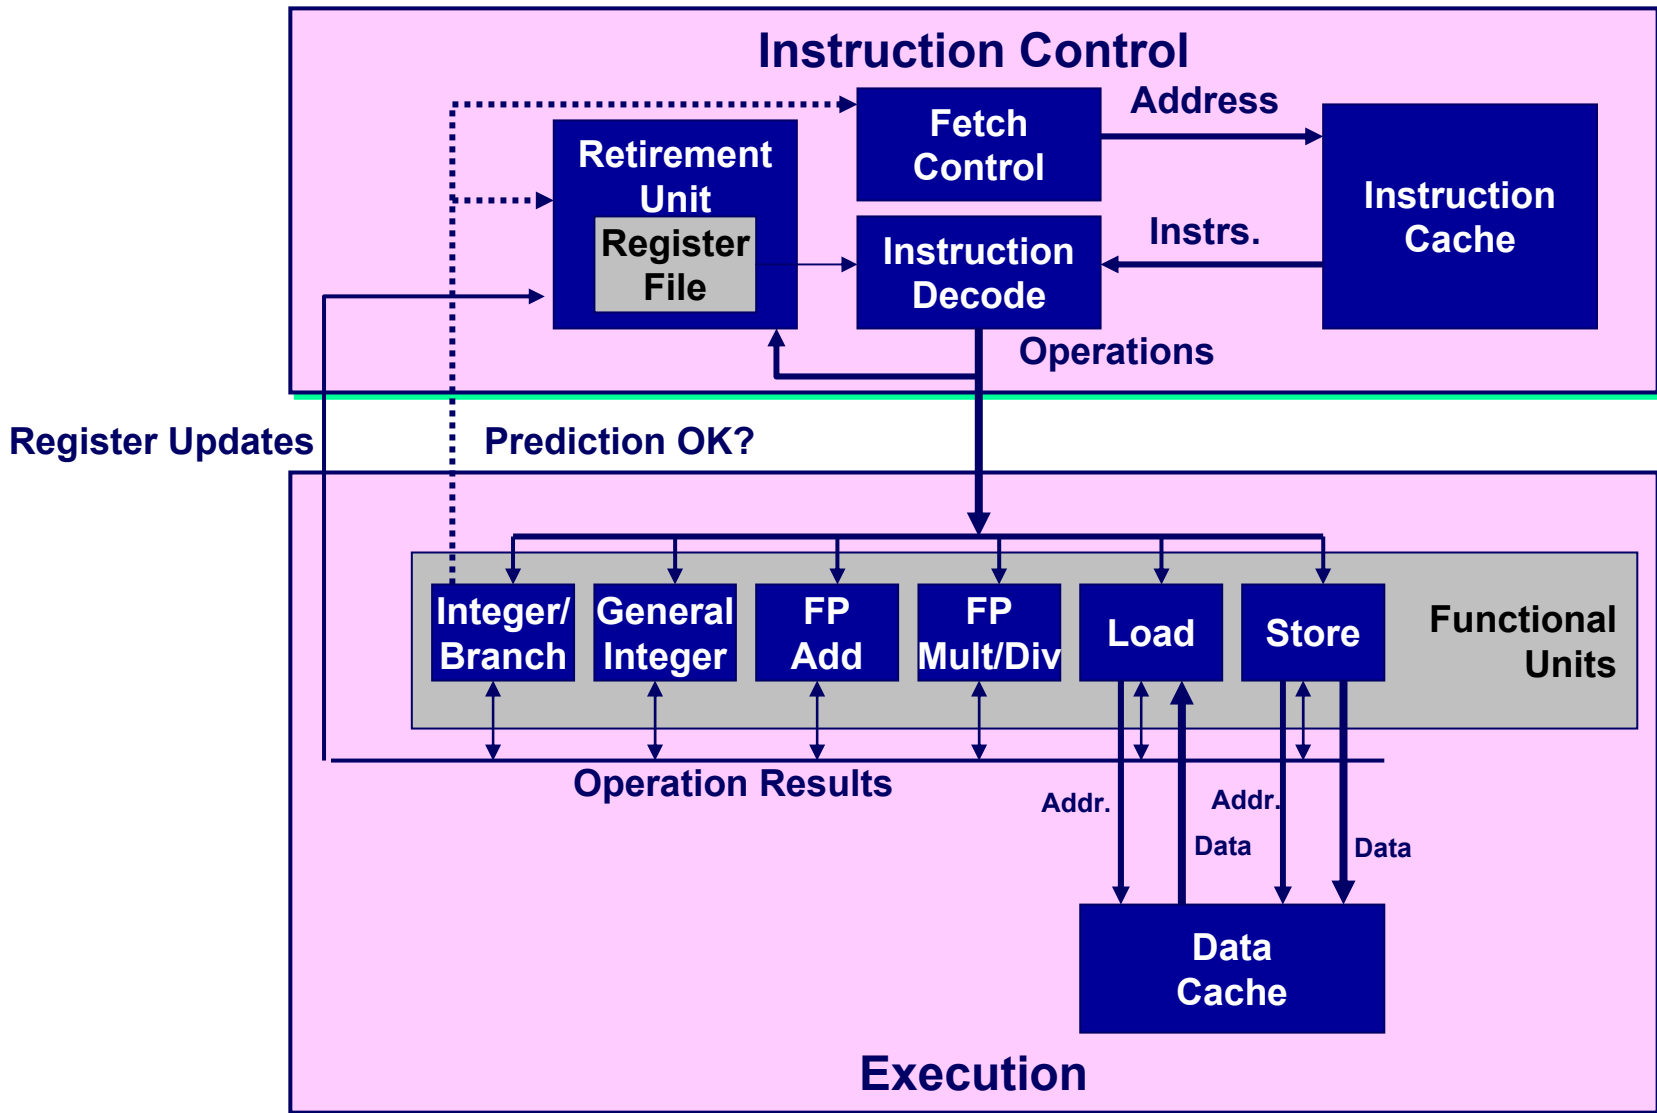
## Don't Do Anything Stupid

- Watch out for hidden algorithmic inefficiencies
- Write compiler-friendly code
  - Help compiler past optimization blockers: function calls & memory references.

## Tune Code For Machine

- Exploit instruction-level parallelism
- Avoid unpredictable branches
- Make code cache friendly
  - Covered later in course

# Modern CPU Design



# CPU Capabilities of Pentium IV


## Multiple Instructions Can Execute in Parallel

- 1 load, with address computation
- 1 store, with address computation
- 2 simple integer (one may be branch)
- 1 complex integer (multiply/divide)
- 1 FP/SSE3 unit
- 1 FP move (does all conversions)

## Some Instructions Take > 1 Cycle, but Can be Pipelined

■ Instruction	Latency	Cycles/Issue
■ Load / Store	5	1
■ Integer Multiply	10	1
■ Integer/Long Divide	36/106	36/106
■ Single/Double FP Multiply	7	2
■ Single/Double FP Add	5	2
■ Single/Double FP Divide	32/46	32/46

# Instruction Control



## Grabs Instruction Bytes From Memory

- Based on current PC + predicted targets for predicted branches
- Hardware dynamically guesses whether branches taken/not taken and (possibly) branch target

## Translates Instructions Into *Operations* (for CISC style CPUs)

- Primitive steps required to perform instruction
- Typical instruction requires 1–3 operations

## Converts Register References Into *Tags*

- Abstract identifier linking destination of one operation with sources of later operations

# Translating into Operations

**Goal: Each Operation Utilizes Single Functional Unit**

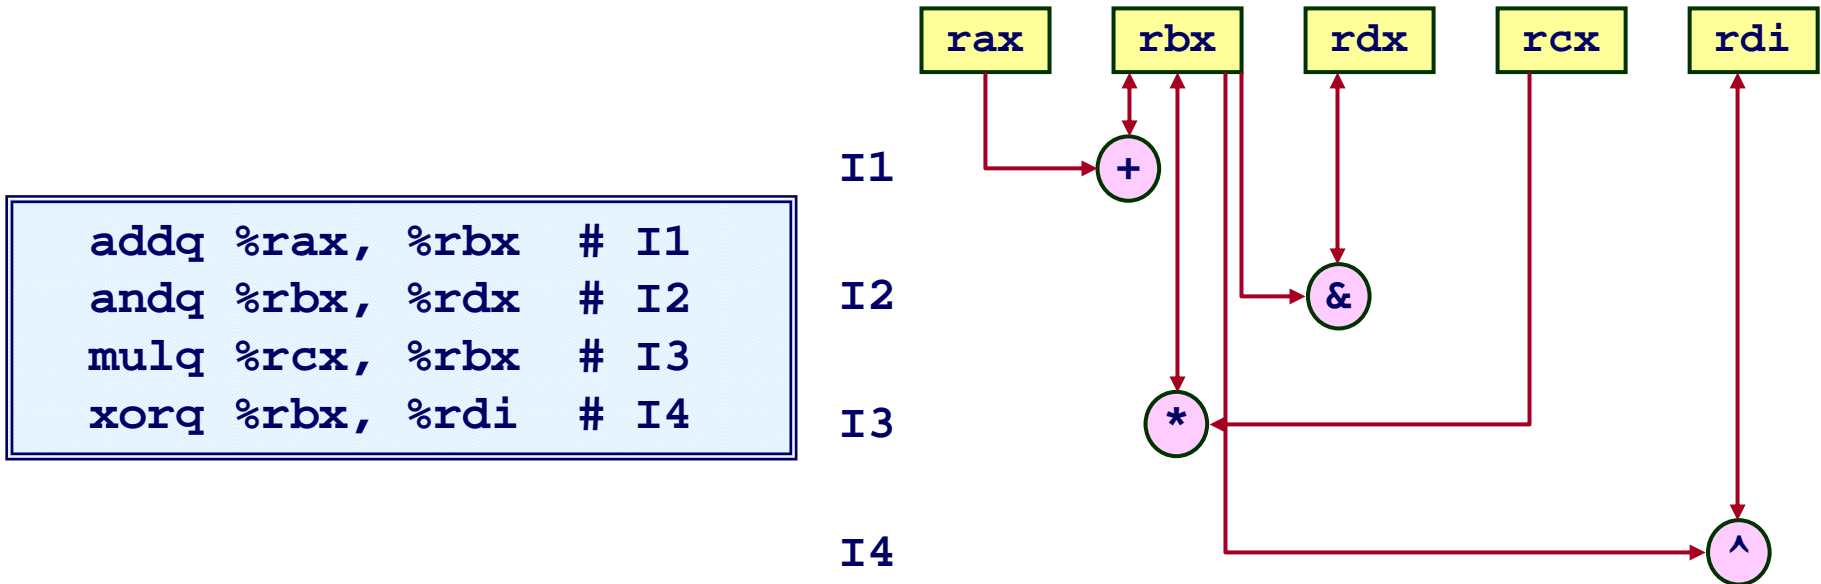
```
addq %rax, 8(%rbx,%rdx,4)
```

- Requires: Load, Integer arithmetic, Store

load 8(%rbx,%rdx,4)	→	temp1
add %rax, temp1	→	temp2
store temp2, 8(%rbx,%rdx,4)		

- Exact form and format of operations is trade secret
- Operations: split up instruction into simpler pieces
- Devise temporary names to describe how result of one operation gets used by other operations

# Traditional View of Instruction Execution



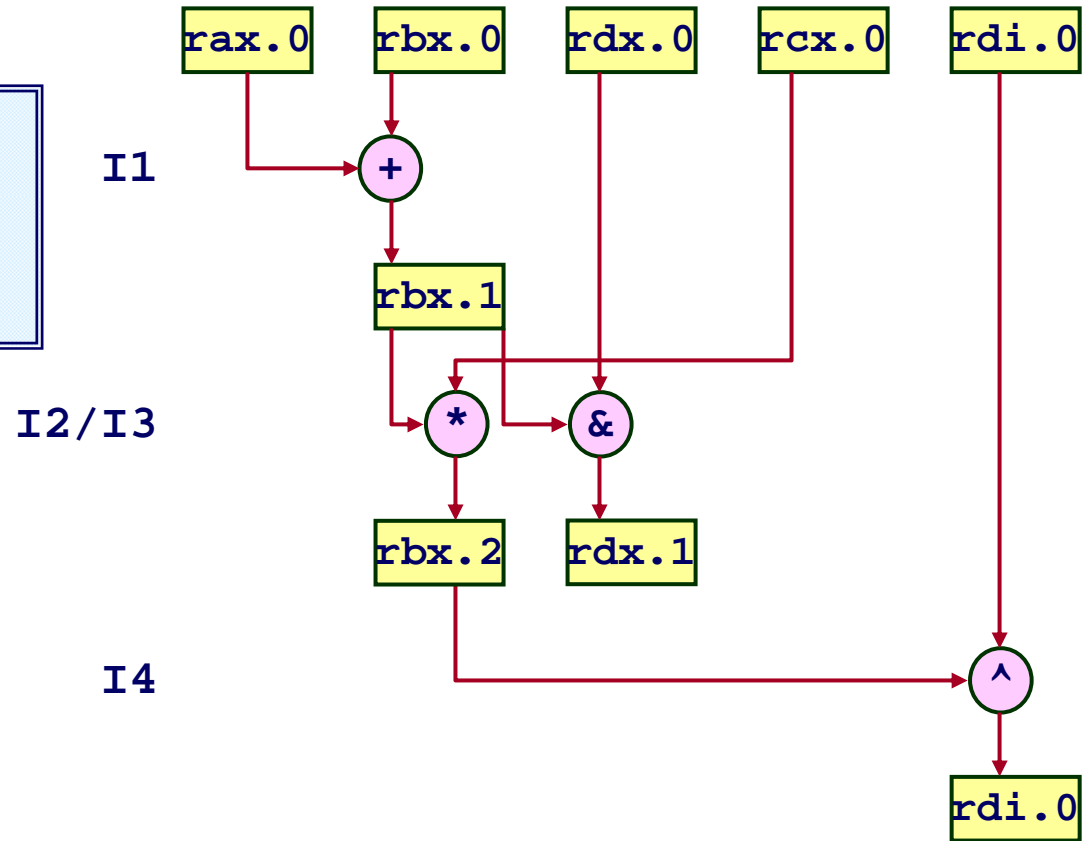
## Imperative View

- Registers are fixed storage locations
  - Individual instructions read & write them
- Instructions must be executed in specified sequence to guarantee proper program behavior



# Dataflow View of Instruction Execution

```
addq %rax, %rbx # I1
andq %rbx, %rdx # I2
mulq %rcx, %rbx # I3
xorq %rbx, %rdi # I4
```



## Functional View

- View each write as creating new instance of value
- Operations can be performed as soon as operands available
- No need to execute in original sequence

# Example Computation

```
void combine4(vec_ptr v, data_t *dest)
{
    int i;
    int length = vec_length(v);
    data_t *d = get_vec_start(v);
    data_t t = IDENT;
    for (i = 0; i < length; i++)
        t = t OP d[i];
    *dest = t;
}
```

## Data Types


- Use different declarations for data\_t
- int
- float
- double

## Operations

- Use different definitions of OP and IDENT
- + / 0
- \* / 1

# Cycles Per Element

- Convenient way to express performance of program that operates on vectors or lists
- Length =  $n$
- $T = CPE * n + \text{Overhead}$



# x86-64 Compilation of Combine4

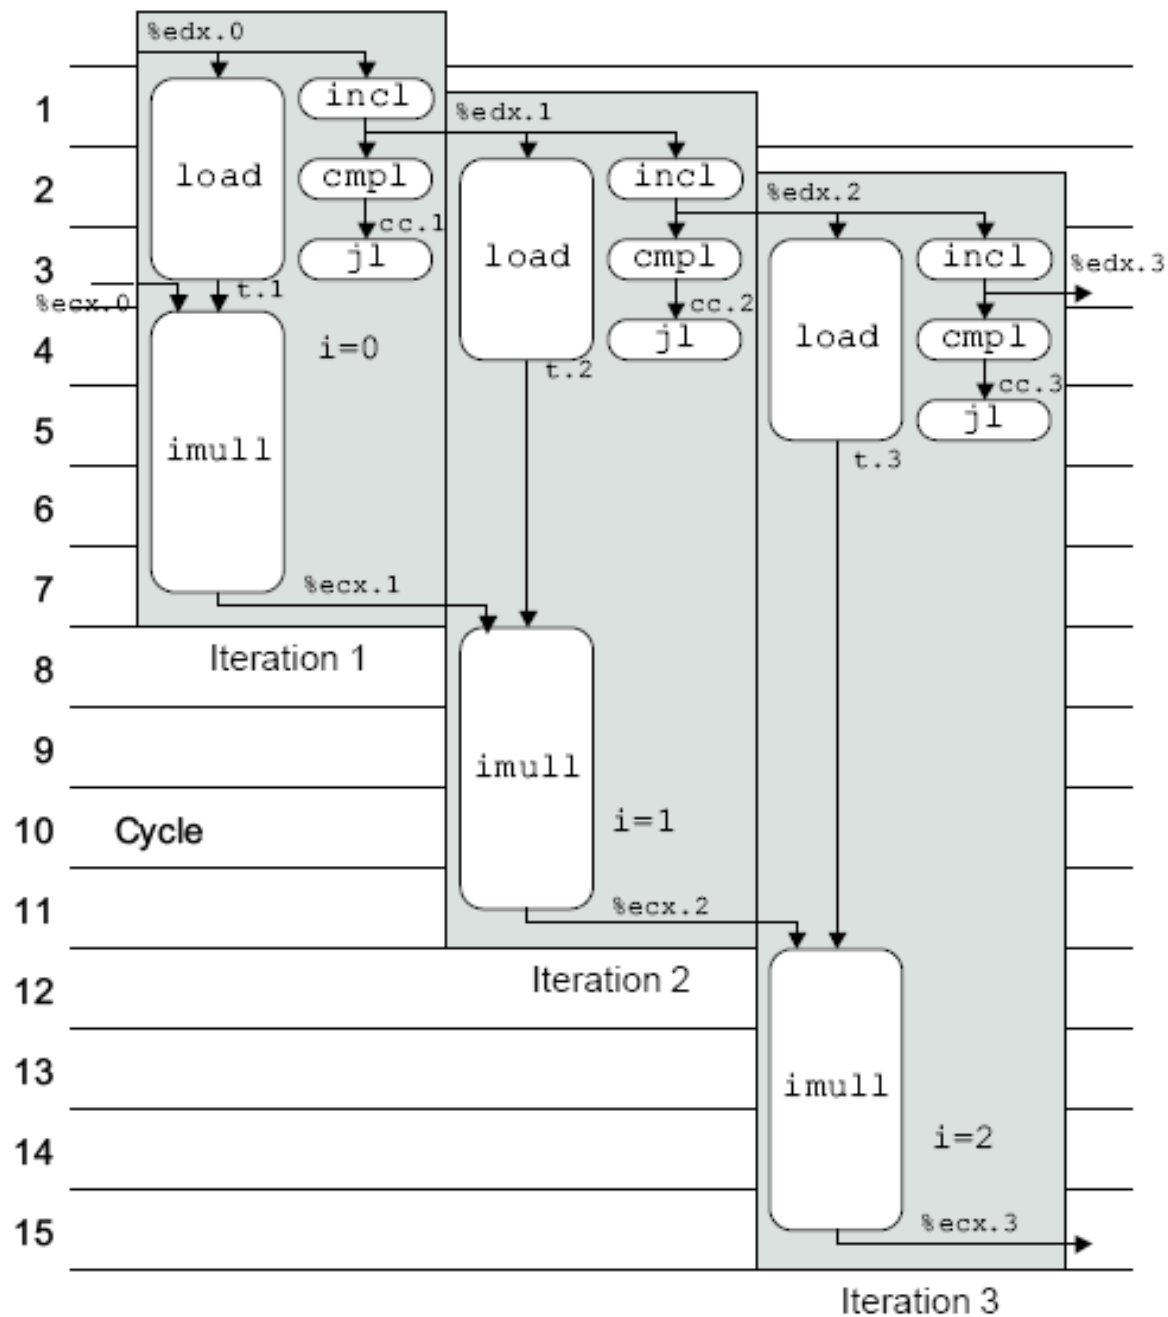
## Inner Loop (Integer Multiply)

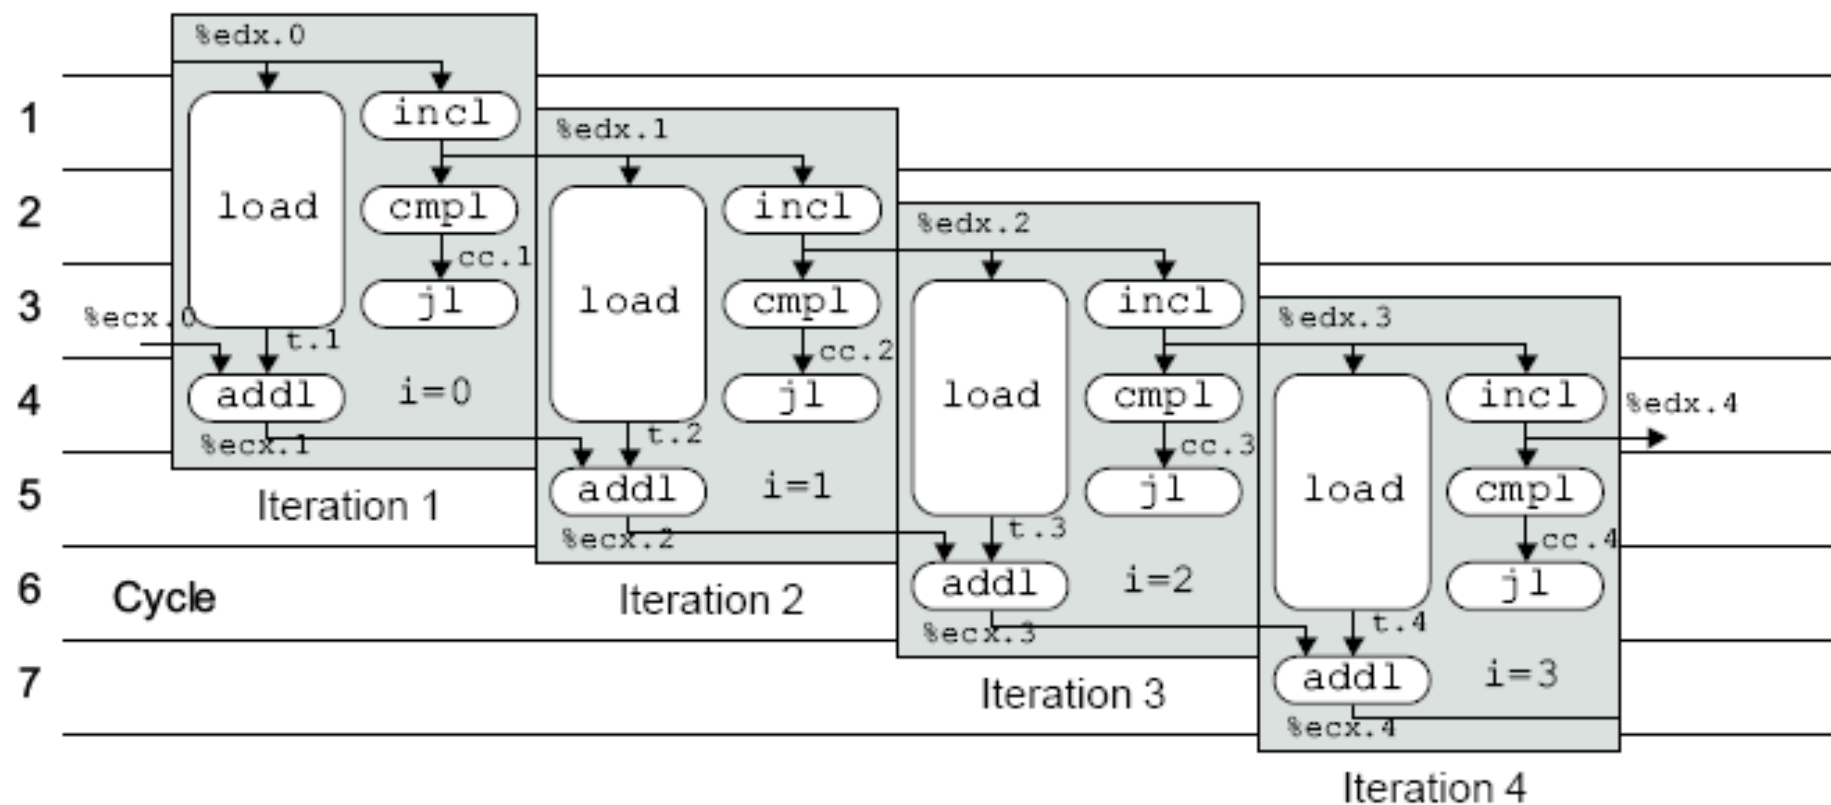
```
L33:                                # Loop:
    movl    (%eax,%edx,4), %ebx    # temp = d[i]
    incl    %edx                    # i++
    imull   %ebx, %ecx              # x *= temp
    cmpl    %esi, %edx              # i:length
    jl      L33                    # if < goto Loop
```

## Performance

- 5 instructions in 2.2 clock cycles for integer addition, in 10 clock cycles for integer multiplication:

Method	Integer		Floating Point	
Combine4	2.20	10.00	5.00	7.00





# Serial Computation

## Computation (length=12)

```
(((((1 * d[0]) *
d[1]) * d[2]) * d[3]) *
d[4]) * d[5]) * d[6]) *
d[7]) * d[8]) * d[9]) *
d[10]) * d[11])
```

## Performance

- N elements, D cycles/operation
- N\*D cycles

Method	Integer		Floating Point	
	Addition	Multiplication	Addition	Multiplication
Combine4	2.20	10.00	5.00	7.00

Addition Multiplication Addition Multiplication

# Loop Unrolling

```
void unroll2a_combine(vec_ptr v, data_t *dest)
{
    int length = vec_length(v);
    int limit = length-1;
    data_t *d = get_vec_start(v);
    data_t x = IDENT;
    int i;
    /* Combine 2 elements at a time */
    for (i = 0; i < limit; i+=2) {
        x = (x OPER d[i]) OPER d[i+1];
    }
    /* Finish any remaining elements */
    for (; i < length; i++) {
        x = x OPER d[i];
    }
    *dest = x;
}
```

- Perform 2x more useful work per iteration



# Effect of Loop Unrolling

Method	Integer		Floating Point	
Combine4	2.20	10.00	5.00	7.00
Unroll 2	1.50	10.00	5.00	7.00

## Helps Integer Sum

- Before: 5 operations per element
- After: 6 operations per 2 elements
  - = 3 operations per element

## Others Don't Improve

- Sequential dependency
  - Each operation must wait until previous one completes

```
x = (x OPER d[i]) OPER d[i+1];
```

# Loop Unrolling with Reassociation

```
void unroll2aa_combine(vec_ptr v, data_t *dest)
{
    int length = vec_length(v);
    int limit = length-1;
    data_t *d = get_vec_start(v);
    data_t x = IDENT;
    int i;
    /* Combine 2 elements at a time */
    for (i = 0; i < limit; i+=2) {
        x = x OPER (d[i] OPER d[i+1]);
    }
    /* Finish any remaining elements */
    for (; i < length; i++) {
        x = x OPER d[i];
    }
    *dest = x;
}
```

- Could change numerical results for FP

# Effect of Reassociation

Method	Integer		Floating Point	
Combine4	2.20	10.00	5.00	7.00
Unroll 2	1.50	10.00	5.00	7.00
2 X 2 reassociate	1.56	5.00	2.75	3.62

**Nearly 2X speedup for Int \*, FP +, FP \***

- Breaks sequential dependency


```
x = x OPER (d[i] OPER d[i+1]);
```

- While computing result for iteration i, can precompute  $d[i+2]*d[i+3]$  for iteration i+2

# Reassociated Computation

## Performance

- N elements, D cycles/operation
- Should be  $(N/2+1)*D$  cycles
  - CPE =  $D/2$
- Measured CPE slightly worse for FP



```
x = x OPER (d[i] OPER d[i+1]);
```

# Loop Unrolling with Separate Accum.

```
void unroll2a_combine(vec_ptr v, data_t *dest)
{
    int length = vec_length(v);
    int limit = length-1;
    data_t *d = get_vec_start(v);
    data_t x0 = IDENT;
    data_t x1 = IDENT;
    int i;
    /* Combine 2 elements at a time */
    for (i = 0; i < limit; i+=2) {
        x0 = x0 OPER d[i];
        x1 = x1 OPER d[i+1];
    }
    /* Finish any remaining elements */
    for (; i < length; i++) {
        x0 = x0 OPER d[i];
    }
    *dest = x0 OPER x1;
}
```

## ■ Different form of reassociation

# Effect of Reassociation

Method	Integer		Floating Point	
Combine4	2.20	10.00	5.00	7.00
Unroll 2	1.50	10.00	5.00	7.00
2 X 2 reassociate	1.56	5.00	2.75	3.62
2 X 2 separate accum.	1.50	5.00	2.50	3.50

**Nearly 2X speedup for Int \*, FP +, FP \***

- Breaks sequential dependency


```
x0 = x0 OPER d[i];
x1 = x1 OPER d[i+1];
```

- Computation of even elements independent of odd ones

# Separate Accum. Computation

## Performance

- N elements, D cycles/operation
- Should be  $(N/2+1)*D$  cycles
  - CPE =  $D/2$
- Measured CPE matches prediction!



```
x0 = x0 OPER d[i];  
x1 = x1 OPER d[i+1];
```

# Unrolling & Accumulating

## Idea

- Can unroll to any degree  $L$
- Can accumulate  $K$  results in parallel
- $L$  must be multiple of  $K$

## Limitations

- Diminishing returns
  - Cannot go beyond pipelining limitations of execution units
- Large overhead
  - Finish off iterations sequentially
  - Especially for shorter lengths



# Unrolling & Accumulating: Intel FP \*

## Case

- Intel Nocoma
- FP Multiplication
- Theoretical Limit: 2.00

FP *	Unrolling Factor L							
K	1	2	3	4	6	8	10	12
1	7.00	7.00		7.01		7.00		
2		3.50		3.50		3.50		
3			2.34					
4				2.01		2.00		
6					2.00			2.01
8						2.01		
10							2.00	
12								2.00

# Unrolling & Accumulating: Intel FP +

## Case

- Intel Nocoma
- FP Addition
- Theoretical Limit: 2.00

FP +	Unrolling Factor L							
K	1	2	3	4	6	8	10	12
1	5.00	5.00		5.02		5.00		
2		2.50		2.51		2.51		
3			2.00					
4				2.01		2.00		
6					2.00			1.99
8						2.01		
10							2.00	
12								2.00

# Unrolling & Accumulating: Intel Int \*

## Case

- Intel Nocoma
- Integer Multiplication
- Theoretical Limit: 1.00

Int *	Unrolling Factor L							
K	1	2	3	4	6	8	10	12
1	10.00	10.00		10.00		10.01		
2		5.00		5.01		5.00		
3			3.33					
4				2.50		2.51		
6					1.67			1.67
8						1.25		
10							1.09	
12								1.14

# Unrolling & Accumulating: Intel Int +

## Case

- Intel Nocoma
- Integer addition
- Theoretical Limit: 1.00 (unrolling enough)

Int +	Unrolling Factor L							
K	1	2	3	4	6	8	10	12
1	2.20	1.50		1.10		1.03		
2		1.50		1.10		1.03		
3			1.34					
4				1.09		1.03		
6					1.01			1.01
8						1.03		
10							1.04	
12								1.11

# Intel vs. AMD

## FP \*

FP *	Unrolling Factor L							
K	1	2	3	4	6	8	10	12
1	7.00	7.00		7.01		7.00		
2		3.50		3.50		3.50		
3			2.34					
4				2.01		2.00		
6					2.00			2.01
8						2.01		
10							2.00	
12								2.00

FP *	Unrolling Factor L							
K	1	2	3	4	6	8	10	12
1	4.00	4.00		4.00		4.01		
2		2.00		2.00		2.00		
3			1.34					
4				1.00		1.00		
6					1.00			1.00
8						1.00		
10							1.00	
12								1.00

## Machines

- Intel Nocomoa
  - 3.2 GHz
- AMD Opteron
  - 2.0 GHz

## Performance

- AMD lower latency & better pipelining
- But slower clock rate

# Intel vs. AMD

Int \*

## Performance

- AMD multiplier much lower latency
- Can get high performance with less work
- Doesn't achieve as good an optimum

Int *	Unrolling Factor L							
K	1	2	3	4	6	8	10	12
1	10.00	10.00		10.00		10.01		
2		5.00		5.01		5.00		
3			3.33					
4				2.50		2.51		
6					1.67			1.67
8						1.25		
10							1.09	
12								1.14

Int *	Unrolling Factor L							
K	1	2	3	4	6	8	10	12
1	3.00	3.00		3.00		3.00		
2		2.33		2.0		1.35		
3			2.00					
4				1.75		1.38		
6					1.50			1.50
8						1.75		
10							1.30	
12								1.33

# Intel vs. AMD

## Int +

Int +	Unrolling Factor L							
K	1	2	3	4	6	8	10	12
1	2.20	1.50		1.10		1.03		
2		1.50		1.10		1.03		
3			1.34					
4				1.09		1.03		
6					1.01			1.01
8						1.03		
10							1.04	
12								1.11

Int +	Unrolling Factor L							
K	1	2	3	4	6	8	10	12
1	2.32	1.50		0.75		0.63		
2		1.50		0.83		0.63		
3			1.00					
4				1.00		0.63		
6					0.83			0.67
8						0.63		
10							0.60	
12								0.85

## Performance

- AMD gets below 1.0
- Even just with unrolling

## Explanation

- Both Intel & AMD can “double pump” integer units
- Only AMD can load two elements / cycle

# Streaming SIMD Extensions (SSE3)

SSE3, also known by its Intel code name Prescott New Instructions (PNI), is the third iteration of the SSE instruction set for the IA-32 architecture. Intel introduced SSE3 in early 2004 with the Prescott revision of their Pentium 4 CPU. In April 2005, AMD introduced a subset of SSE3 in revision E (Venice and San Diego) of their Athlon 64 CPUs. The earlier SIMD instruction sets on the x86 platform, from oldest to newest, are MMX, 3DNow! (developed by AMD), SSE and SSE2.


-- from Wikipedia



# Programming with SSE3

## XMM Registers

- 16 total, each 16 bytes
- 16 single-byte integers



- 8 16-bit integers



- 4 32-bit integers




- 4 single-precision floats



- 2 double-precision floats



- 1 single-precision float




- 1 double-precision float



# Scalar & SIMD Operations


## ■ Scalar Operations: Single Precision

`addss %xmm0,%xmm1`




## ■ SIMD Operations: Single Precision

`addps %xmm0,%xmm1`



## ■ SIMD Operations: Double Precision

`addpd %xmm0,%xmm1`



# Getting GCC to Use SIMD Operations

## Declarations

```
typedef float vec_t __attribute__((mode(V4SF)));  
typedef union {  
    vec_t v;  
    float d[4];  
} pack_t
```

## Accessing Vector Elements

```
pack_t xfer;  
vec_t accum;  
for (i = 0; i < 4; i++)  
    xfer.d[i] = IDENT;  
accum = xfer.v;
```

## Invoking SIMD Operations

```
vec_t chunk = *((vec_t *) d);  
accum = accum OPER chunk;
```

# Implementing Combine

```
void SSEx1_combine(vec_ptr v, float *dest)
{
    pack_t xfer;
    vec_t accum;
    float *d = get_vec_start(v);
    int cnt = vec_length(v);
    float result = IDENT;

    /* Initialize vector of 4 accumulators */

    /* Step until d aligned to multiple of 16 */

    /* Use packed operations with 4X parallelism */

    /* Single step to finish vector */

    /* Combine accumulators */
}
```

# Getting Started

## Create Vector of 4 Accumulators

```
/* Initialize vector of 4 accumulators */  
int i;  
for (i = 0; i < 4; i++)  
    xfer.d[i] = IDENT;  
accum = xfer.v;
```

## Single Step to Meet Alignment Requirements

- Memory address of vector must be multiple of 16

```
/* Step until d aligned to multiple of 16 */  
while (((long) d)%16 && cnt) {  
    result = result OPER *d++;  
    cnt--;  
}
```

# SIMD Loop

- Similar to 4-way loop unrolling
- Express with single arithmetic operation
  - Translates into single `addps` or `mulps` instruction

```
/* Use packed operations with 4X parallelism */  
while (cnt >= 4) {  
    vec_t chunk = *((vec_t *) d);  
    accum = accum OPER chunk;  
    d += 4;  
    cnt -= 4;  
}
```

# Completion

## Finish Off Final Elements

- Similar to standard unrolling

```
/* Single step to finish vector */  
while (cnt) {  
    result = result OPER *d++;  
    cnt--;  
}
```

## Combine Accumulators

- Use union to reference individual elements

```
/* Combine accumulators */  
xfer.v = accum;  
for (i = 0; i < 4; i++)  
    result = result OPER xfer.d[i];  
*dest = result;
```

# SIMD Results

## Intel Nocoma

	Unrolling Factor L			
	4	8	16	32
FP +	1.25	0.82	0.50	0.58
FP *	1.90	1.24	0.90	0.57
Int +	0.84	0.70	0.51	0.58
Int *	39.09	37.65	36.75	37.44

## AMD Opteron

	Unrolling Factor L			
	4	8	16	32
FP +	1.00	0.50	0.50	0.50
FP *	1.00	0.50	0.50	0.50
Int +	0.75	0.38	0.28	0.27
Int *	9.40	8.63	9.32	9.12

## Results

- FP approaches theoretical optimum of 0.50
- Int + shows speed up
- For int \*, compiler does not generate SIMD code

## Portability

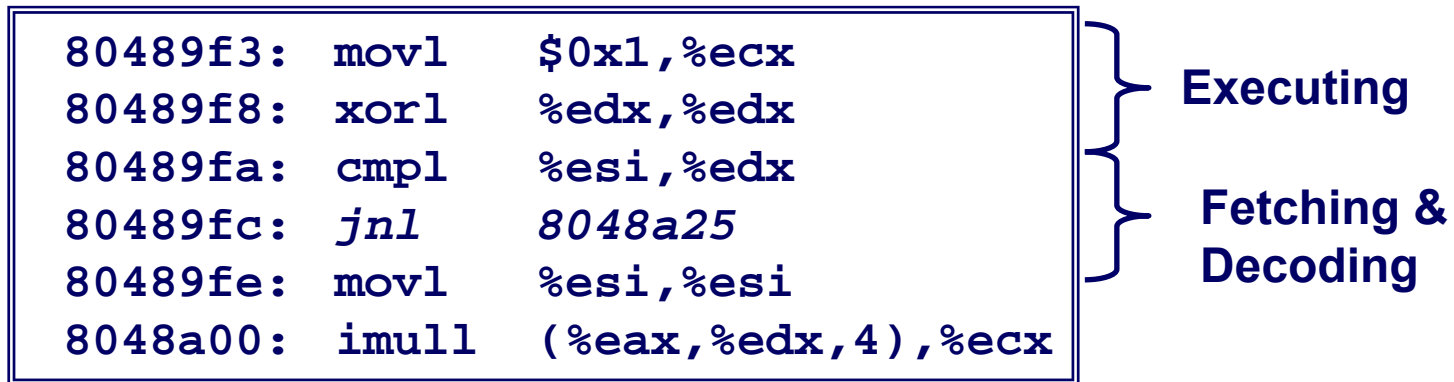
- GCC can target other machines with this code
  - Altivec instructions for PowerPC



# What About Branches?

## Challenge

- Instruction Control Unit must work well ahead of Exec. Unit
  - To generate enough operations to keep EU busy



- When encounters conditional branch, cannot reliably determine where to continue fetching

# Branch Outcomes

- When encounter conditional branch, cannot determine where to continue fetching
  - Branch Taken: Transfer control to branch target
  - Branch Not-Taken: Continue with next instruction in sequence
- Cannot resolve until outcome determined by branch/integer unit

```
80489f3: movl    $0x1,%ecx
80489f8: xorl    %edx,%edx
80489fa: cmpl    %esi,%edx
80489fc: jnl     8048a25
80489fe: movl    %esi,%esi
8048a00: imull   (%eax,%edx,4),%ecx
```

Branch Not-Taken

Branch Taken

```
8048a25: cmpl    %edi,%edx
8048a27: j1      8048a20
8048a29: movl    0xc(%ebp),%eax
8048a2c: leal    0xffffffff8(%ebp),%esp
8048a2f: movl    %ecx,(%eax)
```

# Branch Prediction

## Idea

- Guess which way branch will go
- Begin executing instructions at predicted position
  - But don't actually modify register or memory data

```
80489f3: movl    $0x1,%ecx
80489f8: xorl    %edx,%edx
80489fa: cmpl    %esi,%edx
80489fc: jnl     8048a25
. . .
```

Predict Taken

```
8048a25: cmpl    %edi,%edx
8048a27: jl      8048a20
8048a29: movl    0xc(%ebp),%eax
8048a2c: leal    0xffffffff8(%ebp),%esp
8048a2f: movl    %ecx,(%eax)
```

Begin  
Execution

# Branch Prediction Through Loop

```
80488b1:  movl    (%ecx,%edx,4),%eax
80488b4:  addl    %eax,(%edi)
80488b6:  incl    %edx
80488b7:  cmpl    %esi,%edx      i = 98
80488b9:  jl      80488b1
```

*Assume vector length = 100*

Predict Taken (OK)

```
80488b1:  movl    (%ecx,%edx,4),%eax
80488b4:  addl    %eax,(%edi)
80488b6:  incl    %edx
80488b7:  cmpl    %esi,%edx      i = 99
80488b9:  jl      80488b1
```

Predict Taken  
(Oops)

```
80488b1:  movl    (%ecx,%edx,4),%eax
80488b4:  addl    %eax,(%edi)
80488b6:  incl    %edx
80488b7:  cmpl    %esi,%edx      i = 100
80488b9:  jl      80488b1
```

Read  
invalid  
location

```
80488b1:  movl    (%ecx,%edx,4),%eax
80488b4:  addl    %eax,(%edi)
80488b6:  incl    %edx
80488b7:  cmpl    %esi,%edx      i = 101
80488b9:  jl      80488b1
```



# Branch Misprediction Invalidation

```
80488b1:  movl    (%ecx,%edx,4),%eax
80488b4:  addl    %eax,(%edi)
80488b6:  incl    %edx
80488b7:  cmpl    %esi,%edx      i = 98
80488b9:  jl      80488b1
```

*Assume vector length = 100*

Predict Taken (OK)

```
80488b1:  movl    (%ecx,%edx,4),%eax
80488b4:  addl    %eax,(%edi)
80488b6:  incl    %edx
80488b7:  cmpl    %esi,%edx      i = 99
80488b9:  jl      80488b1
```

Predict Taken (Oops)

```
80488b1:  movl    (%ecx,%edx,4),%eax
80488b4:  addl    %eax,(%edi)
80488b6:  incl    %edx
80488b7:  cmpl    %esi,%edx      i = 100
80488b9:  jl      80488b1
```

Invalidate

```
80488b1:  movl    (%ecx,%edx,4),%eax
80488b4:  addl    %eax,(%edi)
80488b6:  incl    %edx
80488b7:  cmpl    %esi,%edx      i = 101
```

# Branch Misprediction Recovery

*Assume vector length = 100*

```
80488b1:    movl    (%ecx,%edx,4),%eax
80488b4:    addl    %eax,(%edi)
80488b6:    incl    %edx
80488b7:    cmpl    %esi,%edx      i = 99
80488b9:    jl      80488b1
80488bb:    leal    0xffffffffe8(%ebp),%esp
80488be:    popl    %ebx
80488bf:    popl    %esi
80488c0:    popl    %edi
```

**Definitely not taken**

## Performance Cost

- Multiple clock cycles on modern processor
- One of the major performance limiters

# Determining Misprediction Penalty

```
int cnt_gt = 0;
int cnt_le = 0;
int cnt_all = 0;

int choose_cmov(int x, int y)
{
    int result;
    if (x > y) {
        result = cnt_gt;
    } else {
        result = cnt_le;
    }
    ++cnt_all;
    return result;
}
```

## **GCC/x86-64 Tries to Minimize Use of Branches**

- **Generates conditional moves when possible/sensible**

```
choose_cmov:
    cmpl    %esi, %edi          # x:y
    movl    cnt_le(%rip), %eax  # r = cnt_le
    cmovg   cnt_gt(%rip), %eax  # if >= r=cnt_gt
    incl    cnt_all(%rip)       # cnt_all++
    ret                                # return r
```

# Forcing Conditional

```
int cnt_gt = 0;
int cnt_le = 0;

int choose_cond(int x, int y)
{
    int result;
    if (x > y) {
        result = ++cnt_gt;
    } else {
        result = ++cnt_le;
    }
    return result;
}
```

- Cannot use conditional move when either outcome has side effect

```
choose_cond:
    cmpl    %esi, %edi
    jle     .L8
    movl    cnt_gt(%rip), %eax
    incl    %eax
    movl    %eax, cnt_gt(%rip)
    ret
.L8:
    movl    cnt_le(%rip), %eax
    incl    %eax
    movl    %eax, cnt_le(%rip)
    ret
```

If

Then

Else



# Testing Methodology

## Idea

- Measure procedure under two different prediction probabilities
  - $P = 1.0$ : Perfect prediction
  - $P = 0.5$ : Random data

## Test Data

- $x = 0, y = \pm 1$

Case +1:  $y = [+1, +1, +1, \dots, +1, +1]$

Case -1:  $y = [-1, -1, -1, \dots, -1, -1]$

Case A:  $y = [+1, -1, +1, \dots, +1, -1]$

Case R:  $y = [+1, -1, -1, \dots, -1, +1]$  (Random)

# Testing Outcomes

Intel Nocoma

Case	cmov	cond
+1	12.3	18.2
-1	12.3	12.2
A	12.3	15.2
R	12.3	31.2

AMD Opteron

Case	cmov	cond
+1	8.05	10.1
-1	8.05	8.1
A	8.05	9.2
R	8.05	15.7

## Observations:

- Conditional move insensitive to data
- Perfect prediction for regular patterns
  - But, else case requires 6 (Intel) or 2 (AMD) additional cycles
  - Averages to 15.2
- Branch penalties:
  - Intel  $2 * (31.2 - 15.2) = 32$  cycles
  - AMD  $2 * (15.7 - 9.2) = 13$  cycles

# Role of Programmer

*How should I write my programs, given that I have a good, optimizing compiler?*

## Don't: Smash Code into Oblivion

- Hard to read, maintain, & assure correctness

## Do:

- Select best algorithm
- Write code that's readable & maintainable
  - Procedures, recursion, without built-in constant limits
  - Even though these factors can slow down code
- Eliminate optimization blockers
  - Allows compiler to do its job

## Focus on Inner Loops

- Do detailed optimizations where code will be executed repeatedly
- Will get most performance gain here
- Understand how enough about machine to tune effectively

# Complete Example

```
#include <stdio.h>
#include <xmmintrin.h>

/*
 * Compiling:
 * gcc -O3 -march=pentium3 -mfpmath=sse -funroll-loops
 *      -fomit-frame-pointer -o sse sse.c
 */
float inp1[4] = { 1.2, 3.5, 1.7, 2.8 };
float inp2[4] = { -0.7, 2.6, 3.3, -4.0 };
float out[4];

// __m128 : This is 128 bits long.
//          It allows the storage of 4 floating point numbers.
__m128 a, b, c;

// _mm_load_ps and _mm_store_ps: The address must be 16-byte aligned

float inp_sse1[4] __attribute__((aligned(16))) = { 1.2, 3.5, 1.7, 2.8 };
float inp_sse2[4] __attribute__((aligned(16))) = { -0.7, 2.6, 3.3, -4.0 };
float out_sse[4] __attribute__((aligned(16)));
```

# Complete Example

```
int main(void)
{
    int i;

    printf("Ordinal calculation...\n");

    for(i = 0; i < 4; i++)
    {
        out[i] = inp1[i] + inp2[i];
        printf("Result %d: %.2f\n",i,out[i]);
    }
    printf("...with SSE instructions...\n");

    a = _mm_load_ps(inp_sse1);
    b = _mm_load_ps(inp_sse2);
    c = _mm_add_ps(a, b);
    mm_store_ps(out_sse,c);

    for(i = 0; i < 4; i++) {
        printf("Result %d: %.2f\n",i,out_sse[i]);
    }
    return 0;
}
```