

CIS 450 – Computer Architecture and Organization

Lecture 21: Shells and Signals

Mitch Neilsen

(neilsen@ksu.edu)

219D Nichols Hall

Topics

- **Process Tree**
- **Shells**
- **Signals**
 - **Signal Handlers**
- **Non-local Jumps**

The World of Multitasking

System Runs Many Processes Concurrently

- **Process: executing program**
 - State consists of memory image + register values + program counter
- **Continually switches from one process to another**
 - Suspend process when it needs I/O resource or timer event occurs
 - Resume process when I/O available or given scheduling priority
- **Appears to user(s) as if all processes executing simultaneously**
 - Even though most systems can only execute one process at a time
 - Except possibly with lower performance than if running alone

Programmer's Model of Multitasking


Basic Functions

- `fork()` spawns new process
 - Called once, returns twice
- `exit()` terminates own process
 - Called once, never returns
 - Puts the process into “zombie” state
- `wait()` and `waitpid()` wait for and reap terminated children
- `exec1()` and `execve()` run a new program in an existing process
 - Called once, (normally) never returns

Programming Challenge

- Understanding the nonstandard semantics of the functions
- Avoiding improper use of system resources
 - E.g. “fork bombs” can disable a system; e.g., `while(1) fork();`

Unix Process Hierarchy



The ps command

```
cislinux> ps aux -w --forest
```

(output edited to fit slide)

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.0	1516	248	?	Ss	Sep24	0:10	init [3]
root	2	0.0	0.0	0	0	?	S	Sep24	0:02	[migration/0]
root	3	0.0	0.0	0	0	?	SN	Sep24	0:00	[ksoftirqd/0]
root	4	0.0	0.0	0	0	?	S	Sep24	1:44	[migration/1]
root	5	0.0	0.0	0	0	?	SN	Sep24	0:00	[ksoftirqd/1]
root	6	0.0	0.0	0	0	?	S<	Sep24	0:00	[events/0]
root	8	0.0	0.0	0	0	?	S<	Sep24	0:00	[khelper]
root	9	0.0	0.0	0	0	?	S<	Sep24	0:00	[kthread]
root	12	0.0	0.0	0	0	?	S<	Sep24	0:10	_ [kblockd/0]
root	13	0.0	0.0	0	0	?	S<	Sep24	0:00	_ [kblockd/1]
...										
matts	584	0.0	0.0	5788	1376	?	Ss	Sep28	0:01	SCREEN
matts	585	0.0	0.0	2892	196	pts/8	Ss	Sep28	0:00	_ -/bin/bash
matts	589	0.0	0.0	6764	1008	pts/8	S+	Sep28	0:00	_ mysql >
matts	1768	0.0	0.0	2888	548	pts/15	Ss+	Sep28	0:00	_ -/bin/bash
matts	10119	0.0	0.0	2888	876	pts/37	Ss	Sep28	0:00	_ -/bin/bash
matts	25748	0.0	0.0	1920	748	pts/37	S+	Oct26	0:00	_ less F>
eab9844	3178	0.0	0.0	5508	1036	?	Ss	Sep28	0:01	SCREEN
eab9844	3179	0.0	0.0	3556	188	pts/34	Ss	Sep28	0:00	_ -/bin/tcsh
eab9844	3189	0.0	0.0	6816	1748	pts/34	S+	Sep28	0:09	_ irssi >

The ps Command (cont.)

```
USER      PID  TTY      STAT  COMMAND
root      889  tty1     S      /bin/login -- agn
agn       900  tty1     S      \_ xinit -- :0
root      921  ?        SL     \_ /etc/X11/X -auth /usr1/agn/.Xauthority :0
agn       948  tty1     S      \_ /bin/sh /afs/cs.cmu.edu/user/agn/.xinitrc
agn       958  tty1     S      \_ \_ xterm -geometry 80x45+1+1 -C -j -ls -n
agn       966  pts/0    S      \_ \_ -tcsh
agn       1184 pts/0    S      \_ \_ /usr/local/bin/wish8.0 -f /usr
agn       1212 pts/0    S      |   \_ /usr/local/bin/wish8.0 -f
agn       3346 pts/0    S      |   \_ \_ aspell -a -S
agn       1191 pts/0    S      \_ \_ /bin/sh /usr/local/libexec/moz
agn       1204 8 pts/0    S      \_ \_ /usr/local/libexec/mozilla
agn       1207 8 pts/0    S      |   \_ /usr/local/libexec/moz
agn       1208 8 pts/0    S      |   \_ \_ /usr/local/libexec
agn       1209 8 pts/0    S      |   \_ \_ /usr/local/libexec
agn       17814 8 pts/0    S      |   \_ \_ /usr/local/libexec
agn       2469 pts/0    S      \_ \_ usr/local/lib/Acrobat
agn       2483 pts/0    S      \_ \_ java_vm
agn       2484 pts/0    S      \_ \_ java_vm
agn       2485 pts/0    S      \_ \_ java_vm
agn       3042 pts/0    S      \_ \_ java_vm
agn       959  tty1     S      \_ \_ /bin/sh /usr/local/libexec/kde/bin/sta
agn       1020 tty1     S      \_ \_ kwrapper ksmserver
```

Shell Programs

A **shell** is an application program that runs programs on behalf of the user.

- **sh** – Original Unix Bourne Shell
- **csh** – BSD Unix C Shell, **tcsh** – Enhanced C Shell
- **bash** – Bourne-Again Shell

```
int main()
{
    char cmdline[MAXLINE];

    while (1) {
        /* read */
        printf("> ");
        fgets(cmdline, MAXLINE, stdin);
        if (feof(stdin))
            exit(0);

        /* evaluate */
        eval(cmdline);
    }
}
```

Execution is a sequence of
read/evaluate steps

Signals

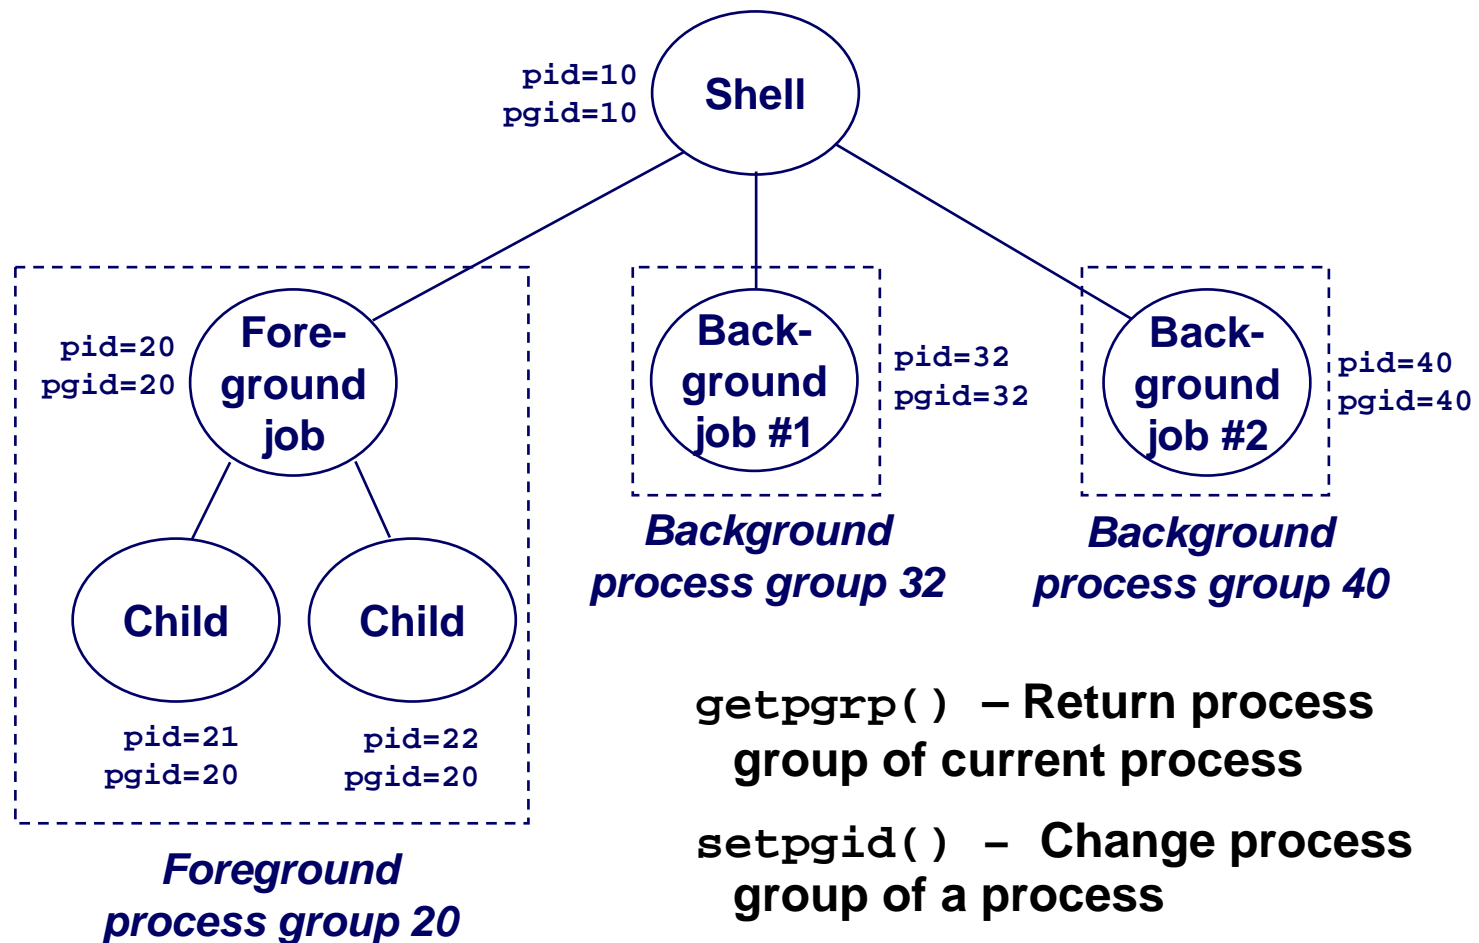
A **signal** is a small message that notifies a process that an event of some type has occurred in the system.

- Kernel abstraction for exceptions and interrupts.
- Sent from the kernel (sometimes at the request of another process) to a process.
- Different signals are identified by small integer ID's (1-30)
- The only information in a signal is its ID and the fact that it arrived.

ID	Name	Default Action	Corresponding Event
2	SIGINT	Terminate	Interrupt from keyboard (ctrl-c)
9	SIGKILL	Terminate	Kill program (cannot override or ignore)
11	SIGSEGV	Terminate & Dump	Segmentation violation
14	SIGALRM	Terminate	Timer signal
17	SIGCHLD	Ignore	Child stopped or terminated

Process Groups

Every process belongs to exactly one process group



Sending Signals with `kill` Program

`kill` program sends arbitrary signal to a process or process group

Examples

- `kill -9 24818`
 - Send SIGKILL to process 24818
- `kill -9 -24817`
 - Send SIGKILL to every process in process group 24817.

```
linux> ./forks 16
linux> Child1: pid=24818 pgrp=24817
Child2: pid=24819 pgrp=24817
```

```
linux> ps
  PID TTY          TIME CMD
 24788 pts/2        00:00:00 tcsh
 24818 pts/2        00:00:02 forks
 24819 pts/2        00:00:02 forks
 24820 pts/2        00:00:00 ps
```


```
linux> kill -9 -24817
```

```
linux> ps
  PID TTY          TIME CMD
 24788 pts/2        00:00:00 tcsh
 24823 pts/2        00:00:00 ps
linux>
```

Sending Signals from the Keyboard

Typing ctrl-c (ctrl-z) sends a SIGINT (SIGTSTP) to every job in the foreground process group.

- SIGINT – default action is to terminate each process
- SIGTSTP – default action is to stop (suspend) each process



Example of `ctrl-c` and `ctrl-z`

```
linux> ./forks 17
Child: pid=24868 pgrp=24867
Parent: pid=24867 pgrp=24867
<typed ctrl-z>
Suspended
linux> ps a
  PID TTY          STAT       TIME COMMAND
 24788 pts/2        S           0:00 -usr/local/bin/tcsh -i
 24867 pts/2        T           0:01 ./forks 17
 24868 pts/2        T           0:01 ./forks 17
 24869 pts/2        R           0:00 ps a
bass> fg
./forks 17
<typed ctrl-c>
linux> ps a
  PID TTY          STAT       TIME COMMAND
 24788 pts/2        S           0:00 -usr/local/bin/tcsh -i
 24870 pts/2        R           0:00 ps a
```

Sending Signals with `kill` Function

```
void fork12()
{
    pid_t pid[N];
    int i, child_status;
    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0)
            while(1); /* Child infinite loop */

    /* Parent terminates the child processes */
    for (i = 0; i < N; i++) {
        printf("Killing process %d\n", pid[i]);
        kill(pid[i], SIGINT);
    }

    /* Parent reaps terminated children */
    for (i = 0; i < N; i++) {
        pid_t wpid = wait(&child_status);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminated abnormally\n", wpid);
    }
}
```

Installing Signal Handlers

The `signal` function modifies the default action associated with the receipt of signal `signum`:

```
■ handler_t *signal(int signum, handler_t *handler)
```

Different values for `handler`:

- **SIG_IGN**: ignore signals of type `signum`
- **SIG_DFL**: revert to the default action on receipt of signals of type `signum`.
- Otherwise, `handler` is the address of a **signal handler**
 - Called when process receives signal of type `signum`
 - Referred to as “**installing**” the handler.
 - Executing handler is called “**catching**” or “**handling**” the signal.
 - When the handler executes its return statement, control passes back to instruction in the control flow of the process that was interrupted by receipt of the signal.

Signal Handling Example

```
void int_handler(int sig)
{
    printf("Process %d received signal %d\n",
           getpid(), sig);
    exit(0);
}

void fork13()
{
    pid_t pid[N];
    int i, child_status;
    signal(SIGINT, int_handler);

    . . .
}
```

```
linux> ./forks 13
Killing process 24973
Killing process 24974
Killing process 24975
Killing process 24976
Killing process 24977
Process 24977 received signal 2
Child 24977 terminated with exit status 0
Process 24976 received signal 2
Child 24976 terminated with exit status 0
Process 24975 received signal 2
Child 24975 terminated with exit status 0
Process 24974 received signal 2
Child 24974 terminated with exit status 0
Process 24973 received signal 2
Child 24973 terminated with exit status 0
linux>
```


Signal Handler Funkiness

```
int ccount = 0;
void child_handler(int sig)
{
    int child_status;
    pid_t pid = wait(&child_status);
    ccount--;
    printf("Received signal %d from process %d\n",
           sig, pid);
}

void fork14()
{
    pid_t pid[N];
    int i, child_status;
    ccount = N;
    signal(SIGCHLD, child_handler);
    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0) {
            /* Child: Exit */
            exit(0);
        }
    while (ccount > 0)
        pause(); /* Suspend until signal occurs */
}
```

Pending signals are not queued

- For each signal type, just have single bit indicating whether or not signal is pending
- Even if multiple processes have sent this signal

Living With Non-Queuing Signals

Must check for all terminated jobs

- Typically loop with `wait`

```
void child_handler2(int sig)
{
    int child_status;
    pid_t pid;
    while ((pid = waitpid(-1, &child_status, WNOHANG)) > 0) {
        ccount--;
        printf("Received signal %d from process %d\n", sig, pid);
    }
}

void fork15()
{
    . . .
    signal(SIGCHLD, child_handler2);
    . . .
}
```

Signal Handler Funkiness (Cont.)

Signal arrival during long system calls (say a `read`)

- Signal handler interrupts `read()` call
- Linux: upon return from signal handler, the `read()` call is restarted automatically
- Some other flavors of Unix can cause the `read()` call to fail with an `EINTR` error number (`errno`); in this case, the application program can restart the slow system call

Subtle differences like these complicate the writing of portable code that uses signals.

A Program That Reacts to Externally Generated Events (ctrl-c)

```
#include <stdlib.h>
#include <stdio.h>
#include <signal.h>

void handler(int sig) {
    printf("You think hitting ctrl-c will stop the bomb?\n");
    sleep(2);
    printf("Well...");
    fflush(stdout);
    sleep(1);
    printf("OK\n");
    exit(0);
}

main() {
    signal(SIGINT, handler); /* installs ctrl-c handler */
    while(1) {
    }
}
```

A Program That Reacts to Internally Generated Events

```
#include <stdio.h>
#include <signal.h>

int beeps = 0;

/* SIGALRM handler */
void handler(int sig) {
    printf("BEEP\n");
    fflush(stdout);

    if (++beeps < 5)
        alarm(1);
    else {
        printf("BOOM!\n");
        exit(0);
    }
}
```

```
main() {
    signal(SIGALRM, handler);
    alarm(1); /* send SIGALRM in
               1 second */

    while (1) {
        /* handler returns here */
    }
}
```

```
cislinux> a.out
BEEP
BEEP
BEEP
BEEP
BEEP
BOOM!
cislinux>
```

Non-local Jumps: `setjmp/longjmp`

Powerful (but dangerous) user-level mechanism for transferring control to an arbitrary location.

- **Controlled to way to break the procedure call / return discipline**
- **Useful for error recovery and signal handling**

```
int setjmp( jmp_buf j )
```

- **Must be called before `longjmp`**
- **Identifies a return site for a subsequent `longjmp`.**
- **Called once, returns one or more times**

Implementation:

- **Remember where you are by storing the current register context, stack pointer, and PC value in `jmp_buf`.**
- **Return 0**

set jmp / long jmp (cont)

```
void longjmp(jmp_buf j, int i)
```

- **Meaning:**

- return from the `set jmp` remembered by jump buffer `j` again...
- ...this time returning `i` instead of 0

- **Called after `set jmp`**

- **Called once, but never returns**

long jmp Implementation:

- **Restore register context from jump buffer `j`**
- **Set `%eax` (the return value) to `i`**
- **Jump to the location indicated by the PC stored in jump buf `j`.**

setjmp/longjmp Example

```
#include <setjmp.h>
jmp_buf buf;

main() {
    if (setjmp(buf) != 0) {
        printf("back in main due to an error\n");
    } else {
        printf("first time through\n");
        p1(); /* p1 calls p2, which calls p3 */
    }
    ...
    p3() {
        <error checking code>
        if (error)
            longjmp(buf, 1)
    }
}
```


Putting It All Together: A Program That Restarts Itself When `ctrl-c`'d

```
#include <stdio.h>
#include <signal.h>
#include <setjmp.h>

sigjmp_buf buf;

void handler(int sig) {
    siglongjmp(buf, 1);
}

main() {
    signal(SIGINT, handler);

    if (!sigsetjmp(buf, 1))
        printf("starting\n");
    else
        printf("restarting\n");
```

```
while(1) {
    sleep(1);
    printf("processing...\n");
}
```

```
bass> a.out
starting
processing...
processing...
restarting
processing...
processing...
restarting
processing...
```

← Ctrl-c

← Ctrl-c

Summary

Signals provide process-level exception handling

- Can generate from user programs
- Can define effect by declaring signal handler

Some caveats

- Very high overhead
 - >10,000 clock cycles
 - Only use for exceptional conditions
- Don't have queues
 - Just one bit for each pending signal type

Non-local jumps provide exceptional control flow within process

- Within constraints of stack discipline