

CIS 450 – Computer Architecture and Organization

Lecture 24: Dynamic Memory Allocation

Mitch Neilsen

(neilsen@ksu.edu)

219D Nichols Hall

Topics

- **From last time ... Memory Mapping**
- **Dynamic Memory Allocation**
- **Simple explicit allocators**
 - **Data structures**
 - **Mechanisms**
 - **Policies**
- **BrickOS Example**

Memory Mapping

Creation of new VM *area* done via “memory mapping”

- Create new `vm_area_struct` and page tables for area
- Area can be backed by (i.e., get its initial values from) :
 - Regular file on disk (e.g., an executable object file)
 - » Initial page bytes come from a section of a file
 - Nothing (e.g., `bss`)
 - » Initial page bytes are zeros
- Dirty pages are swapped back and forth between a special swap file.

Key point: no virtual pages are copied into physical memory until they are referenced!

- Known as “demand paging”
- Crucial for time and space efficiency

User-Level Memory Mapping

```
void *mmap(void *start, int len,  
           int prot, int flags, int fd, int offset)
```

- Map `len` bytes starting at offset `offset` of the file specified by file description `fd`, preferably at address `start` (usually 0 for don't care).
 - `prot`: `MAP_READ`, `MAP_WRITE`
 - `flags`: `MAP_PRIVATE`, `MAP_SHARED`
- Return a pointer to the mapped area.
- Example: fast file copy
 - Useful for applications like Web servers that need to quickly copy files.
 - `mmap` allows file transfers without copying into user space.

mmap() Example: Fast File Copy

```
#include <unistd.h>
#include <sys/mman.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

/*
 * mmap.c - a program that uses mmap
 * to copy itself to stdout
 */
```

```
int main() {
    struct stat stat;
    int i, fd, size;
    char *bufp;

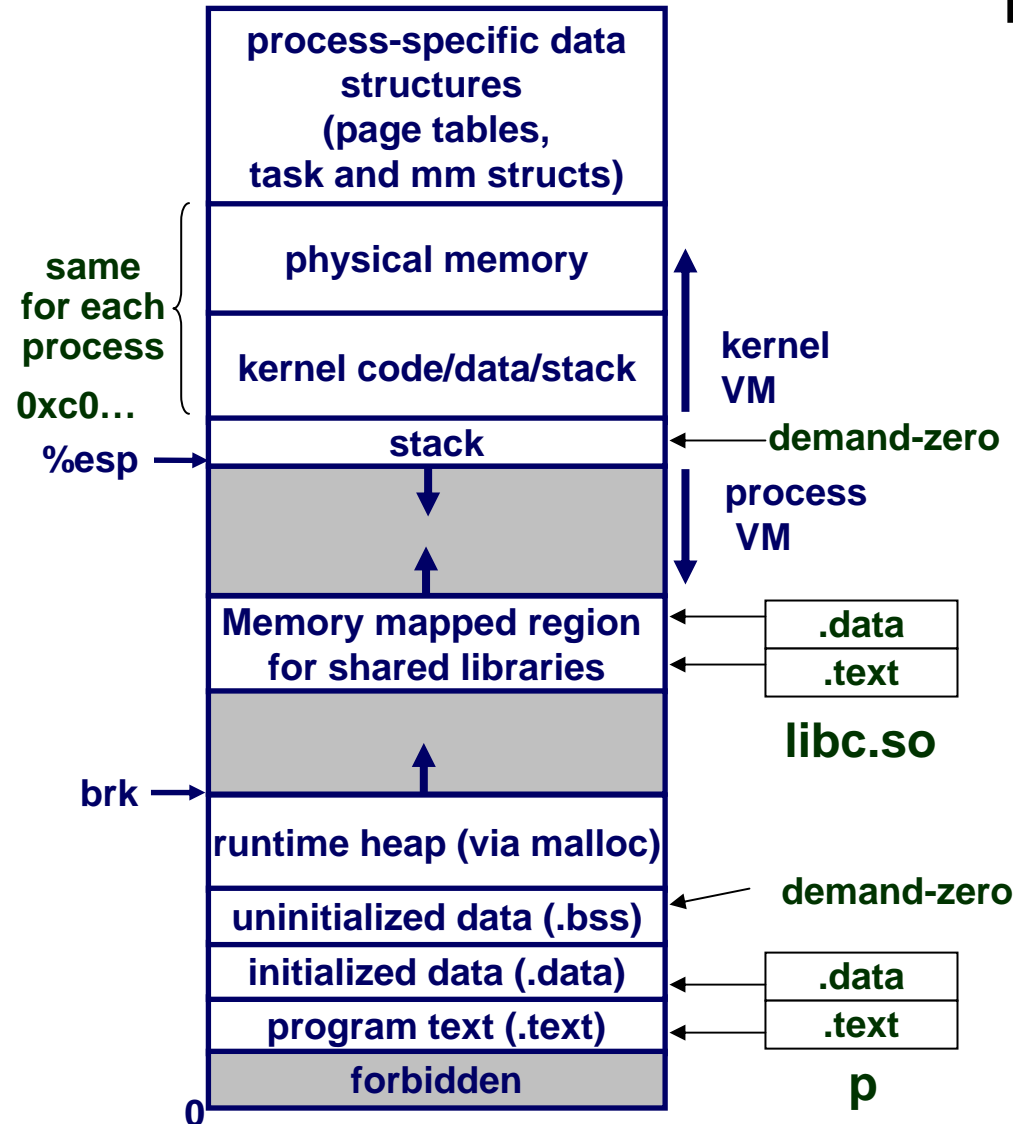
    /* open the file & get its size*/
    fd = open("./mmap.c", O_RDONLY);
    fstat(fd, &stat);
    size = stat.st_size;
    /* map the file to a new VM area */
    bufp = mmap(0, size, PROT_READ,
                MAP_PRIVATE, fd, 0);

    /* write the VM area to stdout */
    write(1, bufp, size);
}
```

Exec() Revisited

To run a new program **p** in the current process using **exec()**:

- Free **vm_area_struct**'s and page tables for old areas.
- Create new **vm_area_struct**'s and page tables for new areas.
 - Stack, bss, data, text, shared libs.
 - Text and data backed by ELF executable object file.
 - bss and stack initialized to zero.
- Set PC to entry point in **.text**
 - Linux will swap in code and data pages as needed.



Fork() Revisited

To create a new process using `fork()`:

- **Make copies of the old process's `mm_struct`, `vm_area_struct`'s, and page tables.**
 - At this point the two processes are sharing all of their pages.
 - How to get separate spaces without copying all the virtual pages from one space to another?
 - » “copy on write” technique.
- **Copy-on-write**
 - Make pages of writeable areas read-only
 - Flag `vm_area_struct`'s for these areas as private “copy-on-write”.
 - Writes by either process to these pages will cause page faults.
 - » Fault handler recognizes copy-on-write, makes a copy of the page, and restores write permissions.

Net result:

- **Copies are deferred until absolutely necessary (i.e., when one of the processes tries to modify a shared page).**

Memory System Summary

Cache Memory

- Purely a speed-up technique
- Behavior invisible to application programmer and (mostly) OS
- Implemented totally in hardware

Virtual Memory

- Supports many OS-related functions
 - Process creation
 - Task switching
 - Protection
- Combination of hardware & software implementation
 - Software management of tables, allocations
 - Hardware access of tables
 - Hardware caching of table entries (TLB)

Harsh Reality

Memory Matters

Memory is not unbounded

- It must be allocated and managed
- Many applications are memory dominated
 - Especially those based on complex, graph algorithms


Memory referencing bugs especially pernicious

- Effects are distant in both time and space

Memory performance is not uniform

- Cache and virtual memory effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

Dynamic Memory Allocation



Explicit vs. Implicit Memory Allocator

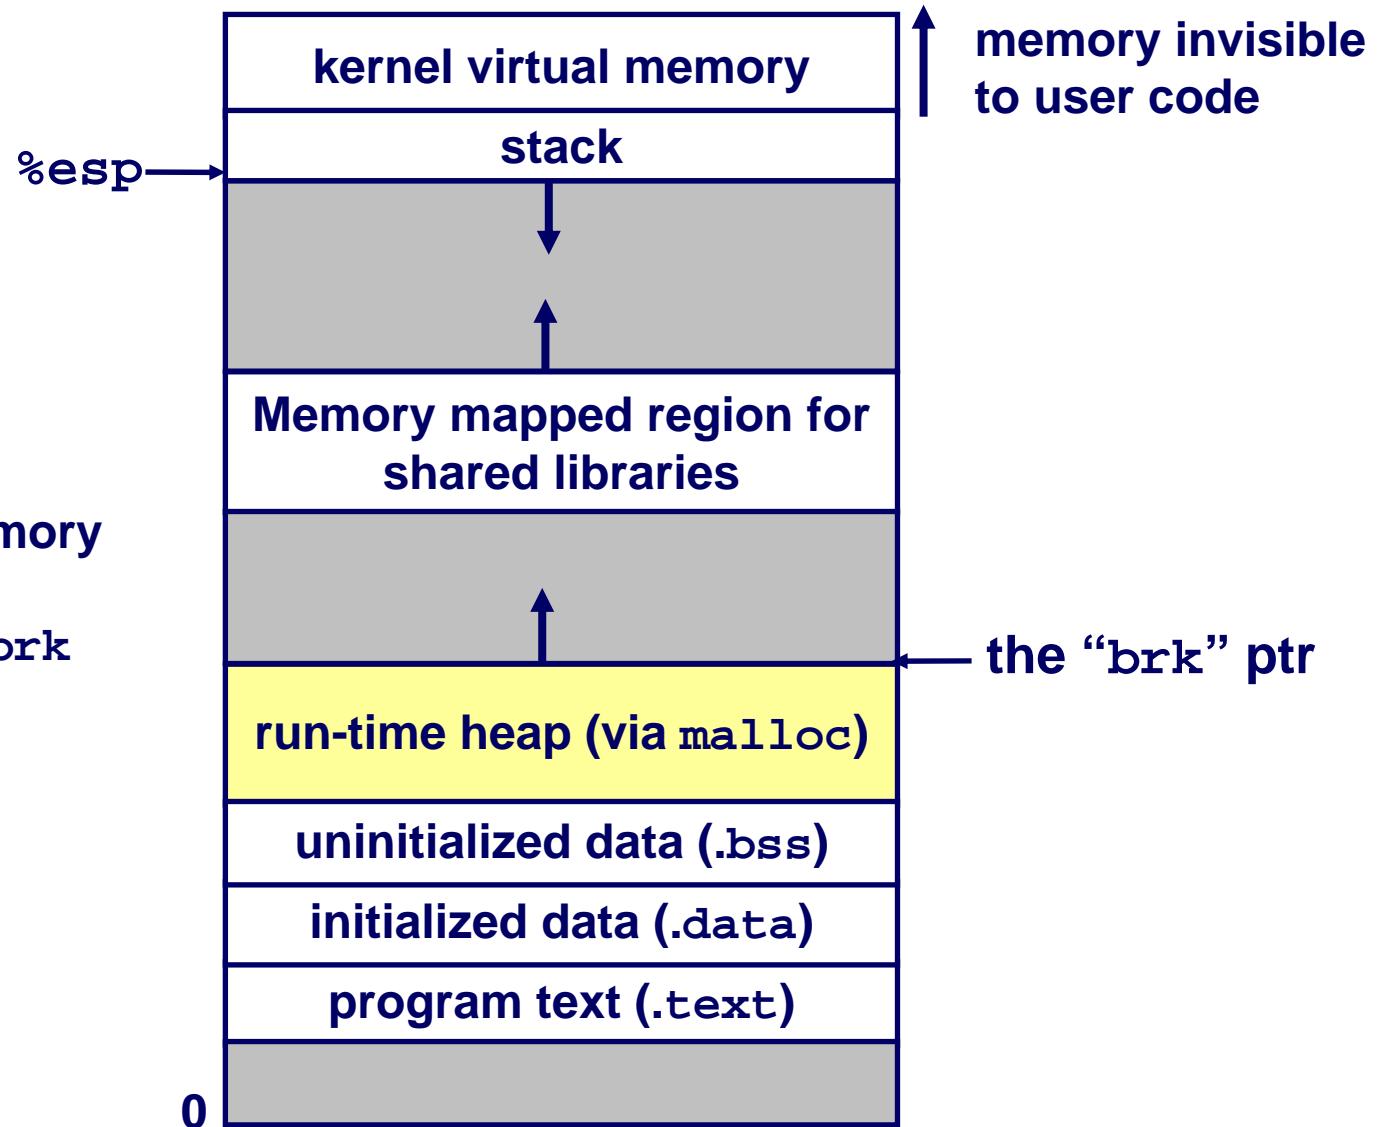
- **Explicit:** application allocates and frees space
 - E.g., `malloc` and `free` in C
- **Implicit:** application allocates, but does not free space
 - E.g. garbage collection in Java, ML or Lisp

Allocation

- In both cases the memory allocator provides an abstraction of memory as a set of blocks
- Doles out free memory blocks to application

Will discuss simple explicit memory allocation today

Process Memory Image



Allocators request additional heap memory from the operating system using the `sbrk` function.

Malloc Package

```
#include <stdlib.h>
```

```
void *malloc(size_t size)
```

- If successful:

- Returns a pointer to a memory block of at least `size` bytes, (typically) aligned to 8-byte boundary.
- If `size == 0`, returns `NULL`

- If unsuccessful: returns `NULL` (0) and sets `errno`.

```
void free(void *p)
```

- Returns the block pointed at by `p` to pool of available memory
- `p` must come from a previous call to `malloc` or `realloc`.

```
void *realloc(void *p, size_t size)
```

- Changes size of block `p` and returns pointer to new block.
- Contents of new block unchanged up to min of old and new size.

Malloc Example

```
void foo(int n, int m) {
    int i, *p;

    /* allocate a block of n ints */
    p = (int *)malloc(n * sizeof(int));
    if (p == NULL) {
        perror("malloc");
        exit(0);
    }
    for (i=0; i<n; i++) p[i] = i;

    /* add m bytes to end of p block */
    if ((p = (int *) realloc(p, (n+m) * sizeof(int))) == NULL) {
        perror("realloc");
        exit(0);
    }
    for (i=n; i < n+m; i++) p[i] = i;


    /* print new array */
    for (i=0; i<n+m; i++)
        printf("%d\n", p[i]);

    free(p); /* return p to available memory pool */
}
```

Assumptions

Assumptions made in this lecture

- Memory is word addressed (each word can hold a pointer)




Allocation Examples


```
p1 = malloc(4)
```



```
p2 = malloc(5)
```



```
p3 = malloc(6)
```



```
free(p2)
```



```
p4 = malloc(2)
```



Constraints

Applications:

- Can issue arbitrary sequence of allocation and free requests
- Free requests must correspond to an allocated block

Allocators

- Can't control number or size of allocated blocks
- Must respond immediately to all allocation requests
 - *i.e.*, can't reorder or buffer requests
- Must allocate blocks from free memory
 - *i.e.*, can only place allocated blocks in free memory
- Must align blocks so they satisfy all alignment requirements
 - 8 byte alignment for GNU malloc (`libc malloc`) on Linux boxes
- Can only manipulate and modify free memory
- Can't move the allocated blocks once they are allocated
 - *i.e.*, compaction is not allowed

Performance Goals: Throughput

Given some sequence of `malloc` and `free` requests:

- $R_0, R_1, \dots, R_k, \dots, R_{n-1}$

Want to maximize throughput and peak memory utilization.

- These goals are often conflicting

Throughput:

- Number of completed requests per unit time

- Example:

- 5,000 `malloc` calls and 5,000 `free` calls in 10 seconds
- Throughput is 1,000 operations/second.

Performance Goals: Peak Memory Utilization

Given some sequence of `malloc` and `free` requests:

- $R_0, R_1, \dots, R_k, \dots, R_{n-1}$

Def: Aggregate payload P_k :

- `malloc(p)` results in a block with a *payload* of p bytes.
- After request R_k has completed, the *aggregate payload* P_k is the sum of currently allocated payloads.

Def: Current heap size is denoted by H_k

- Assume that H_k is monotonically nondecreasing

Def: Peak memory utilization:

- After k requests, *peak memory utilization* is:
 - $U_k = (\max_{i < k} P_i) / H_k$


Internal Fragmentation

Poor memory utilization caused by *fragmentation*.

- Comes in two forms: *internal* and *external* fragmentation

Internal fragmentation

- For some block, *internal fragmentation* is the difference between the block size and the payload size.



- Caused by overhead of maintaining heap data structures, padding for alignment purposes, or explicit policy decisions (e.g., not to split the block).
- Depends only on the pattern of *previous* requests, and thus is easy to measure.


External Fragmentation

Occurs when there is enough aggregate heap memory, but no single free block is large enough


```
p1 = malloc(4)
```



```
p2 = malloc(5)
```



```
p3 = malloc(6)
```



```
free(p2)
```



```
p4 = malloc(6)
```

oops!

External fragmentation depends on the pattern of *future* requests, and thus is difficult to measure.


Implementation Issues

- **How do we know how much memory to free just given a pointer?**
- **How do we keep track of the free blocks?**
- **What do we do with the extra space when allocating a structure that is smaller than the free block it is placed in?**
- **How do we pick a block to use for allocation -- many might fit?**
- **How do we reinsert freed block?**

Knowing How Much to Free


Standard method

- Keep the length of a block in the word preceding the block.
 - This word is often called the *header field* or *header*
- Requires an extra word for every allocated block




Keeping Track of Free Blocks

Method 1: *Implicit list* using lengths -- links all blocks



Method 2: *Explicit list* among the free blocks using pointers within the free blocks



Method 3: *Segregated free list*

- Different free lists for different size classes

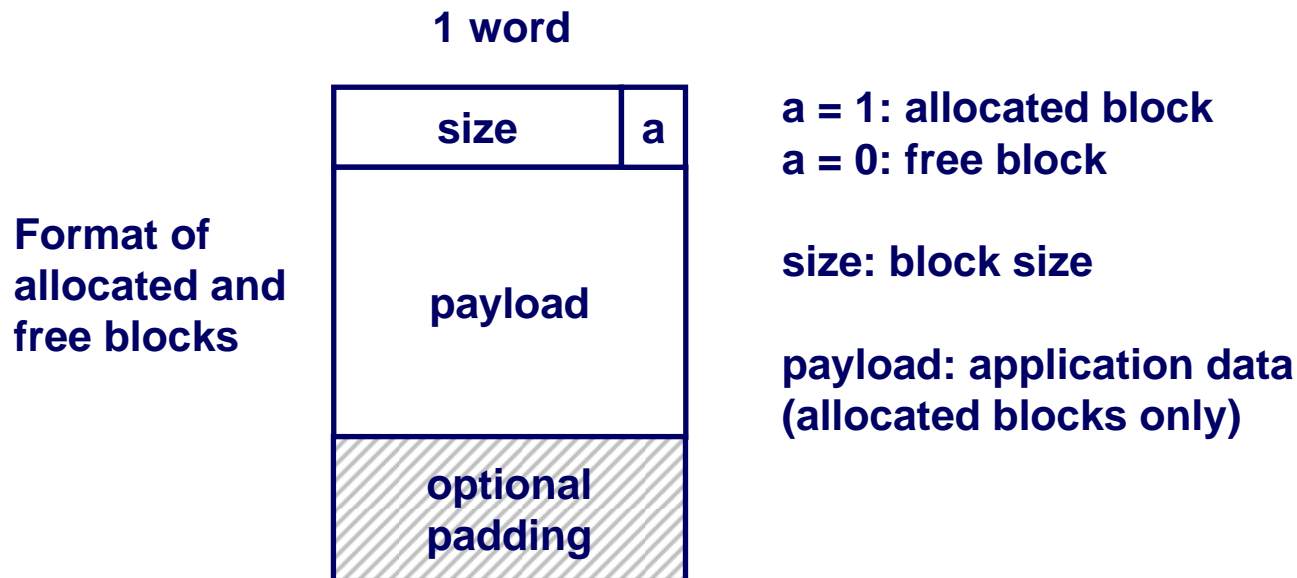
Method 4: Blocks sorted by size

- Can use a balanced tree (e.g. Red-Black tree) with pointers within each free block, and the length used as a key

Method 1: Implicit List

Need to identify whether each block is free or allocated

- Can use extra bit
- Bit can be put in the same word as the size if block sizes are always multiples of two (mask out low order bit when reading size).



Implicit List: Finding a Free Block

First fit:

- Search list from beginning, choose first free block that fits

```
p = start;
while ((p < end) &&          \\ not passed end
      ((*p & 1) ||          \\ already allocated
      (*p <= len)))         \\ too small
    p = p + (*p & -2);      \\ goto next block
```

- Can take linear time in total number of blocks (allocated and free)
- In practice it can cause “splinters” at beginning of list

Next fit:

- Like first-fit, but search list from location of end of previous search
- Research suggests that fragmentation is worse

Best fit:

- Search the list, choose the free block with the closest size that fits
- Keeps fragments small --- usually helps fragmentation
- Will typically run slower than first-fit

Bitfields

How to represent the Header:

● Masks and bitwise operators

```
#define SIZEMASK                (~0x7)

#define PACK(size, alloc)      ((size) | (alloc))

#define GET_SIZE(p)            ((p)->size & SIZEMASK)
```


● Bitfields

```
struct {
    unsigned allocated:1;
    unsigned size:31;
} Header;
```

Implicit List: Allocating in Free Block


Allocating in a free block - *splitting*

- Since allocated space might be smaller than free space, we might want to split the block



```
void addblock(ptr p, int len) {  
    int newsize = ((len + 1) >> 1) << 1; // add 1 and round up  
    int oldsize = *p & -2;                // mask out low bit  
    *p = newsize | 1;                     // set new length  
    if (newsize < oldsize)  
        *(p+newsize) = oldsize - newsize; // set length in remaining  
}
```

addblock(p, 2)




Implicit List: Freeing a Block

Simplest implementation:

- Only need to clear allocated flag

```
void free_block(ptr p) { *p = *p & -2 }
```

- But can lead to “false fragmentation”




There is enough free space, but the allocator won't be able to find it

Implicit List: Coalescing

Join (**coalesce**) with next and/or previous block if they are free

■ Coalescing with next block

```
void free_block(ptr p) {  
    *p = *p & -2;           // clear allocated flag  
    next = p + *p;           // find next block  
    if ((*next & 1) == 0)  
        *p = *p + *next;     // add to this block if  
                               // not allocated  
}
```


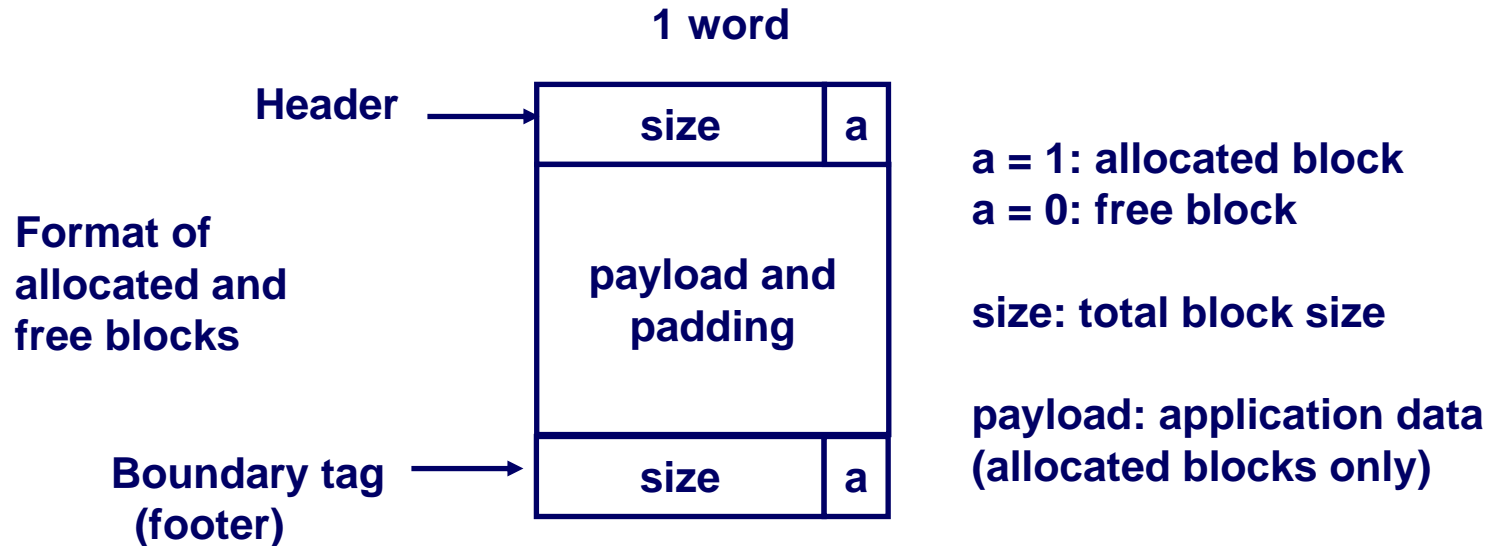


■ But how do we coalesce with previous block?

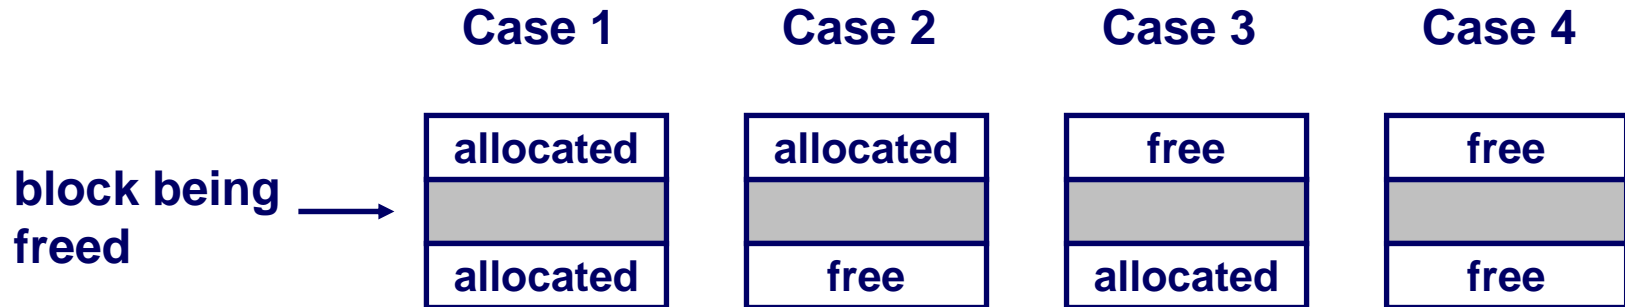
Implicit List: Bidirectional Coalescing

Boundary tags [Knuth73]


- Replicate size/allocated word at bottom of free blocks
- Allows us to traverse the “list” backwards, but requires extra space
- Important and general technique!




Constant Time Coalescing




Constant Time Coalescing (Case 1)




Constant Time Coalescing (Case 2)



Constant Time Coalescing (Case 3)



Constant Time Coalescing (Case 4)



Summary of Key Allocator Policies

Placement policy:

- First fit, next fit, best fit, etc.
- Trades off lower throughput for less fragmentation
 - *Interesting observation: segregated free lists (next lecture) approximate a best fit placement policy without having to search entire free list.*

Splitting policy:

- When do we go ahead and split free blocks?
- How much internal fragmentation are we willing to tolerate?

Coalescing policy:

- *Immediate coalescing*: coalesce each time `free` is called
- *Deferred coalescing*: try to improve performance of `free` by deferring coalescing until needed. e.g.,
 - Coalesce as you scan the free list for `malloc`.
 - Coalesce when the amount of external fragmentation reaches some threshold.

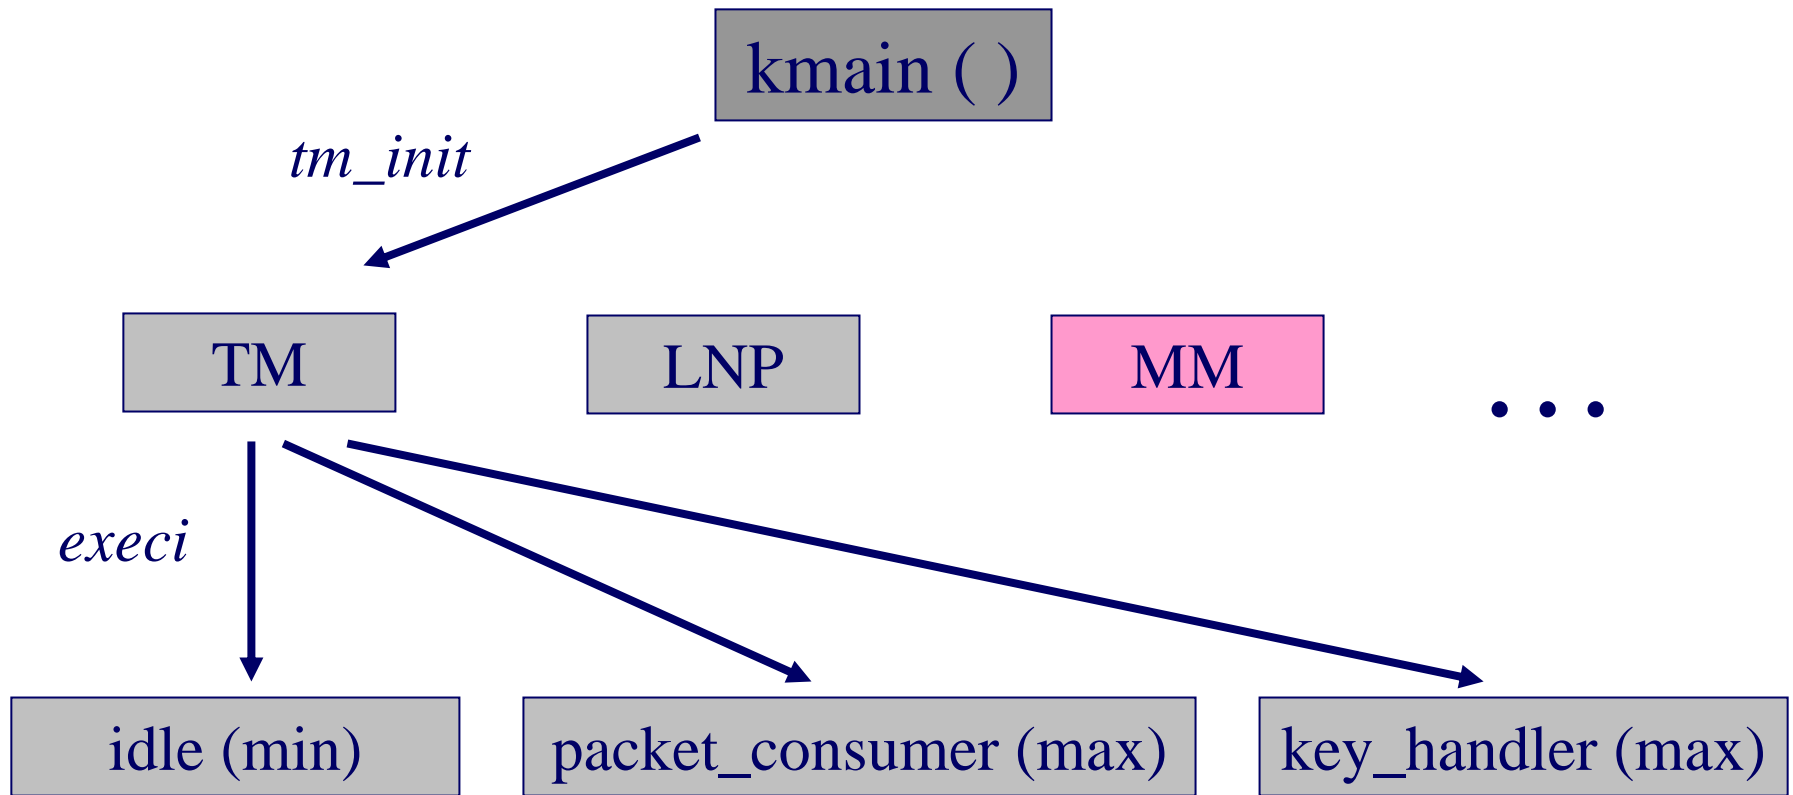
Implicit Lists: Summary

- **Implementation:** very simple
- **Allocate cost:** linear time worst case
- **Free cost:** constant time worst case -- even with coalescing
- **Memory usage:** will depend on placement policy
 - First fit, next fit or best fit

Not used in practice for `malloc/free` because of linear time allocate. Used in many special purpose applications.

However, the concepts of splitting and boundary tag coalescing are general to *all* allocators.

BrickOS Memory Management (mm.c)



BrickOS mm.c

```
#include <sys/mm.h>

#ifdef CONF_MM
#include <stdlib.h>
#include <sys/tm.h>
#include <sys/critsec.h>
#include <string.h>
////////////////////////////////////
// Variables
////////////////////////////////////
size_t  *mm_first_free;          //!< first free block

#ifndef CONF_TM
typedef size_t tid_t;           //!< dummy process ID type
//! current process ID
/*! we need a non-null, non-0xffff current pid even if there is no task
   management. */
const tid_t ctid=0x0001;
#endif
```

BrickOS mm.c

```
////////////////////////////////////  
// Functions  
////////////////////////////////////  
// memory block structure:  
// 0 1      : pid of owner (0=empty)  
// 2 3      : size of data block >> 1  (n 16-bit words)  
// 4 ... 4+2n-1 : data  
  
//! check for free blocks after this one and join them if possible  
/* \param ptr pointer to size field of current block  
   \return size of block */  
size_t mm_try_join(size_t *ptr) {  
    size_t *next=ptr+*ptr+1;  
    size_t increase=0;  
  
    while(*next==MM_FREE && next>=&mm_start) {  
        increase+=*(next+1) + MM_HEADER_SIZE;  
        next      +=*(next+1) + MM_HEADER_SIZE;  
    }  
    return (*ptr)+=increase;  
}
```


BrickOS mm.c

```
//! defragment free blocks
/*! use mm_try_join on each free block of memory
*/

void mm_defrag() {
    size_t *ptr = &mm_start;
#ifdef CONF_TM
    ENTER_KERNEL_CRITICAL_SECTION();
#endif
    while(ptr >= &mm_start) {
        if(*ptr == MM_FREE)
            mm_try_join(ptr+1);
        ptr += *(ptr+1);
        ptr += MM_HEADER_SIZE;
    }
#ifdef CONF_TM
    LEAVE_KERNEL_CRITICAL_SECTION();
#endif
}
```

BrickOS mm.c

```
//! update first free block pointer
/*! \param start pointer to owner field of a memory block to start with.
*/
void mm_update_first_free(size_t *start) {
    size_t *ptr=start;

    while((*ptr!=MM_FREE) && (ptr>=&mm_start))
        ptr+=*(ptr+1)+MM_HEADER_SIZE;

    mm_first_free=ptr;
}
```

BrickOS mm.c

```
#!/ initialize memory management
void mm_init() {
    size_t *current,*next;

    current=&mm_start;

    // memory layout
    MM_BLOCK_FREE      (&mm_start);    // ram
                                         // something at 0xc000 ?
    MM_BLOCK_RESERVED(0xef30);         // lcddata
    MM_BLOCK_FREE      (0xef50);       // ram2
    MM_BLOCK_RESERVED(0xf000);         // motor
    MM_BLOCK_FREE      (0xfe00);       // ram4
    MM_BLOCK_RESERVED(0xff00);         // stack, onchip

    // expand last block to encompass all available memory
    *current=(int)((((int) current)-2)>>1);

    mm_update_first_free(&mm_start);
}
```

BrickOS malloc()

```
//! allocate a block of memory
/*! \param size requested block size, return 0 on error, else pointer to block.*/
void *malloc(size_t size) {
    size_t *ptr,*next;
    size=(size+1)>>1;          // only multiples of 2
#ifdef CONF_TM
    ENTER_KERNEL_CRITICAL_SECTION();
#endif
    ptr=mm_first_free;
    while(ptr>=&mm_start) {
        if(*(ptr++)==MM_FREE) {      // free block?
#ifdef CONF_TM
            mm_try_join(ptr);    // unite with later blocks
#endif
            if(*ptr>=size) {        // big enough?
                *(ptr-1)=(size_t)ctid; // set owner
                if((*ptr-size)>=MM_SPLIT_THRESH) {
                    next=ptr+size+1;
                    *(next++)=MM_FREE;
                    *(next)=*ptr-size-MM_HEADER_SIZE;
                    mm_try_join(next);
                }
            }
        }
    }
}
```

```

        *ptr=size;
    }
    // was it the first free one?
    if(ptr==mm_first_free+1)
        mm_update_first_free(ptr+*ptr+1);
#ifdef CONF_TM
    LEAVE_KERNEL_CRITICAL_SECTION();
#endif
    return (void*) (ptr+1);
}
}
ptr+=(*ptr)+1;        // find next block.
}
#ifdef CONF_TM
    LEAVE_KERNEL_CRITICAL_SECTION();
#endif
    return NULL;
}

```

```

//! free a previously allocated block of memory.
void free(void *the_ptr) {
    size_t *ptr=the_ptr;
#ifdef CONF_TM
        size_t *p2,*next;
#endif
    if(ptr==NULL || (((size_t)ptr)&1) )
        return;
    ptr-=MM_HEADER_SIZE;
    *((size_t*) ptr)=MM_FREE;        // mark as free
#ifdef CONF_TM
    // for task-safe ops, free must be atomic and nonblocking, update mm_first_free
    if(ptr<mm_first_free || mm_first_free<&mm_start)
        mm_first_free=ptr;          // update mm_first_free
#else // without task management, we have time to unite neighboring mem. blocks.
    p2=&mm_start;
    while(p2!=ptr) {                // we could make free
        next=p2+*(p2+1)+MM_HEADER_SIZE;    // O(1) if we included
        if(*p2==MM_FREE && next==ptr)      // a pointer to the
            break;                        // previous block.
        p2=next;                        // I don't want to.
    }
    mm_try_join(p2+1);              // defragment free areas
    if(ptr<mm_first_free || mm_first_free<&mm_start)
        mm_update_first_free(ptr);      // update mm_first_free
#endif
}

```

```

//! allocate adjacent blocks of memory
/*! \param nmemb number of blocks (must be > 0)
    \param size individual block size (must be >0)
    \return 0 on error, else pointer to block
*/
void *calloc(size_t nmemb, size_t size) {
    void *ptr;
    size_t original_size = size;

    if (nmemb == 0 || size == 0)
        return 0;

    size*=nmemb;

    // if an overflow occurred, size/nmemb will not equal original_size
    if (size/nmemb != original_size)
        return 0;

    if((ptr=malloc(size))!=NULL)
        memset(ptr,0,size);

    return ptr;
}

```

```
//! free all blocks allocated by the current process (with TID = ctid).
/*! called by exit() and kmain().
*/
void mm_reaper() {
    size_t *ptr;

    // pass 1: mark as free
    ptr=&mm_start;
    while(ptr>=&mm_start) {
        if(*ptr==(size_t)ctid)
            *ptr=MM_FREE;
        ptr+=*(ptr+1)+MM_HEADER_SIZE;
    }

    // pass 2: defragment free areas
    // this may alter free blocks
    mm_defrag();
}
```



```
//! return the number of bytes of unallocated memory
int mm_free_mem(void) {
    int free = 0;
    size_t *ptr;

#ifdef CONF_TM
    ENTER_KERNEL_CRITICAL_SECTION();
#endif

    // Iterate through the free list
    for (ptr = mm_first_free; ptr >= &mm_start; ptr += *(ptr+1) + MM_HEADER_SIZE)
        free += *(ptr+1);

#ifdef CONF_TM
    LEAVE_KERNEL_CRITICAL_SECTION();
#endif
    return free*2;
}
#endif
```