

Lecture 3: Processes and Threads

Instructor: Mitch Neilsen

Office: N219D

Quote of the Day

“Talent is cheaper than table salt.

What separates the talented individual from the successful one is a lot of hard work.”

-- Stephen King

Project 0: Notes

■ Add BXSHARE environment variable

- `export BXSHARE=/<home dir>/cis520/usr/local/share/bochs`

■ `$ pintos run alarm-multiple`

- with Unity 3d – causes system freeze :-(
- `<ctrl>-alt F1`, kill offending processes or reset

■ Workarounds:

- Select Unity 2d on login to lab machine
- Run command-line version of output:

```
$ pintos -v -- run alarm-multiple
```

Project 0: Notes

■ Goals for this week:

- Finish installing Pintos
- Add new test program
 - ▶ Hint: first, use **grep -r alarm-multiple *** in the pintos/src/tests folder to see what changes are needed
 - ▶ Then, rebuild the operating system back in the pintos/src/threads folder; e.g., **make clean** and **make**


Processes

- **A *process* is an instance of a running program**
- **Modern OSes run multiple processes simultaneously**
- **Examples (all can run simultaneously):**
 - gcc fileA.c – compiler running on fileA
 - gcc fileB.c – compiler running on fileB
 - emacs – text editor
 - firefox – web browser
- **Non-examples (implemented as one process):**
 - Multiple firefox tabs or emacs frames (one process, multiple threads)
- **Why processes?**
 - Simplicity of programming
 - Higher throughput (better CPU utilization), lower latency

Speed


- **Multiple processes can increase CPU utilization**

- Overlap one process's computation with another's wait



- **Multiple processes can reduce latency**

- Running *A* then *B* requires 100 sec for *B* to complete



- Running *A* and *B* concurrently allows *B* to finish faster



- *A* is slightly slower, but less than 100 sec unless *A* and *B* are both completely CPU-bound

Processes in the real world

- **Processes, parallelism fact of life much longer than OSes have been around**
 - E.g., say takes 1 worker 10 months to make 1 widget
 - Company may hire 100 workers to make 10,000 widgets
 - Latency for first widget $\gg 1/10$ month
 - Throughput may be < 10 widgets per month (if we can't perfectly parallelize tasks)
 - Or > 10 widgets per month if we get better utilization (e.g., 100 workers on 10,000 widgets never idly waiting for paint to dry)
- **You will see this with Pintos**
 - BUT, don't expect labs to take $1/3$ time with three people ;-)

A process's view of the world

- **Each process has its own view of the machine**

- Its own address space
- Its own open files
- Its own virtual CPU (through preemptive multitasking)


- **`*(char *)0xc000` different in P_1 & P_2**

- **Greatly simplifies programming model**

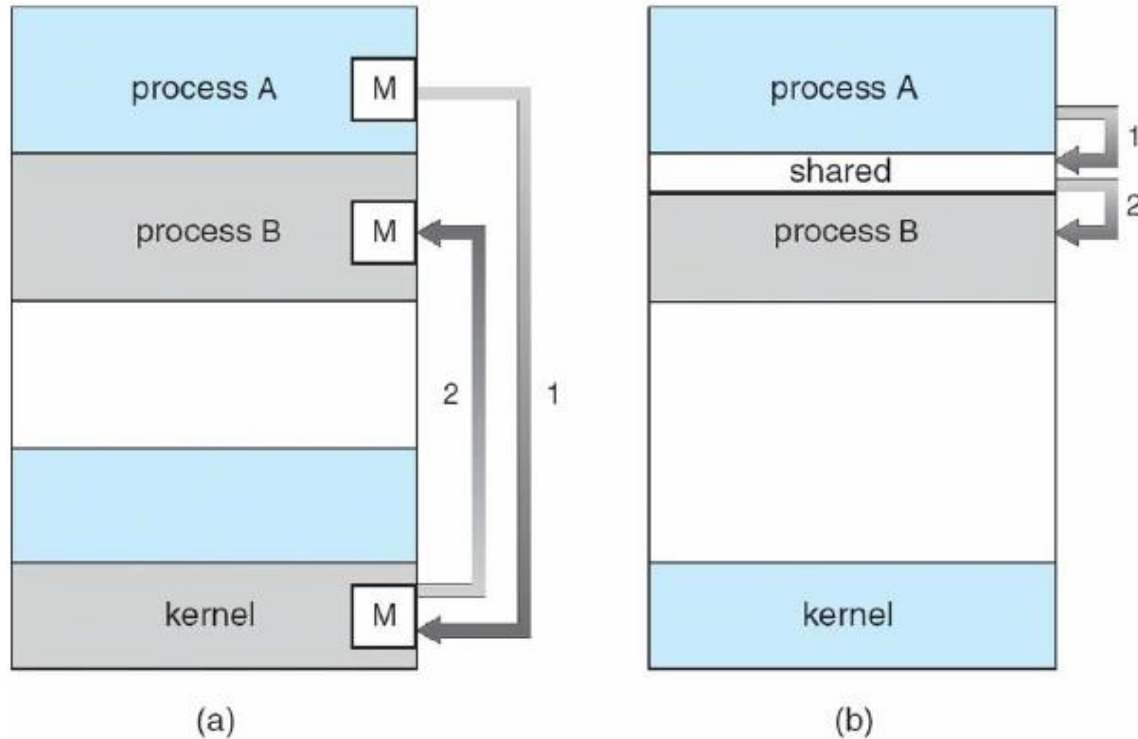
- gcc does not care that firefox is running

- **Sometimes want interaction between processes**

- Simplest is through files: emacs edits file, gcc compiles it
- More complicated: Shell/command, Window manager/app.



Inter-Process Communication



- **How can processes interact in real time?**

- (a) By passing messages through the kernel
- (b) By sharing a region of physical memory
- (c) Through asynchronous signals or alerts

Rest of lecture

- **User view of processes**
 - Crash course in basic Unix/Linux system call interface
 - How to create, kill, and communicate between processes
- **Kernel view of processes**
 - Implementing processes in the kernel
- **Threads**
- **How to implement threads**

UNIX files I/O

- **Applications “open” files (or devices) by name**
 - I/O happens through open files
- `int open(char *path, int flags, /*mode*/...);`
 - flags: `O_RDONLY`, `O_WRONLY`, `O_RDWR`
 - `O_CREAT`: create the file if non-existent
 - `O_EXCL`: (w/ `O_CREAT`) create even if file exists already
 - `O_TRUNC`: Truncate the file to length 0
 - `O_APPEND`: Start writing from end of file
 - mode: final argument with `O_CREAT` – set r,w,x permissions
- **Returns file descriptor—used for all I/O to file**

Error returns

- **What if open fails? Returns -1 (invalid fd)**
- **Most system calls return -1 on failure**
 - Specific kind of error stored in global int **errno**
- **#include <sys/errno.h> for possible values**
 - `errno = 2 = ENOENT` "No such file or directory"
 - `errno = 13 = EACCES` "Permission Denied"
- **perror function prints human-readable message**
 - `perror ("initfile");`
→ "initfile: No such file or directory"

Operations on file descriptors

- `int read (int fd, void *buf, int nbytes);`
 - Returns number of bytes read
 - Returns 0 bytes at end of file, or -1 on error
- `int write (int fd, void *buf, int nbytes);`
 - Returns number of bytes written, -1 on error
- `off_t lseek (int fd, off t_pos, int whence);`
 - whence: 0 – start, 1 – current, 2 – end
 - ◁ Returns previous file offset, or -1 on error
- `int close (int fd);`

File descriptor numbers

- **File descriptors are inherited by processes**
 - When one process spawns another, same fds by default
- **Descriptors 0, 1, and 2 have special meaning**
 - 0 – “standard input” (stdin in ANSI C)
 - 1 – “standard output” (stdout, printf in ANSI C)
 - 2 – “standard error ” (stderr, perror in ANSI C)
 - Normally all three attached to terminal
- **Example: type.c**
 - Prints the contents of a file to stdout

type.c

```
void
typefile (char *filename)
{
    int fd, nread;
    char buf[1024];

    fd = open (filename, O_RDONLY);
    if (fd == -1) {
        perror (filename);
        return;
    }

    while ((nread = read (fd, buf, sizeof (buf))) > 0)
        write (1, buf, nread);

    close (fd);
}
```

Creating processes

- `int fork (void);`
 - Create new process that is exact copy of current one
 - Returns *process ID* of new process in "parent"
 - Returns 0 in "child"
- `int waitpid (int pid, int *stat, int opt);`
 - pid – process to wait for, or -1 for any
 - stat – will contain exit value, or signal
 - opt – usually 0 or WNOHANG
 - Returns process ID or -1 on error

Deleting processes

- `void exit (int status);` // modern Linux replace w/ `return(..);`
 - Current process ceases to exist
 - status shows up in `waitpid` (shifted)
 - By convention, status of 0 is success, non-zero error
- `int kill (int pid, int sig);`
 - Sends signal `sig` to process `pid`
 - `SIGTERM` most common value, kills process by default (but application can catch it for "cleanup")
 - `SIGKILL` stronger, kills process always, and cannot be ignored

Running programs

- `int execve (char *prog, char **argv, char **envp);`
 - prog – full pathname of program to run
 - argv – argument vector that gets passed to main
 - envp – environment variables, e.g., PATH, HOME
- **Generally called through a wrapper functions**
 - `int execlp (char *prog, char **argv);`
Search PATH for prog, use current environment
 - `int execlp (char *prog, char *arg, ...);`
List arguments one at a time, finish with NULL
- **Example:** minish.c
 - Loop that reads a command, then executes it
- **Warning: Pintos exec more like combined fork/exec**

minish.c (simplified)

```
pid_t pid; char **av;
```

```
void doexec () {  
    execvp (av[0], av);  
    perror (av[0]);  
    exit (1);  
}
```

```
/* ... main loop: */
```

```
for (;;) {  
    parse_next_line_of_input (av, stdin);
```

```
    switch (pid = fork ()) {
```

```
    case -1:
```

```
        perror ("fork"); break;
```

```
    case 0:
```

```
        doexec (); break;
```

```
    default:
```

```
        waitpid (pid, NULL, 0); break;
```


```
    }
```

```
}
```

minish.c (simplified)

```
pid_t pid; char **av;
```

```
void doexec () {  
    execvp (av[0], av);  
    perror (av[0]);  
    exit (1);  
}  
  
/* ... main loop: */  
for (;;) {  
    parse_next_line_of_input (av, stdin);  
  
    switch (pid = fork ()) {  
    case -1:  
        perror ("fork"); break;  
    case 0:  
        doexec (); break;  
    default:  
        waitpid (pid, NULL, 0); break;  
    }  
}
```



Manipulating file descriptors

- `int dup2 (int oldfd, int newfd);`
 - Closes `newfd`, if it was a valid descriptor
 - Makes `newfd` an exact copy of `oldfd`
 - Two file descriptors will share same offset (lseek on one will affect both)
- `int fcntl (int fd, F_SETFD, int val)`
 - Sets *close on exec* flag if `val = 1`, clears if `val = 0`
 - Makes file descriptor non-inheritable by spawned programs
- **Example:** `redirsh.c`
 - Loop that reads a command and executes it
 - Recognizes: `$ command < input > output 2> errlog`

redirsh.c

```
void doexec (void) {
    int fd;

    /* infile non-NULL if user typed "command < infile" */
    if (infile) {
        if ((fd = open (infile, O_RDONLY)) < 0) {
            perror (infile);
            exit (1);
        }
        if (fd != 0) {
            dup2 (fd, 0);
            close (fd);
        }
    }

    /* ... Do same for outfile -> fd 1, errfile -> fd 2 ... */

    execvp (av[0], av);
    perror (av[0]);
    exit (1);
}
```

Pipes

- `int pipe (int fds[2]);`
 - Returns two file descriptors in `fds[0]` and `fds[1]`
 - Writes to `fds[1]` will be read on `fds[0]`
 - When last copy of `fds[1]` closed, `fds[0]` will return EOF
 - Returns 0 on success, -1 on error
- **Operations on pipes**
 - read/write/close – as with files
 - When `fds[1]` closed, `read(fds[0])` returns 0 bytes
 - When `fds[0]` closed, `write(fds[1])`:
 - ◁ Kills process with SIGPIPE, or if blocked
 - ◁ Fails with EPIPE
- **Example:** `pipesh.c`
 - Sets up pipeline: `$ command1 | command2 | command3 ...`


pipesh.c (simplified)

```
void doexec (void) {
    int pipefds[2];
    while (outcmd) {
        pipe (pipefds);
        switch (fork ()) {
            case -1:
                perror ("fork"); exit (1);
            case 0:
                dup2 (pipefds[1], 1);
                close (pipefds[0]); close (pipefds[1]);
                outcmd = NULL;
                break;
            default:
                dup2 (pipefds[0], 0);
                close (pipefds[0]); close (pipefds[1]);
                parse_command_line (&av, &outcmd, outcmd);
                break;
        }
    }
} /* ... */
```


pipesh.c (simplified)

```
void doexec (void) {  
    int pipefds[2];  
    while (outcmd) {  
        pipe (pipefds);  
        switch (fork ()) {  
            case -1:  
                perror ("fork"); exit (1);  
            case 0:  
                dup2 (pipefds[1], 1);  
                close (pipefds[0]); close (pipefds[1]);  
                outcmd = NULL;  
                break;  
            default:  
                dup2 (pipefds[0], 0);  
                close (pipefds[0]); close (pipefds[1]);  
                parse_command_line (&av, &outcmd, outcmd);  
                break;  
        }  
    }  
}
```

/* ... */



Why fork?

- **Most calls to fork followed by** `execve`
- **Could also combine into one *spawn* system call**
 - This is what Pintos `exec` does
- **Occasionally useful to fork one process**
 - Unix *dump* utility backs up file system to tape
 - If tape fills up, must restart at some logical point
 - Implemented by forking to revert to old state if tape ends
- **Real win is simplicity of interface**
 - Tons of things you might want to do to child:
Manipulate file descriptors, environment, resource limits, etc.
 - Yet fork requires *no* arguments at all

Spawning process w/o fork

- Without fork, require tons of different options
- Example: Windows CreateProcess system call

```
BOOL WINAPI CreateProcess(  
    __in_opt LPCTSTR lpApplicationName,  
    __inout_opt LPTSTR lpCommandLine,  
    __in_opt LPSECURITY_ATTRIBUTES lpProcessAttributes,  
    __in_opt LPSECURITY_ATTRIBUTES lpThreadAttributes,  
    __in BOOL bInheritHandles,  
    __in DWORD dwCreationFlags,  
    __in_opt LPVOID lpEnvironment,  
    __in_opt LPCTSTR lpCurrentDirectory,  
    __in LPSTARTUPINFO lpStartupInfo,  
    __out LPPROCESS_INFORMATION lpProcessInformation );
```


Implementing processes

- **OS keeps data structure for each proc**
 - Process Control Block (PCB)
 - Called proc in Unix, task_struct in Linux
- **Tracks *state* of the process**
 - Running, runnable, blocked, etc.
- **Includes information necessary to run**
 - Registers, virtual memory mappings, etc.
 - Open files (including memory mapped files)
- **Various other data about the process**
 - Credentials (user/group ID), signal mask, controlling terminal, priority, accounting statistics, whether being debugged, which system call binary emulation in use, ...

Process state
Process ID
User id, etc.
Program counter
Registers
Address space (VM data structs)
Open files

PCB

Process states



- **Process can be in one of several states**
 - *new* & *terminated* at beginning & end of life
 - *running* – currently executing (or will execute on kernel return)
 - *ready* – can run, but kernel has chosen different process to run
 - *waiting* – needs async event (e.g., disk operation) to proceed
- **Which process should kernel run?**
 - if 0 runnable, run idle loop, if 1 runnable, run it
 - if >1 runnable, must make scheduling decision

Processes

- **A *process* is an instance of a running program**
- **Modern OSes run multiple processes simultaneously**
- **Why processes?**
 - Simplicity of programming
 - Higher throughput (better CPU utilization), lower latency
 - But, relatively expensive to create a new process, next time we will turn our attention to scheduling and threads

Summary

- Course web page via K-State OnLine has all lecture notes, assignments, handouts, etc.
- Read Ch. 1-4: Processes and Threads
- Friday: Finish Project 0 and set up base revision of Pintos on version control system; e.g., using subversion, git, etc.
- Watch YouTube Video on Git with Scott Chacon:
<http://www.youtube.com/watch?v=ZDR433b0HJY>



Upload

0:06:48 / 1:22:12

CC Settings

Introduction to Git with Scott Chacon of GitHub