

Lecture 6: Synchronization

Instructor: Mitch Neilsen

Office: N219D

Outline

■ Reading:

- Ch. 4 - Threads
- Ch. 5 - CPU Scheduling
- Ch. 6 - Synchronization



■ Project 1: Scheduling and Synchronization

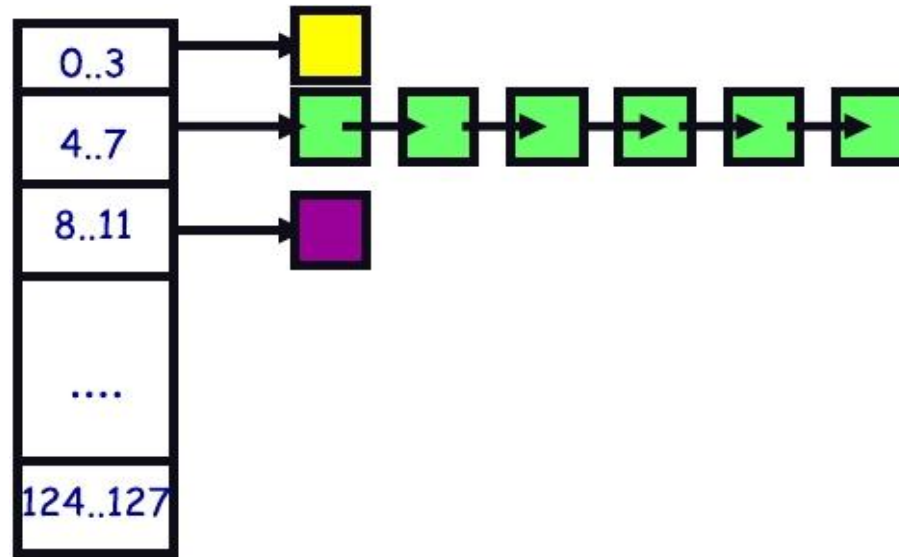
- Alarm Clock
- Priority-based Scheduler
- Synchronization and Priority Inheritance
- [Extra Credit] MLFQ Scheduler

Quote of the Day

"Sometimes you are in sync with the times, sometimes you are in advance, sometimes you are late. "

-- Bernardo Bertolucci

Multilevel feedback queues (BSD)



- **Every runnable process on one of 32 run queues**
 - Kernel runs process on highest-priority non-empty queue
 - Round-robins among processes on same queue
- **Process priorities dynamically computed**
 - Processes moved between queues to reflect priority changes
 - If a process gets higher priority than running process, run it
- **Idea: Favor interactive jobs that use less CPU**

Process priority (BSD model)

- **p_nice** – user-settable weighting factor
- **p_estcpu** – per-process estimated CPU usage
 - Incremented whenever timer interrupt found proc. running
 - Decayed every second when process runnable

$$p_estcpu \leftarrow \left(\frac{2 \cdot \text{load}}{2 \cdot \text{load} + 1} \right) p_estcpu$$

- Load is sampled average of length of run queue plus short-term sleep queue over last minute
- **Set process priority by** (lower p_usrpri = higher priority)

$$p_usrpri \leftarrow 50 + \left(\frac{p_estcpu}{4} \right) + 2 \cdot p_nice$$

(value clipped if over 127)

Sleeping process increases priority

- **p_estcpu not updated while asleep**
 - Instead p_slptime keeps count of sleep time
- **When process becomes runnable**

$$p_estcpu \leftarrow \left(\frac{2 \cdot \text{load}}{2 \cdot \text{load} + 1} \right)^{p_slptime} \cdot p_estcpu$$

- Approximates decay ignoring nice and past loads
- **Previous description based on [McKusick]^a**

^aSee library.stanford.edu for off-campus access

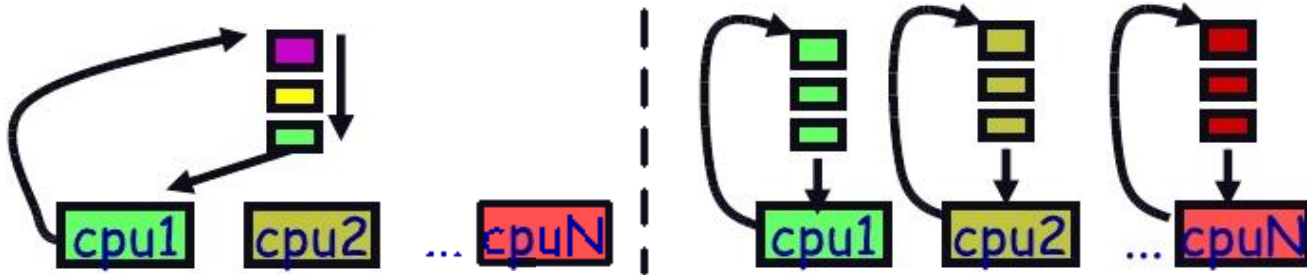
Pintos notes

- **Same basic idea for second half of project 1**
 - But 64 priorities, not 128
 - Higher numbers mean higher priority
 - Okay to have only one run queue if you prefer (less efficient, but we won't deduct points for it)
- **Have to negate priority equation:**

$$\text{priority} = 63 - \left(\frac{\text{recent_cpu}}{4} \right) - 2 \cdot \text{nice}$$

Multiprocessor scheduling issues

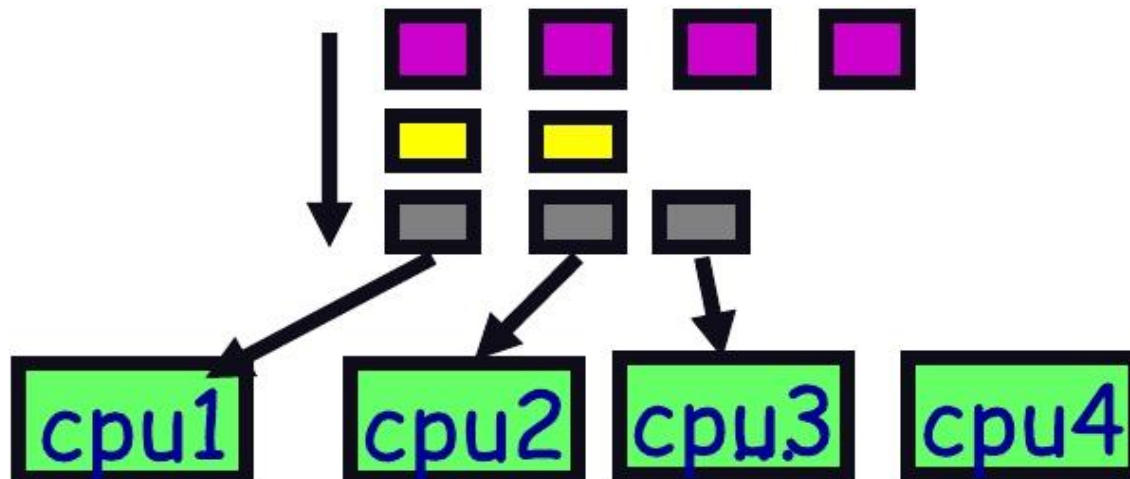
- **Must decide on more than which processes to run**
 - Must decide on **which CPU** to run which process
- **Moving between CPUs has costs**
 - More cache misses, depending on arch more TLB misses too
- ***Affinity scheduling*** - try to keep threads on same CPU



- But also prevent load imbalances
- Do *cost-benefit* analysis when deciding to migrate

Multiprocessor scheduling (cont)

- **Want related processes scheduled together**
 - Good if threads access same resources (e.g., cached files)
 - Even more important if threads communicate often, otherwise must context switch to communicate
- ***Gang scheduling* - schedule all CPUs synchronously**
 - With synchronized quanta, easier to schedule related processes/threads together



Thread scheduling

- **With thread library, have two scheduling decisions:**
 - ***Local Scheduling*** – Thread library decides which user thread to put onto an available kernel thread
 - ***Global Scheduling*** – Kernel decides which kernel thread to run next
- **Can expose to the user**
 - E.g., `pthread_attr_setscope` allows two choices
 - `PTHREAD_SCOPE_SYSTEM` – thread scheduled like a process (effectively one kernel thread bound to user thread – Will return `ENOTSUP` in user-level pthreads implementation)
 - `PTHREAD_SCOPE_PROCESS` – thread scheduled within the current process (may have multiple user threads multiplexed onto kernel threads)

Thread dependencies

- **Say H at high priority, L at low priority**
 - L acquires lock l .
 - Scene 1: H tries to acquire l , fails, spins. L never gets to run.
 - Scene 2: H tries to acquire l , fails, blocks. M enters system at medium priority. L never gets to run.
 - Both scenes are examples of **priority inversion**
- **Scheduling = deciding who should make progress**
 - Obvious: a thread's importance should increase with the importance of those that depend on it.
 - Naive priority schemes violate this.

Priority donation

- **Say higher number = higher priority**
- **Example 1: L (prio 2), M (prio 4), H (prio 8)**
 - L holds lock l
 - M waits on l , L 's priority raised to $L' = \max(M, L) = 4$
 - Then H waits on l , L 's priority raised to $\max(H, L') = 8$
- **Example 2: Same threads**
 - L holds lock l , M holds lock l_2
 - M waits on l , L 's priority now $L' = 4$ (as before)
 - Then H waits on l_2 . M 's priority goes to $M' = \max(H, M) = 8$, and L 's priority raised to $\max(M', L') = 8$
- **Example 3: L (prio 2), M_1, \dots, M_{1000} (all prio 4)**
 - L has l , and M_1, \dots, M_{1000} all block on l . L 's priority is $\max(L, M_1, \dots, M_{1000}) = 4$.

Review: Thread Package API

- `tid thread_create (void (*fn) (void *), void *arg);`
 - Create a new thread that calls function `fn` with `arg`
- `void thread_exit ();`
- `void thread_join (tid thread);`
- **The execution of multiple threads is interleaved**
- **Can have *non-preemptive threads*:**
 - One thread executes exclusively until it makes a blocking call.
- **Or *preemptive threads*:**
 - May switch to another thread between any two instructions.
- **Using multiple CPUs is inherently preemptive**
 - Even if you don't take `CPU0` away from thread `T`, another thread on `CPU1` can execute between any two instructions of `T`.

Program A

```
int flag1 = 0, flag2 = 0;
```

```
void p1 ( ) {  
    flag1 = 1;  
    if (!flag2) { critical_section_1 (); }  
}
```

```
void p2 ( ) {  
    flag2 = 1;  
    if (!flag1) { critical_section_2 (); }  
}
```

```
int main ( ) {  
    tid id = thread_create (p1, NULL);  
    p2 (); thread_join (id);  
}
```

Even though the threads might deadlock, can both critical sections run simultaneously?

Program B

```
int data = 0, ready = 0;
```

```
void p1 ( ) {  
    data = 2000;  
    ready = 1;  
}
```

```
void p2 ( ) {  
    while (!ready)  
        ;  
    use (data);  
}
```

```
int main ( ) { ... }
```

**Can use() be called
with value 0?**

Program C

```
int a = 0, b = 0;
```

```
void p1 ( ) { a = 1; }
```

```
void p2 ( ) {  
    if (a == 1)  
        b = 1;  
}
```

```
void p3 ( ) {  
    if (b == 1)  
        use (a);  
}
```

```
int main ( ) { ... }
```

**Can use() be called
with value 0?**

Correct answers

- Program A: I don't know
- Program B: I don't know
- Program C: I don't know
- Why?
 - It depends on your hardware
 - If it provides *sequential consistency*, then answers are all **No**
 - But not all hardware provides sequential consistency
- Note: Examples and other slide content from [Adve & Gharachorloo]

Sequential Consistency

- ***Sequential consistency***: The result of execution is as if all operations were executed in some sequential order, and the operations of each processor occurred in the order specified by the program. **[Lamport]**
- Boils down to two requirements:
 1. Maintaining ***program order*** on individual processors
 2. Ensuring ***write atomicity***
- **Without SC, multiple CPUs can be “worse” than preemptive threads**
 - May see results that cannot occur with any interleaving on 1 CPU
- **Why doesn't all hardware support sequential consistency?**

SC thwarts hardware optimizations

- **Complicates write buffers**
 - E.g., read $\text{flag}(n)$ before $\text{flag}(3 - n)$ written through in Program A
- **Can't re-order overlapping write operations**
 - Concurrent writes to different memory modules
 - Coalescing writes to same cache line
- **Complicates non-blocking reads**
 - E.g., speculatively prefetch data in Program B
- **Makes cache coherence more expensive**
 - Must delay write completion until invalidation/update (Program B)
 - Can't allow overlapping updates if no globally visible order (Program C)

SC thwarts compiler optimizations

- **Code motion**
- **Caching value in register**
 - E.g., ready flag in Program B
- **Common subexpression elimination**
 - Could cause memory location to be read fewer times
- **Loop blocking**
 - Re-arrange loops for better cache performance
- **Software pipelining**
 - Move instructions across iterations of a loop to overlap instruction latency with branch cost

x86 consistency

- **x86 supports multiple consistency/caching models**
 - Memory Type Range Registers (MTRR) specify consistency for ranges of physical memory (e.g., frame buffer)
 - Page Attribute Table (PAT) allows control for each 4K page
- **Choices include:**
 - **WB**: Write-back caching (**the default**)
 - **WT**: Write-through caching (all writes go to memory)
 - **UC**: Uncacheable (for device memory)
 - **WC**: Write-combining – weak consistency & no caching
- **Some instructions have weaker consistency**
 - String instructions
 - Special “non-temporal” instructions that bypass cache

x86 WB consistency

- **Old x86s (e.g, 486, Pentium 1) had almost SC**
 - Exception: A read could finish before an earlier write to a different location
 - Which of Programs **A**, **B**, **C** might be affected?
- **Newer x86s let a processor read its own writes early**

x86 WB consistency

- **Old x86s (e.g, 486, Pentium 1) had almost SC**
 - Exception: A read could finish before an earlier write to a different location
 - Which of Programs **A**, **B**, **C** might be affected? **Just A**
- **Newer x86s let a processor read its own writes early**

- E.g., both of these functions can return 2:

```
int flag1 = 0, flag2 = 0;
```

```
int p1 (void *ignored)
{
    register int f, g;
    flag1 = 1;
    f = flag1;
    g = flag2;
    return 2*f + g;
}
```

```
int p2 (void *ignored)
{
    register int f, g;
    flag2 = 1;
    f = flag2;
    g = flag1;
    return 2*f + g;
}
```

- Older CPUs would wait at "f = ..." until store completes

x86 atomicity

- lock **prefix makes a memory instruction atomic**
 - Usually locks bus for duration of instruction (expensive!)
 - Can avoid locking if memory already exclusively cached
 - All lock instructions totally ordered
 - Other memory instructions cannot be re-ordered w. locked ones
- xchg **instruction is always locked (even w/o prefix)**
- **Special fence instructions can prevent re-ordering**
 - LFENCE – can't be reordered w. reads (or later writes)
 - SFENCE – can't be reordered w. writes
 - MFENCE – can't be reordered w. reads or writes

Assuming sequential consistency

- Let's for now say we have sequential consistency
- Example concurrent code: Producer/Consumer
 - buffer stores **BUFFER_SIZE** items
 - **count** is number of used slots
 - **in** is next empty buffer slot to fill (if any)
 - **out** is oldest filled slot to consume (if any)

```

void producer (void *ignored) {
    for (;;) {
        /* produce an item and put in nextProduced */
        while (count == BUFFER_SIZE)
            ; // do nothing
        buffer [in] = nextProduced;
        in = (in + 1) % BUFFER_SIZE;
        count++;
    }
}

```

```

void consumer (void *ignored) {
    for (;;) {
        while (count == 0)
            ; // do nothing
        nextConsumed = buffer[out];
        out = (out + 1) % BUFFER_SIZE;
        count--;
        /* consume the item in nextConsumed */
    }
}

```

- **What can go wrong here?**

Data races

- count **may have wrong value**
- **Possible implementation of** `count++` **and** `count--`

`register ← count`

`register ← register + 1`

`count ← register`

`register ← count`

`register ← register - 1`

`count ← register`

- **Possible execution (count one less than correct):**

`register ← count`

`register ← register + 1`

`register ← count`

`register ← register - 1`

`count ← register`

`count ← register`

Data races (continued)

- **What about a single-instruction add?**
 - E.g., i386 allows single instruction `addl $1, count`
 - So implement `count++/--` with one instruction
 - Now are we safe?

Data races (continued)

- **What about a single-instruction add?**
 - E.g., i386 allows single instruction `addl $1, count`
 - So implement `count++/--` with one instruction
 - Now are we safe?
- **Not atomic on multiprocessor!**
 - Will experience exact same race condition
 - Can potentially make atomic with lock prefix
 - But lock very expensive
 - Compiler won't generate it, assumes you don't want penalty
- **Need solution to *critical section* problem**
 - Place `count++` and `count--` in critical section
 - Protect critical sections from concurrent execution

Critical Section: Desired solution

- ***Mutual Exclusion***

- Only one thread can be in critical section at a time

- ***Progress***

- Say no process currently in critical section (C.S.)
- One of the processes trying to enter will eventually get in

- ***Bounded waiting***

- Once a thread T starts trying to enter the critical section, there is a bound on the number of times other threads get in

- **Note progress vs. bounded waiting**

- If no thread can enter C.S., don't have progress
- If thread A waiting to enter C.S. while B repeatedly leaves and re-enters C.S. *ad infinitum*, don't have bounded waiting

Peterson's solution

- **Still assuming sequential consistency**
- **Assume two threads, T_0 and T_1**
- **Variables**
 - `int not_turn;` – not this thread's turn to enter C.S.
 - `bool wants[2];` – `wants[i]` indicates if T_i wants to enter C.S.

- **Code:**

```
for (;;) { /* code in thread i */
    wants[i] = true;
    not_turn = i;
    while (wants[1-i] && not_turn == i)
        /* other thread wants in and not our turn, so loop */;
    Critical_section ();
    wants[i] = false;
    Remainder_section ();
}
```

Does Peterson's solution work?

```
for (;;) { /* code in thread i */
    wants[i] = true;
    not_turn = i;
    while (wants[1-i] && not_turn == i)
        /* other thread wants in and not our turn, so loop */
    Critical_section ();
    wants[i] = false;
    Remainder_section ();
}
```

- **Mutual exclusion – can't both be in C.S.**
 - Would mean $wants[0] == wants[1] == true$, so not_turn would have blocked one thread from C.S.
- **Progress – If T_{1-i} not in C.S., can't block T_i**
 - Means $wants[1-i] == false$, so T_1 won't loop
- **Bounded waiting – similar argument to progress**
 - If T_i wants lock and T_{1-i} tries to re-enter, T_{1-i} will set $not_turn = 1 - i$, allowing T_i in

Mutexes

- **Peterson expensive, only works for 2 processes**
 - Can generalize to n , but for some fixed n
- **Want to insulate programmer from implementing synchronization primitives**
- **Thread packages typically provide *mutexes*:**

```
void mutex_init (mutex_t *m, ..);  
void mutex_lock (mutex_t *m);  
int mutex_trylock (mutex_t *m);  
void mutex_unlock (mutex_t *m);
```

 - Only one thread acquires m at a time, others wait
 - **All global data should be protected by a mutex!**
- **OS kernels also need synchronization**
 - May or may not look like mutexes

Same concept, many names

- **Most popular application-level thread API: *pthread*s**
 - Function names in this lecture all based on *pthread*s
 - Just add pthread_ prefix
 - E.g., pthread_mutex_t, pthread_mutex_lock, . . .
- **Same abstraction in **Pintos** under different name**
 - Data structure is struct lock
 - void lock_init (struct lock *);
 - void lock_acquire (struct lock *);
 - bool lock_try_acquire (struct lock *);
 - void lock_release (struct lock *);
- **Extra Pintos feature:**
 - Release checks that the lock was acquired by the same thread
 - bool lock_held_by_current_thread (struct lock *lock);

Improved producer

```
mutex_t mutex = MUTEX_INITIALIZER;

void producer (void *ignored) {
    for (;;) {
        /* produce an item and put in nextProduced */

        mutex_lock (&mutex);
        while (count == BUFFER_SIZE) {
            mutex_unlock (&mutex); // <--- Why?
            thread_yield ();
            mutex_lock (&mutex);
        }

        buffer [in] = nextProduced;
        in = (in + 1) % BUFFER_SIZE;
        count++;
        mutex_unlock (&mutex);
    }
}
```

Improved consumer

```
void consumer (void *ignored) {
    for (;;) {
        mutex_lock (&mutex);
        while (count == 0) {
            mutex_unlock (&mutex);
            thread_yield ();
            mutex_lock (&mutex);
        }

        nextConsumed =  buffer[out];
        out = (out + 1) % BUFFER_SIZE;
        count--;
        mutex_unlock (&mutex);

        /* consume the item in nextConsumed */
    }
}
```

Condition variables

- **Busy-waiting in application is a bad idea**
 - Thread consumes CPU even when can't make progress
 - Unnecessarily slows other threads and processes
- **Better to inform scheduler of which threads can run**
- **Typically done with condition variables**
- `void cond_init (cond_t *, ...);`
 - Initialize
- `void cond_wait (cond_t *c, mutex_t *m);`
 - Atomically unlock m and sleep until c signaled
 - Then re-acquire m and resume executing
- `void cond_signal (cond_t *c);`
`void cond_broadcast (cond_t *c);`
 - Wake one/all threads waiting on c

Improved producer

```
mutex_t mutex = MUTEX_INITIALIZER;
cond_t nonempty = COND_INITIALIZER;
cond_t nonfull = COND_INITIALIZER;

void producer (void *ignored) {
    for (;;) {
        /* produce an item and put in nextProduced */

        mutex_lock (&mutex);
        while (count == BUFFER_SIZE)
            cond_wait (&nonfull, &mutex);

        buffer [in] = nextProduced;
        in = (in + 1) % BUFFER_SIZE;
        count++;
        cond_signal (&nonempty);
        mutex_unlock (&mutex);
    }
}
```

Improved consumer

```
void consumer (void *ignored) {
    for (;;) {
        mutex_lock (&mutex);
        while (count == 0)
            cond_wait (&nonempty, &mutex);

        nextConsumed = buffer[out];
        out = (out + 1) % BUFFER_SIZE;
        count--;
        cond_signal (&nonfull);
        mutex_unlock (&mutex);

        /* consume the item in nextConsumed */
    }
}
```

Condition variables (continued)

- **Why must** `cond_wait` **both release mutex & sleep?**
- **Why not separate mutexes and condition variables?**

```
while (count == BUFFER_SIZE) {  
    mutex_unlock (&mutex);  
    cond_wait (&nonfull);  
    mutex_lock (&mutex);  
}
```


Condition variables (continued)

- **Why must** `cond_wait` **both release mutex & sleep?**
- **Why not separate mutexes and condition variables?**

```
while (count == BUFFER_SIZE) {  
    mutex_unlock (&mutex);  
    cond_wait (&nonfull);  
    mutex_lock (&mutex);  
}
```

- **Can end up stuck waiting when bad interleaving**

PRODUCER

```
while (count == BUFFER_SIZE);  
mutex_unlock (&mutex);
```

```
cond_wait (&nonfull);
```

CONSUMER

```
mutex_lock (&mutex);
```

...

```
count--;
```

```
cond_signal (&nonfull);
```

Other thread package features

- **Alerts** – cause exception in a thread
- **Timedwait** – timeout on condition variable
- **Shared locks** – concurrent read accesses to data
- **Thread priorities** – control scheduling policy
 - Mutex attributes allow various forms of *priority donation*
(will be familiar concept after lab 1)
- **Thread-specific global data**
- **Different synchronization primitives** (in a few slides)
 - Monitors
 - Semaphores

Implementing synchronization

- **User-visible mutex is straight-forward data structure**

```
typedef struct mutex {  
    bool is_locked;           /* true if locked */  
    thread_id_t owner;        /* thread holding lock, if locked */  
    thread_list_t waiters;    /* threads waiting for lock */  
  
    lower_level_lock_t lk;    /* Protect above fields */  
};
```

- **Need lower-level lock lk for mutual exclusion**
 - Internally, mutex_* functions bracket code with
lock(mutex->lk) . . . unlock(mutex->lk)
 - Otherwise, data races! (E.g., two threads manipulating waiters)
- **How to implement lower_level_lock_t?**
 - Could use Peterson's algorithm, but typically a bad idea
(too slow and don't know maximum number of threads)

Approach #1: Disable interrupts

- **Only for apps with $n : 1$ threads (1 kthread)**
 - Cannot take advantage of multiprocessors
 - But sometimes most efficient solution for uniprocessors
- **Have per-thread “do not interrupt” (DNI) bit**
- **lock (lk): sets thread’s DNI bit**
- **If timer interrupt arrives**
 - Check interrupted thread’s DNI bit
 - If DNI clear, preempt current thread
 - If DNI set, set “interrupted” (I) bit & resume current thread
- **unlock (lk): clears DNI bit *and* checks I bit**
 - If I bit is set, immediately yields the CPU

Approach #2: Spinlocks

- **Most CPUs support atomic read-[modify-]write**
- **Example:** `int test_and_set (int *lockp);`
 - Atomically sets `*lockp = 1` and returns old value
 - Special instruction – can't be implemented in portable C
- **Use this instruction to implement *spinlocks*:**

```
#define lock(lockp)    while (test_and_set (lockp))
#define trylock(lockp) (test_and_set (lockp) == 0)
#define unlock(lockp) *lockp = 0
```
- **Spinlocks implement mutex's `lower_level_lock_t`**
- **Can you use spinlocks instead of mutexes?**
 - Wastes CPU, especially if thread holding lock not running
 - Mutex functions have short C.S., less likely to be preempted
 - On multiprocessor, sometimes good to spin for a bit, then yield

Synchronization on x86

- **Test-and-set only one possible atomic instruction**
- **x86 xchg instruction, exchanges reg with mem**
 - Can use to implement test-and-set

```
_test_and_set:
    movl    8(%esp), %edx    # %edx = lockp
    movl    $1, %eax        # %eax = 1
    xchgl   %eax, (%edx)     # swap (%eax, *lockp)
    ret
```

- **CPU locks memory system around read and write**
 - Recall xchgl always acts like it has lock prefix
 - Prevents other uses of the bus (e.g., DMA)
- **Usually runs at memory bus speed, not CPU speed**
 - Much slower than cached read/buffered write

Kernel Synchronization

- **Should kernel use locks or disable interrupts?**
- **Old UNIX had non-preemptive threads, no mutexes**
 - Interface designed for single CPU, so count++ etc. not data race
 - ... *Unless* memory shared with an interrupt handler

```
int x = splhigh (); // Disable interrupts
// Touch data shared with interrupt handler
splx (x);           // Restore previous state
```
 - C.f., **Pintos** `intr_disable` / `intr_set_level`
- **Used arbitrary pointers like condition variables**
 - `int [t]sleep (void *ident, int priority, ...);`
put thread to sleep; will wake up at priority (`~cond_wait`)
 - `int wakeup (void *ident);`
wake up all threads sleeping on `ident` (`~cond_broadcast`)

Kernel locks

- **Nowadays, should design for multiprocessors**
 - Even if first version of OS is for uniprocessor
 - Someday may want multiple CPUs and need *preemptive* threads
 - That's why Pintos uses locks
- **Multiprocessor performance needs fine-grained locks**
 - Want to be able to call into the kernel on multiple CPUs
- **If kernel has locks, should it ever disable interrupts?**

Kernel locks

- **Nowadays, should design for multiprocessors**
 - Even if first version of OS is for uniprocessor
 - Someday may want multiple CPUs and need *preemptive* threads
 - That's why Pintos uses locks
- **Multiprocessor performance needs fine-grained locks**
 - Want to be able to call into the kernel on multiple CPUs
- **If kernel has locks, should it ever disable interrupts?**
 - Yes! Can't sleep in interrupt handler, so can't wait for lock
 - So even modern OSes have support for disabling interrupts
 - Often uses DNI trick, which is cheaper than masking interrupts in hardware

Semaphores [Dijkstra]

- **A *Semaphore* is initialized with an integer N**
- **Provides two functions:**
 - `sem_wait (S)` (originally called P , called ***sema_down*** in Pintos)
 - `sem_signal (S)` (originally called V , called ***sema_up*** in Pintos)
- **Guarantees `sem_wait` will return only N more times than `sem_signal` called**
 - Example: If $N == 1$, then semaphore is a mutex with `sem_wait` as lock and `sem_signal` as unlock
- **Semaphores allow elegant solutions to some problems**

Semaphore

A semaphore is a structure consisting of 2 parts:

```
struct semaphore {  
    int count; // number of resources available  
    queue Q; // queue of process/thread ids of blocked  
}
```

Shorthand notation:

semaphore $S = 1 \rightarrow S.count = 1, S.Q = \{ \}$

Operations on Semaphores

There are two basic semaphore operations:

`sem_wait(S):`

- if ($S.count > 0$) then $S.count = S.count - 1$;
- else block calling process in $S.Q$;

`sem_signal(S):`

- if ($S.Q$ is non-empty) then wakeup a process in $S.Q$;
- else $S.count = S.count + 1$;

Semaphore Example: Mutual Exclusion

Semaphore S = 1;

Thread A:

sem_wait(S);

(do work in critical section CS);

sem_signal(S);

Thread B:

sem_wait(S);

(do work in CS);

sem_signal(S);

Semaphore Example: Order Execution

Semaphore $S = 0$;

Thread A \rightarrow Thread B:

Thread A:

(do work);

sem_signal(S);

Thread B:

sem_wait(S);

(do work);



Semaphore producer/consumer

- **Can re-write producer/consumer to use three semaphores**
- **Semaphore** mutex **initialized to 1**
 - Used as mutex, protects buffer, in, out. . .
- **Semaphore** full **initialized to 0**
 - To block consumer when buffer empty
- **Semaphore** empty **initialized to N**
 - To block producer when queue full

```
void producer (void *ignored) {  
    for (;;) {  
        /* produce an item and put in nextProduced */  
        sem_wait (&empty);  
        sem_wait (&mutex);  
        buffer [in] = nextProduced;  
        in = (in + 1) % BUFFER_SIZE;  
        sem_signal (&mutex);  
        sem_signal (&full);  
    }  
}
```

```
void consumer (void *ignored) {  
    for (;;) {  
        sem_wait (&full);  
        sem_wait (&mutex);  
        nextConsumed = buffer[out];  
        out = (out + 1) % BUFFER_SIZE;  
        sem_signal (&mutex);  
        sem_signal (&empty);  
        /* consume the item in nextConsumed */  
    }  
}
```


Summary

- Read Ch. 1-6
- Processes and Threads (Ch. 4)
- Process Scheduling (Ch. 5)
- Synchronization (Ch. 6)
- Project 1 – Scheduling and Synchronization