

# Lecture 7: Synchronization (cont.)

**Instructor: Mitch Neilsen**

**Office: N219D**

# Outline

## ■ Reading:

- Ch. 4 - Threads
- Ch. 5 - CPU Scheduling
- Ch. 6 - Synchronization




## ■ Project 1: Scheduling and Synchronization

- Alarm Clock
- Priority-based Scheduler
- Synchronization and Priority Inheritance
- [Extra Credit] MLFQ Scheduler

# Quote of the Day

"Only those who attempt the absurd will achieve the impossible."

-- M. C. Escher



# Improved producer

```
mutex_t mutex = MUTEX_INITIALIZER;

void producer (void *ignored) {
    for (;;) {
        /* produce an item and put in nextProduced */

        mutex_lock (&mutex);
        while (count == BUFFER_SIZE) {
            mutex_unlock (&mutex); // <--- Why?
            thread_yield ();
            mutex_lock (&mutex);
        }

        buffer [in] = nextProduced;
        in = (in + 1) % BUFFER_SIZE;
        count++;
        mutex_unlock (&mutex);
    }
}
```

# Improved consumer

```
void consumer (void *ignored) {
    for (;;) {
        mutex_lock (&mutex);
        while (count == 0) {
            mutex_unlock (&mutex);
            thread_yield ();
            mutex_lock (&mutex);
        }

        nextConsumed =  buffer[out];
        out = (out + 1) % BUFFER_SIZE;
        count--;
        mutex_unlock (&mutex);

        /* consume the item in nextConsumed */
    }
}
```

# Condition variables

- **Busy-waiting in application is a bad idea**
  - Thread consumes CPU even when can't make progress
  - Unnecessarily slows other threads and processes
- **Better to inform scheduler of which threads can run**
- **Typically done with condition variables**
- `void cond_init (cond_t *, ...);`
  - Initialize
- `void cond_wait (cond_t *c, mutex_t *m);`
  - Atomically unlock m and sleep until c signaled
  - Then re-acquire m and resume executing
- `void cond_signal (cond_t *c);`  
`void cond_broadcast (cond_t *c);`
  - Wake one/all threads waiting on c

# Improved producer

```
mutex_t mutex = MUTEX_INITIALIZER;
cond_t nonempty = COND_INITIALIZER;
cond_t nonfull = COND_INITIALIZER;

void producer (void *ignored) {
    for (;;) {
        /* produce an item and put in nextProduced */

        mutex_lock (&mutex);
        while (count == BUFFER_SIZE)
            cond_wait (&nonfull, &mutex);

        buffer [in] = nextProduced;
        in = (in + 1) % BUFFER_SIZE;
        count++;
        cond_signal (&nonempty);
        mutex_unlock (&mutex);
    }
}
```

# Improved consumer

```
void consumer (void *ignored) {  
    for (;;) {  
        mutex_lock (&mutex);  
        while (count == 0)  
            cond_wait (&nonempty, &mutex);  
  
        nextConsumed = buffer[out];  
        out = (out + 1) % BUFFER_SIZE;  
        count--;  
        cond_signal (&nonfull);  
        mutex_unlock (&mutex);  
  
        /* consume the item in nextConsumed */  
    }  
}
```



# Condition variables (continued)

- **Why must** `cond_wait` **both release mutex & sleep?**
- **Why not separate mutexes and condition variables?**

```
while (count == BUFFER_SIZE) {  
    mutex_unlock (&mutex);  
    cond_wait (&nonfull);  
    mutex_lock (&mutex);  
}
```

# Condition variables (continued)

- **Why must** `cond_wait` **both release mutex & sleep?**
- **Why not separate mutexes and condition variables?**

```
while (count == BUFFER_SIZE) {  
    mutex_unlock (&mutex);  
    cond_wait (&nonfull);  
    mutex_lock (&mutex);  
}
```

- **Can end up stuck waiting when bad interleaving**

PRODUCER

```
while (count == BUFFER_SIZE);  
mutex_unlock (&mutex);
```

```
cond_wait (&nonfull);
```

CONSUMER

```
mutex_lock (&mutex);
```

```
...
```

```
count--;
```

```
cond_signal (&nonfull);
```

# Implementing synchronization

- **User-visible mutex is straight-forward data structure**

```
typedef struct mutex {  
    bool is_locked;           /* true if locked */  
    thread_id_t owner;        /* thread holding lock, if locked */  
    thread_list_t waiters;    /* threads waiting for lock */  
  
    lower_level_lock_t lk; /* Protect above fields */  
};
```

- **Need lower-level lock lk for mutual exclusion**
  - Internally, mutex\_\* functions bracket code with  
lock(mutex->lk) . . . unlock(mutex->lk)
  - Otherwise, data races! (E.g., two threads manipulating waiters)
- **How to implement lower\_level\_lock\_t?**
  - Could use Peterson's algorithm, but typically a bad idea  
(too slow and don't know maximum number of threads)

# Approach #1: Disable interrupts

- **Only for apps with  $n : 1$  threads (1 kthread)**
  - Cannot take advantage of multiprocessors
  - But sometimes most efficient solution for uniprocessors
- **Have per-thread “do not interrupt” (DNI) bit**
- **lock (lk): sets thread’s DNI bit**
- **If timer interrupt arrives**
  - Check interrupted thread’s DNI bit
  - If DNI clear, preempt current thread
  - If DNI set, set “interrupted” (I) bit & resume current thread
- **unlock (lk): clears DNI bit *and* checks I bit**
  - If I bit is set, immediately yields the CPU

# Approach #2: Spinlocks

- **Most CPUs support atomic read-[modify-]write**
- **Example:** `int test_and_set (int *lockp);`
  - Atomically sets `*lockp = 1` and returns old value
  - Special instruction – can't be implemented in portable C
- **Use this instruction to implement *spinlocks*:**

```
#define lock(lockp)    while (test_and_set (lockp))
#define trylock(lockp) (test_and_set (lockp) == 0)
#define unlock(lockp) *lockp = 0
```
- **Spinlocks implement mutex's `lower_level_lock_t`**
- **Can you use spinlocks instead of mutexes?**
  - Wastes CPU, especially if thread holding lock not running
  - Mutex functions have short C.S., less likely to be preempted
  - On multiprocessor, sometimes good to spin for a bit, then yield

# Synchronization on x86

- **Test-and-set only one possible atomic instruction**
- **x86 xchg instruction, exchanges reg with mem**
  - Can use to implement test-and-set

```
_test_and_set:
    movl    8(%esp), %edx    # %edx = lockp
    movl    $1, %eax         # %eax = 1
    xchgl   %eax, (%edx)     # swap (%eax, *lockp)
    ret
```

- **CPU locks memory system around read and write**
  - Recall xchgl always acts like it has lock prefix
  - Prevents other uses of the bus (e.g., DMA)
- **Usually runs at memory bus speed, not CPU speed**
  - Much slower than cached read/buffered write

# Kernel Synchronization

- **Should kernel use locks or disable interrupts?**
- **Old UNIX had non-preemptive threads, no mutexes**
  - Interface designed for single CPU, so count++ etc. not data race
  - ... *Unless* memory shared with an interrupt handler

```
int x = splhigh (); // Disable interrupts
// Touch data shared with interrupt handler
splx (x);           // Restore previous state
```
  - C.f., **Pintos** `intr_disable` / `intr_set_level`
- **Used arbitrary pointers like condition variables**
  - `int [t]sleep (void *ident, int priority, ...);`  
put thread to sleep; will wake up at priority (`~cond_wait`)
  - `int wakeup (void *ident);`  
wake up all threads sleeping on `ident` (`~cond_broadcast`)

# Kernel locks

- **Nowadays, should design for multiprocessors**
  - Even if first version of OS is for uniprocessor
  - Someday may want multiple CPUs and need *preemptive* threads
  - That's why Pintos uses locks
- **Multiprocessor performance needs fine-grained locks**
  - Want to be able to call into the kernel on multiple CPUs
- **If kernel has locks, should it ever disable interrupts?**



# Kernel locks

- **Nowadays, should design for multiprocessors**
  - Even if first version of OS is for uniprocessor
  - Someday may want multiple CPUs and need *preemptive* threads
  - That's why Pintos uses locks
- **Multiprocessor performance needs fine-grained locks**
  - Want to be able to call into the kernel on multiple CPUs
- **If kernel has locks, should it ever disable interrupts?**
  - Yes! Can't sleep in interrupt handler, so can't wait for lock
  - So even modern OSes have support for disabling interrupts
  - Often uses DNI trick, which is cheaper than masking interrupts in hardware

# Semaphores [Dijkstra]

- **A *Semaphore* is initialized with an integer  $N$**
- **Provides two functions:**
  - `sem_wait (S)` (originally called  $P$ , called ***sema\_down*** in **Pintos**)
  - `sem_signal (S)` (originally called  $V$ , called ***sema\_up*** in **Pintos**)
- **Guarantees `sem_wait` will return only  $N$  more times than `sem_signal` called**
  - Example: If  $N == 1$ , then semaphore is a mutex with `sem_wait` as lock and `sem_signal` as unlock
- **Semaphores allow elegant solutions to some problems**

# Semaphore

A semaphore is a structure consisting of 2 parts:

```
struct semaphore {  
    int count; // number of resources available  
    queue Q; // queue of process/thread ids of blocked  
}
```

Shorthand notation:

semaphore  $S = 1 \rightarrow S.count = 1, S.Q = \{ \}$

# Operations on Semaphores

There are two basic semaphore operations:

`sem_wait(S):`

- if ( $S.count > 0$ ) then  $S.count = S.count - 1$ ;
- else block calling process in  $S.Q$ ;

`sem_signal(S):`

- if ( $S.Q$  is non-empty) then wakeup a process in  $S.Q$ ;
- else  $S.count = S.count + 1$ ;

# Semaphore Example: Mutual Exclusion

Semaphore S = 1;

Thread A:

`sem_wait(S);`

`(do work in critical section CS);`

`sem_signal(S);`

Thread B:

`sem_wait(S);`

`(do work in CS);`

`sem_signal(S);`

# Semaphore Example: Order Execution

Semaphore  $S = 0$ ;

Thread A  $\rightarrow$  Thread B:

Thread A:

(do work);

sem\_signal(S);

Thread B:

sem\_wait(S);

(do work);



# Semaphore producer/consumer

- **Can re-write producer/consumer to use three semaphores**
- **Semaphore** mutex **initialized to 1**
  - Used as mutex, protects buffer, in, out. . .
- **Semaphore** full **initialized to 0**
  - To block consumer when buffer empty
- **Semaphore** empty **initialized to N**
  - To block producer when queue full

```
void producer (void *ignored) {  
    for (;;) {  
        /* produce an item and put in nextProduced */  
        sem_wait (&empty);  
        sem_wait (&mutex);  
        buffer [in] = nextProduced;  
        in = (in + 1) % BUFFER_SIZE;  
        sem_signal (&mutex);  
        sem_signal (&full);  
    }  
}
```

```
void consumer (void *ignored) {  
    for (;;) {  
        sem_wait (&full);  
        sem_wait (&mutex);  
        nextConsumed = buffer[out];  
        out = (out + 1) % BUFFER_SIZE;  
        sem_signal (&mutex);  
        sem_signal (&empty);  
        /* consume the item in nextConsumed */  
    }  
}
```



# Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let **S** and **Q** be two semaphores initialized to 1

$P_0$		$P_1$
wait (S);		wait (Q);
wait (Q);		wait (S);
	.	.
signal (S);		signal (Q);
signal (Q);		signal (S);

- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended
- **Priority Inversion** - Scheduling problem when lower-priority process holds a lock needed by higher-priority process

# Classical Synchronization Problems

- Bounded-Buffer (Producer-Consumer) Problem
- Readers and Writers Problem
- Dining-Philosophers Problem

# Readers-Writers Problem

- A data set is shared among a number of concurrent processes
  - Readers – only read the data set; they do **not** perform any updates
  - Writers – can both read and write
- Problem – allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time
- Shared Data
  - Data set
  - Semaphore **mutex** initialized to 1
  - Semaphore **wrt** initialized to 1
  - Integer **readcount** initialized to 0

# Readers-Writers Problem (cont.)

- The structure of a writer process


```
do {  
    wait (wrt) ;  
  
    //  writing is performed  
  
    signal (wrt) ;  
} while (TRUE);
```

# Readers-Writers Problem (cont.)

- The structure of a reader process

```
do {  
    wait (mutex) ;  
    readcount ++ ;  
    if (readcount == 1)  
        wait (wrt) ;  
    signal (mutex)  
        // reading is performed  
    wait (mutex) ;  
    readcount - - ;  
    if (readcount == 0)  
        signal (wrt) ;  
    signal (mutex) ;  
} while (TRUE);
```

# Dining-Philosophers Problem



- Shared data
  - Bowl of rice (data set)
  - Semaphore **chopstick [5]** each initialized to 1

# Dining-Philosophers Problem (cont.)

- The structure of Philosopher  $i$ :

```
do {  
    wait ( chopstick[i] );  
    wait ( chopStick[ (i + 1) % 5] );  
  
    // eat  
  
    signal ( chopstick[i] );  
    signal ( chopstick[ (i + 1) % 5] );  
  
    // think  
  
} while (TRUE);
```

# Problems with Semaphores

- Correct use of semaphore operations:
  - signal (mutex) .... wait (mutex)
  - wait (mutex) ... wait (mutex)
  - Omitting of wait (mutex) or signal (mutex) (or both)




# Monitors

- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- Only one process may be active within the monitor at a time

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { .... }
    ...
    procedure Pn (...) {.....}

    initialization code ( ....) { ... }
    ...
}
```


# Schematic view of a Monitor



# Condition Variables

- `condition x, y;`
- Two operations on a condition variable:
  - `x.wait ()` – a process that invokes the operation is suspended.
  - `x.signal ()` – resumes one of processes (if any) that invoked `x.wait ()`

# Monitor with Condition Variables



# Solution to Dining Philosophers

monitor DP

```
{  
    enum { THINKING; HUNGRY, EATING) state [5] ;  
    condition self [5];  
  
    void pickup (int i) {  
        state[i] = HUNGRY;  
        test(i);  
        if (state[i] != EATING) self [i].wait;  
    }  
  
    void putdown (int i) {  
        state[i] = THINKING;  
        // test left and right neighbors  
        test((i + 4) % 5);  
        test((i + 1) % 5);  
    }  
}
```

# Solution to Dining Philosophers (cont)

```
void test (int i) {  
    if ( (state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) ) {  
        state[i] = EATING ;  
        self[i].signal () ;  
    }  
}
```

```
    initialization_code() {  
        for (int i = 0; i < 5; i++)  
            state[i] = THINKING;  
    }  
}
```

# Solution to Dining Philosophers (cont)

- Each philosopher *i* invokes the operations `pickup()` and `putdown()` in the following sequence:

`DiningPhilosophers.pickup (i);`

`EAT`

`DiningPhilosophers.putdown (i);`

# Monitor Implementation Using Semaphores

- Variables

```
semaphore mutex; // (initially = 1)
semaphore next;  // (initially = 0)
int next-count = 0;
```

- Each procedure  $F$  will be replaced by

```
wait(mutex);
...
body of  $F$ ;
...
if (next_count > 0)
    signal(next)
else
    signal(mutex);
```

- Mutual exclusion within a monitor is ensured.



# Monitor Implementation

- For each condition variable  $x$ , we have:

```
semaphore x_sem; // (initially = 0)  
int x-count = 0;
```

- The operation  $x.wait$  can be implemented as:

```
x-count++;  
if (next_count > 0)  
    signal(next);  
else  
    signal(mutex);  
wait(x_sem);  
x-count--;
```

# Monitor Implementation

- The operation `x.signal` can be implemented as:

```
if (x-count > 0) {  
    next_count++;  
    signal(x_sem);  
    wait(next);  
    next_count--;  
}
```

# Linux Synchronization

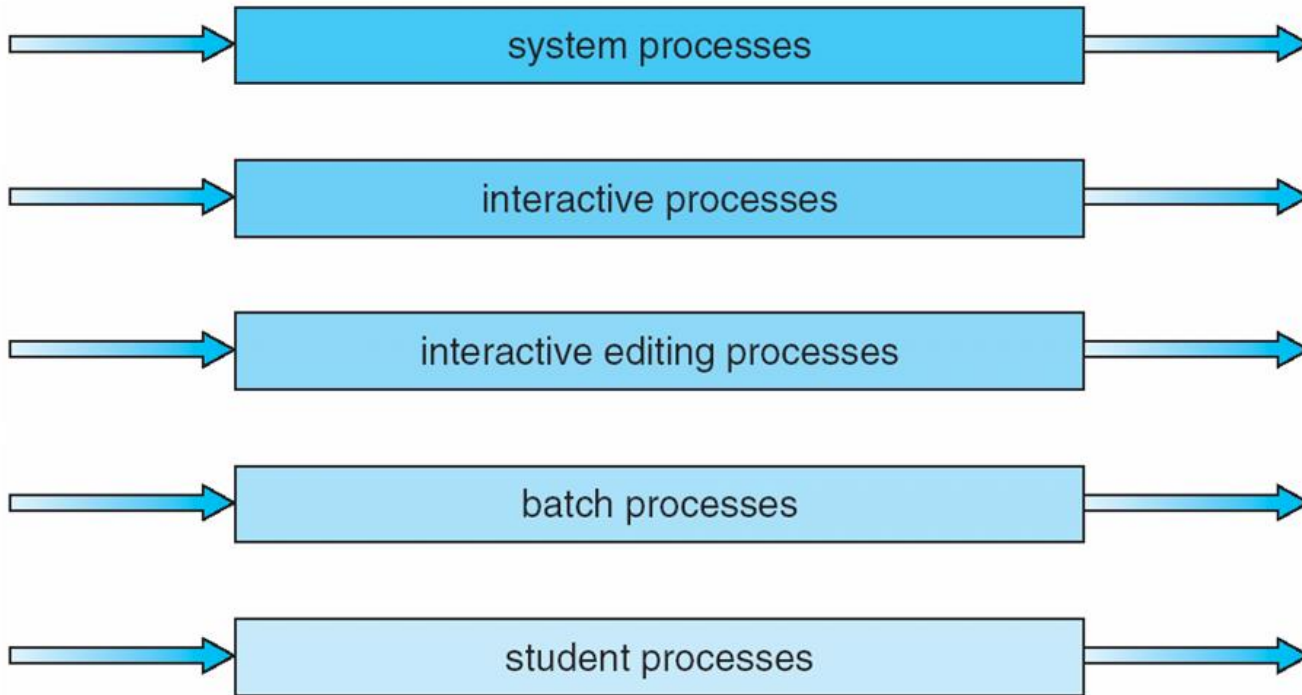
- Linux:
  - Prior to kernel Version 2.6, disables interrupts to implement short critical sections
  - Version 2.6 and later, fully preemptive
- Linux provides:
  - semaphores
  - spin locks

# Pthreads Synchronization

- Pthreads API is OS-independent
- It provides:
  - mutex locks
  - condition variables
- Non-portable extensions include:
  - read-write locks
  - spin locks

# Multilevel Queue Scheduling

highest priority



lowest priority


# Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
  - number of queues
  - scheduling algorithms for each queue
  - method used to determine when to upgrade a process
  - method used to determine when to demote a process
  - method used to determine which queue a process will enter when that process needs service

# Example of Multilevel Feedback Queue

- Three queues:
  - $Q_0$  – RR with time quantum 8 milliseconds
  - $Q_1$  – RR time quantum 16 milliseconds
  - $Q_2$  – FCFS
- Scheduling
  - A new job enters queue  $Q_0$  which is served RR. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue  $Q_1$ .
  - At  $Q_1$  job is again served RR and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue  $Q_2$ .

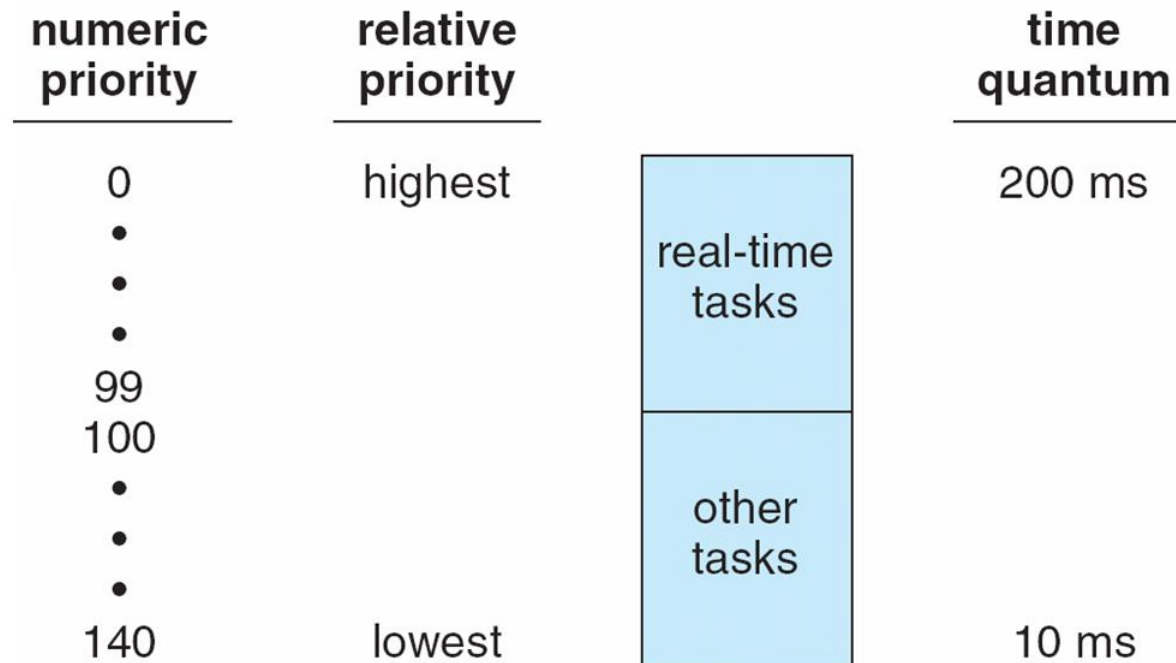
# Multilevel Feedback Queues





# Linux Scheduling

- Constant order  $O(1)$  scheduling time
- Two priority ranges: time-sharing and real-time
- **Real-time** range from 0 to 99 and **nice** value from 100 to 140



# Summary

- Read Ch. 1-6
- Processes and Threads (Ch. 4)
- Process Scheduling (Ch. 5)
- Synchronization (Ch. 6)
- Project 1 – Scheduling and Synchronization