

# Lecture 8: Deadlocks

**Instructor: Mitch Neilsen**

**Office: N219D**

# Outline

## ■ Reading:

- Ch. 4 - Threads
- Ch. 5 - CPU Scheduling
- Ch. 6 - Synchronization
- Ch. 7 - Deadlocks



## ■ Project 1: Scheduling and Synchronization

- Alarm Clock
- Priority-based Scheduler
- Synchronization and Priority Inheritance
- [Extra Credit] MLFQ Scheduler

## Quote of the Day

"Man's greatest asset is the unsettled mind. "

- Isaac Asimov

# Scheduling Criteria

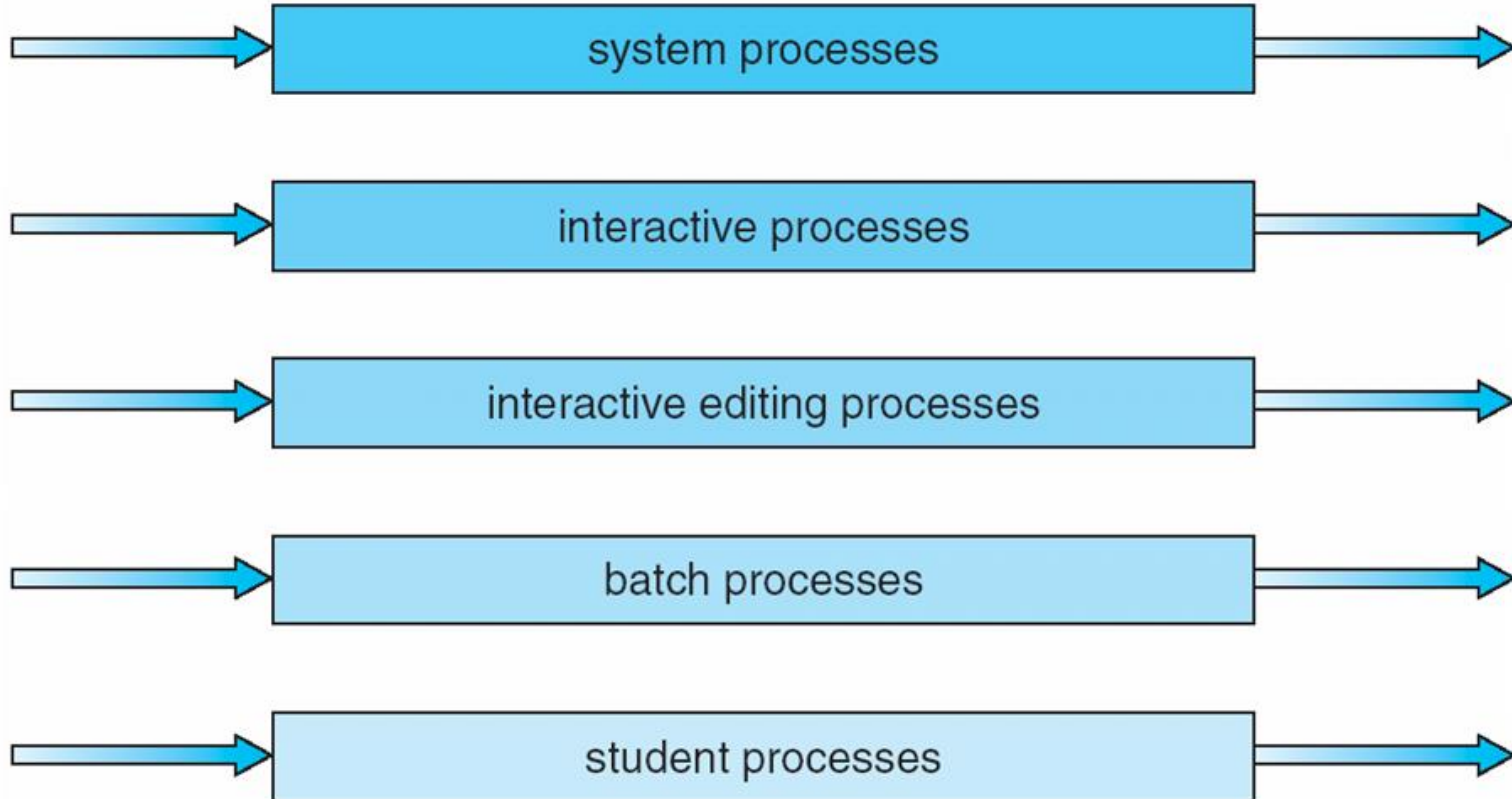
- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – # of processes that complete their execution per time unit
- **Turnaround time** – amount of time to execute a particular process
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

# Multilevel Queue

- Ready queue is partitioned into separate queues:
  - foreground (interactive)
  - background (batch)
- Each queue has its own scheduling algorithm
  - foreground – RR
  - background – FCFS
- Scheduling must be done between the queues
  - Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
  - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR, and 20% to background in FCFS

# Multilevel Queue Scheduling

highest priority



lowest priority

# Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
  - number of queues
  - scheduling algorithms for each queue
  - method used to determine when to upgrade a process
  - method used to determine when to demote a process
  - method used to determine which queue a process will enter when that process needs service

# Example of Multilevel Feedback Queue

## ■ Three queues:


- $Q_0$  – RR with time quantum 8 milliseconds
- $Q_1$  – RR time quantum 16 milliseconds
- $Q_2$  – FCFS

## ■ Scheduling

- A new job enters queue  $Q_0$  which is served RR. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue  $Q_1$ .
- At  $Q_1$  job is again served RR and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue  $Q_2$ .



# Multilevel Feedback Queues



# Synchronization

- Processes A and B complete before C?

```
semaphore AtoC = 0;
```

```
semaphore BtoC = 0;
```

Process A:

- do work

```
sem_signal(AtoC);
```

Process B:

- do work


```
sem_signal(BtoC);
```

Process C:

```
sem_wait(AtoC);
```

```
sem_wait(BtoC);
```

- do work



# Chapter 7: Deadlocks

- The Deadlock Problem
- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock

# The Deadlock Problem

- A set of blocked processes each holds a resource and waits to acquire a resource held by another process in the set.
- The following program can fail to make progress – how?

```
mutex_t m1, m2;

void p1 (void *ignored) {
    lock (m1);
    lock (m2);
    /* critical section */
    unlock (m2);
    unlock (m1);
}

void p2 (void *ignored) {
    lock (m2);
    lock (m1);
    /* critical section */
    unlock (m1);
    unlock (m2);
}
```

- Do all deadlocks involve semaphores?

# More Deadlocks

- **Same problem with condition variables**

- Suppose resource 1 managed by  $c1$ , resource 2 by  $c2$
- A has resource 1, waits on  $c2$ , B has resource 2, waits on  $c1$


- **Or have combined mutex/condition variable deadlock:**

- lock (a); lock (b); while (!ready) wait (b, c);  
unlock (b); unlock (a);
- lock (a); lock (b); ready = true; signal (c);  
unlock (b); unlock (a);

- **One lesson: Dangerous to hold locks when crossing abstraction barriers!**

- i.e., lock (a) then call function that uses condition variable

# Bridge Crossing Example



- Traffic only in one direction
- Each section of a bridge can be viewed as a resource
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback)
- Several cars may have to be backed up if a deadlock occurs
- Starvation is possible
- Note – Most OSes do not prevent or deal with deadlocks

# System Model

- Resource types  $R_1, R_2, \dots, R_m$   
*CPU cycles, memory space, I/O devices*
- Each resource type  $R_i$  has  $W_i$  instances.
- Each process utilizes a resource as follows:
  - **request**
  - **use**
  - **release**

# Deadlock Characterization

Deadlock can arise if four conditions hold simultaneously.

- **Mutual exclusion:** only one process at a time can use a resource
- **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes
- **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task
- **Circular wait:** there exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .



# Resource-Allocation Graph

A set of vertices  $V$  and a set of edges  $E$ .

- $V$  is partitioned into two types:
  - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system
  - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system
- **request edge** – directed edge  $P_i \rightarrow R_j$
- **assignment edge** – directed edge  $R_j \rightarrow P_k$

# Resource-Allocation Graph (Cont.)


- Process




- Resource Type with 4 instances




- $P_i$  requests instance of  $R_j$




- $P_i$  is holding an instance of  $R_j$




# Example of a Resource Allocation Graph



# Resource Allocation Graph With A Deadlock



# Is This Deadlock?



# Basic Facts

- If graph contains no cycles  $\Rightarrow$  no deadlock
- If graph contains a cycle  $\Rightarrow$ 
  - if only one instance per resource type, then deadlock
  - if several instances per resource type, possibility of deadlock

# Methods for Handling Deadlocks

- Ensure that the system will *never* enter a deadlock state
- Allow the system to enter a deadlock state and then recover
- Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX

# Deadlock Prevention

Restrain the ways request can be made

- **Mutual Exclusion** – not required for sharable resources; must hold for non-sharable resources
- **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
  - Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none
  - Low resource utilization; starvation possible



# Deadlock Prevention (Cont.)

## ■ **No Preemption** –

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released
- Preempted resources are added to the list of resources for which the process is waiting
- Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting

## ■ **Circular Wait** – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration

# Deadlock Avoidance

Requires that the system has some additional ***a priori*** information available

- Simplest and most useful model requires that each process declare the ***maximum number*** of resources of each type that it may need
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition
- Resource-allocation ***state*** is defined by the number of available and allocated resources, and the maximum demands of the processes


# Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state
- System is **in safe state** if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL the processes in the systems such that for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , with  $j < i$
- That is:
  - If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished
  - When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on

# Basic Facts

- If a system is in safe state  $\Rightarrow$  no deadlocks
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock
- Deadlock Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state; i.e., conservatively prohibit system from entering a state that may lead to deadlock.

# Safe, Unsafe , Deadlock State




# Deadlock Avoidance Algorithms

- Single instance of a resource type
  - Use a resource-allocation graph
  - Check for cycles
- Multiple instances of a resource type
  - Use the Banker's Algorithm

# Resource-Allocation Graph Scheme

- **Claim edge**  $P_i \text{-----} \rightarrow R_j$  indicates that process  $P_i$  may request resource  $R_j$ ; represented by a dashed line
- Claim edge converts to request edge when a process requests a resource
- Request edge converted to an assignment edge when the resource is allocated to the process
- When a resource is released by a process, assignment edge reconverts to a claim edge
- Resources must be claimed *a priori* in the system


# Resource-Allocation Graph



Dashed line is a *claim edge*, the processes may request  $R_2$



# Unsafe State In Resource-Allocation Graph



## **Note cycle in graph**

- $P_1$  might request  $R_2$  before relinquishing  $R_1$
- Would cause deadlock

# Resource-Allocation Graph Algorithm

- Suppose that process  $P_i$  requests a resource  $R_j$
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph

# Banker's Algorithm

- Multiple instances
- Each process must *a priori* report maximum possible resource use for each resource type
- When a process requests a resource it may have to wait even if the resource is available
- When a process gets all its resources it must return them in a finite amount of time

# Data Structures for Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resources types.

- **Available:** Vector of length  $m$ . If  $available[j] = k$ , there are  $k$  instances of resource type  $R_j$  available
- **Max:**  $n \times m$  matrix. If  $Max[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$
- **Allocation:**  $n \times m$  matrix. If  $Allocation[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$
- **Need:**  $n \times m$  matrix. If  $Need[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task

$$Need[i,j] = Max[i,j] - Allocation[i,j]$$

# Safety Algorithm

1. Let *Work* and *Finish* be vectors of length *m* and *n*, respectively. Initialize:

*Work* = *Available*

*Finish* [*i*] = *false* for *i* = 0, 1, ..., *n*- 1

2. Find and *i* such that both:

(a) *Finish* [*i*] = *false*

(b)  $Need_i \leq Work$

If no such *i* exists, go to step 4

3. *Work* = *Work* + *Allocation*<sub>*i*</sub>

*Finish*[*i*] = *true*

go to step 2

4. If *Finish* [*i*] == *true* for all *i*, then the system is in a safe state

# Resource-Request Algorithm for Process $P_i$

$Request$  = request vector for process  $P_i$ .

If  $Request_i[j] = k$  then process  $P_i$  wants  $k$  instances of resource type  $R_j$

1. If  $Request_i \leq Need_i$  go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If  $Request_i \leq Available$ , go to step 3. Otherwise  $P_i$  must wait, since resources are not available
3. Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:

$$Available = Available - Request_i;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

- If safe  $\Rightarrow$  the resources are allocated to  $P_i$
- If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored

# Example of Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$ ; 3 resource types:  $A$ ,  $B$ , and  $C$  ( $A$  (10 instances),  $B$  (5 instances), and  $C$  (7 instances))

Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
$P_0$	0 1 0	7 5 3	3 3 2
$P_1$	2 0 0	3 2 2	
$P_2$	3 0 2	9 0 2	
$P_3$	2 1 1	2 2 2	
$P_4$	0 0 2	4 3 3	

## Example (Cont.)

- The content of the matrix *Need* is defined to be  $Need = Max - Allocation$

	<u>Need</u>		
	A	B	C
$P_0$	7	4	3
$P_1$	1	2	2
$P_2$	6	0	0
$P_3$	0	1	1
$P_4$	4	3	1

- The system is in a safe state since the sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria



## Example: $P_1$ Request (1,0,2)

- Check that Request  $\leq$  Available (that is,  $(1,0,2) \leq (3,3,2) \Rightarrow$  true, after the request remaining available is  $(2,3,0)$ )

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	2 3 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 1	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

- Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement
- Can a request for  $(3,3,0)$  by  $P_4$  be granted safely?
- Can a request for  $(0,2,0)$  by  $P_0$  be granted safely?

## Example: $P_1$ Request (1,0,2)

- Check that Request  $\leq$  Available (that is,  $(1,0,2) \leq (3,3,2) \Rightarrow$  true, after the request remaining available is  $(2,3,0)$ )

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	<b>2 3 0</b>
<b><math>P_1</math></b>	<b>3 0 2</b>	<b>0 2 0</b>	
$P_2$	3 0 1	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

- Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement
- Can a request for  $(3,3,0)$  by  $P_4$  be granted safely?
- Can a request for  $(0,2,0)$  by  $P_0$  be granted safely?

## Example: $P_1$ Request (1,0,2)

- Check that Request  $\leq$  Available (that is,  $(1,0,2) \leq (3,3,2) \Rightarrow$  true, after the request remaining available is  $(2,3,0)$ )

	<u>Allocation</u>			<u>Need</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	4	3	<b>5</b>	<b>3</b>	<b>2</b>
<b><math>P_1</math></b>	<b>0</b>	<b>0</b>	<b>0</b>	<b>3</b>	<b>2</b>	<b>2</b>			
$P_2$	3	0	1	6	0	0			
$P_3$	2	1	1	0	1	1			
$P_4$	0	0	2	4	3	1			

- Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement
- Can a request for  $(3,3,0)$  by  $P_4$  be granted safely?
- Can a request for  $(0,2,0)$  by  $P_0$  be granted safely?

## Example: $P_1$ Request (1,0,2)

- Check that Request  $\leq$  Available (that is,  $(1,0,2) \leq (3,3,2) \Rightarrow$  true, after the request remaining available is  $(2,3,0)$ )

	<u>Allocation</u>			<u>Need</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	7	4	3	<b>5</b>	<b>3</b>	<b>2</b>
<b><math>P_1</math></b>	<b>0</b>	<b>0</b>	<b>0</b>	<b>3</b>	<b>2</b>	<b>2</b>			
$P_2$	3	0	1	6	0	0			
<b><math>P_3</math></b>	<b>2</b>	<b>1</b>	<b>1</b>	<b>0</b>	<b>1</b>	<b>1</b>			
$P_4$	0	0	2	4	3	1			

- Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement
- Can a request for  $(3,3,0)$  by  $P_4$  be granted safely?
- Can a request for  $(0,2,0)$  by  $P_0$  be granted safely?

## Example: $P_1$ Request (1,0,2)

- Check that Request  $\leq$  Available (that is,  $(1,0,2) \leq (3,3,2) \Rightarrow$  true, after the request remaining available is  $(2,3,0)$ )

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	<b>7 4 3</b>
<b><math>P_1</math></b>	<b>0 0 0</b>	<b>3 2 2</b>	
$P_2$	3 0 1	6 0 0	
<b><math>P_3</math></b>	<b>0 0 0</b>	<b>2 2 2</b>	
$P_4$	0 0 2	4 3 1	

- Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement
- Can a request for  $(3,3,0)$  by  $P_4$  be granted safely?
- Can a request for  $(0,2,0)$  by  $P_0$  be granted safely?


# Deadlock Detection

- Allow system to enter a deadlock state
- Apply a Deadlock Detection Algorithm to find it
- Then, apply a Recovery Scheme to recover from it


# Single instance of each Resource Type

- Maintain *wait-for* graph
  - Nodes are processes
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$
- Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock
- An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph

# Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph



Corresponding Wait-for Graph



# Several Instances of a Resource Type

- **Available:** A vector of length  $m$  indicates the number of available resources of each type.
- **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process.
- **Request:** An  $n \times m$  matrix indicates the current request of each process. If  $Request[i, j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .

# Detection Algorithm

1. Let *Work* and *Finish* be vectors of length  $m$  and  $n$ , respectively  
Initialize:
  - (a) *Work* = *Available*
  - (b) For  $i = 1, 2, \dots, n$ , if  $Allocation_i \neq 0$ , then  
*Finish*[ $i$ ] = false; otherwise, *Finish*[ $i$ ] = true
2. Find an index  $i$  such that both:
  - (a) *Finish*[ $i$ ] == false
  - (b)  $Request_i \leq Work$

If no such  $i$  exists, go to step 4

## Detection Algorithm (Cont.)

3.  $Work = Work + Allocation_i$   
 $Finish[i] = true$   
go to step 2
4. If  $Finish[i] == false$ , for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state. Moreover, if  $Finish[i] == false$ , then  $P_i$  is deadlocked

**Algorithm requires on the order of  $O(m \times n^2)$  operations to detect whether the system is in deadlocked state**

# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time  $T_0$ :

	<u>Allocation</u>				<u>Request</u>			<u>Available</u>		
	A	B	C		A	B	C	A	B	C
$P_0$	0	1	0		0	0	0	0	0	0
$P_1$	2	0	0		2	0	2			
$P_2$	3	0	3		0	0	0			
$P_3$	2	1	1		1	0	0			
$P_4$	0	0	2		0	0	2			

- Sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$

# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 0 0	0 0 0	0 1 0
$P_1$	2 0 0	2 0 2	
$P_2$	3 0 3	0 0 0	
$P_3$	2 1 1	1 0 0	
$P_4$	0 0 2	0 0 2	

- Sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$

# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 0 0	0 0 0	<b>3 1 3</b>
$P_1$	2 0 0	2 0 2	
<b><math>P_2</math></b>	<b>0 0 0</b>	<b>0 0 0</b>	
$P_3$	2 1 1	1 0 0	
$P_4$	0 0 2	0 0 2	

- Sequence  $\langle P_0, \mathbf{P_2}, P_3, P_1, P_4 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$

# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 0 0	0 0 0	<b>5 2 4</b>
$P_1$	2 0 0	2 0 2	
$P_2$	0 0 0	0 0 0	
<b><math>P_3</math></b>	<b>0 0 0</b>	<b>0 0 0</b>	
$P_4$	0 0 2	0 0 2	

- Sequence  $\langle P_0, P_2, \mathbf{P_3}, P_1, P_4 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$

# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 0 0	0 0 0	<b>7 2 4</b>
<b><math>P_1</math></b>	<b>0 0 0</b>	<b>0 0 0</b>	
$P_2$	0 0 0	0 0 0	
$P_3$	0 0 0	0 0 0	
$P_4$	0 0 2	0 0 2	

- Sequence  $\langle P_0, P_2, P_3, \mathbf{P_1}, P_4 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$



# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 0 0	0 0 0	<b>7 2 6</b>
$P_1$	0 0 0	0 0 0	
$P_2$	0 0 0	0 0 0	
$P_3$	0 0 0	0 0 0	
$P_4$	<b>0 0 0</b>	<b>0 0 0</b>	

- Sequence  $\langle P_0, P_2, P_3, P_1, \mathbf{P_4} \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$

# Example (Cont.)

- $P_2$  requests an additional instance of type C

Request

A B C

$P_0$  0 0 0

$P_1$  2 0 1

$P_2$  0 0 1

$P_3$  1 0 0

$P_4$  0 0 2

- State of system?
  - Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes; requests
  - Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$

# Example (cont.)

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	0 0 0	0 0 0
$P_1$	2 0 0	2 0 2	
$P_2$	3 0 3	0 0 1	
$P_3$	2 1 1	1 0 0	
$P_4$	0 0 2	0 0 2	

- Sequence  $\langle P_0 \rangle$  will result in  $Finish[i] = \text{false}$  for all  $i=1,2,3,4$

# Example (cont.)

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances)
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 0 0	0 0 0	0 1 0
$P_1$	2 0 0	2 0 2	
$P_2$	3 0 3	0 0 1	
$P_3$	2 1 1	1 0 0	
$P_4$	0 0 2	0 0 2	

- Sequence  $\langle P_0 \rangle$  will result in  $Finish[i] = \text{false}$  for all  $i=1,2,3,4$

# Detection-Algorithm Usage

- When, and how often, to invoke depends on:
  - How often a deadlock is likely to occur?
  - How many processes will need to be rolled back?
    - ▶ one for each disjoint cycle
- If detection algorithm is invoked arbitrarily, there may be many cycles in the resource graph and so we would not be able to tell which of the many deadlocked processes “caused” the deadlock

# Recovery from Deadlock: Process Termination

- Abort all deadlocked processes
- Abort one process at a time until the deadlock cycle is eliminated
- In which order should we choose to abort?
  - Priority of the process
  - How long process has computed, and how much longer to completion
  - Resources the process has used
  - Resources process needs to complete
  - How many processes will need to be terminated
  - Is process interactive or batch?

# Recovery from Deadlock: Resource Preemption

- Selecting a victim(s) – minimize cost
- Rollback – return to some safe state, restart process for that state
- Starvation – same process may always be picked as victim, include number of rollbacks in cost factor

# Summary

- Read Ch. 1-6
- Processes and Threads (Ch. 4)
- Process Scheduling (Ch. 5)
- Synchronization (Ch. 6)
- Project 1 – Scheduling and Synchronization