

Lecture 14: Dynamic Memory

Instructor: Mitch Neilsen

Office: N219D

Quote of the Day

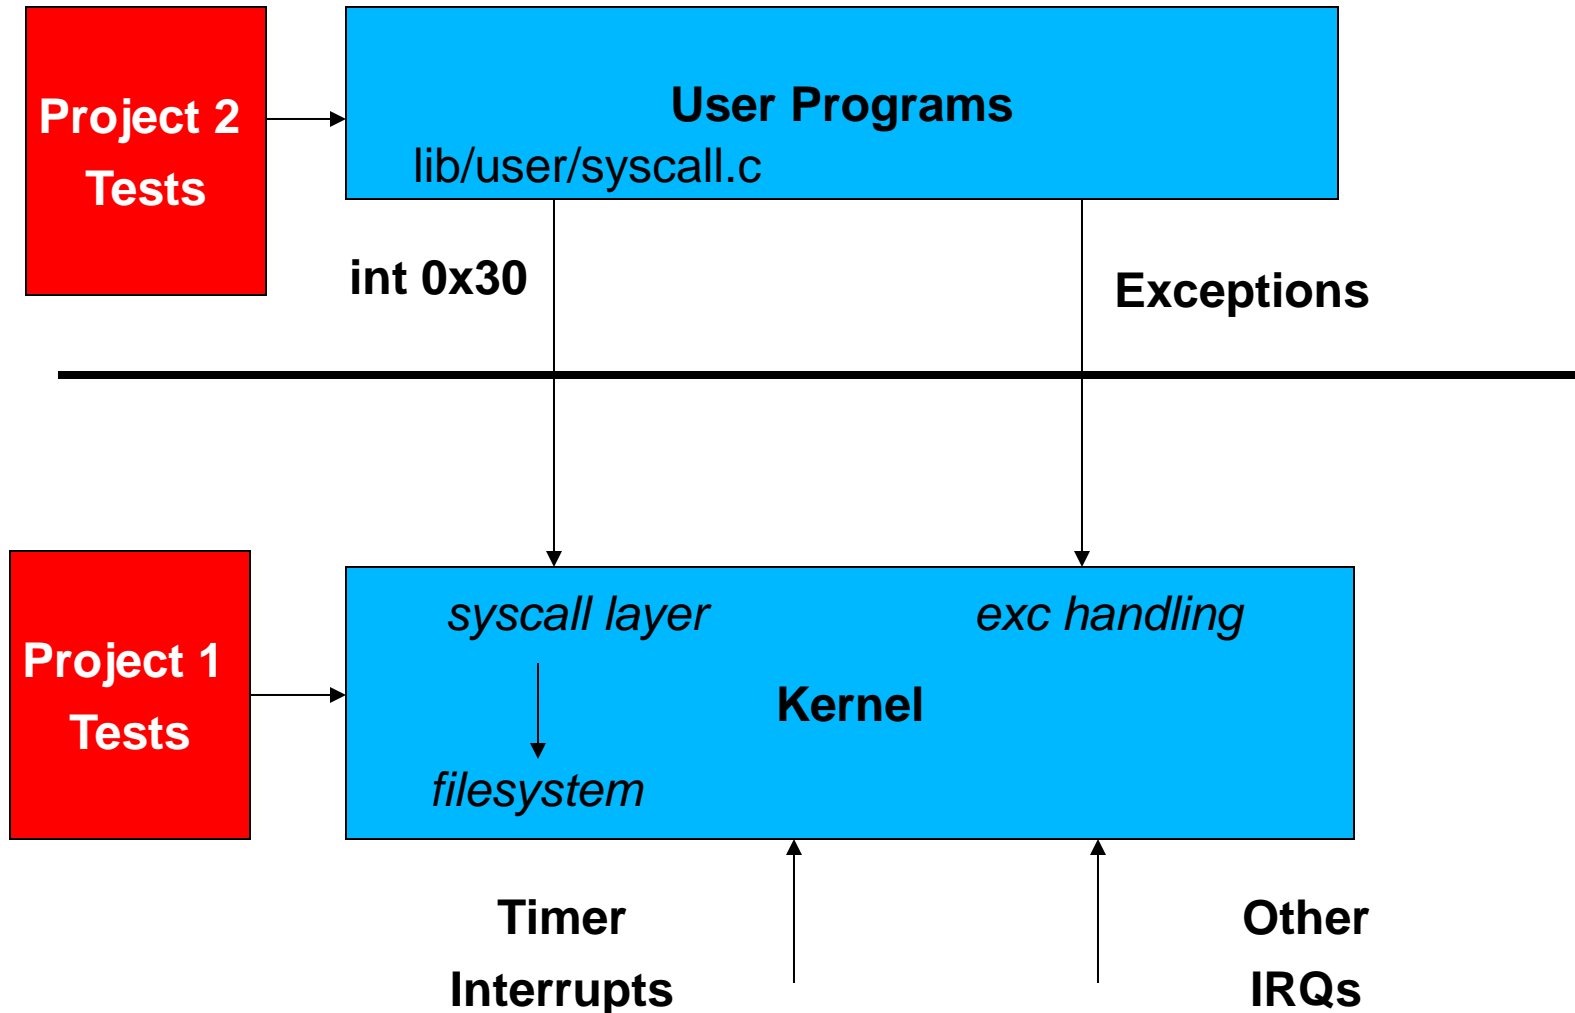
"I have not failed. I've just found 10,000 ways that won't work."

-- Thomas Edison

Project 2

- System Call: `void halt(void)`
- TYPO: `power_off()` should be `shutdown_power_off()` .

Project 1 and Project 2



When does a process need to access Operating System functionality?

- Here are several examples:
 - Reading a file. The OS must perform the file system operations required to read the data off of disk.
 - Creating a child process. The OS must set stuff up for the child process.

How do processes invoke OS functionality?

- By making a system call.
 - Conceptually, processes call a subroutine that goes off and performs the required functionality. But OS must execute in supervisor mode, which allows it to do things like manipulate the disk directly.
 - To switch from normal user mode to supervisor mode, most machines provide a system call instruction.
 - ▶ **This instruction causes an exception to take place.**
 - ▶ **The hardware switches from user mode to supervisor mode and invokes the exception handler inside the operating system.**
 - There is typically some kind of convention that the process uses to interact with the OS.

Let's do an example - Open() system call

- System calls typically start out with a normal subroutine call.

```
int handle = open("sample.txt");
```

- **open()** executes a syscall instruction, which generates a system call exception 0x30.

```
syscall11 (SYS_OPEN, file);
```

- By convention, the Open subroutine puts a number on the stack to indicate which routine (SYS_OPEN = 6) should be invoked.
 - ▶ Inside the exception handler the OS looks at the stack to figure out what system call it should perform.
- The Open system call also takes a parameter. By convention, the compiler also puts this (e.g., a pointer to the filename) on the stack.
 - ▶ More conventions: return values are put into the %EAX register.
- Inside the exception handler, the OS figures out what action to take, performs the action, then returns back to the user program.

SYS_OPEN is defined in lib/syscall-nr.h

```
#ifndef __LIB_SYSCALL_NR_H
#define __LIB_SYSCALL_NR_H

/* System call numbers. */
enum
{
    /* Projects 2 and later. */
    SYS_HALT,          /* Halt the operating system. */
    SYS_EXIT,          /* Terminate this process. */
    SYS_EXEC,          /* Start another process. */
    SYS_WAIT,          /* Wait for a child process to die. */
    SYS_CREATE,        /* Create a file. */
    SYS_REMOVE,        /* Delete a file. */
    SYS_OPEN,          /* Open a file. */
    ...
}
```

Thus, SYS_OPEN = 6.

Open() system call details

In pintos/src/lib/user/syscall.c:

```
int
open (const char *file)
{
    return syscall1 (SYS_OPEN, file);
}

.. (and above) ..

/* Invokes syscall NUMBER, passing argument ARG0, and returns the
   return value as an `int'. */
#define syscall1(NUMBER, ARG0) \
    ({ \
        int retval; \
        asm volatile \
            ("pushl %[arg0]; pushl %[number]; int $0x30; addl $8, %%esp" \
             : "=a" (retval) \
             : [number] "i" (NUMBER), \
               [arg0] "g" (ARG0) \
             : "memory"); \
        retval;
```

Initialize syscall handler

```
void  
syscall_init (void)  
{  
    intr_register_int (0x30, 3, INTR_ON, syscall_handler, "syscall");  
    lock_init (&fs_lock);  
}
```

Add sys_open to syscall_handler

```
static void
syscall_handler (struct intr_frame *f)
{
    typedef int syscall_function (int, int, int);

    /* A system call. */
    struct syscall
    {
        size_t arg_cnt;           /* Number of arguments. */
        syscall_function *func;   /* Implementation. */
    };

    /* Table of system calls. */
    static const struct syscall syscall_table[] =
    {
        {0, (syscall_function *) sys_halt},
        {1, (syscall_function *) sys_exit},
        {1, (syscall_function *) sys_exec},
        {1, (syscall_function *) sys_wait},
        {2, (syscall_function *) sys_create},
        {1, (syscall_function *) sys_remove},
        {1, (syscall_function *) sys_open},
        ...
    }
```

Add sys_open to syscall_handler

```
const struct syscall *sc;
unsigned call_nr;
int args[3];
/* Get the system call. */
copy_in (&call_nr, f->esp, sizeof call_nr);
if (call_nr >= sizeof syscall_table / sizeof *syscall_table)
    thread_exit ();
sc = syscall_table + call_nr;

/* Get the system call arguments. */
ASSERT (sc->arg_cnt <= sizeof args / sizeof *args);
memset (args, 0, sizeof args);
copy_in (args, (uint32_t *) f->esp + 1, sizeof *args * sc->arg_cnt);

/* Execute the system call,
   and set the return value. */
f->eax = sc->func (args[0], args[1], args[2]);
}
```

Add sys_open function to syscall.c

```
sys_open (const char *ufile)
{
    char *kfile = copy_in_string (ufile);
    struct file_descriptor *fd;
    int handle = -1;

    fd = malloc (sizeof *fd);
    if (fd != NULL)
    {
        lock_acquire (&fs_lock);
        fd->file = filesys_open (kfile);
        if (fd->file != NULL)
        {
            struct thread *cur = thread_current ();
            handle = fd->handle = cur->next_handle++;
            list_push_front (&cur->fds, &fd->elem);
        }
        else
            free (fd);
        lock_release (&fs_lock);
    }
    palloc_free_page (kfile);
    return handle;
}
```

```

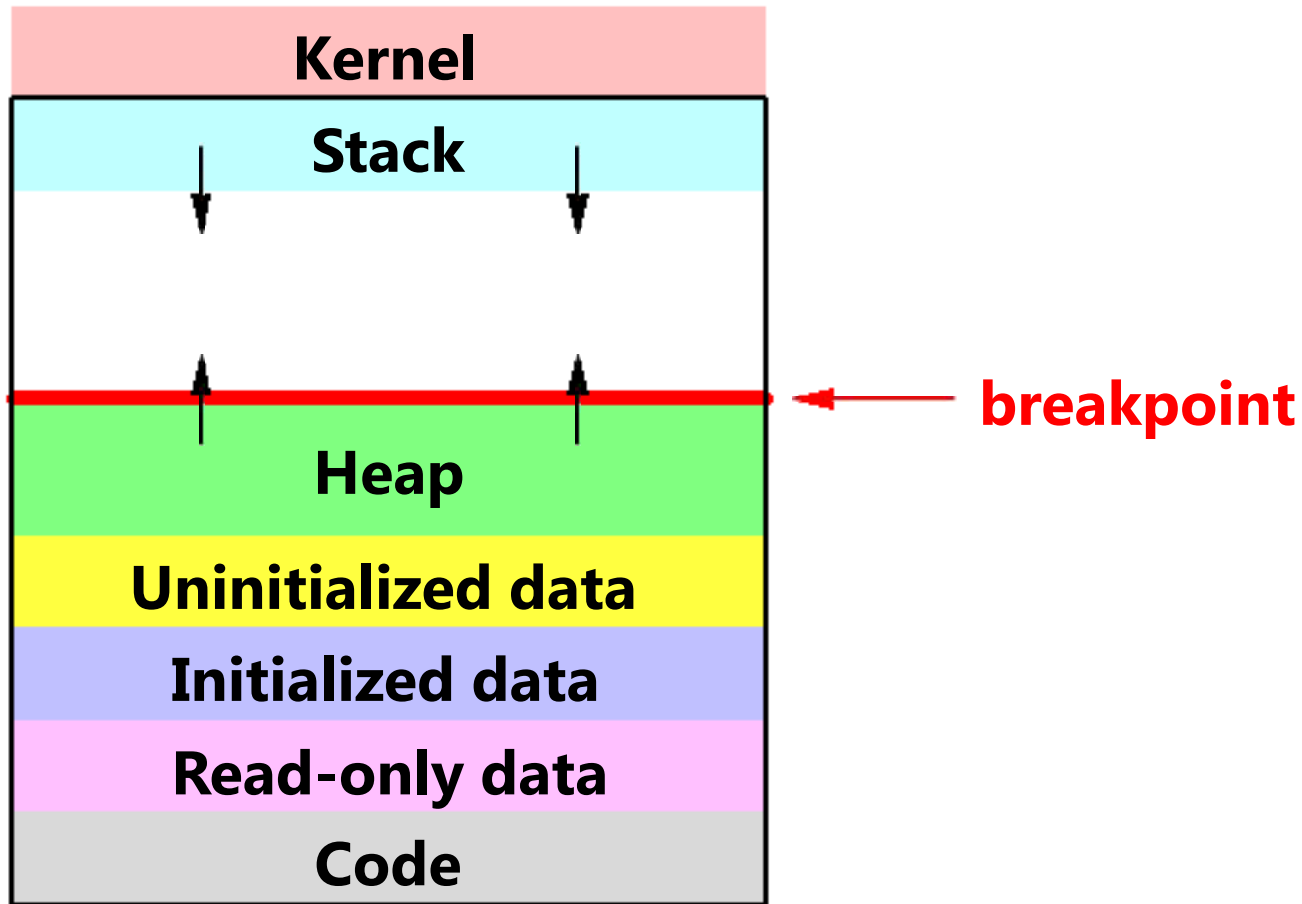
/* Creates a copy of user string US in kernel memory and returns it
   as a page that must be freed with palloc_free_page().
   Truncates the string at PGSIZE bytes in size. */
static char *
copy_in_string (const char *us)
{
    char *ks;
    size_t length;
    ks = palloc_get_page (PAL_ASSERT | PAL_ZERO);
    if (ks == NULL)
        thread_exit ();
    for (length = 0; length < PGSIZE; length++)
    {
        if (us >= (char *) PHYS_BASE || !get_user (ks + length, us++))
        {
            palloc_free_page (ks);
            thread_exit ();
        }
        if (ks[length] == '\\0')
            return ks;
    }
    ks[PGSIZE - 1] = '\\0';
    return ks;
}

```

Chapter 9: Virtual Memory

- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- **Memory-Mapped Files**
- **Allocating Kernel Memory**
- **Other Considerations**
- **Operating-System Examples**

Recall typical virtual address space

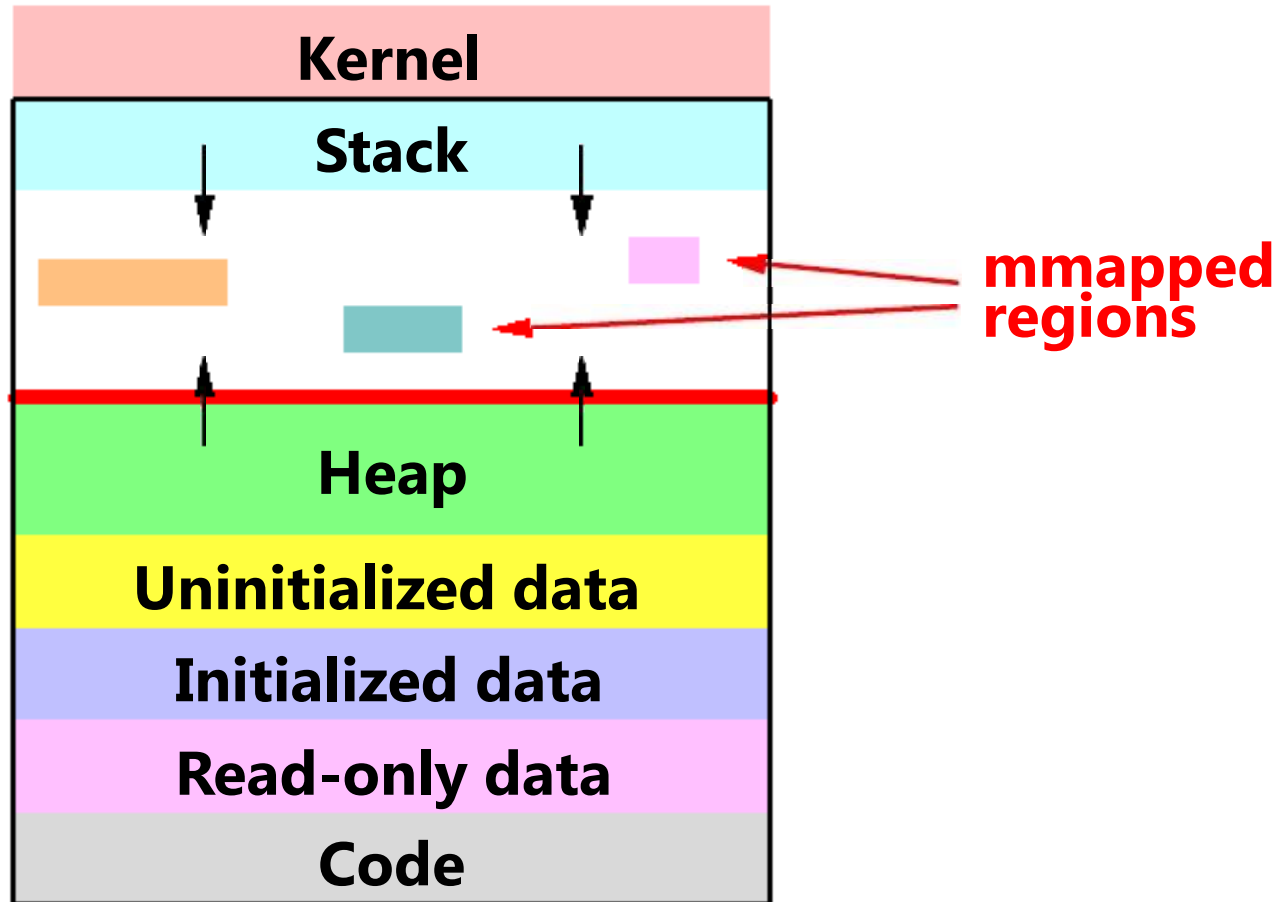


- Dynamically allocated memory goes in heap
- Top of heap called *breakpoint*
 - Addresses between breakpoint and stack all invalid

Early VM system calls

- **OS keeps “Breakpoint” – top of heap**
 - Memory regions between breakpoint & stack fault on access
- `char *brk (const char *addr);`
 - Set and return new value of breakpoint
- `char *sbrk (int incr);`
 - Increment value of the breakpoint & return old value
- **Can implement malloc in terms of sbrk**
 - But hard to “give back” physical memory to system

Memory mapped files



- Other memory objects between heap and stack

mmap system call

- `void *mmap (void *addr, size_t len, int prot, int flags, int fd, off_t offset)`
 - Map file specified by `fd` at virtual address `addr`
 - If `addr` is `NULL`, let kernel choose the address
- `prot` – **protection of region**
 - OR of `PROT_EXEC`, `PROT_READ`, `PROT_WRITE`, `PROT_NONE`
- `flags`
 - `MAP_ANON` – anonymous memory (`fd` should be `-1`)
 - `MAP_PRIVATE` – modifications are private
 - `MAP_SHARED` – modifications seen by everyone

More VM system calls

- `int msync(void *addr, size_t len, int flags);`
 - Flush changes of mmapped file to backing store
- `int munmap(void *addr, size_t len)`
 - Removes memory-mapped object
- `int mprotect(void *addr, size_t len, int prot)`
 - Changes protection on pages
- `int mincore(void *addr, size_t len, char *vec)`
 - Returns in vec which pages are present

Exposing page faults

```
struct sigaction {  
    union {                                /* signal handler */  
        void (*sa_handler)(int);  
        void (*sa_sigaction)(int, siginfo_t *, void *);  
    };  
    sigset_t sa_mask;    /* signal mask to apply */  
    int sa_flags;  
};
```

```
int sigaction (int sig, const struct sigaction *act,  
              struct sigaction *oact)
```

- **Can specify function to run on SIGSEGV
(Unix signal raised on invalid memory access)**

Example: OpenBSD/i386 siginfo

```
struct sigcontext {  
    int sc_gs; int sc_fs; int sc_es; int sc_ds;  
    int sc_edi; int sc_esi; int sc_ebp; int sc_ebx;  
    int sc_edx; int sc_ecx; int sc_eax;  
  
    int sc_eip; int sc_cs; /* instruction pointer */  
    int sc_eflags; /* condition codes, etc. */  
    int sc_esp; int sc_ss; /* stack pointer */  
  
    int sc_onstack; /* sigstack state to restore */  
    int sc_mask; /* signal mask to restore */  
  
    int sc_trapno;  
    int sc_err;  
};
```

VM tricks at user level

- **Combination of** mprotect/sigaction **very powerful**
 - Can use OS VM tricks in user-level programs [Appel]
 - E.g., fault, unprotect page, return from signal handler
- **Technique used in object-oriented databases**
 - Bring in objects on demand
 - Keep track of which objects may be dirty
 - Manage memory as a cache for much larger object DB
- **Other interesting applications**
 - Useful for some garbage collection algorithms
 - Snapshot processes (copy on write)

Dynamic memory allocation

- **Almost every useful program uses it**
 - Gives wonderful functionality benefits
 - ◁ Don't have to statically specify complex data structures
 - ◁ Can have data grow as a function of input size
 - ◁ Allows recursive procedures (stack growth)
 - But, can have a huge impact on performance
- **Today: how to implement it**
 - Lecture draws on [\[Wilson\]](#) (good survey from 1995)
- **Some interesting facts:**
 - Two or three line code change can have huge, non-obvious impact on how well an allocator works (examples to come)
 - Proven: impossible to construct an "always good" allocator
 - Surprising result: after 35 years, memory management still poorly understood

Why is it hard?

- Satisfy arbitrary set of allocation and free's.
- Easy without free: set a pointer to the beginning of some big chunk of memory ("heap") and increment on each allocation:

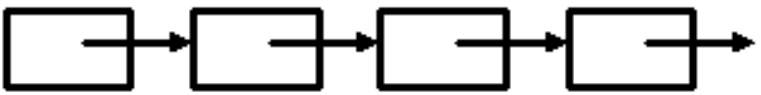



- Problem: free creates holes ("fragmentation") Result? Lots of free space but cannot satisfy request!



More abstractly

freelist

- **What an allocator must do:** 
 - Track which parts of memory in use, which parts are free
 - Ideal: no wasted space, no time overhead
- **What the allocator cannot do:**
 - Control order of the number and size of requested blocks
 - Change user ptrs \Rightarrow (bad) placement decisions permanent



- **The core fight: minimize fragmentation**
 - App frees blocks in any order, creating holes in "heap"
 - Holes too small? cannot satisfy future requests

What is fragmentation really?

- **Inability to use memory that is free**
- **Two factors required for fragmentation**
 - Different lifetimes—if adjacent objects die at different times, then fragmentation:



- If they die at the same time, then no fragmentation:




- Different sizes: If all requests the same size, then no fragmentation (that's why no external fragmentation w. paging):



Important decisions

- **Placement choice: where in free memory to put a requested block?**
 - Freedom: can select any memory in the heap
 - Ideal: put block where it won't cause fragmentation later (impossible in general: requires future knowledge)
- **Split free blocks to satisfy smaller requests?**
 - Fights internal fragmentation
 - Freedom: can chose any larger block to split
 - One way: chose block with smallest remainder (best fit)
- **Coalescing free blocks to yield larger blocks**



- Freedom: when to coalesce (deferring can be good) fights external fragmentation

Impossible to “solve” fragmentation

- **If you read allocation papers to find the best allocator**
 - All discussions revolve around tradeoffs
 - The reason? There cannot be a best allocator
- **Theoretical result:**
 - For any possible allocation algorithm, there exist streams of allocation and deallocation requests that defeat the allocator and force it into severe fragmentation.
- **How much fragmentation should we tolerate?**
 - Let M = bytes of live data, n_{\min} = smallest allocation, n_{\max} = largest – How much gross memory required?
 - Bad allocator: $M \cdot (n_{\max} / n_{\min})$
(only ever uses a memory location for a single size)
 - Good allocator: $\sim M \cdot \log(n_{\max} / n_{\min})$

Pathological examples

- **Given allocation of 7 20-byte chunks**



- What's a bad stream of frees and then allocates?
- **Given a 128-byte limit on malloced space**
 - What's a really bad combination of mallocs & frees?
- **Next: two allocators (best fit, first fit) that, in practice, work pretty well**
 - "pretty well" = ~20% fragmentation under many workloads

Pathological examples

- **Given allocation of 7 20-byte chunks**



- What's a bad stream of frees and then allocates?
 - **Free every other chunk, then alloc 21 bytes**
- **Given a 128-byte limit on malloced space**
 - What's a really bad combination of mallocs & frees?
- **Next: two allocators (best fit, first fit) that, in practice, work pretty well**
 - "pretty well" = ~20% fragmentation under many workloads

Pathological examples

- **Given allocation of 7 20-byte chunks**



- What's a bad stream of frees and then allocates?
 - **Free every other chunk, then alloc 21 bytes**
- **Given a 128-byte limit on malloced space**
 - What's a really bad combination of mallocs & frees?
 - **Malloc 128 1-byte chunks, free every other**
 - **Malloc 32 2-byte chunks, free every other (1- & 2-byte) chunk**
 - **Malloc 16 4-byte chunks, free every other chunk. . .**
- **Next: two allocators (best fit, first fit) that, in practice, work pretty well**
 - "pretty well" = ~20% fragmentation under many workloads

Best fit

- **Strategy: minimize fragmentation by allocating space from block that leaves smallest fragment**

- Data structure: heap is a list of free blocks, each has a header holding block size and pointers to next



- Code: Search freelist for block closest in size to the request. (Exact match is ideal)
- During free (usually) coalesce adjacent blocks

- **Problem: Sawdust**

- Remainder so small that over time left with "sawdust" everywhere
- Fortunately not a problem in practice

Best fit gone wrong

- **Simple bad case: allocate n, m ($n < m$) in alternating orders, free all the n s, then try to allocate an $n + 1$**
- **Example: start with 100 bytes of memory**

- alloc 19, 21, 19, 21, 19



- free 19, 19, 19:



- alloc 20? Fails! (wasted space = 57 bytes)

- **However, doesn't seem to happen in practice (though the way real programs behave suggest it easily could)**

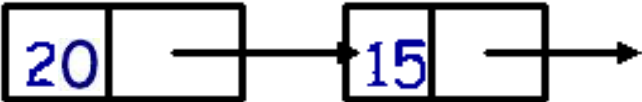
First fit

- **Strategy: pick the first block that fits**
 - Data structure: free list, sorted lifo, fifo, or by address
 - Code: scan list, take the first one
- **LIFO: put free object on front of list.**
 - Simple, but causes higher fragmentation
 - Potentially good for cache locality
- **Address sort: order free blocks by address**
 - Makes coalescing easy (just check if next block is free)
 - Also preserves empty/idle space (locality good when paging)
- **FIFO: put free object at end of list**
 - Gives similar fragmentation as address sort, but unclear why

Subtle pathology: LIFO FF

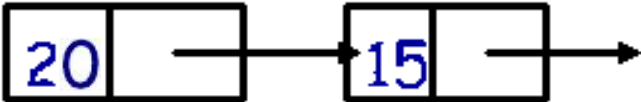
- **Storage management example of subtle impact of simple decisions**
- **LIFO first fit seems good:**
 - Put object on front of list (cheap), hope same size used again (cheap + good locality)
- **But, has big problems for simple allocation patterns:**
 - E.g., repeatedly intermix short-lived $2n$ -byte allocations, with long-lived $(n + 1)$ -byte allocations
 - Each time large object freed, a small chunk will be quickly taken, leaving useless fragment. Pathological fragmentation

First fit: Nuances

- **First fit sorted by address order, in practice:**
 - Blocks at front preferentially split, ones at back only split when no larger one found before them
 - Result? Seems to roughly sort free list by size
 - So? Makes first fit operationally similar to best fit: a first fit of a sorted list = best fit!
- **Problem: sawdust at beginning of the list**
 - Sorting of list forces a large requests to skip over many small blocks. Need to use a scalable heap organization
- **Suppose memory has free blocks:** 
 - If allocation ops are 10 then 20, best fit wins
 - When is FF better than best fit?

First fit: Nuances

- **First fit sorted by address order, in practice:**
 - Blocks at front preferentially split, ones at back only split when no larger one found before them
 - Result? Seems to roughly sort free list by size
 - So? Makes first fit operationally similar to best fit: a first fit of a sorted list = best fit!
- **Problem: sawdust at beginning of the list**
 - Sorting of list forces a large requests to skip over many small blocks. Need to use a scalable heap organization

- **Suppose memory has free blocks:** 
 - If allocation ops are 10 then 20, best fit wins
 - When is FF better than best fit?
 - **Suppose allocation ops are 8, 12, then 12 \Rightarrow first fit wins**

First/best fit: weird parallels

- **Both seem to perform roughly equivalently**
- **In fact the placement decisions of both are roughly identical under both randomized and real workloads!**
 - No one knows why
 - Pretty strange since they seem pretty different
- **Possible explanations:**
 - First fit like best fit because over time its free list becomes sorted by size: the beginning of the free list accumulates small objects and so fits tend to be close to best
 - Both have implicit “open space heuristic” try not to cut into large open spaces: large blocks at end only used when have to be (e.g., first fit: skips over all smaller blocks)

Some worse ideas

- **Worst-fit:**

- Strategy: fight against sawdust by splitting blocks to maximize leftover size
- In real life seems to ensure that no large blocks around

- **Next fit:**

- Strategy: use first fit, but remember where we found the last thing and start searching from there
- Seems like a good idea, but tends to break down entire list


- **Buddy systems:**

- Round up allocations to power of 2 to make management faster
- Result? Heavy internal fragmentation

Known patterns of real programs

- So far we've treated programs as black boxes.
- Most real programs exhibit 1 or 2 (or all 3) of the following patterns of alloc/dealloc:

- *Ramps*: accumulate data monotonically over time




- *Peaks*: allocate many objects, use briefly, then free all



- *Plateaus*: allocate many objects, use for a long time




Pattern 1: ramps



- **In a practical sense: ramp = no free!**
 - Implication for fragmentation?
 - What happens if you evaluate allocator with ramp programs only?


Pattern 2: peaks




- **Peaks: allocate many objects, use briefly, then free all**
 - Fragmentation a real danger
 - What happens if peak allocated from contiguous memory?
 - Interleave peak & ramp? Interleave two different peaks?

Exploiting peaks

- **Peak phases: alloc a lot, then free everything**
 - So have new allocation interface: alloc as before, but only support free of everything
 - Called "arena allocation", "obstack" (object stack), or alloca/procedure call (by compiler people)
- **Arena = a linked list of large chunks of memory**
 - Advantages: alloc is a pointer increment, free is "free"
No wasted space for tags or list pointers



Pattern 3: Plateaus




trace of perl running a string processing script

- **Plateaus: allocate many objects, use for a long time**
 - What happens if overlap with peak or different plateau?

Fighting fragmentation

- **Segregation = reduced fragmentation:**


- Allocated at same time ~ freed at same time
- Different type ~ freed at different time



- **Implementation observations:**

- Programs allocate small number of different sizes
- Fragmentation at peak use more important than at low
- Most allocations small (< 10 words)
- Work done with allocated memory increases with size
- Implications?

Simple, fast segregated free lists




- **Array of free lists for small sizes, tree for larger**
 - Place blocks of same size on same page
 - Have count of allocated blocks: if goes to zero, can return page
- **Pro: segregate sizes, no size tag, fast small alloc**
- **Con: worst case waste: 1 page per size even w/o free, after pessimal free waste 1 page per object**

Typical space overheads

- **Free list bookkeeping + alignment determine minimum allocatable size:**


- Store size of block
- Pointers to next and previous freelist element



- Machine enforced overhead: alignment. Allocator doesn't know type. Must align memory to conservative boundary
- Minimum allocation unit? Space overhead when allocated?

Getting more space from OS

- **On Unix, can use** `sbrk`
 - E.g., to activate a new zero-filled page:



`sbrk(4096)`

```
/* add nbytes of valid virtual address space */  
void *get_free_space(unsigned nbytes) {  
    void *p;  
    if(!(p = sbrk(nbytes)))  
        error("virtual memory exhausted");  
    return p;  
}
```

- **For large allocations, sbrk a bad idea**
 - May want to give memory back to OS
 - Can't with `sbrk` unless big chunk last thing allocated
 - So allocate large chunk using `mmap`'s `MAP_ANON`

Faults + resumption = power

- **Resuming after fault lets us emulate many things**
 - “every problem can be solved with layer of indirection”
- **Example: sub-page protection**
- **To protect sub-page region in paging system:**




- Set entire page to weakest permission; record in PT



- Any access that violates perm will cause an access fault
- Fault handler checks if page special, and if so, if access allowed. Continue or raise error, as appropriate

More fault resumption examples


- **Emulate accessed bits:**
 - Set page permissions to "invalid".
 - On any access will get a fault: Mark as accessed
- **Avoid save/restore of FP registers**
 - Make first FP operation fault to detect usage
- **Emulate non-existent instructions:**
 - Give inst an illegal opcode; OS fault handler detects and emulates fake instruction
- **Run OS on top of another OS!**
 - Slam OS into normal process
 - When does something "privileged," real OS gets woken up with a fault.
 - If op allowed, do it, otherwise kill.
 - IBM's VM/370. VMware (sort of)



Not just for kernels

- **User-level code can resume after faults, too**
- **mprotect – protects memory**
- **sigaction – catches signal after page fault**
 - Return from signal handler restarts faulting instruction
- **Many applications detailed by Appel & Li**
- **Example: concurrent snapshotting of process**
 - Mark all of process's memory read-only w. mprotect
 - One thread starts writing all of memory to disk
 - Other thread keeps executing
 - On fault – write that page to disk, make writable, resume

Distributed shared memory



- **Virtual memory allows us to go to memory or disk**
 - But, can use the same idea to go anywhere! Even to another computer. Page across network rather than to disk. Faster, and allows network of workstations (NOW)

Persistent stores


- **Idea: Objects that persist across program invocations**
 - E.g., object-oriented database; useful for CAD/CAM type apps
- **Achieve by memory-mapping a file**
- **But only write changes to file at end if commit**
 - Use dirty bits to detect which pages must be written out
 - Or emulate dirty bits with *mprotect/sigaction* (using write faults)
- **On 32-bit machine, store can be larger than memory**
 - But single run of program won't access > 4GB of objects
 - Keep mapping betw. 32-bit mem ptrs and 64-bit disk offsets
 - Use faults to bring in pages from disk as necessary
 - After reading page, translate pointers—known as *swizzling*

Garbage collection

- **In safe languages, run time knows about all pointers**
 - So can move an object if you change all the pointers
- **What memory locations might a program access?**
 - Any objects whose pointers are currently in registers
 - Recursively, any pointers in objects it might access
 - Anything else is *unreachable*, or *garbage*; memory can be re-used
- **Example: stop-and-copy garbage collection**
 - Memory full? Temporarily pause program, allocate new heap
 - Copy all objects pointed to by registers into new heap
 - ◁ Mark old copied objects as copied, record new location
 - Start scanning through new heap. For each pointer:
 - ◁ Copied already? Adjust pointer to new location
 - ◁ Not copied? Then copy it and adjust pointer
 - Free old heap—program will never access it—and continue

Concurrent garbage collection

- **Idea: Stop & copy, but without the stop**
 - *Mutator* thread runs program, *collector* concurrently does GC
- **When collector invoked:**
 - Protect from space & unscanned to space from mutator
 - Copy objects in registers into *to space*, resume mutator
 - All pointers in scanned *to space* point to *to space*
 - If mutator accesses unscanned area, fault, scan page, resume



Heap overflow detection

- **Many GCed languages need fast allocation**
 - E.g., in lisp, constantly allocating cons cells
 - Allocation can be as often as every 50 instructions
- **Fast allocation is just to bump a pointer**

```
char *next_free;  
char *heap_limit;  
  
void *alloc (unsigned size) {  
    if (next_free + size > heap_limit)          /* 1 */  
        invoke_garbage_collector ();           /* 2 */  
    char *ret = next_free;  
    next_free += size;  
    return ret;  
}
```

- **But would be even faster to eliminate lines 1 & 2!**


Heap overflow detection 2

- **Mark page at end of heap inaccessible**
 - `mprotect (heap_limit, PAGE_SIZE, PROT_NONE);`
- **Program will allocate memory beyond end of heap**
- **Program will use memory and fault**
 - Note: Depends on specifics of language
 - But many languages will touch allocated memory immediately
- **Invoke garbage collector**
 - Must now put just allocated object into new heap
- **Note: requires more than just resumption**
 - Faulting instruction must be resumed
 - But must resume with different target virtual address
 - Doable on most architectures since GC updates registers

Reference counting

- **Seemingly simpler GC scheme:**

- Each object has "ref count" of pointers to it
- Increment when pointer set to it
- Decrement when pointer killed (C++ destructors handy for such "smart pointers")



```
void foo(bar c) {  
    bar a, b;  
    a = c; ..... c->refcnt++;  
    b = a; ..... a->refcnt++;  
    a = 0; ..... c->refcnt--;  
    return; ..... b->refcnt--;  
}
```


- ref count == 0? Free object

- **Works well for hierarchical data structures**

- E.g., pages of physical memory

Reference counting pros/cons

- **Circular data structures always have ref count > 0**
 - No external pointers means **lost memory**



- **Can do manually w/o PL support, but error-prone**
- **Potentially more efficient than real GC**
 - No need to halt program to run collector
 - Avoids weird unpredictable latencies
- **Potentially less efficient than real GC**
 - With real GC, copying a pointer is cheap
 - With reference counting, must write ref count each time

Summary

- Read Ch. 1-8
- Processes and Threads (Ch. 4)
- Process Scheduling (Ch. 5)
- Synchronization (Ch. 6)
- Deadlock (Ch. 7)
- Memory Management (Ch. 8)
- Virtual Memory (Ch. 9)
- Project #2 – System Calls and User-Level Processes