

Lecture 23: Protection and Security

Instructor: Mitch Neilsen

Office: N219D

Quote of the Day

“I needed a password eight characters long, so I picked Snow White and the Seven Dwarves.”

-- Nick Helm, comedian

Chapter 14: Protection

- Goals of Protection
- Principles of Protection
- Domain of Protection
- Access Matrix
- Implementation of Access Matrix
- Access Control
- Revocation of Access Rights
- Capability-Based Systems
- Language-Based Protection

Objectives

- Discuss the goals and principles of protection in a modern computer system
- Explain how protection domains combined with an access matrix are used to specify the resources a process may access
- Examine capability and language-based protection systems

Core Components

- Authentication
 - Verify that we really know who we are talking to
- Authorization
 - Check that user X is allowed to do Y
- Access Enforcement
 - Ensure that authorization decision is respected
- Hard to do because every system has holes
 - Social vs. technical enforcement

Authentication

- Passwords
 - Weakest form, and most common
 - Subject to dictionary attacks
 - Passwords should not be stored in clear text, instead, use a **one-way hash function** to generate shadow passwords.
- Badge or Keycard
 - Should not be (easily) forgeable
 - Problem: how to invalidate?
- Biometrics
 - Problem: ensure trusted path to device

Now that we've authenticated, Authorization

- Operating system consists of a collection of objects, hardware or software
- Each object has a unique name and can be accessed through a well-defined set of operations.
- Protection problem - ensure that each object is accessed correctly and only by those processes/users that are authorized to do so.

Principles of Protection

Guiding principle – **principle of least privilege**

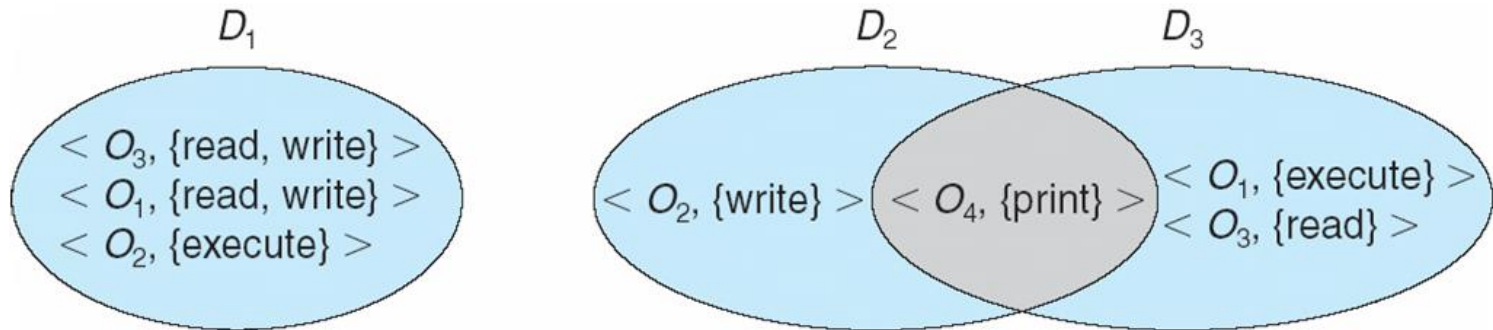
- Programs, users and systems should be given just enough privileges to perform their tasks
- Example: SELinux - Security-Enhanced Linux

Domain Structure

Access-right = $\langle \textit{object-name}, \textit{rights-set} \rangle$

where *rights-set* is a subset of all valid operations that can be performed on the object.

Domain = set of access-rights



Domain Implementation (UNIX)

System consists of 2 domains:

- User
- Supervisor

UNIX

```
ls -l /usr/bin/passwd  
-rwsr-xr-x  root  root  ..  passwd
```

- Domain = user-id
- Domain switch accomplished via file system.


Each file has associated with it a domain bit (setuid bit (**s**)).

When file is executed and setuid = on, then user-id is set to owner of the file being executed. When execution completes user-id is reset.

Domain Implementation (MULTICS)

Let D_i and D_j be any two domain rings.

If $j < i \Rightarrow D_i \subseteq D_j$



Access Matrix

- View protection as a matrix (*access matrix*)
- Rows represent domains
- Columns represent objects
- **Access(i, j)** is the set of operations that a process executing in Domain_i can invoke on Object_j

Access Matrix

object domain	F_1	F_2	F_3	printer
D_1	read		read	
D_2				print
D_3		read	execute	
D_4	read write		read write	

Use of Access Matrix

If a process in Domain D_i tries to do “op” on object O_j , then “op” must be in the access matrix.

Can be expanded to dynamic protection.

- Operations to add, delete access rights.
- Special access rights:

owner of O_i

copy op from O_i to O_j

control – D_i can modify D_j access rights

transfer – switch from domain D_i to D_j

Use of Access Matrix (Cont.)

Access matrix design separates mechanism from policy.

- **Policy - what should be done**

User dictates policy.

Who can access what object and in what mode.

- **Mechanism - how it should be done**

Operating system provides access-matrix + rules.

If ensures that the matrix is only manipulated by authorized agents and that rules (policy) are strictly enforced.

Implementation of Access Matrix

Each column = **Access Control List** for one object --
Defines who can perform what operation.

Domain 1 = read

Domain 4 = read, write

object \ domain	F_1	F_2	F_3	printer
D_1	read		read	
D_2				print
D_3		read	execute	
D_4	read write		read write	

Each Row = **Capability List** (like keys or tickets)
For each domain, what operations are allowed on
what objects. E.g., domain D_3 :

Object F_2 – read

Object F_3 – execute

object \ domain	F_1	F_2	F_3	printer
D_1	read		read	
D_2				print
D_3		read	execute	
D_4	read write		read write	

Access Control List

- Store list of <user (domain), set of privileges> for each object
 - Example: files. For each file store who is allowed to access it (and how: read, write, execute)
 - Most contemporary file systems support it.
- Groups can be used to compress the information:
 - Traditional Unix permissions `rw-r-xr-x`; `chmod 755 myDir`.
- Where in the filesystem should you store ACLs/permissions?

Capabilities

- Store (capability) list of <object, set of privileges> for each user (domain)
 - Typically used in systems that must be very secure
 - Default is empty capability list
- Capabilities also often function as names
 - Can access something if you know the name
- Must make names unforgeable, or must have system monitor who holds what capabilities each user (domain) has (e.g., by storing them in a protected area)

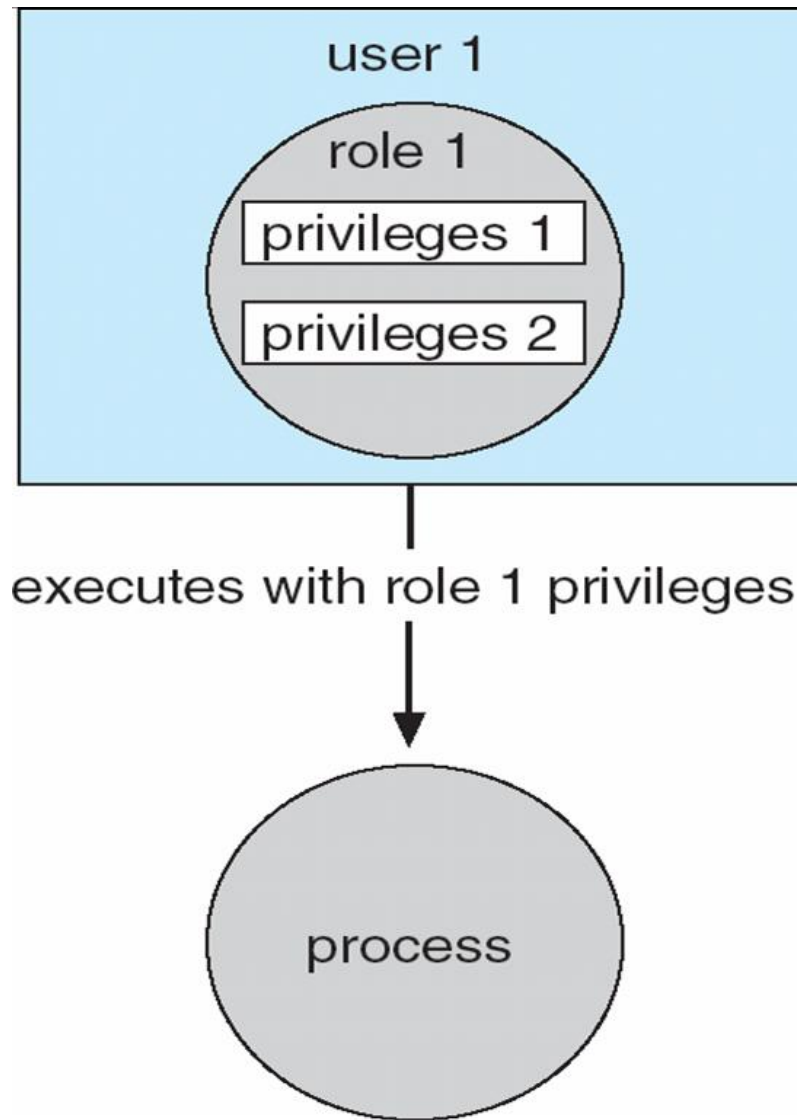
Access Control

Protection can be applied to non-file resources, too

Solaris 10 provides **role-based access control** to implement least privilege

- Privilege is right to execute system call or use an option within a system call
- Can be assigned to processes
- Users assigned roles granting access to privileges and programs

Role-based Access Control in Solaris 10



Revocation of Access Rights

Access List – Delete access rights from access list. (e.g., chmod)

- Simple
- Immediate

Capability List – Scheme required to locate capability in the system before the capability can be revoked.

- Reacquisition
- Back-pointers
- Indirection
- Keys

Capability-Based Systems

Hydra

- Fixed set of access rights known to and interpreted by the system.
- Interpretation of user-defined rights performed solely by user's program; system provides access protection for use of these rights.

Cambridge CAP System

- Data capability - provides standard read, write, execute of individual storage segments associated with object.
- Software capability - interpretation left to the subsystem, through its protected procedures.

Language-Based Protection

- Specification of protection in a programming language allows the high-level description of policies for the allocation and use of resources.
- Language implementation can provide software for protection enforcement when automatic hardware-supported checking is unavailable.
- Interpret protection specifications to generate calls on whatever protection system is provided by the hardware and the operating system.
- Based on type-safe languages (Java, C#, etc.)
- Do not allow direct memory access.
- Include access modifiers (private/public, etc.)
- Verify code before they execute it with respect to these safety properties.
- Build security systems on top of type-safe language run-times which associate code with sets of privileges.

Example: Protection in Java

- Protection is handled by the Java Virtual Machine (JVM)
- A class is assigned a protection domain when it is loaded by the JVM.
- The protection domain indicates what operations the class can (and cannot) perform.
- If a library method is invoked that performs a privileged operation, the stack is inspected to ensure the operation can be performed by the library.

Stack Inspection

protection domain:	untrusted applet	URL loader	networking
socket permission:	none	*.lucent.com:80, connect	any
class:	gui: ... get(url); open(addr); ...	get(URL u): ... doPrivileged { open('proxy.lucent.com:80'); } <request u from proxy> ... }	open(Addr a): ... checkPermission (a, connect); connect (a); ... }

The diagram illustrates the stack inspection process. A blue arrow points from the `open(addr);` call in the `gui` class to the `open('proxy.lucent.com:80');` call inside the `doPrivileged` block of the `URL loader` class. Another blue arrow points from the `open('proxy.lucent.com:80');` call to the `open(Addr a):` method in the `networking` class. A red lightning bolt icon is placed next to the `checkPermission(a, connect);` call in the `networking` class, indicating a security check or exception.

Note that a process must not be allowed to modify annotations on its own stack frame or do other manipulations to bypass stack inspection.

Chapter 15: Security

- The Security Problem
- Program Threats
- System and Network Threats
- Cryptography as a Security Tool
- User Authentication
- Implementing Security Defenses
- Firewalls to Protect Systems and Networks
- Computer-Security Classifications
- An Example: Windows XP

Objectives

- To discuss security threats and attacks
- To explain the fundamentals of encryption, authentication, and hashing
- To examine the uses of cryptography in computing
- To describe the various countermeasures to security attacks

The Security Problem

- Security must consider external environment of the system, and protect the system resources.
- Intruders (crackers) attempt to breach security.
- A **threat** is potential security violation.
- An **attack** is attempt to breach security.
- Attacks can be accidental or malicious.
- Easier to protect against accidental than malicious misuse.

Security Violations




Categories

- Breach of confidentiality
- Breach of integrity
- Breach of availability
- Theft of service
- Denial of service

Methods

- Masquerading (breach authentication)
- Replay attack
- Message modification
- Man-in-the-middle attack
- Session hijacking

Standard Security Attacks



Security Measure Levels

- Security must occur at four levels to be effective:
 1. Physical
 2. Human
 3. Operating System
 4. Network
- Security is as weak as the weakest chain

Program Threats

Stack and Buffer Overflow

- Exploits a bug in a program (overflow either the stack or memory buffers)

Trojan Horse

- Code segment that misuses its environment
- Exploits mechanisms for allowing programs written by users to be executed by other users
- **Spyware, pop-up browser windows, covert channels**

Trap Door

- Specific user identifier or password that circumvents normal security procedures; could be included in a compiler

Logic Bomb

- Program that initiates a security incident under certain circumstances

Buffer Overflow Attacks

November, 1988

- First Internet Worm spread over then-new Internet
- Many university machines compromised
- No malicious effect

Today

- Buffer overflow is still the initial entry for over 50% of network-based attacks

String Library Code

- Implementation of Unix function `gets()`

No way to specify limit on number of characters to read

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getc();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getc();
    }
    *p = '\0';
    return dest;
}
```

Similar problems exist with other Unix functions:

`strcpy`: Copies string of arbitrary length

`scanf`, `fscanf`, `sscanf`, when given `%s` conversion specification

Vulnerable Buffer Code

```
/* Echo Line */  
void echo()  
{  
    char buf[4];    /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
int main()  
{  
    printf("Type a string:");  
    echo();  
    return 0;  
}
```

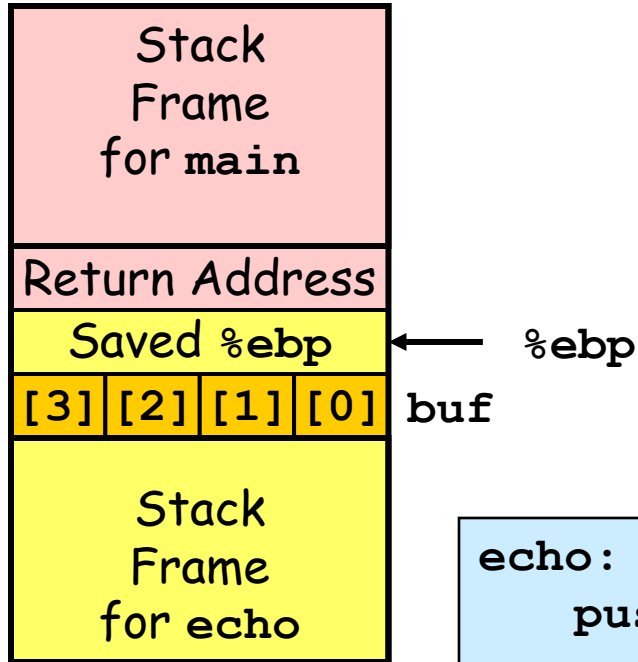
Buffer Overflow Executions

```
unix>./bufdemo  
Type a string:123  
123
```

```
unix>./bufdemo  
Type a string:12345  
12345  
→ note valid output, bad input
```

```
unix>./bufdemo  
Type a string:12345678  
Segmentation Fault
```

Buffer Overflow Stack (IA32)

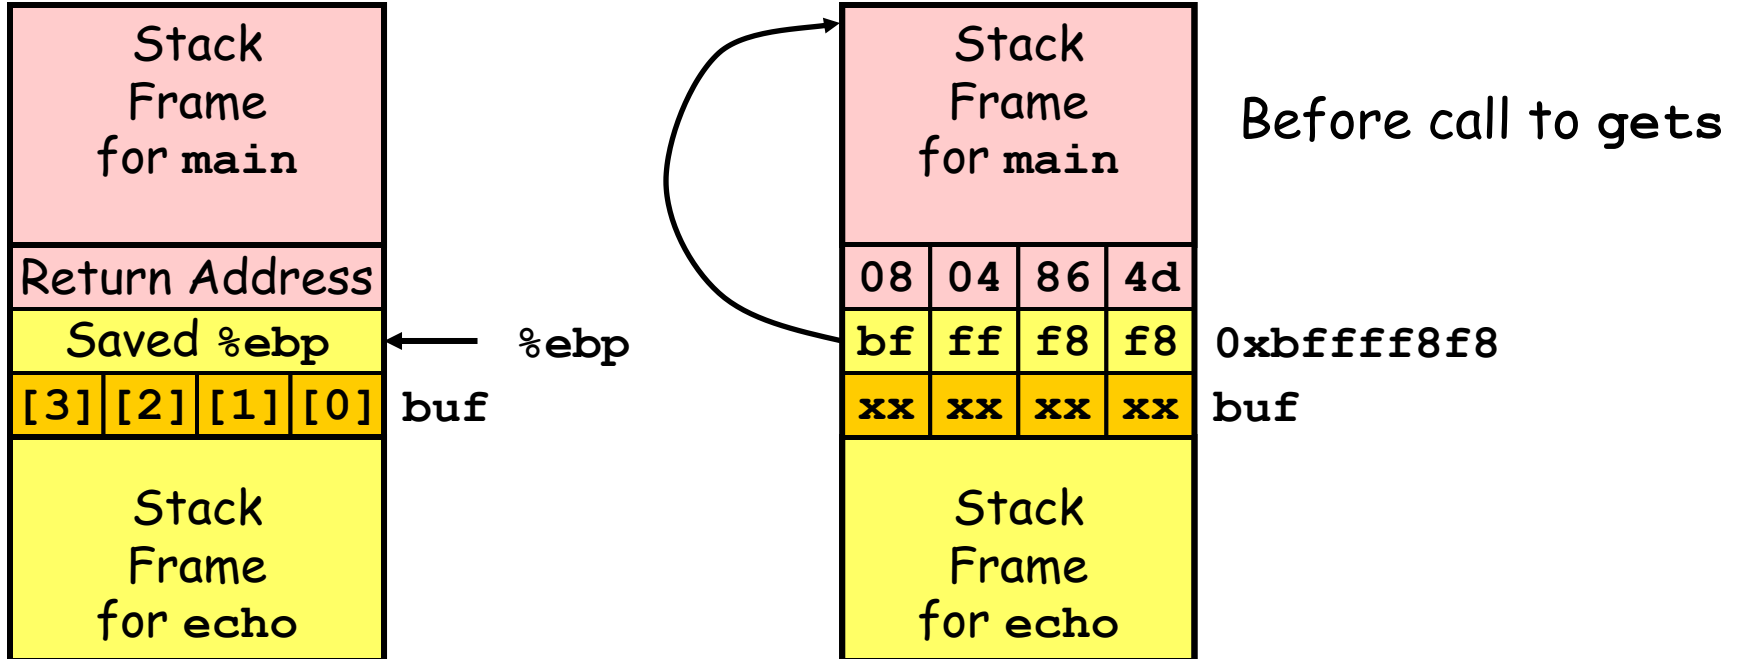


```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

```
echo:
    pushl %ebp                # Save %ebp on stack
    movl %esp,%ebp
    subl $20,%esp            # Allocate stack space
    pushl %ebx                # Save %ebx
    addl $-12,%esp            # Allocate stack space
    leal -4(%ebp),%ebx        # Compute buf as %ebp-4
    pushl %ebx                # Push buf on stack
    call gets                 # Call gets
    . . .
```

Buffer Overflow Stack Example

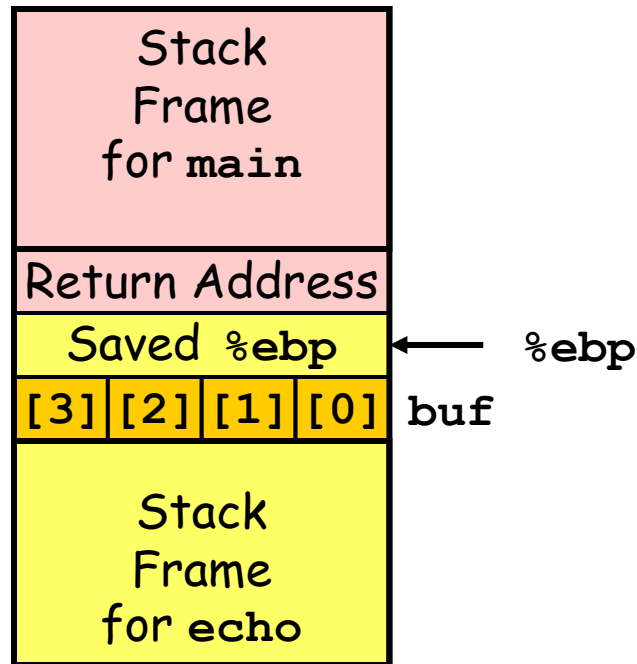
```
unix> gdb bufdemo
(gdb) break echo
Breakpoint 1 at 0x8048583
(gdb) run
Breakpoint 1, 0x8048583 in echo ()
(gdb) print /x *(unsigned *)$ebp
$1 = 0xbffff8f8
(gdb) print /x *((unsigned *)$ebp + 1)
$3 = 0x804864d
```




```
8048648: call 804857c <echo>
804864d: mov 0xffffffe8(%ebp),%ebx # Return Point
```

Buffer Overflow Example #1

Before Call to gets

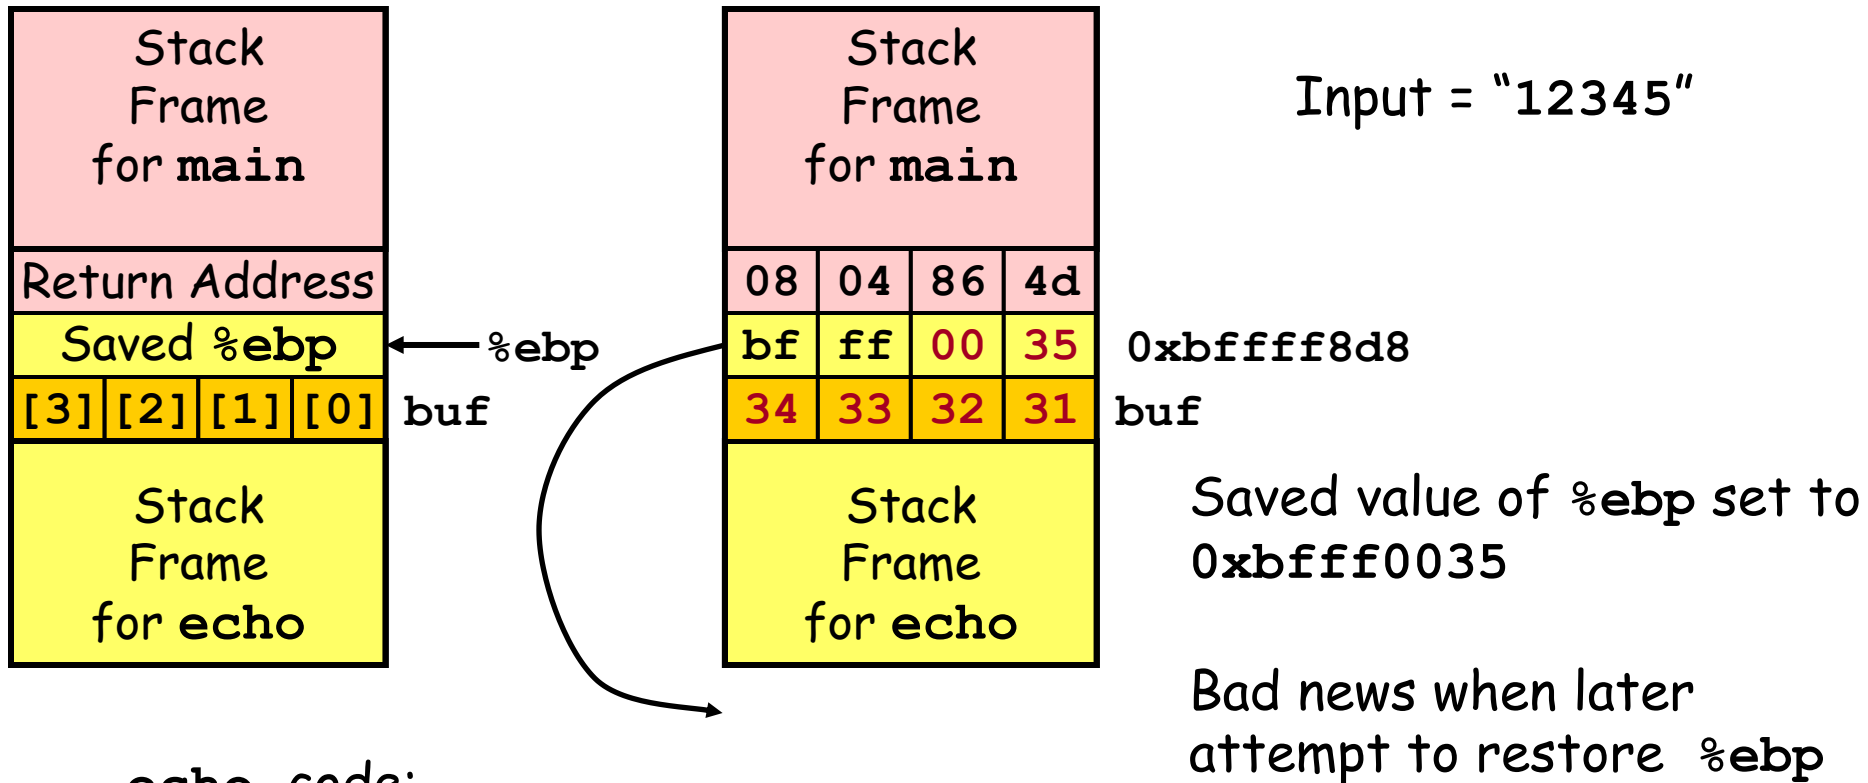


Input = "123"



No Problem

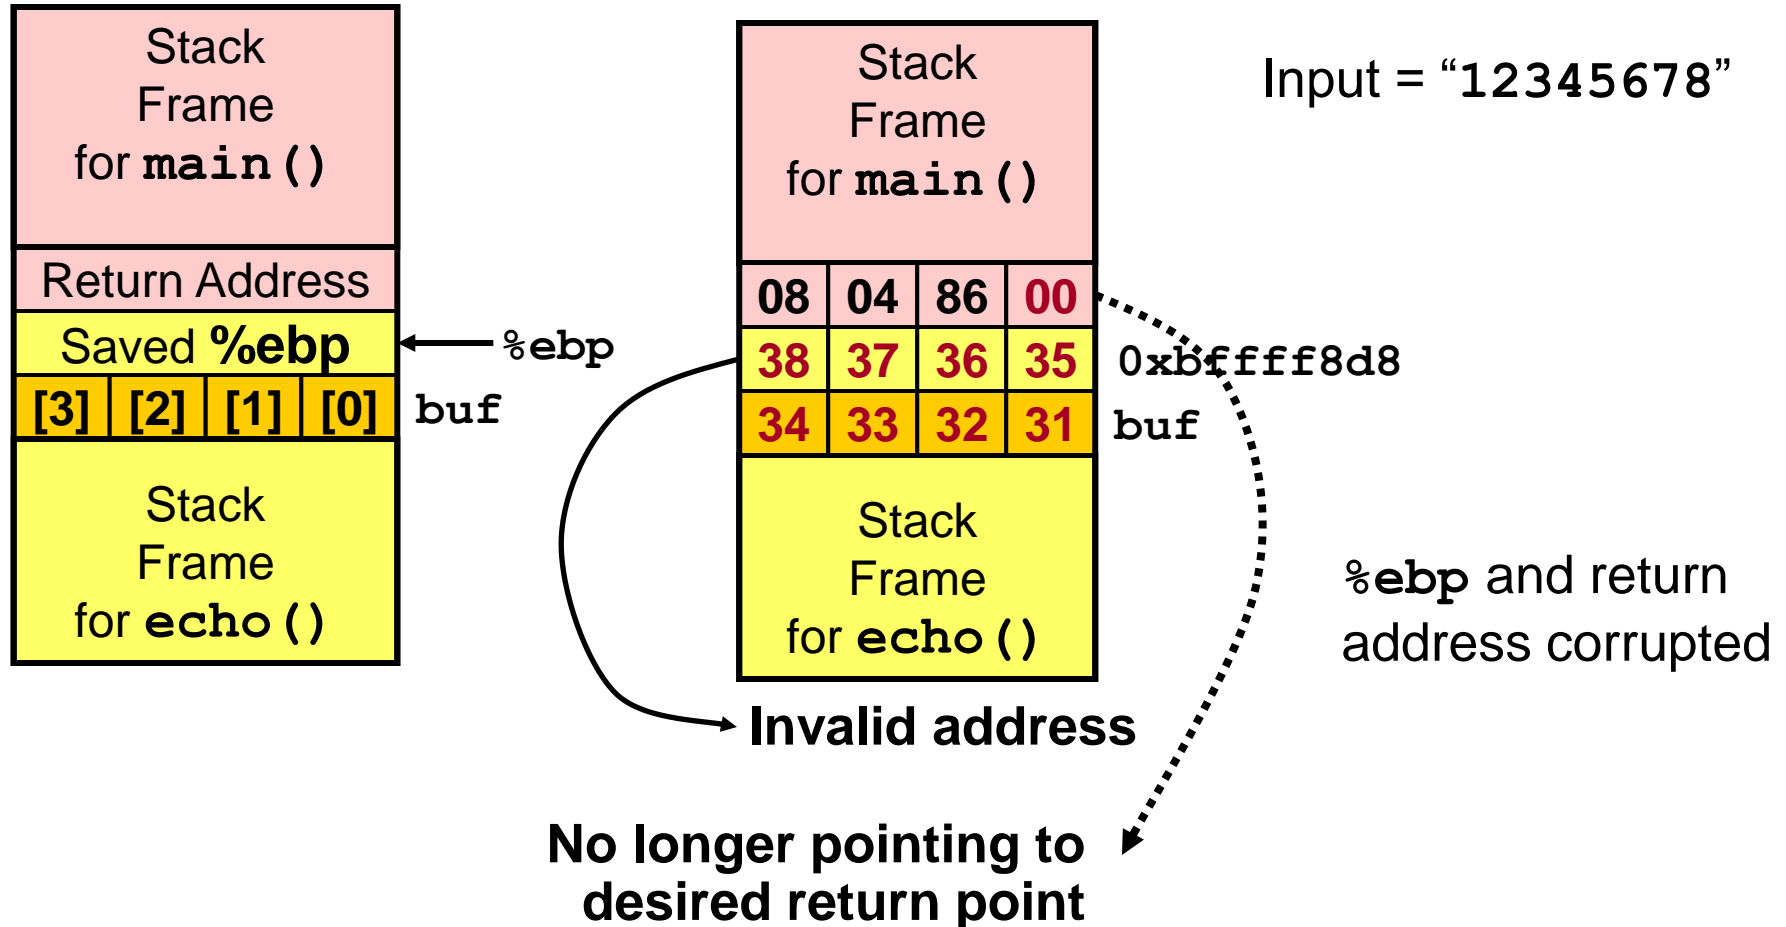
Buffer Overflow Stack Example #2



echo code:

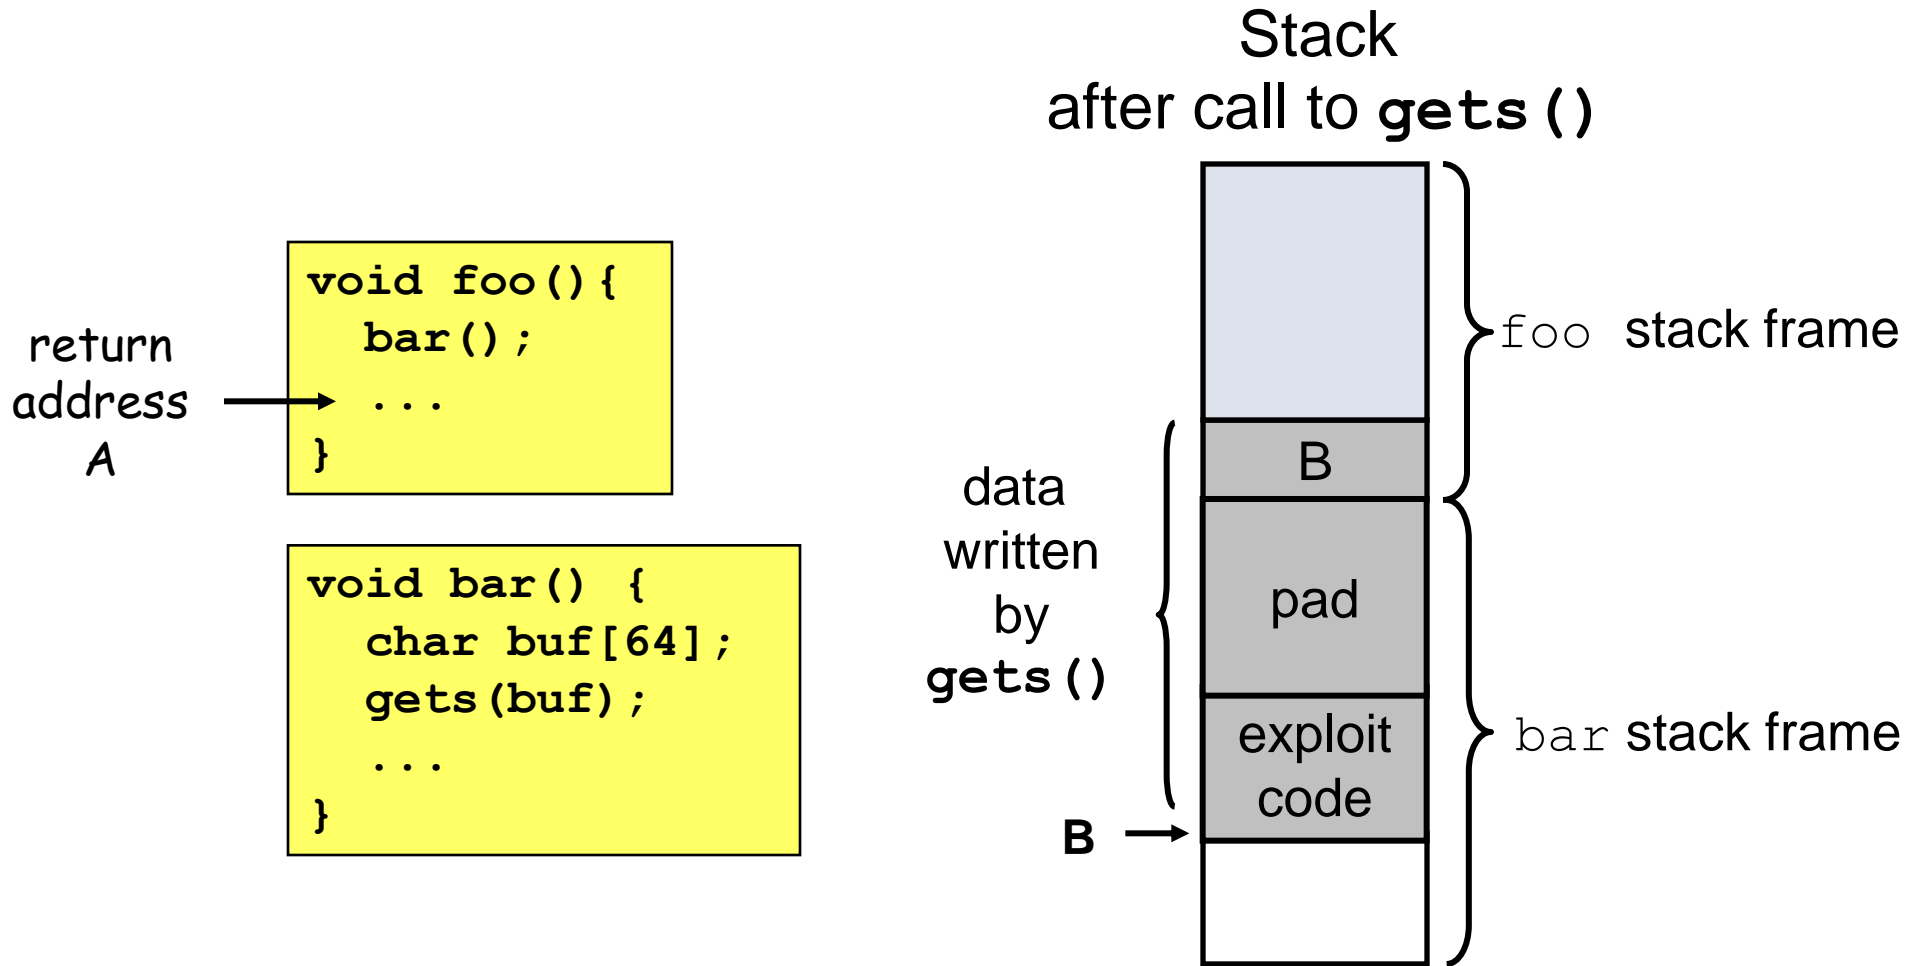
```
8048592: push    %ebx
8048593: call    80483e4 <_init+0x50> # gets
8048598: mov     0xffffffffe8(%ebp),%ebx
804859b: mov     %ebp,%esp
804859d: pop    %ebp      # %ebp gets set to invalid value
804859e: ret
```


Buffer Overflow Stack Example #3



```
8048648: call 804857c <echo>
804864d: mov 0xffffffe8(%ebp),%ebx # Return Point
```

Malicious Use of Buffer Overflow



- Input string contains byte representation of executable code
- Overwrite return address with address of buffer
- When **bar()** executes **ret**, will jump to exploit code

Exploits Based on Buffer Overflows

Buffer overflow bugs allow remote machines to execute arbitrary code on victim machines.

Internet worm

- Early versions of the finger server (fingerd) used `gets()` to read the argument sent by the client:

`finger neilsen@cis.ksu.edu`

- Worm attacked fingerd server by sending phony argument:

`finger "exploit-code padding new-return-address"`

exploit code: executed a root shell on the victim machine with a direct TCP connection to the attacker.

C Call Stack

- When a function call is made, the return address is pushed onto the stack.
- Often the values of parameters passed to the function are put onto the stack (call-by-value).
- Usually the function saves the stack frame pointer (old %ebp) on the stack.
- Local variables are placed on the stack.

Stack Direction

- On Linux (x86) the stack grows from high addresses to low.
- Pushing something onto the stack moves the Top Of the Stack (%esp) towards address 0.

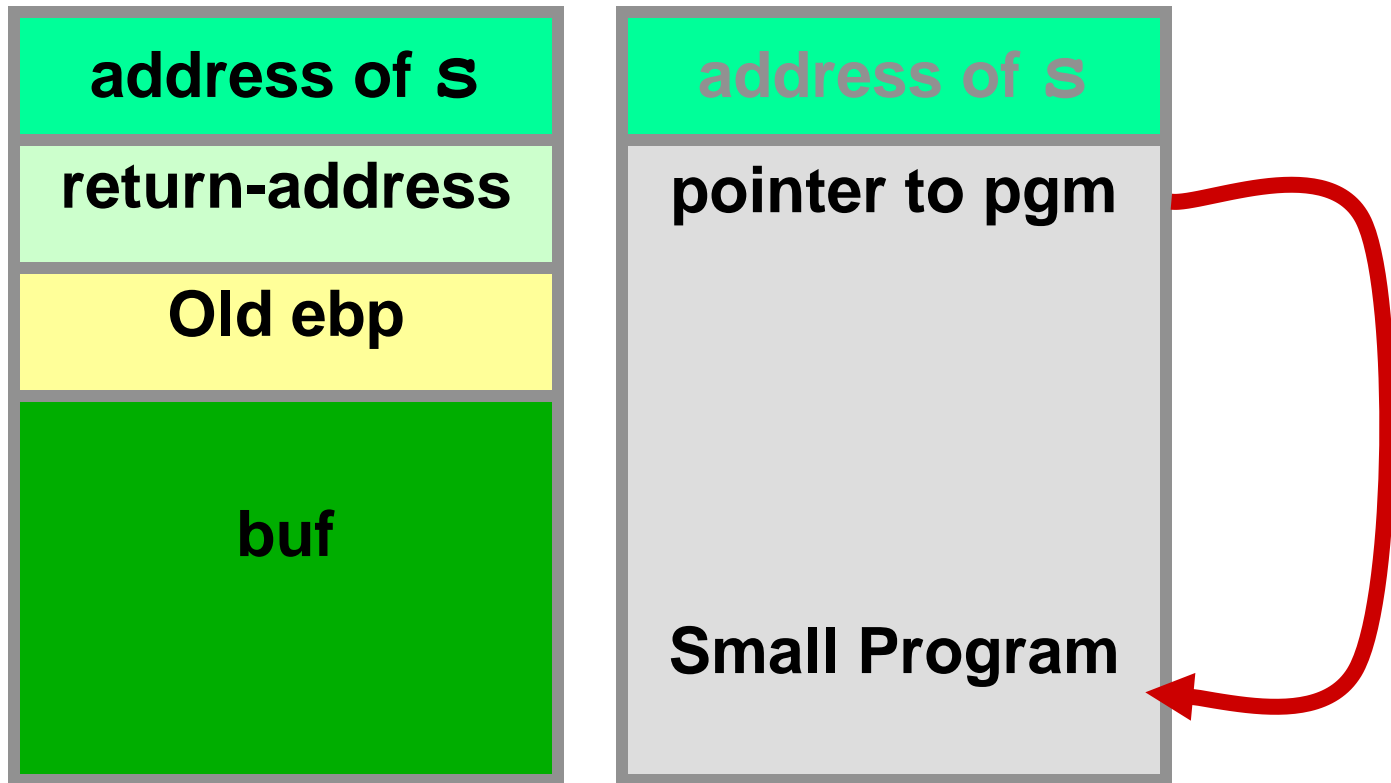
“Smashing the Stack”*

- The general idea is to overflow a buffer so that it overwrites the return address.
- When the function is done it will jump to whatever address is on the stack.
- We put some code in the buffer and set the return address to point to it!

*taken from the title of an article in Phrack 49-7

Before and After

```
void foo(char *s) {  
    char buf[100];  
    strcpy(buf, s);  
    ...  
}
```



Issues

- How do we know what value the pointer should have (the new “return address”).
- It’s the address of the buffer, but how do we know what address this is?
- How do we build the “small program” and put it in a string?

Guessing Addresses

- Typically you need the source code so you can *estimate* the address of both the buffer and the return-address.
- An estimate is often good enough!

Building the small program

- Typically, the small program stuffed in to the buffer does an **exec()** .
- Sometimes it changes the password file or other files...

exec ()

- In Unix, the way to run a new program is with an **exec ()** system call.
 - There is actually a *family* of **exec ()** system calls...
 - This doesn't create a new process, it changes the current process to run a new program.
 - To create a new process you need another system call (e.g., **fork ()**).

exec () example

```
#include <stdio.h>
```

```
void execl_s(void) {  
    execl("/bin/ls", "ls", NULL);  
    printf("Line not printed if execl is  
           successful.\n");  
}
```

Generating a String

- You can take code like the previous slide, and generate machine language.
- Copy down the individual byte values and build a string.
- To do a simple `exec()` requires less than 100 bytes.

Some important issues

- The small program should be position-independent – able to run at any memory location.
- Statically link the libraries to see the code generated for the `exec()` system call; e.g., `gcc execExample.c`, to see how the `exec()` system call is made. To statically link the libraries, use:

`gcc -static execExample.c`

- It can't be too large, or we can't fit the program and the new return-address on the stack!

A Sample Program/String

Does an exec() of /bin/ls:

```
unsigned char cde[] =  
"\xeb\x1f\x5e\x89\x76\x08\x31\xc0"  
"\x88\x46\x07\x89\x46\x0c\xb0\x0b"  
"\x89\xf3\x8d\x4e\x08\x8d\x56\x0c"  
"\xcd\x80\x31\xdb\x89\xd8\x40xcd"  
"\x80\xe8\xdc\xff\xff\xff/bin/ls";
```

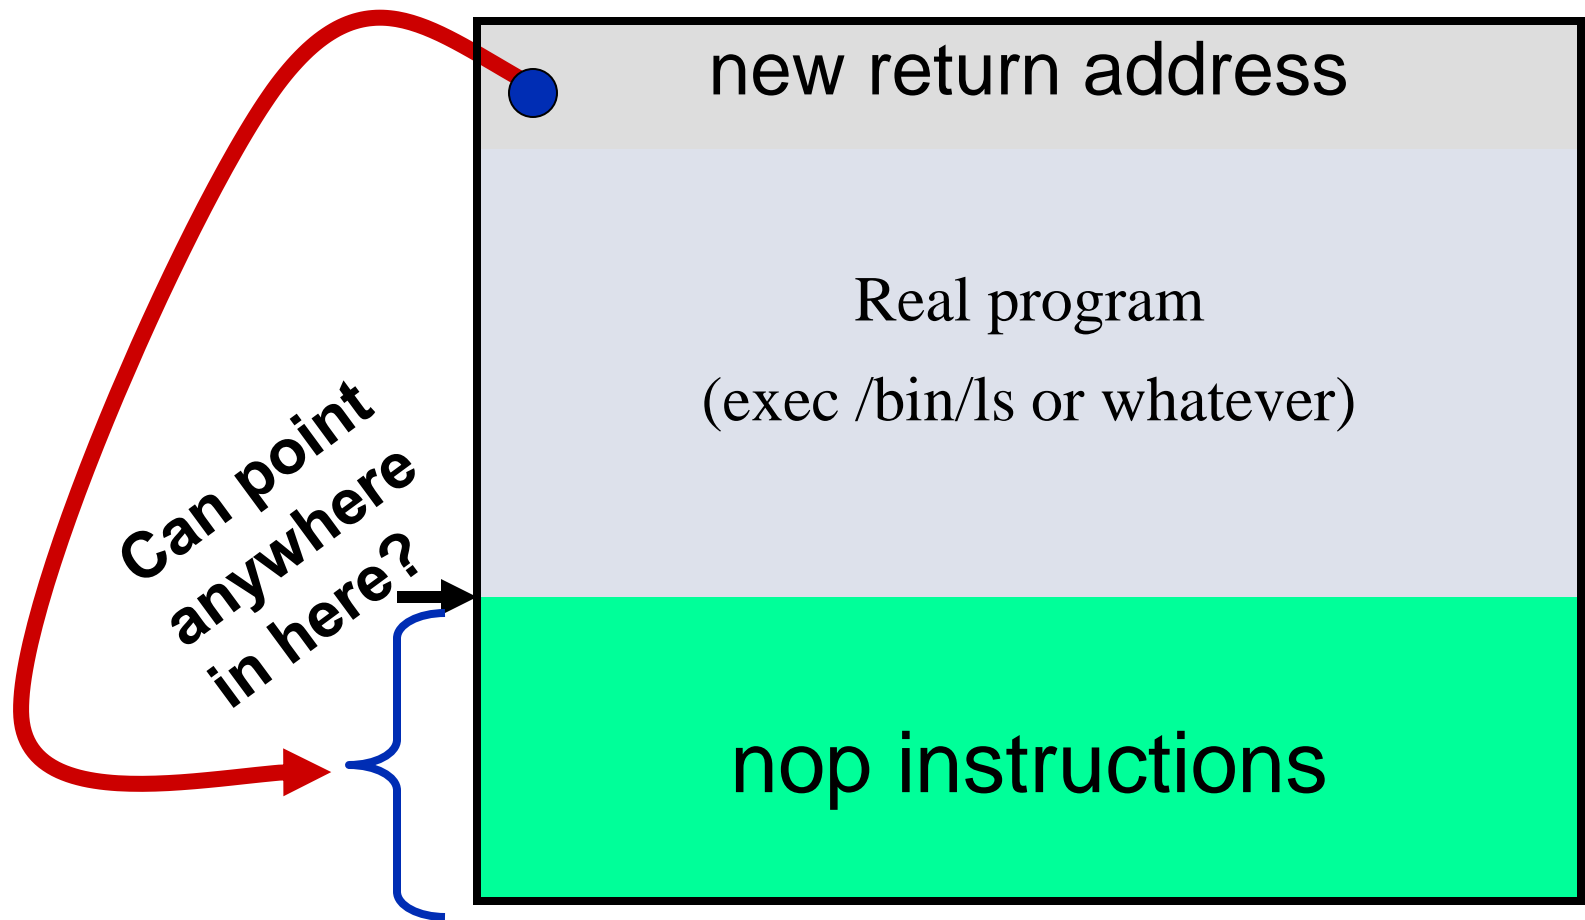
Attacking a real program

- Recall that the idea is to feed a server a string that is too big for a buffer.
- This string overflows the buffer and overwrites the return address on the stack.
- Assuming we put our small program in the string, we need to know it's address.

NOPs

- Most CPUs have a *No-Operation* instruction – it does nothing but advance the instruction pointer.
- Usually we can put a bunch of these ahead of our program (in the string).
- As long as the new return-address points to a NOP we are OK.

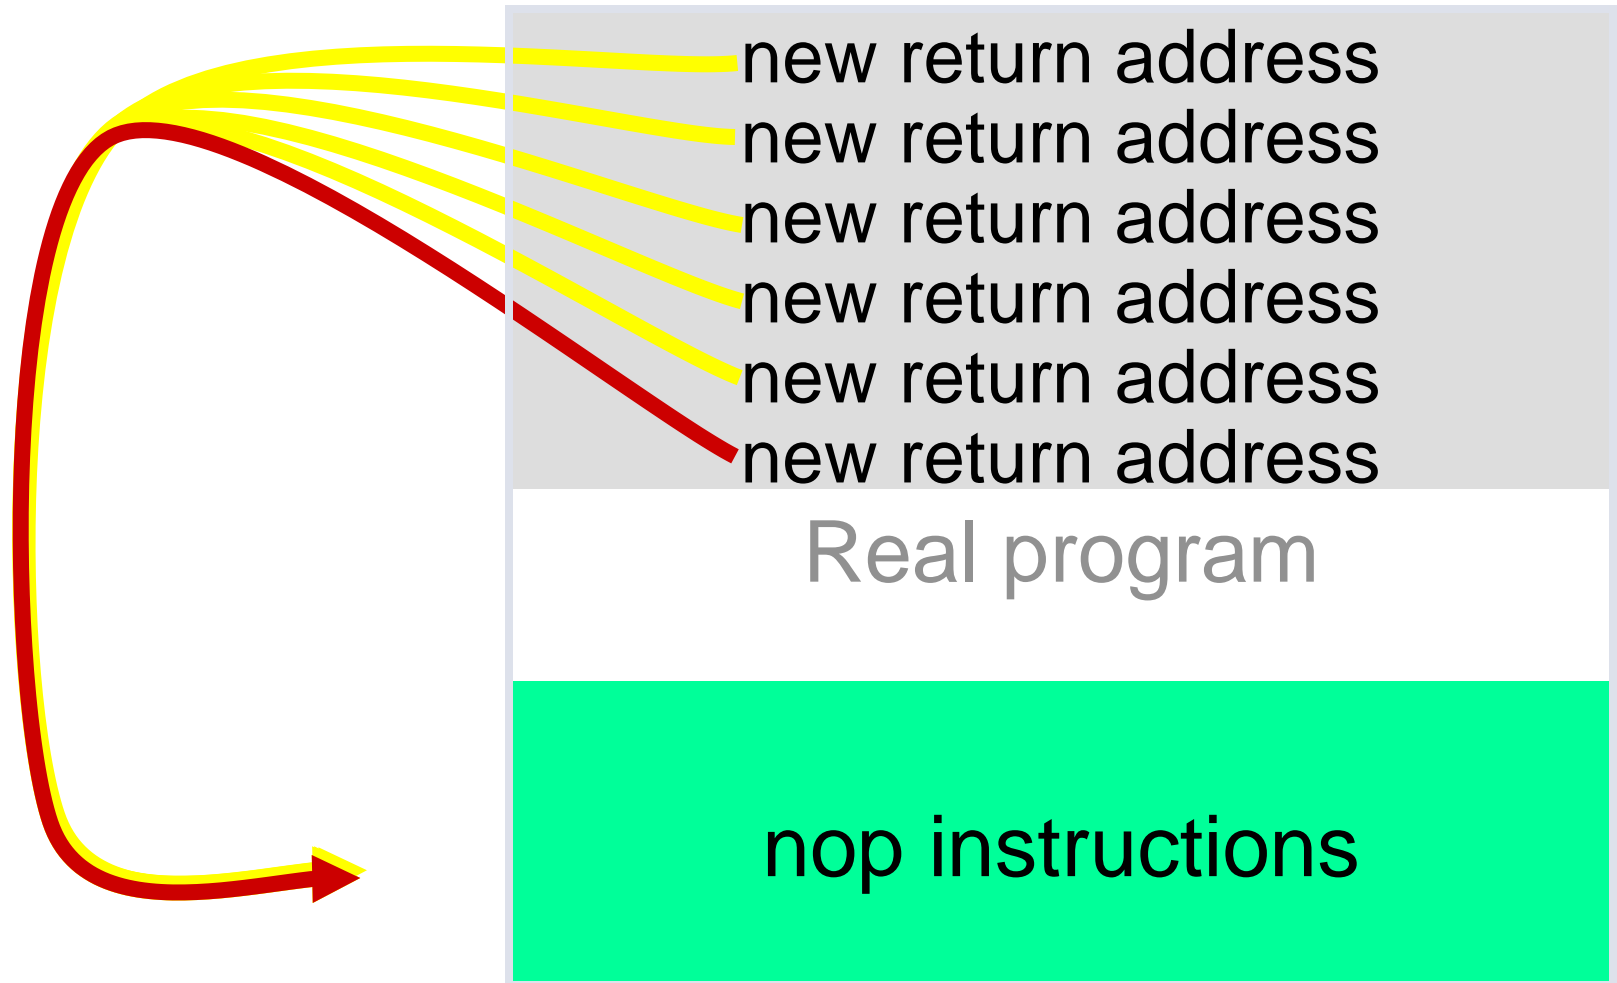
Using NOPs



Estimating the stack size

- We can also guess at the location of the return address relative to the overflowed buffer.
- Put in a bunch of new return addresses!

Estimating the Location



Summary

- Read Ch 13-15
- Quiz #2 – Wed., Nov. 20, open-book, open-notes
- Project #3
 - Application
 - Kernel modifications – add new system call