

Using Synchronization

Dr. Daniel Andresen

CIS520

Avoiding Races #1

1. Identify *critical sections*, code sequences that:

- rely on an invariant condition being true;
- temporarily violate the invariant;
- transform the data structure from one legal state to another;
- or make a sequence of actions that assume the data structure will not “change underneath them”.

2. *Never sleep or yield in a critical section.*



Voluntarily relinquishing control may allow another thread to run and “trip over your mess” or modify the structure while the operation is in progress.

Critical Sections in the Color Stack

```
InitColorStack() {  
    push(blue);  
    push(purple);  
}
```

```
PushColor() {  
    if (s[top] == purple) {  
        ASSERT(s[top-1] == blue);  
        push(blue);  
    } else {  
        ASSERT(s[top] == blue);  
        ASSERT(s[top-1] == purple);  
        push(purple);  
    }  
}
```



Avoiding Races #2

Is caution with *yield* and *sleep* sufficient to prevent races?

No!

Concurrency races may also result from:

- involuntary context switches (timeslicing)
e.g., caused by the Nachos thread scheduler with **-rs** flag
- external events that asynchronously change the flow of control
interrupts (inside the kernel) or signals/APCs (outside the kernel)
- physical concurrency (on a multiprocessor)

How to ensure atomicity of critical sections in these cases?

Synchronization primitives!

Synchronization 101

Synchronization constrains the set of possible interleavings:

- Threads can't prevent the scheduler from switching them out, but they can “agree” to stay out of each other's way.
 - voluntary blocking or spin-waiting on entrance to critical sections
 - notify blocked or spinning peers on exit from the critical section
- If we're “inside the kernel” (e.g., the Nachos kernel), we can *temporarily* disable interrupts.
 - no races from interrupt handlers or involuntary context switches
 - a blunt instrument to use as a last resort
 - Disabling interrupts is not an accepted synchronization mechanism!*
 - insufficient on a multiprocessor

Mutual Exclusion

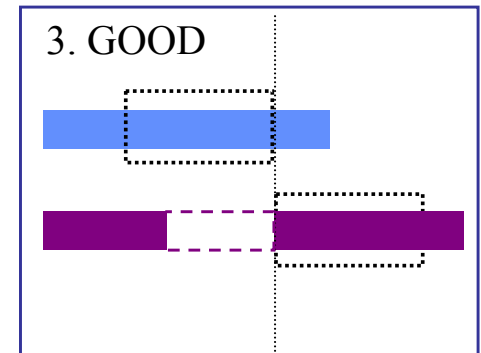
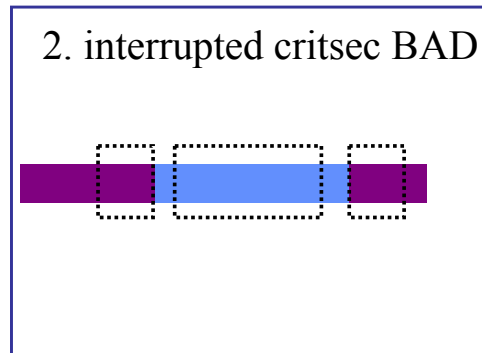
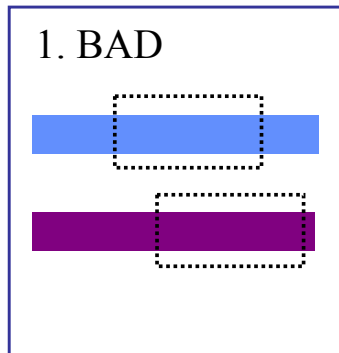
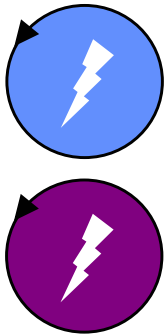
Race conditions can be avoided by ensuring *mutual exclusion* in critical sections.

- Critical sections are code sequences that contribute to races.

Every race (possible incorrect interleaving) involves two or more threads executing related critical sections concurrently.

- To avoid races, we must *serialize* related critical sections.

Never allow more than one thread in a critical section at a time.



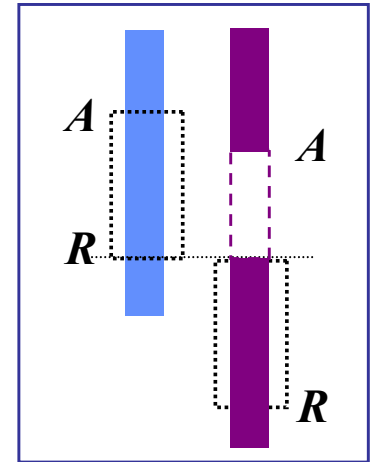
Locks

Locks can be used to ensure mutual exclusion in conflicting critical sections.

- A lock is an object, a data item in memory.

Methods: *Lock::Acquire* and *Lock::Release*.

- Threads pair calls to *Acquire* and *Release*.
- *Acquire* before entering a critical section.
- *Release* after leaving a critical section.
- Between *Acquire/Release*, the lock is *held*.
- *Acquire* does not return until any previous holder releases.
- Waiting locks can spin (a *spinlock*) or block (a *mutex*).

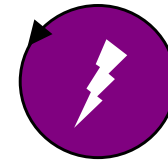
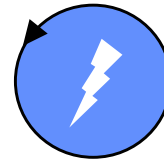


Example: Per-Thread Counts and Total

```
/* shared by all threads */
```

```
int counters[N];
```

```
int total;
```



```
/*
```

```
 * Increment a counter by a specified value, and keep a running sum.
```

```
 * This is called repeatedly by each of N threads.
```

```
 * tid is an integer thread identifier for the current thread.
```

```
 * value is just some arbitrary number.
```

```
 */
```

```
void
```

```
TouchCount(int tid, int value)
```

```
{
```

```
    counters[tid] += value;
```

```
    total += value;
```

```
}
```


Using Locks: An Example

```
int counters[N];
```

```
int total;
```

```
Lock *lock;
```

```
/*
```

```
 * Increment a counter by a specified value, and keep a running sum.
```

```
*/
```

```
void
```

```
TouchCount(int tid, int value)
```

```
{
```

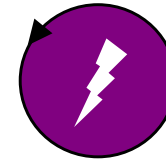
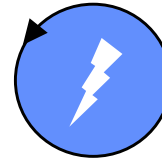
```
    lock->Acquire();
```

```
    counters[tid] += value;
```

```
    total += value;
```

```
    lock->Release();
```

```
}
```



```
/* critical section code is atomic...*/
```

```
/* ...as long as the lock is held */
```

Relativity of Critical Sections

1. If a thread is executing a critical section, never permit another thread to enter the same critical section.

Two executions of the same critical section on the same data are *always* “mutually conflicting” (assuming it modifies the data).

2. If a thread is executing a critical section, never permit another thread to enter a *related* critical section.

Two different critical sections may be mutually conflicting.

E.g., if they access the same data, and at least one is a writer.

E.g., *List::Add* and *List::Remove* on the same list.

3. Two threads may safely enter *unrelated* critical sections.

If they access different data or are reader-only.

Semaphores

Semaphores handle all of your synchronization needs with one elegant but confusing abstraction.

- controls allocation of a resource with multiple instances
- a non-negative integer with special operations and properties

initialize to arbitrary value with *Init* operation

“souped up” increment (*Up* or *V*) and decrement (*Down* or *P*)

- atomic sleep/wakeup behavior implicit in *P* and *V*

P does an atomic *sleep*, if the semaphore value is zero.

P means “probe”; it cannot decrement until the semaphore is positive.

V does an atomic *wakeup*.

$\text{num}(P) \leq \text{num}(V) + \text{init}$

Semaphores as Mutexes

```
semaphore->Init(1);  
  
void Lock::Acquire()  
{  
    semaphore->P();  
}  
  
void Lock::Release()  
{  
    semaphore->V();  
}
```

Semaphores must be initialized with a value representing the number of free resources: mutexes are a single-use resource.

$P()$ to acquire a resource; blocks if no resource is available.

$V()$ to release a resource; wakes up one waiter, if any.

P and V are *atomic*.

Mutexes are often called *binary semaphores*.

However, “real” mutexes have additional constraints on their use.

Ping-Pong with Semaphores



blue->Init(0);
purple->Init(1);

```
void  
PingPong() {  
    while(not done) {  
        blue->P();  
        Compute();  
        purple->V();  
    }  
}
```



```
void  
PingPong() {  
    while(not done) {  
        purple->P();  
        Compute();  
        blue->V();  
    }  
}
```



Ping-Pong with One Semaphore?



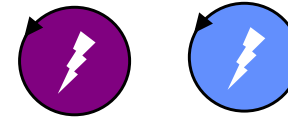
```
sem->Init(0);
```

```
blue: { sem->P(); PingPong(); }
```

```
purple: { PingPong(); }
```

```
void
```

```
PingPong() {  
    while(not done) {  
        Compute();  
        sem->V();  
        sem->P();  
    }  
}
```



Nachos semaphores have Mesa-like semantics:

They do not guarantee that a waiting thread wakes up “in time” to consume the count added by a *V()*.

- semaphores are not “fair”
- no count is “reserved” for a waking thread
- uses “passive” vs. “active” implementation

Another Example With Dual Semaphores

blue->Init(0);
purple->Init(0);

```
void Blue() {  
    while(not done) {  
        Compute();  
        purple->V();  
        blue->P();  
    }  
}
```



```
void Purple() {  
    while(not done) {  
        Compute();  
        blue->V();  
        purple->P();  
    }  
}
```



Basic Producer/Consumer

```
empty->Init(1);  
full->Init(0);  
int buf;
```

```
void Produce(int m) {  
    empty->P();  
    buf = m;  
    full->V();  
}
```

```
int Consume() {  
    int m;  
    full->P();  
    m = buf;  
    empty->V();  
    return(m);  
}
```

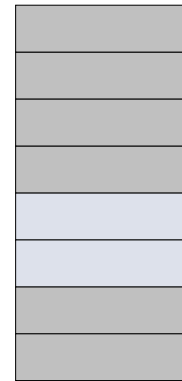
This use of a semaphore pair is called a *split binary semaphore*: the sum of the values is always one.

A Bounded Resource

```
int AllocateEntry() {
    int i;
    while (!FindFreeItem(&i))
        block and wait for a free slot
    slot[i] = 1;      /* grab free slot */
    return(i);
}

void ReleaseEntry(int i) {
    slot[i] = 0;
    wakeup waiter, if any
}

boolean FindFreeItem(int* index) {
    for (i = 0; i < TableSize; i++)
        if (slot[i] == 0) return it;
    return (FALSE);
}
```



A Bounded Resource with a Counting Semaphore

```
semaphore->Init(N);
```

A semaphore for an N-way resource is called a *counting semaphore*.

```
int AllocateEntry() {  
    int i;  
    semaphore->Down();  
    ASSERT(FindFreeItem(&i));  
    slot[i] = 1;  
    return(i);  
}
```

A caller that gets past a *Down* is guaranteed that a resource instance is reserved for it.

```
void ReleaseEntry(int i) {  
    slot[i] = 0;  
    semaphore->Up();  
}
```

Problems?

Note: the current value of the semaphore is the number of resource instances free to allocate.

But semaphores do not allow a thread to read this value directly. Why not?

Spin-Yield: Just Say No

```
void  
Thread::Await() {  
    awaiting = TRUE;  
    while(awaiting)  
        Yield();  
}
```

```
void  
Thread::Awake() {  
    if (awaiting)  
        awaiting = FALSE;  
}
```