


Project 2: User Programs and System Calls

P2: Project 2 – System Calls




Support Code

Students Create

Public Tests

Project 1 and Project 2



Using the File system

- ❑ Interfacing with the file system
- ❑ No need to modify the file system
- ❑ Certain limitations
 - No internal synchronization
 - File size fixed
 - File data allocated as a single extent
 - No subdirectories
 - File names limited to 14 chars
 - System crash might corrupt the file system
- ❑ Files to take a look at: 'filesys.h' & 'file.h'

Some commands

- ❑ In **userprog/build**, create a new simulated disk
 - `pintos-mkdisk fs.dsk --filesystem-size=2`
- ❑ Format the disk
 - `pintos -f -q`
 - This will only work after your disk is created and kernel is built!
- ❑ Copy the program "echo" onto the disk
 - `pintos -p ../../examples/echo -a echo -- -q`
- ❑ Run the program
 - `pintos -q run 'echo x'`
- ❑ All in a single command:

```
pintos --filesystem-size=2 -p ../../examples/echo -a echo -- -f  
-q run 'echo x'
```
- `pintos ... -v -- ..` - to from terminal without X11 server
- ❑ **\$ make check** or **make grade** – builds the disk automatically
 - You can just copy & paste the commands that make check does!

Various directories

- ❑ Few user programs:
 - src/examples
- ❑ Relevant files:
 - userprog/
- ❑ Other files:
 - threads/

Project 2 Requirements

- ❑ Process Termination Messages
- ❑ Argument Passing
- ❑ System Calls
- ❑ Deny writes to executables

1. Process Termination

- ❑ When a Process Terminates
 - `printf ("%s: exit(%d)\n", ...);`
- ❑ Name: Full name passed to `process_execute()` in `process.c`
- ❑ Exit Code
- ❑ Do not print any other message!

2. Argument Passing

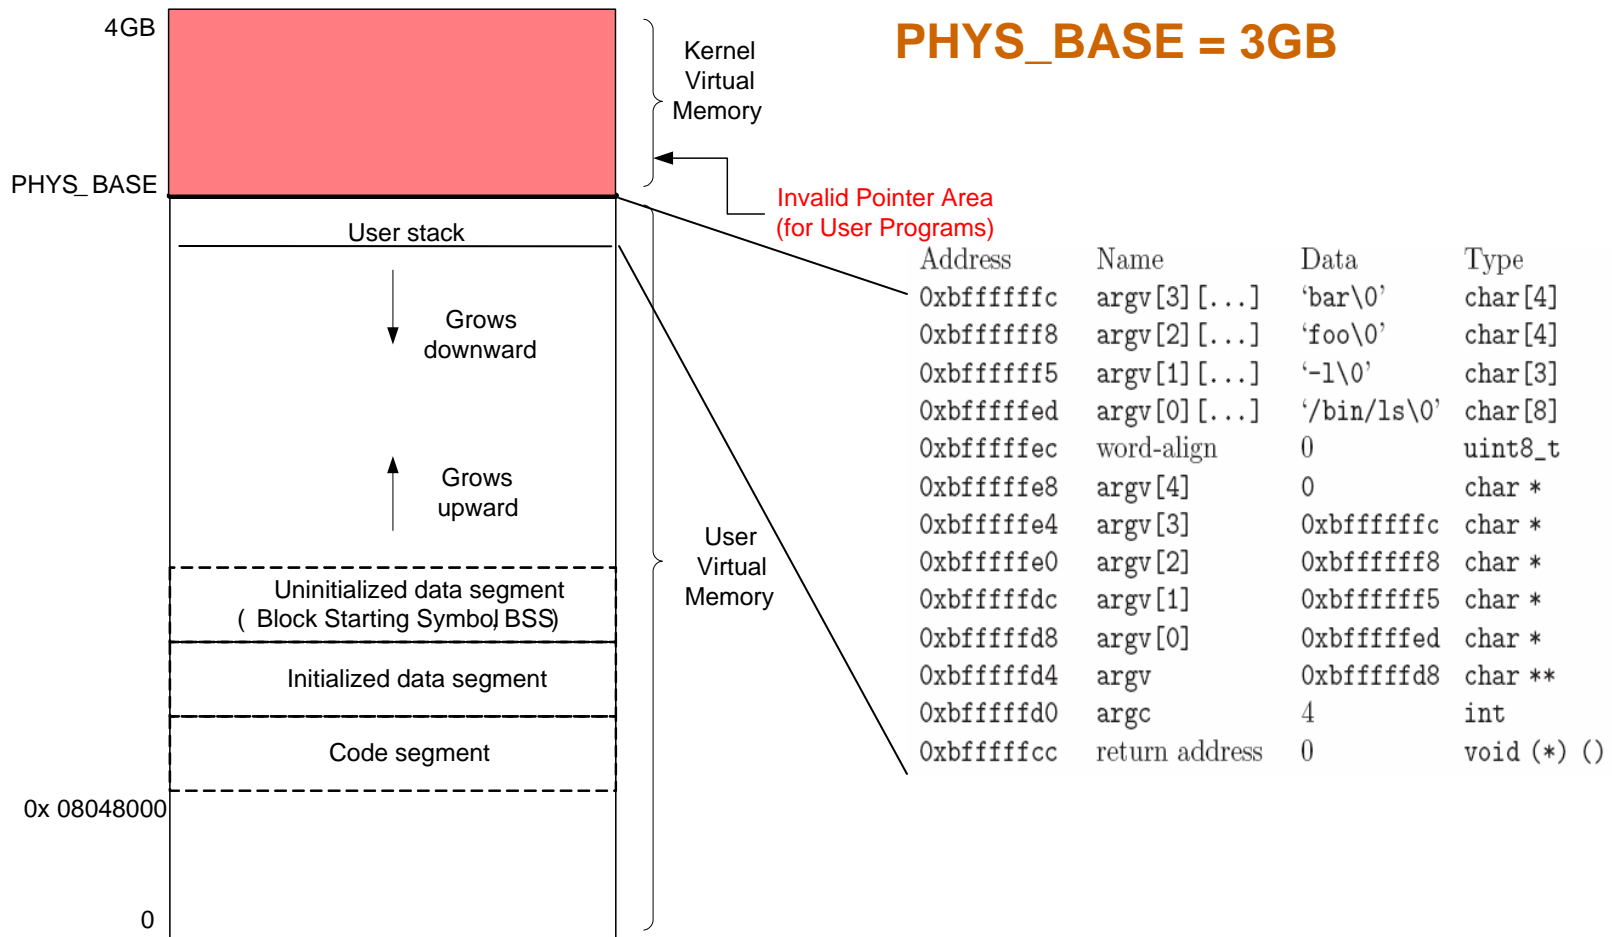
- ❑ No support currently for argument passing
- ❑ Change `*esp = PHYS_BASE` to `*esp = PHYS_BASE - 12` in `setup_stack()` to get started
- ❑ Change `process_execute()` to process multiple arguments
- ❑ Can limit the arguments to fit in a page(4 kb)
- ❑ String Parsing: `strtok_r()` in `lib/string.h`

```
pgm.c
main(int argc,
      char *argv[]) {
    ...
}

$ pintos run 'pgm alpha beta'
argc = 3
argv[0] = "pgm"
argv[1] = "alpha"
argv[2] = "beta"
```

Memory Layout

PHYS_BASE = 3GB



Setting up the Stack

How to setup the stack for the program - `/bin/ls -l foo bar`

Address	Name	Data	Type
0xbfffffffcc	argv[3] [...]	'bar\0'	char[4]
0xbfffffff8	argv[2] [...]	'foo\0'	char[4]
0xbffffff5	argv[1] [...]	'-l\0'	char[3]
0xbffffffed	argv[0] [...]	'/bin/ls\0'	char[8]
0xbffffffec	word-align	0	uint8_t
0xbffffffe8	argv[4]	0	char *
0xbffffffe4	argv[3]	0xbfffffffcc	char *
0xbffffffe0	argv[2]	0xbfffffff8	char *
0xbffffffdc	argv[1]	0xbffffff5	char *
0xbffffffd8	argv[0]	0xbffffffed	char *
0xbffffffd4	argv	0xbffffffd8	char **
0xbffffffd0	argc	4	int
0xbffffffcc	return address	0	void (*) ()

3. System Calls – already discussed some; e.g. `open()` system call

- ❑ No Support for system calls currently
- ❑ Implement the system call handler in `userprog/syscall.c`
- ❑ System call numbers defined in `lib/syscall-nr.h`
- ❑ Process Control: `exit`, `exec`, `wait`
- ❑ File system: `create`, `remove`, `open`, `filesize`, `read`, `write`, `seek`, `tell`, `close`
- ❑ Others: `halt`

System Call Details


- ❑ Types of Interrupts – External and Internal
- ❑ System calls – Internal Interrupts or Software Exceptions
- ❑ 80x86 – ‘int’ instruction to invoke system calls
- ❑ Pintos – ‘int \$0x30’ to invoke system call

Continued...

- A system call has:
 - System call number
 - (possibly) arguments
- When `syscall_handler()` gets control:

```
syscall_handler (struct intr_frame *f) {  
    f->esp _____  
    ....  
    f->eax = ... ;  
}
```

- System calls that return a value () must modify **f->eax**



Caller's User Stack

System calls – File system

- ❑ Decide on how to implement the file descriptors – keep it simple -- $O(n)$ data structures will entail no deduction
- ❑ Access granularity is the entire file system – add one global lock
- ❑ `write()` – fd 1 writes to console – use `putbuf()` to write entire buffer to console
- ❑ `read()` – fd 0 reads from console – use `input_getc()` to get input from keyboard
- ❑ Implement the rest of the system calls

System calls – Process Control

- ❑ `wait(pid)` – Waits for process `pid` to die and returns the status `pid` returned from `exit`
- ❑ Returns -1 if
 - `pid` was terminated by the kernel
 - `pid` does not refer to child of the calling thread
 - `wait()` has already been called for the given `pid`
- ❑ `exec(cmd)` – runs the executable whose name is given in command line
 - returns -1 if the program cannot be loaded
- ❑ `exit(status)` – terminates the current user program, returns `status`
 - status of 0 indicates success, non zero otherwise

Process Control – continued...

- ❑ Parent may or may not wait for its child
- ❑ Parent may call wait() after child terminates!
- ❑ Implement process_wait() in process.c
- ❑ Then, implement wait() in terms of process_wait()
- ❑ Cond variables and/or semaphores will help
 - Think about what semaphores may be used for and how they must be initialized

```
main() {  
    int i; pid_t p;  
    p = exec("pgm a b");  
    i = wait (p);  
    ... /* i must be 5 */  
}
```

```
main() {  
    int status;  
    ... status = 5;  
    exit(status);  
}      pgm.c
```

Memory Access

- ❑ Invalid pointers must be rejected. Why?
 - Kernel has access to all of physical memory including that of other processes
 - Kernel like user process would fault when it tries to access unmapped addresses
- ❑ User process cannot access kernel virtual memory
- ❑ User Process after it has entered the kernel can access kernel virtual memory and user virtual memory
- ❑ How to handle invalid memory access?

Memory Access – contd...

- ❑ Two methods to handle invalid memory access
 - Verify the validity of user provided pointer and then dereference it
 - ❑ Look at functions in `userprog/pagedir.c`, `threads/vaddr.h`
 - ❑ Strongly recommended!
 - Check if user pointer is below `PHYS_BASE` and dereference it
 - ❑ Could cause page fault
 - ❑ Handle the page fault by modifying the `page_fault()` code in `userprog/exception.c`
 - Make sure that resources are not leaked

Some Issues to look at...

- ❑ Check the validity of the system call parameters
- ❑ Every single location should be checked for validity before accessing it; e.g. not only for `f->esp`, but also the locations `f->esp + 1`, `f->esp+2` and `f->esp+3` should be checked
- ❑ Read system call parameters into kernel memory (except for long buffers) – write a `copy_in` function for this purpose.

Denying writes to Executables

- ❑ Use `file_deny_write()` to prevent writes to an open file
- ❑ Use `file_allow_write()` to re-enable write
- ❑ Closing a file will re-enable writes

Suggested Order of Implementation

- ❑ Change `*esp = PHYS_BASE` to `*esp = PHYS_BASE - 12` to get started
- ❑ Implement the system call infrastructure
- ❑ Change `process_wait()` to a infinite loop to prevent pintos getting powered off before the process gets executed
- ❑ Implement `exit` system call
- ❑ Implement `write` system call
- ❑ Start making other changes

Pintos Project 2 Sample Test

- ❑ Example Open System Call
- ❑ Test: tests/userprog/open-normal.c

```
/* Open a file. */
#include <syscall.h>
#include "tests/lib.h"
#include "tests/main.h"

void
test_main (void)
{
    int handle = open ("sample.txt");
    if (handle < 2)
        fail ("open() returned %d", handle);
}
```

userprog/ - make check

```
gcc -c ../../tests/userprog/open-normal.c -o tests/userprog/open-normal.o
-g -msoft-float -O -fno-stack-protector -nostdinc -I../../ -I../../lib
-I../../lib/user -I. -Wall -W -Wstrict-prototypes -Wmissing-prototypes
-Wsystem-headers -MMD -MF tests/userprog/open-normal.d
gcc -Wl,--build-id=none -nostdlib -static -Wl,-T,../../lib/user/user.lds
tests/userprog/open-normal.o tests/main.o tests/lib.o lib/user/entry.o
libc.a -o tests/userprog/open-normal
pintos -v -k -T 60 --qemu --filesystem-size=2 -p tests/userprog/open-normal
-a open-normal -p ../../tests/userprog/sample.txt -a sample.txt -- -q
-f run open-normal < /dev/null 2> tests/userprog/open-normal.errors
> tests/userprog/open-normal.output
perl -I../../ ../../tests/userprog/open-normal.ck
tests/userprog/open-normal tests/userprog/open-normal.result
pass tests/userprog/open-normal
-----
ar r libc.a lib/debug.o lib/random.o lib/stdio.o lib/stdlib.o lib/string.o
lib/arithmetic.o lib/ustar.o lib/user/debug.o lib/user/syscall.o
lib/user/console.o
ranlib libc.a
```


\$ objdump -D open-normal

080480a0 <test_main>:

80480a0:	55	push	%ebp
80480a1:	89 e5	mov	%esp,%ebp
80480a3:	83 ec 18	sub	\$0x18,%esp
80480a6:	c7 04 24 6a a7 04 08	movl	\$0x804a76a, (%esp)
80480ad:	e8 1f 21 00 00	call	804a1d1 <open>
80480b2:	83 f8 01	cmp	\$0x1,%eax

..

0804a1d1 <open>:

804a1d1:	55	push	%ebp
804a1d2:	89 e5	mov	%esp,%ebp
804a1d4:	ff 75 08	pushl	0x8(%ebp)
804a1d7:	6a 06	push	\$0x6
804a1d9:	cd 30	int	\$0x30
804a1db:	83 c4 08	add	\$0x8,%esp
804a1de:	5d	pop	%ebp
804a1df:	c3	ret	

..

0804a76a <.rodata.str1.1>:

804a76a:	73 61 6d 70 6c 65 2e 74 78 74 00
	s a m p l e . t x t

src/userprog/syscall.c

```
void
syscall_init (void)
{
    intr_register_int (0x30, 3, INTR_ON, syscall_handler,
        "syscall");
    lock_init (&fs_lock);
}
```

src/userprog/syscall.c

```
/* System call handler. */  
  
static void  
syscall_handler (struct intr_frame *f)  
{  
    typedef int syscall_function (int, int, int);  
    /* A system call. */  
    struct syscall  
    {  
        size_t arg_cnt;          /* Number of arguments. */  
        syscall_function *func;  /* Implementation. */  
    };  
    /* Table of system calls. */  
    static const struct syscall syscall_table[] =  
    {  
        {0, (syscall_function *) sys_halt},  
        {1, (syscall_function *) sys_exit},  
        {1, (syscall_function *) sys_exec},  
        {1, (syscall_function *) sys_wait},  
        {2, (syscall_function *) sys_create},  
        {1, (syscall_function *) sys_remove},  
        {1, (syscall_function *) sys_open},  <-- call number 6
```

src/userprog/syscall.c

```
...
const struct syscall *sc;
unsigned call_nr;
int args[3];

/* Get the system call. */
copy_in (&call_nr, f->esp, sizeof call_nr);
if (call_nr >= sizeof syscall_table / sizeof *syscall_table)
    thread_exit ();
sc = syscall_table + call_nr;

/* Get the system call arguments. */
ASSERT (sc->arg_cnt <= sizeof args / sizeof *args);
memset (args, 0, sizeof args);
copy_in (args, (uint32_t *) f->esp + 1, sizeof *args * sc->arg_cnt);

/* Execute the system call,
   and set the return value. */
f->eax = sc->func (args[0], args[1], args[2]);
}
```

src/userprog/syscall.c

```
/* Open system call. */
static int
sys_open (const char *ufile)
{
    char *kfile = copy_in_string (ufile);
    struct file_descriptor *fd;
    int handle = -1;

    fd = malloc (sizeof *fd);
    if (fd != NULL)
    {
        lock_acquire (&fs_lock);
        fd->file = filesys_open (kfile);
        if (fd->file != NULL)
        ...
    }
```

Copy a byte from user space to kernel space

```
/* Copies a byte from user address USRC to kernel address DST. */
static inline bool
get_user (uint8_t *dst, const uint8_t *usrc)
{
    int eax;
    asm ("movl $1f, %%eax; movb %2, %%al; movb %%al, %0; 1:"
        : "=m" (*dst), "=&a" (eax) : "m" (*usrc));
    return eax != 0;
}

/* Writes BYTE to user address UDST. */
static inline bool
put_user (uint8_t *udst, uint8_t byte)
{
    int eax;
    asm ("movl $1f, %%eax; movb %b2, %0; 1:"
        : "=m" (*udst), "=&a" (eax) : "q" (byte));
    return eax != 0;
}
```

Copy a string from user space to kernel space

```
static char *
copy_in_string (const char *us)
{
    char *ks;
    size_t length;
    ks = pallocc_get_page (0);
    if (ks == NULL)
        thread_exit ();
    for (length = 0; length < PGSIZE; length++)
    {
        if (us >= (char *) PHYS_BASE || !get_user (ks + length, us++))
        {
            pallocc_free_page (ks);
            thread_exit ();
        }
        if (ks[length] == '\\0')
            return ks;
    }
    ks[PGSIZE - 1] = '\\0';
    return ks;
}
```

Run a single test

□ From userprog/build:

- **rm tests/userprog/open-normal.output**
- **make test/userprog/open-normal.output**

```
pintos -v -k -T 60 --qemu --filesys-size=2 -p tests/userprog/open-normal  
-a open-normal -p ../../tests/userprog/sample.txt -a sample.txt -- -q -f run  
open-normal < /dev/null 2> tests/userprog/open-normal.errors >  
tests/userprog/open-normal.output
```


examples/shell.c

```
int main (void)
{
    printf ("Shell starting..\n");
    for (;;)
    {
        char command[80];
        /* Read command. */
        printf ("--");
        read_line (command, sizeof command);

        /* Execute command. */
        if (!strcmp (command, "exit"))
            break;
        else if (!memcmp (command, "cd ", 3))
        {
            if (!chdir (command + 3))
                printf ("\"%s\": chdir failed\n", command + 3);
        }
        else if (command[0] == '\0')
```

examples/shell.c

```
else if (command[0] == '\\0')
{
    /* Empty command. */
}
else
{
    pid_t pid = exec (command);
    if (pid != PID_ERROR)
        printf ("\"%s\": exit code %d\\n", command, wait (pid));
    else
        printf ("exec failed\\n");
}

printf ("Shell exiting.");
return EXIT_SUCCESS;
}
```

\$ objdump -D examples/shell

080480cc <main>:

80480cc:	55	push	%ebp
80480cd:	89 e5	mov	%esp,%ebp
80480cf:	83 e4 f0	and	\$0xfffffffff0,%esp
80480d2:	57	push	%edi
80480d3:	56	push	%esi
80480d4:	53	push	%ebx
80480d5:	83 ec 74	sub	\$0x74,%esp
80480d8:	c7 04 24 3e a1 04 08	movl	\$0x804a13e, (%esp)
80480df:	e8 23 1d 00 00	call	8049e07 <puts>
80480e4:	c7 04 24 5a a1 04 08	movl	\$0x804a15a, (%esp)
80480eb:	e8 2f 0c 00 00	call	8048d1f <printf>
80480f0:	8d 44 24 18	lea	0x18(%esp), %eax
80480f4:	89 44 24 68	mov	%eax, 0x68(%esp)
80480f8:	89 c3	mov	%eax,%ebx
80480fa:	c7 44 24 08 01 00 00	movl	\$0x1, 0x8(%esp)
8048101:	00		
8048102:	8d 44 24 6f	lea	0x6f(%esp), %eax
8048106:	89 44 24 04	mov	%eax, 0x4(%esp)
804810a:	c7 04 24 00 00 00 00	movl	\$0x0, (%esp)
8048111:	e8 a9 1a 00 00	call	8049bbf <read>

\$ objdump -D examples/shell

```
..
80481f1:      80 7c 24 18 00      cmpb    $0x0,0x18(%esp)
80481f6:      0f 84 e8 fe ff ff   je      80480e4 <main+0x18>
80481fc:      8d 44 24 18      lea     0x18(%esp),%eax
8048200:      89 04 24      mov     %eax,(%esp)
8048203:      e8 50 19 00 00     call    8049b58 <exec>
..
08049bbf <read>:
8049bbf:      55      push    %ebp
8049bc0:      89 e5     mov     %esp,%ebp
8049bc2:      ff 75 10   pushl   0x10(%ebp)
8049bc5:      ff 75 0c   pushl   0xc(%ebp)
8049bc8:      ff 75 08   pushl   0x8(%ebp)
8049bcb:      6a 08     push    $0x8
8049bcd:      cd 30     int     $0x30
8049bcf:      83 c4 10   add     $0x10,%esp
8049bd2:      5d      pop     %ebp
8049bd3:      c3      ret
```

syscall exec() calls process_execute

```
/* Starts a new thread running a user program loaded from
   FILENAME.  The new thread may be scheduled (and may even exit)
   before process_execute() returns.  Returns the new process's
   thread id, or TID_ERROR if the thread cannot be created. */
tid_t process_execute (const char *file_name)
{
    struct exec_info exec;
    char thread_name[15];
    char *save_ptr;
    tid_t tid;
    /* Initialize exec_info. */
    exec.file_name = file_name;
    sema_init (&exec.load_done, 0);
    /* Create a new thread to execute FILE_NAME. */
    strcpy (thread_name, file_name, sizeof thread_name);
    strtok_r (thread_name, " ", &save_ptr);
    tid = thread_create (thread_name, PRI_DEFAULT, start_process, &exec);
```

process_execute

```
/* Starts a new thread running a user program loaded from
   FILENAME.  The new thread may be scheduled (and may even exit)
   before process_execute() returns.  Returns the new process's
   thread id, or TID_ERROR if the thread cannot be created. */
tid_t process_execute (const char *file_name)
{
    ..

    if (tid != TID_ERROR)
    {
        sema_down (&exec.load_done);
        if (exec.success)
            list_push_back (&thread_current ()->children,
                           &exec.wait_status->elem);
        else
            tid = TID_ERROR;
    }
    return tid;
}
```

start_process

```
static void start_process (void *exec_)
{
    struct exec_info *exec = exec_;
    struct intr_frame if_;
    bool success;

    /* Initialize interrupt frame and load executable. */
    memset (&if_, 0, sizeof if_);
    if_.gs = if_.fs = if_.es = if_.ds = if_.ss = SEL_UDSEG;
    if_.cs = SEL_UCSEG;
    if_.eflags = FLAG_IF | FLAG_MBS;
    success = load (exec->file_name, &if_.eip, &if_.esp);

    /* Allocate wait_status. */
    if (success)
    {
        exec->wait_status = thread_current ()->wait_status
            = malloc (sizeof *exec->wait_status);
        success = exec->wait_status != NULL;
    }
}
```

start_process

```
/* Initialize wait_status. */
if (success)
{
    ..
}
/* Notify parent thread and clean up. */
exec->success = success;
sema_up (&exec->load_done);
if (!success)
    thread_exit ();
/* Start the user process by simulating a return from an
interrupt, implemented by intr_exit (in
threads/intr-stubs.S). Because intr_exit takes all of its
arguments on the stack in the form of a `struct intr_frame',
we just point the stack pointer (%esp) to our stack frame
and jump to it. */
asm volatile ("movl %0, %%esp; jmp intr_exit" : : "g" (&if_) : "memory");
}
```


Cycle Counters

- Most modern systems have built in registers that are incremented every clock cycle
 - Very fine grained
 - Maintained as part of process state
 - In Linux, counts elapsed global time
- Special assembly code instruction to access
- On (recent model) Intel machines:
 - 64 bit counter.
 - RDTSC instruction sets `%edx` to high order 32-bits, `%eax` to low order 32-bits
- Aside: Is this a security issue?

Cycle Counter Period

❑ Wrap Around Times for 550 MHz machine

- Low order 32 bits wrap around every $2^{32} / (550 * 10^6) = 7.8$ seconds
- High order 64 bits wrap around every $2^{64} / (550 * 10^6) = 33539534679$ seconds
 - ❑ 1065 years

❑ For 2 GHz machine

- Low order 32 bits every 2.1 seconds
- High order 64 bits every 293 years

Measuring with Cycle Counter

■ Idea

- Get current value of cycle counter
 - store as pair of unsigned's `cyc_hi` and `cyc_lo`
- Compute something
- Get new value of cycle counter
- Perform double precision subtraction to get elapsed cycles

```
/* Keep track of most recent reading of cycle counter */
static unsigned cyc_hi = 0;
static unsigned cyc_lo = 0;

void start_counter()
{
    /* Get current value of cycle counter */
    access_counter(&cyc_hi, &cyc_lo);
}
```

Accessing the Cycle Counter

- GCC allows inline assembly code with mechanism for matching registers with program variables
- Code only works on x86 machine compiling with GCC

```
void access_counter(unsigned *hi, unsigned *lo)
{
    /* Get cycle counter */
    asm("rdtsc; movl %%edx,%0; movl %%eax,%1"
        : "=r" (*hi), "=r" (*lo)
        : /* No input */
        : "%edx", "%eax");
}
```

- Emit assembly with `rdtsc` and two `movl` instructions

Closer Look at Extended ASM

```
asm("Instruction String"  
    : Output List  
    : Input List  
    : Clobbers List);  
}
```

```
void access_counter  
    (unsigned *hi, unsigned *lo)  
{  
    /* Get cycle counter */  
    asm("rdtsc; movl %%edx,%0; movl %%eax,%1"  
        : "=r" (*hi), "=r" (*lo)  
        : /* No input */  
        : "%edx", "%eax");  
}
```

■ Instruction String

□ Series of assembly commands

- Separated by “;” or “\n”
- Use “%%” where normally would use “%”

Closer Look at Extended ASM

```
asm("Instruction String"  
    : Output List  
    : Input List  
    : Clobbers List  
    )
```

```
void access_counter  
(unsigned *hi, unsigned *lo)  
{  
    /* Get cycle counter */  
    asm("rdtsc; movl %%edx,%0; movl %%eax,%1"  
        : "=r" (*hi), "=r" (*lo)  
        : /* No input */  
        : "%edx", "%eax");  
}
```

□ Output List

- Expressions indicating destinations for values %0, %1, ..., %j
 - Enclosed in parentheses
 - Must be *lvalue*
 - Value that can appear on LHS of assignment
- Tag "=r" indicates that symbolic value (%0, etc.), should be replaced by a register

Closer Look at Extended ASM

```
asm("Instruction String"  
    : Output List  
    : Input List  
    : Clobbers List)  
}
```

```
void access_counter  
    (unsigned *hi, unsigned *lo)  
{  
    /* Get cycle counter */  
    asm("rdtsc; movl %%edx,%0; movl %%eax,%1"  
        : "=r" (*hi), "=r" (*lo)  
        : /* No input */  
        : "%edx", "%eax");  
}
```

■ Input List

- Series of expressions indicating sources for values $\%j+1$, $\%j+2$, ...
 - Enclosed in parentheses
 - Any expression returning value
- Tag "r" indicates that symbolic value ($\%0$, etc.) will come from register

Closer Look at Extended ASM

```
asm("Instruction String"  
    : Output List  
    : Input List  
    : Clobbers List);  
}
```

```
void access_counter  
    (unsigned *hi, unsigned *lo)  
{  
    /* Get cycle counter */  
    asm("rdtsc; movl %%edx,%0; movl %%eax,%1"  
        : "=r" (*hi), "=r" (*lo)  
        : /* No input */  
        : "%edx", "%eax");  
}
```

□ Clobbers List

- List of register names that get altered by assembly instruction
- Compiler will make sure doesn't store something in one of these registers that must be preserved across asm
 - Value set before & used after

Completing Measurement

- Get new value of cycle counter
- Perform double precision subtraction to get elapsed cycles
- Express as double to avoid overflow problems

```
double get_counter()
{
    unsigned ncyc_hi, ncyc_lo
    unsigned hi, lo, borrow;
    /* Get cycle counter */
    access_counter(&ncyc_hi, &ncyc_lo);
    /* Do double precision subtraction */
    lo = ncyc_lo - cyc_lo;
    borrow = lo > ncyc_lo;
    hi = ncyc_hi - cyc_hi - borrow;
    return (double) hi * (1 << 30) * 4 + lo;
}
```

Timing With Cycle Counter

□ Determine Clock Rate of Processor

- Count number of cycles required for some fixed number of seconds

```
double MHZ;  
int sleep_time = 10;  
start_counter();  
sleep(sleep_time);  
MHZ = get_counter() / (sleep_time * 1e6);
```

□ Time Function P()

- First attempt: Simply count cycles for one execution of P

```
double tsecs;  
start_counter();  
P();  
tsecs = get_counter() / (MHZ * 1e6);
```

Example – testClock.c

```
#include <stdio.h>
#include "clock.h"
```

```
int main()
{
```

Processor Clock Rate ~= 2673.5 MHz

cycles = 5343976388.000000, MHz = 2673.526339, cycles/Mhz = 1998849.351153
elapsed time = 1.998849 seconds

```
    double cycles, Mhz;
```

```
    Mhz = mhz(1);
```

```
    start_counter();
```

```
    sleep(2);
```

```
    cycles = get_counter();
```

```
    printf("cycles = %f, MHz = %f, cycles/Mhz = %f\n", cycles, Mhz,
cycles/Mhz);
```

```
    printf("elapsed time = %f seconds \n", cycles/(1.0e6*Mhz));
```

```
    return 0;
```

```
}
```

Measurement Pitfalls

▣ Overhead

- Calling `get_counter()` incurs small amount of overhead
- Want to measure long enough code sequence to compensate