

# The Pintos Instructional Operating System Kernel

Ben Pfaff

blp@nicira.com

Nicira Networks

Anthony Romano

ajromano@stanford.edu

Stanford University

Godmar Back

gback@cs.vt.edu

Virginia Tech

# Overview

- Tool paper
- Series of 4 projects that provide backbone of lab component that accompanies OS course
- Suitable for Junior/Senior/1<sup>st</sup> Grad students
- Used by several institutions
  - Stanford (4 years+), Virginia Tech (3 years), University of San Francisco, William and Mary, University of Salzburg, Linköping Universitet, KAIST, Seoul National University, POSTECH

# Teaching OS

- Internal Perspective
  - Teaches how an OS works from the inside, specifically, the kernel
  - Places students in the perspective of OS designer, rather than OS user
- Concrete Approach
  - Design and create realistic artifacts
  - Internalize abstractions by seeing concrete incarnation

# Pintos Features

- Small enough so entire code can be read and understood by students
  - Unlike Linux or Windows
- Runs and debugs in simulated environment
  - 100% reproducible, like single-threaded user code
- Runs and debugs in emulated environment
  - facilitates non-intrusive analysis tools
- Runs on real hardware
  - boots on student's PC/laptop

```
USB Device 1: Fingerprint Sensor (      )
UHCI: Enabling 2 root ports
USB: scanning devices...
UHCI: Enabling 2 root ports
USB: scanning devices...
USB Device 1: Flashdrive 383B (Memorex )
uda: 247,616 sectors (120 MB), USB
uda1: 945 sectors (472 kB), Pintos OS kernel (20)
uda2: 9,072 sectors (4 MB), Pintos file system (21)
uda3: 1,000 sectors (504 kB), Pintos scratch (22)
fileys: using uda2
scratch: using uda3
Boot complete.
Executing 'shell':
Shell starting...The best operating system!
--echo Hello World
echo Hello World
echo: exit(0)
"echo Hello World": exit code 0
--shell
Shell starting...The best operating system!
--exit
Shell exiting.shell: exit(0)
"shell": exit code 0
--
_
```

# Project Principles (1)

- Read Before You Code
  - Provide well-documented code that serves as example of what we expect from students
  - Between 0-600 lines per project

# Project Principles (2)

- **Maximize Creative Freedom**
  - Specify requirements
  - Don't prescribe solution approaches

# Project Principles (3)

- **Practice Test-driven Development**
  - All tests are public, reading tests makes requirements concrete
  - Student can add their own tests

Project	Functionality	Robustness	Regression
1	27		
2	41	35	
3	20	14	75
4	39	7	75



# Project Principles (4)

- **Justify Your Design**

- Provide structured questionnaires that students use to describe and justify their design rationale

# Project Principles (5)

- **Work In Teams**
  - 2-4 students
  - Allows for brainstorming and mutual support
  - Mimics industrial setting, e.g., use of shared source code repository and versioning
  - Design questionnaires still submitted individually

# Pintos Project Themes

1. Threads
2. User Programs
3. Virtual Memory
4. File Systems

Priority Scheduling

MLFQS Scheduling

Alarm  
Clock

P1: Kernel-mode Test Cases

Threading  
Simple Scheduler

Device Support  
Keyboard, VGA, USB, Serial Port, Timer, PCI, IDE

Boot Support

Pintos Kernel

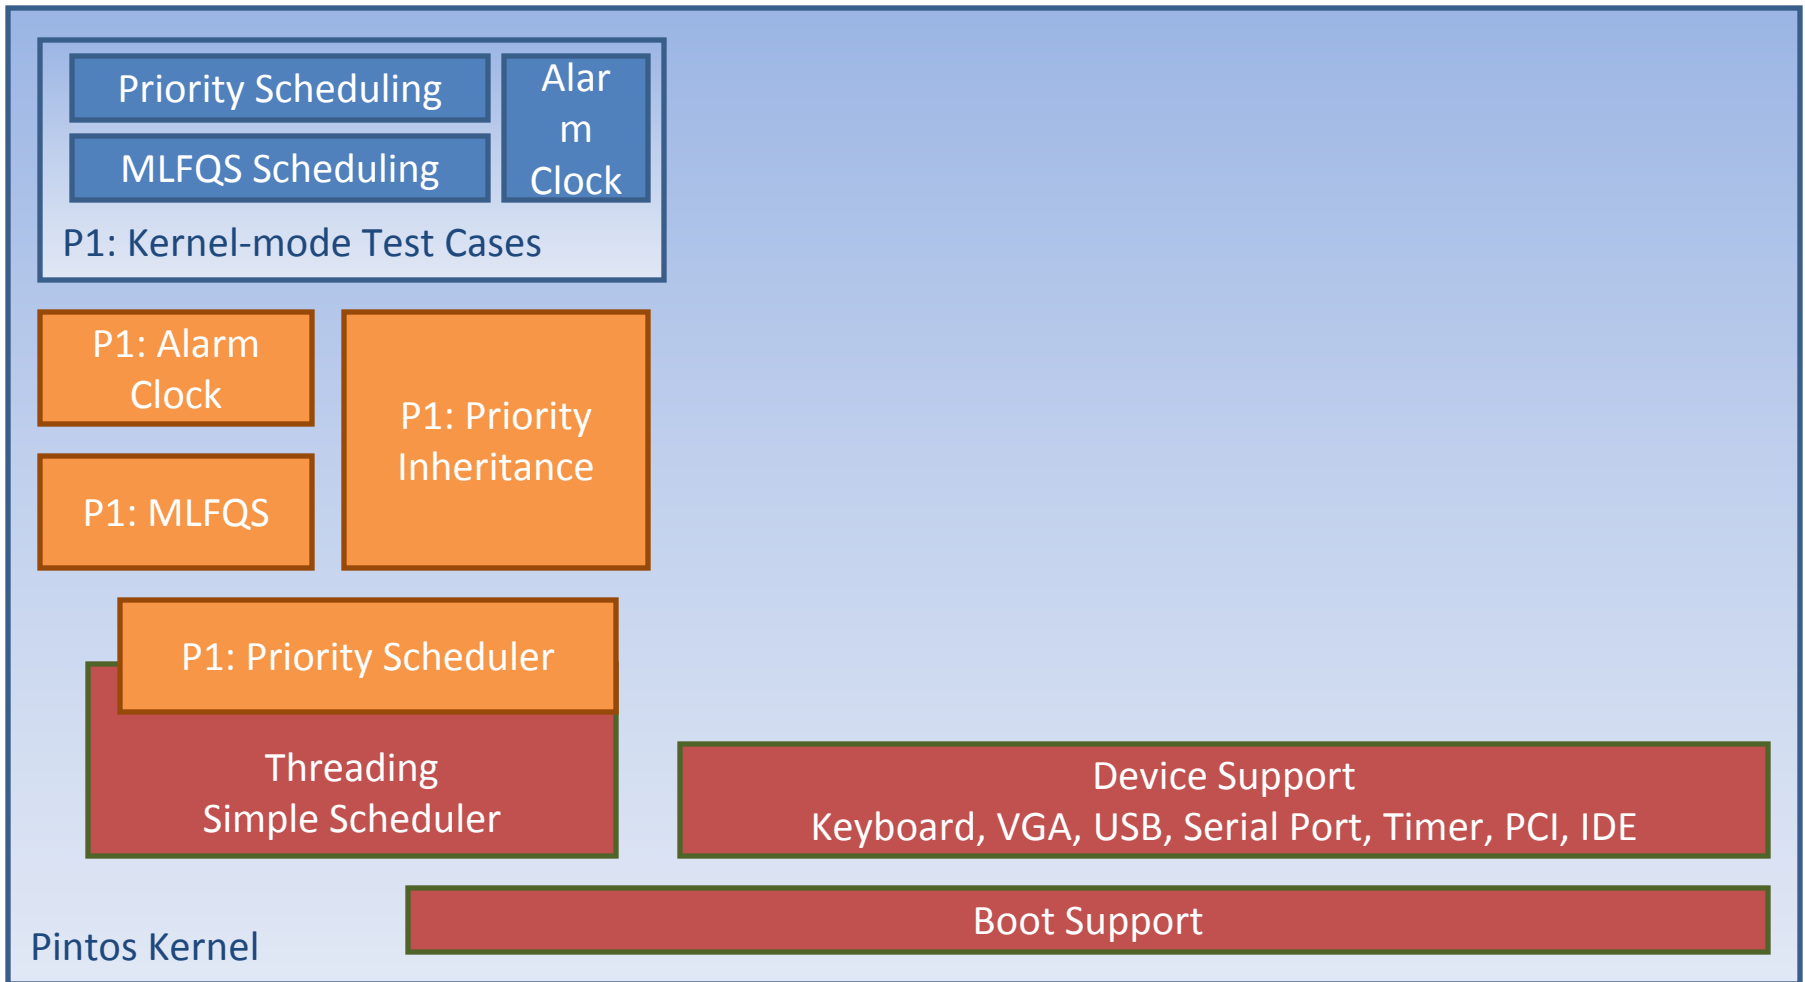
3/7/2009

Support Code

SIGCSE 2009

Public Tests

Pre Project 1<sup>12</sup>



# Example of Project 1 Test

```
void
test_priority_change (void)
{
    msg ("Creating a high-priority thread 2.");
    thread_create ("thread 2", PRI_DEFAULT + 1, changing_thread, NULL);
    msg ("Thread 2 should have just lowered its priority.");
    thread_set_priority (PRI_DEFAULT - 2);
    msg ("Thread 2 should have just exited.");
}
```

```
static void
changing_thread (void *aux UNUSED)
{
    msg ("Thread 2 now lowering priority.");
    thread_set_priority (PRI_DEFAULT - 1);
    msg ("Thread 2 exiting.");
}
```

*Expected output:*

Creating a high-priority thread 2.  
Thread 2 now lowering priority.  
Thread 2 should have just lowered its priority.  
Thread 2 exiting.  
Thread 2 should have just exited.

# make grade (1)

TOTAL TESTING SCORE: 100.0%

ALL TESTED PASSED -- PERFECT SCORE

---

## SUMMARY BY TEST SET

Test Set	Pts	Max	% Ttl	% Max
tests/threads/Rubric.alarm	18/	18	20.0%/	20.0%
tests/threads/Rubric.priority	38/	38	40.0%/	40.0%
tests/threads/Rubric.mlfqs	37/	37	40.0%/	40.0%
Total			100.0%/	100.0%

---

*Pintos include fully automated grading scripts, students see score before submission*

# make grade (2)

## SUMMARY OF INDIVIDUAL TESTS

Functionality and robustness of alarm clock (tests/threads/Rubric.alarm):

4/ 4 tests/threads/alarm-single  
4/ 4 tests/threads/alarm-multiple  
4/ 4 tests/threads/alarm-simultaneous  
4/ 4 tests/threads/alarm-priority  
  
1/ 1 tests/threads/alarm-zero  
1/ 1 tests/threads/alarm-negative

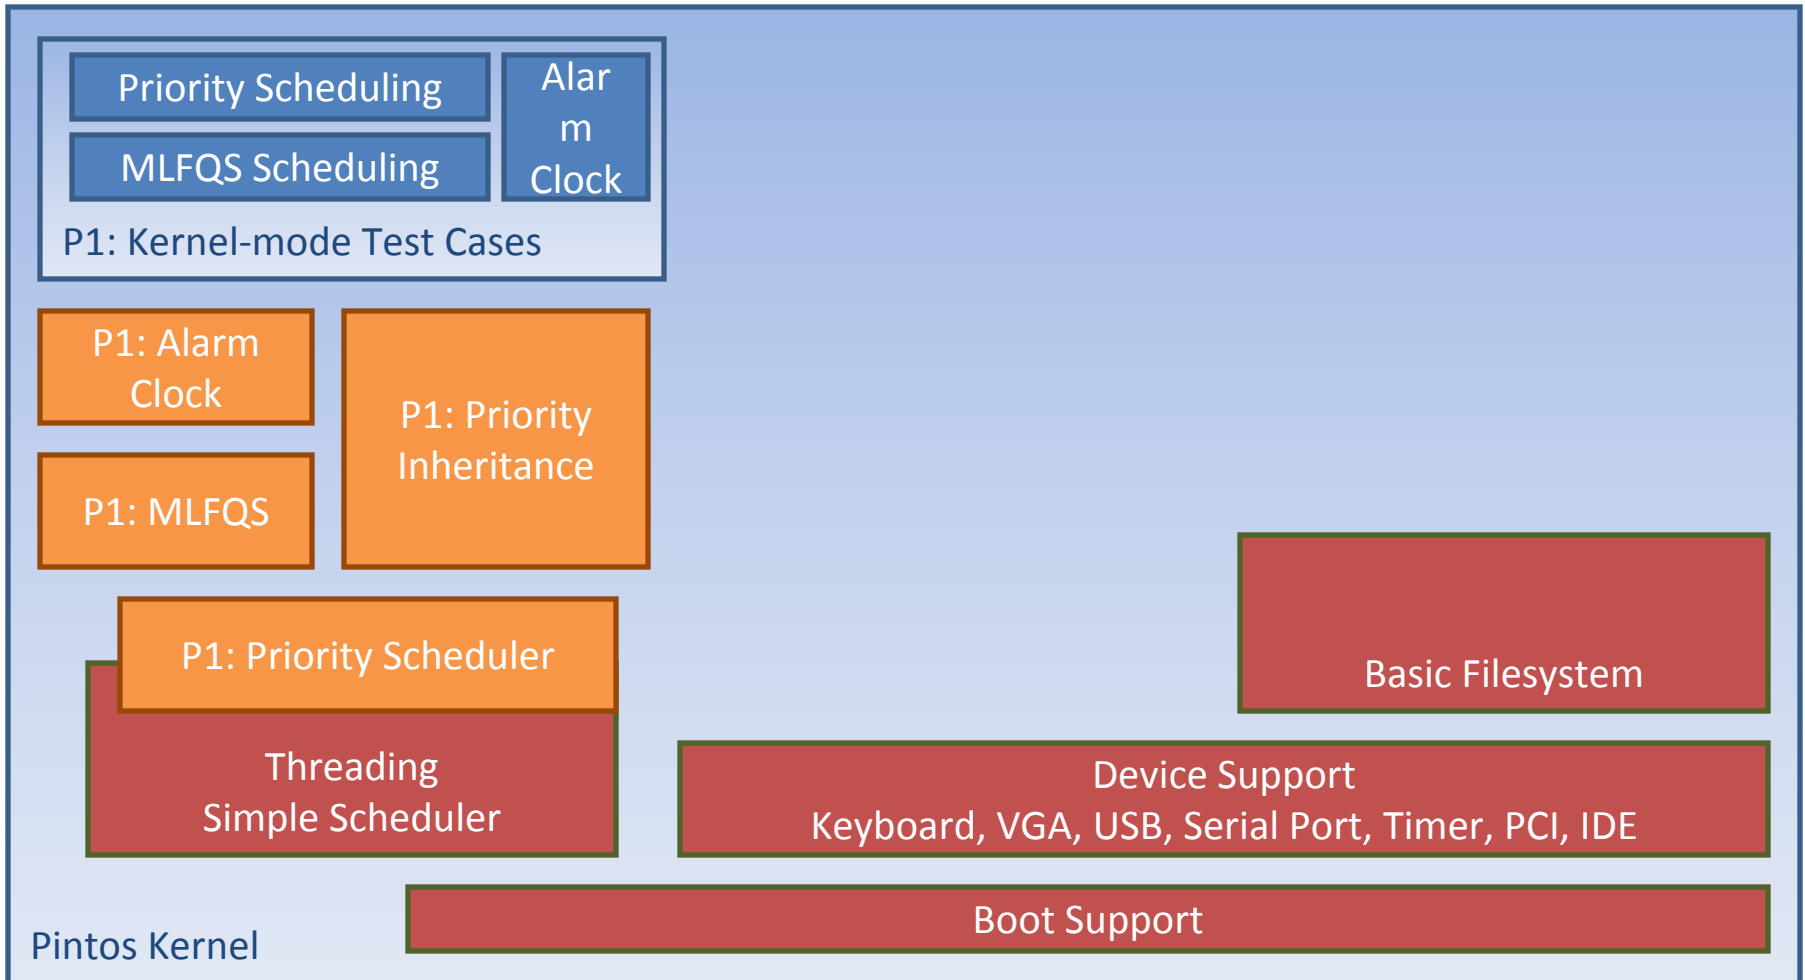
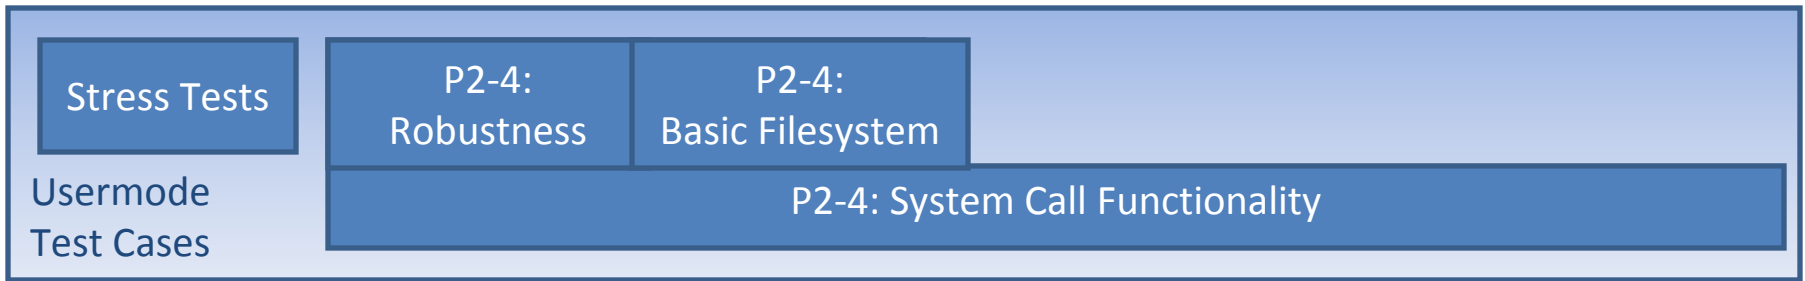
- Section summary.

6/ 6 tests passed  
18/ 18 points subtotal

Functionality of priority scheduler (tests/threads/Rubric.priority):

**3/ 3 tests/threads/priority-change**  
3/ 3 tests/threads/priority-preempt





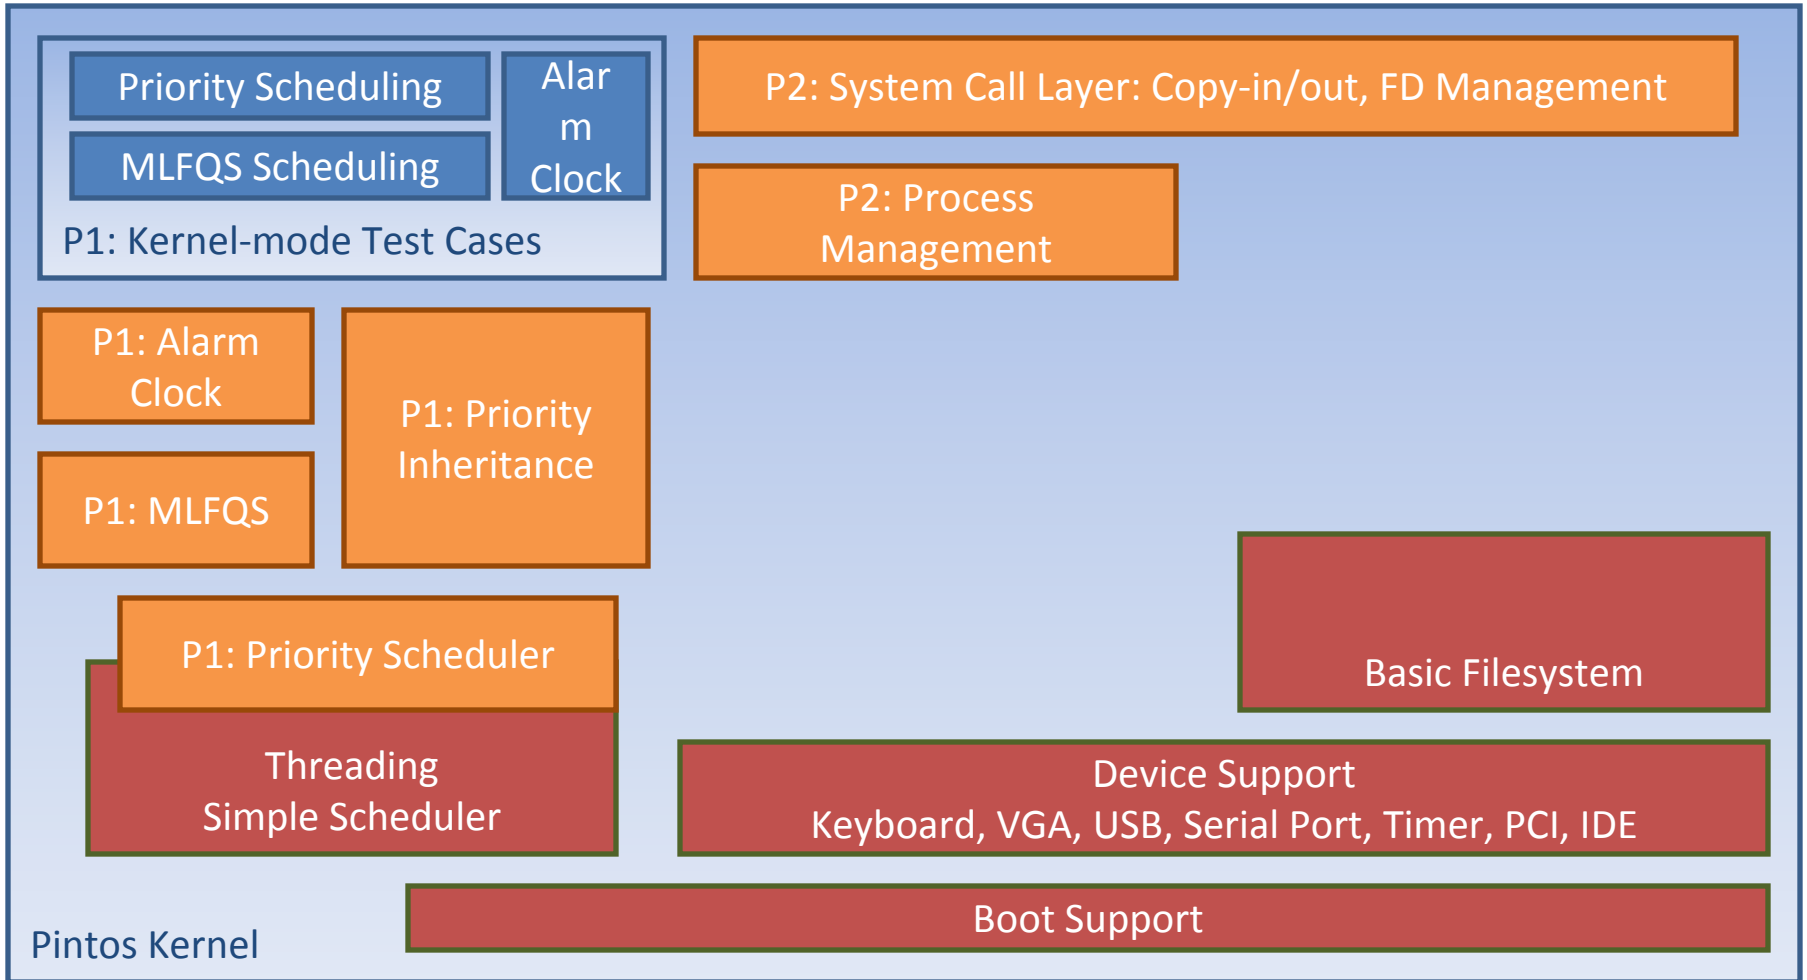
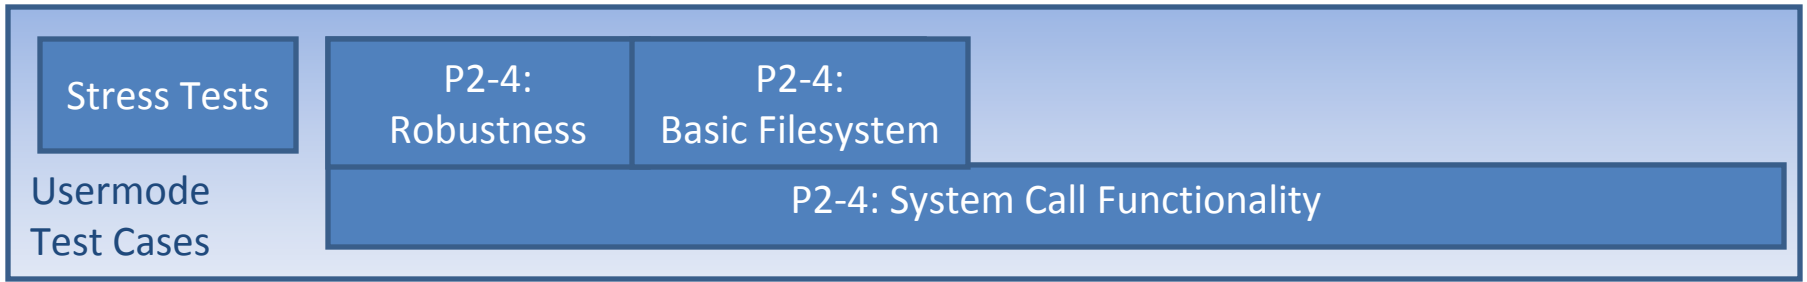
3/7/2009

Support Code

Students Create

Public Tests

Pre Project 2



# Project 2 Functionality Test

```
/* This program echoes its command-line arguments */
```

```
int  
main (int argc, char *argv[])  
{  
    int i;  
    msg ("begin");  
    msg ("argc = %d", argc);  
    for (i = 0; i <= argc; i++)  
        if (argv[i] != NULL)  
            msg ("argv[%d] = '%s'", i, argv[i]);  
        else  
            msg ("argv[%d] = null", i);  
    msg ("end");  
    return 0;  
}
```

*Expected output for 'args 1 2'*

```
begin  
argc=3  
argv[0] = 'args'  
argv[1] = '1'  
argv[2] = '2'  
argv[3] = null  
end
```

# Project 2 Robustness Test

*/\* This program attempts to read memory at an address that is not mapped.*

*This should terminate the process with a -1 exit code. \*/*

```
#include "tests/lib.h"
```

```
#include "tests/main.h"
```

*Expected output:*  
bad-read: exit(-1)

```
void
```

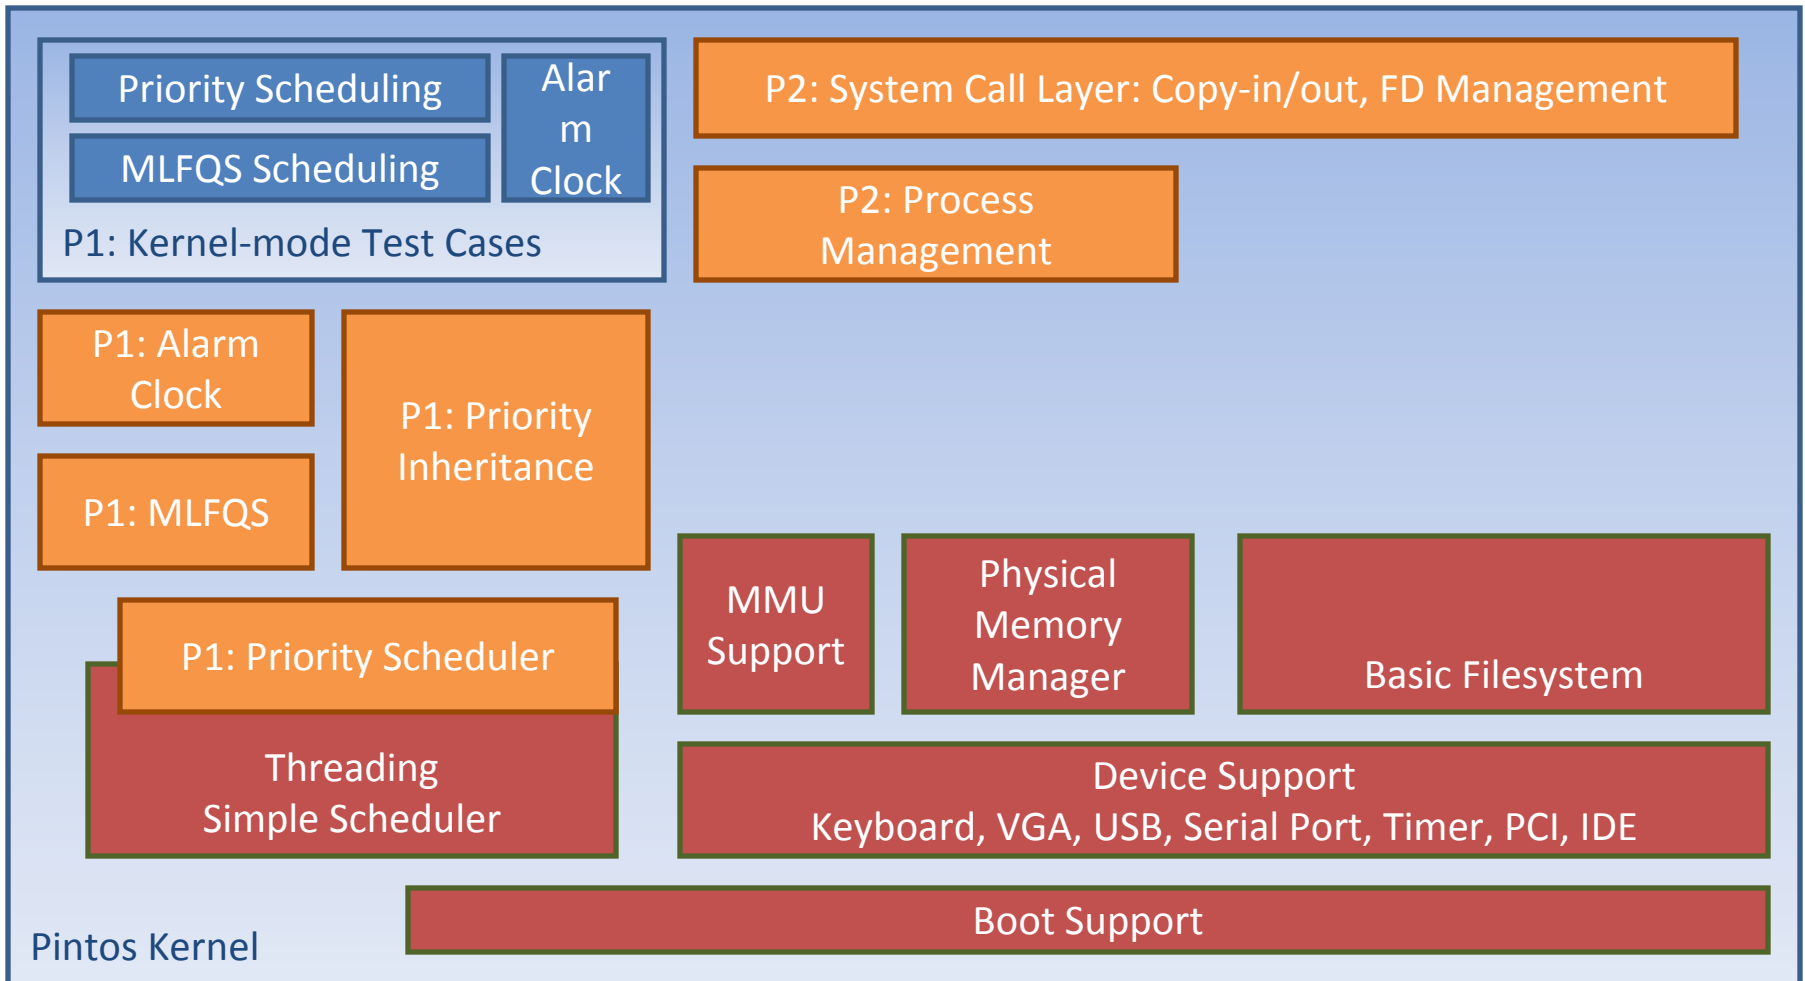
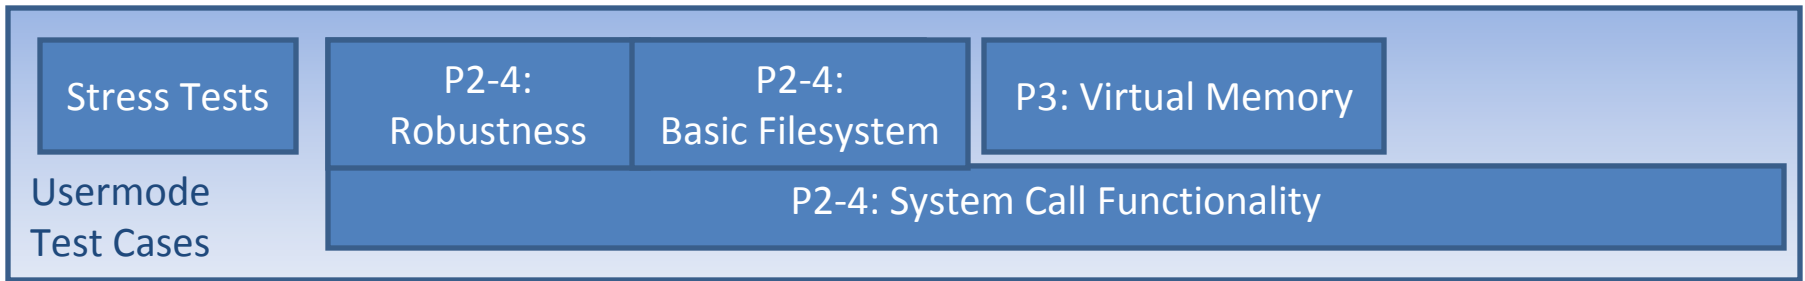
```
test_main (void)
```

```
{
```

```
    msg ("Congratulations - you have successfully dereferenced NULL: %d",  
         *(int *)NULL);
```

```
    fail ("should have exited with -1");
```

```
}
```



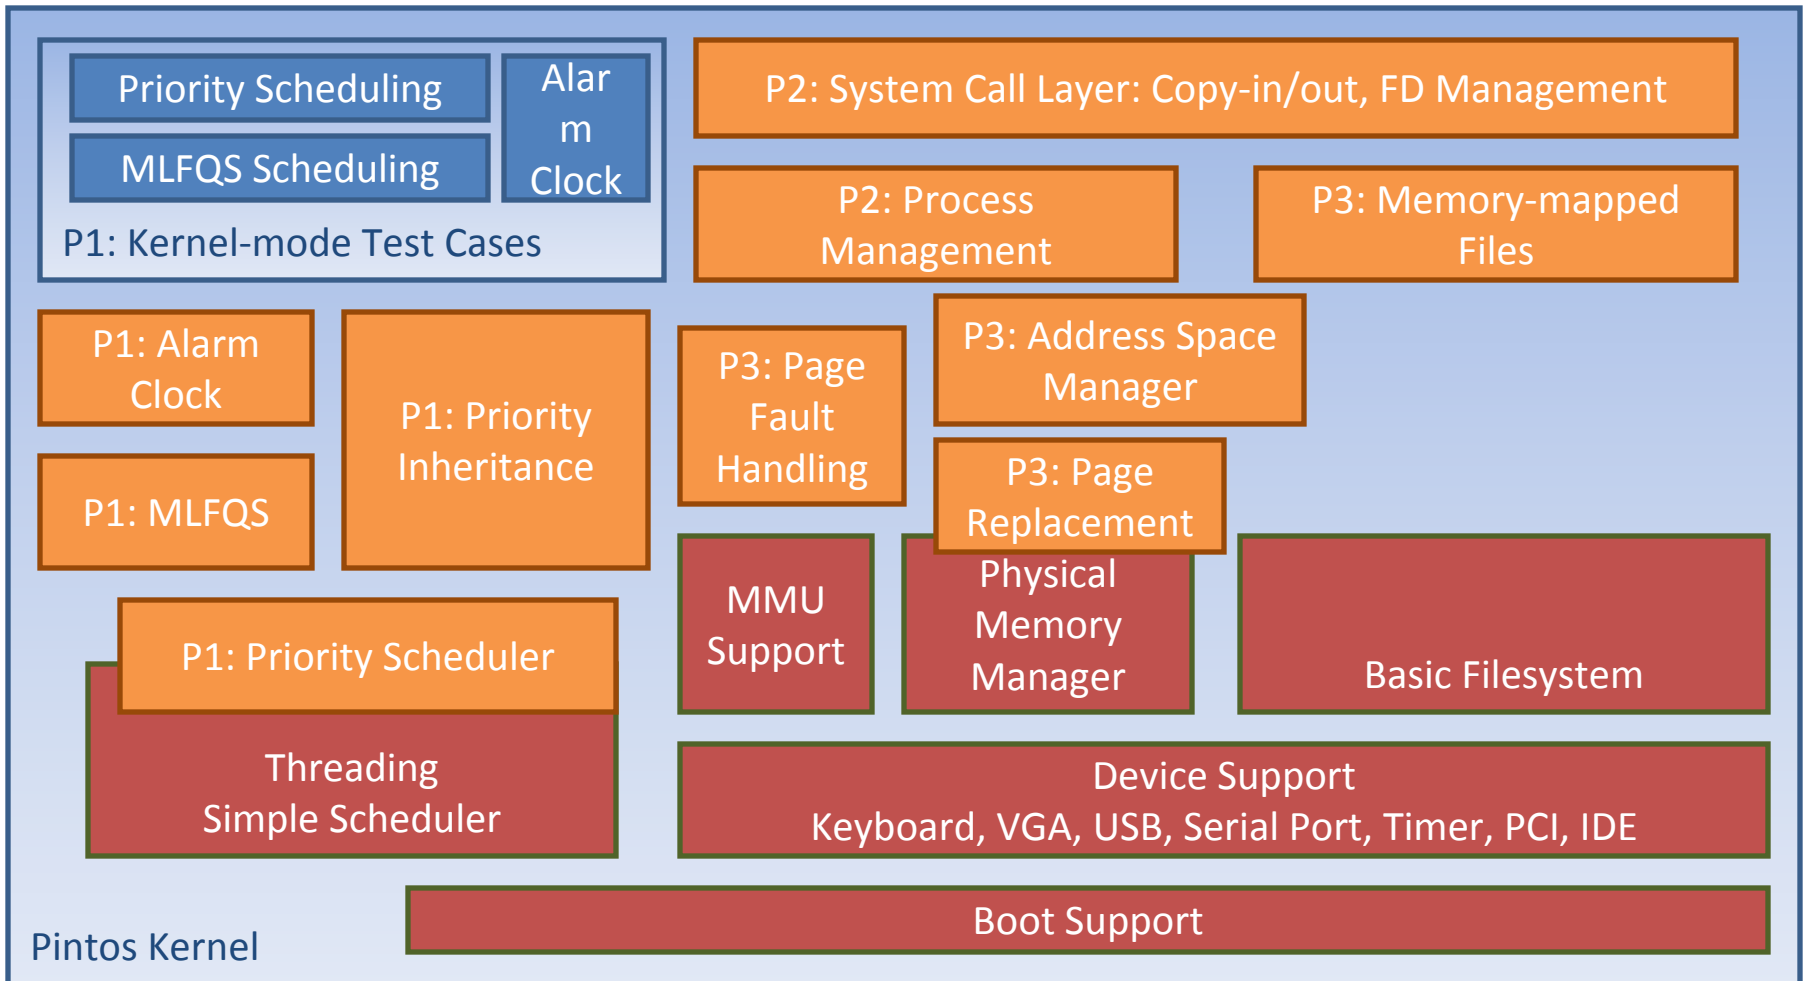

3/7/2009

Support Code

Students Create

Public Tests

Pre Project 3



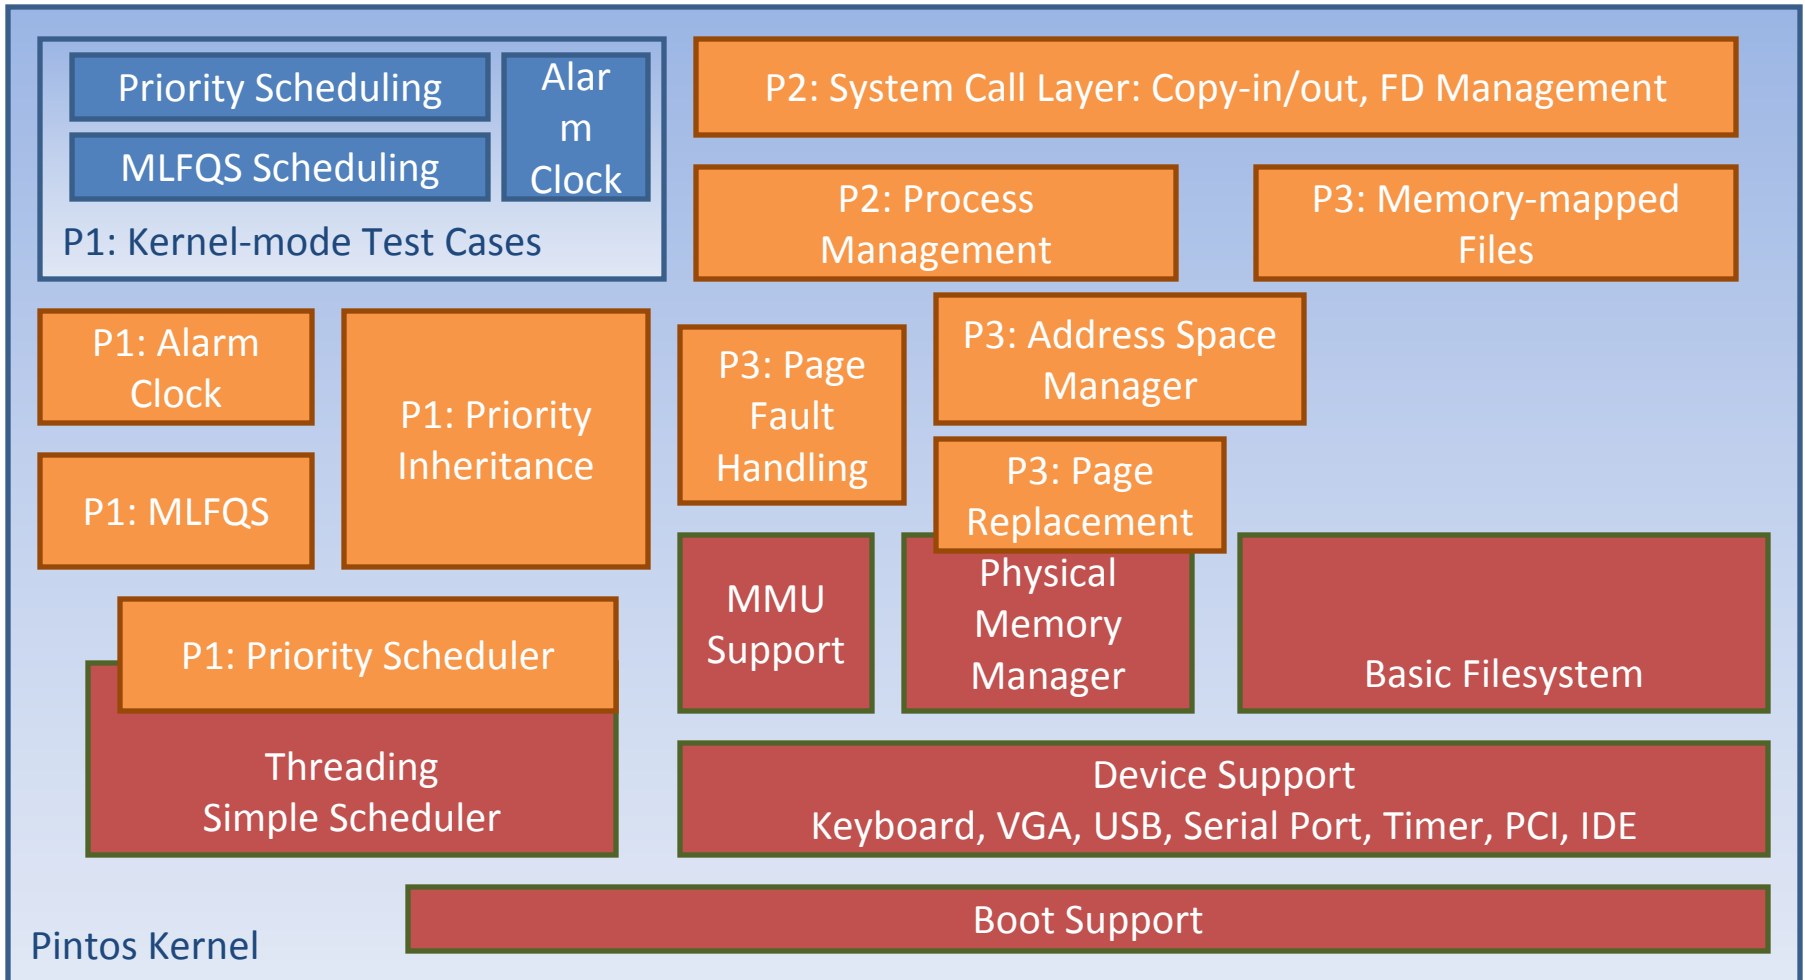
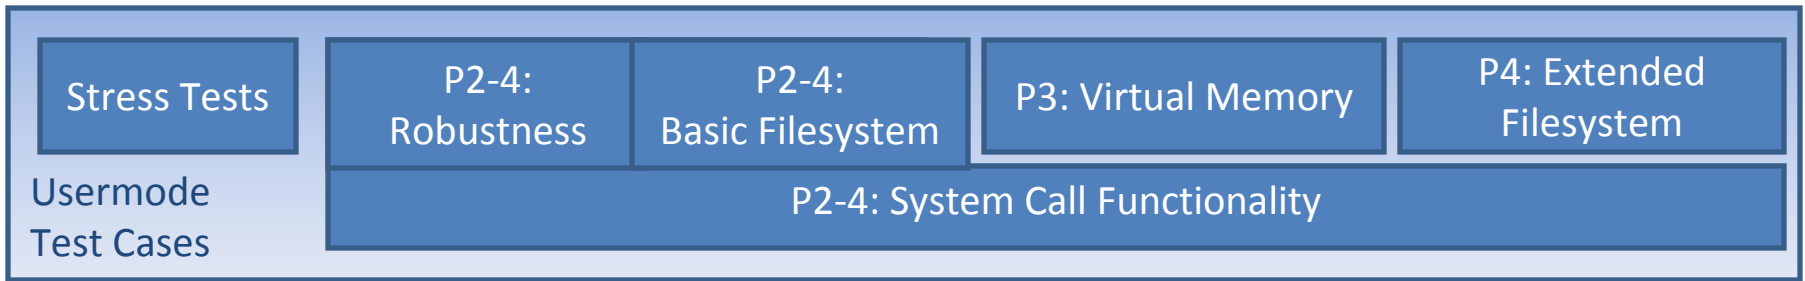
3/7/2009

Support Code

Students Create

Public Tests

Post Project 3



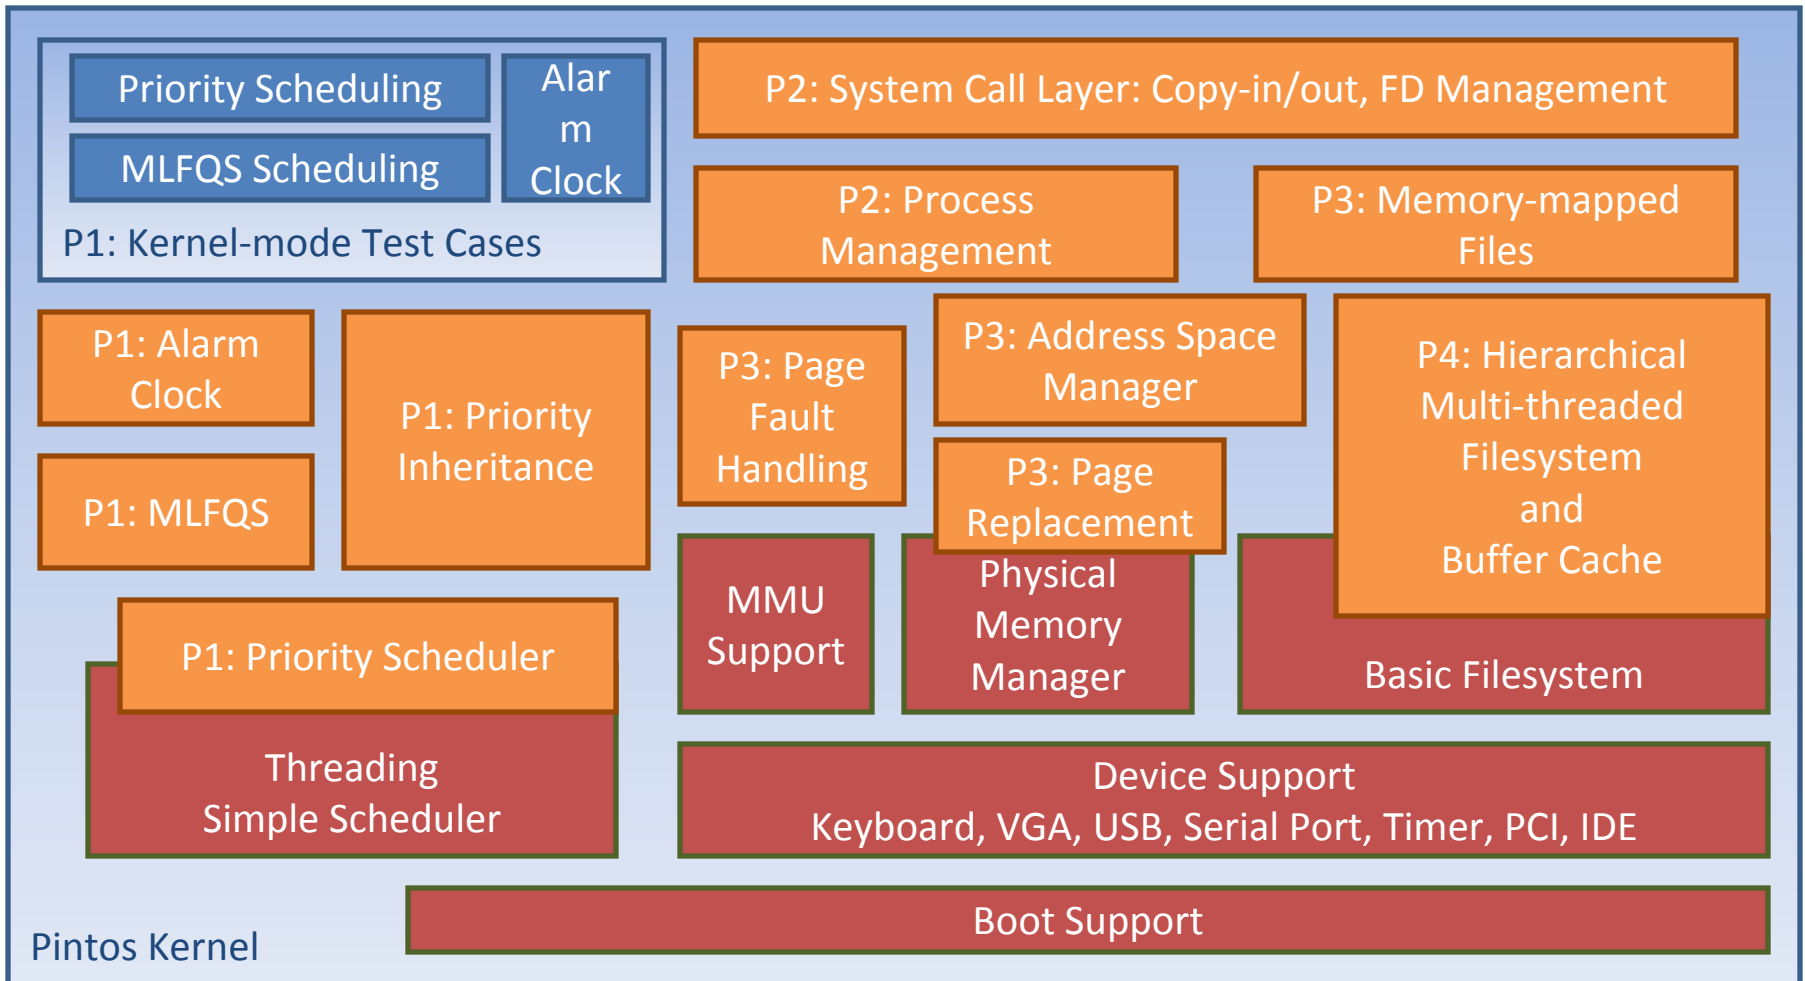

3/7/2009

Support Code

Students Create

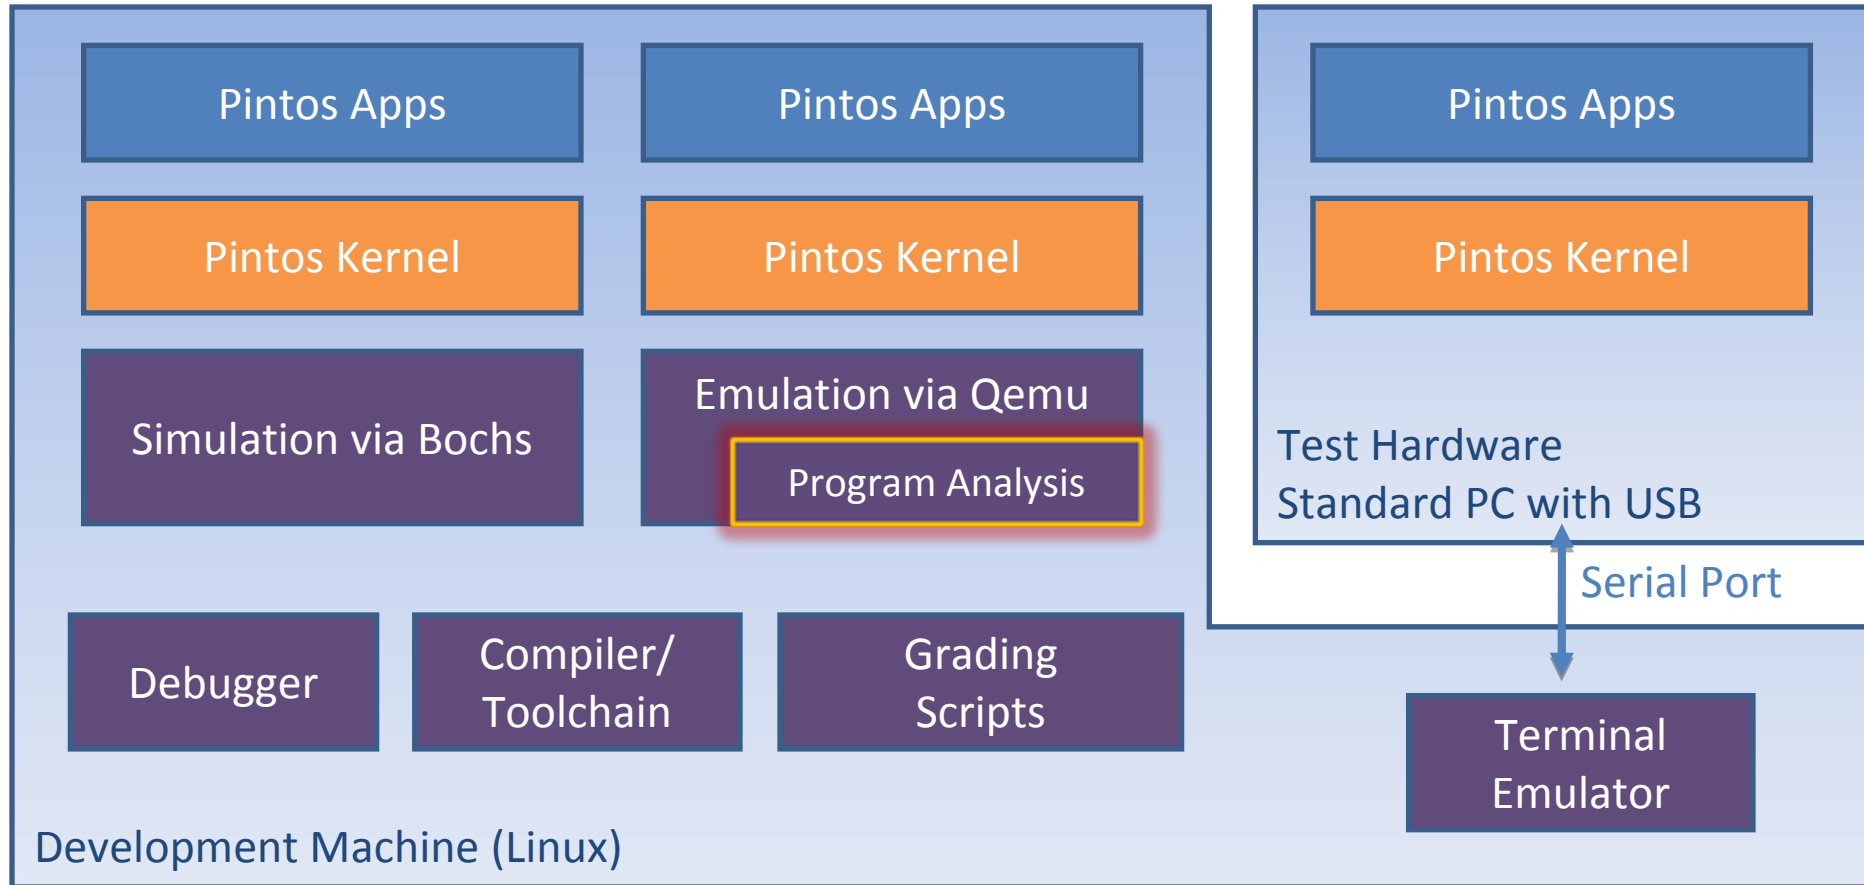
Public Tests

Pre Project 4





# Program Analysis



# Race Detection Example

\*\*\* Race #1 \*\*\*

- Fault Point -

IP: c002da7d

Function: `list_begin`

Memory address at which race occurred: c003afc4

Memory base of object in which race occurred: c003afc0

This race affects global variable: `open_inodes`

Lockset:

- Threads involved in race -

\* Backtrace (thread #1) \*

```
list_remove (c002d565)(lib/kernel/list.c:260)
inode_close (c0032c1f)(filesys/inode.c:177)
file_close (c0032224)(filesys/file.c:52)
syscall_handler (c003175c)(userprog/syscall.c:288)
intr_handler (c0021f47)(threads/interrupt.c:377)
??? (c0022107)(.././threads/intr-stubs.S:38)
```

\* Lockset (thread #1) \*

*In this example, students forgot to protect the list of open inodes, which is accessed concurrently by an exiting process (left backtrace) trying to close its files and a starting process (right backtrace) trying to open and read its executable*


\* Backtrace (thread #2) \*

```
list_begin (c002da7d)(lib/kernel/list.c:74)
inode_open (c0032c83)(filesys/inode.c:118)
dir_open_root (c00327c0)(filesys/directory.c:57)
filesys_open (c0031b27)(filesys/filesys.c:69)
start_process (c002f7fb)(userprog/process.c:358)
kernel_thread (c002170f)(threads/thread.c:538)
```

\* Lockset (thread #2) \*

# Evaluation (Fall 2008)

- How confident are you in your ability to understand the output of the race condition checker?
  1. Not at all confident, the output was very confusing. (2/44)
  2. I sort of understood what it was trying to tell me, but my understanding was vague. (17/44)
  3. After careful analysis of the output, I understood the causes leading to the displayed race and was able to fix it. (16/44)
  4. Once I learned the general format of the output, I quickly found the underlying race condition that was flagged. (4/44)
  5. No answer (5/44)
- Based on survey given during final exam
- In addition, more than 50% of students reported that the race condition checker helped them find actual bugs that made them pass project tests!



# Setting Up Pintos

- Requires simple Linux server
  - 1 quad core machine can support 8-10 students easily
  - All work can be done using remote ssh access, or an IDE can be used
  - No root user access required
  - Uses mostly host tools (gcc, binutils) and packages (bochs, qemu)
- Includes texinfo manual (HTML, 129-page PDF)
  - Documentation separates generic and institution-specific parts in separate files, e.g.
    - Stanford: @set coursenumber CS140
    - Virginia Tech: @set coursenumber CS 3204

# Placement in Curriculum

- Cannot be a first course in C
- Should probably be 4<sup>th</sup> or 5<sup>th</sup> programming course
- Can be a first or second course in OS
- Pintos projects can stretch over 10-15 weeks
- Satisfies a “deep design” requirement

# Related Work

- Systems that provide internal kernel perspective
- Simulated architecture only:
  - Nachos, ToyOS, OS/161, Yalnx
- Emulated:
  - GeekOS, JOS
- Real hardware:
  - GeekOS, Xinu, PortOS, JOS, Minix, Windows CRK, adapted versions of Linux

# Future Work

- Educational:
  - Introduce modular assignment structure to allow instructor to tailor assignments with reduced or varied scope
  - Integrate assessment tools
  - Integrate static analysis tools
  - Integrate performance measures
- Technological:
  - Introduce multi-core/multi-processor support

# Thank You!

- Ben Pfaff
- Anthony Romano
- Godmar Back
- Many Instructors, TA's, and students who have contributed with tests and suggestions
- URL: [www.pintos-os.org](http://www.pintos-os.org)
- Mailing list: [pintos-os@googlegroups.com](mailto:pintos-os@googlegroups.com)



