

**CIS 730 Artificial Intelligence**  
**CIS 530 Principles of Artificial Intelligence**  
Fall 2013

**Homework 8 of 8: Machine Problem (MP8)**  
Reasoning and Learning, Part II:  
Machine Learning Basics

Assigned: Friday 15 Nov 2013  
Due: Fri 22 Nov 2013 (before midnight)

The purpose of this assignment is to help you exercise some basic concepts of machine learning.

This homework assignment is worth a total of 100%.

Include a copy of your solution *with* the final project submission before the due date.

Preliminaries:

- In your web browser, open the URL <http://www.cs.waikato.ac.nz/~ml/weka/>
- Download the *Waikato Environment for Knowledge Analysis (WEKA)* v3.6.9 (GUI version © 2013 I. H. Witten, E. Frank, *et al.*) to your local system (this can be a Windows, Unix, Mac, or other system, but the binaries are precompiled for ix86 Linux. Follow the instructions in the *WEKA3* manual for installing it into your home directory.
- Download the University of California Irvine (UCI) Machine Learning and KDD Repository data sets, available at <http://prdownloads.sourceforge.net/weka/uci-20070111.tar.gz>. (See also [http://www.cs.waikato.ac.nz/~ml/weka/index\\_datasets.html](http://www.cs.waikato.ac.nz/~ml/weka/index_datasets.html) for the complete list of data collections curated by the Waikato Machine Learning Group.)
- Create a subdirectory “UCI” of the “data” directory for your installed copy of WEKA and unpack the UCI ML/KDD data there.
- Launch WEKA and start the Explorer application from the Applications menu.
- Click “Open File” on the Preprocess tab. Load each of the specified data sets (in Windows, browse to your “Program Files (x86)/Weka-3-6/data” directory), one at a time. The data sets are:
  - Jergen’s audiology data set from Baylor College of Medicine, donated by Porter (audiology.arff, <http://bit.ly/150PsCc>)
  - Quinlan’s credit approval data set (credit-a.arff, <http://bit.ly/150PnOU>)
  - Nakai’s *E. coli* data set from the Institute of Molecular and Cellular Biology at Osaka University, donated by Horton (ecoli.arff, <http://bit.ly/150PIGS>)
  - The Hayes-Roths’ data set, donated by Aha (hayes-roth\_test.arff and hayes-roth\_train.arff, <http://bit.ly/150Piyl>)
  - Schlimmer’s U.S. Congressional voting record classification data set (vote.arff, <http://bit.ly/150PxWK>)
- Apply the specified *inducers* (inductive learning algorithms).
- Refer to the new Weka wiki (<http://weka.wikispaces.com>) for documentation.

1. (40% for CIS 530, 30% for CIS 730) **Supervised classification learning: running experiments and reporting results.** Create a table with one row per inducer and three columns per data set. The table should look like this:

Inducer	Data Set 1			Data Set 2		
	Accuracy	Precision	Recall	Accuracy	Precision	Recall
1						
2						
3						

After loading a data set using Open File in the Preprocess tab, click Choose in the Classify tab to select an inducer. If an inducer is not applicable to a particular data set, it will be shown in red. Indicate "N/A" for all results in such cases.

Report the mean accuracy, precision, and recall for the following inducers on each of the five data sets above.

- trees → J48 (Decision Trees)
- bayes → Naïve Bayes (Simple Bayes)
- functions → RBFNetwork (Radial Basis Function network)
- lazy → IBk (k-Nearest Neighbor)

2. (40% for CIS 530, 30% for CIS 730) **Interpreting results from classification learning.**
- Accuracy, precision, and recall.** For each data set, which inducer achieved the highest accuracy? Did it also achieve the highest precision and recall? If not, what can you conclude from the performance measures?
  - Evaluating inducers.** Is one inducer giving you consistently superior results across all data sets? Why or why not?
3. **Experiment Replication Script (20% for CIS 530, 40% for CIS 730).** Write a script in any scripting language (e.g., Unix shell script, Tcl/Tk, Perl, Ruby, Python, Lua, or VB.net) to automate the collection of experimental results for your final project report. Alternatively, write a data generator or preprocessor using a scripting language or a high-level imperative language such as C++ or Java. Examples include:
- A Perl script to run *Angband* using different levels of training (the basic APWborg, trained with 1000 turns of combat data, trained with 2000, 4000, 8000, 16000)
  - A Python driver for training a *TAC-SCM* agent and then bringing them into new games on a running server
  - A Java program to generate test data for benchmarking the import module you write to map data from a particular protein database format into the unified protein interaction format. Measure the ontology reasoner and I/O costs separately.

Turn in a standalone table of results produced by this program along with your source code.

**Class participation (required).** Think about the following question and post your answer to the class participation thread for MP8 in the KSOL message board by Fri 06 Dec 2013.

Collect a *baseline* result and compare it to your system's. Compute *percent decrease (or increase) versus baseline* for a chosen metric. This can be test set accuracy, test set precision/recall/*F*-score (aka *F1* score or *f*-measure), or area under the ROC curve.

Report the baseline and new values, and the difference (incremental gain or loss due to your improvement) over one or more experimental *replications* – i.e., repetitions of the experiment under the same task specification. List results in your final project report.

**References:**

[http://en.wikipedia.org/wiki/Precision\\_\(information\\_retrieval\)](http://en.wikipedia.org/wiki/Precision_(information_retrieval))

[http://en.wikipedia.org/wiki/Recall\\_\(information\\_retrieval\)](http://en.wikipedia.org/wiki/Recall_(information_retrieval))

[http://en.wikipedia.org/wiki/F1\\_Score](http://en.wikipedia.org/wiki/F1_Score)

[http://en.wikipedia.org/wiki/Sensitivity\\_and\\_specificity](http://en.wikipedia.org/wiki/Sensitivity_and_specificity)

[http://en.wikipedia.org/wiki/ROC\\_Curve](http://en.wikipedia.org/wiki/ROC_Curve)