

LECTURE 32 OF 42

Machine Learning: Basic Concepts Discussion: Inductive Bias, Decision Trees

William H. Hsu

Department of Computing and Information Sciences, KSU

KSOL course page: http://snipurl.com/v9v3
Course web site: http://www.kddresearch.org/Courses/CIS730
Instructor home page: http://www.cis.ksu.edu/~bhsu

Reading for Next Class:

Section 18.1 – 18.2, Russell and Norvig



CIS 530 / 730

LECTURE 32 OF 42



LECTURE OUTLINE

- Learning Algorithms and Models
 - * Models: decision trees, winnow, artificial neural networks, naïve Bayes, genetic algorithms (GAs) and genetic programming (GP), instance-based learning (nearest-neighbor), inductive logic programming (ILP)
 - * Algorithms: for decision trees (ID3/C4.5/J48), ANNs (backprop), etc.
 - * Methodologies: supervised, unsupervised, reinforcement; knowledge-guided
- Theory of Learning
 - * Computational learning theory (COLT): complexity, limitations of learning
 - * Probably Approximately Correct (PAC) learning
 - * Probabilistic, statistical, information theoretic results
- Multistrategy Learning: Combining Techniques, Knowledge Sources
- Data: Time Series, Very Large Databases (VLDB), Text Corpora
- Applications
 - * Performance element: classification, decision support, planning, control
 - * Database mining and knowledge discovery in databases (KDD)
 - * Computer inference: learning to reason





WHY MACHINE LEARNING?

- New Computational Capability
 - * Database mining: converting records into knowledge
 - * Self-customizing programs: learning news filters, adaptive monitors
 - * Learning to act: robot planning, control optimization, decision support
 - * Applications that are hard to program: automated driving, speech recognition
- Better Understanding of Human Learning and Teaching
 - * Cognitive science: theories of knowledge acquisition (e.g., through practice)
 - * Performance elements: reasoning (inference) and recommender systems
- Time is Right
 - * Recent progress in algorithms and theory
 - * Rapidly growing volume of online data from various sources
 - * Available computational power
 - * Growth, interest in learning-based industries (e.g., data mining/KDD)

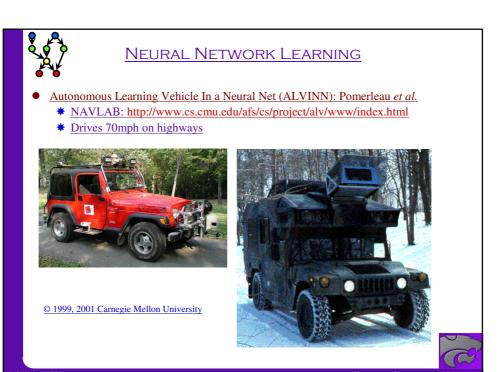


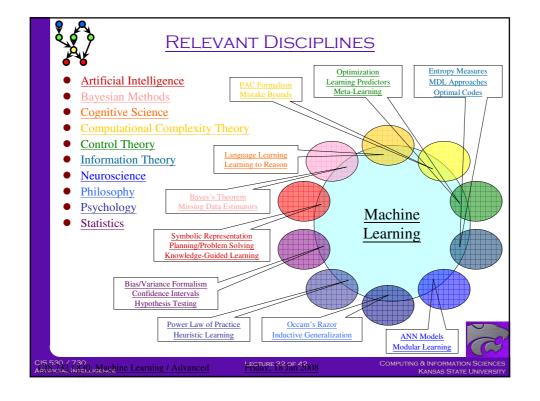
CIS 530 / 730 ARTIFICIAL INTELLIGENCE LECTURE 32 OF 42

RULE AND DECISION TREE LEARNING

- Example: Rule Acquisition from Historical Data
- Data
 - * Patient 103 (time = 1): Age 23, First-Pregnancy: no, Anemia: no, Diabetes: no, Previous-Premature-Birth: no, Ultrasound: unknown, Elective C-Section: unknown, Emergency-C-Section: unknown
 - * Patient 103 (time = 2): Age 23, First-Pregnancy: no, Anemia: no, Diabetes: yes, Previous-Premature-Birth: no, Ultrasound: unknown, Elective C-Section: no, Emergency-C-Section: unknown
 - * Patient 103 (time = n): Age 23, First-Pregnancy: no, Anemia: no, Diabetes: yes, Previous-Premature-Birth: no, Ultrasound: abnormal, Elective C-Section: no, Emergency-C-Section: YES
- Learned Rule
 - * IF no previous vaginal delivery, AND abnormal 2nd trimester ultrasound, AND malpresentation at admission, AND no elective C-Section
 THEN probability of emergency C-Section is 0.6
 - ***** Training set: 26/41 = 0.634
 - ***** Test set: 12/20 = 0.600



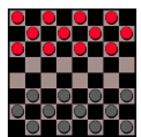






SPECIFYING A LEARNING PROBLEM

- Learning = Improving with Experience at Some Task
 - \star Improve over task T,
 - * with respect to performance measure P,
 - * based on experience E.
- Example: Learning to Play Checkers
 - **★** *T*: play games of checkers
 - * P: percent of games won in tournament play
 - * E: opportunity to play against self
- Refining the Problem Specification: Issues
 - **★** What experience?
 - **★** What *exactly* should be learned?
 - * How shall it be represented?
 - **★** What specific algorithm to learn it?
- Defining the Problem Milieu
 - **☀** Performance element: How shall results of learning be applied?
 - **★** How shall performance element be evaluated? Learning system?





CIS 732 / 730; Machine Learning / Advanced

ECTURE 32 OF 42₀₀

EXAMPLE: LEARNING TO PLAY CHECKERS

- Type of Training Experience
 - **☀** Direct or indirect?
 - **★** Teacher or not?
 - * Knowledge about the game (e.g., openings/endgames)?
- Problem: Is Training Experience *Representative* (of Performance Goal)?
- Software Design
 - * Assumptions of the learning system: *legal* move generator exists
 - * Software requirements: generator, evaluator(s), parametric target function
- Choosing a Target Function
 - **★** *ChooseMove: Board* → *Move* action selection function, or *policy*
 - * $\underline{V: Board \rightarrow R \text{evaluation function for game tree search (minimax / }\alpha \beta)}$
 - **★** Ideal target *V*; approximated target
 - * Goal of learning process: operational description (approximation) of V
- Chinook: Checkers Solved by Game Tree Search (July 2007)
- Reference: http://en.wikipedia.org/wiki/Checkers





A TARGET FUNCTION FOR LEARNING TO PLAY CHECKERS

- Possible Definition
 - ***** If b is final board state that is won, then V(b) = +100 (or MAXINT)
 - ***** If b is final board state that is lost, then V(b) = -100 (or MAXINT)
 - ***** If b is final board state that is drawn, then V(b) = 0
 - * If b is not final board state in the game, then V(b) = V(b') where b' is best final board state that can be achieved starting from b and playing optimally until end
 - * Correct values, but not operational
- Choosing a Representation for the Target Function
 - * Collection of rules?
 - ★ Neural network?
 - * Polynomial function (e.g., linear, quadratic combination) of board features?
 - * Other?
- A Representation for Learned Function

★ bp/rp = number of black/red pieces; bk/rk = number of black/red kings $\hat{\mathbf{W}}(\mathbf{b}) = \mathbf{w}_{b} + \mathbf{w}$



FFGTURE 18 Jan 2008



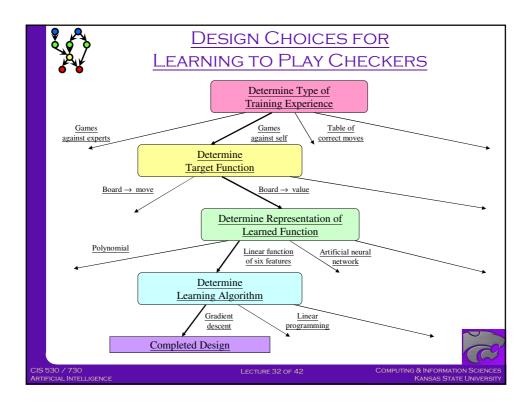
A TRAINING PROCEDURE FOR LEARNING TO PLAY CHECKERS

- **Obtaining Training Examples**
 - * <u>V(b)</u> target function
 - learned function $-\hat{V}(b)$
 - V_{train}(b) training value ("signal")
- Rule For Estimating Training Values
- $V_{train}(\mathbf{b}) \leftarrow \hat{\mathbf{V}}(Successor(\mathbf{b}))$
- Rule for Training (Weight Tuning)
 - ★ Least Mean Square (LMS) weight update rule
 - * REPEAT
 - Select training example b at random
 - Compute the *error*(*b*) for this training example
 - For each board feature f_i , update weight w_i as follows **error** $(\mathbf{b}) = \mathbf{V}_{train}(\mathbf{b}) \mathbf{V}(\mathbf{b})$

where c is small, constant factor to adjust learning rate

 $\mathbf{w}_i \leftarrow \mathbf{w}_i + \mathbf{c} \cdot \mathbf{f}_i \cdot \mathbf{error}(\mathbf{b})$



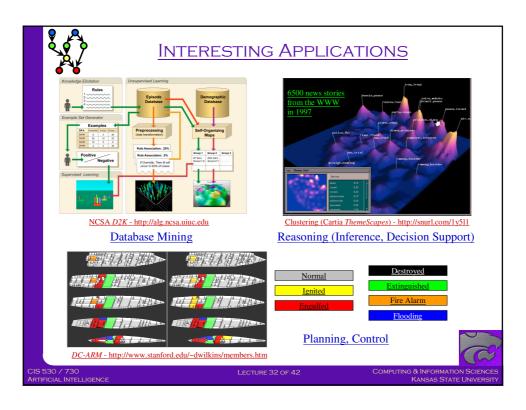




SOME ISSUES IN MACHINE LEARNING

- What Algorithms Can Approximate Functions Well? When?
- How Do Learning System Design Factors Influence Accuracy?
 - * Number of training examples
 - **★** Complexity of *hypothesis representation*
- How Do Learning Problem Characteristics Influence Accuracy?
 - * Noisy data
 - * Multiple data sources
- What Are Theoretical Limits of Learnability?
- How Can Prior Knowledge of Learner Help?
- What Clues Can We Get From Biological Learning Systems?
- How Can Systems Alter Their Own Representation?







WHAT TO LEARN?

- Classification Functions
 - **★** Learning hidden functions: estimating ("fitting") parameters
 - * Concept learning (e.g., chair, face, game)
 - * Diagnosis, prognosis: risk assessment, medical monitoring, security, ERP
- Models
 - **★** Map (for navigation)
 - * Distribution (query answering, aka QA)
 - * Language model (e.g., automaton/grammar)
- Skills
 - * Playing games
 - * Planning
 - * Reasoning (acquiring representation to use in reasoning)
- Cluster Definitions for Pattern Recognition
 - * Shapes of objects
 - * Functional or taxonomic definition
- Many Problems Can Be Reduced to Classification





HOW TO LEARN IT?

- Supervised
 - **★** What is learned? Classification function; other models
 - ***** Inputs and outputs? Learning: examples $\langle x, f(x) \rangle \rightarrow$ approximation $\hat{f}(x)$
 - **★** How is it learned? Presentation of examples to learner (by teacher)
- Unsupervised
 - **★** Cluster definition, or *vector quantization* function (*codebook*)
 - ***** Learning: observations $x \times$ distance metric $d(x_1, x_2) \rightarrow$ discrete codebook f(x)
 - * Formation, segmentation, labeling of clusters based on observations, metric
- Reinforcement
 - * Control policy (function from states of the world to actions)
 - ***** Learning: state/reward sequence $\{\langle s_i, r_i \rangle : 1 \le i \le n\} \rightarrow \text{policy } p : s \rightarrow a$
 - **★** (Delayed) feedback of reward values to agent based on actions
 - **★** Model updated based on reward, (partially) observable state



CIS 530 / 730 Artificial Intelligence ECTURE 32 OF 42

ARTIFICIAL INTELLIGENCE

SUPERVISED INDUCTIVE LEARNING: CLASSIFICATION AND REGRESSION

- Given: Training Examples $\langle x, f(x) \rangle$ of Some Unknown Function f
- Find: A Good Approximation to f
- Examples (besides Concept Learning)
 - Disease diagnosis
 - x = properties of patient (medical history, symptoms, lab tests)
 - f = disease (or recommended therapy)
 - * Risk assessment
 - x = properties of consumer, policyholder (demographics, accident history)
 - f = risk level (expected cost)
 - * Automatic steering
 - x =bitmap picture of road surface in front of vehicle
 - f = degrees to turn the steering wheel
 - ★ Part-of-speech tagging
 - * Computer security: fraud/intrusion detection, attack graphs
 - * Information extraction: clusters of documents
 - * Social networks and weblogs: predicting links, sentiment analysis
 - **★** Multisensor integration and prediction





LEARNING AND TYPES[1].

A GENERIC SUPERVISED LEARNING PROBLEM



Example	X ₁	\mathbf{X}_2	X ₃	X_4	у
0	0	1	1	0	0
1	0	0	0	0	0
2	0	0	1	1	1
3	1	0	0	1	1
4	0	1	1	0	0
5	1	1	0	0	0
6	0	1	0	1	0

- Input x_i : t_i , desired output y: t, "target" function f: $(t_1 \times t_2 \times t_3 \times t_4) \to t$
- Learning function: Vector $(t_1 \times t_2 \times t_3 \times t_4 \times t) \rightarrow (t_1 \times t_2 \times t_3 \times t_4) \rightarrow t$



CIS 530 / 730 ARTIFICIAL INTELLIGENCE ECTURE 32 OF 42

COMPUTING & INFORMATION SCIENCE



LEARNING AND TYPES [2]: UNRESTRICTED HYPOTHESIS SPACE

- $|A \rightarrow B| = |B|^{|A|}$ proof?
- $|H^4 \rightarrow H| = |\{0,1\} \times \{0,1\} \times \{0,1\} \times \{0,1\} \rightarrow \{0,1\}| = 2^{2^4} = 65536$ functions
- Complete Ignorance: Is Learning Possible?
 - ★ Need to see every possible input/output pair
 - * After 7 examples, still have $2^9 = 512$ possibilities (out of 65536) for f

Example	\mathbf{x}_1	\mathbf{X}_{2}	X ₃	\mathbf{X}_4	y
0	0	0	0	0	?
1	0	0	0	1	?
2	0	0	1	0	0
3	0	0	1	1	1
4	0	1	0	0	0
5	0	1	0	1	0
6	0	1	1	0	0
7	0	1	1	1	?
8	1	0	0	0	?
9	1	0	0	1	1
10	1	0	1	0	?
11	1	0	1	1	?
12	1	1	0	0	0
13	1	1	0 1		?
14	1	1	1	0	?
15	1	1	1	1	?



CIS 530 / 730 ARTIFICIAL INTELLIGENCE LECTURE 32 OF 42

COMPUTING & INFORMATION SCIENCES

KANSAS STATE LINIVERSITY



TRAINING EXAMPLES FOR CONCEPT ENJOYSPORT

- Specification for Examples
 - * Similar to a data type definition
 - * 6 attributes: Sky, Temp, Humidity, Wind, Water, Forecast
 - * Nominal-valued (symbolic) attributes enumerative data type
- Binary (Boolean-Valued or H Valued) Concept
- Supervised Learning Problem: Describe the General Concept

Example	Sky	Air Temp	Humidity	Wind	Water	Forecast	Enjoy Sport
0	Sunny	Warm	Normal	Strong	Warm	Same	Yes
1	Sunny	Warm	High	Strong	Warm	Same	Yes
2	Rainy	Cold	High	Strong	Warm	Change	No
3	Sunny	Warm	High	Strong	Cool	Change	Yes



CIS 530 / 730

ECTURE 32 OF 42

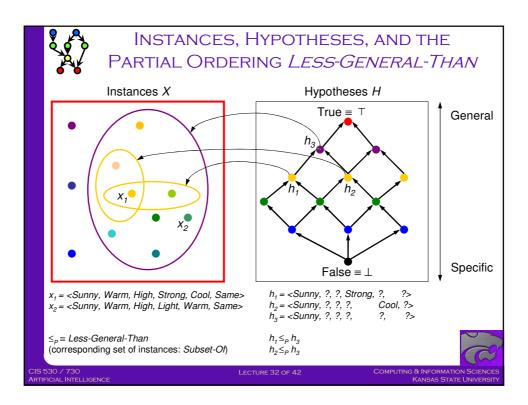
COMPUTING & INFORMATION SCIENCE



REPRESENTING HYPOTHESES

- Many Possible Representations
- Hypothesis *h*: Conjunction of Constraints on Attributes
- Constraint Values
 - * Specific value (e.g., Water = Warm)
 - **★** Don't care (e.g., "Water = ?")
 - **★** No value allowed (e.g., "Water = Ø")
- Example Hypothesis for *EnjoySport*
 - * Sky AirTemp Humidity Wind Water Forecast <Sunny? ? Strong ? Same>
 - * <u>Is this consistent with the training examples?</u>
 - **★** What are some hypotheses that are consistent with the examples?







FIND-SALGORITHM

1. Initialize h to the most specific hypothesis in H

H: the hypothesis space

(partially ordered set under relation Less-Specific-Than)

2. For each positive training instance x

For each attribute constraint a_i in h

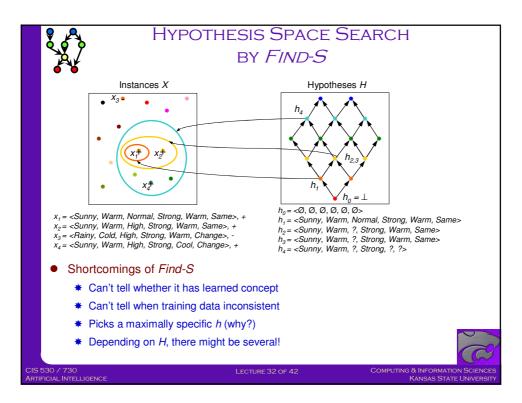
IF constraint a_i in h is satisfied by x

THEN do nothing

ELSE replace a_i in h by next more general constraint satisfied by x

3. Output hypothesis h







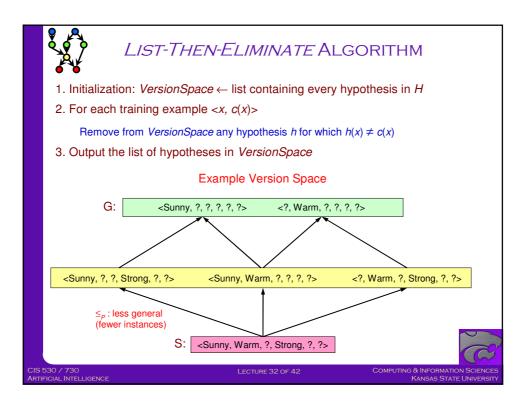
VERSION SPACES

- Definition: Consistent Hypotheses
 - * A hypothesis h is <u>consistent</u> with a set of training examples D of target concept c if and only if h(x) = c(x) for each training example < x, c(x) > in D.
 - * Consistent $(h, D) \equiv \forall \langle x, c(x) \rangle \in D \cdot h(x) = c(x)$
- Given
 - * Hypothesis space H
 - **☀** Data set *D*: set of training examples
- Definition
 - ★ Version space VS_{H,D} with respect to H, D
 - **☀** Subset of hypotheses from *H* consistent with all training examples in *D*
 - * $VS_{H,D} \equiv \{ h \in H \mid Consistent(h, D) \}$



CIS 530 / 730 Artificial Intelligence LECTURE 32 OF 42

COMPUTING & INFORMATION SCIENCES KANSAS STATE UNIVERSITY





HYPOTHESIS SPACES AS LATTICES

- Meet Semilattice
 - * Every pair of hypotheses h_i and h_i has greatest lower bound (GLB) $h_i \vee h_i$
 - * Is H meet semilattice?
 - ***** ⊥≡ some ø
- Join Semilattice
 - ***** Every pair of hypotheses h_i and h_j has *least upper bound* (LUB) $h_i \wedge h_j$ Is H join semilattice?
 - * ^T≡ all ?
- (Full) Lattice
 - * Every pair of hypotheses has GLB $h_i \vee h_i$ and LUB $h_i \wedge h_i$
 - * Both meet semilattice and join semilattice
 - * Partial ordering Less-General-Than





REPRESENTING VERSION SPACES AS LATTICES

- Definition: General (Upper) Boundary
 - **★** General boundary *G* of version space *VS_{H,D}*: set of most general members
 - * Most general \equiv maximal elements of $VS_{H,D}$ \equiv "set of necessary conditions"
- Definition: Specific (Lower) Boundary
 - * Specific boundary S of version space VS_{H,D}: set of least general members
 - * Most specific \equiv minimal elements of $VS_{H,D}$ \equiv "set of sufficient conditions"
- Version Space
 - * $VS_{H,D}$ = consistent poset (partially-ordered subset of H)
 - **★** Every member of version space lies between S and G
 - * $VS_{H,D} \equiv \{ h \in H \mid \exists s \in S, g \in G : s \leq_P h \leq_P g \}, \leq_P \equiv \text{Less-General-Than} \}$

"Version space is defined as set of hypotheses sandwiched between specific *s* and general *g* (given data)"

CIS 530 / 730 ARTIFICIAL INTELLIGENCE LECTURE 32 OF 42

COMPUTING & INFORMATION SCIENCES
KANSAS STATE LINIVERSITY



CANDIDATE ELIMINATION ALGORITHM [1]

1. Initialization

 $G_0 \leftarrow \top \equiv \text{most general hypothesis in } H, \text{ denoted } \{<?, \dots, ?>\}$

 $S_0 \leftarrow \bot \equiv \text{least general hypotheses in } H, \text{ denoted } \{<\emptyset, \dots, \emptyset>\}$

2. For each training example d

If *d* is a positive example (*Update-S*) // generalize

Remove from G any hypotheses inconsistent with d

For each hypothesis *s* in *S* that is not consistent with *d*

Remove s from S // "move S upwards"

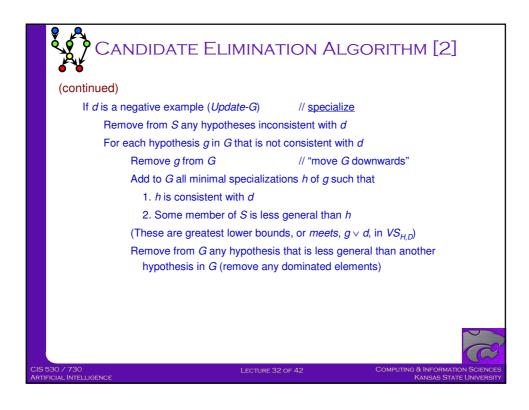
Add to S all minimal generalizations h of s such that

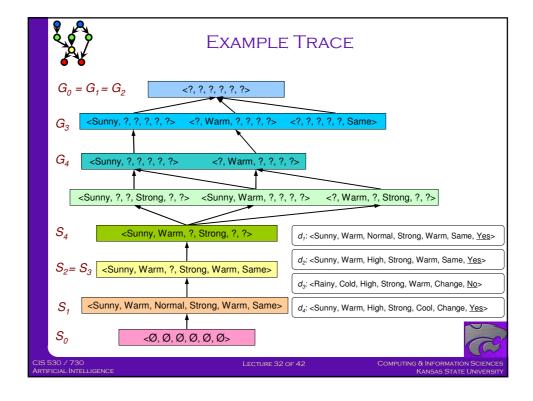
- 1. h is consistent with d
- 2. Some member of G is more general than h

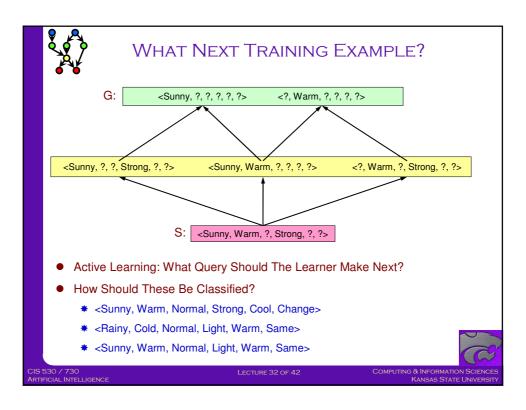
(These are least upper bounds, or joins, $s \land d$, in $VS_{H,D}$)

Remove from *S* any hypothesis that is more general than another hypothesis in *S* (remove any dominating elements)











WHAT JUSTIFIES THIS INDUCTIVE LEAP?

- Example: Inductive Generalization
 - * Positive example: <Sunny, Warm, Normal, Strong, Cool, Change, Yes>
 - * Positive example: <Sunny, Warm, Normal, Light, Warm, Same, Yes>
 - * Induced S: <Sunny, Warm, Normal, ?, ?, ?>
- Why Believe We Can Classify The Unseen?
 - * e.g., <Sunny, Warm, Normal, Strong, Warm, Same>
 - * When is there enough information (in a new case) to make a prediction?





AN UNBIASED LEARNER

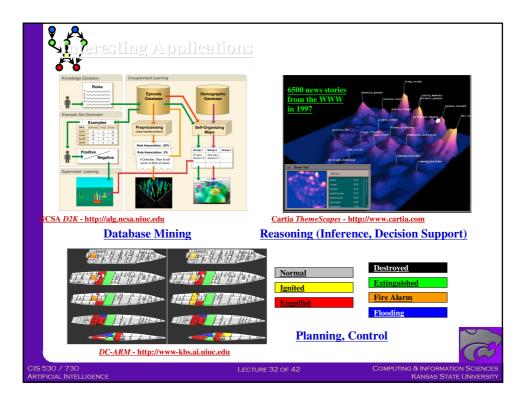
- **Inductive Bias**
 - Any preference for one hypothesis over another, besides consistency
 - Example: $H \equiv conjunctive$ concepts with don't cares
 - What concepts can *H* not express? (Hint: what are its syntactic limitations?)
- Idea
 - Choose <u>unbiased</u> H': expresses every teachable concept (i.e., power set of X)
 - Recall: $|A \rightarrow B| = |B|^{|A|} (A = X; B = \{labels\}; H = A \rightarrow B)$
 - {{Rainy, Sunny, Cloudy} \times {Warm, Cold} \times {Normal, High} \times {None-Mild, Strong} \times $\{Cool,\,Warm\}\times\{Same,\,Change\}\}\rightarrow\{0,\,1\}$
- An Exhaustive Hypothesis Language
 - Consider: H' = disjunctions (\vee), conjunctions (\wedge), negations (\neg) over H
 - $|H'| = 2^{(2 \cdot 2 \cdot 2 \cdot 3 \cdot 2 \cdot 2)} = 2^{96}; |H| = 1 + (3 \cdot 3 \cdot 3 \cdot 4 \cdot 3 \cdot 3) = 973$
- What Are S, G For The Hypothesis Language H??
 - S ← disjunction of all positive examples
 - $G \leftarrow conjunction of all negated negative examples$





- **Example: Inductive Generalization**
 - Positive example: <Sunny, Warm, Normal, Strong, Cool, Change, Yes>
 - Positive example: <Sunny, Warm, Normal, Light, Warm, Same, Yes>
 - Induced S: <Sunny, Warm, Normal, ?, ?, ?>
- Why Believe We Can Classify The Unseen?
 - e.g., <Sunny, Warm, Normal, Strong, Warm, Same>
 - When is there enough information (in a new case) to make a prediction?







Example of A Biased H

- Conjunctive concepts with don't cares
- What concepts can H not express? (Hint: what are its syntactic limitations?)

Idea

- Choose H' that expresses every teachable concept
- i.e., H' is the power set of X
- Recall: $| \mathbf{A} \rightarrow \mathbf{B} | = | \mathbf{B} |^{|\mathbf{A}|} (\mathbf{A} = X; \mathbf{B} = \{labels\}; H' = \mathbf{A} \rightarrow \mathbf{B})$
- $\quad \{\{Rainy, Sunny\} \times \{Warm, Cold\} \times \{Normal, High\} \times \{None, Mild, Strong\} \times \{Cool, \\ \underline{Warm} \times \{Same, Change\}\} \rightarrow \{0, 1\}$

An Exhaustive Hypothesis Language

- Consider: H' = disjunctions (\vee), conjunctions (\wedge), negations (\neg) over previous H
- $|H'| = 2^{(2 \cdot 2 \cdot 2 \cdot 3 \cdot 2 \cdot 2)} = 2^{96}; |H| = 1 + (3 \cdot 3 \cdot 3 \cdot 4 \cdot 3 \cdot 3) = 973$

What Are S, G For The Hypothesis Language H'?

- $\underline{S} \leftarrow disjunction \ of \ all \ positive \ examples$
- $G \leftarrow conjunction \ of \ all \ negated \ negative \ examples$





Components of An Inductive Bias Definition

- Concept learning algorithm L
- Instances X, target concept c
- Training examples $D_c = \{ \langle x, c(x) \rangle \}$
- $L(x_i, D_c)$ = classification assigned to instance x_i by L after training on D_c

Definition

- The inductive bias of L is any minimal set of assertions B such that, for any target concept c and corresponding training examples D_c ,

$$\forall x_i \in X . [(B \land D_c \land x_i) \vdash L(x_i, D_c)]$$
 where $A \vdash$

B means A logically entails B

- Informal idea: preference for (i.e., restriction to) certain hypotheses by structural (syntactic) means

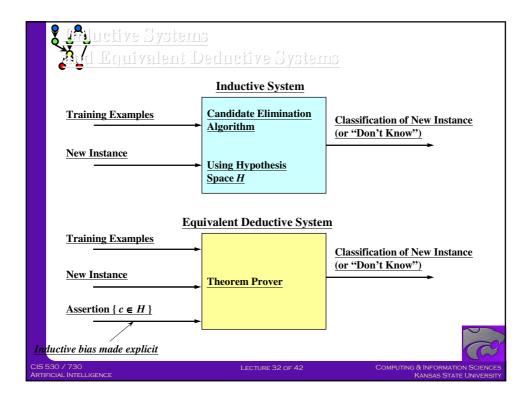
Rationale

- Prior assumptions regarding target concept
- Basis for inductive generalization

ION SCIENCES

CIS 530 / 730 Artificial Intelligence LECTURE 32 OF 42

COMPUTING & INFORMATION SCIENCES
KANSAS STATE UNIVERSITY





<u>ee Learners with Different Biases</u>

Rote Learner

- Weakest bias: anything seen before, i.e., no bias
- Store examples
- Classify x if and only if it matches previously observed example

Version Space Candidate Elimination Algorithm

- Stronger bias: concepts belonging to conjunctive H
- Store extremal generalizations and specializations
- Classify *x if and only if* it "falls within" *S* and *G* boundaries (all members agree)

Find-S

- Even stronger bias: most specific hypothesis
- Prior assumption: any instance not observed to be positive is negative
- Classify x based on S set



CIS 530 / 730 ARTIFICIAL INTELLIGENCE LECTURE 32 OF 42

Views of Learning

Removal of (Remaining) Uncertainty

- Suppose unknown function was known to be m-of-n Boolean function
- Could use training data to infer the function

Learning and Hypothesis Languages

- Possible approach to guess a good, small hypothesis language:
 - Start with a very small language
 - Enlarge until it contains a hypothesis that fits the data
- Inductive bias
 - Preference for certain languages
 - Analogous to data compression (removal of redundancy)
 - Later: coding the "model" versus coding the "uncertainty" (error)

We Could Be Wrong!

- Prior knowledge could be wrong (e.g., $y = x_4 \land \text{one-of } (x_1, x_3)$ also consistent)
- If guessed language was wrong, errors will occur on new cases





Strategies for Machine Learning

- Develop Ways to Express Prior Knowledge
 - Role of prior knowledge: guides search for hypotheses / hypothesis languages
 - Expression languages for prior knowledge
 - Rule grammars; stochastic models; etc.
 - Restrictions on computational models; other (formal) specification methods
- Develop Flexible Hypothesis Spaces
 - Structured collections of hypotheses
 - Agglomeration: nested collections (hierarchies)
 - Partitioning: decision trees, lists, rules
 - Neural networks; cases, etc.
 - Hypothesis spaces of adaptive size
- Either Case: Develop Algorithms for Finding A Hypothesis That Fits Well
 - Ideally, will generalize well
- Later: Bias Optimization (Meta-Learning, Wrappers)



CIS 530 / 730 Artificial Intelligence ECTURE 32 OF 42

Inputational Learning Theory

- What General Laws Constrain Inductive Learning?
- What Learning Problems Can Be Solved?
- When Can We Trust The Output of A Learning Algorithm?
- We Seek Theory To Relate:
 - Probability of successful learning
 - Number of training examples
 - Complexity of hypothesis space
 - Accuracy to which target concept is approximated
 - Manner in which training examples are presented





SUMMARY

- Reading: Chapters 1-2, Mitchell
- Suggested Exercises: 2.2, 2.3, 2.4, 2.6
- Taxonomy of Learning Systems
- Today: Overview, Learning from Examples
 - * (Supervised) concept learning framework
 - * Simple approach: assumes no noise; illustrates key concepts
- Today & Next Class: Hypothesis Learning and Inductive Bias
 - **★** Sources: Mitchell (1997) online notes and handout
 - **★** Wednesday: inductive learning, version space
 - * Friday: candidate elimination algorithm, active learning, inductive bias
 - **★** Background concepts: partially-ordered set (poset) formalism
- Next Class: Decision Trees, Intro Computational Learning Theory (COLT)
 - **★** First paper review due Wed 30 Jan 2008
 - **★** Sources: Kearns & Vazirani (1994), Han & Kamber 2nd edition (2006)



CIS 530 / 730 Artificial Intelligence LECTURE 32 OF 42



TERMINOLOGY

- Learning: Improving at Task given Performance Measure, Experience
- Performance Element: Part of System that Applies Result of Learning
- Types of Learning
 - **★** Supervised: with "teacher" (often, classification from labeled examples)
 - **★** <u>Unsupervised: from data, using similarity measure (unlabeled instances)</u>
 - * Reinforcement: "by doing", with reward/penalty signal
- Supervised Learning: Target Functions
 - * Target function function c or f to be learned
 - **★** Target desired value y to be predicted (sometimes "target function")
 - * Example / labeled instance tuples of the form $\langle x, f(x) \rangle$
 - * Classification function, classifier nominal-valued f (enumerated return type)
- Clustering: Application of Unsupervised Learning
- Concepts and Hypotheses
 - * Concept function c from observations to TRUE or FALSE (membership)
 - * Class label output of classification function
 - * Hypothesis proposed function h believed to be similar to c (or f)

