

**Homework Assignment 6 [15 points] – due October 25 (in class)**

1. [5 points] After a system's crash, the undo-log using nonquiescent checkpointing contains the following data.

<START T1>
<T1, X1, 1>
<START CKPT ????>
<START T2>
<T2, X2, 2>
<T1, X1, 3>
<START T3>
<COMMIT T1>
<END CKPT>
<START CKPT ????>
<T2, X2, 4>
<T3, X3, 5>
<START T4>
<COMMIT T2>
<T4, X4, 6>
<COMMIT T3>
<END CKPT>
<START T5>
<T5, X5, 7>
<START CKPT ????>
<T4, X4, 8>
CRASH !!!

- i. What are the correct values of the three <START CKPT ????> records? You have to provide three correct values for the three “????”s.
  
- ii. Assuming that the three <START CKPT ????> records are correctly stored in the log, according to your answer at i., show which elements are recovered by the undo recovery manager and compute their values after recovery.
  
- iii. Indicate what fragment of the log the recovery manager needs to read.

2. [5 points] After a system's crash, the redo-log using nonquiescent checkpointing contains the following data.

<START T1>
<T1, A, 10>
<START T2>
<T2, B, 5>
<T1, C, 7>
<START T3>
<T3, D, 12>
<COMMIT T1>
<START CKPT ????>
<START T4>
<T2, E, 5>
<COMMIT T2>
<T3, F, 1>
<T4, G, 15>
<END CKPT>
<COMMIT T3>
<START T5>
<T5, H, 3>
<START CKPT ????>
<COMMIT T5>
CRASH !!!

- i. What are the correct values of the two <START CKPT ????> records? You have to provide two correct values for the two ????.
- ii. Indicate and explain what fragment of the log the recovery manager needs to read.

- iii. Assuming that the two < START CKPT ??? > records are correctly stored in the log, according to your answer above, show which elements are recovered by the redo recovery manager and compute their values after recovery.

3. [5 points] The SuperSQL database system stores its undo log file in a table, with the following schema:

**Log(N, T, E, V)**

where **N** is the entry number (0, 1, 2, 3, ...), **T** is the transaction id, **E** is the element id, and **V** is the old value. A log entry of the form <T, E, V> is simply represented by the tuple (N, T, E, V), where N is the entry number and E>0 for an ordinary element id. The log entries <START T>, <COMMIT T>, and <ABORT T> are represented by a tuple (N, T, E, null), where E=-1 for START, E=-2 for COMMIT, and E=-3 for ABORT. For example, the log:

<START T1>
<T1, X1, 55>
<START T2>
<T2, X2, 99>
<COMMIT T1>
. . . .

is represented by the table:

<b>N</b>	<b>T</b>	<b>E</b>	<b>V</b>
0	T1	-1	
1	T1	X1	55
2	T2	-1	
3	T2	X2	99
4	T1	-2	
. . .			

Recall that each transaction starts and ends at most once; for example, a sequence `<START T> ... <COMMIT T> ... <START T> ...` will not occur in the log. Moreover, any update by the transaction T will occur between its `<START T>` and `<COMMIT T>`, or between `<START T>` and `<ABORT T>` respectively. Finally, once a transaction has ended in COMMIT or ABORT and the corresponding log record is on disk, the transaction has completed and does not need to be undone.

Write a SQL query that can be run during database recovery, after a system crash. The answer to your query should return a table with two attributes, **E**, and **V**, indicating which elements have to be updated with what values. You should include each element **E** at most once in your answer: otherwise it would be ambiguous how to update it.