

Homework Assignment 7 [20 points] – due November 1st (by midnight)

1. [3 points] Consider a database with objects X and Y and assume that there are two transactions T1 and T2. Transaction T1 reads objects X and Y and then writes X. Transaction T2 reads objects X and Y and then writes objects X and Y. Give three examples of schedules for the transactions T1, T2 such that:

- i. Your first schedule should contain a write-read conflict.

The following schedule results in a write-read conflict:

`r2(X), r2(Y), w2(X), r1(X) ...`

`r1(X)` is a dirty read here.

- ii. Your second schedule should contain a write-write conflict.

The following schedule results in a write-write conflict:

`r2(X), r2(Y), r1(X), r1(Y), w1(X), w2(X) ...`

Now, T2 has overwritten uncommitted data.

- iii. Your third schedule should contain a read-write conflict.

The following schedule results in a read-write conflict:

`r2(X), r2(Y), r1(X), r1(Y), w1(X), r2(X) ...`

Now, T2 will get an unrepeatable read on X.

In each case your schedule may contain additional conflicts, but should contain at least one conflict of the type indicated. (In particular you may give a single schedule, which illustrates all three conflicts!) In each case, indicate the conflict of the type you are illustrating.

2. [5 points] For the following schedules:


(a) $r_2(A); r_1(C); r_2(B); w_2(B); r_3(B); r_1(A); r_3(C); w_3(C); w_1(A)$

(b) $r_2(A); r_1(C); r_2(B); r_3(B); w_2(B); r_1(A); r_3(C); w_3(C); w_1(A)$


Answer the following questions:

- What is the precedence graph for the schedule?
- Is the schedule conflict-serializable? If so, what are the equivalent serial schedules?

Schedule (a) is conflict-serializable because the precedence graph has no cycles. There is an arrow $T_2 \rightarrow T_1$ because of the conflict $R_2(A) \dots W_1(A)$. There is an arrow $T_2 \rightarrow T_3$ because of the conflict $W_2(B) \dots R_3(B)$. There is an arrow $T_1 \rightarrow T_3$ because of the conflict $R_1(C) \dots W_3(C)$. The only possible conflict-equivalent serial schedule is (T_2, T_1, T_3) .



Schedule (b) is not conflict-serializable because the precedence graph has a cycle. There is an arrow $T_2 \rightarrow T_1$ because of the conflict $R_2(A) \dots W_1(A)$. There is an arrow $T_3 \rightarrow T_2$ because of the conflict $R_3(B) \dots W_2(B)$. There is an arrow $T_1 \rightarrow T_3$ because of the conflict $R_1(C) \dots W_3(C)$.




3. [6 points]

- i. Consider the following two transactions and schedule (time goes from top to bottom). Is this schedule conflict-serializable? Explain why or why not.

T0	T1
R ₀ (A)	
W ₀ (A)	
	R ₁ (A)
	R ₁ (B)
	C ₁
R ₀ (B)	
W ₀ (B)	
C ₀	

The schedule is not conflict serializable because the precedence graph contains a cycle. The graph has an edge $T0 \rightarrow T1$ because the schedule contains $w_0[A] \rightarrow r_1[A]$. The graph has an edge $T1 \rightarrow T0$ because the schedule contains $r_1[B] \rightarrow w_0[B]$.



- ii. Show how 2PL can ensure a conflict-serializable schedule for the transactions above. Use the notation $L_i(A)$ to indicate that transaction i acquires the lock on element A and $U_i(A)$ to indicate that transaction i releases its lock on A .

Multiple solutions are possible. Two examples are shown below:

T0	T1
L ₀ (A)	
L ₀ (B)	
R ₀ (A)	
W ₀ (A)	
U ₀ (A)	
	L ₁ (A)
	R ₁ (A)
	L ₁ (B); DENIED...
R ₀ (B)	
W ₀ (B)	
U ₀ (B)	
c0	...GRANTED; R ₁ (B)
	U ₁ (A)
	U ₁ (B)
	c1

T0	T1
L ₀ (A)	
R ₀ (A)	
W ₀ (A)	
	L ₁ (A) DENIED...
L ₀ (B)	
R ₀ (B)	
W ₀ (B)	
U ₀ (A)	
U ₀ (B)	
c0	
	...GRANTED; R ₁ (A) L ₁ (B); R ₁ (B); U ₁ (A) U ₁ (B) c1

iii. Show how the use of locks without 2PL can lead to a schedule that is NOT conflict-serializable.

T0	T1
L ₀ (A)	
R ₀ (A)	
W ₀ (A)	
U ₀ (A)	
	L ₁ (A)
	R ₁ (A)
	U ₁ (A)
	L ₁ (B)
	R ₁ (B)
	U ₁ (B)
	c1
L ₀ (B)	
R ₀ (B)	
W ₀ (B)	
U ₀ (B)	
c0	

4. [6 points] The following schedules are presented to a timestamp-based scheduler. Assume that the read and write timestamps of each element start at 0 ($RT(X) = WT(X) = 0$), and the commit bits for each element are set ($C(X)=1$). Explain what happens as each schedule executes.

a. $st_1, st_2, st_3, r_1(A), r_2(B), w_1(C), r_3(B), r_3(C), w_2(B), w_3(A)$

<u>T₁</u>	<u>T₂</u>	<u>T₃</u>	<u>A</u>	<u>B</u>	<u>C</u>
			RT=0 WT=0 C=1	RT=0 WT=0 C=1	RT=0 WT=0 C=1
st ₁ . Assume TS(T ₁) = 1	st ₂ . Assume TS(T ₂)=2	st ₃ . Assume TS(T ₃)=3			
r ₁ (A)			RT=1		
	r ₂ (B)			RT=2	
w ₁ (C)					WT=1 C=0
		r ₃ (B)		RT=3	
		r ₃ (C) delay			
	w ₂ (B) TS(T ₂)<RT(B) abort				
		w ₃ (A) wait			

More details about how the scheduler proceeds:

- st₁, st₂, st₃: We assign timestamps in the order the transactions start: TS(T₁) := 1, TS(T₂) := 2, and TS(T₃) := 3.
- r₁(A): OK because element A hasn't yet been written (WT(A) = 0). Sets RT(A) := (TS(T₁) = 1).
- r₂(B): OK because WT(B) = 0. Sets RT(B) := 2.
- w₁(C): OK because C hasn't yet been read or written (RT(C) = WT(C) = 0). Sets WT(C) := 1 and C(C) := 0.
- r₃(B): OK because WT(B) = 0. Sets RT(B) := 3.
- r₃(C): We have (WT(C) = 1) ≤ 3, but C(C) = 0 because T₁ has not yet committed or aborted. Delay T₃ until C(C) = 1 or T₁ aborts, then recheck timestamps and retry this action.
- w₂(B): We have (RT(B) = 3) > 2, so allowing this write would cause a write too late anomaly, because T₃ should have read the value about to be written by T₂, which comes earlier in the serialization order. Rollback T₂.
- w₃(A): Wait until and unless T₃ is unblocked and r₃(C) above succeeds. If T₃ is later aborted, then do not execute this action.

b. $st_1, st_2, r_1(A), r_2(B), w_2(A), com_2, w_1(B)$

<u>T₁</u>	<u>T₂</u>	<u>A</u>	<u>B</u>
		RT=0 WT=0 C=1	RT=0 WT=0 C=1
st_1 . Assume TS(T ₁) = 1	st_2 . Assume TS(T ₂)=2		
$r_1(A)$		RT=1	
	$r_2(B)$		RT=2
	$w_2(A)$	WT=2 C=0	
	com_2	C=1	
$w_1(B)$ abort			

More details about how the scheduler proceeds:

- st_1, st_2 : Let TS(T₁) := 1 and TS(T₂) := 2.
- $r_1(A)$: OK because WT(A) = 0. Sets RT(A) := 1.
- $r_2(B)$: OK because WT(B) = 0. Sets RT(B) := 2.
- $w_2(A)$: OK because (RT(A)=1) ≤ 2 and (WT(A)=0) ≤ 2. Sets WT(A):=2 and C(A):=0.
- com_2 : Set C(A) := 1 because T2 was the last transaction to write A, as determined by WT(A).
- $w_1(B)$: Rollback T1, because (RT(B) = 2) > 1.