

CIS 560 – Database System Concepts

Lecture 3

SQL

August 30, 2013

Credits for slides: Suciu, Chang, Ullman.

Copyright: Caragea, 2013

Announcements

- CIS MySQL accounts
- First SQL assignment to be posted tonight

Outline

Last time:

- Data in SQL
- Simple Queries in SQL (Section 6.1)
- Queries with more than one relation (Section 6.2)

Today:

- Subqueries (Section 6.3)
- Aggregations (Sections 6.4.3 - 6.4.6)

Next time:

- Nulls (Sections 6.1.6 - 6.1.7)
- Outer joins (Section 6.3.8)
- Views (Sections 8.1, 8.2, 8.3)

3

Review

Tables in SQL - terminology

Most general form of an SQL query

DISTINCT

ORDER BY

Keys and foreign keys

Joins

Tuple variables

4

Subqueries

- A *subquery* is another SQL query nested inside a larger query
- Such inner-outer queries are called *nested queries*
- A subquery may occur in:
 1. A SELECT clause
 2. A FROM clause
 3. A WHERE clause

Rule of thumb: avoid writing nested queries when possible; keep in mind that sometimes it's impossible

5

1. Subqueries in SELECT

Product (pname, price, company)

Company(cname, city)

For each product return the city where it is manufactured

```
SELECT X.pname, (SELECT Y.city
                  FROM Company Y
                  WHERE Y.cname=X.company)
FROM Product X
```

What happens if the subquery returns more than one city ?

6

1. Subqueries in SELECT

Product (pname, price, company)

Company(cname, city)

Whenever possible, don't use a nested queries:

```
SELECT pname, (SELECT city FROM Company WHERE cname=company)
FROM Product
```

=

```
SELECT pname, city
FROM Product, Company
WHERE cname=company
```

We have
“unnested”
the query

7

2. Subqueries in FROM

Product (pname, price, company)

Company(cname, city)

Find all products whose prices is > 20 and < 30

```
SELECT X.pname
FROM (SELECT * FROM Product AS Y WHERE Y.price > 20) AS X
WHERE X.price < 30
```

Unnest this query !

8

3. Subqueries in WHERE

Product (pname, price, company) Existential quantifiers
Company(cname, city)

Find all cities that make some products with price < 100

Using **EXISTS**:

```
SELECT DISTINCT Company.city
FROM   Company
WHERE  EXISTS (SELECT *
               FROM   Product
               WHERE  company = cname and Product.price < 100)
```

9

3. Subqueries in WHERE

Product (pname, price, company) Existential quantifiers
Company(cname, city)

Find all cities that make some products with price < 100

Predicate Calculus (a.k.a. First Order Logic)

```
{ y | ∃ x. Company(x,y) ∧ ( ∃ z. ∃ p. Product(z,p,x) ∧ p < 100) }
```

10

3. Subqueries in WHERE

Product (pname, price, company) Existential quantifiers
Company(cname, city)

Find all cities that make some products with price < 100

Using **IN**

```
SELECT DISTINCT Company.city
FROM   Company
WHERE  Company.cname IN (SELECT Product.company
                        FROM   Product
                        WHERE  Product.price < 100)
```

11

3. Subqueries in WHERE

Product (pname, price, company) Existential quantifiers
Company(cname, city)

Find all cities that make some products with price < 100

Using **ANY**:

```
SELECT DISTINCT Company.city
FROM   Company
WHERE  100 > ANY (SELECT price
                 FROM   Product
                 WHERE  company = cname)
```

12

3. Subqueries in WHERE

Product (pname, price, company) Existential quantifiers
Company(cname, city)

Find all cities that make some products with price < 100

Now let's unnest it:

13

3. Subqueries in WHERE

Product (pname, price, company) Existential quantifiers
Company(cname, city)

Find all cities that make some products with price < 100

Now let's unnest it:

```
SELECT DISTINCT Company.cname
FROM   Company, Product
WHERE  Company.cname = Product.company and Product.price < 100
```

Existential quantifiers are easy ! 😊

14

3. Subqueries in WHERE

Product (pname, price, company)

Universal quantifiers

Company(cname, city)

Find all cities with companies
that make only products with price < 100

Universal quantifiers are hard ! ☹

15

3. Subqueries in WHERE

Product (pname, price, company)

Universal quantifiers

Company(cname, city)

Find all cities with companies
that make only products with price < 100

Predicate Calculus (a.k.a. First Order Logic)

$\{ y \mid \exists x. \text{Company}(x,y) \wedge (\forall z. \forall p. \text{Product}(z,p,x) \rightarrow p < 100) \}$

16

3. Subqueries in WHERE

De Morgan's Laws:

$$\begin{aligned}\neg(A \wedge B) &= \neg A \vee \neg B \\ \neg(A \vee B) &= \neg A \wedge \neg B \\ \neg \forall x. P(x) &= \exists x. \neg P(x) \\ \neg \exists x. P(x) &= \forall x. \neg P(x)\end{aligned}$$

$$\neg(A \rightarrow B) = A \wedge \neg B$$

$\{ y \mid \exists x. \text{Company}(x,y) \wedge (\forall z. \forall p. \text{Product}(z,p,x) \rightarrow p < 100) \}$

=

$\{ y \mid \exists x. \text{Company}(x,y) \wedge \neg (\exists z \exists p. \text{Product}(z,p,x) \wedge p \geq 100) \}$

=

$\{ y \mid \exists x. \text{Company}(x,y) \} -$

$\{ y \mid \exists x. \text{Company}(x,y) \wedge (\exists z \exists p. \text{Product}(z,p,x) \wedge p \geq 100) \}$

3. Subqueries in WHERE

1. Find *the other* companies: i.e. s.t. some product ≥ 100

```
SELECT DISTINCT Company.city
FROM Company
WHERE Company.cname IN (SELECT Product.company
                        FROM Product
                        WHERE Product.price >= 100)
```

2. Find all companies s.t. all their products have price < 100

```
SELECT DISTINCT Company.city
FROM Company
WHERE Company.cname NOT IN (SELECT Product.company
                           FROM Product
                           WHERE Product.price >= 100)
```

3. Subqueries in WHERE

Product (pname, price, company)

Universal quantifiers

Company(cname, city)

Find all cities with companies
that make only products with price < 100

Using **EXISTS**:

```
SELECT DISTINCT Company.city
FROM   Company
WHERE  NOT EXISTS (SELECT *
                   FROM Product
                   WHERE company = cname and Product.price >= 100)
```

19

3. Subqueries in WHERE

Product (pname, price, company)

Universal quantifiers

Company(cname, city)

Find all cities with companies
that make only products with price < 100

Using **ALL**:

```
SELECT DISTINCT Company.city
FROM   Company
WHERE  100 > ALL (SELECT price
                  FROM Product
                  WHERE company = cname)
```

20

Question for Database Fans and their Friends

- Can we unnest the *universal quantifier* query ?

21

Monotone Queries

- A query Q is **monotone** if:
 - Whenever we add tuples to one or more of the tables...
 - ... the answer to the query cannot contain fewer tuples
- **Fact**: all unnested queries are monotone
 - Proof: using the “nested for loops” semantics
- **Fact**: A query with universal quantifier is not monotone
- **Consequence**: we cannot unnest a query with a universal quantifier

22

Queries that must be nested

- Queries with universal quantifiers or with negation
- The drinkers-bars-beers example next
- This is a famous example by Ullman

23

The drinkers-bars-beers example

Likes(drinker, beer)
Frequents(drinker, bar)
Serves(bar, beer)

Challenge: write these in SQL

Find drinkers that frequent some bar that serves some beer they like.

$x: \exists y. \exists z. \text{Frequents}(x, y) \wedge \text{Serves}(y, z) \wedge \text{Likes}(x, z)$

Find drinkers that frequent only bars that serve some beer they like.

$x: \forall y. \text{Frequents}(x, y) \Rightarrow (\exists z. \text{Serves}(y, z) \wedge \text{Likes}(x, z))$

Find drinkers that frequent some bar that serves only beers they like.

$x: \exists y. \text{Frequents}(x, y) \wedge \forall z. (\text{Serves}(y, z) \Rightarrow \text{Likes}(x, z))$

Find drinkers that frequent only bars that serve only beer they like.

$x: \forall y. \text{Frequents}(x, y) \Rightarrow \forall z. (\text{Serves}(y, z) \Rightarrow \text{Likes}(x, z))$

24

Aggregation

```
SELECT avg(price)
FROM Product
WHERE maker='Toyota'
```

```
SELECT count(*)
FROM Product
WHERE year > 1995
```

SQL supports several aggregation operations:

sum, count, min, max, avg

Except count, all aggregations apply to a single attribute

25

Aggregation: Count

COUNT applies to duplicates, unless otherwise stated:

```
SELECT Count(category)
FROM Product
WHERE year > 1995
```

same as Count(*)

We probably want:

```
SELECT Count(DISTINCT category)
FROM Product
WHERE year > 1995
```

26

More Examples

Purchase(product, date, price, quantity)

```
SELECT Sum(price * quantity)
FROM Purchase
```

```
SELECT Sum(price * quantity)
FROM Purchase
WHERE product = 'bagel'
```

What do
they mean ?


27

Simple Aggregations

Purchase

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

```
SELECT Sum(price * quantity)
FROM Purchase
WHERE product = 'Bagel'
```



90 (= 60+30)

28

Grouping and Aggregation

Purchase(product, price, quantity)

Find total quantity for all sales.

Find total quantity for all sales over \$1.

Find total quantities for all sales over \$1, by product.

29

Grouping and Aggregation

Purchase(product, price, quantity)

Find total quantities for all sales over \$1, by product.

```
SELECT    product, Sum(quantity) AS TotalSales
FROM      Purchase
WHERE     price > 1
GROUP BY  product
```

Let's see what this means...

30

Grouping and Aggregation

1. Compute the **FROM** and **WHERE** clauses.
2. Group by the attributes in the **GROUP BY**
3. Compute the **SELECT** clause: grouped attributes and aggregates.

31

1&2. FROM-WHERE-GROUPBY

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

32

3. SELECT

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10



???

```
SELECT product, Sum(quantity) AS TotalSales
FROM Purchase
WHERE price > 1
GROUP BY product
```

33

3. SELECT

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10



Product	TotalSales
Bagel	40
Banana	20

```
SELECT product, Sum(quantity) AS TotalSales
FROM Purchase
WHERE price > 1
GROUP BY product
```

34

GROUP BY vs. Nested Queries

Find total quantities for all sales over \$1, by product.

```
SELECT    product, Sum(quantity) AS TotalSales
FROM      Purchase
WHERE     price > 1
GROUP BY  product
```

```
SELECT DISTINCT x.product, (SELECT Sum(y.quantity)
                             FROM   Purchase y
                             WHERE  x.product = y.product
                             AND    price > 1)
                             AS TotalSales
FROM      Purchase x
WHERE     price > 1
```

Why twice ?

35

Another Example

What does
it mean ?

```
SELECT    product,
           sum(quantity) AS SumSales
           max(price) AS MaxQuantity
FROM      Purchase
GROUP BY  product
```

36

HAVING Clause

Same query, except that we consider only products that had the overall quantity sold at least 30.

```
SELECT    product, Sum(quantity)
FROM      Purchase
WHERE     price > 1
GROUP BY  product
HAVING    Sum(quantity) > 30
```

HAVING clause contains conditions on aggregates.

37

General form of Grouping and Aggregation

```
SELECT    S
FROM      R1, ..., Rn
WHERE     C1
GROUP BY  a1, ..., ak
HAVING    C2
```

Why ?

S = may contain attributes a_1, \dots, a_k and/or any aggregates but
NO OTHER ATTRIBUTES

C1 = is any condition on the attributes in R_1, \dots, R_n

C2 = is any condition on aggregate expressions

38

General form of Grouping and Aggregation

```
SELECT S  
FROM R1,...,Rn  
WHERE C1  
GROUP BY a1,...,ak  
HAVING C2
```

Evaluation steps:

1. Evaluate FROM-WHERE, apply condition C1
2. Group by the attributes a_1, \dots, a_k
3. Apply condition C2 to each group (may have aggregates)
4. Compute aggregates in S and return the result

39