

CIS 560 – Database System Concepts

Lecture 4

SQL

September 4, 2013

Credits for slides: Suciu, Chang, Ullman.

Copyright: Caragea, 2013

Announcements

- First SQL assignment due September 6th

Outline

Last time:

- Subqueries (Section 6.3)
- Aggregations (Sections 6.4.3 - 6.4.6)

Today:

- Unnesting aggregates and finding witnesses
- Nulls (Sections 6.1.6 - 6.1.7)
- Outer joins (Section 6.3.8)

Next:

- Views (Sections 8.1, 8.2, 8.3)
- Constraints (Sections 2.3, 7.1, 7.2)

3

Review

- EXISTS
- IN
- ANY
 - *operand comparison_operator ANY (subquery)*
- ALL
 - *operand comparison_operator ALL(subquery)*
- What kind of subqueries can't be unnested?
- Aggregation operators?
- Most general form of a query?
- How is the query evaluated?

General form of Grouping and Aggregation

```
SELECT  S  
FROM    R1,...,Rn  
WHERE   C1  
GROUP BY a1,...,ak  
HAVING  C2
```

S = may contain attributes a_1, \dots, a_k and/or any aggregates but
NO OTHER ATTRIBUTES

C1 = is any condition on the attributes in R_1, \dots, R_n

C2 = is any condition on aggregate expressions

5

General form of Grouping and Aggregation

```
SELECT  S  
FROM    R1,...,Rn  
WHERE   C1  
GROUP BY a1,...,ak  
HAVING  C2
```

Evaluation steps:

1. Evaluate FROM-WHERE, apply condition C1
2. Group by the attributes a_1, \dots, a_k
3. Apply condition C2 to each group (may have aggregates)
4. Compute aggregates in S and return the result

6

Advanced SQLizing

1. Unnesting Aggregates
2. Finding witnesses

7

Unnesting Aggregates

Product (pname, price, company)
Company(cname, city)

Find the number of companies in each city

```
SELECT DISTINCT city, (SELECT count(*)  
                        FROM Company Y  
                        WHERE X.city = Y.city)  
FROM Company X
```

```
SELECT city, count(*)  
FROM Company  
GROUP BY city
```

Equivalent queries

Note: no need for DISTINCT
(DISTINCT *is the same* as GROUP BY)

8

Unnesting Aggregates

Product (pname, price, company)
Company(cname, city)

Find the number of products made in each city

```
SELECT DISTINCT X.city, (SELECT count(*)  
                          FROM Product Y, Company Z  
                          WHERE Y.cname=Z.company  
                          AND Z.city = X.city)  
FROM Company X
```

```
SELECT X.city, count(*)  
FROM Company X, Product Y  
WHERE X.cname=Y.company  
GROUP BY X.city
```

They are NOT
equivalent !
(WHY?)

9

More Unnesting

Author(login,name)

Wrote(login,url)

- Find authors who wrote ≥ 10 documents.
- Attempt 1: with nested queries

```
SELECT DISTINCT Author.name  
FROM Author  
WHERE count(SELECT Wrote.url  
              FROM Wrote  
              WHERE Author.login=Wrote.login)  
       > 10
```

This is
SQL by
a novice

10

More Unnesting

- Find all authors who wrote at least 10 documents:
- Attempt 2: SQL style (with GROUP BY)

```
SELECT Author.name  
FROM Author, Wrote  
WHERE Author.login=Wrote.login  
GROUP BY Author.name  
HAVING count(wrote.url) > 10
```

This is
SQL by
an expert

11

Finding Witnesses

Store(sid, sname)

Product(pid, pname, price, sid)

For each store,
find its most expensive products

12

Finding Witnesses

Finding the maximum price is easy...

```
SELECT Store.sid, max(Product.price)
FROM   Store, Product
WHERE  Store.sid = Product.sid
GROUP BY Store.sid
```

But we need the *witnesses*, i.e. the products with max price

13

Finding Witnesses

Store(sid, sname)

Product(pid, pname, price, sid)

```
SELECT Store.sname, x.pname
FROM   Store, Product x
WHERE  Store.sid = x.sid and
       x.price >=
       ALL (SELECT y.price
            FROM Product y
            WHERE Store.sid = y.sid)
```

14

NULLS in SQL

- Whenever we don't have a value, we can put a NULL
- Can mean many things:
 - Value does not exist
 - Value exists but is unknown
 - Value not applicable
 - Value is withheld
 - Etc.
- The schema specifies for each attribute if it can be null (*nullable* attribute) or not
- How does SQL cope with tables that have NULLs ?

15

Null Values

- If x is NULL then $4*(3-x)/7$ is still NULL
- If x is NULL then $x='Joe'$ is UNKNOWN
- In SQL there are three Boolean values:

FALSE	=	0
UNKNOWN	=	0.5
TRUE	=	1

16

Null Values

- $C1 \text{ AND } C2 = \min(C1, C2)$
- $C1 \text{ OR } C2 = \max(C1, C2)$
- $\text{NOT } C1 = 1 - C1$

```
SELECT *  
FROM Person  
WHERE (age < 25) AND  
      (height > 6 OR weight > 190)
```

E.g.
age=20
height=NULL
weight=200

Rule in SQL: include only tuples that yield TRUE

17

Null Values

Unexpected behavior:

```
SELECT *  
FROM Person  
WHERE age < 25 OR age >= 25
```

Some Persons are not included !

18

Null Values

Can test for NULL explicitly:

- x IS NULL
- x IS NOT NULL

```
SELECT *  
FROM Person  
WHERE age < 25 OR age >= 25 OR age IS NULL
```

Now it includes all Persons

19

Patterns

- WHERE clauses can have conditions in which a string is compared with a pattern, to see if it matches.
- General form:
 <Attribute> LIKE <pattern> or
 <Attribute> NOT LIKE <pattern>
- Pattern is a quoted string with
 % = any string
 _ = any character

20

Example

- From Drinkers (name, addr, phone)
find the drinkers with exchange 555:

```
SELECT name
FROM Drinkers
WHERE phone LIKE '%555- _ _ _ _';
```

21

Union, Intersection, and Difference

- Union, intersection, and difference of relations are expressed by the following forms, each involving subqueries:

(subquery) UNION (subquery)

(subquery) INTERSECT (subquery)

(subquery) EXCEPT (subquery)

22

Example

- From relations

Likes(drinker, beer)

Sells(bar, beer, price) and

Frequents(drinker, bar)

find the drinkers and beers such that:

1. The drinker likes the beer, and
2. The drinker frequents at least one bar that sells the beer.

23

Solution

The drinker frequents
a bar that sells the
beer.

(SELECT * FROM Likes)

```
INTERSECT
(SELECT drinker, beer
 FROM Sells, Frequents
 WHERE Frequents.bar = Sells.bar
```

);

24

Bag Semantics for Set Operations in SQL

- Although the SELECT-FROM-WHERE statement uses *bag* semantics, the default for union, intersection, and difference is set semantics.
 - That is, duplicates are eliminated as the operation is applied.

25

Motivation: Efficiency

- When doing projection, it is easier to avoid eliminating duplicates.
 - Just work tuple-at-a-time.
- When doing intersection or difference, it is most efficient to sort the relations first.
 - At that point you may as well eliminate the duplicates anyway.

26

Controlling Duplicate Elimination

- Force the result to be a set by
SELECT DISTINCT . . .
- Force the result to be a bag (i.e., don't eliminate
duplicates) by ALL, as in . . . UNION ALL . . .

27

Example: ALL

```
Frequents(drinker, bar)  
Likes(drinker, beer)
```

```
(SELECT drinker FROM Frequents)  
  EXCEPT ALL  
(SELECT drinker FROM Likes);
```

28