

CIS 560 – Database System Concepts

Lecture 5

# SQL

September 6, 2013

Credits for slides: Suciu, Chang, Ullman.

Copyright: Caragea, 2013

## Announcements

- SQL1 assignment due by midnight
  - phpMyAdmin known to hang when executing too many poorly formed queries, e.g.

```
SELECT movie_name, COUNT(aid.actor_id), aid.gender
FROM movie_info, actor_ids AS aid, actor_movies AS amov
WHERE aid.gender = 'female'LIMIT 0, 30;
```
  - Connecting from Linux might work better
  - <http://support.cis.ksu.edu/FrequentlyAskedQuestions#FrequentlyAskedQuestions.2BAC8-Misc.ConnectingfromLinux-1>
- GTA office hours on Fridays moved to 12:30pm (he also holds office hours on Mondays at 9:30am)
- SQL2 assignment will be posted tonight

## Outline

Last time:

- Unnesting aggregates and finding witnesses
- Nulls (Sections 6.1.6 - 6.1.7)

Today:

- Outer joins (Section 6.3.8)
- Views (Sections 8.1, 8.2, 8.3)

Next:

- Constraints (Sections 2.3, 7.1, 7.2)
- E/R Diagrams (Sections 4.1-4.5)

3

## Review

- Example of a witness query
- What does a NULL value mean?
- How many Boolean values in SQL?
- How do we evaluate expressions containing NULL?
- Set operations?

4

## Outerjoins

Product(name, category)  
Purchase(prodName, store)

An “inner join”:

```
SELECT Product.name, Purchase.store  
FROM   Product, Purchase  
WHERE  Product.name = Purchase.prodName
```

Same as:

```
SELECT Product.name, Purchase.store  
FROM   Product JOIN Purchase ON  
        Product.name = Purchase.prodName
```

But Products that never sold will be lost !

5

## Outerjoins

Product(name, category)  
Purchase(prodName, store)

If we want the never-sold products, need an “outerjoin”:

```
SELECT Product.name, Purchase.store  
FROM   Product LEFT OUTER JOIN Purchase ON  
        Product.name = Purchase.prodName
```

6

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Name	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz
OneClick	NULL

7

## Application

Compute, for each product, the total number of sales in 'September'

Product(name, category)

Purchase(prodName, month, store)

```
SELECT Product.name, count(*)  
FROM   Product, Purchase  
WHERE  Product.name = Purchase.prodName  
       and Purchase.month = 'September'  
GROUP BY Product.name
```

What's wrong ?

8

## Application

Compute, for each product, the total number of sales in 'September'

Product(name, category)

Purchase(prodName, month, store)

```
SELECT Product.name, count(store)
FROM   Product LEFT OUTER JOIN Purchase ON
        Product.name = Purchase.prodName
        and Purchase.month = 'September'
GROUP BY Product.name
```

Now we also get the products that sold in 0 quantity

9

## Outer Joins

- Left outer join:
  - Include the left tuple even if there's no match
- Right outer join:
  - Include the right tuple even if there's no match
- Full outer join:
  - Include both the left and right tuples even if there's no match

10

# Views

Views are relations, except that they may not be physically stored.

Useful for presenting different information to different users

[Employee](#)(ssn, name, department, project, salary)

```
CREATE VIEW Developers AS
SELECT name, project
FROM Employee
WHERE department = 'Development'
```

Payroll has access to [Employee](#), others only to [Developers](#)

11

## Example

[Purchase](#)(customer, product, store)

[Product](#)(pname, price)

```
CREATE VIEW CustomerPrice AS
SELECT x.customer, y.price
FROM Purchase x, Product y
WHERE x.product = y.pname
```

[CustomerPrice](#)(customer, price) “virtual table”

12

Purchase(customer, product, store)  
Product(pname, price)  
CustomerPrice(customer, price)

We can later use the view:

```
SELECT u.customer, v.store  
FROM   CustomerPrice u, Purchase v  
WHERE  u.customer = v.customer AND  
        u.price > 100
```

13

## Types of Views

- Virtual views:
  - Used in databases
  - Computed only on-demand – slow at runtime
  - Always up to date
- Materialized views
  - Used in data warehouses
  - Pre-computed offline – fast at runtime
  - May have stale data
  - Indexes *are* materialized views (Section 8.3)

We discuss  
only virtual  
views in class

14

## Applications of Virtual Views


- Physical data independence, e.g.
  - Vertical data partitioning
  - Horizontal data partitioning
- Security
  - The view reveals only what the users are allowed to know

15

## Vertical Partitioning

**Resumes**

<b>SSN</b>	<b>Name</b>	<b>Address</b>	<b>Resume</b>	<b>Picture</b>
234234	Mary	Huston	Clob1...	Blob1...
345345	Sue	Seattle	Clob2...	Blob2...
345343	Joan	Seattle	Clob3...	Blob3...
234234	Ann	Portland	Clob4...	Blob4...



**T1**

<b>SSN</b>	<b>Name</b>	<b>Address</b>
234234	Mary	Huston
345345	Sue	Seattle
...		

**T2**

<b>SSN</b>	<b>Resume</b>
234234	Clob1...
345345	Clob2...

**T3**

<b>SSN</b>	<b>Picture</b>
234234	Blob1...
345345	Blob2...

16



## Vertical Partitioning

```
CREATE VIEW Resumes AS
  SELECT T1.ssn, T1.name, T1.address,
         T2.resume, T3.picture
  FROM   T1,T2,T3
  WHERE  T1.ssn=T2.ssn and T2.ssn=T3.ssn
```

17

## Vertical Partitioning

```
SELECT address
FROM   Resumes
WHERE  name = 'Sue'
```

Which of the tables T1, T2, T3 will  
be queried by the system ?

18

# Vertical Partitioning

When to do this:

- When some fields are large, and rarely accessed
  - E.g. Picture
- In distributed databases
  - Customer personal info at one site, customer profile at another
- In data integration
  - T1 comes from one source
  - T2 comes from a different source

19

# Horizontal Partitioning

## Customers

SSN	Name	City	Country
234234	Mary	Huston	USA
345345	Sue	Seattle	USA
345343	Joan	Seattle	USA
234234	Ann	Portland	USA
--	Frank	Calgary	Canada
--	Jean	Montreal	Canada



## CustomersInHuston

SSN	Name	City	Country
234234	Mary	Huston	USA

## CustomersInSeattle

SSN	Name	City	Country
345345	Sue	Seattle	USA
345343	Joan	Seattle	USA

## CustomersInCanada

SSN	Name	City	Country
--	Frank	Calgary	Canada
--	Jean	Montreal	Canada

## Horizontal Partitioning

```
CREATE VIEW Customers AS  
  CustomersInHuston  
  UNION ALL  
  CustomersInSeattle  
  UNION ALL  
  ...
```

21

## Horizontal Partitioning

```
SELECT name  
FROM Customers  
WHERE city = 'Seattle'
```

Which tables are inspected by the system ?

22

## Horizontal Partitioning


Better:

```
CREATE VIEW Customers AS
  (SELECT * FROM CustomersInHuston
   WHERE city = 'Huston')
  UNION ALL
  (SELECT * FROM CustomersInSeattle
   WHERE city = 'Seattle')
  UNION ALL
  . . .
```

23

## Horizontal Partitioning

```
SELECT name
FROM Customers
WHERE city = 'Seattle'
```



```
SELECT name
FROM CustomersInSeattle
```

24

# Horizontal Partitioning

Applications:

- Optimizations:
  - E.g. archived applications and active applications
- Distributed databases
- Data integration

25

# Views and Security

**Customers:**

Name	Address	Balance
Mary	Huston	450.99
Sue	Seattle	-240
Joan	Seattle	333.25
Ann	Portland	-520

Fred is allowed to see this

```
CREATE VIEW PublicCustomers
SELECT Name, Address
FROM Customers
```

Fred is not allowed to see this

26

## Views and Security

### Customers:

Name	Address	Balance
Mary	Huston	450.99
Sue	Seattle	-240
Joan	Seattle	333.25
Ann	Portland	-520

John is  
not allowed  
to see >0  
balances

```
CREATE VIEW BadCreditCustomers
SELECT *
FROM Customers
WHERE Balance < 0
```

27

## Constraints and Triggers

- A *constraint* is a relationship among data elements that the DBMS is required to enforce.
  - Example: key constraints.
- *Triggers* are only executed when a specified condition occurs, e.g., insertion of a tuple.
  - Easier to implement than many constraints.

28

# Constraints in SQL

- Keys
- Foreign keys (referential integrity)
- Attribute-level constraints
  - Constrain values of a particular attribute
- Tuple-level constraints
  - Relationships among components
- Global constraints: assertions

simplest

Most  
complex

The more complex the constraint, the harder it is to check and to enforce

29

## Single Attribute Keys

Product(name, category, price)

```
CREATE TABLE Product (  
    name CHAR(30) PRIMARY KEY,  
    category VARCHAR(20),  
    price INT)
```

```
CREATE TABLE Product (  
    name CHAR(30),  
    category VARCHAR(20),  
    price INT,  
    PRIMARY KEY (name))
```

30

## Keys with Multiple Attributes

```
CREATE TABLE Product (  
    name CHAR(30),  
    category VARCHAR(20),  
    price INT,  
    PRIMARY KEY (name, category))
```

Name	Category	Price
Gizmo	Gadget	10
Camera	Photo	20
Gizmo	Photo	30
<del>Gizmo</del>	<del>Gadget</del>	<del>40</del>

Product(name, category, price)

31

## Other Keys

```
CREATE TABLE Product (  
    productID CHAR(10),  
    name CHAR(30),  
    category VARCHAR(20),  
    price INT,  
    PRIMARY KEY (productID),  
    UNIQUE (name, category))
```

There is at most one **PRIMARY KEY**;  
there can be many **UNIQUE**

32



# Foreign Key Constraints

Referential  
integrity  
constraints

```
CREATE TABLE Purchase (  
  prodName CHAR(30)  
  REFERENCES Product(name),  
  date DATETIME)
```

prodName is a **foreign key** to Product(name)  
name must be a **key** in Product

May write  
just Product  
(why ?)

33

Product

<u>Name</u>	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

<u>ProdName</u>	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

34

## Foreign Key Constraints

```
CREATE TABLE Purchase (  
    prodName CHAR(30),  
    category VARCHAR(20),  
    date DATETIME,  
    FOREIGN KEY (prodName, category)  
    REFERENCES Product(name, category))
```


- (name, category) must be a key in Product

35

## What happens during updates ?

Types of “problematic” updates:

- In Purchase: insert/update
- In Product: delete/update



36