

CIS 560 – Database System Concepts

Lecture 11

Decomposition Goals & Transactions in SQL

September 20, 2013

Credits for slides: Suciu, Chang, Ullman.

Copyright: Caragea, 2013

Announcements

- HW4 will be posted tonight, due September 27th
- Exam 1 – October 7th (sample exam posted on KSOL)
- Project information posted on KSOL
- Project proposals
 - Information about team members and English description due September 24th
 - E/R diagram and relational schema due October 4th
- Proposal presentations October 9th and 11th

Review

- Closure of a set of attributes?
- Key/superkey?
- What is a “bad” functional dependency?
- What does it mean for a relation to be in BCNF?

BCNF Decomposition Algorithm

BCNF_Decompose(R)

find X s.t.: $X \neq X^+ \neq [\text{all attributes}]$

if (not found) **then** “R is in BCNF”

let $Y = X^+ - X$

let $Z = [\text{all attributes}] - X^+$

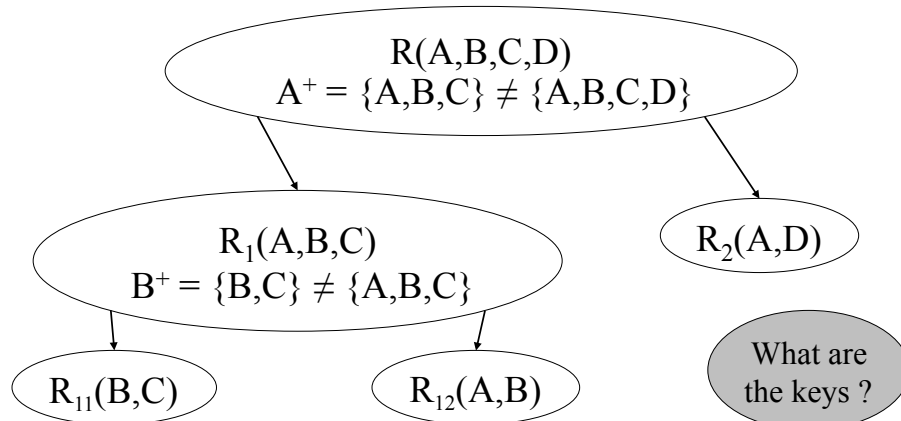
decompose R into R1($X \cup Y$) and R2($X \cup Z$)

continue to decompose recursively R1 and R2

$R(A,B,C,D)$


$A \rightarrow B$
 $B \rightarrow C$

Example



What happens if in R we first pick B^+ ? Or $\{A, B\}^+$?

Decompositions in General

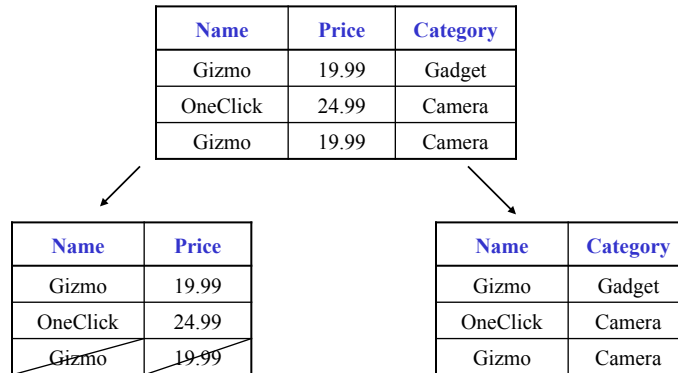


R_1 = projection of R on $A_1, \dots, A_n, B_1, \dots, B_m$

R_2 = projection of R on $A_1, \dots, A_n, C_1, \dots, C_p$

Correct Decomposition

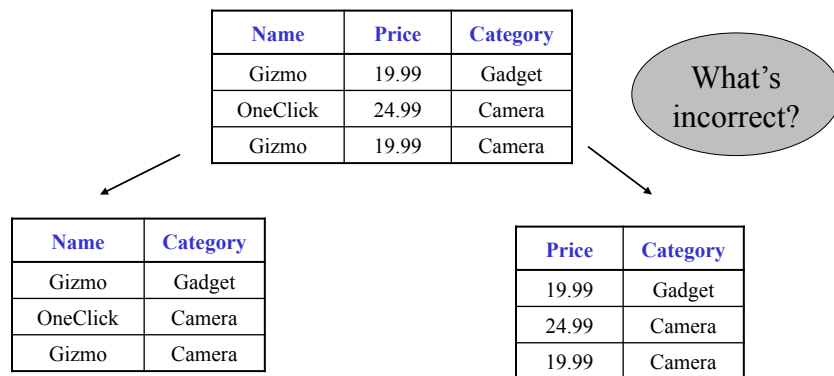
Sometimes it is correct:



Lossless decomposition

Incorrect Decomposition


Sometimes it is not:



Lossy decomposition


Lossless Decompositions

A decomposition is *lossless* if we can recover:



R' is in general larger than R. Must ensure $R' = R$

Decompositions in General



Note: don't need $A_1, \dots, A_n \rightarrow C_1, \dots, C_p$

BCNF decomposition is always lossless. WHY ?

Decomposition Based on BCNF is Necessarily Lossless

$R(A, B, C), \quad A \rightarrow B$

BCNF: $R_1(A, B), \quad R_2(A, C)$

Some tuple (a, b, c) in R
decomposes into (a, b) in R_1
and (a, c) in R_2

Recover tuple (a, b, c) in R ?

Decomposition Based on BCNF is Necessarily Lossless

$R(A, B, C), \quad A \rightarrow B$

BCNF: $R_1(A, B), \quad R_2(A, C)$

Some tuple (a, b, c) in R (a, b', c') also in R
decomposes into (a, b) in R_1 (a, b') also in R_1
and (a, c) in R_2 (a, c') also in R_2

Recover tuples in R : $(a, b, c), \quad (a, b, c'), (a, b', c), (a, b', c')$ also in R ?

Can (a, b', c) be a bogus tuple?

General Decomposition Goals

- Elimination of anomalies
 - Recoverability of information
 - Can we get the original relation back?
 - Preservation of dependencies
 - Want to enforce FDs without performing joins
- BCNF
- 3NF

Sometimes cannot decompose into BCNF without losing ability to check some FDs

A Problem with BCNF

Unit	Company	Product
------	---------	---------

FD's: Unit \rightarrow Company; Company, Product \rightarrow Unit

So, there is a BCNF violation, and we decompose.

A Problem with BCNF

Unit	Company	Product
------	---------	---------

FD's: $\text{Unit} \rightarrow \text{Company}$; $\text{Company, Product} \rightarrow \text{Unit}$

So, there is a BCNF violation, and we decompose.

Unit	Company
------	---------

$\text{Unit} \rightarrow \text{Company}$

Unit	Product
------	---------

No FDs

In BCNF we lose the FD: $\text{Company, Product} \rightarrow \text{Unit}$

So What's the Problem?

Unit	Company	Unit	Product
Galaga99	UI	Galaga99	databases
Bingo	UI	Bingo	databases

No problem so far. All *local* FD's are satisfied.

Let's put all the data back into a single table again:

Unit	Company	Product
Galaga99	UI	databases
Bingo	UI	databases

Violates the dependency: $\text{company, product} \rightarrow \text{unit}$!

Preserving FDs

- What if, when a relation is decomposed, the X of an $X \rightarrow Y$ ends up only in one of the new relations and the Y ends up only in another?
- Such a decomposition is not “dependency-preserving.”

Solution: 3rd Normal Form (3NF)

A relation R is in 3rd normal form if :

Whenever there is a nontrivial dependency $A_1, A_2, \dots, A_n \rightarrow B$ for R ,
then $\{A_1, A_2, \dots, A_n\}$ is a super-key for R,
or B is part of a key.

3NF vs. BCNF

- R is in **BCNF** if whenever $X \rightarrow A$ holds, then X is a superkey.
- Slightly stronger than 3NF.
- Example: R(A,B,C) with $\{A,B\} \rightarrow C$, $C \rightarrow A$
 - 3NF but not BCNF

Trade-offs

BCNF = no anomalies, but may lose some FDs

3NF = keeps all FDs, but may have some anomalies

- Everyday relational DBs
 - aim for BCNF, settle for 3NF

Caveat

- Normalization is not the be-all and end-all of DB design.
- Example: suppose attributes A and B are always used together, but normalization theory says they should be in different tables.
 - decomposition might produce unacceptable performance loss (extra disk reads)

Where We Are

What we have already learned so far

- Relational model of data
- Data manipulation language: SQL
- Views and constraints
- Database design (E/R diagrams & normalization)

But what if I want to update my data?

- Need to worry about transactions in SQL

21

Transactions

- The problem: An application must perform *several* writes and reads to the database, as a unit.
 - Example: Two people attempt to book the last seat on a flight.
- Solution: multiple actions of the application are bundled into one unit called *Transaction*.
 - Transactions guarantee certain properties to hold that prevent problems.

22

Turing Awards to Database Researchers

- Charles Bachman 1973 for CODASYL
- Edgar Codd 1981 for relational databases
- Jim Gray 1998 for transactions

23

The World Without Transactions

- Just write applications that talk to databases
- Rely on operating systems for scheduling, and for concurrency control
- What can go wrong?
 - Three famous anomalies
 - Other anomalies are possible (but not so famous)

24