

CIS 560 – Database System Concepts

Lecture 12

Transactions in SQL and Relational Algebra

September 23, 2013

Credits for slides: Suciu, Chang.

Copyright: Caragea, 2013.

Outline

Last time:

- DB Design: Normalization (3.3-3.4)

Today:

- Transactions in SQL (6.6)
- Relational algebra (Sections 2.4 and 5.1-5.2)

Next:

- Relational algebra (Sections 2.4 and 5.1-5.2)
- Introduction to Database Programming (Ch. 9)
 - Connect to DB and call SQL from Java

2

Review

- BCNF
- 3NF
- Lossless decomposition
- Preserving functional dependencies

3

Transactions

- The problem: An application must perform *several* writes and reads to the database, as a unit.
 - Example: Two people attempt to book the last seat on a flight.
- Solution: multiple actions of the application are bundled into one unit called *Transaction*.
 - Transactions guarantee certain properties to hold that prevent problems.

4

The World Without Transactions

- Just write applications that talk to databases
- Rely on operating systems for scheduling, and for concurrency control
- What can go wrong?
 - Three famous anomalies
 - Other anomalies are possible (but not so famous)

5

Lost Updates

Client 1:

```
UPDATE Customer  
SET rentals= rentals + 1  
WHERE cname= 'Fred'
```

Client 2:

```
UPDATE Customer  
SET rentals= rentals + 1  
WHERE cname= 'Fred'
```

Two people attempt to rent two movies for Fred, from two different terminals. What happens?

6

Inconsistent Read

Client 1: rent-a-movie
 x = **SELECT** rentals **FROM** Cust
 WHERE cname= 'Fred'

```
if(x < 5)
{ UPDATE Cust
  SET rentals= rentals + 1
  WHERE cname= 'Fred' }
else println("Denied !")
```

What's wrong ?

Client 2: rent-a-movie
 x = **SELECT** rentals **FROM** Cust
 WHERE cname= 'Fred'

```
if(x < 5)
{ UPDATE Cust
  SET rentals= rentals + 1
  WHERE cname= 'Fred' }
else println("Denied !")
```

7

Inconsistent Read

Client 1: rent-two-movies
 x = **SELECT** rentals **FROM** Cust
 WHERE cname= 'Fred'

```
if(x < 4) { /* movie 1...*/
  UPDATE Cust
  SET rentals= rentals + 1
  WHERE cname= 'Fred'

  /* ...and movie 2...*/
  UPDATE Cust
  SET rentals= rentals + 1
  WHERE cname= 'Fred'
}
else println("Denied !")
```

Client 2: rent-a-movie
 x = **SELECT** rentals **FROM** Cust
 WHERE cname= 'Fred'

```
if(x < 5)
{ UPDATE Cust
  SET rentals= rentals + 1
  WHERE cname= 'Fred' }
else println("Denied !")
```

What's wrong ?

8

Inconsistent Read

Client 1: move from gizmo → gadget

```
UPDATE Products
SET quantity = quantity + 5
WHERE product = 'gizmo'
```

```
UPDATE Products
SET quantity = quantity - 5
WHERE product = 'gadget'
```

Client 2: inventory....

```
SELECT sum(quantity)
FROM Product
```

What's wrong ?

9

Dirty Reads

Client 1: transfer \$100 acc1 → acc2
 X = Account1.balance
 Account2.balance += 100

```
If (X >= 100) Account1.balance -= 100
else { /* rollback ! */
    account2.balance -= 100
    println("Denied !")
}
```

Client 2: transfer \$100 acc2 → acc3
 Y = Account2.balance
 Account3.balance += 100

```
If (Y >= 100) Account2.balance -= 100
else { /* rollback ! */
    account3.balance -= 100
    println("Denied !")
}
```

What's wrong ?

The Three Famous Anomalies

- Lost update
 - Two tasks T and T' both modify the same data
 - T and T' both commit
 - Final state shows effects of only T, but not of T'
- Inconsistent read:
 - One task T might see some but not all changes made by T'
- Dirty read
 - T reads data written by T' while T' has not committed
 - What can go wrong: T' writes more data (which T has already read), or T' aborts

11

Protection against crashes

Client 1:

```
UPDATE Accounts
SET balance= balance - 500
WHERE name= 'Fred'
```

```
UPDATE Accounts
SET balance = balance + 500
WHERE name= 'Joe'
```

Crash !

What's wrong ?

12

Transactions

Ensure two things:

- Concurrency control
 - Making sure simultaneous transactions don't interfere with one another
- Recovery
 - Taking action to restore the DB to a consistent state

13

Definition

- **A transaction** = one or more operations, which reflects a single real-world transition
 - Happens completely or not at all
- Examples
 - Transfer money between accounts
 - Rent a movie; return a rented movie
 - Purchase a group of products
 - Register for a class (either waitlisted or allocated)
- By using transactions, all previous problems disappear

14

Transactions in Applications

- Default: each statement = one transaction
- Multi-statement transactions:

START TRANSACTION

[SQL statements]

COMMIT or **ROLLBACK (=ABORT)**

15

Revised Code

Client 1: rent-a-movie

START TRANSACTION

x = **SELECT** rentals **FROM** Cust
WHERE cname= 'Fred'

if(x < 5)

{ **UPDATE** Cust
SET rentals= rentals + 1
WHERE cname= 'Fred' }

else println("Denied !")

COMMIT

Client 2: rent-a-movie

START TRANSACTION

x = **SELECT** rentals **FROM** Cust
WHERE cname= 'Fred'

if(x < 5)

{ **UPDATE** Cust
SET rentals= rentals + 1
WHERE cname= 'Fred' }

else println("Denied !")

COMMIT

16

Revised Code

Client 1: transfer \$100 acc1 → acc2

START TRANSACTION

X = Account1.balance; Account2.balance += 100

If (X >= 100) { Account1.balance -= 100; COMMIT }
else { println("Denied!"); ROLLBACK }

Client 2: transfer \$100 acc2 → acc3

START TRANSACTION

X = Account2.balance; Account3.balance += 100

If (X >= 100) { Account2.balance -= 100; COMMIT }
else { println("Denied!"); ROLLBACK }

17

Using Transactions

Very easy to use:

- START TRANSACTION
- COMMIT
- ROLLBACK

But what EXACTLY do they mean ?

- Popular culture: ACID
- Underlying theory: serializability

18

Transaction Properties

ACID

- **A**tomic
 - State shows either all the effects of a transaction, or none of them
- **C**onsistent
 - A transaction moves from a state where integrity holds, to another where integrity holds
- **I**solated
 - Effect of transactions is the same as transactions running one after another (i.e., looks like batch mode)
- **D**urable
 - Once a transaction has committed, its effects remain in the database

19

ACID: Atomicity

- Two possible outcomes for a transaction
 - It *commits*: all the changes are made
 - It *aborts*: no changes are made
- That is, transaction's activities are all or nothing

20

ACID: Isolation

- A transaction executes concurrently with other transaction
- Isolation: the effect is as if each transaction executes in isolation of the others

21

ACID: Consistency

- The database satisfies integrity constraints
 - Account number is unique
 - Stock amount can't be negative
 - Sum of *debits* and of *credits* is 0
- Constraints may be explicit or implicit
- How consistency is achieved:
 - Applications preserve consistency, assuming they run atomically, and they run in isolation
 - The system ensures atomicity and isolation

22

ACID: Durability

- The effect of a transaction must continue to exist after the transaction, or the whole program has terminated
- Means: write data to disk
- Sometimes also means recovery

23

ROLLBACK

- If the application gets to a place where it can't complete the transaction successfully, it can execute ROLLBACK
- This causes the system to "abort" the transaction
 - The database returns to the state without any of the previous changes made by activity of the transaction

24

Reasons for Rollback

- Users change their minds (“ctl-C”/cancel)
- Explicit in program, when application program finds a problem
 - E.g. when the # of rented movies > max # allowed
 - Use it freely in your next programming assignment!
- System-initiated abort
 - System crash
 - Housekeeping, e.g. due to timeouts

25

Relational Algebra

26

The WHAT and the HOW

- In SQL we write **WHAT** we want to get from the data
- The database system needs to figure out **HOW** to get the data we want
- The passage from **WHAT** to **HOW** goes through the **Relational Algebra**

27

SQL = WHAT

Product(pid, name, price)

Purchase(pid, cid, store)

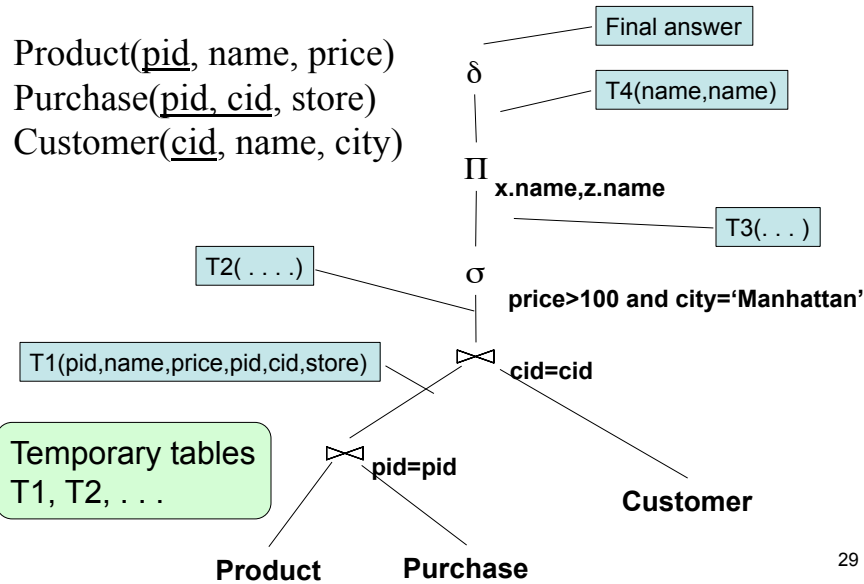
Customer(cid, name, city)

```
SELECT DISTINCT x.name, z.name
FROM Product x, Purchase y, Customer z
WHERE x.pid = y.pid and y.cid = z.cid and
      x.price > 100 and z.city = 'Manhattan'
```

It's clear WHAT we want, unclear HOW to get it

28

Relational Algebra = HOW

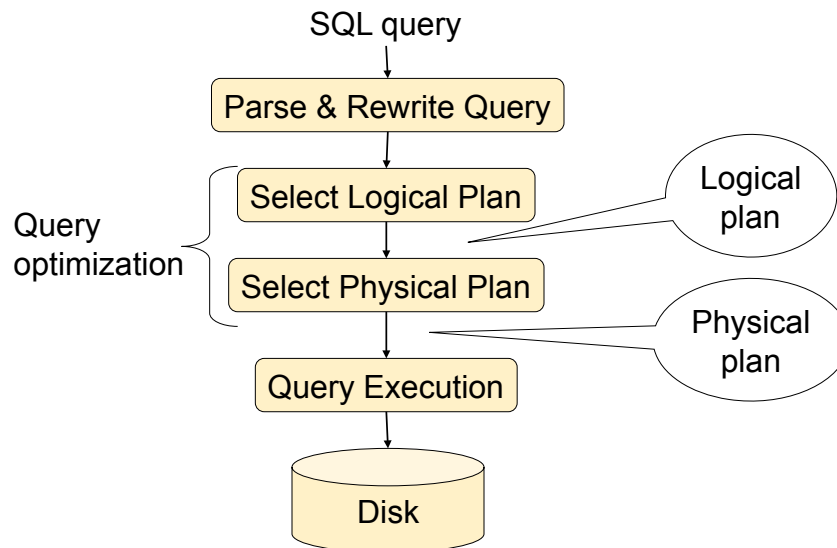


Relational Algebra = HOW

The order is now clearly specified:

Iterate over PRODUCT...
 ...join with PURCHASE...
 ...join with CUSTOMER...
 ...select tuples with Price > 100 and City = 'Manhattan'...
 ...eliminate duplicates...
 ...and that's the final answer !

Steps of the Query Processor



Query Plans

- **Logical query plan:** an extended relational algebra tree
- **Physical query plan:** with additional annotations at each node
 - Access method to use for each relation
 - Implementation to use for each relational operator

Relations

- A relation is a set of tuples
 - Sets: {a,b,c}, {a,d,e,f}, { }, . . .
- But, commercial DBMSs implement **relations that are bags** rather than sets
 - Bags: {a, a, b, c}, {b, b, b, b, b}, . . .

33

Sets versus Bags

- Relational Algebra has two flavors:
 - **Over sets**: theoretically elegant but limited
 - **Over bags**: needed for SQL queries + more efficient
 - Example: Compute average price of all products
- We discuss set semantics
 - We mention bag semantics only where needed

34

Relational Algebra

- Query language associated with the relational model
- Queries specified in an operational manner
 - A query gives a step-by-step procedure
- Relational operators
 - Take one or two relation instances as argument
 - Return one relation instance as result
 - Easy to compose into relational algebra expressions

35

Relational Algebra (1/3)

The Basic Five operators:

- Union: \cup
- Set difference: $-$
- Selection: $\sigma_{\text{condition}}(S)$
 - Condition is a Boolean combination (\wedge, \vee) of terms
 - Term is: attr op const, attr op attr
 - Op is: $<$, $<=$, $=$, $!=$, $>=$, or $>$
- Projection: $\pi_{\text{list-of-attributes}}(S)$
- Cross-product or Cartesian product: \times

36

Relational Algebra (2/3)

Derived or auxiliary operators:

- Intersection (\cap)
- Join $R \bowtie_{\theta} S = \sigma_{\theta}(R \times S)$
- Variations of joins
 - natural, equijoin, theta join
 - outer-join and semi-join
- Rename $\rho_{B1, \dots, Bn}(S)$

37

Relational Algebra (3/3)

Extensions for bags:

- Duplicate elimination: δ
- Group by: γ [Same symbol as aggregation]
 - Partitions tuples of a relation into “groups”
- Sorting: τ

Other extensions:

- Aggregation: γ (min, max, sum, average, count)

38

Union and Difference

$R1 \cup R2$

Example: ActiveEmployees \cup RetiredEmployees

$R1 - R2$

Example: AllEmployees – RetiredEmployees

Be careful when applying to bags!

39

What about Intersection?

- Derived operator using minus

$$R1 \cap R2 = R1 - (R1 - R2)$$

Example: UnionizedEmployees \cap RetiredEmployees

40

Selection

- Returns all tuples which satisfy a condition

$$\sigma_{\text{condition}}(R)$$

- Examples

- $\sigma_{\text{Salary} > 40000}(\text{Employee})$
- $\sigma_{\text{name} = \text{"Smith"}}(\text{Employee})$

Maps to the WHERE clause in SQL

41

Employee

SSN	Name	Salary
1234545	John	200000
5423341	Smith	600000
4352342	Fred	500000

$\sigma_{\text{Salary} > 40000}(\text{Employee})$

SSN	Name	Salary
5423341	Smith	600000
4352342	Fred	500000

42

Projection

- Eliminates columns

$$\Pi_{A_1, \dots, A_n}(R)$$

- Example: project on Name and Salary:
 - $\Pi_{\text{Name, Salary}}(\text{Employee})$
 - Answer(Name, Salary)

Semantics differs over set or over bags

43

Employee

SSN	Name	Salary
1234545	John	200000
5423341	John	600000
4352342	John	200000

$\Pi_{\text{Name, Salary}}(\text{Employee})$

Name	Salary
John	20000
John	60000
John	20000

Bag semantics

Name	Salary
John	20000
John	60000

Set semantics

Which is more efficient to implement?

44

Selection & Projection Examples

Patient

no	name	zip	disease
1	p1	98125	flu
2	p2	98125	heart
3	p3	98120	lung
4	p4	98120	heart

$\pi_{\text{zip,disease}}(\text{Patient})$

zip	disease
98125	flu
98125	heart
98120	lung
98120	heart

$\sigma_{\text{disease}='heart'}(\text{Patient})$

no	name	zip	disease
2	p2	98125	heart
4	p4	98120	heart

$\pi_{\text{zip}}(\sigma_{\text{disease}='heart'}(\text{Patient}))$

zip
98120
98125

Cartesian Product

- Each tuple in R1 with each tuple in R2

$$R1 \times R2$$

- Very rare in practice; mainly used to express joins

Employee

Name	SSN
John	999999999
Tony	777777777

Dependent

EmpSSN	DepName
999999999	Emily
777777777	Joe

Employee × Dependent

Name	SSN	EmpSSN	DepName
John	999999999	999999999	Emily
John	999999999	777777777	Joe
Tony	777777777	999999999	Emily
Tony	777777777	777777777	Joe

47

Renaming

- Changes the schema, not the instance

$$\rho_{B_1, \dots, B_n}(R)$$

- Example:
 - $\rho_{\text{FirstName}, \text{SocSecNo}}(\text{Employee})$
 $\rightarrow \text{Employee}(\text{FirstName}, \text{SocSecNo})$

48

Renaming

Employee

Name	SSN
John	999999999
Tony	777777777

$\rho_{\text{FirstName, SocSecNo}}(\text{Employee})$

Employee

FirstName	SocSecNo
John	999999999
Tony	777777777

49