

**CIS 560 – Database System Concepts**

**Lecture 13**

# Relational Algebra

September 25, 2013

Credits for slides: Suciu, Chang.

Copyright: Caragea, 2013.

## Outline

Last time:

- Transactions in SQL (6.6)

Today:

- Relational algebra (Sections 2.4 and 5.1-5.2)

Next:

- Introduction to Database Programming (Ch. 9)
  - Connect to DB and call SQL from Java

## Review

- Transaction anomalies
- ACID properties

3

## The WHAT and the HOW

- In SQL we write **WHAT** we want to get from the data
- The database system needs to figure out **HOW** to get the data we want
- The passage from **WHAT** to **HOW** goes through the **Relational Algebra**

4

## SQL = WHAT

Product(pid, name, price)

Purchase(pid, cid, store)

Customer(cid, name, city)

```
SELECT DISTINCT x.name, z.name
FROM Product x, Purchase y, Customer z
WHERE x.pid = y.pid and y.cid = z.cid and
      x.price > 100 and z.city = 'Manhattan'
```

It's clear WHAT we want, unclear HOW to get it

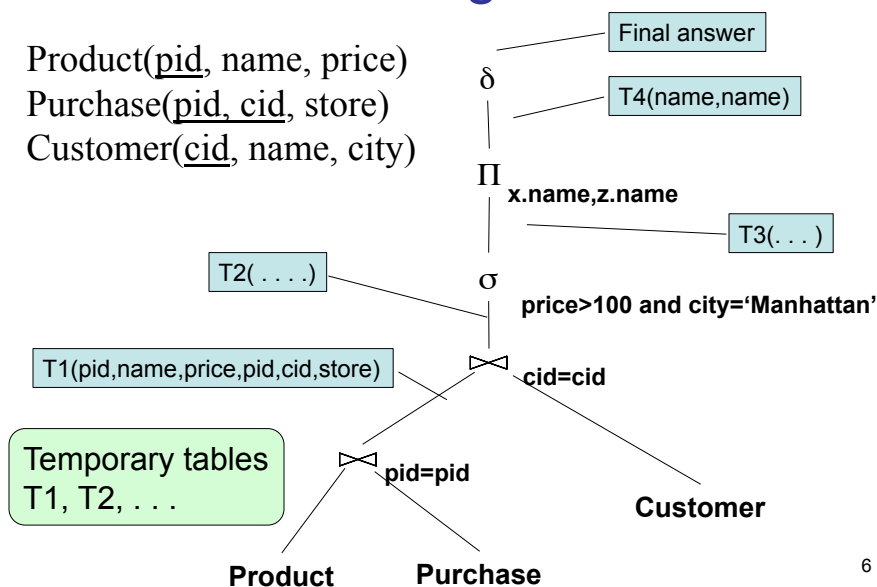
5

## Relational Algebra = HOW

Product(pid, name, price)

Purchase(pid, cid, store)

Customer(cid, name, city)



6

## Relational Algebra = HOW

The order is now clearly specified:

Iterate over PRODUCT...  
 ...join with PURCHASE...  
 ...join with CUSTOMER...  
 ...select tuples with Price>100 and City='Manhattan'...  
 ...eliminate duplicates...  
 ...and that's the final answer !

7

## Relational Algebra (1/3)

The Basic Five operators:

- Union:  $\cup$
- Set difference:  $-$
- Selection:  $\sigma_{\text{condition}}(S)$ 
  - Condition is a Boolean combination ( $\wedge, \vee$ ) of terms
  - Term is: attr op const, attr op attr
  - Op is:  $<$ ,  $<=$ ,  $=$ ,  $!=$ ,  $>=$ , or  $>$
- Projection:  $\pi_{\text{list-of-attributes}}(S)$
- Cross-product or Cartesian product:  $\times$

8

## Relational Algebra (2/3)

Derived or auxiliary operators:

- Intersection ( $\cap$ )
- Join  $R \bowtie_{\theta} S = \sigma_{\theta}(R \times S)$
- Variations of joins
  - natural, equijoin, theta join
  - outer-join and semi-join
- Rename  $\rho_{B1, \dots, Bn}(S)$

9

## Relational Algebra (3/3)

Extensions for bags:

- Duplicate elimination:  $\delta$
- Group by:  $\gamma$  [Same symbol as aggregation]
  - Partitions tuples of a relation into “groups”
- Sorting:  $\tau$

Other extensions:

- Aggregation:  $\gamma$  (min, max, sum, average, count)

10

## Different Types of Join

- **Theta-join:**  $R \bowtie_{\theta} S = \sigma_{\theta}(R \times S)$ 
  - Join of R and S with a join condition  $\theta$
  - Cross-product followed by selection  $\theta$
- **Equijoin:**  $R \bowtie_{\theta} S = \pi_A(\sigma_{\theta}(R \times S))$ 
  - Join condition  $\theta$  consists only of equalities
  - Projection  $\pi_A$  drops all redundant attributes
- **Natural join:**  $R \bowtie S = \pi_A(\sigma_{\theta}(R \times S))$ 
  - Equijoin
  - Equality on all fields with same name in R and in S

11

## Theta-Join Example

AnonPatient P

age	zip	disease
54	98125	heart
20	98120	flu

AnonJob J

age	zip	job
54	98125	lawyer
20	98120	cashier

$$P \bowtie_{P.age=J.age \wedge P.zip=J.zip \wedge P.age<50} J$$

P.age	P.zip	disease	job	J.age	J.zip
20	98120	flu	cashier	20	98120

12

## Equijoin Example

AnonPatient P

age	zip	disease
54	98125	heart
20	98120	flu

AnonJob J

age	zip	job
54	98125	lawyer
20	98120	cashier

$$P \bowtie_{P.age=J.age} J$$

P.age	P.zip	disease	job	J.zip
54	98125	heart	lawyer	98125
20	98120	flu	cashier	98120

13

## Natural Join Example

AnonPatient P

age	zip	disease
54	98125	heart
20	98120	flu

AnonJob J

age	zip	job
54	98125	lawyer
20	98120	cashier

$$P \bowtie J$$

P.age	P.zip	disease	job
54	98125	heart	lawyer
20	98120	flu	cashier

14

## So Which Join Is It?

- When we write  $R \bowtie S$  we usually mean an equijoin - we often omit the equality predicate when it is clear from the context.

15

## More Joins

- Outer join
  - Include tuples with no matches in the output
  - Use NULL values for missing attributes
- Variants
  - Left outer join
  - Right outer join
  - Full outer join

16



## Outer Join Example

AnonPatient P

age	zip	disease
54	98125	heart
20	98120	flu
33	98120	lung

AnonJob J

age	zip	job
54	98125	lawyer
20	98120	cashier

 $P \bowtie J$ 

P.Age	P.zip	disease	job
54	98125	heart	lawyer
20	98120	flu	cashier
33	98120	lung	NULL

17

## Joins

- The join operation and all its variants (equijoin, natural join, outer-join) are at the heart of relational database systems
- WHY?

18

## Example of Algebra Queries

- Q1: Jobs of patients who have heart disease

19

## Example of Algebra Queries

- Q1: Jobs of patients who have heart disease
  - $\pi_{\text{job}}(\text{AnonJob} \bowtie (\sigma_{\text{disease}='heart'}(\text{AnonPatient})))$

20

## More Examples

Supplier(sno,sname,scity,sstate)  
 Part(pno,pname,psize,pcolor)  
 Supply(sno,pno,qty,price)

Q2: Name of supplier of parts with size greater than 10

21

## More Examples

Supplier(sno,sname,scity,sstate)  
 Part(pno,pname,psize,pcolor)  
 Supply(sno,pno,qty,price)

Q2: Name of supplier of parts with size greater than 10

$\pi_{\text{sname}}(\text{Supplier} \bowtie \text{Supply} \bowtie (\sigma_{\text{psize} > 10}(\text{Part})))$

22

## More Examples

Supplier(sno,sname,scity,sstate)  
 Part(pno,pname,psize,pcolor)  
 Supply(sno,pno,qty,price)

Q2: Name of supplier of parts with size greater than 10

$\pi_{\text{sname}}(\text{Supplier} \bowtie \text{Supply} \bowtie (\sigma_{\text{psize} > 10}(\text{Part})))$

Q3: Name of supplier of red parts or parts with size greater than 10

23

## More Examples

Supplier(sno,sname,scity,sstate)  
 Part(pno,pname,psize,pcolor)  
 Supply(sno,pno,qty,price)

Q2: Name of supplier of parts with size greater than 10

$\pi_{\text{sname}}(\text{Supplier} \bowtie \text{Supply} \bowtie (\sigma_{\text{psize} > 10}(\text{Part})))$

Q3: Name of supplier of red parts or parts with size greater than 10

$\pi_{\text{sname}}(\text{Supplier} \bowtie \text{Supply} \bowtie ((\sigma_{\text{psize} > 10}(\text{Part}) \cup \sigma_{\text{pcolor} = \text{'red'}}(\text{Part})))$

24

## Regular Algebra Expressions versus Programs

- An Algebra Expression is like a program
  - Several operations
  - Strictly specified order
- But relational algebra expressions have limitations

25

## Limitations of Relational Algebra

- Cannot compute “transitive closure”

Name1	Name2	Relationship
Fred	Mary	Father
Mary	Joe	Cousin
Mary	Bill	Spouse
Nancy	Lou	Sister

- Find all direct and indirect relatives of Fred
- Cannot express in RA!!! Need to write Java program

26

## From SQL to Relational Algebra

Product(pid, name, price)

Purchase(pid, cid, store)

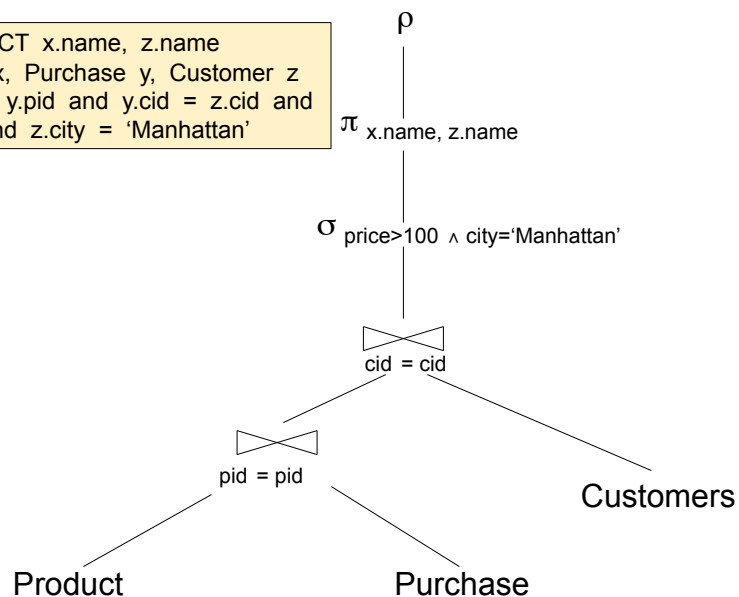
Customer(cid, name, city)

```
SELECT DISTINCT x.name, z.name
FROM Product x, Purchase y, Customer z
WHERE x.pid = y.pid and y.cid = z.cid and
x.price > 100 and z.city = 'Manhattan'
```

27

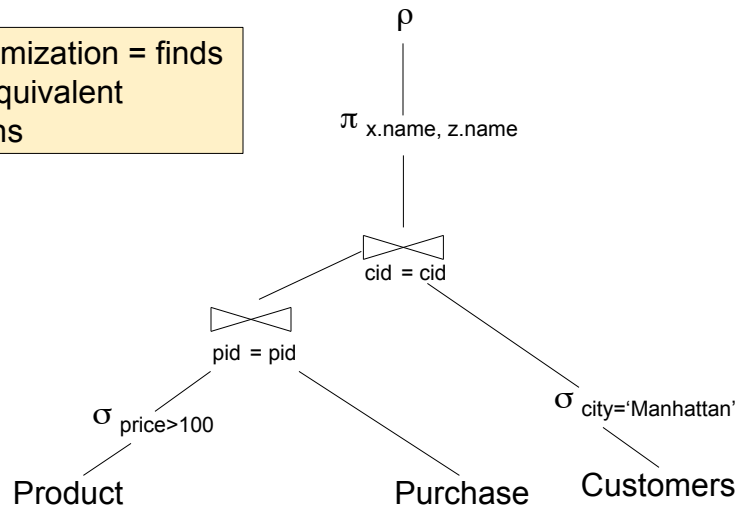
## From SQL to Relational Algebra

```
SELECT DISTINCT x.name, z.name
FROM Product x, Purchase y, Customer z
WHERE x.pid = y.pid and y.cid = z.cid and
x.price > 100 and z.city = 'Manhattan'
```

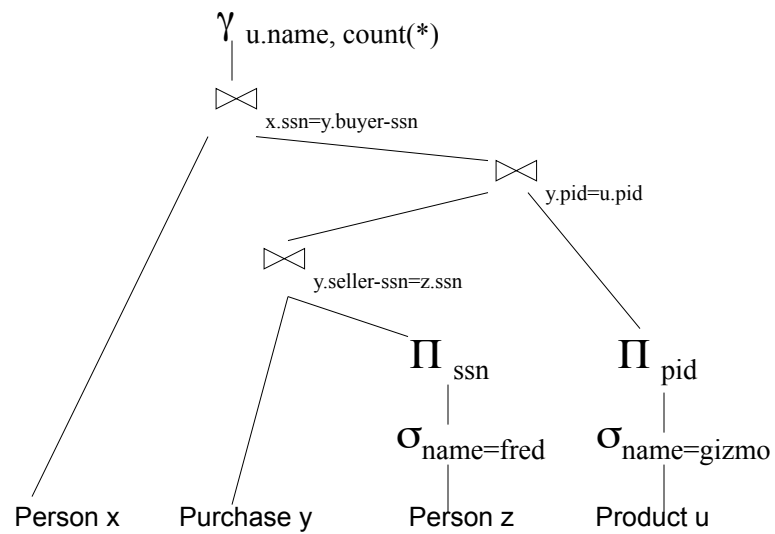


## An Equivalent Expression

Query optimization = finds cheaper equivalent expressions



## Complex RA Expressions – SQL?



30

## Example Database Schema

```
Supplier(sno,sname,scity,sstate)
Part(pno,pname,psize,pcolor)
Supplies(sno,pno,price)
```

### View: Suppliers in Manhattan, KS

```
CREATE VIEW NearbySupp AS
SELECT sno, sname
FROM Supplier
WHERE scity='Manhattan' AND sstate='KS'
```

```
Supplier(sno,sname,scity,sstate)
Part(pno,pname,psize,pcolor)
Supplies(sno,pno,price)
```

## Exercise – Translate to RA

Find the names of all suppliers in Manhattan who supply part number 2

```
SELECT sname FROM NearbySupp
WHERE sno IN ( SELECT sno
                FROM Supplies
                WHERE pno = 2 )
```



Supplier(sno,sname,scity,sstate)  
 Part(pno,pname,psize,pcolor)  
 Supplies(sno,pno,price)

## Rewritten Version of Our Query

### Original query:

```
SELECT sname
FROM NearbySupp
WHERE sno IN ( SELECT sno
                FROM Supplies
                WHERE pno = 2 )
```

### View:

```
CREATE VIEW NearbySupp AS
SELECT sno, sname
FROM Supplier
WHERE scity='Manhattan' AND
sstate='KS'
```

Supplier(sno,sname,scity,sstate)  
 Part(pno,pname,psize,pcolor)  
 Supplies(sno,pno,price)

## Rewritten Version of Our Query

### Original query:

```
SELECT sname
FROM NearbySupp
WHERE sno IN ( SELECT sno
                FROM Supplies
                WHERE pno = 2 )
```

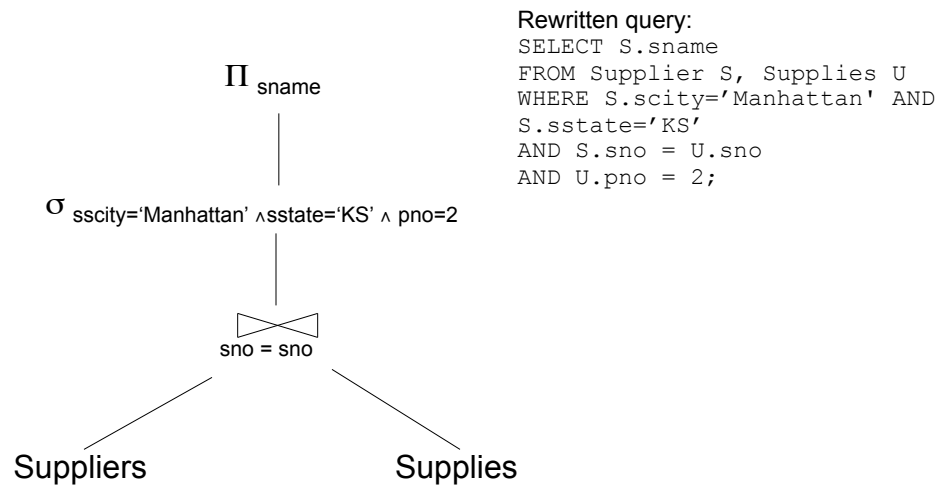
### View:

```
CREATE VIEW NearbySupp AS
SELECT sno, sname
FROM Supplier
WHERE scity='Manhattan' AND
sstate='KS'
```

### Rewritten query:

```
SELECT S.sname
FROM Supplier S, Supplies U
WHERE S.scity='Manhattan' AND S.sstate='KS'
AND S.sno = U.sno
AND U.pno = 2;
```

## Logical Query Plan



## Exercise – Translate to RA

Supplier(sno,sname,scity,sstate)  
 Part(pno,pname,psize,pcolor)  
 Supplies(sno,pno,price)

```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'KS'
and not exists
  SELECT *
  FROM Supplies P
  WHERE P.sno = Q.sno
  and P.price > 100
```

Supplier(sno,sname,scity,sstate)  
Part(pno,pname,psize,pcolor)  
Supplies(sno,pno,price)

## How about Subqueries?

```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'KS'
and not exists
  SELECT *
  FROM Supplies P
  WHERE P.sno = Q.sno
  and P.price > 100
```

Supplier(sno,sname,scity,sstate)  
Part(pno,pname,psize,pcolor)  
Supplies(sno,pno,price)

## How about Subqueries?

```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'KS'
and not exists
  SELECT *
  FROM Supplies P
  WHERE P.sno = Q.sno
  and P.price > 100
```

Correlation!

Supplier(sno,sname,scity,sstate)  
 Part(pno,pname,psize,pcolor)  
 Supplies(sno,pno,price)

## How about Subqueries?

```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'KS'
and not exists
  SELECT *
  FROM Supplies P
  WHERE P.sno = Q.sno
  and P.price > 100
```

De-Correlation

```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'KS'
and Q.sno not in
  SELECT P.sno
  FROM Supplies P
  WHERE P.price > 100
```

Supplier(sno,sname,scity,sstate)  
 Part(pno,pname,psize,pcolor)  
 Supplies(sno,pno,price)

## How about Subqueries?

Un-nesting

```
(SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'KS')
EXCEPT
(SELECT P.sno
FROM Supplies P
WHERE P.price > 100)
```

```
SELECT Q.sno
FROM Supplier Q
WHERE Q.sstate = 'KS'
and Q.sno not in
  SELECT P.sno
  FROM Supplies P
  WHERE P.price > 100
```

Supplier(sno,sname,scity,sstate)  
Part(pno,pname,psize,pcolor)  
Supplies(sno,pno,price)

## How about Subqueries?

```
(SELECT Q.sno  
FROM Supplier Q  
WHERE Q.sstate = 'KS')  
EXCEPT  
(SELECT P.sno  
FROM Supplies P  
WHERE P.price > 100)
```

Finally...

