

## Introduction to Database Programming (Chapter 9)

September 27, 2013

Credits for slides: Chang, Ullman.

Copyright: Caragea & Milliner, 2013.

### Where We Are

- User perspective
  - How to use a database system
    - Conceptual data modeling, database schema design (normalization), the SQL query language, **database programming**
- System perspective
  - How database systems work
    - Data storage and indexing, query optimization and processing, transaction management

## SQL in Real Programs

- We have seen how SQL is used at the **generic query interface** --- an environment where we sit at a terminal and ask queries of a database.
- SQL is a very high-level language - **not** intended for general-purpose computations.
- Most of the time, we need to write conventional programs that interact with SQL.

## Three-Tier Architecture

A common environment for using a database has three tiers of processes:

1. *Web servers* - talk to the user.
2. *Application servers* - execute the business logic.
3. *Database servers* - get what the application servers need from the database.


## Example: Amazon

- Database holds the information about products, customers, etc.
- Business logic includes things like “what do I do after someone clicks ‘checkout’?”
  - **Answer:** Show the “how will you pay for this?” screen.

## Environments, Connections, Statements

- The database is, in many DB-access languages, an **environment**.
- Database servers maintain some number of **connections**, so application servers can ask queries or perform modifications.
- The application server issues **statements**: queries and modifications, usually.

## Diagram to Remember



## DB Programming Options

- **Outside DBMS:** use SQL together with general-purpose programming languages.
  - SQL statements are embedded in a host language (e.g., C).
    - Directly embedded SQL, not widely used [70's - ]
  - Connection tools are used to allow a conventional language to access a database (e.g., CLI, JDBC, PHP/DB).
    - API approach - Call-Level Interface [90's - ]
- **Inside DBMS:** augment SQL with constructs from general-purpose programming languages.
  - Code in a specialized language is stored in the database itself (e.g., PSM, PL/SQL).
    - Stored procedures/functions [80's - ]

## Basic Paradigm

- Connect to a DB server.
- Say what database you want to use.
- Assemble a string containing an SQL statement.
- Get the DBMS to prepare a plan for executing the statement.
- Execute the statement.
- Extract the results into variables in the local programming language.

## Programming Outside the DBMS

## Embedded SQL

- A standard for combining SQL with different languages.
- Key idea: Use a preprocessor to turn SQL statements into procedure calls that fit with the host-language code surrounding.
- All embedded SQL statements begin with EXEC SQL, so the preprocessor can find them easily.

```
EXEC SQL INSERT Studio(name, address)
VALUES (:studioName, :studioAddr);
```

## API Approach


- Instead of making use of a preprocessor to translate EXEC SQL ... statements into function calls (as in embedded SQL), the programmer uses a library of functions or classes and calls them as part of an ordinary C or Java program.
- SQL commands are sent to the DBMS at runtime.
- These API's are based on a standard called SQL/CLI = "Call-Level Interface."
  - C + CLI
  - Java + JDBC
  - PHP + PEAR/DB

## Open Interface: JDBC

- Database specific API makes the program dependent to one DBMS.
- Open interfaces solve this problem:
  - Designed by a third party, e.g. the creator of the language.
  - Implemented by DBMS vendors.
  - Used by DB programmers.
- Programmers do not have to change their code to work with another DBMS.

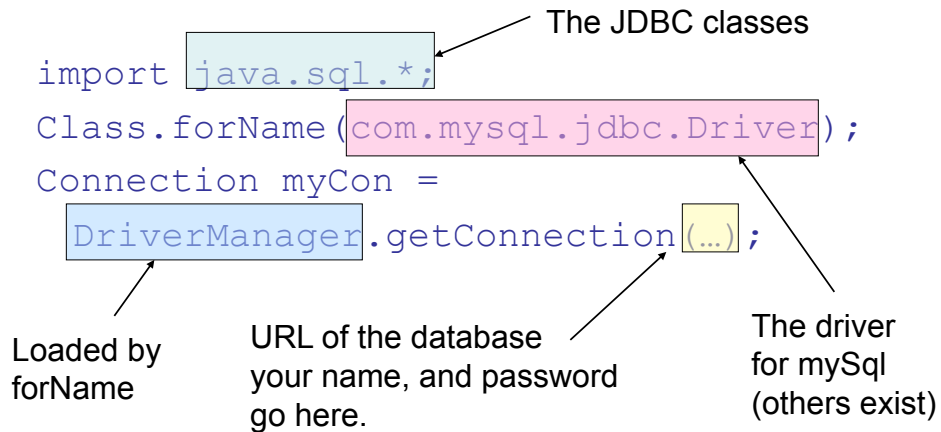
<http://java.sun.com/products/jdbc/overview.html>

## Environment, Connection, Statements



## Making a Connection

A *connection object* is obtained from the environment in a somewhat implementation-dependent way.



## Connecting to CIS MySQL DB

```
Connection myCon =
    DriverManager.getConnection
        (jdbc:mysql://mysql.cis.ksu.edu/dcaragea,
         dcaragea, mypassword);
```



## JDBC Statements

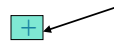
- JDBC provides two classes:
  - `Statement` - an object that can accept a string that is an SQL statement and can execute such a string.
  - `PreparedStatement` - an object that has an associated SQL statement ready to execute.
- **Why `PreparedStatement`? Performance!**
  - The SQL command is stored in DBMS after the first call.

## Creating Statements

The `Connection` class has methods to create `Statement` and `PreparedStatement`.

```
Statement stat1 = myCon.createStatement();
PreparedStatement stat2 =
    myCon.prepareStatement(
        "SELECT Address FROM Supplier "
        "WHERE Supplier_Name = 'John Smith'"
    );
```

String concatenation



`Supplier(Supplier_ID, Supplier_Name, Address)`

## Executing SQL Statements

- JDBC distinguishes **queries** from **modifications** (which it calls “updates”).
- `Statement` and `PreparedStatement` have methods `executeQuery` and `executeUpdate`.
  - For `Statement`, these methods have one argument: the query or modification to be executed.
  - For `PreparedStatement`: no argument.
- Programmer handles errors using `SQLException` class.

## Example: Update

- `stat1` is a `Statement`.
- We can use it to insert a tuple as:

```
stat1.executeUpdate(  
    "INSERT INTO Supplier " +  
    "VALUES('S4', 'Mary', '12 Goodwin St.')" ) ;
```

## Example: Query

- `stat2` is a `PreparedStatement` holding the query  
“`SELECT Address FROM Supplier WHERE  
Supplier_Name = 'John Smith'`”
- `executeQuery` returns an object of class `ResultSet`.  
`ResultSet List = stat2.executeQuery();`

## Accessing the ResultSet

- An object of type `ResultSet` is a cursor.
- Method `next()` advances the “cursor” to the next tuple.
  - The first time `next()` is applied, it gets the first tuple.
  - If there are no more tuples, `next()` returns the value `FALSE`.

## Accessing Components of Tuples

- When a `ResultSet` is referring to a single tuple, we can get the components of that tuple by applying certain methods to the `ResultSet`.
- Method `getX(i)`, where `X` is some type, and `i` is the component number, returns the value of that component.
  - The value must have type `X`.

## Example: Accessing Components

- `List` is the `ResultSet` for the query “`SELECT Address FROM Supplier WHERE Supplier_Name = 'John Smith'`”.
- Access the address from each tuple by:

```
while(List.Next()) {  
    theAddress = List.getString(1);  
    /*do something with the address */  
}
```

## Example: Accessing Components

- Menu is the ResultSet for the query

```
"SELECT beer, price FROM Sells WHERE  
bar = 'Joe' 's Bar'"
```

- Access beer and price from each tuple by:

```
while (Menu.next()) {  
    theBeer = Menu.getString(1);  
    thePrice = Menu.getFloat(2);  
    /*something with theBeer and thePrice*/  
}
```

## Parameter Passing

We can use “?” as “parameters” of a query and bind values to those parameters.

```
PreparedStatement stat3 =  
    myCon.prepareStatement(  
        "INSERT INTO Supplier " +  
        "VALUES (?, ?, ?)");  
/* get id, name, address, e.g. from user*/  
stat3.setInt(1, id);  
stat3.setString(2, name);  
stat3.setString(3, address);  
stat3.executeUpdate();
```

# SQL Injection

SQL Injection Attacks by Example

<http://www.unixwiz.net/techtips/sql-injection.html>

## What is an SQL Injection Attack?

- Many web applications take user input from a form
- Often this user input is used literally in the construction of an SQL query submitted to a database. For example:
  - `SELECT productdata FROM table WHERE productname = 'user input product name';`
- An SQL injection attack involves placing SQL statements in the user input

## An Example SQL Injection Attack

Product Search:

- The application
  - Reads the input into a variable `prodname`
  - The input is put directly into the SQL statement

```
String query = "SELECT prodinfo FROM proddtable WHERE prodname = '"
+ prodname + "'";
```
- Creates the following SQL query:

```
SELECT prodinfo FROM proddtable WHERE prodname = 'blah' OR 'x' = 'x';
```

  - Attacker has now successfully caused the entire database to be returned.

## A More Malicious Example

- What if the attacker had instead entered:

```
blah'; DROP TABLE prodinfo; --
```
- Results in the following SQL:

```
SELECT prodinfo FROM proddtable WHERE prodname = 'blah';
DROP TABLE prodinfo; --'
```

  - Note how the comment (--) consumes the final quote.
- Causes the entire table to be deleted
  - Depends on knowledge of table name
  - This is sometimes exposed to the user in debug code called during a database error
  - Use non-obvious table names, and never expose them to user

## Other Injection Possibilities

- Using SQL injections, attackers can:
  - Add new data to the database
    - Perform an INSERT in the injected SQL
  - Modify data currently in the database
    - Perform an UPDATE in the injected SQL
  - Often can gain access to other user's system capabilities by obtaining their password

## Defenses

- Check syntax of input for validity
  - Rather than "remove known bad data", it's better to "remove everything but known good data"
  - Many classes of input have fixed languages
    - Email addresses, dates, part numbers, etc.
    - Verify that the input is a valid string in the language
    - If you can exclude quotes and semicolons that's good
  - Not always possible: consider the name Bill O'Reilly
    - Want to allow the use of single quotes in names
- Have length limits on input
  - Many SQL injection attacks depend on entering long strings



## More Defenses

- Scan query string for undesirable word combinations that indicate SQL statements
  - INSERT, DROP, etc.
  - If you see these, you can check against SQL syntax to see if they represent a statement or valid user input
- Limit database permissions and segregate users
  - If you're only reading the database, connect to database as a user that only has read permissions
  - Never connect as a database administrator in your web application
- Configure database error reporting
  - Default error reporting often gives away information that is valuable for attackers (table name, field name, etc.)
  - Configure so that this information is never exposed to a user

## Even More Defenses

- Use provided functions for escaping strings
  - MySQL's library function `mysql_real_escape_string` prepends backslashes to the following characters: `\x00`, `\n`, `\r`, `\`, `'`, `"` and `\x1a`.
- Not a silver bullet!
  - Consider:
    - `SELECT fields FROM table WHERE id = 23 OR 1=1`
    - No quotes here!

## Best Defense

If possible, use bound variables (and prepared statements)

Insecure version:

```
Statement s = connection.createStatement();  
ResultSet rs = s.executeQuery("SELECT email FROM member WHERE name = "  
    + formField);
```

Secure version:

```
PreparedStatement ps = connection.prepareStatement(  
    "SELECT email FROM member WHERE name = ?");  
ps.setString(1, formField);  
ResultSet rs = ps.executeQuery();
```