

CIS 560 – Database System Concepts

Lecture 20

Concurrency Control

October 21, 2013

Credits for slides: Chang, Ullman, Whitehead.

Copyright: Caragea, 2013.

Outline

Last:

- Undo logging 17.2
- Redo logging 17.3
- Redo/undo 17.4

Today:

- Serial and serializable schedules 18.1
- Conflict serializability 18.2

Next:

- Locks 18.3

2

Concurrency Control: The Problem

- Multiple concurrent transactions T_1, T_2, \dots
- They read/write common elements A_1, A_2, \dots
- How can we prevent unwanted interference ?

The SCHEDULER is responsible for that

3

Some Famous Anomalies

- Dirty reads
- Inconsistent reads
- Unrepeatable reads
- Lost updates

Many other things can go wrong too

4

Conflicts

- Write-Read – WR
- Read-Write – RW
- Write-Write – WW

5

Dirty Reads

Write-Read Conflict

T_1 : WRITE(A)

T_1 : ABORT

T_2 : READ(A)

6

Inconsistent Read

Write-Read Conflict

T_1 : A := 20; B := 20;
 T_1 : WRITE(A)

T_1 : WRITE(B)

T_2 : READ(A);
 T_2 : READ(B);

7

Unrepeatable Read

Read-Write Conflict

T_1 : WRITE(A)

T_2 : READ(A);

T_2 : READ(A);

8

Lost Update

Write-Write Conflict

T_1 : READ(A)

T_1 : $A := A + 5$

T_1 : WRITE(A)

T_2 : READ(A);

T_2 : $A := A * 1.3$

T_2 : WRITE(A);

9

Schedules

- Given multiple transactions

A *schedule* is a sequence of interleaved actions from all transactions

10

Example

T1	T2
READ(A, t)	READ(A, s)
t := t+100	s := s*2
WRITE(A, t)	WRITE(A,s)
READ(B, t)	READ(B,s)
t := t+100	s := s*2
WRITE(B,t)	WRITE(B,s)

11

A Serial Schedule

T1	T2	A	B
READ(A, t)		25	25
t := t+100			
WRITE(A, t)		125	
READ(B, t)			
t := t+100			125
WRITE(B,t)			
	READ(A,s)		
	s := s*2		
	WRITE(A,s)	250	
	READ(B,s)		
	s := s*2		
	WRITE(B,s)		250

Serial schedule: (T1, T2)

12

A Serial Schedule

T1	T2	A	B
	READ(A,s)	25	25
	s := s*2		
	WRITE(A,s)	50	
	READ(B,s)		
	s := s*2		50
	WRITE(B,s)		
READ(A, t)			
t := t+100			
WRITE(A, t)		150	
READ(B, t)			
t := t+100			
WRITE(B,t)			150

Serial schedule: (T2, T1)

13

Serializable Schedule

A schedule is serializable if it is equivalent to a serial schedule

A schedule S is serializable, if there is a serial schedule S', such that for every initial database state, the effects of S and S' are the same

14

A Serializable Schedule

T1	T2	A	B
READ(A, t)		25	25
t := t+100			
WRITE(A, t)		125	
	READ(A,s)		
	s := s*2		
	WRITE(A,s)	250	
READ(B, t)			
t := t+100			
WRITE(B,t)			125
	READ(B,s)		
	s := s*2		
	WRITE(B,s)		250

This is NOT a serial schedule

15

A Non-Serializable Schedule

T1	T2	A	B
READ(A, t)		25	25
t := t+100			
WRITE(A, t)		125	
	READ(A,s)		
	s := s*2		
	WRITE(A,s)	250	
	READ(B,s)		
	s := s*2		
	WRITE(B,s)		50
READ(B, t)			
t := t+100			
WRITE(B,t)			150

16

Transaction Semantics

T1	T2	A	B
READ(A, t)		25	25
t := t+100			
WRITE(A, t)		125	
	READ(A,s)		
	s := s+200		
	WRITE(A,s)	325	
	READ(B,s)		
	s := s+200		
	WRITE(B,s)		225
READ(B, t)			
t := t+100			
WRITE(B,t)			325

Is this serializable?

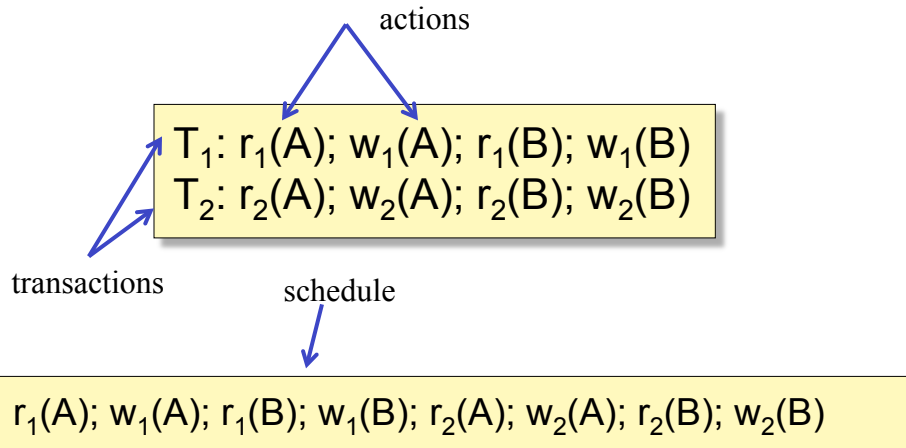
17

Ignoring Details

- Serializability is undecidable !
- Scheduler should not look at transaction details
- Assume worst case updates
 - Only care about reads $r(A)$ and writes $w(A)$
 - Not the actual values involved

18

Notation



19

Conflict Serializability

Conflict: pair of consecutive actions in schedule
s.t. if swapped, then behavior changes

Two actions by same transaction T_i :

$r_i(X); w_i(Y)$

Two writes by T_i, T_j to same element

$w_i(X); w_j(X)$

Read/write by T_i, T_j to same element

$w_i(X); r_j(X)$

$r_i(X); w_j(X)$

20

Conflict Serializability

- A schedule is conflict serializable if it can be transformed into a serial schedule by a series of swappings of adjacent non-conflicting actions

Example:

$r_1(A); w_1(A); r_2(A); w_2(A); r_1(B); w_1(B); r_2(B); w_2(B)$



$r_1(A); w_1(A); r_1(B); w_1(B); r_2(A); w_2(A); r_2(B); w_2(B)$

21

Example

$r_1(A); w_1(A); r_2(A); w_2(A); r_1(B); w_1(B); r_2(B); w_2(B)$

$r_1(A); w_1(A); r_2(A); r_1(B); w_2(A); w_1(B); r_2(B); w_2(B)$

$r_1(A); w_1(A); r_1(B); r_2(A); w_2(A); w_1(B); r_2(B); w_2(B)$

$r_1(A); w_1(A); r_1(B); r_2(A); w_1(B); w_2(A); r_2(B); w_2(B)$

$r_1(A); w_1(A); r_1(B); w_1(B); r_2(A); w_2(A); r_2(B); w_2(B)$

22

The Precedence Graph Test

Is a schedule conflict-serializable?


Simple test:

- Build a graph of all transactions T_i
- Edge from T_i to T_j if T_i makes an action that conflicts with one of T_j and comes first
 - T_i must *precede* T_j iff an action from T_i conflicts with a later action from T_j
- The test: if the graph has no cycles, then it is conflict serializable!

23

Example 1

$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$




This schedule is conflict-serializable

Any transaction order which follows the precedences shown in the graph gives an equivalent serial schedule.

24

Example 2

$r_2(A); r_1(B); w_2(A); r_2(B); r_3(A); w_1(B); w_3(A); w_2(B)$



This schedule is NOT conflict-serializable

25

Exercise: construct precedence graph

$w_3(A)$

$r_1(A)$

$w_1(B)$


$r_2(B)$

$w_3(C)$

$r_2(C)$

26


Exercise: construct precedence graph



Is conflict-serializable?

27

Exercise: construct precedence graph




T_3, T_1, T_2

28

Serial Equivalent

T_3, T_1, T_2



29


Home exercise

$r_1(A)$
 $r_2(A)$
 $r_1(B)$
 $r_2(B)$
 $r_3(A)$
 $r_4(B)$
 $w_1(A)$
 $w_2(B)$

30

Conflict Serializability

- A serializable schedule need not be conflict serializable, even under the “worst case update” assumption



So, conflict serializability is a stronger condition.

31

Scheduler

- The scheduler is the module that schedules the transaction's actions, ensuring serializability.
- How? Three main techniques:
 - Locks
 - Timestamps
 - Validation

32

Locking Scheduler

Simple idea:

- Each element has a unique lock
- Each transaction must first acquire the lock before reading/writing that element
- If the lock is taken by another transaction, then wait
- The transaction must release the lock(s)

33

Notation

$L_i(A)$ = transaction T_i acquires lock for element A

$U_i(A)$ = transaction T_i releases lock for element A

34

Example

T1	T2
$L_1(A)$; READ(A, t) $t := t+100$ WRITE(A, t); $U_1(A)$; $L_1(B)$	$L_2(A)$; READ(A, s) $s := s*2$ WRITE(A, s); $U_2(A)$; $L_2(B)$; DENIED...
READ(B, t) $t := t+100$ WRITE(B, t); $U_1(B)$;	... GRANTED ; READ(B, s) $s := s*2$ WRITE(B, s); $U_2(B)$;

Is this a conflict-serializable schedule?

35

Example

T1	T2
$L_1(A)$; READ(A, t) $t := t+100$ WRITE(A, t); $U_1(A)$; $L_1(B)$	$L_2(A)$; READ(A, s) $s := s*2$ WRITE(A, s); $U_2(A)$; $L_2(B)$; DENIED...
READ(B, t) $t := t+100$ WRITE(B, t); $U_1(B)$;	... GRANTED ; READ(B, s) $s := s*2$ WRITE(B, s); $U_2(B)$;

Scheduler has ensured a conflict-serializable schedule

36

Example

T1	T2
$L_1(A)$; READ(A, t) $t := t+100$ WRITE(A, t); $U_1(A)$; $L_1(B)$; READ(B, t) $t := t+100$ WRITE(B,t); $U_1(B)$;	$L_2(A)$; READ(A,s) $s := s*2$ WRITE(A,s); $U_2(A)$; $L_2(B)$; READ(B,s) $s := s*2$ WRITE(B,s); $U_2(B)$;

Is this a conflict-serializable schedule?

37

Example

T1	T2
$L_1(A)$; READ(A, t) $t := t+100$ WRITE(A, t); $U_1(A)$; $L_1(B)$; READ(B, t) $t := t+100$ WRITE(B,t); $U_1(B)$;	$L_2(A)$; READ(A,s) $s := s*2$ WRITE(A,s); $U_2(A)$; $L_2(B)$; READ(B,s) $s := s*2$ WRITE(B,s); $U_2(B)$;

Locks did not enforce conflict-serializability!!!

38

Two Phase Locking (2PL)

The 2PL rule:

- In every transaction, all lock requests must precede all unlock requests
- This ensures conflict serializability! (why?)

39

Example: 2PL transactions

T1	T2
$L_1(A); L_1(B); \text{READ}(A, t)$ $t := t+100$ $\text{WRITE}(A, t); U_1(A)$	
	$L_2(A); \text{READ}(A, s)$ $s := s*2$ $\text{WRITE}(A, s);$ $L_2(B); \text{DENIED...}$
$\text{READ}(B, t)$ $t := t+100$ $\text{WRITE}(B, t); U_1(B);$	
	$\text{...GRANTED}; \text{READ}(B, s)$ $s := s*2$ $\text{WRITE}(B, s); U_2(A); U_2(B);$

Now it is conflict-serializable

40