

CIS 560 – Database System Concepts

Lecture 21

# Concurrency Control

October 23, 2013

Credits for slides: Chang, Ullman, Whitehead.

Copyright: Caragea, 2013.

## Outline

Last:

- Serial and serializable schedules 18.1
- Conflict serializability 18.2
- Locks 18.3

Today:

- Locks 18.3
- Timestamps 18.8

Next:

- Indexes and B-trees 14.1-14.2

## Review

- Schedule
- Serial schedule
- Serializable schedule
- Conflict serializable schedule
- Precedence graph
- Locks
- Two phase locking (2PL)

3

## What about Aborts?

- 2PL enforces conflict-serializable schedules
- But what if a transaction releases its locks and then aborts?
- Serializable schedule definition only considers transactions that commit
  - Relies on assumptions that aborted transactions can be undone completely

4

## Example with Abort

T1	T2
L <sub>1</sub> (A); L <sub>1</sub> (B); READ(A, t)	
t := t+100	
WRITE(A, t); U <sub>1</sub> (A)	
	L <sub>2</sub> (A); READ(A,s)
	s := s*2
	WRITE(A,s);
	L <sub>2</sub> (B); <b>DENIED...</b>
READ(B, t)	
t := t+100	
WRITE(B,t); U <sub>1</sub> (B);	
	<b>...GRANTED;</b> READ(B,s)
	s := s*2
	WRITE(B,s); U <sub>2</sub> (A); U <sub>2</sub> (B);
<b>Abort</b>	<b>Commit</b>

5

## Recoverable Schedules

- A schedule is *recoverable* if whenever a transaction T commits, all transactions who have written elements read by T have already committed.

6

## Is this schedule recoverable?

T1	T2
$L_1(A); L_1(B); \text{READ}(A, t)$ $t := t+100$ $\text{WRITE}(A, t); U_1(A)$	
	$L_2(A); \text{READ}(A, s)$ $s := s*2$ $\text{WRITE}(A, s);$ $L_2(B); \text{DENIED}...$
$\text{READ}(B, t)$ $t := t+100$ $\text{WRITE}(B, t); U_1(B);$	
	$... \text{GRANTED}; \text{READ}(B, s)$ $s := s*2$ $\text{WRITE}(B, s); U_2(A); U_2(B);$
Commit	Commit

7


## Cascading Aborts

- If a transaction T aborts, then we need to abort any other transaction T' that has read an element written by T.
- A schedule is said to *avoid cascading aborts* if whenever a transaction read an element, the transaction that has last written it has already committed.

8

## Strict 2PL

- Strict 2PL: All locks held by a transaction are released when the transaction is completed.
- Ensures that schedules are **recoverable**
  - Transactions commit only after all transactions whose changes they read also commit.
- **Avoids cascading rollbacks.**



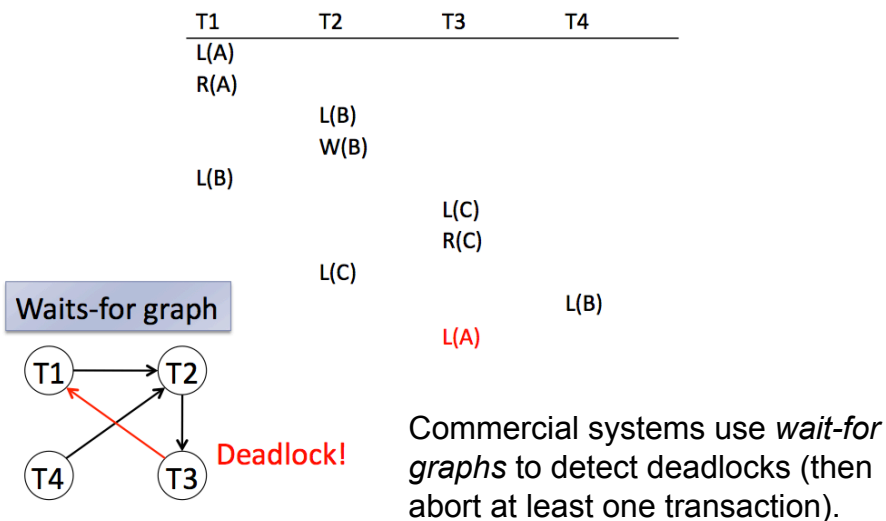
9

## Deadlock

- Transaction  $T_1$  waits for a lock held by  $T_2$ ;
- But  $T_2$  waits for a lock held by  $T_3$ ;
- While  $T_3$  waits for . . . .
- . . .
- . . .and  $T_3$  waits for a lock held by  $T_1$ !
- Could be prevented/detected (see textbook 19.2);

10

## Deadlock Example



11

## Lock Modes

- S = shared lock (for READ)
- X = exclusive lock (for WRITE)
- U = update lock
  - Initially like S
  - Later may be upgraded to X
- I = increment lock (for  $A := A + \text{something}$ )
  - Increment operations commute

Recommended reading 18.4!

12

## Lock Granularity

- **Fine granularity locking** (e.g., tuples)
  - High concurrency
  - High overhead in managing locks
- **Coarse grain locking** (e.g., tables)
  - Many false conflicts
  - Less overhead in managing locks
- Alternative techniques
  - Hierarchical locking (and intentional locks) [commercial DBMSs]

Recommended reading 18.6!

13

## The Locking Scheduler

Task 1:

Add lock/unlock requests to transactions

- Examine all READ(A) or WRITE(A) actions
- Add appropriate lock requests
- Ensure Strict 2PL!

14

## The Locking Scheduler

### Task 2:

Execute the locks accordingly

- Lock table: a big, critical data structure in a DBMS!
- When a lock is requested, check the lock table
  - Grant, or add the transaction to the element's wait list
- When a lock is released, re-activate a transaction from its wait list
- When a transaction aborts, release all its locks
- Check for deadlocks occasionally

Recommended reading 18.5!

15

## Phantom Problem [18.6.3]

- So far we have assumed the database to be a *static* collection of elements (=tuples)
- If tuples are inserted/deleted then the *phantom problem* appears

16



## Phantom Problem

T1	T2
SELECT *	
FROM Product	
WHERE color='blue'	
	INSERT INTO Product(name, color)
	VALUES ('gizmo','blue')
SELECT *	
FROM Product	
WHERE color='blue'	

Suppose there are two blue products, X1, X2:

R1(X1),R1(X2),W2(X3),R1(X1),R1(X2),R1(X3)

Conflict serializable! But not serializable due to phantoms

## Dealing with Phantoms

- In a **static** database:
  - Conflict serializability implies serializability
- In a **dynamic** database, this may fail due to phantoms
- Strict 2PL guarantees conflict serializability, but not serializability
- Expensive ways of dealing with phantoms:
  - Lock the entire table, or
  - Lock the index entry for 'blue' (if index is available)

Serializable transactions are very expensive <sup>8</sup>

## Concurrency Control Mechanisms

- Pessimistic:
  - *Intuition: assume that things will go wrong unless transactions are prevented in advance from engaging in nonserializable behavior.*
    - Locks
- Optimistic
  - *Intuition: assume that no unserializable behavior will occur and only fix things up when a violation is apparent!*
    - Timestamp based: basic, multiversion
    - Validation
    - Snapshot isolation: a variant of both

19

## Timestamps

- Each transaction receives a unique timestamp  $TS(T)$ , when the transaction first notifies the scheduler that it is beginning.
- Could be:
  - The system's clock
  - A unique counter, incremented by the scheduler

20

## Timestamps

Main invariant:

The timestamp order defines the serialization order of the transactions.

21

## Main Idea

- For any two conflicting actions, ensure that their order is the serialized order:

In each of these cases

- $w_U(X) \dots r_T(X)$
  - $r_U(X) \dots w_T(X)$
  - $w_U(X) \dots w_T(X)$
- } Possible conflicts

When T requests  $r_T(X)$  or  $w_T(X)$ , need to check  $TS(U) \leq TS(T)$

22

## Timestamps

With each element  $X$ , associate

- $RT(X)$  = the highest timestamp of any transaction  $U$  that read  $X$
- $WT(X)$  = the highest timestamp of any transaction  $U$  that wrote  $X$
- $C(X)$  = the commit bit: true when transaction with highest timestamp that wrote  $X$  committed

If 1 element = 1 page, then these are associated with each page  $X$  in the buffer pool

23

## Ensuring Recoverable Schedules

- Recall the definition: if a transaction reads an element, then the transaction that wrote it must have already committed
- Use the commit bit  $C(X)$  to keep track if the transaction that last wrote  $X$  has committed

24

## Simplified Timestamp-based Scheduling

Note: simple version that ignores the commit bit

- Only for transactions that do not abort
- Otherwise, may result in non-recoverable schedule

### Transaction wants to read element X

If  $TS(T) < WT(X)$  then ROLLBACK  
Else READ and update  $RT(X)$  to larger of  $TS(T)$  or  $RT(X)$

### Transaction wants to write element X

If  $TS(T) < RT(X)$  then ROLLBACK  
Else if  $TS(T) < WT(X)$  ignore write & continue (Thomas Write Rule)  
Otherwise, WRITE and update  $WT(X) = TS(T)$

25

## Details

Read too late:

- T wants to read X, and  $TS(T) < WT(X)$

START(T) ... START(U) ...  $w_U(X)$  ...  $r_T(X)$

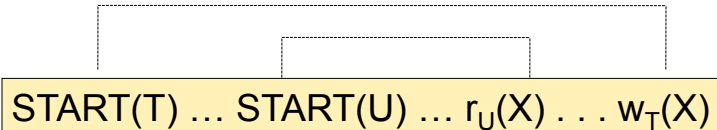
Need to rollback T!

26

## Details

Write too late:

- T wants to write X, and  $TS(T) < RT(X)$



START(T) ... START(U) ...  $r_U(X)$  ...  $w_T(X)$

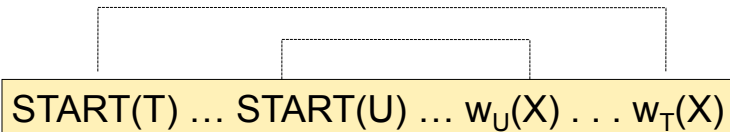
Need to rollback T!

27

## Details

Write too late, but we can still handle it:

- T wants to write X, and  
 $TS(T) \geq RT(X)$  but  $WT(X) > TS(T)$



START(T) ... START(U) ...  $w_U(X)$  ...  $w_T(X)$

Don't write X at all !  
 (Thomas' rule)

28

T1	T2	T3	A	B	C
200	150	175	RT =0 WT=0	RT =0 WT=0	RT =0 WT=0
$r_1(B)$					
	$r_2(A)$				
		$r_3(C)$			
$w_1(B)$					
$w_1(A)$					
	$w_2(C)$				
		$w_3(A)$			

29

T1	T2	T3	A	B	C
200	150	175	RT =0 WT=0	RT =0 WT=0	RT =0 WT=0
$r_1(B)$				RT=200	
	$r_2(A)$		RT=100		
		$r_3(C)$			RT=175
$w_1(B)$				WT=200	
$w_1(A)$			WT=200		
	$w_2(C)$ <b>rollback</b>				
		$w_3(A)$			

30