

CIS 560 – Database System Concepts

Lecture 23

# Indexes and B<sup>+</sup> Trees

October 28, 2013

Credits for slides: Chang, Ullman, Whitehead.

Copyright: Caragea, 2013.

## Outline

Last:

- Timestamps 18.8
- Indexes

Today:

- Indexes and B-trees 14.1-14.2

Next:

- Query execution 15.1-15.6
- Query optimization 16

## Index Classification


- Clustered/unclustered
  - Clustered = records close in index are close in data
  - Unclustered = records close in index might be far in data
- Primary/secondary:
  - Primary = on primary key
  - Secondary = on any other key
- Dense/sparse
  - Dense = each record has an entry in the index
  - Sparse = only some records have
- Organization: B+ tree or Hashtable

## Clustered/Unclustered

- Clustered
  - Index determines the location of indexed records
  - Typically, a clustered index is one where values are data records (but not necessary)
- Unclustered
  - Index does not determine data location
  - In these indexes: value = pointer to data record


## Clustered, Dense Index

- File is sorted on the index attribute
- Only one per table




## Clustered, Sparse Index

- Sparse index: one key per data block




## Unclustered Index

- Several per table
- Often for indexing other attributes than primary key
- Always dense (why?)



## Clustered vs. Unclustered Index




In a clustered B+ tree index,  
data entries can be data records

## Q: Our textbook as example: Indexes?



- How many indexes? Where?
- What are keys? What are records?
- Clustered?
- Dense?
- Primary?



## Example

| id  | age | salary | other |
|-----|-----|--------|-------|
| 006 | 19  | 50k    | ...   |
| 005 | 20  | 55k    | ...   |
| 004 | 25  | 50k    | ...   |
| 007 | 30  | 80k    | ...   |
| 002 | 35  | 75k    | ...   |
| 003 | 35  | 70k    | ...   |
| 001 | 40  | 65k    | ...   |

| id  | age | salary | other |
|-----|-----|--------|-------|
| 006 | 19  | 50k    | ...   |
| 004 | 25  | 50k    | ...   |
| 005 | 20  | 55k    | ...   |
| 001 | 40  | 65k    | ...   |
| 003 | 35  | 70k    | ...   |
| 002 | 35  | 75k    | ...   |
| 007 | 30  | 80k    | ...   |

data file = index file  
clustered, dense index

separate index file  
unclustered, dense index


## Hash-Based Index Example 1

Hash-based index on id

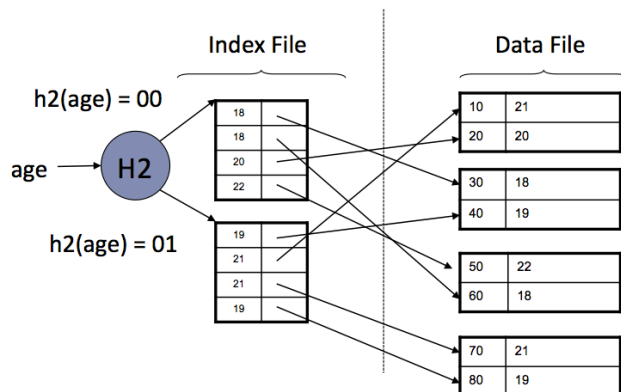
This is a **primary** index

In this case, data entries in the index are actual data records - there is no separate data file

This index is also **clustered** because it determines the order of indexed records



## Hash-Based Index Example 2




**Secondary** index

Data entries in index are (key, recordID) pairs

**Unclustered** index

## Hash-Based Indexes




Good for point queries but not range queries

## B+ Trees


- Search trees
- Idea in B Trees:
  - make 1 node = 1 block
  - keep tree **balanced** in height
- Idea in B+ Trees:
  - make leaves into a linked list to facilitate range queries

## B+ Trees Basics

- Parameter  $d$  = the degree
- Each internal node has  $d \leq m \leq 2d$  keys (except root)

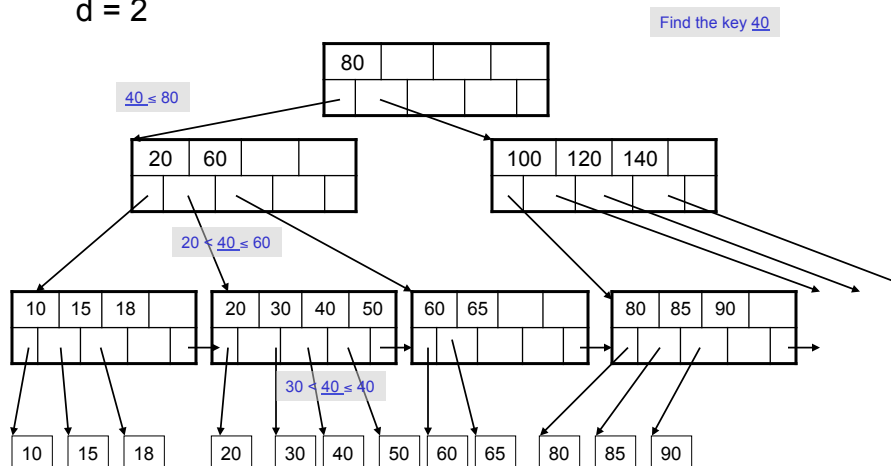


- Each leaf has  $d \leq m \leq 2d$  keys:



## B+ Tree Example

$d = 2$





## Using a B+ Tree

- Exact key values:
  - Start at the root
  - Proceed down, to the leaf

Index on People(age)

```
Select name
From People
Where age = 25
```

- Range queries:
  - As above
  - Then sequential traversal

```
Select name
From People
Where 20 <= age
and age <= 30
```

## Which queries can use this index?

Index on People(name, zipcode)

```
Select *
From People
Where name = 'Smith'
and zipcode = 12345
```

```
Select *
From People
Where name = 'Smith'
```

```
Select *
From People
Where zipcode = 12345
```

## B+ Tree Design

- How large d?
- Example:
  - Key size = 4 bytes
  - Pointer size = 8 bytes
  - Block size = 4096 bytes

## B+ Tree Design

- How large d?
- Example:
  - Key size = 4 bytes
  - Pointer size = 8 bytes
  - Block size = 4096 bytes
- $2d \times 4 + (2d+1) \times 8 \leq 4096$
- $d = 170$


## B+ Trees in Practice

- Average degree (fanout) = 133. Typical order (max number of children): 100. Typical fill-factor: 67%.
- Typical capacities:
  - Height 4:  $133^4 = 312,900,700$  records
  - Height 3:  $133^3 = 2,352,637$  records
- Can often hold top levels in buffer pool:
  - Level 1 = 1 page ~ 8 Kbytes
  - Level 2 = 133 pages ~ 1 Mbyte
  - Level 3 = 17,689 pages ~133 MBytes

## Insertion in a B+ Tree

Insert (K, P)


- Find leaf where K belongs, insert
- If no overflow ( $2d$  keys or less), halt
- If overflow ( $2d+1$  keys), split node, insert in parent:



- If leaf, keep K3 too in right node
- When root splits, new root has 1 key only

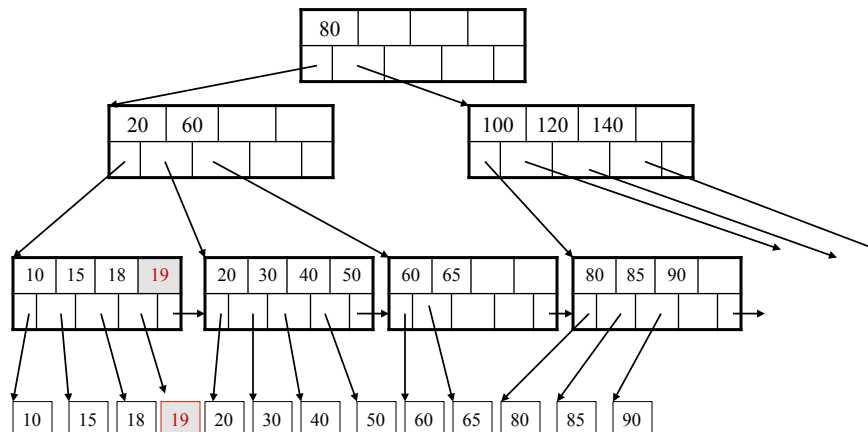
## Insertion in a B+ Tree

Insert K=19




## Insertion in a B+ Tree

After insertion



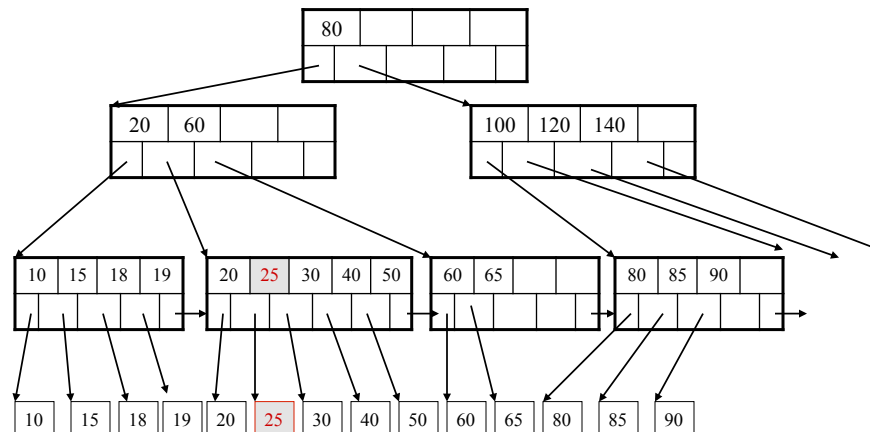
## Insertion in a B+ Tree

Now insert 25



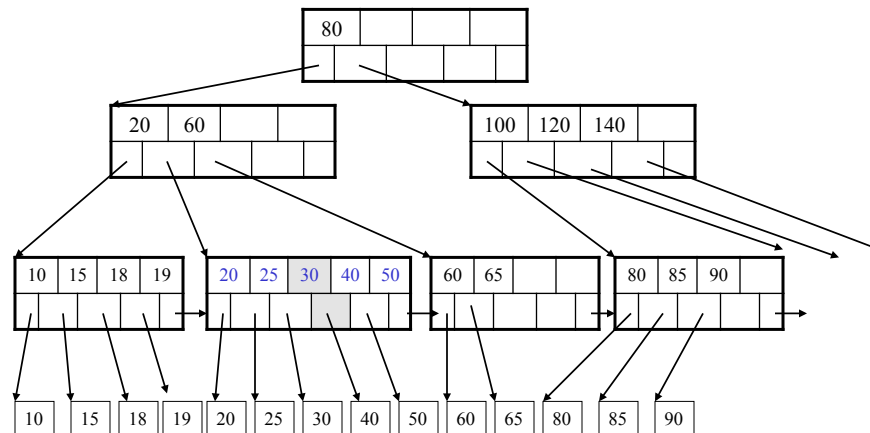
## Insertion in a B+ Tree

After insertion



## Insertion in a B+ Tree

But now have to split!



## Insertion in a B+ Tree

After the split

