

**CIS 560 – Database System Concepts**

**Lecture 25**

# Query Processing

November 1, 2013

Credits for slides: Chang, Ullman, Whitehead.

Copyright: Caragea, 2013.

## Outline

Last:

- Indexes and B-trees 14.1-14.2

Today:

- Query processing

Next:

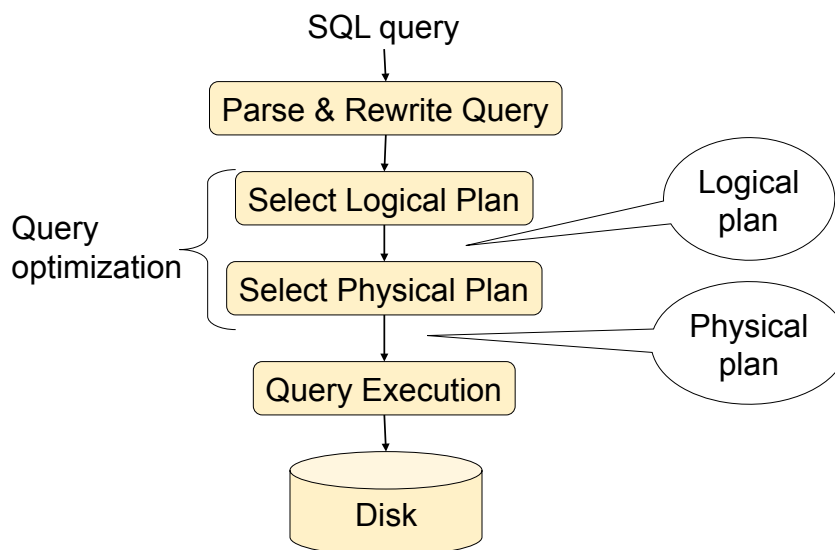
- Query execution 15.1-15.6
- Query optimization 16

## Planning

- Assignment 7 (concurrency control) due 11/1
- Assignment 8 (indexes) due 11/8
- Assignment 9 (query optimization) due 11/15
- Exam 2 (assignments 6-9) – 11/20
- Project assignment due 11/22
- Quiz from special topics – 12/06
- Project presentations 12/9, 12/11, 12/13
- Project reports – finals weeks

3

## Steps of the Query Processor



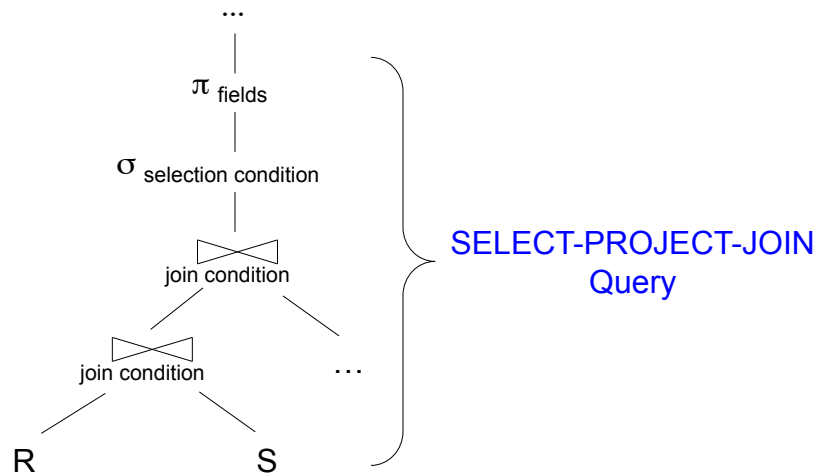
## Query Optimization

- **Step 3: Query optimization**
  - Finds an efficient query plan for executing the query
- **A query plan is**
  - **Logical query plan:** an extended relational algebra tree
  - **Physical query plan:** with additional annotations at each node
    - Access method to use for each relation
    - Implementation to use for each relational operator


## Query Block

- **Most optimizers operate on individual query blocks**
- A query block is an SQL query with **no nesting**
  - **Exactly one**
    - SELECT clause
    - FROM clause
  - **At most one**
    - WHERE clause
    - GROUP BY clause
    - HAVING clause

## Typical Plan for Block (1/2)



## Typical Plan For Block (2/2)




## Physical Query Plan

- Logical query plan with extra annotations
- **Access path selection** for each relation
  - Use a file scan or use an index
- **Implementation choice** for each operator
- **Scheduling decisions** for operators

Supplier(sno,sname,scity,sstate)  
 Part(pno,pname,psize,pcolor)  
 Supplies(sno,pno,price)

## Physical Query Plan

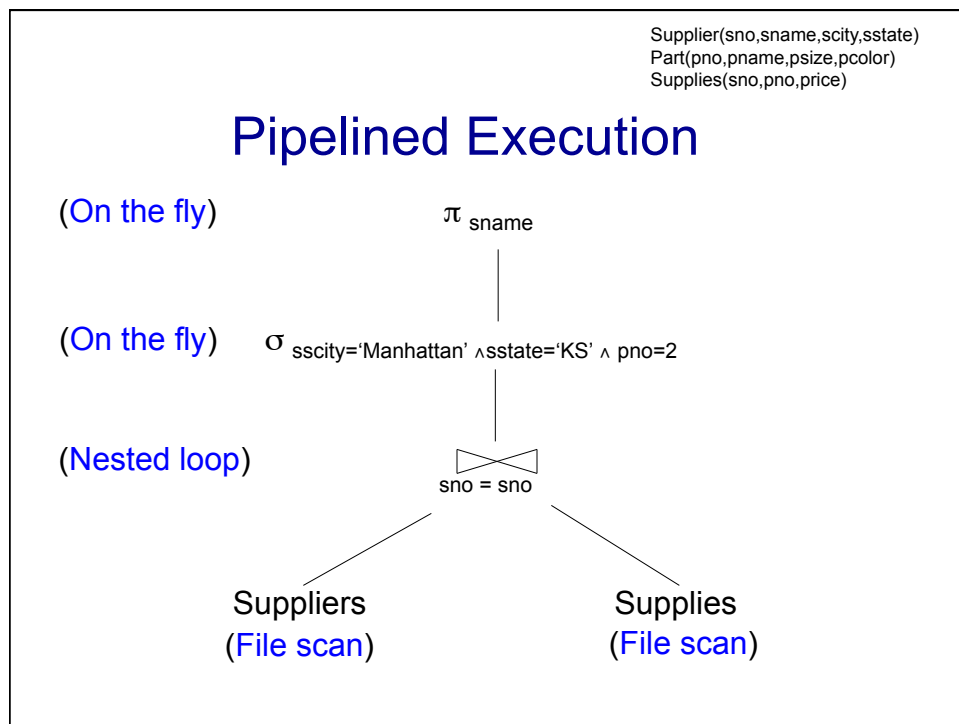


## Final Step in Query Processing

- **Step 4: Query execution**
  - How to **synchronize operators**?
  - How to **pass data between operators**?
- What techniques are possible?
  - One thread per query
  - **Iterator interface**
  - **Pipelined execution**
  - **Intermediate result materialization**

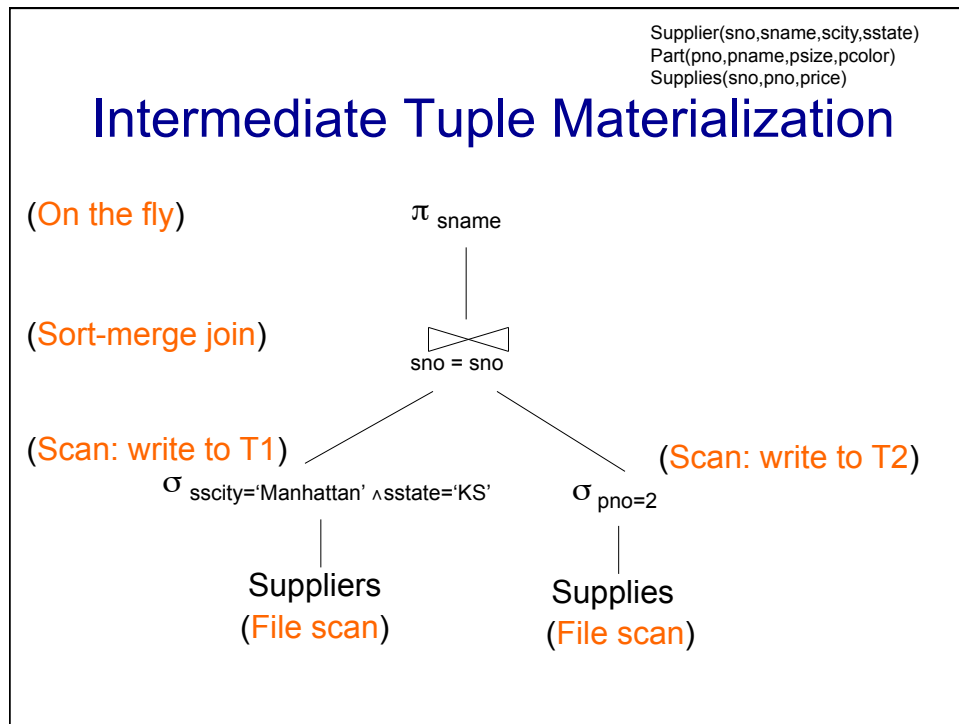
## Iterator Interface

- Each **operator implements this interface**
- Interface has only three methods
  - **open()**
    - Initializes operator state
    - Sets parameters such as selection condition
  - **get\_next()**
    - Operator invokes get\_next() recursively on its inputs
    - Performs processing and produces an output tuple
  - **close()**: cleans-up state



## Pipelined Execution

- Applies parent operator to tuples directly as they are produced by child operators
- Benefits
  - No operator synchronization issues
  - Saves cost of writing intermediate data to disk
  - Saves cost of reading intermediate data from disk
  - Good resource utilizations on single processor
- This approach is used whenever possible



## Intermediate Tuple Materialization

- Writes the results of an operator to an intermediate table on disk

No direct benefit but:

- Necessary when data is larger than main memory
- Necessary when operator needs to examine the same tuples multiple times



## Physical Operators

Each of the logical operators may have one or more implementations = physical operators

Will discuss several basic physical operators (operator algorithms), with a focus on join

## Why Learn About Operator Algorithms?

- Implemented in commercial DBMSs
  - DBMSs implement different subsets of known algorithms
- Good algorithms can greatly improve performance
- Need to know about physical operators to understand query optimization

Supplier(sno,sname,scity,sstate)  
 Part(pno,pname,psize,pcolor)  
 Supplies(sno,pno,price)

## Join Physical Operators

Logical operator:

**Supplies(sno,pno,price)  $\bowtie_{pno=pno}$  Part(pno,pname,psize,pcolor)**

Three physical operators for the join, assuming the tables are in main memory.

Nested Loop Join

Merge join

Hash join

Supplier(sno,sname,scity,sstate)  
 Part(pno,pname,psize,pcolor)  
 Supplies(sno,pno,price)

## 1. Nested Loop Join

```
for S in Supplies do {
  for P in Part do {
    if (S.pno == P.pno) output(S,P);
  }
}
```

Supplies = *outer relation*

Part = *inner relation*

Note: sometimes terminology  
is switched

20

## It's more complicated...

- Each **operator** implements this interface
- **open()**
- **get\_next()**
- **close()**

## Nested Loop Join Revisited

Supplier(sno,sname,scity,sstate)  
 Part(pno,pname,psize,pcolor)  
 Supplies(sno,pno,price)

```
open ( ) {
  Supplies.open( );
  Part.open( );
  S = Supplies.get_next( );
}
```

```
close ( ) {
  Supplies.close ( );
  Part.close ( );
}
```

```
get_next( ) {
  repeat {
    P= Part.get_next( );
    if (P== NULL)
      { Part.close();
        S= Supplies.get_next( );
        if (S== NULL) return NULL;
        Part.open( );
        P= Part.get_next( );
      }
  } until (S.pno == P.pno);
  return (S, P)
}
```

ALL operators need to be implemented this way !

## Hash-based Join

- $R \bowtie S$
- Main memory *hash-based join*:
  - Scan S, build buckets in main memory
  - Then scan R and join

Supplier(sno,sname,scity,sstate)  
 Part(pno,pname,psize,pcolor)  
 Supplies(sno,pno,price)

## 2. Hash Join (main memory)

Build  
phase

```
for S in Supplies do insert(S.pno, S);

for P in Part do {
  LS = find(P.pno);
  for S in LS do { output(S, P); }
}
```

Probing

Supplies=outer  
Part=inner

Recall: need to rewrite as open, get\_next, close

Supplier(sno,sname,scity,sstate)  
 Part(pno,pname,psize,pcolor)  
 Supplies(sno,pno,price)

### 3. Merge Join (main memory)

```
Part1 = sort(Part, pno);
Supplies1 = sort(Supplies, pno);
P=Part1.get_next(); S=Supplies1.get_next();

While (P!=NULL and S!=NULL) {
  case:
    P.pno > S.pno:  P = Supplies1.get_next( );
    P.pno < S.pno:  S = Part1.get_next();
    P.pno == S.pno { output(P,S);
                    S = Supplies1.get_next();
                  }
}
```

## Summary

- Three algorithms for main memory join:
  - Nested loop join
  - Hash join
  - Merge join

If  $|R| = m$  and  $|S| = n$ ,  
 what is the asymptotic  
 complexity for  
 computing  $R \bowtie S$  ?

## Other Main Memory Algorithms

- Grouping:  $\gamma(R)$ 
  - Nested loop
  - Hash table
  - Sorting
- Duplicate elimination:  $\delta(R)$ 
  - *Exactly* the same algorithms (why?)

How do these algorithms work, and what are their complexities?