

CIS 560 – Database System Concepts

Lecture 27

Operator Algorithms

November 6, 2013

Credits for slides: Chang, Ullman, Whitehead.

Copyright: Caragea, 2013.

Outline

- Query execution 15.1-15.6
- Query optimization 16

Planning

- Assignment 8 (indexes) due 11/8
- Project - DB design revision due 11/11
 - No class that day – use the time to work on project
- Assignment 9 (query optimization) due 11/15
- Exam 2 (assignments 6-9) – 11/20
- Project - DB implementation and queries due 11/22
- Quiz from special topics – 12/06
- Project presentations 12/9, 12/11, 12/13
- Project reports – finals week

3

Types of Algorithms

- One-pass algorithms (Sec. 15.2 and 15.3)
- Index-based algorithms (Sec 15.6)
- Two-pass algorithms (Sec 15.4 and 15.5)

Hash Join

Hash join: $R \bowtie S$

- Scan R, build buckets in main memory
- Then scan S and join
- Cost: $B(R) + B(S)$
- One-pass algorithm when $B(R) \leq M$

Sort-Merge Join

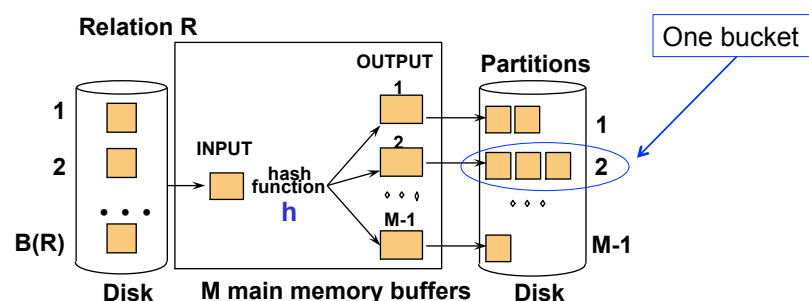
- Sort-merge join: $R \bowtie S$
 - Scan R and sort in main memory
 - Scan S and sort in main memory
 - Merge R and S
- Cost: $B(R) + B(S)$
- One pass algorithm when $B(S) + B(R) \leq M$
- Typically, this is NOT a one pass algorithm

Two-Pass Algorithms

- What if data does not fit in memory?
- Need to process it in multiple passes
- Two key techniques
 1. Hashing
 2. Sorting

Partitioned Hash Algorithms

- Idea: partition a relation R into buckets, on disk
- Each bucket has size $\approx B(R)/M$



Does each bucket fit in main memory?
 Yes, if $B(R)/M \leq M$, i.e. $B(R) \leq M^2$

Partitioned Hash Join

$R \bowtie S$

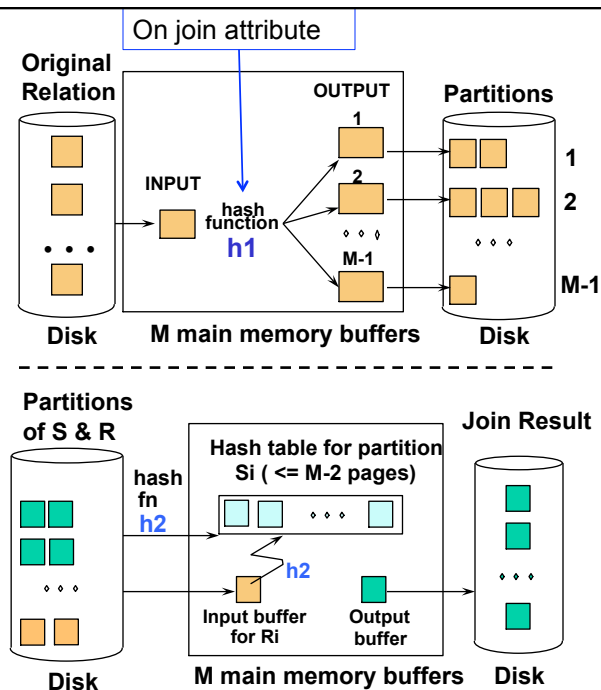
- Step 1:
 - Hash S into M-1 buckets
 - Send all buckets to disk
- Step 2
 - Hash R into M-1 buckets
 - Send all buckets to disk
- Step 3
 - Join every pair of buckets

Cost: $3B(R) + 3B(S)$

Assumption: $\min(B(R), B(S)) \leq M^2$

Partitioned Hash-Join

- Partition both relations using hash function h_1 :
R tuples in partition i will only match S tuples in partition i .
- Read in a partition of S, hash it using h_2 ($\neq h_1$). Scan matching partition of R, search for matches.



Cost of Partitioned Hash Join


- Read+write+read = $3B(R) + 3B(S)$
- Assumption: $\min(B(R), B(S)) \leq M^2$

External Sorting (Two-Phase, Multiway Merge Sort)

- Problem:
 - Sort a file of size B with memory M
- Where we need this:
 - ORDER BY in SQL queries
 - Several physical operators
 - Bulk loading of B+-tree indexes.
- Sorting is two-pass when $B \leq M^2$


External Merge-Sort: Step 1

- Load M pages in memory, sort



External Merge-Sort: Step 2

- Merge $M - 1$ runs (sorted sublists) into a new run (sorted list)
 - $M-1$ runs of length M : $M(M - 1) \approx M^2 \Rightarrow B \leq M^2$



Cost of External Merge Sort

- Read+write+read = $3B(R)$
- Assumption: $B(R) \leq M^2$
- Other considerations
 - In general, a lot of optimizations are possible.

Two-Pass Join Algorithm Based on Sorting

Join $R \bowtie S$

- Step1: Sort both R and S on the join attribute:
 - Cost: $4B(R)+4B(S)$ (because need to write to disk)
- Step2: Read both relations in sorted order, match tuples
 - Cost: $B(R)+B(S)$
- Total cost: $5B(R)+5B(S)$
- Assumption: $B(R) \leq M^2, B(S) \leq M^2$
- All tuples with a common value for the join attribute must fit in the memory.

Two-Pass Algorithms Based on Sorting

Join $R \bowtie S$

- If the number of tuples in R matching those in S is small (or vice versa), we can compute the join during the second phase of sort (sort-merge-join)
- Total cost: $3B(R) + 3B(S)$
- Assumption: $B(R) + B(S) \leq M^2$

Summary of Join Algorithms

- Nested Loop Join: $B(R) + B(R)B(S)/M$
 - Assuming block-at-a-time refinement, with one block-at-a time, the cost is: $B(R) + B(R)B(S)$
- Hash Join: $3B(R) + 3B(S)$
 - Assuming: $\min(B(R), B(S)) \leq M^2$
- Sort-Merge Join: $3B(R) + 3B(S)$
 - Assuming $B(R) + B(S) \leq M^2$
- Index Nested Loop Join: $B(R) + T(R)B(S)/V(S,a)$
 - Assuming S has clustered index on attribute a

Exercise

Consider joining two relations $R(x, y)$ and $S(x, z)$ on their common attribute x . The size of relation R is 1000 blocks and the size of relation S is 500 blocks. Attribute x of relation R has 50 different values and the values are evenly distributed in R . Attribute x of relation S also has the same 50 different values and the values are evenly distributed in S . Suppose that neither relation is sorted by attribute x . Furthermore, suppose that the memory buffer has 101 blocks.

- Briefly explain if/how the two relations R and S can be joined using a two-pass sort-merge join algorithm and compute the cost of join.
- Similarly, explain how the two relations R and S can be joined using a two-pass partitioned hash-join algorithm and compute the cost of join.