

CIS 560 – Database System Concepts

Lecture 28

# Query Optimization

November 8, 2013

Credits for slides: Chang, Ullman, Whitehead.

Copyright: Caragea, 2013.

## Planning

- Assignment 8 (indexes) due 11/8
- Project - DB design revision due 11/11
  - No class that day – use the time to work on project
- Assignment 9 (query optimization) due 11/15
- Exam 2 (assignments 6-9) – 11/20
- Project - DB implementation and queries due 11/22
- Quiz from special topics – 12/06
- Project presentations 12/9, 12/11, 12/13
- Project reports – finals week

2

## Summary of Join Algorithms

- Nested Loop Join:  $B(R) + B(R)B(S)/M$ 
  - Assuming block-at-a-time refinement, with one block-at-a time, the cost is:  $B(R) + B(R)B(S)$
- Hash Join:  $3B(R) + 3B(S)$ 
  - Assuming:  $\min(B(R), B(S)) \leq M^2$
- Sort-Merge Join:  $3B(R) + 3B(S)$ 
  - Assuming  $B(R) + B(S) \leq M^2$
- Index Nested Loop Join:  $B(R) + T(R)B(S)/V(S,a)$ 
  - Assuming S has clustered index on attribute a

## Query Optimization Goal

- For a query
  - There exists many logical and physical query plans
  - Query optimizer needs to pick a good one

## Example

Supplier(sid, sname, scity, sstate)  
Supply(sid, pno, quantity)

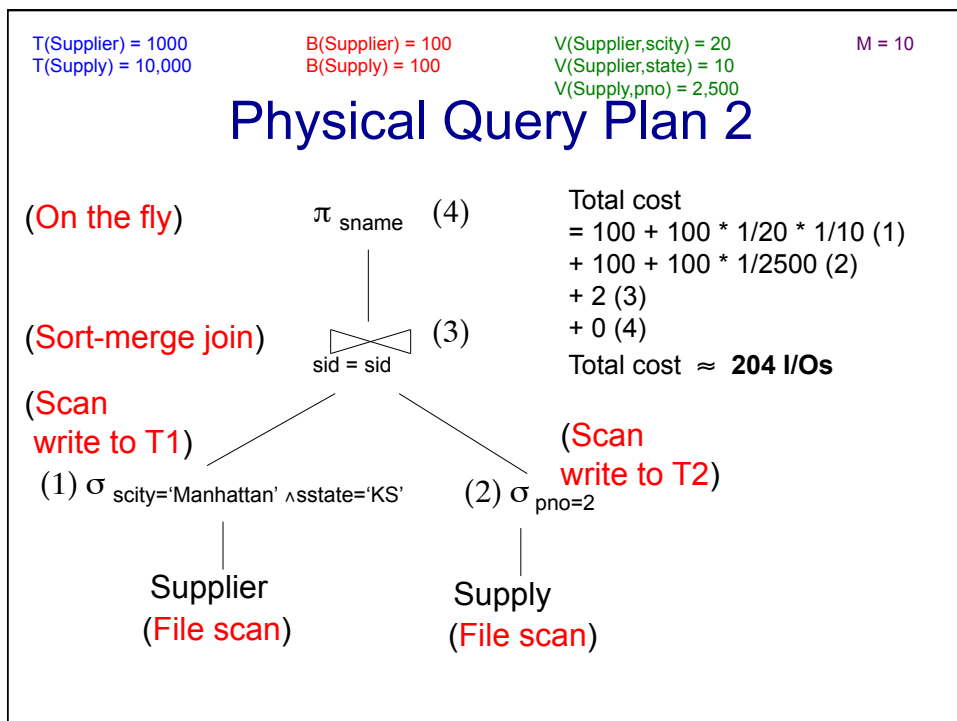
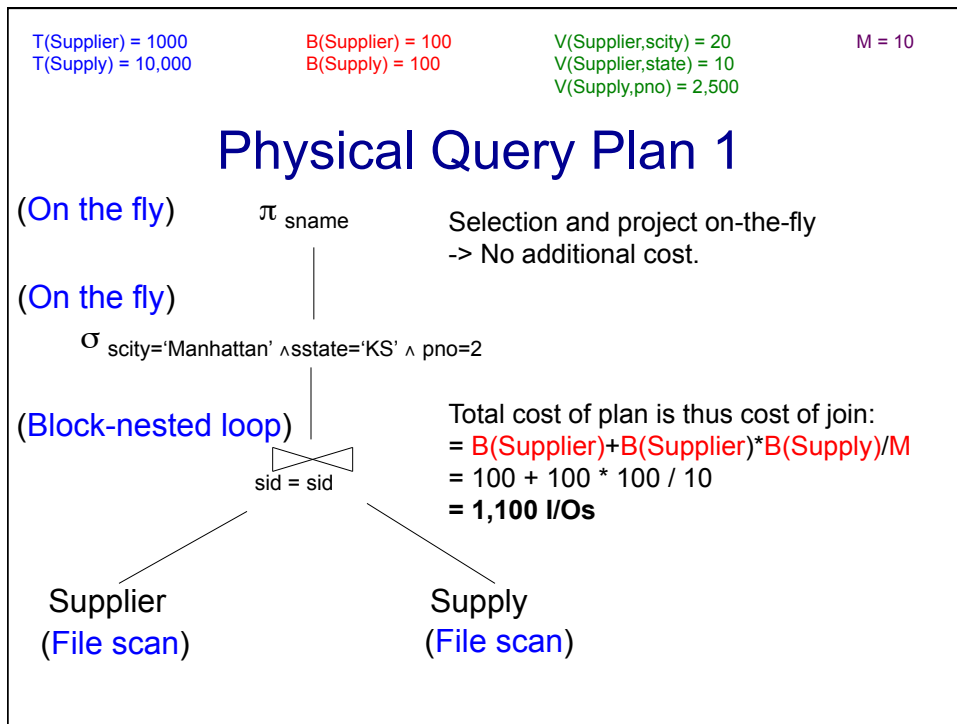
- Some statistics
  - T(Supplier) = 1000 records
  - T(Supply) = 10,000 records
  - B(Supplier) = 100 pages
  - B(Supply) = 100 pages
  - V(Supplier,scity) = 20, V(Supplier,state) = 10
  - V(Supply,pno) = 2,500
  - Both relations are clustered
- M = 10

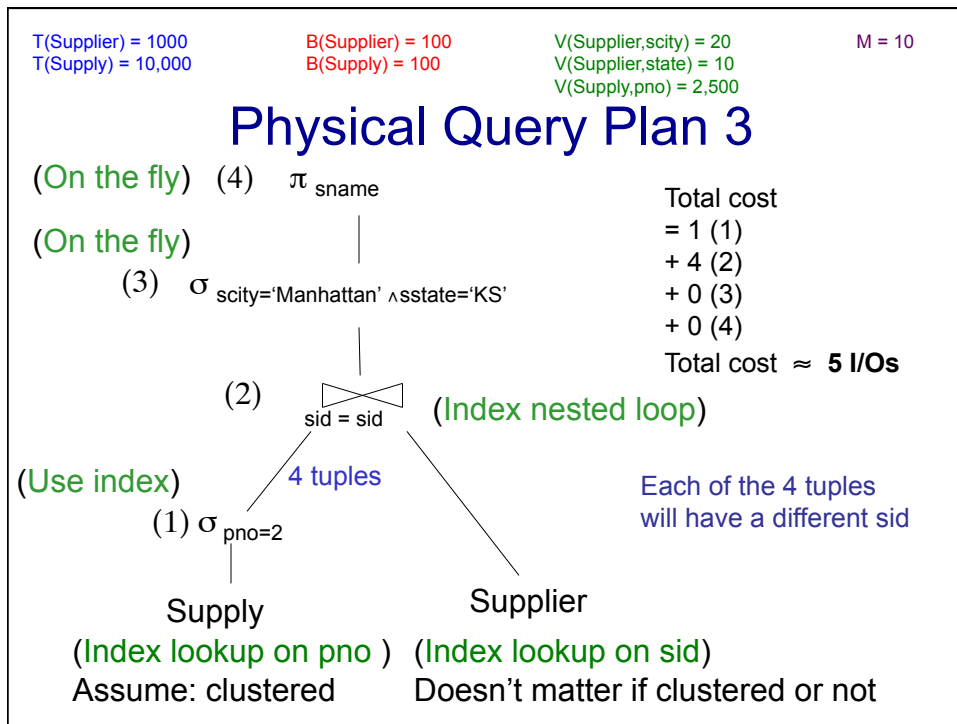
```
SELECT sname
FROM Supplier x, Supply y
WHERE x.sid = y.sid
      and y.pno = 2
      and x.scity = 'Manhattan'
      and x.sstate = 'KS'
```

## Relational Algebra Expressions

$$\pi_{\text{sname}}(\sigma_{\text{scity}='Manhattan' \wedge \text{sstate}='KS' \wedge \text{pno}=2} (\text{Supplier} \bowtie_{\text{sid}=\text{sid}} \text{Supply}))$$

$$\pi_{\text{sname}}((\sigma_{\text{scity}='Manhattan' \wedge \text{sstate}='KS'}(\text{Supplier})) \bowtie_{\text{sid}=\text{sid}} (\sigma_{\text{pno}=2}(\text{Supply})))$$





## Simplifications

- In the previous examples, we assumed that all index pages were in memory
- When this is not the case, we need to add the cost of fetching index pages from disk

## Lessons

- Need to consider several physical plans
  - even for one, simple logical plan
- No magic “best” plan: depends on the data
  - In order to make the right choice
    - need to have **statistics** over the data
    - the B’s, the T’s, the V’s

## Query Optimization Algorithm

- Enumerate alternative plans
- Compute estimated cost of each plan
  - Compute number of I/Os
- Choose plan with lowest cost
  - This is called cost-based optimization

## Components of an optimizer

We need three things in an optimizer:

- Search space (algebraic laws – relational algebra equivalences)
- Algorithm for enumerating query plans
- A cost estimator for a plan

## Relational Algebra Equivalences



- We can commute and combine operators
- We just have to be careful that the fields we need are available when we apply the operator

## Commutativity, Associativity, Distributivity

$$\begin{aligned} R \cup S &= S \cup R, & R \cup (S \cap T) &= (R \cup S) \cap T \\ R \cap S &= S \cap R, & R \cap (S \cup T) &= (R \cap S) \cup T \\ R \bowtie S &= S \bowtie R, & R \bowtie (S \bowtie T) &= (R \bowtie S) \bowtie T \end{aligned}$$

$$R \bowtie (S \cup T) = (R \bowtie S) \cup (R \bowtie T)$$

## Left-Deep Plans and Bushy Plans





## Example

Which plan is more efficient?

$R \bowtie (S \bowtie T)$  or  $(R \bowtie S) \bowtie T$ ?

- Assumptions:
  - Every join selectivity is 10%
    - That is:  $T(R \bowtie S) = 0.1 * T(R) * T(S)$  etc.
  - $B(R)=100$ ,  $B(S) = 50$ ,  $B(T)=500$
  - All joins are main memory joins
  - All intermediate results are materialized

## Laws involving selection:

$$\begin{aligned} \sigma_{C_1}(\sigma_{C_2}(R)) &= \sigma_{C_2}(\sigma_{C_1}(R)) \\ \sigma_{C \text{ AND } C'}(R) &= \sigma_C(\sigma_{C'}(R)) = \sigma_C(R) \cap \sigma_{C'}(R) \\ \sigma_{C \text{ OR } C'}(R) &= \sigma_C(R) \cup \sigma_{C'}(R) \\ \sigma_C(R \cup S) &= \sigma_C(R) \cup \sigma_C(S) \end{aligned}$$

$$\begin{aligned} \sigma_C(R \cup S) &= \sigma_C(R) \cup S \\ \sigma_C(R - S) &= \sigma_C(R) - S \\ \sigma_C(R \bowtie S) &= \sigma_C(R) \bowtie S \end{aligned}$$

When C involves  
only attributes of R

## Example: Simple Algebraic Laws

- Example:  $R(A, B, C, D), S(E, F, G)$

$$\sigma_{F=3}(R \bowtie_{D=E} S) = ?$$

$$\sigma_{A=5 \text{ AND } G=9}(R \bowtie_{D=E} S) = ?$$

## Example: Simple Algebraic Laws

- Example:  $R(A, B, C, D), S(E, F, G)$

$$\sigma_{F=3}(R \bowtie_{D=E} S) = R \bowtie_{D=E} (\sigma_{F=3}(S))$$

$$\begin{aligned} \sigma_{A=5 \text{ AND } G=9}(R \bowtie_{D=E} S) &= \sigma_{A=5}(\sigma_{G=9}(R \bowtie_{D=E} S)) \\ &= (\sigma_{A=5}(R)) \bowtie_{D=E} (\sigma_{G=9}(S)) \end{aligned}$$

## Laws Involving Projections

$$\begin{aligned}\Pi_M(R \bowtie S) &= \Pi_M(\Pi_P(R) \bowtie \Pi_Q(S)) \\ \Pi_M(\Pi_N(R)) &= \Pi_M(R) \quad /* \text{note that } M \subseteq N */\end{aligned}$$

- Example  $R(A,B,C,D), S(E, F, G)$

$$\Pi_{A,B,G}(R \bowtie_{D=E} S) = \Pi_{?}(\Pi_{?}(R) \bowtie_{D=E} \Pi_{?}(S))$$

## Laws Involving Projections

$$\begin{aligned}\Pi_M(R \bowtie S) &= \Pi_M(\Pi_P(R) \bowtie \Pi_Q(S)) \\ \Pi_M(\Pi_N(R)) &= \Pi_M(R) \quad /* \text{note that } M \subseteq N */\end{aligned}$$

- Example  $R(A,B,C,D), S(E, F, G)$

$$\Pi_{A,B,G}(R \bowtie_{D=E} S) = \Pi_{A,B,G}(\Pi_{A,B,D}(R) \bowtie_{D=E} \Pi_{E,G}(S))$$

## Search Space Challenges

- Search space is huge!
  - Many possible equivalent trees
  - Many implementations for each operator
  - Many access paths for each relation
    - File scan or index + matching selection condition
- Cannot consider ALL plans
  - Heuristics: only partial plans with “low” cost

## Algorithms for enumerating plans: key decisions

- Logical plan
  - What logical plans do we consider (left-deep, bushy?)
    - *Search space*
  - Which algebraic laws do we apply, and in which context(s)?
    - *Optimization rules*
  - In what order do we explore the search space?
    - *Optimization algorithm*
- Physical plan
  - What join algorithms to use?
  - What access paths to use (file scan or index)?

## Types of Optimizers

- Rule-based optimizers:
  - Apply greedily rules that always improve
    - Typically: push selections down, pull projections up
  - Very limited: no longer used today
- Cost-based optimizers
  - Use a cost model to estimate the cost of each plan
  - Select the “cheapest” plan



## The Search Space

- Complete plans
- Bottom-up plans
- Top-down plans

## Complete Plans

R(A,B)  
S(B,C)  
T(C,D)

SELECT \*  
FROM R, S, T  
WHERE R.B=S.B and S.C=T.C and R.A<40




Why is this  
search space  
inefficient?

## Bottom-up Partial Plans

R(A,B)  
S(B,C)  
T(C,D)

SELECT \*  
FROM R, S, T  
WHERE R.B=S.B and S.C=T.C and R.A<40


Why is this  
better?



## Top-down Partial Plans

R(A,B)  
S(B,C)  
T(C,D)

```
SELECT *
FROM R, S, T
WHERE R.B=S.B and S.C=T.C and R.A<40
```



## Search Strategies

- **Branch-and-bound:**
  - Remember the cheapest complete plan P seen so far and its cost C
  - Stop generating partial plans whose cost is > C
  - If a cheaper complete plan is found, replace P, C
- **Hill climbing:**
  - Remember only the cheapest partial plan seen so far
- **Dynamic programming:**
  - Remember all cheapest partial plans

## Dynamic Programming

Originally proposed in System R [1979]

- Limited to joins: *join reordering algorithm*
- Bottom-up
- Only handles single block queries:

```
SELECT list  
FROM   R1, ..., Rn  
WHERE  cond1 AND cond2 AND ... AND condk
```


## Dynamic Programming

- Search space = join trees
- Algebraic laws = commutativity, associativity
- Algorithm = dynamic programming ☺



## Join Trees


- $R1 \bowtie R2 \bowtie \dots \bowtie Rn$
- Join tree:



- A plan = a join tree
- A partial plan = a subtree of a join tree


## Types of Join Trees

- Left deep:




## Types of Join Trees

- Right deep:



## Types of Join Trees

- Bushy:



```
SELECT list
FROM R1, ..., Rn
WHERE cond1 AND cond2 AND ... AND condk
```

## Dynamic Programming

Join ordering:

- Given: a query  $R1 \bowtie R2 \bowtie \dots \bowtie Rn$
- Find optimal order
- Assume we have a function  $\text{cost}()$  that gives us the cost of every join tree

```
SELECT list
FROM R1, ..., Rn
WHERE cond1 AND cond2 AND ... AND condk
```

## Dynamic Programming

- Idea: for each subset of  $\{R1, \dots, Rn\}$ , compute the best plan for that subset
- In increasing order of set cardinality:
  - Step 1: for  $\{R1\}, \{R2\}, \dots, \{Rn\}$
  - Step 2: for  $\{R1, R2\}, \{R1, R3\}, \dots, \{Rn-1, Rn\}$
  - ...
  - Step n: for  $\{R1, \dots, Rn\}$
- It is a bottom-up strategy
- A subset of  $\{R1, \dots, Rn\}$  is also called a *subquery*

```
SELECT list
FROM R1, ..., Rn
WHERE cond1 AND cond2 AND ... AND condk
```

## Dynamic Programming

- For each subquery  $Q \subseteq \{R_1, \dots, R_n\}$  compute the following:
  - $\text{Size}(Q)$  = the estimated size of  $Q$
  - $\text{Plan}(Q)$  = a best plan for  $Q$
  - $\text{Cost}(Q)$  = the estimated cost of that plan

```
SELECT list
FROM R1, ..., Rn
WHERE cond1 AND cond2 AND ... AND condk
```

## Dynamic Programming

- **Step 1:** For each  $\{R_i\}$  do:
  - $\text{Size}(\{R_i\}) = B(R_i)$
  - $\text{Plan}(\{R_i\}) = R_i$
  - $\text{Cost}(\{R_i\}) = (\text{cost of scanning } R_i)$

```
SELECT list
FROM R1, ..., Rn
WHERE cond1 AND cond2 AND ... AND condk
```

## Dynamic Programming

- **Step i:** For each  $Q \subseteq \{R_1, \dots, R_n\}$  of cardinality  $i$  do:
  - $\text{Size}(Q)$  = estimate it recursively
  - For every pair of subqueries  $Q', Q''$  s.t.  $Q = Q' \cup Q''$  compute  $\text{cost}(\text{Plan}(Q') \bowtie \text{Plan}(Q''))$ 
    - $\text{Cost}(Q)$  = the smallest such cost
    - $\text{Plan}(Q)$  = the corresponding plan

```
SELECT list
FROM R1, ..., Rn
WHERE cond1 AND cond2 AND ... AND condk
```

## Dynamic Programming

- **After step n:** Return  $\text{Plan}(\{R_1, \dots, R_n\})$

## Example

To illustrate, ad-hoc cost model (from the book ☺):

- In practice: more realistic size/cost estimations
- $\text{Cost}(P_1 \bowtie P_2) = \text{Cost}(P_1) + \text{Cost}(P_2) + \text{size}(\text{intermediate results for } P_1, P_2)$ 
  - Intermediate results:
    - If  $P_1$  is a join, then the size of the intermediate result is  $\text{size}(P_1)$ , otherwise the size is 0
    - Similarly for  $P_2$
- Cost of a scan = 0

## Dynamic Programming

Example:

- $\text{Cost}(R5 \bowtie R7) = 0$  (no intermediate results)
- $\text{Cost}((R2 \bowtie R1) \bowtie R7)$ 
  - $= \text{Cost}(R2 \bowtie R1) + \text{Cost}(R7) + \text{size}(R2 \bowtie R1)$
  - $= \text{size}(R2 \bowtie R1)$

```
SELECT *
FROM R, S, T, U
WHERE cond1 AND cond2 AND ...
```

## Example

- $R \bowtie S \bowtie T \bowtie U$
- Assumptions:

All join selectivities = 1%

$T(R) = 2000$   
 $T(S) = 5000$   
 $T(T) = 3000$   
 $T(U) = 1000$

$T(R \bowtie S) = 0.01 * T(R) * T(S)$   
 $T(S \bowtie T) = 0.01 * T(S) * T(T)$   
 etc.

$T(R) = 2000$   
 $T(S) = 5000$   
 $T(T) = 3000$   
 $T(U) = 1000$

$T(R \bowtie S) =$   
 $0.01 * T(R) * T(S)$   
 $T(S \bowtie T) =$   
 $0.01 * T(S) * T(T)$   
 etc.

Subquery	Size	Cost	Plan
RS			
RT			
RU			
ST			
SU			
TU			
RST			
RSU			
RTU			
STU			
RSTU			

<div> <math>T(R) = 2000</math>  <math>T(S) = 5000</math>  <math>T(T) = 3000</math>  <math>T(U) = 1000</math> </div> <div> <math>T(R \bowtie S) = 0.01 * T(R) * T(S)</math>  <math>T(S \bowtie T) = 0.01 * T(S) * T(T)</math>            etc.         </div>	Subquery	Size	Cost	Plan
	RS	100k	0	RS
	RT	60k	0	RT
	RU	20k	0	RU
	ST	150k	0	ST
	SU	50k	0	SU
	TU	30k	0	TU
	RST	3M	60k	(RT)S
	RSU	1M	20k	(RU)S
	RTU	0.6M	20k	(RU)T
	STU	1.5M	30k	(TU)S
	RSTU	30M	60k +50k=110k	(RT)(SU)

## Reducing the Search Space

- Restriction 1: only left linear trees (no bushy)
- Restriction 2: no trees with cartesian product

$R(A,B) \bowtie S(B,C) \bowtie T(C,D)$

Plan:  $(R(A,B) \bowtie T(C,D)) \bowtie S(B,C)$   
has a cartesian product.

Most query optimizers will not consider it



## Dynamic Programming: Summary

- Handles only join queries:
  - Selections are pushed down (i.e. early)
  - Projections are pulled up (i.e. late)
- Takes exponential time in general, BUT:
  - Left linear joins may reduce time
  - Non-cartesian products may reduce time further

## Completing the Physical Query Plan

- Choose algorithm for each operator
  - How much memory do we have?
  - Are the input operand(s) sorted?
- Access path selection for base tables
- Decide for each intermediate result:
  - To materialize
  - To pipeline

## Summary of Query Optimization

- Three parts:
  - search space, algorithms, size/cost estimation
- Ideal goal: find optimal plan. But
  - Impossible to estimate accurately
  - Impossible to search the entire space
- Goal of today's optimizers:
  - Avoid very bad plans