
CIS 721 - Real-Time Systems

Lecture 2: Real-Time Systems

Reference Model

Mitch Neilsen
neilsen@ksu.edu

Outline

- **Terms and Concepts (Ch. 1-2)**
 - **Real-Time Systems Reference Model (Ch. 3)**
 - **Commonly Used Approaches For Real-Time Scheduling (Ch. 4)**
 - Clock-Driven Scheduling (Ch. 5)
 - Priority-Driven Scheduling (Ch. 6-7)
-

Definitions

- A **job** is a unit of work that is scheduled and executed by the system ($J_{i,k}$).
- A **task** is a set of related jobs that provide some system function $\tau_i = \{ J_{i,1}, J_{i,2}, \dots, J_{i,n} \}$; e.g., the reception of a data frame could be a job that is part of a task that provides time service.
- The **deadline** of a job is the time at which a job must be completed.

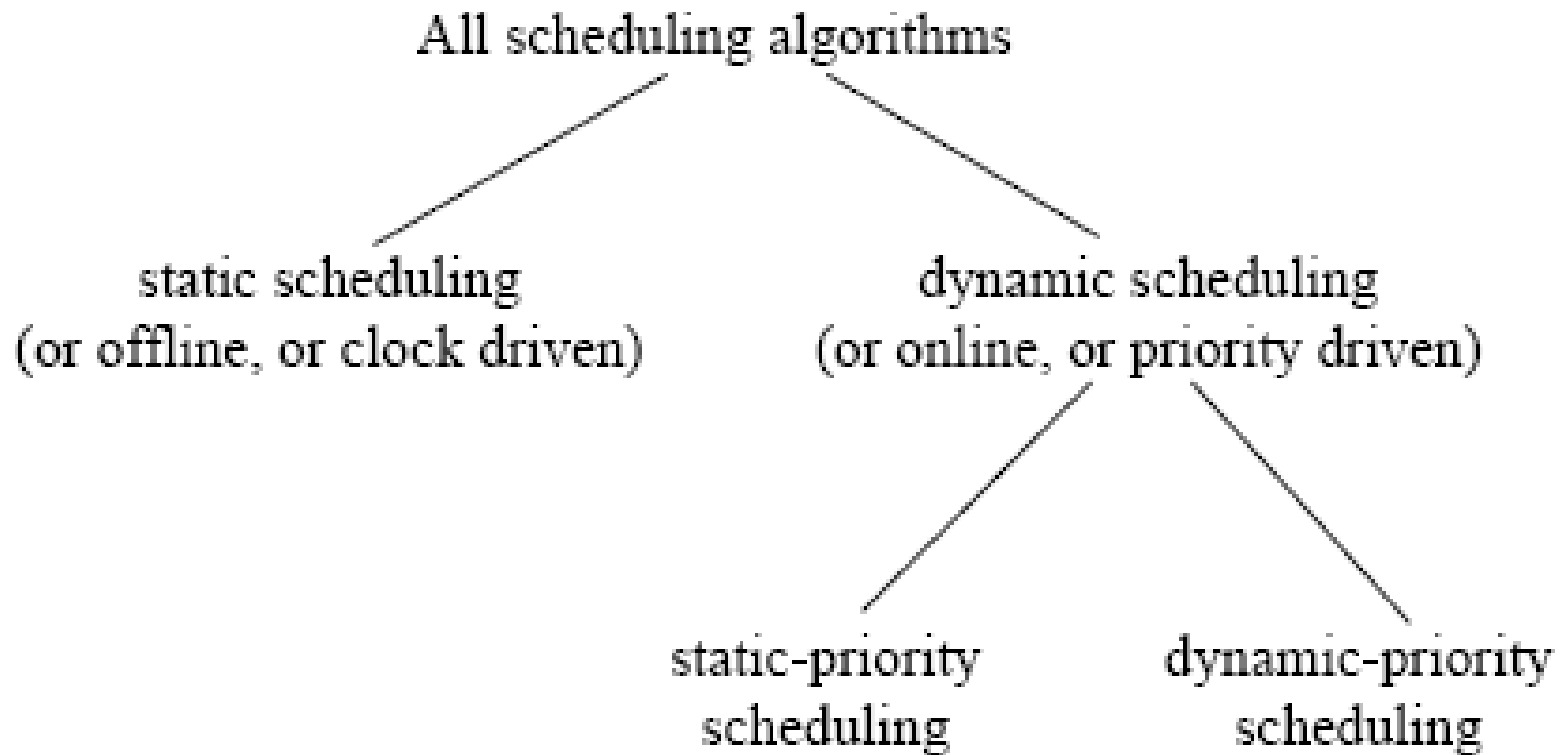
Deadlines

- The **release time** (or **arrival time**) of a job is the time at which the job becomes available for execution (r_i or R_i).
- The **response time** of a job is the length of time between the release time of the job and the time instant when it completes.
- The **relative deadline** of a job is the maximum allowable response time of a job (D_i).
- The **absolute deadline** of a job is the time at which a job must be completed ($d_i = r_i + D_i$).

Algorithms

- We're interested in two types of algorithms:
 - **Scheduling algorithms** are used to generate to a schedule or priority assignment that can be used by a scheduler to schedule tasks at run time.
 - **Feasibility/schedulability analysis algorithms** are used to determine if a given task set is schedulable.
 - Normally, the second class of algorithms are much more complex than the first class.
-

Classification of Scheduling Algorithms



Static Scheduling Algorithms

- **Static scheduling algorithms** can be used if the scheduling algorithm has complete knowledge of the task set and all timing constraints such as deadlines, execution times, precedence, and **future** arrival times.
- The static algorithm operates on the set of tasks and constraints to generate a single, fixed schedule.

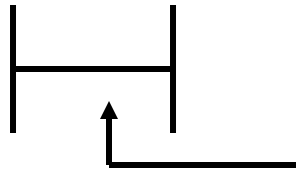
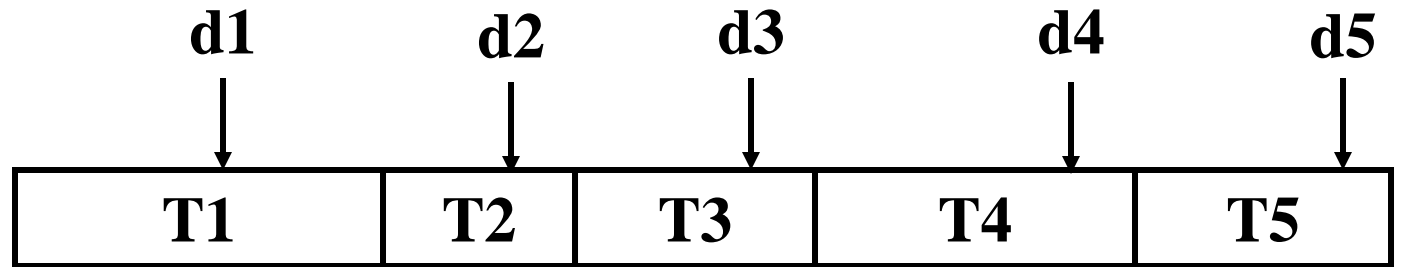
Dynamic Scheduling Algorithms

- **Dynamic scheduling algorithms** have complete knowledge of **currently active** jobs, but new jobs may arrive at any time in the future.
 - Dynamic scheduling is performed at run-time (online); however, offline analysis is usually performed to constrain the dynamic schedule; e.g., assign priorities, etc.
-

Metrics used to evaluate scheduling algorithms

- processor utilization
- throughput
- weighted sum of task completion times
- schedule length
- number of processors required
- maximum lateness
- **missed deadlines**

Minimize maximum lateness



**Maximum lateness is minimized,
but all deadlines are missed.**



Only one deadline missed.

Missed Deadlines

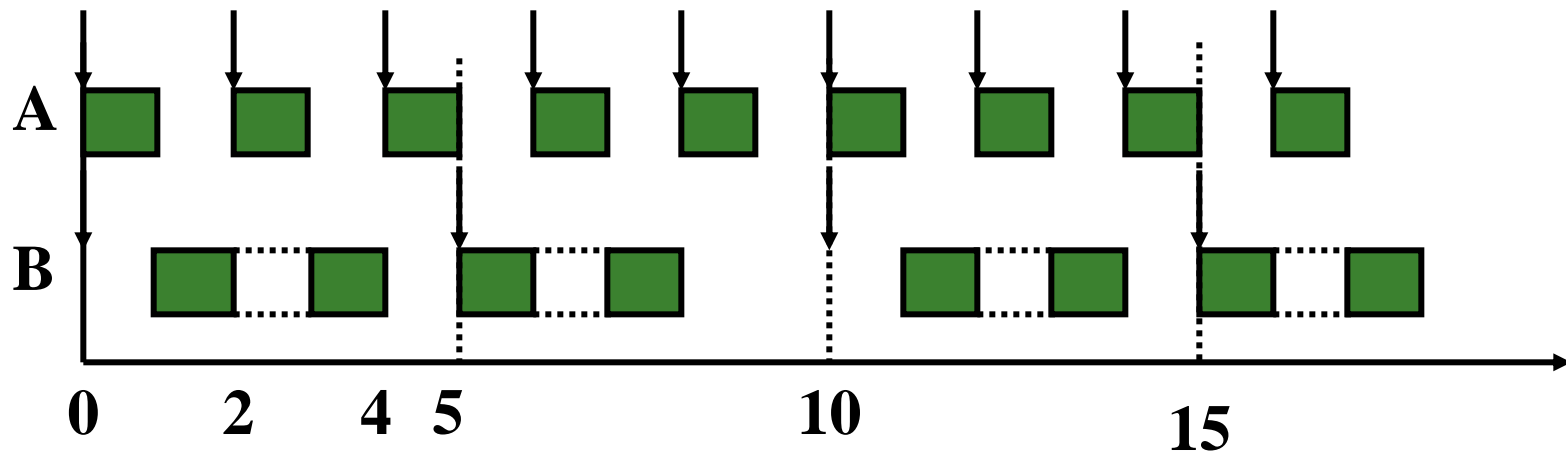
- Much real-time work is only concerned with missed deadlines; e.g., for **hard real-time systems all deadlines must be met**.
- In which case, an **optimal** scheduling algorithm is one that will fail to meet a deadline for any given task set only if no other scheduling algorithm can meet the deadlines.

Periodic Task

- A periodic task $\tau_i = \{ J_{i,1}, J_{i,2}, \dots, J_{i,n} \}$ is a sequence of jobs with identical parameters with:
 - a **period** (p_i or T_i) equal to the length of time between the release times of consecutive jobs,
 - an **execution time** (e_i or C_i) equal to the maximum execution time of any job in the task, and
 - a **phase** (φ_i) equal to the release time of the first job in τ_i .

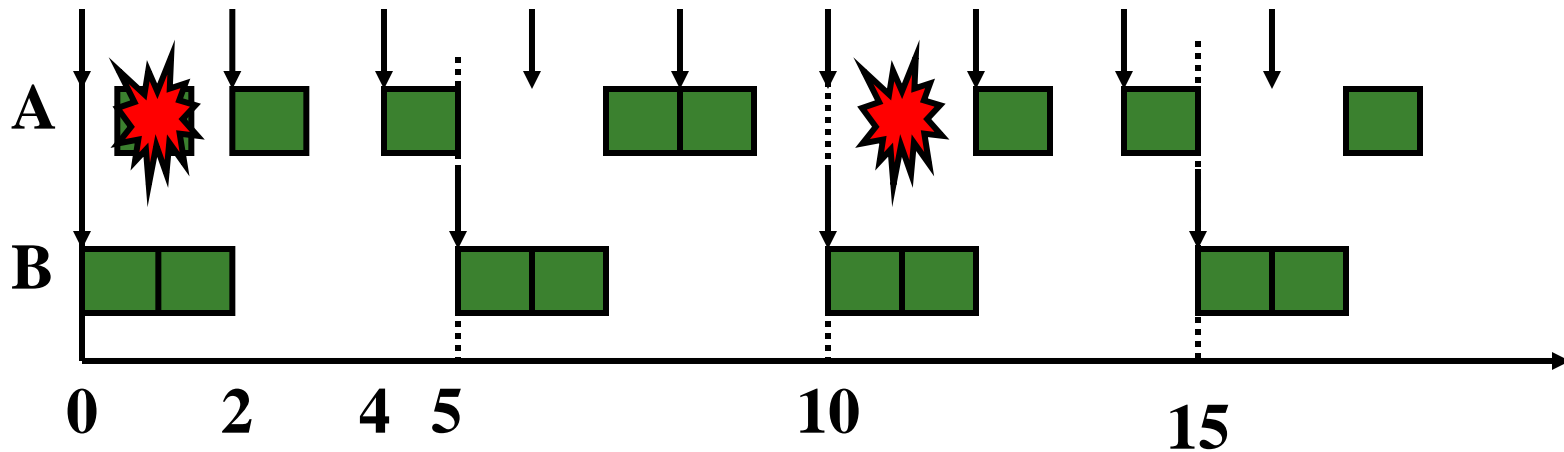
Example #1 - Priority-Driven Scheduler

| Task | Period | Deadline | Run-Time | Phase |
|--------------------------|----------|----------|----------|----------|
| τ_i | T_i | D_i | C_i | ϕ_i |
| <hr/> | | | | |
| A (High Priority) | 2 | 2 | 1 | 0 |
| B (Low Priority) | 5 | 5 | 2 | 0 |



Example #2

| Task | Period | Deadline | Run-Time | Phase |
|--------------------------|----------|----------|----------|----------|
| τ_i | T_i | D_i | C_i | ϕ_i |
| <hr/> | | | | |
| A (Low Priority) | 2 | 2 | 1 | 0 |
| B (High Priority) | 5 | 5 | 2 | 0 |



Example #3

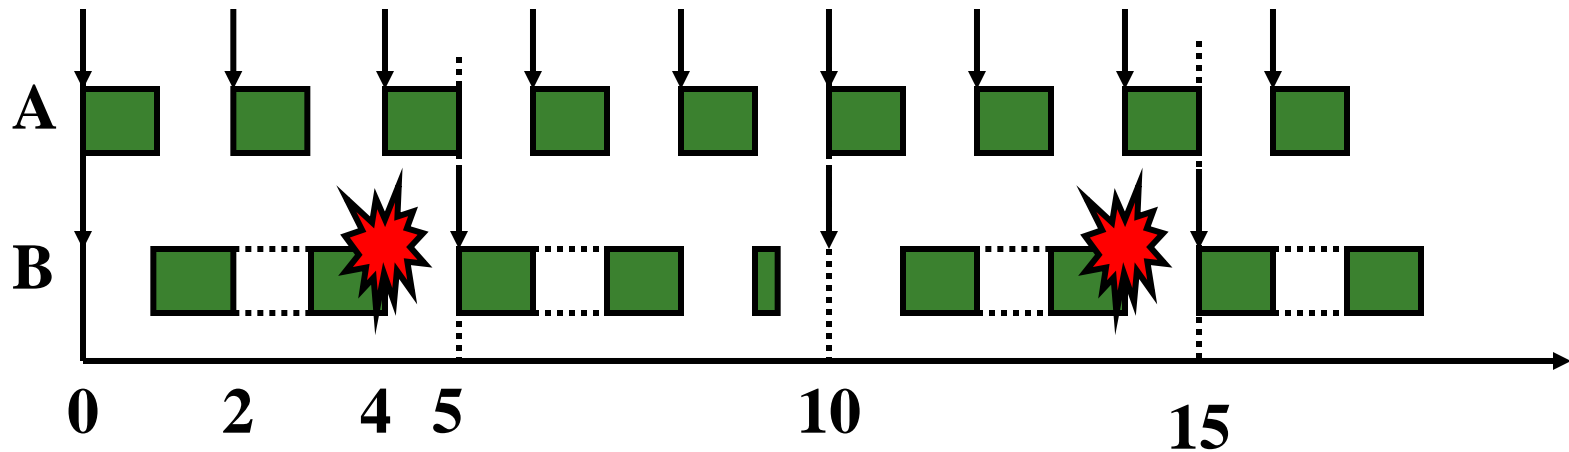
| Task | Period | Deadline | Run-Time | Phase |
|-------------|---------------|-----------------|-----------------|--------------|
| τ_i | T_i | D_i | C_i | ϕ_i |
| <hr/> | | | | |
| A | 2 | 2 | 1 | 0 |
| B | 5 | 5 | 2.1 | 0 |

$$U = C_1 / T_1 + C_2 / T_2 = 1 / 2 + 2.1 / 5 = 0.92$$

Even if $U < 1$, a task set **may not be schedulable** using fixed priority scheduling.

Example #3

| Task | | Period | Deadline | Run-Time | Phase |
|----------|-----------------|--------|----------|----------|----------|
| τ_i | | T_i | D_i | C_i | ϕ_i |
| <hr/> | | | | | |
| A | (High Priority) | 2 | 2 | 1 | 0 |
| B | (Low Priority) | 5 | 5 | 2.1 | 0 |



Observations

- The schedulability of a task set depends on priority assignment (Example 1 is schedulable, but Example 2 is not).
 - Even if the utilization of a task set is less than one, it may not be schedulable by any fixed priority assignment (Example 3 is not).
-

Priority-Driven Scheduling Algorithms for Periodic Tasks

- **Fixed-Priority** - assigns the same priority to all jobs in a task.
- **Dynamic-Priority** - assigns different priorities to the individual jobs in each task.
- After looking at Static Scheduling Algorithms (next time), we will start our investigation of Dynamic Scheduling Algorithms by considering **fixed-priority algorithms**.

Issues in Fixed Priority Assignment

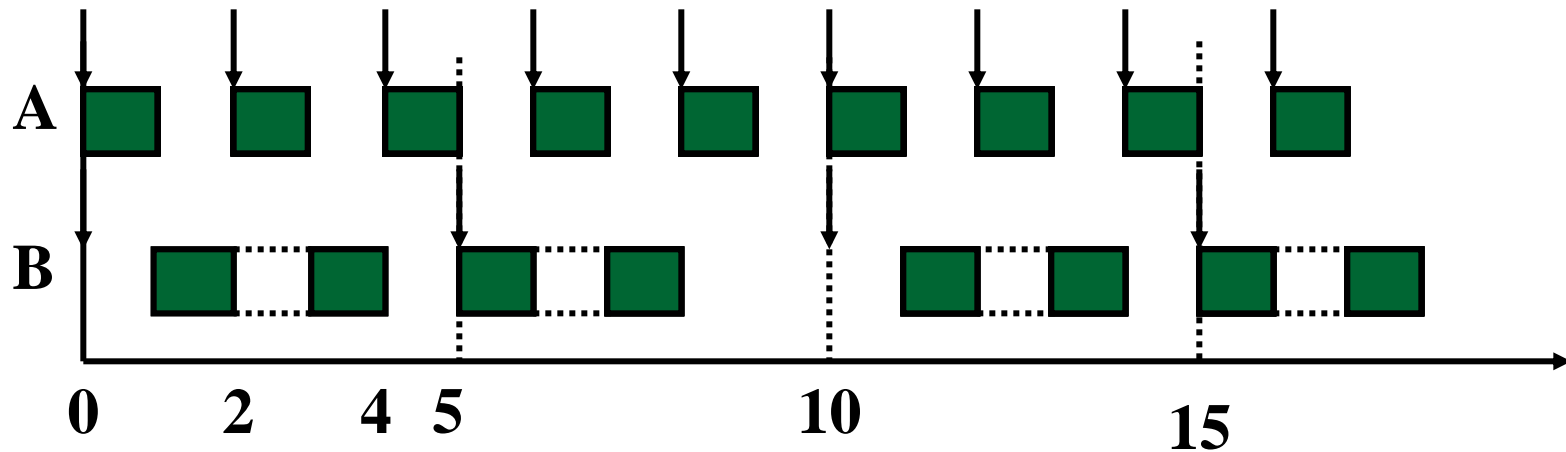
- How to assign priorities?
 - How to determine which assignment is the best; e.g., how to evaluate a priority assignment algorithm (method)?
 - How to compare different priority assignment algorithms?
-

Rate-Monotonic Algorithm (RM)

- The **rate** of a task is the inverse of its period.
 - Task with **higher rates** are assigned **higher priorities**.
 - C. L. Liu and J. W. Layland, “Scheduling Algorithms for Multiprogramming in a Hard Real-Time Environment”, JACM, Vol. 20, No. 1, pages 46-61, 1973.
-

Example #1 - Rate Monotonic Assignment

| Task | Period | Deadline | Run-Time | Phase |
|--------------------------|----------|----------|----------|----------|
| τ_i | T_i | D_i | C_i | ϕ_i |
| <hr/> | | | | |
| A (High Priority) | 2 | 2 | 1 | 0 |
| B (Low Priority) | 5 | 5 | 2 | 0 |



Real-Time Reference Model (Ch. 3)

- Idea: Abstract away functional characteristics and focus on **timing requirements and resource requirements**.
- **Reference Model Components**
 - **Resource Graph** – identify available system resources, resource types, and dependencies
 - **Task Graph** – identify task dependencies
 - **Scheduling and Resource Management** – identify algorithms for scheduling and resource management

Processors and Resources

- **Processors (P_i)** are **active** system resources, such as computers, transmission (tx) links, and database servers
 - **Resources (R_i)** are **passive** system resources, such as memory, mutexes, semaphores, and database locks
 - **Example: Sliding Window Protocol**
 - Job = transmit a message
 - Processor = data link
 - Resource = valid sequence number
-

Types of Resources

- **Reusable** – most resources are reusable; e.g., they can be reused by subsequent jobs after being released.
 - Ex: a mutex is a serially reusable resource
- **Plentiful** – a resource is plentiful if no job is ever prevented from executing due to a lack of this resource.
 - Ex: a read-only (immutable) configuration file
 - Plentiful resources are typically removed from the model.

Resource or Processor?

- For some problems, it is hard to classify system resources as processors or resources.
 - This is where experience and the “art” of modeling comes in to play.
 - **Example: I/O Bus**
 - In many cases the I/O Bus is viewed as a plentiful resource and ignored in the model
 - However, if we want to study how I/O activities impact real-time performance of an I/O arbitration scheme, then the bus must be modeled as a resource or processor.
-

Temporal Parameters

- J_i : **job** – a unit of work
 - T_i (or τ_i) : **task** - a set of related jobs
 - A **periodic task** is sequence of invocations of jobs with identical parameters.
 - r_i : **release time** of job J_i
 - d_i : **absolute deadline** of job J_i
 - D_i : **relative deadline** (or just **deadline**) of job J_i
 - e_i : (Maximum) **execution time** of job J_i
-

Periodic Task Model

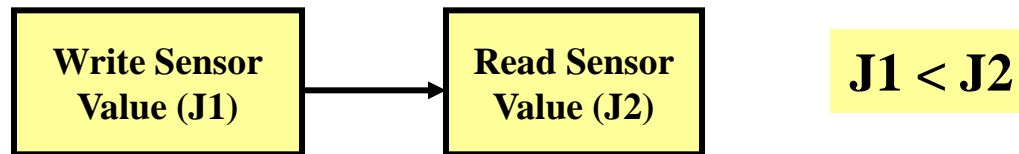
- **Tasks:** T_1, \dots, T_n
- Each consists of a set of **jobs**: $T_i = \{J_{i1}, J_{i2}, \dots\}$
- φ_i : **phase** of task T_i = time when its first job is released
- p_i : **period** of T_i = inter-release time
- H : **hyperperiod** $H = \text{lcm}(p_1, \dots, p_n)$
- e_i : **execution time** of T_i
- u_i : **utilization** of task T_i is given by $u_i = e_i / p_i$
- D_i : (relative) **deadline** of T_i , typically $D_i = p_i$

Types of Release Times

- **Fixed** – release times are known values (**periodic**)
- **Jittered**: $r_i \in [r_i^-, r_i^+]$: release time of job J_i falls within a known interval
- **Sporadic or aperiodic** – release times are unknown
 - $A(x)$ = interarrival time (time between two consecutive jobs) probability distribution
 - $B(x)$ = execution time distribution
- **Definitions**
 - **Sporadic tasks** have jobs with hard relative deadlines, but **aperiodic tasks** have either soft or no deadlines

Precedence Constraints/Graphs

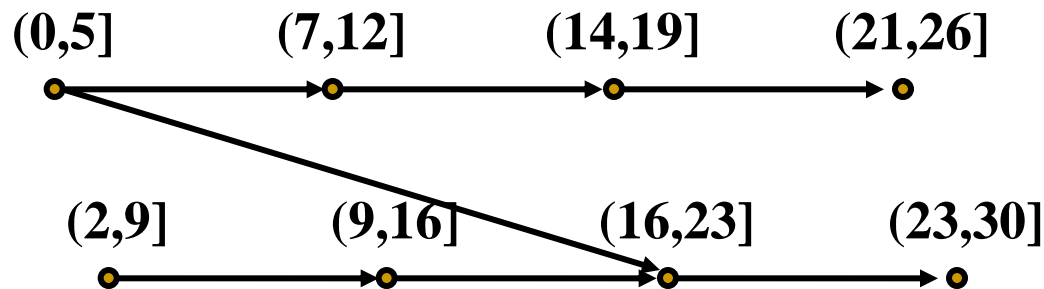
- A **precedence graph** reflects data/control dependencies
- **Example: Sensor/actuator (producer/consumer)**



- A precedence relation, denoted $<$, defines a partial order on the set of jobs.
- $J_i < J_k$ if J_i is a predecessor of J_k
- Precedence graph: $G = (\mathbf{J}, <)$, $\mathbf{J} = \{J_1, J_2, \dots\}$
- Precedence constraints can include AND/OR constraints.
- Some dependencies **cannot** be captured by task graphs
 - Example: access to shared data

Task Graph

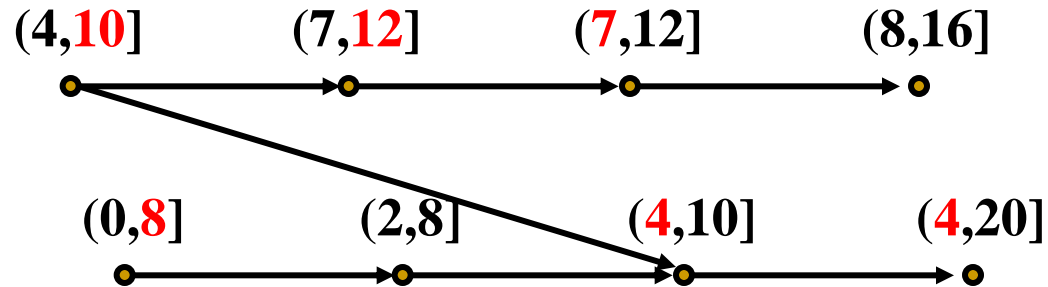
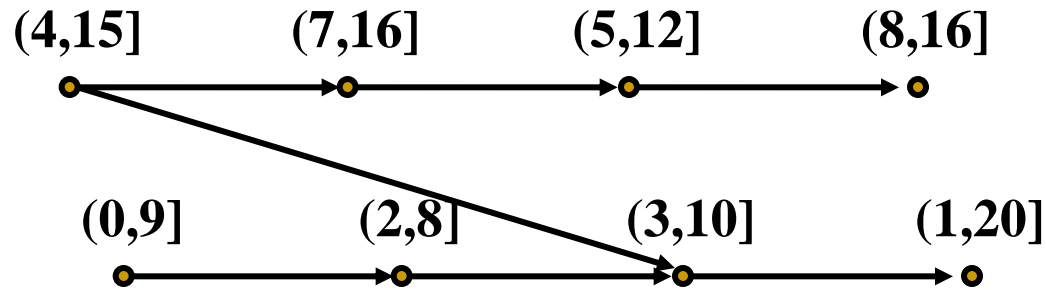
- A **task graph** is an extended precedence graph:
 - Vertices denote jobs
 - Edges denote dependencies
 - The label in brackets above each job give its feasible interval $(r_i, d_i] = (\text{release time}, \text{absolute deadline}]$.



Effective Timing Constraints

- Timing constraints are often inconsistent with precedence constraints; e.g., $d_1 > d_2$, but $J_1 < J_2$
- Effective timing constraints on a single processor:
 - **Effective release time:**
$$r_i^{\text{eff}} = \max(r_i , \{ r_k^{\text{eff}} \mid J_k < J_i \})$$
 - **Effective deadline:**
$$d_i^{\text{eff}} = \min(d_i , \{ d_k^{\text{eff}} \mid J_i < J_k \})$$
- **Theorem:** A set of jobs **J** can be feasibly scheduled on a processor iff it can be feasibly scheduled to meet all effective release times and effective deadlines.

Effective Release Times and Deadlines



Note

- Unless otherwise specified, we will use the terms release time and effective release time interchangeably; likewise, we will use the terms deadline and effective deadline interchangeably.
-

System Characterization

- **Preemptivity** - are the jobs preemptable; e.g., can the current task be suspended to assign the processor to a more urgent task?
- **Context-switching time** - is the time required to switch between tasks negligible?
- **Laxity type** - are deadlines hard or soft?
- **Resource requirements** - are any resources required by the job to execute, and for what time interval are these resources required (e.g., critical sections).
- **Criticalness** – can jobs be assigned weights to indicate their importance relative to other jobs? If so, algorithms can be used to optimize weighted performance metrics.

Schedules

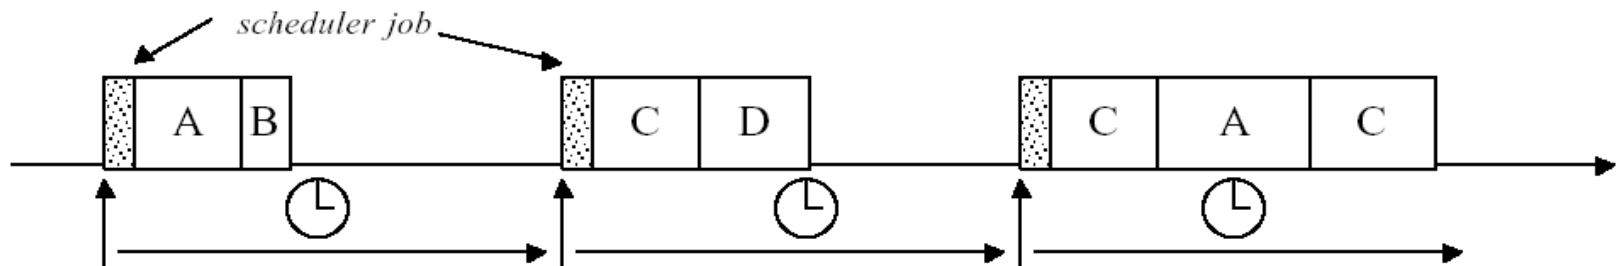
- A **schedule** is an assignment of jobs to available processors. In a **feasible schedule**, every job starts at or after its release time and completes by its deadline.
- In a hard real-time system, a scheduling algorithm is **optimal** if it always produces a feasible schedule if such a schedule exists.
- In a soft real-time system, we can consider different **performance metrics**:
 - Number of **missed deadlines** (tardy jobs).
 - Maximum (or average) **tardiness or lateness**.
 - Maximum (or average) **response time**.

Common Approaches For Real-Time Scheduling

- **Clock-Driven (Time-Driven) Approach** – scheduling decisions are made at specific time instants.
 - **Weighted Round-Robin Approach** - every job joins a FIFO queue; when a job reaches the front of the queue, its weight refers to the fraction of processor time (number of time slices) allocated to the job.
 - **Priority-Driven (Event-Driven) Approach** - ready jobs with highest priorities are scheduled for execution first.
 - Scheduling decisions are made when particular events occur; e.g., a job is released or a processor becomes idle. A **work-conserving** processor is busy whenever there is work to be done.
-

Clock-Driven Scheduling (Ch. 5)

- **Scheduling decision time:** point in time when scheduler decides which job to execute next.
- Scheduling decision time in clock-driven system is defined *a priori*; e.g., the scheduler periodically wakes up and generates a portion of the schedule.



- When job parameters are known *a priori*, the schedule can be precomputed off-line, and stored as a table (called a **table-driven scheduler**).

Priority-Driven Scheduling (Ch. 4)

- **Work-conserving** schedulers never leave the processor idle when there is work to be done; e.g., **priority-driven schedulers** typically apply a work-conserving, list-driven, greedy approach for scheduling.
- **Examples:** FIFO, LIFO, SETF (Shortest Execution Time First), LETF, EDF (Earliest Deadline First).
- Possible **implementation** of preemptive, priority-driven scheduling:
 - Assign priorities to jobs.
 - Scheduling decisions are made when: (a) a job becomes ready, (b) a processor becomes idle, or (c) priorities of jobs change
- At each scheduling decision time, choose ready task with highest priority.

Preemptive vs. Nonpreemptive

- In the **non-preemptive** case, scheduling decisions are only made when the processor becomes idle (not when a job becomes ready or when priorities change).

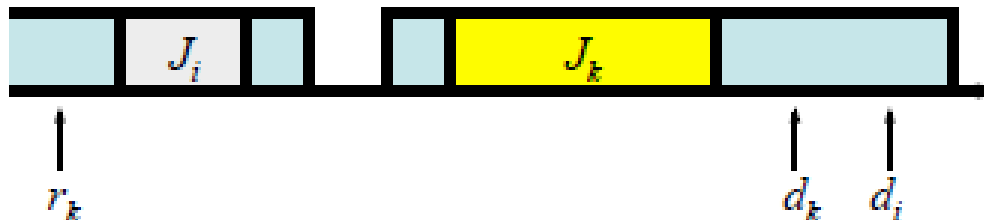
EDF Algorithm

- **Earliest-Deadline-First (EDF) algorithm:**
 - At any time, execute the available job with the **earliest deadline**.
- **Theorem: (Optimality of EDF):** In a system with **one processor** and **preemption** allowed, EDF is optimal; that is, EDF can produce a feasible schedule for a given job set J with arbitrary release times and deadlines, **if** a feasible schedule exists.
- **Proof:** Suppose that a feasible schedule S exists, then apply schedule transformations to S to generate an EDF schedule that is also feasible.

EDF proof (schedule transformations)

1. Any feasible schedule can be transformed into an EDF schedule

- If J_i is scheduled to execute before J_k , but J_i 's deadline is later than J_k 's either:
 - The release time of J_k is after the J_i completes \Rightarrow they're already in EDF order
 - The release time of J_k is before the end of the interval in which J_i executes:



- Swap J_i and J_k (this is always possible, since J_i 's deadline is later than J_k 's)



- Move any jobs following idle periods forward into the idle period



\Rightarrow the result is an EDF schedule

2. So, if EDF fails to produce a feasible schedule, no feasible schedule exists

- If a feasible schedule existed it could be transformed into an EDF schedule, contradicting the statement that EDF failed to produce a feasible schedule

EDF may not be optimal

- When preemption is not allowed:

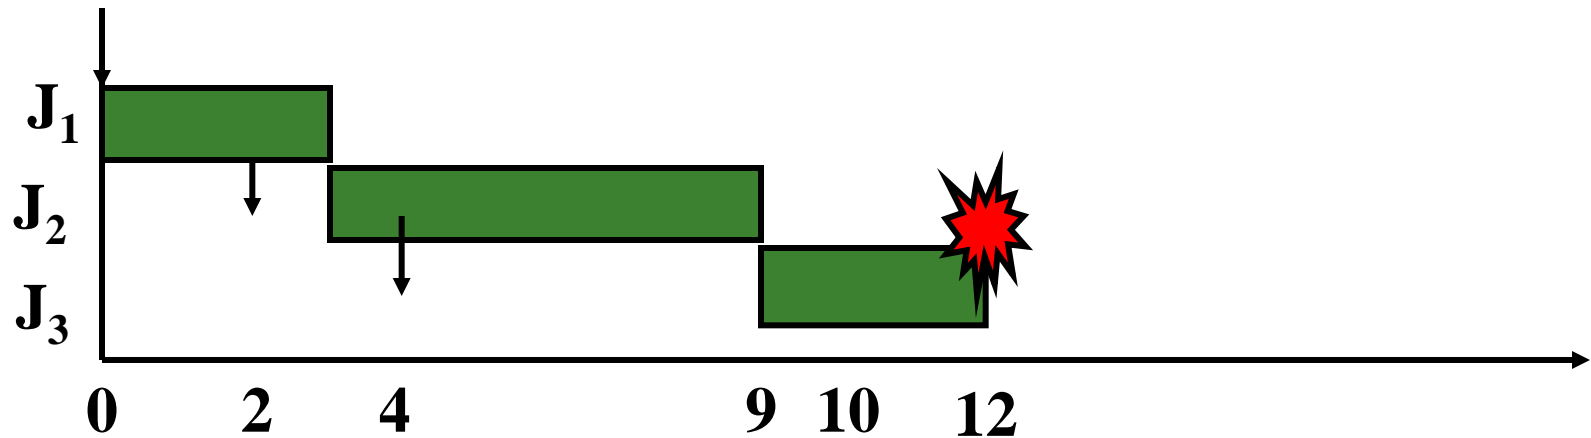
$$\begin{array}{rcl} & r_i & d_i & e_i \\ J_1 & = & (0, & 10, & 3) \\ J_2 & = & (2, & 14, & 6) \\ J_3 & = & (4, & 12, & 4) \end{array}$$

- When more than one processor is used:

$$\begin{array}{rcl} & r_i & d_i & e_i \\ J_1 & = & (0, & 4, & 1) \\ J_2 & = & (0, & 4, & 1) \\ J_3 & = & (0, & 5, & 5) \end{array}$$

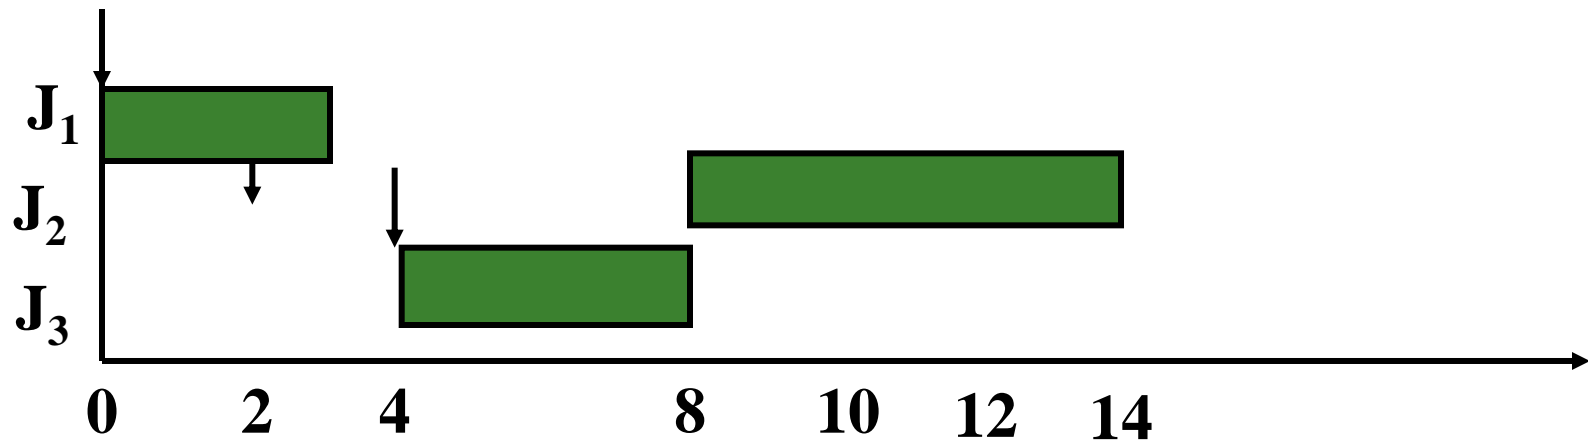
EDF + Preemption is not allowed

| | r_i | d_i | e_i |
|-------|-------|-------|-------|
| J_1 | 0 | 10 | 3 |
| J_2 | 2 | 14 | 6 |
| J_3 | 4 | 12 | 4 |



Optimal + Preemption is not allowed

| | r_i | d_i | e_i |
|-------|-------|-------|-------|
| J_1 | 0 | 10 | 3 |
| J_2 | 2 | 14 | 6 |
| J_3 | 4 | 12 | 4 |



EDF may not be optimal

- When preemption is not allowed:

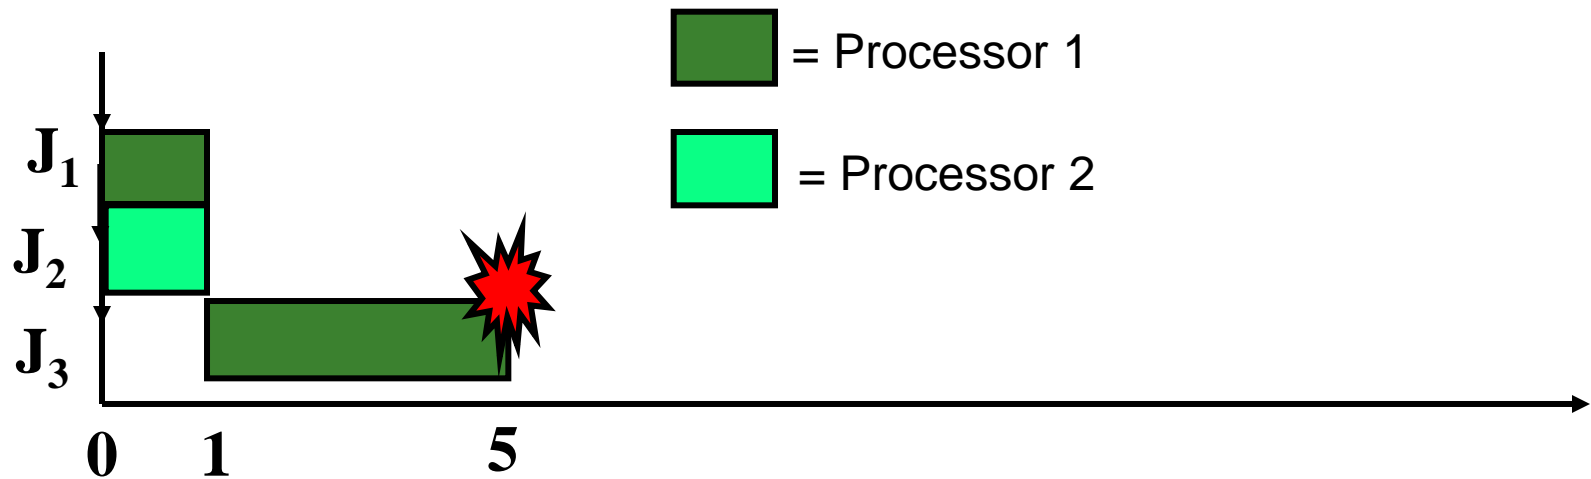
$$\begin{array}{rcl} & r_i & d_i & e_i \\ J_1 & = & (0, & 10, & 3) \\ J_2 & = & (2, & 14, & 6) \\ J_3 & = & (4, & 12, & 4) \end{array}$$

- When more than one processor is used:

$$\begin{array}{rcl} & r_i & d_i & e_i \\ J_1 & = & (0, & 4, & 1) \\ J_2 & = & (0, & 4, & 1) \\ J_3 & = & (0, & 5, & 5) \end{array}$$

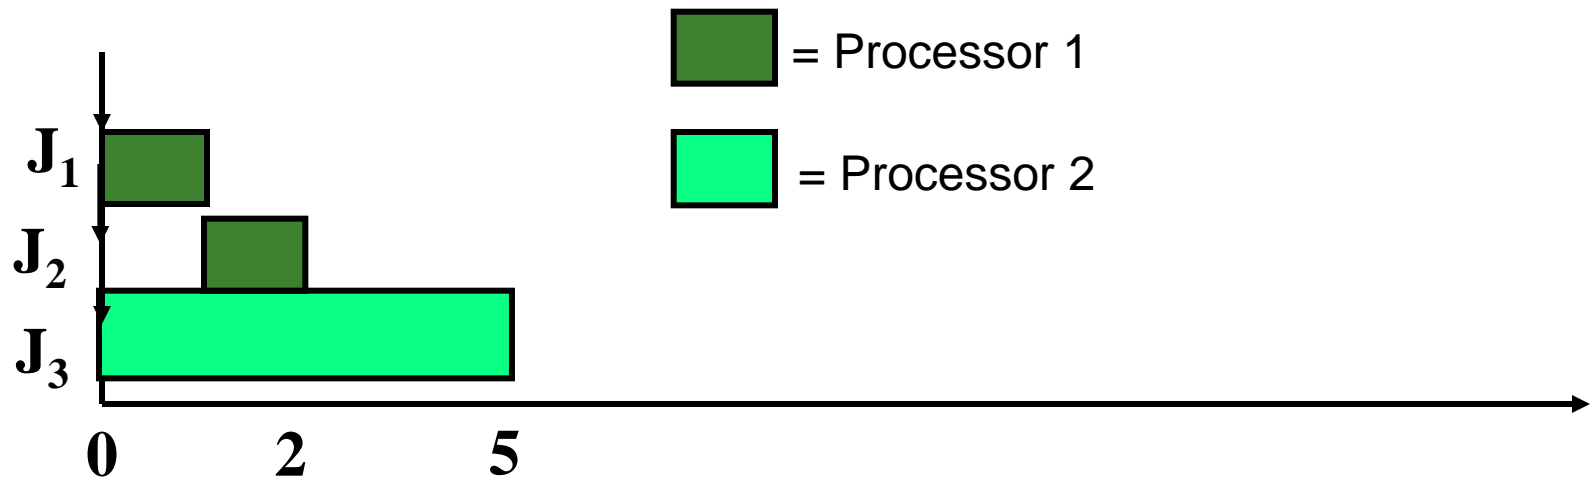
EDF + More than one processor

$$\begin{array}{lcl} & r_i & d_i \quad e_i \\ J_1 & = & (0, 4, 1) \\ J_2 & = & (0, 4, 1) \\ J_3 & = & (0, 5, 5) \end{array}$$



Optimal + More than one processor

$$\begin{array}{lcl} & r_i & d_i \quad e_i \\ J_1 & = & (0, 4, 1) \\ J_2 & = & (0, 4, 1) \\ J_3 & = & (0, 5, 5) \end{array}$$



For Next Time

- Read Ch. 1-3, 5.
 - Static Cyclic Scheduling (Ch. 5)
 - After that, Real-Time Scheduling – Commonly Used Approaches (Ch. 4)
-