

---

CIS 721 - Real-Time Systems

# Lecture 8: Preemption Thresholds

---

Mitch Neilsen  
**neilsen@ksu.edu**

---

# Outline

- Commonly Used Approaches For Real-Time Scheduling (Ch. 4)
    - Clock-Driven Scheduling (Ch. 5)
    - **Priority-Driven Scheduling**
      - **Periodic Tasks (Ch. 6)**
        - **Priority Assignment**
        - **Preemption Thresholds**
      - Aperiodic and Sporadic Tasks (Ch. 7)
-

# Periodic Task Model

- **Periodic task set:**  $\{T_1, \dots, T_n\}$ , each task consists of a set of **jobs**:  $T_i = \{J_{i1}, J_{i2}, \dots\}$
- $\phi_i$ : **phase** of task  $T_i$  = time when its first job is released
- $p_i$ : **period** of  $T_i$  = inter-release time
- $e_i$  or  $C_i$ : **execution time** of  $T_i$
- $u_i$ : **utilization** of task  $T_i$  is given by  $u_i = e_i / p_i$
- $D_i$ : (relative) **deadline** of  $T_i$ , typically  $D_i = p_i$

# Preemption Thresholds

- Y. Wang and M. Saksena, “Scheduling Fixed-Priority Tasks with Preemption Threshold”, In Proceedings of the IEEE Intl. Conf. on Real-Time Computing Systems and Applications, Dec. 1999.
- Scheduling with Preemption Thresholds
  - Task Model and Run-Time Model
  - Response Time Analysis
  - Priority and Preemption Threshold Assignment Algorithms
  - Example: ThreadX Real-Time Operating System

# Task Model

- Task Set  $\Gamma = \{\tau_1, \tau_2, \tau_3, \dots, \tau_n\}$ 
  - Each task  $\tau_i$  is characterized by  $(C_i, T_i, D_i)$ , denoted  $\tau_i \sim (C_i, T_i, D_i)$ .
  - Each task  $\tau_i$  is assigned a priority  $\pi_i \in \{1, 2, \dots, n\}$
  - and a preemption threshold  $\gamma_i \in \{\pi_i, \pi_i + 1, \dots, n\}$ .
- Notes:
  - 1 = lowest priority, n = highest priority.
  - $\pi_i$  = static priority.
  - $\gamma_i$  = dynamic priority.

# Run-Time Model

- Modified fixed-priority, preemptive scheduling.
- When task  $\tau_i$  is released, it is scheduled using its static priority  $\pi_i$ .
- After task  $\tau_i$  starts executing, another task  $\tau_j$  can preempt  $\tau_i$  only if  $\pi_j > \gamma_i \geq \pi_i$ .

# Extremes


- If  $\gamma_i = \pi_i$  for each  $i$ , then the result is **preemptive**, priority-based scheduling.
- If  $\gamma_i = n$  (max. priority) for each  $i$ , then the result is **non-preemptive**, priority-based scheduling.

# Example – assume arbitrary phasing

Task	$C_i$	$T_i$	$D_i$	$\pi_i$	WCRT Preemptive	WCRT Non-Preemptive
$\tau_1$	20	70	50	3	20	<b>55</b>
$\tau_2$	20	80	80	2	40	75
$\tau_3$	35	200	100	1	<b>115</b>	75

Task	Priority	Preemption Threshold	WCRT
$\tau_1$	3	3	40
$\tau_2$	2	3	75
$\tau_3$	1	2	95





**Figure 1. Run-time Behavior with Preemption Threshold**

# Problem Statement

- Given a task set  $\Gamma = \{\tau_1, \tau_2, \tau_3, \dots, \tau_n\}$ , determine if there exists an assignment  $\{(\pi_i, \gamma_i) \mid i = 1, 2, \dots, n\}$  such that  $\Gamma$  is schedulable.
- In other words, determine if there exists an **optimal** assignment of task priorities and preemption thresholds.

# Solutions

- **Brute Force** - try all possible assignments of priorities and preemption thresholds.
  - Time Complexity in  $O(n! n!)$   $\Rightarrow$  not feasible for large  $n$ .
- Use a **Branch and Bound Algorithm** to perform an efficient search for priorities and preemption thresholds.

# Three Step Process

## ■ Response Time Analysis


- Given assignment  $\{ (\pi_i, \gamma_i) \mid i = 1, 2, \dots, n \}$ , compute the worst-case response time ( $R_i$  or  $WCRT_i$ ) for each task  $\tau_i$ .
- A task set  $\Gamma$  is schedulable iff  $R_i \leq D_i$  for all  $i$ .

- Given a priority assignment  $\{ \pi_i \mid i = 1, \dots, n \}$ , determine a feasible set of **preemption thresholds**, if such a set exists.
- Use a branch and bound algorithm to search for a feasible assignment set of **priorities** (and preemption thresholds).

# Response Time Analysis


- The **blocking time** of task  $\tau_i$  is denoted  $B(\tau_i)$ . Blocking occurs if a lower priority task is running and task  $\tau_i$  cannot preempt it.

$$B(\tau_i) = \max_j \{ C_j / \gamma_j \mid \pi_i > \pi_j \}$$



# Busy Period Analysis

- A **critical instant** occurs when all higher priority tasks arrive at the same time, and the task that contributes to the maximum blocking arrives at the critical instant -  $\varepsilon$ .



# Divide Busy Period

- Divide the busy period for  $\tau_i$  into two parts:
  - the length of time from the critical instant (time 0) to the point when  $\tau_i$  starts executing its  $q^{\text{th}}$  job (  $\mathbf{S}_i(\mathbf{q})$  ).
  - the length of time from the time  $\tau_i$  starts executing its  $q^{\text{th}}$  job until it finishes executing its  $q^{\text{th}}$  job (  $\mathbf{F}_i(\mathbf{q}) - \mathbf{S}_i(\mathbf{q})$  ).
- Let  $q = 1, 2, \dots, m$  until we reach  $q = m$  s.t.  $F_i(m) \leq m T_i$  that is, the  $m^{\text{th}}$  job completes before the next job is released.
- Then,

$$R_i = \max_{q \in \{1, \dots, m\}} \{ F_i(q) - (q - 1)T_i \}$$

# Worst-Case Start Time ( $S_i(q)$ )

$$S_i(q) = B(\tau_i) + (q-1)C_i + \sum_{\substack{j \in \{1, \dots, n\} \\ \pi_j > \pi_i}} (1 + \left\lfloor \frac{S_i(q)}{T_j} \right\rfloor) C_j$$



# Worst-Case Finish Time ( $F_i(q)$ )

$$F_i(q) = S_i(q) + C_i + \sum_{\substack{j \in \{1, \dots, n\} \\ \pi_j > \gamma_i}} \left( \left\lceil \frac{F_i(q)}{T_j} \right\rceil - \left( 1 + \left\lfloor \frac{S_i(q)}{T_j} \right\rfloor \right) \right) C_j$$

# $WCRT(\pi_i, \gamma_i)$

Algorithm to compute  $R_i$

Input:  $C_1, \dots, C_m, T_1, \dots, T_m, \pi_1, \dots, \pi_m, \gamma_1, \dots, \gamma_m$

Output:  $R_1, R_2, \dots, R_m$

done = FALSE

q = 1

while (not done)

    compute  $S_i(q)$  and  $F_i(q)$

    if  $F_i(q) \leq q T_i$  then

        done = TRUE

        m = q

    else

        q = q + 1

    end if

end while

$R_i = \max_{q \in \{1, \dots, m\}} (F_i(q) - (q - 1) T_i)$

# Example

Task	$C_i$	$T_i$	$D_i$	$\pi_i$	WCRT Preemptive	WCRT Non-Preemptive
$\tau_1$	20	70	50	3	20	<b>55</b>
$\tau_2$	20	80	80	2	40	75
$\tau_3$	35	200	100	1	<b>115</b>	75

Task	Priority	Preemption Threshold	WCRT
$\tau_1$	3	3	40
$\tau_2$	2	3	75
$\tau_3$	1	2	95


# Preemption Threshold Assignment

- **Lemma 5.1:** Changing the preemption threshold of task  $\tau_i$  from  $\gamma_1$  to  $\gamma_2$  may only affect the worst-case response time of task  $\tau_i$  and those tasks whose priority is between  $\gamma_1$  and  $\gamma_2$ .
- **Corollary 5.1:** The worst-case response time of task  $\tau_i$  will not be affected by the preemption threshold assignment of any higher priority task; e.g., any task  $\tau_j$  with  $\pi_j > \pi_i$ .
- **This implies that we should assign preemption thresholds from lowest to highest priority.**

# Preemption Threshold



- **Theorem 5.1:** Start with a schedulable system with  $n$  tasks. If decreasing the value of  $\gamma_j$  does not change the schedulability of task  $\tau_j$ , then the whole system is still schedulable.
- **Idea:** Keep  $\gamma_j$  as small as possible for each task.
- **Lemma 5.2:** (Quick Test) If setting  $\gamma_j = n$  cannot make task  $\tau_j$  schedulable, then the task set is not schedulable.

# Preemptive Scheduling

 Preemption Threshold [-] [x]

File

Task	C	T	D	Prio	PT	B	S	F	WCRT
1	5	50	15	9	9	0	0	5	5
2	5	60	25	8	8	0	5	10	10
3	7	80	30	7	7	0	10	17	17
4	7	200	40	6	6	0	17	24	24
5	10	200	50	5	5	0	24	34	34
6	8	200	60	4	4	0	34	42	42
7	12	220	70	3	3	0	42	59	59
8	10	230	70	2	2	0	59	74	74
9	15	240	100	1	1	0	74	96	96

Compute

# Non-Preemptive Scheduling

Preemption Threshold

File

Task	C	T	D	Prio	PT	B	S	F	WCRT
1	5	50	15	9	9	15	15	20	20
2	5	60	25	8	9	15	20	25	25
3	7	80	30	7	9	15	25	32	32
4	7	200	40	6	9	15	32	39	39
5	10	200	50	5	9	15	39	49	49
6	8	200	60	4	9	15	49	57	57
7	12	220	70	3	9	15	67	79	79
8	10	230	70	2	9	15	79	89	89
9	15	240	100	1	9	0	74	89	89

Compute

## Algorithm: Assign Preemption Thresholds

*// Assumes that task priorities are already known*

- (1) **for** ( $i := 1$  to  $n$ )
- (2)      $\gamma_i = \pi_i$   
      *// Calculate worst-case response time of  $\tau_i$*
- (3)      $\mathcal{R}_i = \text{WCRT}(\tau_i, \gamma_i)$  ;
- (4)     **while** ( $\mathcal{R}_i > D_i$ ) **do**     *// while not schedulable*
- (5)          $\gamma_i++$  ; *// increase threshold*
- (6)         **if**  $\gamma_i > n$  **then**
- (7)             **return** FAIL; *// system not schedulable.*
- (8)         **endif**
- (9)          $\mathcal{R}_i = \text{WCRT}(\tau_i, \gamma_i)$  ;
- (10)    **end**
- (11) **end**
- (12) **return** SUCCESS



# Preemption Thresholds

The screenshot shows a window titled "Preemption Threshold". Below the title bar is a menu bar with "File". The main area contains a table with 10 columns: Task, C, T, D, Prio, PT, B, S, F, and WCRT. The first 9 rows are filled with numerical data. The 10th row is partially filled. The 'D' column and the 'WCRT' column are highlighted with green boxes. At the bottom center is a button labeled "Compute".

Task	C	T	D	Prio	PT	B	S	F	WCRT
2	5	60	25	8	9	12	17	22	22
3	7	80	30	7	8	12	22	29	29
4	7	200	40	6	7	12	29	36	36
5	10	200	50	5	6	12	36	46	46
6	8	200	60	4	5	12	46	59	59
7	12	220	70	3	8	10	57	69	69
8	10	230	70	2	8	0	59	69	69
9	15	240	100	1	1	0	74	96	96
10									

Compute


---

# Finding Optimal Assignment

- **Problem:** How to find an optimal assignment of task priorities and preemption thresholds.
  - **Solution:**
    - Arrange tasks into two sets -- unsorted (remaining higher priority tasks) and sorted (lower priority tasks).
    - Recursively add tasks from unsorted list to sorted list based on “lateness” heuristic.
    - Tasks are added in priority order, from lowest to highest priority.
-

# Priorities and Preemption Thresholds

- Search  $(\{\tau_1, \dots, \tau_n\}, 1)$



- Initially determine if the task set is schedulable using preemptive priority-based scheduling without preemption thresholds (e.g., can priorities be assigned using RM?).

---

# “Lateness” Heuristic

- From the unsorted list, select the task with the smallest lateness, and add it to the sorted list.
  - Since the task selected has the smallest “lateness” (delay), it should need a lower priority and smaller preemption threshold, leaving more time for higher priority tasks.
-

# Greedy Assignment Algorithm

**Algorithm: GreedyAssignment(RemainingTasks, nextPriority)**

*/\* Terminating Condition \*/*

- (1) **if** (RemainingTasks == NULL) **then**
  - /\* Call the algorithm in Figure 1 for optimal preemption threshold assignment \*/*
  - (2) **if** (AssignThresholds() == SUCCESS) **then return** SUCCESS
  - (3) **else return** FAIL
  - (4) **endif**
- (5) **endif**

# Greedy Assignment Algorithm (cont.)

*/\* Assign Heuristic Value to Each Task \*/*

- (6)   **foreach**  $\tau_k$  in RemainingTasks **do**
- (7)        $\pi_k := \text{nextPriority}$  ;   */\* tentative assignment \*/*
- (8)        $\mathcal{R}_k := \text{WCRT}(\tau_k)$ ;   */\* compute response time \*/*
- (9)       **if**  $\mathcal{R}_k \geq D_k$  **then**  $h\_val_k := D_k - \mathcal{R}_k$
- (10)      **else**  $h\_val_k := \text{GetBlockingLimit}(\tau_k)$  ;
- (11)      **endif**
- (12)       $\pi_k := \pi$  ;   */\* reset, to allow computing heuristic value for other tasks \*/*
- (13) **end**

# Get Blocking Limit Function

Input :  $\tau_k, D_k$

Output : Blocking limit of  $\tau_k$

$R_k = WCRT(\tau_k)$

$Max = D_k - R_k$

$Limit = 0$

**For**  $B(\tau_k) = 1$  **to**  $Max$

$R_k = WCRT(\tau_k, B(\tau_k))$

**If**  $R_k > D_k$  **Then** *Break*

**Else**  $Limit = B(\tau_k)$

**End For**

**Return**  $Limit$

# Greedy Assignment Algorithm (cont.)

```
(14)  /* Select the task with the largest heuristic value next */
(15)   $\tau_k := \text{max\_heuristic\_val}(\text{RemainingTasks})$  ;
(16)   $\pi_k := \text{nextPriority}$  ;   /* final priority assignment */
(17)  /* Recursively Assign Priorities to Remaining Tasks */
(18)  if GreedyAssignment(RemainingTasks  $\leftarrow \tau_k, \text{nextPriority}+1$ ) == SUCCESS then
(19)      return SUCCESS ;
(20)  return FAIL ;
```

---



# Note

- There are cases when this heuristic algorithm is not able to find a feasible assignment, even though a non-preemptive priority assignment algorithm is able to find a solution.
- Thus, we could apply a non-preemptive assignment algorithm first, before using this heuristic algorithm.

---

# Depth-First Search

- Perform a depth-first search to find an optimal priority assignment.
- When a leaf is reached, call `AssignThresholds( )` to see if an optimal preemption threshold assignment exists.

---

# Summary

## ■ To Do:

- ❑ Read Ch. 5-7 + Wang and Saksena's paper on scheduling with preemption thresholds.
  - ❑ Homework #1, Homework #2.
-