
CIS 721 - Real-Time Systems

Lecture 13: VxWorks RTOS

Mitch Neilsen
neilsen@ksu.edu

Outline

- **Priority Inheritance Protocols**
 - Utilization-Based Test
 - **Response Time Analysis**
 - **Real-Time Operating System (VxWorks)**
 - **System Development Environment**
 - **Priority Inheritance Protocols**
 - **Priority Ceiling Protocols**
-

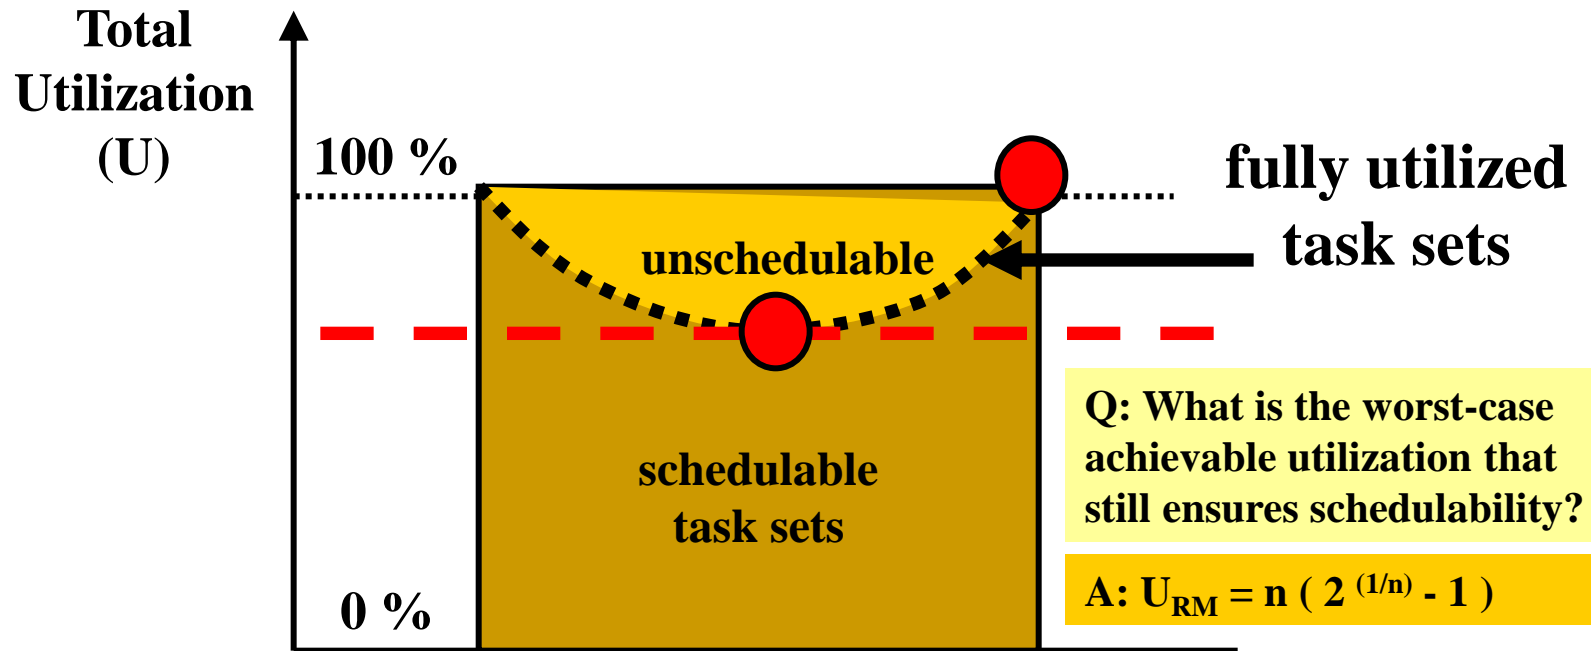
Utilization-Based Test

- Given a periodic task τ_i , the ratio $u_i = C_i / T_i$ is called the **utilization of task τ_i** .
- The **total utilization U** of all tasks in a system is the sum of the utilizations of all individual tasks:

$$U = \sum_{i=1}^n \frac{C_i}{T_i}$$

Utilization-Based Test

If $U + \max\{B_1/T_1, B_2/T_2, \dots, B_n/T_n\} \leq n(2^{1/n} - 1)$,
then the task set is feasible.



Example

Task	Period	Blocking Time	Run-Time
τ_i	T_i	B_i	C_i
<hr/>			
τ_1	100	20	40
τ_2	150	30	40
τ_3	350	0	100

$$\begin{aligned} U &= 40 / 100 + 40 / 150 + 100 / 350 \\ &= 0.4 + 0.267 + 0.286 = 0.953, \text{ so} \\ U + \max\{B_1/T_1, B_2/T_2, \dots, B_n/T_n\} \\ &= 0.953 + 0.2 > 1.0. \end{aligned}$$

Consequently, the test is **inconclusive**.

Response Time Analysis

- The **response time** (R_i) for task τ_i is given by the implicit equation:

$$R_i = C_i + B_i + \sum_{j \in hp(i)} \left\lceil \frac{R_i}{T_j} \right\rceil * C_j$$

- Solve by forming a recurrence relation:

$$w_i^{n+1} = C_i + B_i + \sum_{j \in hp(i)} \left\lceil \frac{w_i^n}{T_j} \right\rceil * C_j$$

$$w_i^0 = C_i + B_i$$

Note: Arbitrary phasing is assumed.

Solving The Recurrence

■ The sequence $w_i^0, w_i^1, w_i^2, \dots, w_i^n$

is clearly non-decreasing:

- If $w_i^{n+1} = w_i^n$, then a fixed point (solution) has been found.
- If $w_i^{n+1} > T_i$, then no solution exists.

Algorithm

Input: $C_1, C_2, \dots, C_m, B_1, B_2, \dots, B_m, T_1, T_2, \dots, T_m$

Output: R_1, R_2, \dots, R_m

for $i = 1$ to m

$n = 0$

$w_i^n = C_i + B_i$

loop

$$w_i^{n+1} = C_i + B_i + \sum_{j \in hp(i)} \left\lceil \frac{w_i^n}{T_j} \right\rceil * C_j$$

if $w_i^{n+1} = w_i^n$ then

$R_i = w_i^n$

break out of loop { solution found }

if $w_i^{n+1} > T_i$ then

break out of loop { no solution }

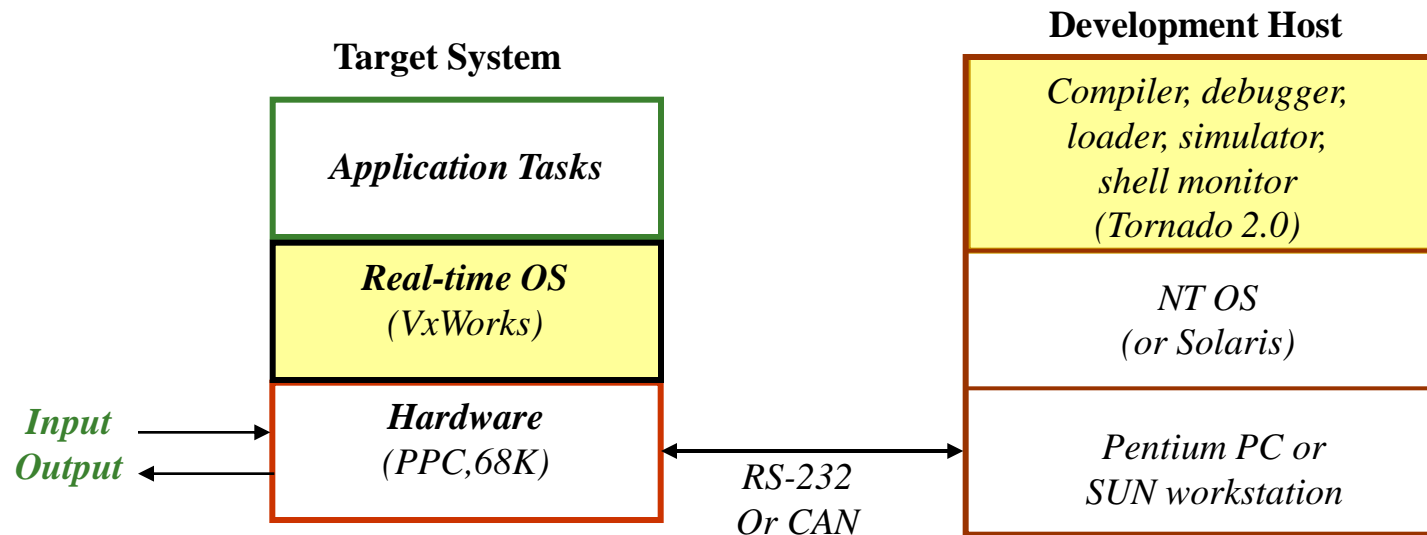
$n = n + 1$

end loop

end for

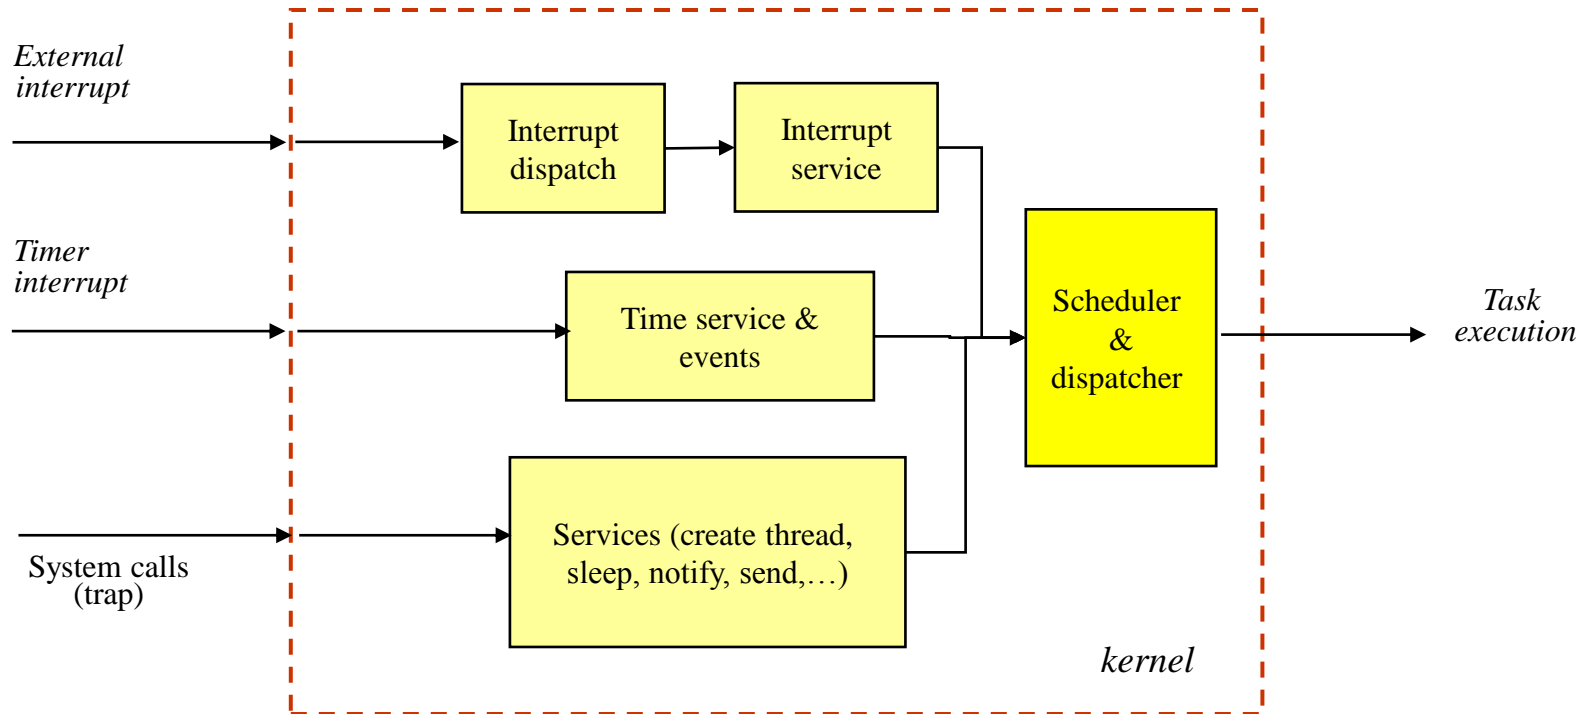
Real-Time System Development

- Real-Time Operating System (VxWorks)
- Real-Time Development Environment (Tornado 2.0)




Real-Time OS (RTOS)

- Functions: **task management**, time service, device drivers, and interrupt service.



Tornado Architecture



VxWorks Task

- A **task** is a *context of execution*. It has:
 - ❑ a **program counter** (PC).
 - ❑ private copies of other CPU **registers**.
 - ❑ a **stack** for local variables, function arguments, and function call chain information.
 - Single processor system: only one task is scheduled for execution at any given time.
 - ❑ When a task is not executing, its context is stored in its *Task Control Block* (TCB) and stack. The TCB is the data structure which the kernel uses to represent and control the task.
-

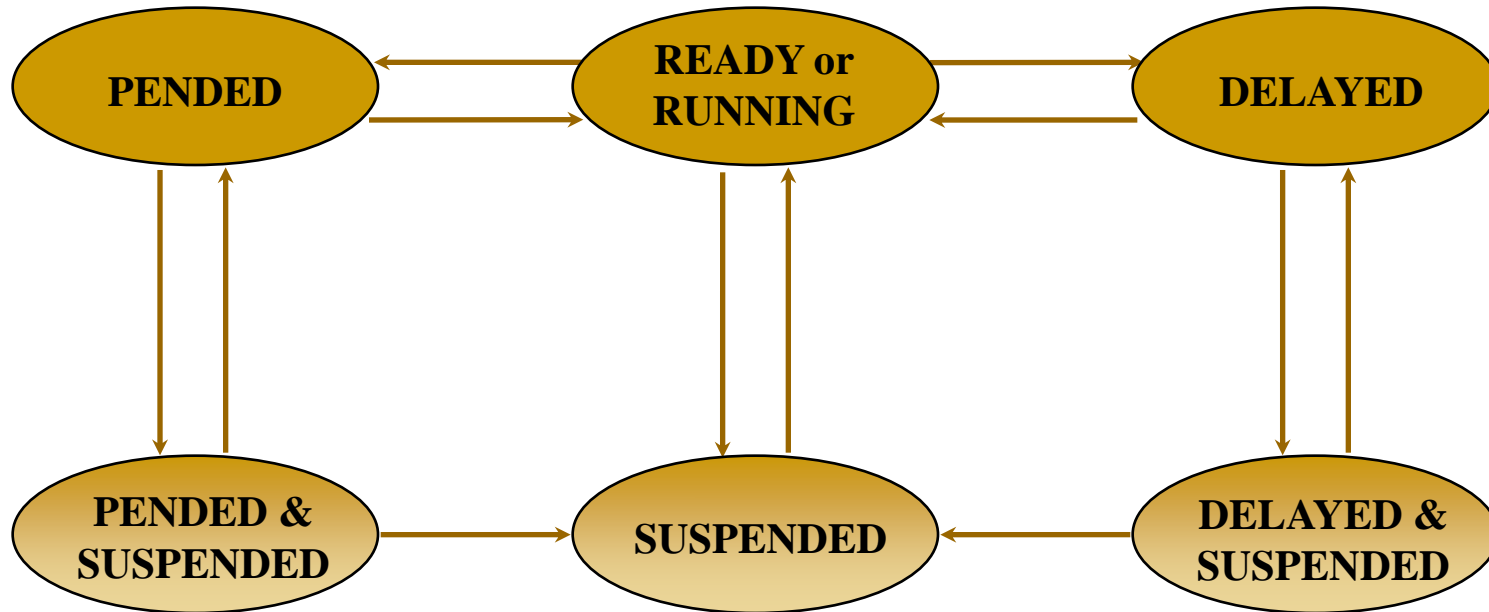
Task Creation Routines

- **taskSpawn()** - spawn (create and activate) a new task.
 - **taskInit()** - initialize (create and suspend) a new task.
 - **taskActivate()** - activate an initialized task.
-

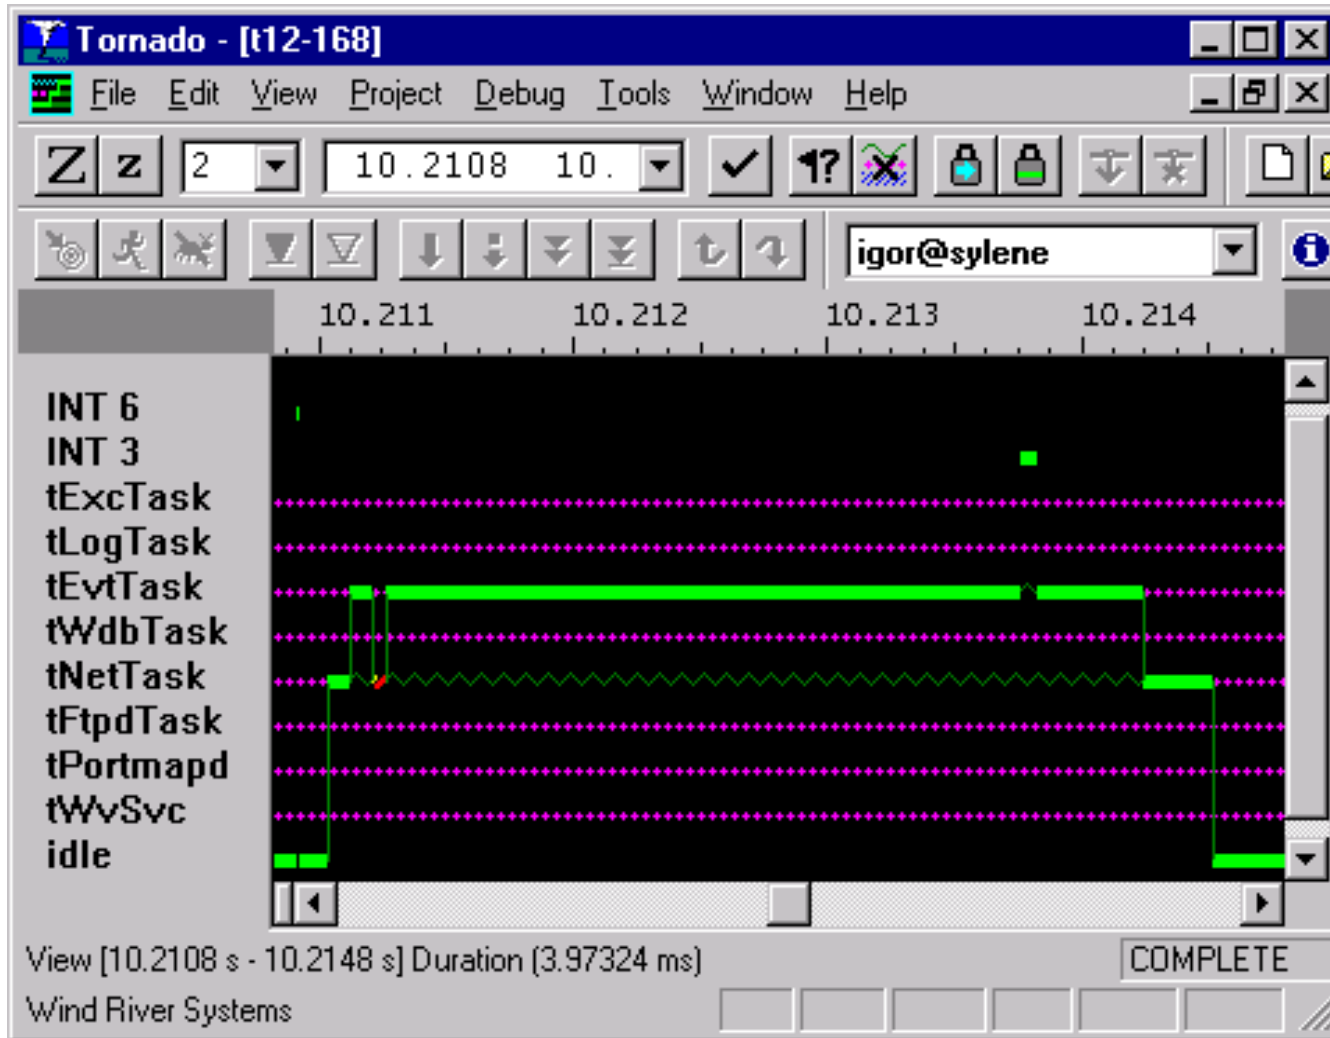
Task State Transitions

- **READY** - The state of a task that is not waiting for any resource other than the CPU.
 - **PEND** - The state of a task that is blocked due to the unavailability of some resource.
 - **DELAY** - The state of a task that is asleep for some duration.
 - **SUSPEND** - The state of a task that is unavailable for execution. This state is used primarily for debugging. Suspension does not inhibit state transition, only task execution. Thus pended-suspended tasks can still unblock and delayed-suspended tasks can still awaken.
-

Task State Transition Diagram



Preemptive Priority Scheduling



At any time, the highest priority task *ready* to run, executes!

A higher priority task that becomes ready preempts an executing task.

Interrupts preempt any task.

Notes: On this board, INT 6 is the system clock, INT 3 is the network interrupt.

Round Robin Scheduling

- Preemptive priority scheduling can be augmented with round-robin scheduling.
 - Round-robin scheduling is used to share the CPU fairly among all ready tasks at the **same priority level** using time slicing.
 - Each task of a group of tasks executes for a defined interval, or time slice; then another task executes for its interval in the cycle.
 - The allocation is fair in that no task in the group gets a second slice of time before the other tasks in the group are given a slice.
-

Kernel Timeslice

- Round-robin scheduling can be enabled with the routine **kernelTimeSlice()**, which takes a parameter for a timeslice or interval length.
- This interval (a variable-length quantum) is the amount of time each task is allowed to run before relinquishing the processor to another equal-priority task.

taskLib

- taskLib contains the kernel functions for creating, destroying, starting and stopping tasks.
- Example routines:
 - To create and start a new task:
 - int **taskSpawn** (name, priority, options, stackSize, entryPt, arg0, ... , arg9)
 - To delay the executing task for a certain number of system clock ticks:
 - STATUS **taskDelay** (ticks)

Semaphores

- Semaphores are VxWorks kernel objects that allow blocking and unblocking of tasks; they are used to coordinate actions of a given task with those of other tasks and with external events.
- VxWorks provides three varieties of semaphores:
 - Binary (synchronization) semaphores.
 - Counting semaphores.
 - Mutex (mutual exclusion) semaphores.
- Each type of semaphore is intended primarily for a particular kind of programming problem.

Binary Semaphores

- Binary semaphores allow tasks to wait for an event without taking up CPU time polling for the event to occur.
- The event could be generated by an interrupt or another task.
- Usage:
 - Create the binary semaphore using **semBCreate()**
 - Call **semTake()** to block until the event occurs
 - Signal a blocked task using **semGive()**

Example

```
#include "vxWorks.h"
```

```
#include "semLib.h"
```

```
SEM_ID semMutex;
```

```
/* Create binary semaphore initialized to 1, enqueue blocked tasks in  
priority order */
```

```
semMutex = semBCreate (SEM_Q_PRIORITY, SEM_FULL);
```

```
semTake (semMutex, WAIT_FOREVER);
```

```
.. critical section ..
```

```
semGive (semMutex);
```

Mutex Semaphores

- Mutex semaphores are used when multiple tasks share a resource (data structure, file, hardware).
- When used correctly, mutex semaphores prevent multiple tasks from accessing the resource at the same time, and possibly corrupting it. Usage:
 - ❑ Create mutex for the resource with **semMCreate()**.
 - ❑ A task wanting to use the resource calls **semTake()** to block until the resource is available (or time-out).
 - ❑ When done with the resource, a task calls **semGive()** to allow other tasks to use the resource. Any task locking a resource must also unlock the resource.

Mutex Semaphores

While binary or counting semaphores may be used for mutual exclusion, mutex semaphores are designed to deal with three common problems which arise in mutual exclusion:

- Mutual exclusion semaphores can be locked in a **recursive fashion**. A task that owns a mutex can acquire it again without blocking. The task must release the mutex as many times as it has acquired it.
 - If a task tA owns a **delete-safe** mutex semaphore, and another task tB tries to delete tA , then tB will block until tA gives up ownership of the mutex. Thus, the mutex owner is safe from deletion while operating on the shared resource.
 - Mutual exclusion semaphores can use a priority inversion protocol to be **inversion-safe**.
-

Priority Inheritance Protocol

- *Inversion-safe* mutexes prevent the priority inversion problem. A task which owns such a mutex is temporarily boosted to the priority of the highest priority task waiting for the mutex. It falls back to its regular priority when it no longer owns an inversion-safe mutex. This prevents a high priority task from being blocked for an indeterminate amount of time waiting for a mutex owned by a low priority task which cannot execute because it has been preempted by an intermediate priority task!
- The following example creates a mutual-exclusion semaphore that uses the priority inheritance algorithm:

```
semId = semMCreate (SEM_Q_PRIORITY | SEM_INVERSION_SAFE);
```

Counting Semaphores

- **Binary semaphores** keep track of whether or not an event has occurred, but not *how many times* the event has occurred (since the last time the event was serviced).
- **Counting semaphores** keep a count of how many times the event has occurred, but not serviced.
 - ❑ May be used to ensure that the event is serviced as many times as it occurs.
 - ❑ May also be used to maintain an atomic count of multiple equivalent available resources.

Semaphore Operations

- **semTake(S.semId, timeout)**
 - if (S.count > 0) then S.count --;
 - else block calling task in S.waitingQ;
- **semGive(S.semId)**
 - if (S.waitingQ is non-empty) then
 - wakeup a process in S.waitingQ;
 - else S.count ++;

Semaphore Options

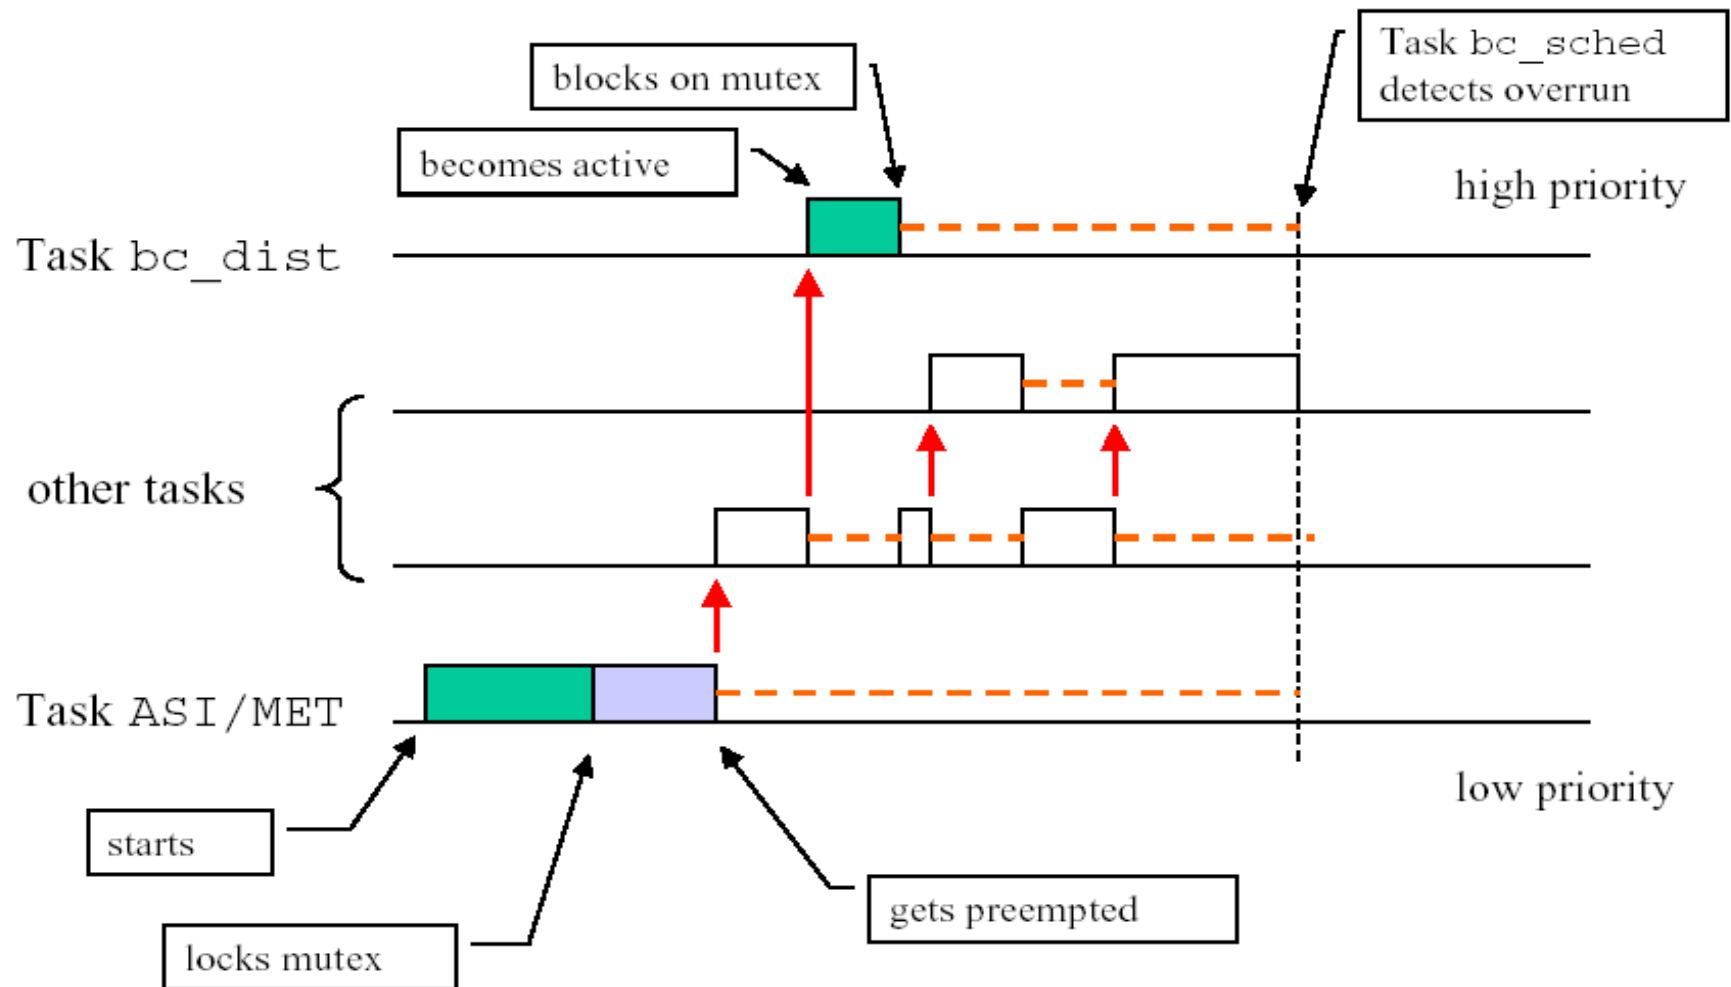
- Tasks pending on a semaphore can be queued in priority order (**SEM_Q_PRIORITY**) or in first-in first-out order (**SEM_Q_FIFO**).
- When tasks try to lock (take) a semaphore, they can specify a specific length of time to wait or **WAIT_FOREVER**.
- Priority inheritance protocols can be enabled using **SEM_INVERSION_SAFE**, this option can only be used with a priority queue; e.g., **SEM_Q_PRIORITY**.

Example: Mars Pathfinder



- Based on VxWorks.
- Developed a mysterious communications problem shortly after successfully landing on Mars on July 4, 1997.
- The problem was due to a classical case of priority inversion which caused a high priority thread to detect a missed deadline and reset the system.
- The design team dynamically modified the system through the debugging software which was still enabled.

Priority Inversion Problem



Example: myproject\cobble.c

```
STATUS progStart (void)
{
    syncSemId = semBCreate (SEM_Q_FIFO, SEM_EMPTY);

    mutexSemId = semMCreate ( SEM_Q_PRIORITY
                              | SEM_DELETE_SAFE);
    runState = 0;

    taskOneId = taskSpawn("taskOne", 200, 0, STACK_SIZE,
                          (FUNCPTR) taskOne,0,0,0,0,0,0,0,0,0,0);

    taskTwoId = taskSpawn("taskTwo", 220, 0, STACK_SIZE,
                          (FUNCPTR) taskTwo,0,0,0,0,0,0,0,0,0,0);

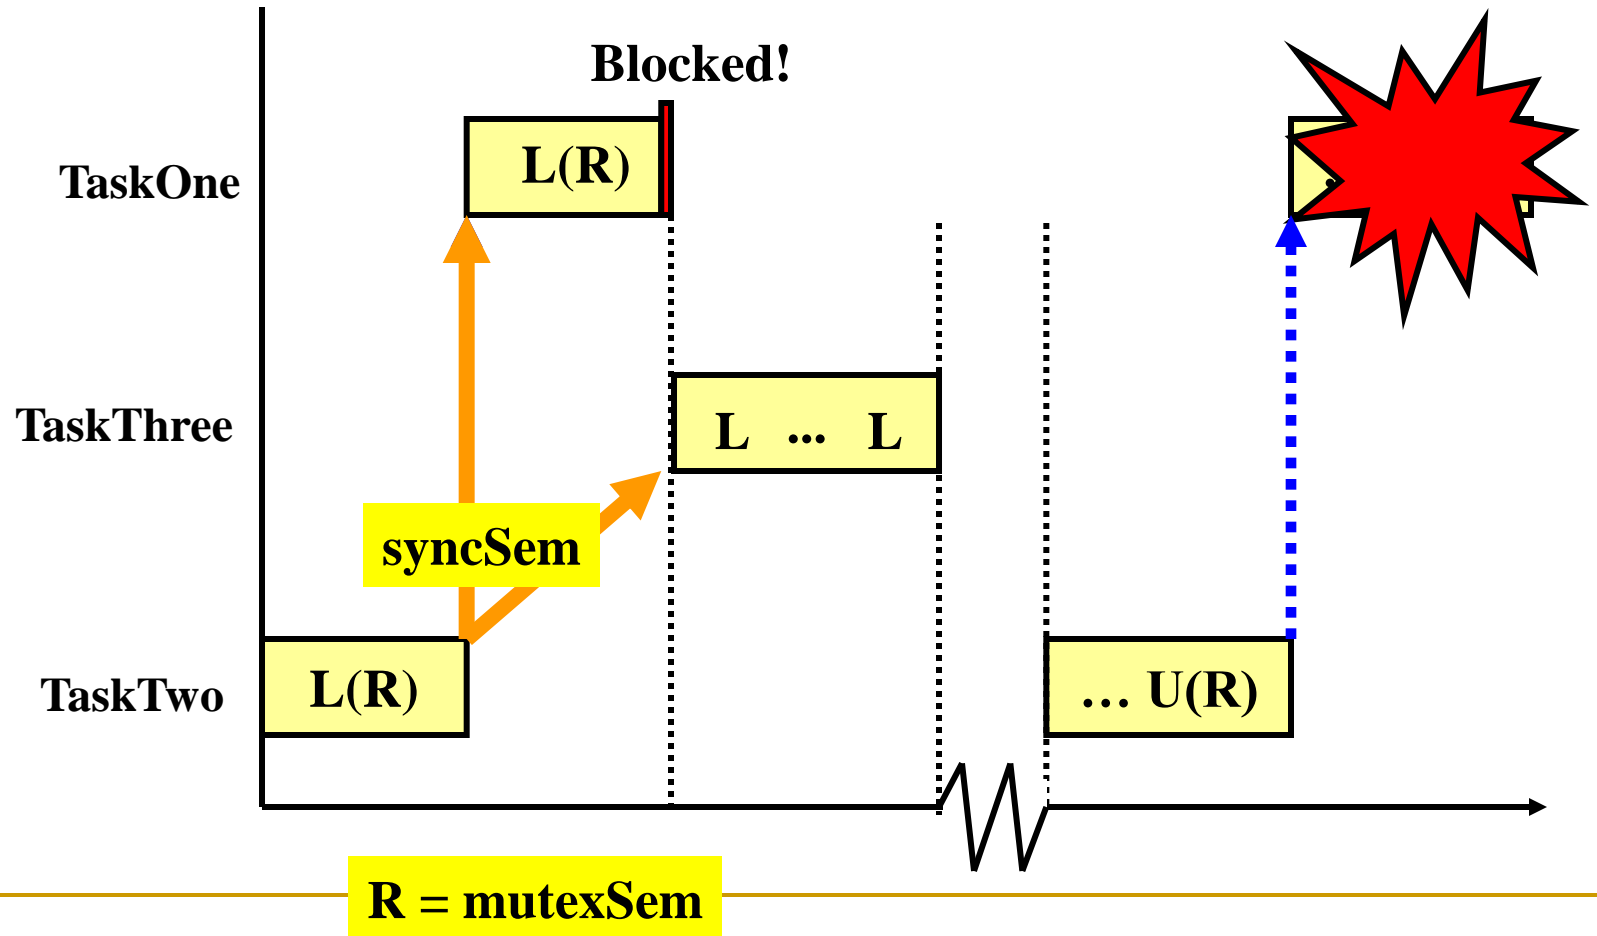
    taskThreeId = taskSpawn("taskThree", 210, 0, STACK_SIZE,
                           (FUNCPTR) taskThree,0,0,0,0,0,0,0,0,0,0);
    return (OK);
}
```

**High
Priority**

**Low
Priority**

**Medium
Priority**

Priority Inversion



TaskOne – High Priority

```
void taskOne()  
{  
    int i,j;  
  
    semTake (syncSemId, WAIT_FOREVER);  
    semTake (mutexSemId, WAIT_FOREVER);  
    for (i=1; i<=10; i++)  
    {  
        printf("TaskOne in CS, i = %d\n",i); critical  
        for (j=1; j<=10000; j++); section  
    }  
    semGive (mutexSemId);  
    for (i=11; i<=20; i++)  
    {  
        printf("TaskOne out of CS, i = %d\n",i);  
        for (j=1; j<=10000; j++);  
    }  
    runState++;  
}
```

TaskTwo – Low Priority

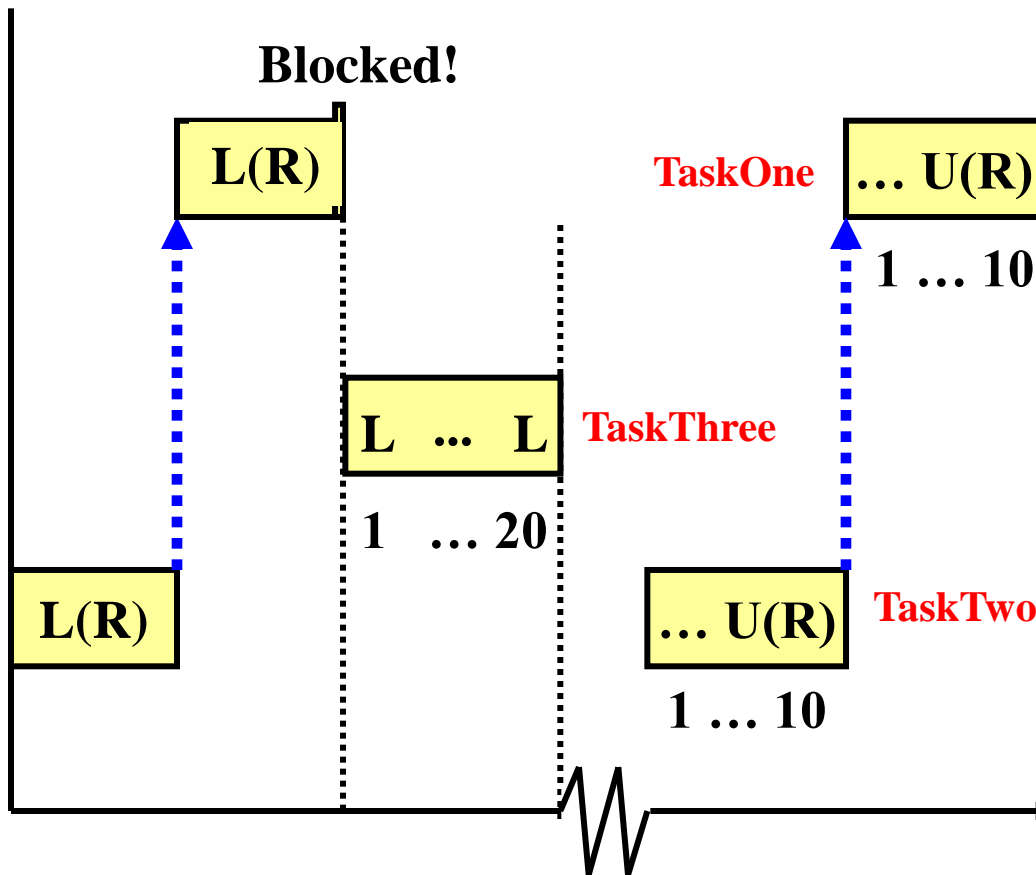
```
void taskTwo()
{
    int i,j;

    semTake (mutexSemId, WAIT_FOREVER);
    semGive (syncSemId);
    semGive (syncSemId,);
    for (i=1; i<=10; i++)
    {
        printf("TaskTwo in CS, i = %d\n",i);
        for (j=1; j<=10000; j++);
    }
    semGive (mutexSemId);
    for (i=11; i<=20; i++)
    {
        printf("TaskTwo out of CS, i = %d\n",i);
    }
    runState++;
}
```

TaskThree – Middle Priority

```
void taskThree()  
{  
    int i,j;  
  
    semTake (syncSemId, WAIT_FOREVER);  
    for (i=1; i<=20; i++)  
    {  
        printf("TaskThree running, i = %d\n",i);  
        for (j=1; j<=10000; j++);  
    }  
    runState++;  
}
```

Output



TaskThree running, $i = 1$
TaskThree running, $i = 2$
...
TaskThree running, $i = 20$
TaskTwo in CS, $i = 1$
...
TaskTwo in CS, $i = 10$
TaskOne in CS, $i = 1$
...
TaskOne in CS, $i = 10$
TaskOne out of CS, $i = 11$
...
TaskOne out of CS, $i = 20$
TaskTwo out of CS, $i = 11$
...
TaskTwo out of CS, $i = 20$

Fix: myproject\cobble.c

```
STATUS progStart (void)
{
    syncSemId = semBCreate (SEM_Q_FIFO, SEM_EMPTY);

    mutexSemId = semMCreate (    SEM_Q_PRIORITY
                                | SEM_INVERSION_SAFE
                                | SEM_DELETE_SAFE);

    runState = 0;

    taskOneId = taskSpawn("taskOne", 200, 0, STACK_SIZE,
                          (FUNCPTR) taskOne, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0);

    taskTwoId = taskSpawn("taskTwo", 220, 0, STACK_SIZE,
                          (FUNCPTR) taskTwo, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0);

    taskThreeId = taskSpawn("taskThree", 210, 0, STACK_SIZE,
                            (FUNCPTR) taskThree, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0);

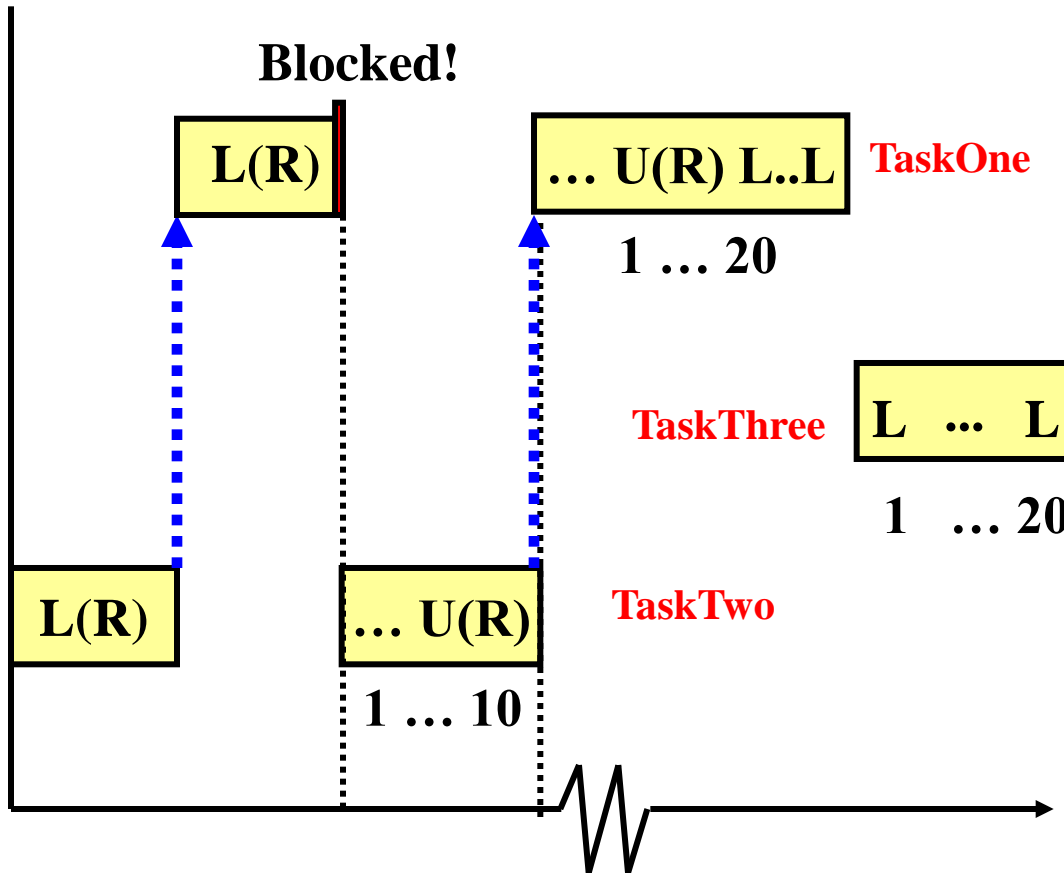
    return (OK);
}
```

**High
Priority**

**Low
Priority**

**Medium
Priority**

Output



TaskTwo in CS, i = 1

...

TaskTwo in CS, i = 10

TaskOne in CS, i = 1

...

TaskOne in CS, i = 10

TaskOne out of CS, i = 11

...

TaskOne out of CS, i = 20

TaskThree running, i = 1

TaskThree running, i = 2

...

TaskThree running, i = 20

TaskTwo out of CS, i = 11


...

TaskTwo out of CS, i = 20

Priority Inheritance – Can Still Deadlock

LOW Priority Task

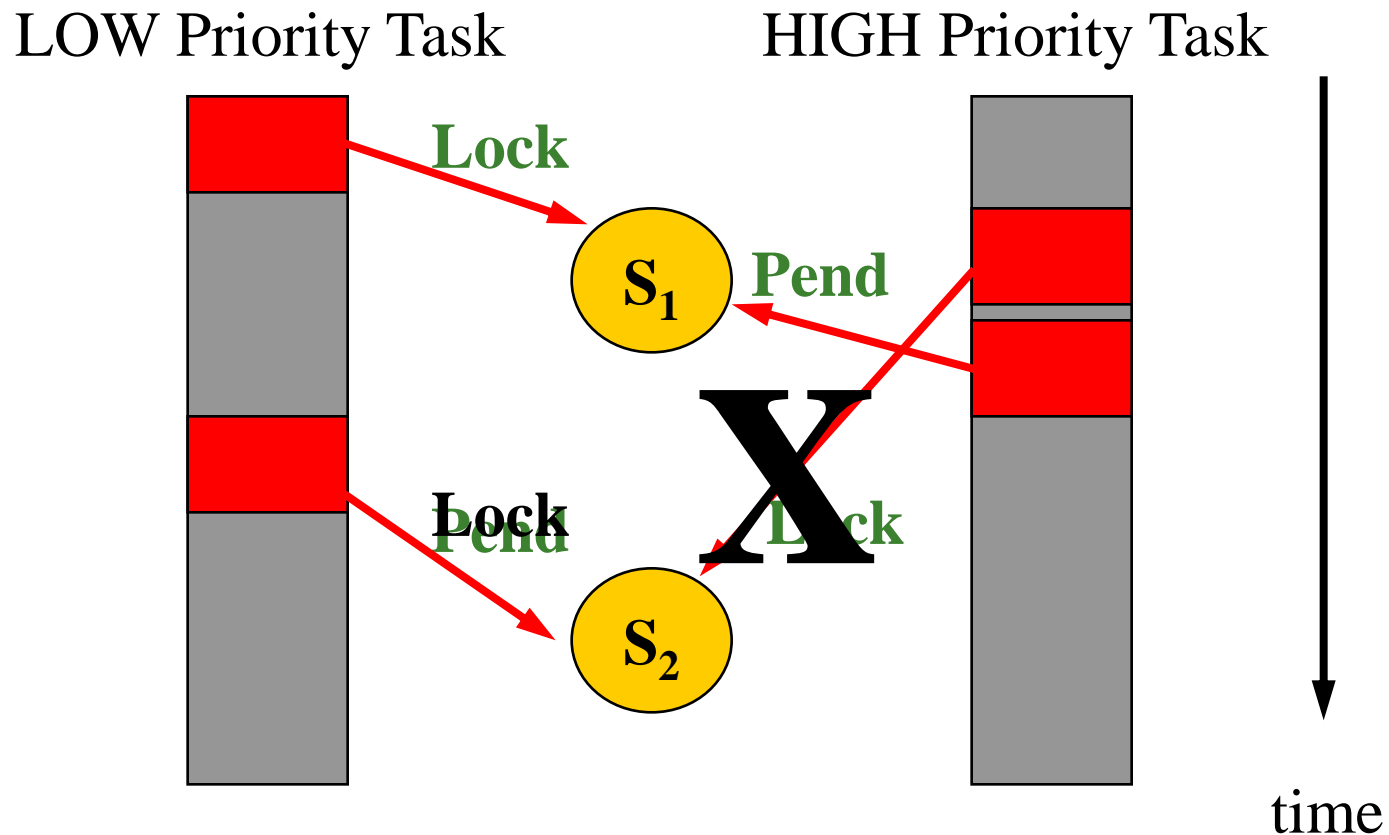
HIGH Priority Task



Priority Ceiling Protocols

- Priority Ceiling Protocols prevent: Deadlock, Transitive Blocking, and Priority Inversion.
- Each task has a static priority.
- Each resource has a static ceiling value = the maximum priority of any task that may use it.
- Each task also has a dynamic priority = the maximum of its own static priority and the ceiling value of any resource it has locked.
- Max. blocking time $B_i = \max \text{usage}(k,i) * CS(k)$

Priority Ceiling Protocols – No Deadlocks



VxWorks System Functions

❑ Create and Activate a Task:

```
taskSpawn ("tctrlSW", PRIO, 0, STACK_SIZE,  
           (FUNCPTR) controlSoftware,0,0,0,0,0,0,0,0,0,0);
```

❑ Create a Mutual Exclusion Semaphore:

```
inverSemId = semMCreate  
            (SEM_Q_PRIORITY | SEM_INVERSION_SAFE);
```

❑ Set Task Priority:

```
taskPrioritySet(taskID, task_priority);
```

❑ Measure the Execution Time of a Task:

```
timexN((FUNCPTR) controlSoftware, NULL,NULL,  
       NULL,NULL,NULL,NULL,NULL,NULL);
```

Example 2: Priority Inversion

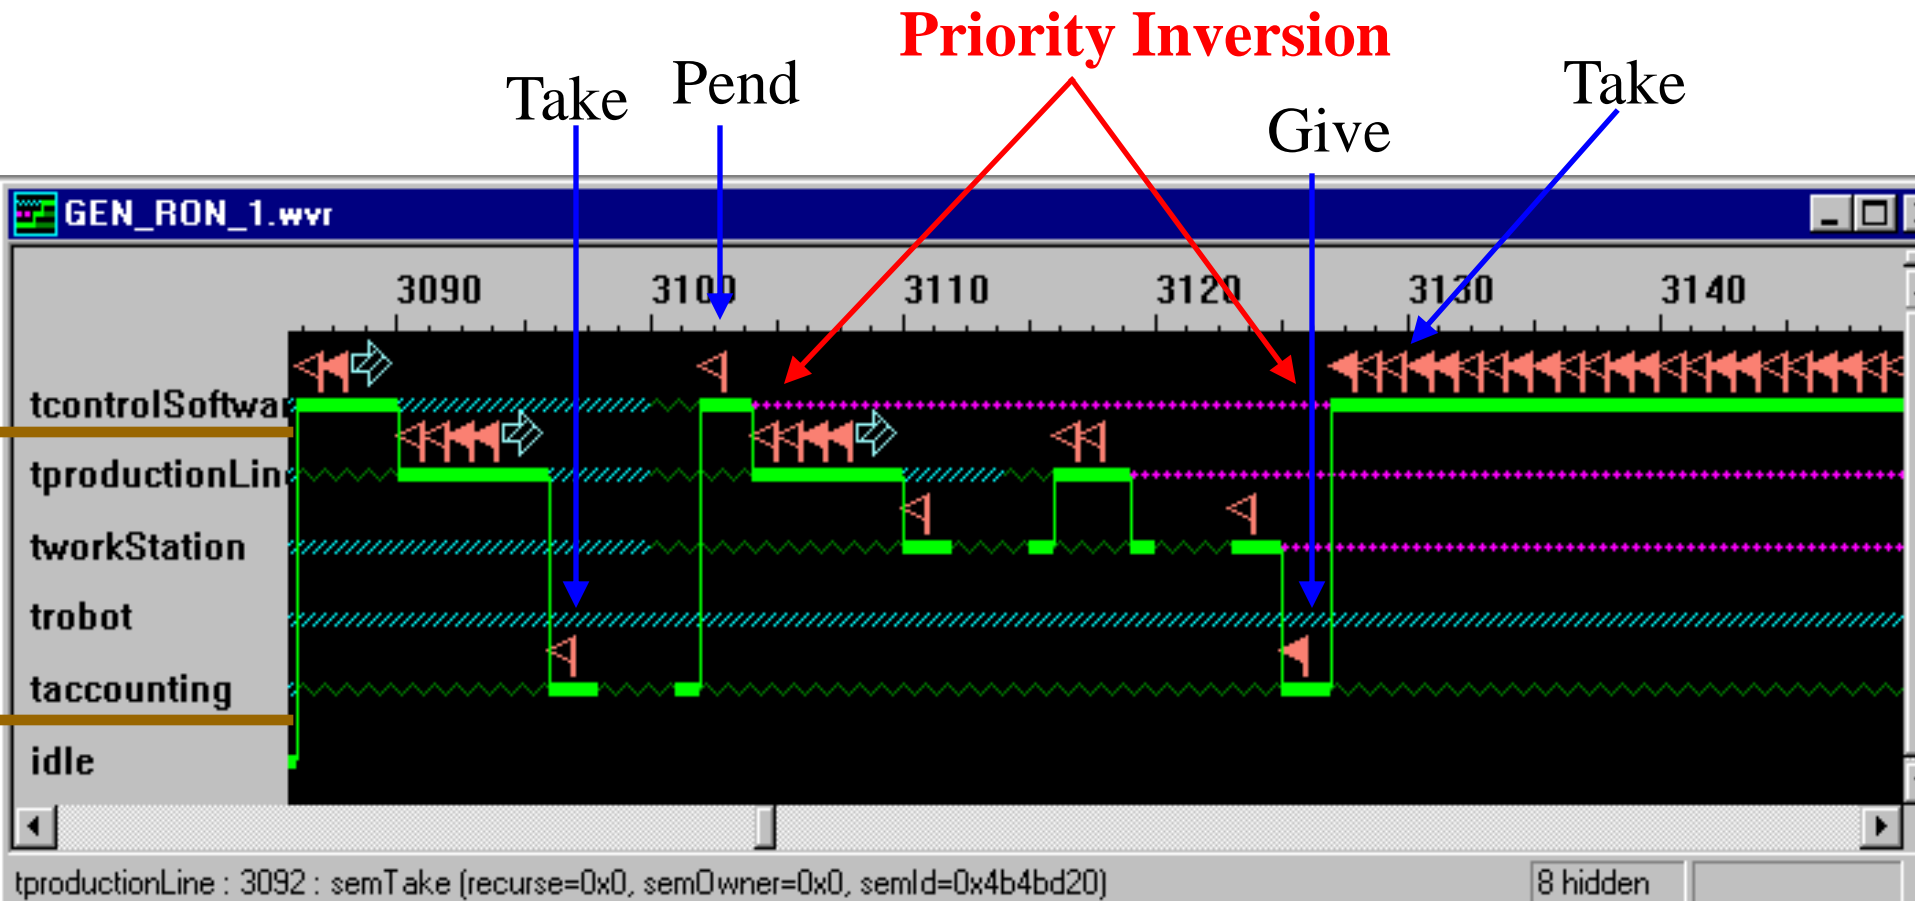
```
void controlSoftware(void) { //HIGH priority
    FOREVER {
        semTake(inverSemId, WAIT_FOREVER);
        for (long i=0; i < LONG_TIME; i++); //wait for a while
        semGive(inverSemId);
    }
}

void accounting(void) { //LOW priority
    FOREVER {
        semTake(inverSemId, WAIT_FOREVER);
        for (long i=0; i < 6 * LONG_TIME; i++);
        semGive(inverSemId);
    }
}
```

Suppose that the HIGH and LOW priority tasks compete for the same semaphore, and MEDIUM priority tasks run frequently.

Example 2: Priority Inversion

Semaphore Operations:



Example 2: Priority Inheritance

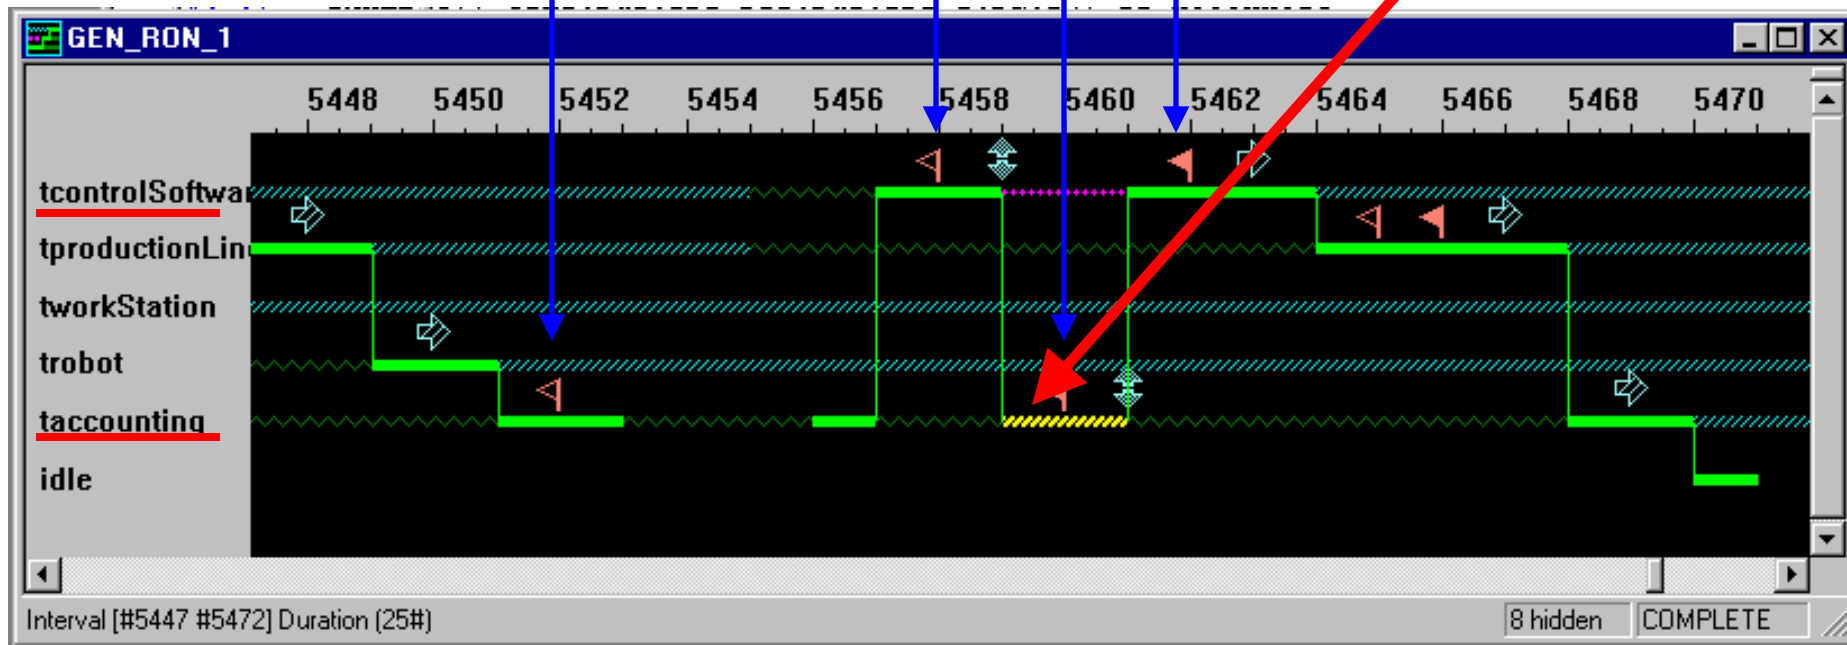
Example: `inverSemId = semMCreate (SEM_Q_PRIORITY
| SEM_INVERSION_SAFE);`

Semaphore: Take

Pend

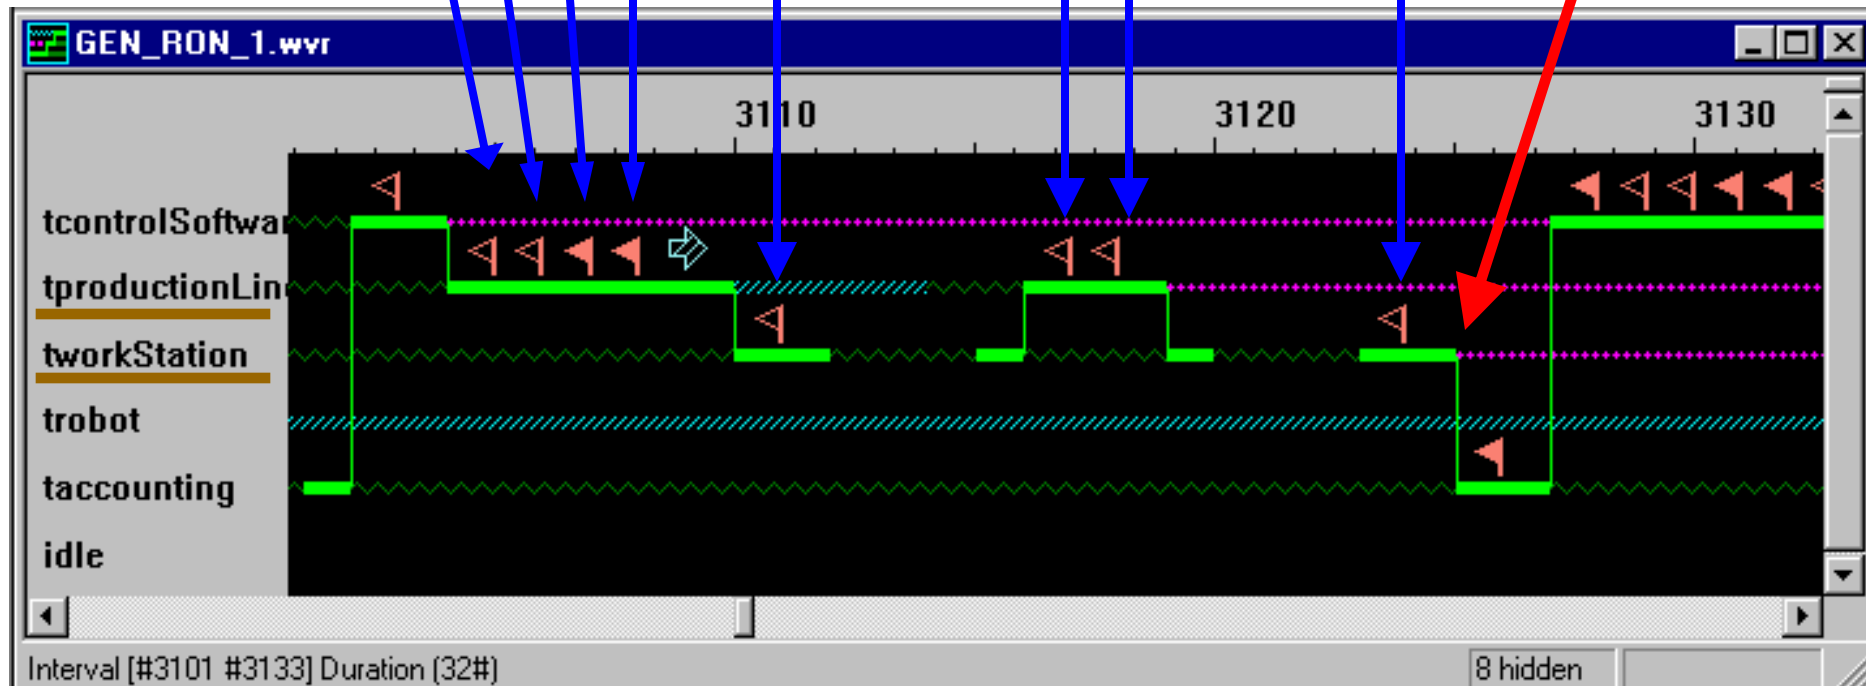
Give Give

**Priority
Inheritance**



Priority Inheritance and Deadlock

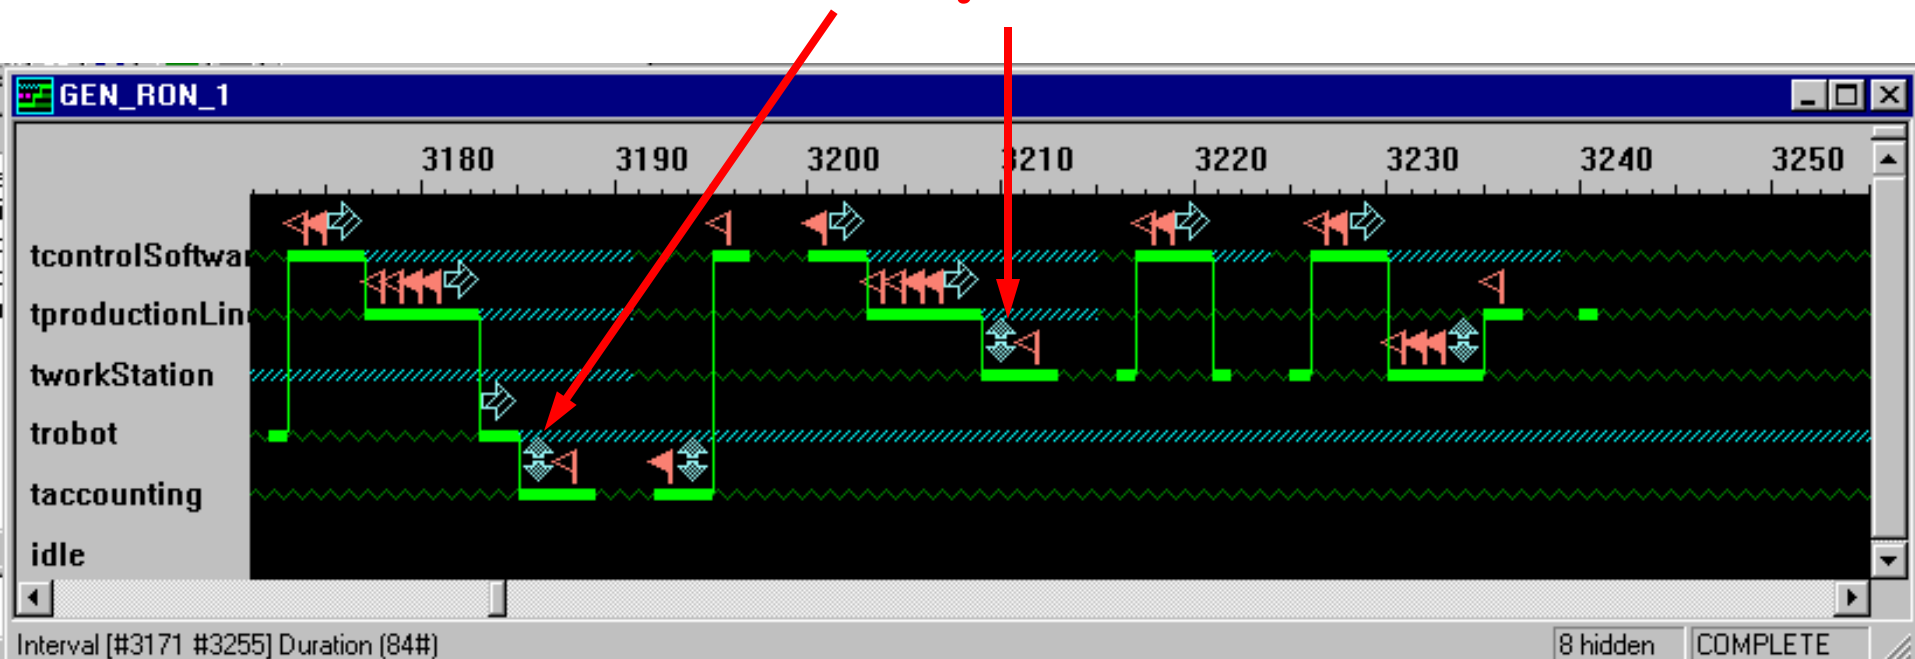
Sem 1 Take Give Take Pend
Sem 2 Take Give Take Pend Deadlock



Avoiding Deadlock: Emulating IPCP

```
void workStation(void) {
    taskPrioritySet(tidWs,MID1); /* IPCP dynamic priority */
    semTake(dlock2SemId, WAIT_FOREVER); ...
    semTake(dlock1SemId, WAIT_FOREVER); ...
    taskPrioritySet(tidWs,MID2); /* IPCP dynamic priority reset to original */ ...}
```

Priority Set



Priority Ceiling Protocols

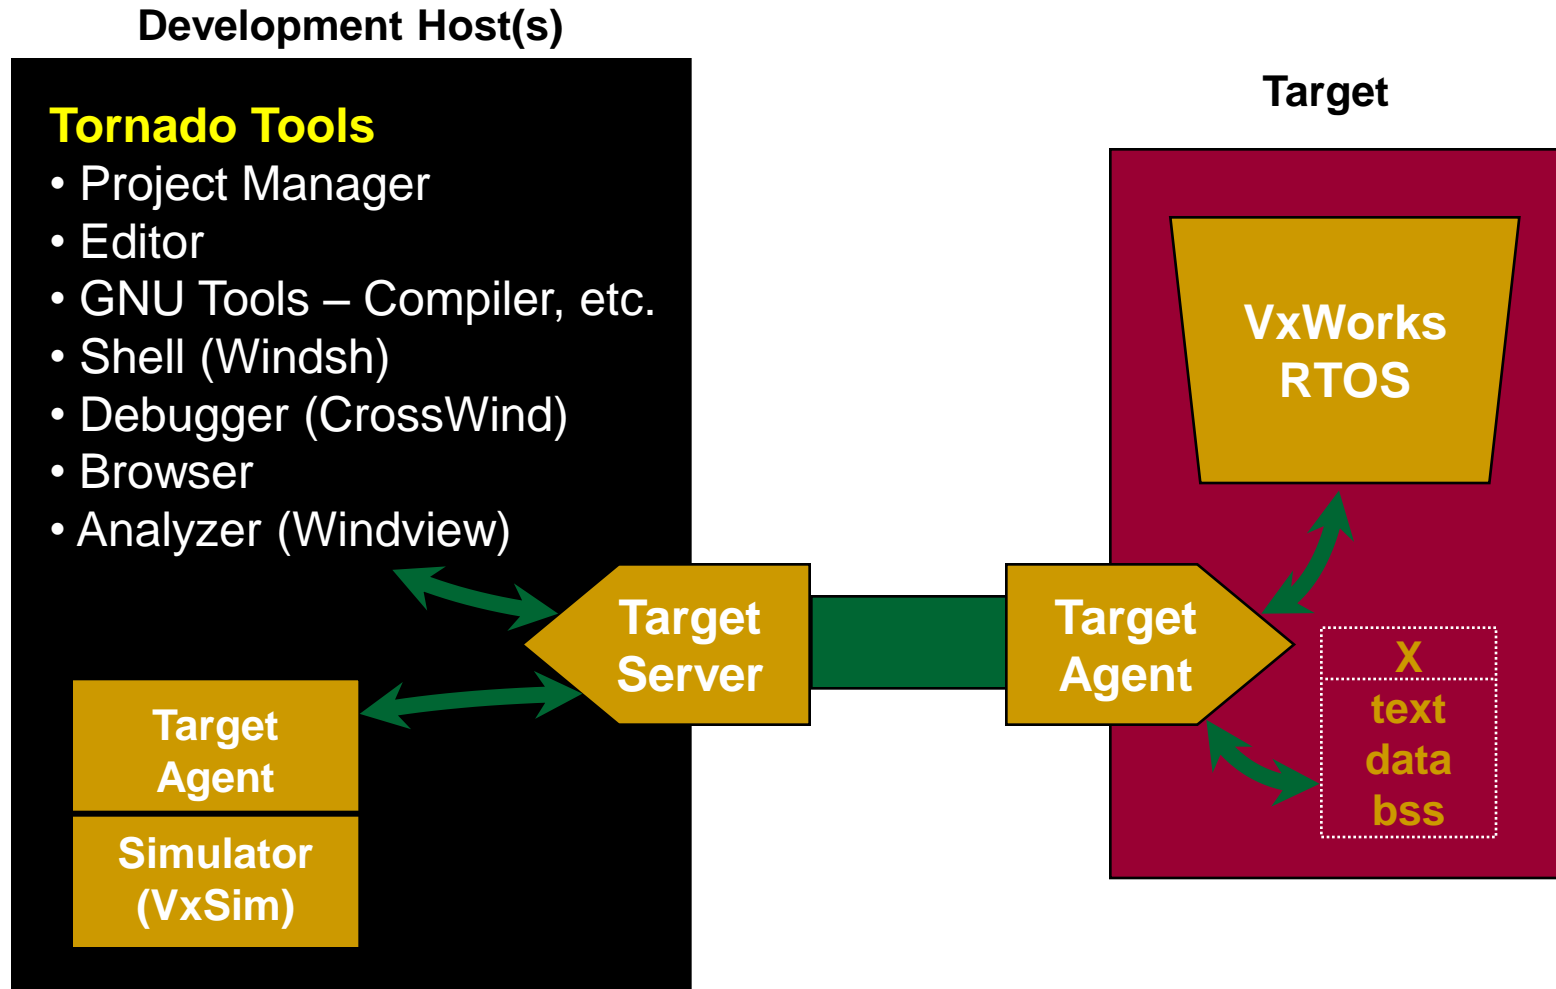
■ Assumptions

- ❑ Fixed priorities
- ❑ Set of resources to be accessed is fixed in advance

■ Advantages

- ❑ No deadlocks
 - ❑ No transitive blocking
 - ❑ A higher priority task can be blocked at most once by a lower priority task
-

Tornado Development Tools



Launching a Target Server

Configure Target Servers [?] [X]

Target Server Descriptions

Target Server Name	Description
igor	igor

[New] [Copy] [Remove]

Change Property: Back End

Available Back Ends	Timeout (sec)	Re-try (Count)
netrom		
wdbrpc		
wdbserial		


Target IP Name/Address: t12-168

Log File: [] [...]

Command Line

tgtsvr.exe t12-168 -n igor -V

[OK] [Launch] [Cancel] [Help]





D:\Tornado\host\ix86-win32\bin\tgtsvr.exe [] [] [X]

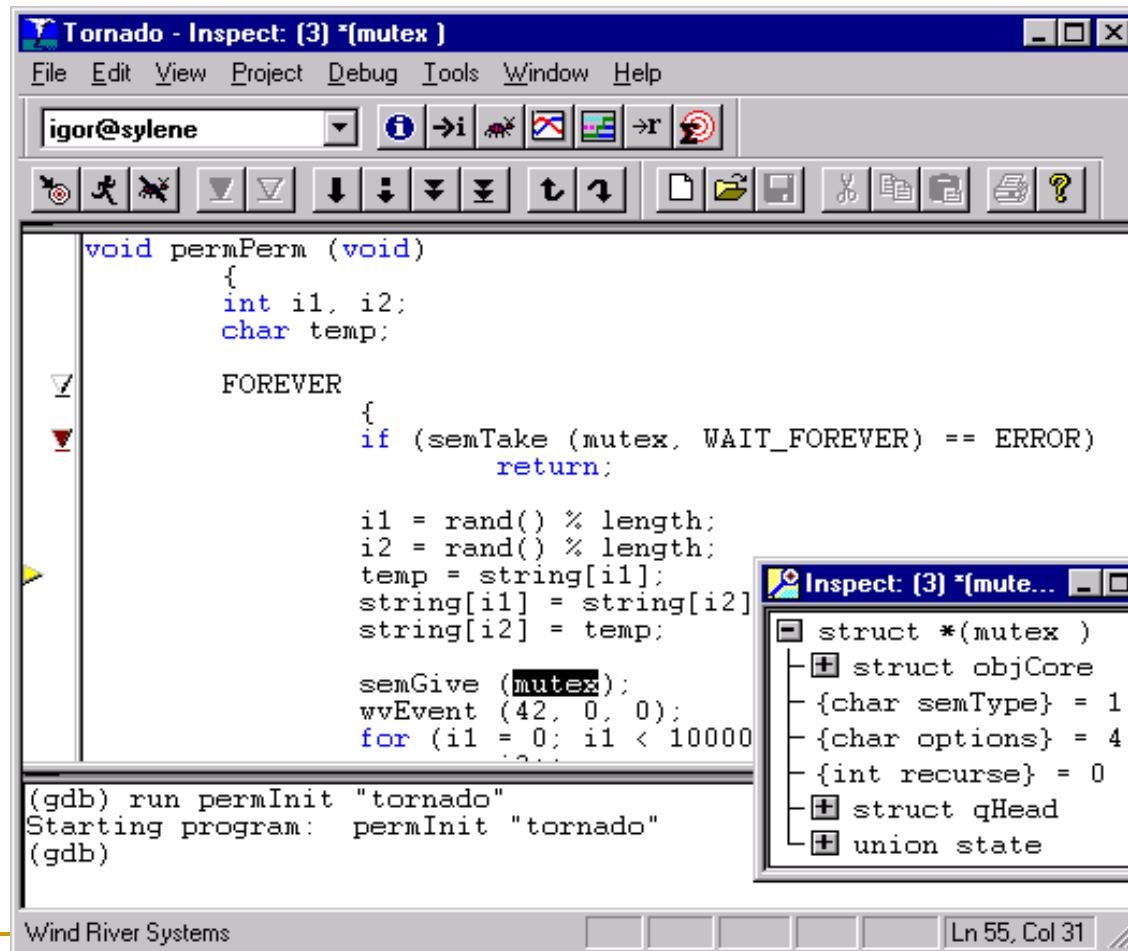
```
tgtsvr.exe (igor@sylene): Thu Jan 15 17:18:02 1998
License request... authorized on Local Host.
Wind River Systems Target Server: NT/Win95 version
Attaching backend... succeeded.
Connecting to target agent... succeeded.
Attaching C++ interface... succeeded.
Attaching a.out OMF reader... succeeded.
```

Starting Tornado Tools

1. Select target server from registry drop-down list.
2. Click on tool icon in launch toolbar.
3. Select de-"bug".

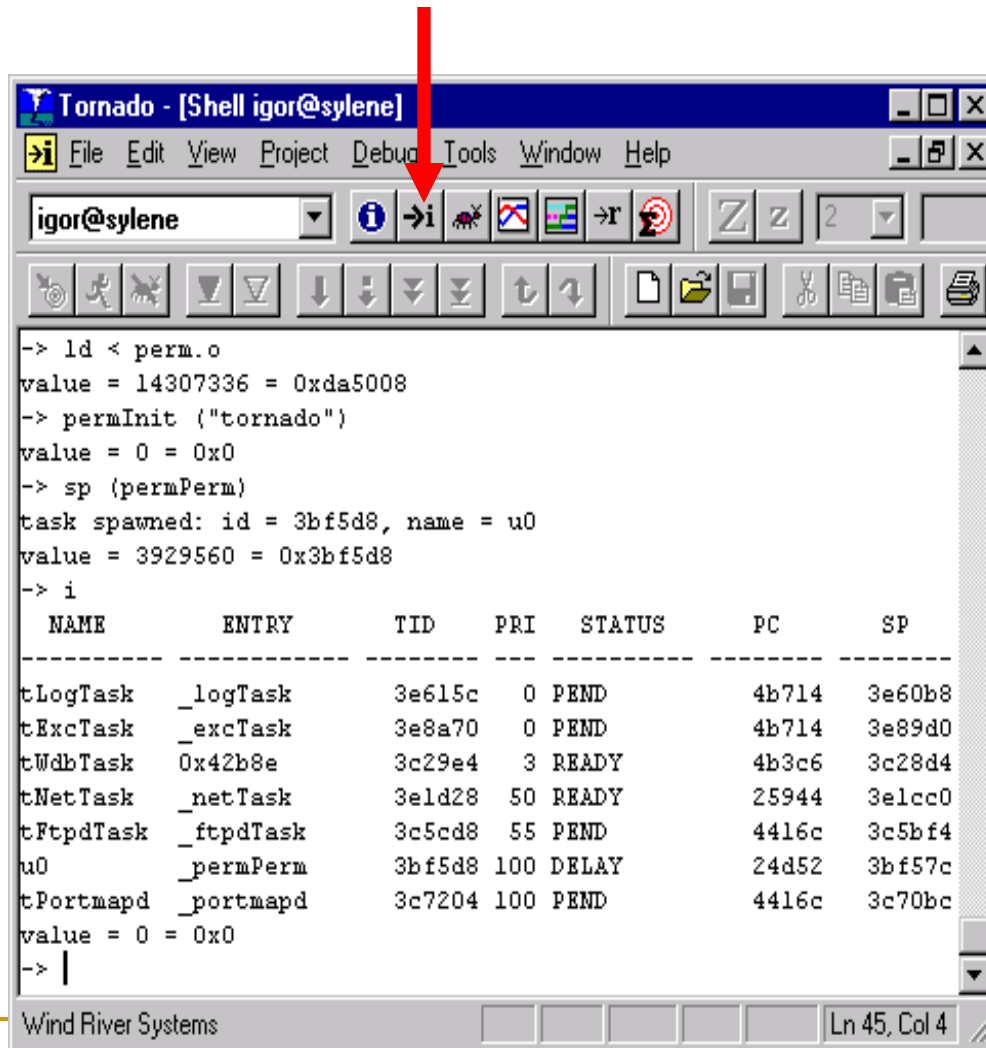


CrossWind Debugger



A graphical, source-level debugger built on the GNU debugger GDB. Provides task level and system level debugging. Use either the graphical or the command line interface.

WindSh - The Tornado Shell



The screenshot shows the Tornado Shell window titled "Tornado - [Shell igor@sylene]". The menu bar includes File, Edit, View, Project, Debug, Tools, Window, and Help. The toolbar contains various icons, including a red arrow pointing to the 'i' icon. The command prompt shows the following commands and output:

```
-> ld < perm.o
value = 14307336 = 0xda5008
-> permInit ("tornado")
value = 0 = 0x0
-> sp (permPerm)
task spawned: id = 3bf5d8, name = u0
value = 3929560 = 0x3bf5d8
-> i
```

NAME	ENTRY	TID	PRI	STATUS	PC	SP
tLogTask	_logTask	3e615c	0	PEND	4b714	3e60b8
tExcTask	_excTask	3e8a70	0	PEND	4b714	3e89d0
tWdbTask	0x42b8e	3c29e4	3	READY	4b3c6	3c28d4
tNetTask	_netTask	3e1d28	50	READY	25944	3e1cc0
tFtpdTask	_ftpdTask	3c5cd8	55	PEND	4416c	3c5bf4
u0	_permPerm	3bf5d8	100	DELAY	24d52	3bf57c
tPortmapd	_portmapd	3c7204	100	PEND	4416c	3c70bc

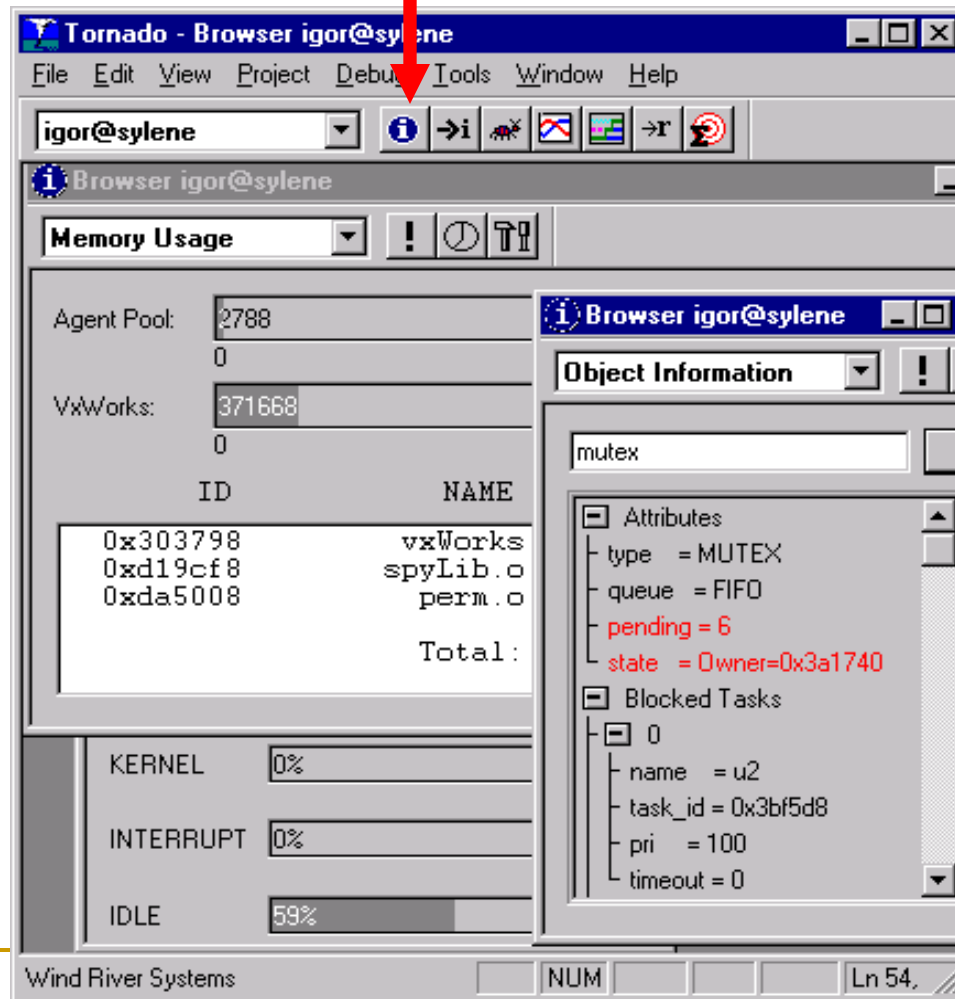
value = 0 = 0x0
-> |

Wind River Systems



A C expression interpreter and an associated Tcl interpreter.

Download code to the target; spawn tasks to execute functions; modify existing variables or create new ones.

The Browser



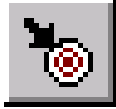

Graphical tool which lets you monitor the state of the target, and display information on particular VxWorks system objects.

Information displayed may be updated on demand  or periodically .


Downloading Object Modules

- Object modules may be downloaded dynamically and linked with modules already present on the target.
 - To download from the shell:
 - `ld < myProg.o`
 - **Debug => Download** also loads debugging information used by CrossWind.
-

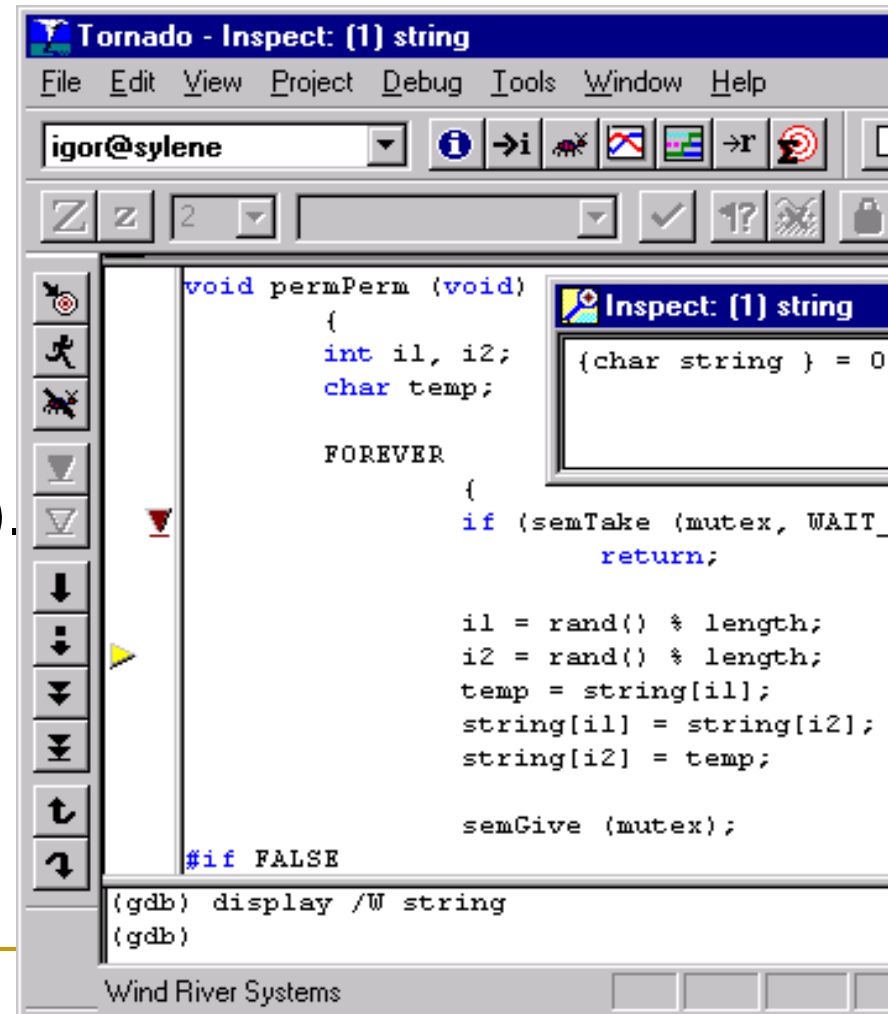
Debugging Tasks

- To download additional modules:
Debug => Download or  Toolbar button
- To create a new task to run a loaded function:
Debug => Run or  Toolbar button
- To debug an already running task:
Debug => Attach
- To debug multiple tasks independently, you may start multiple Tornado sessions.

CrossWind Debugger GUI

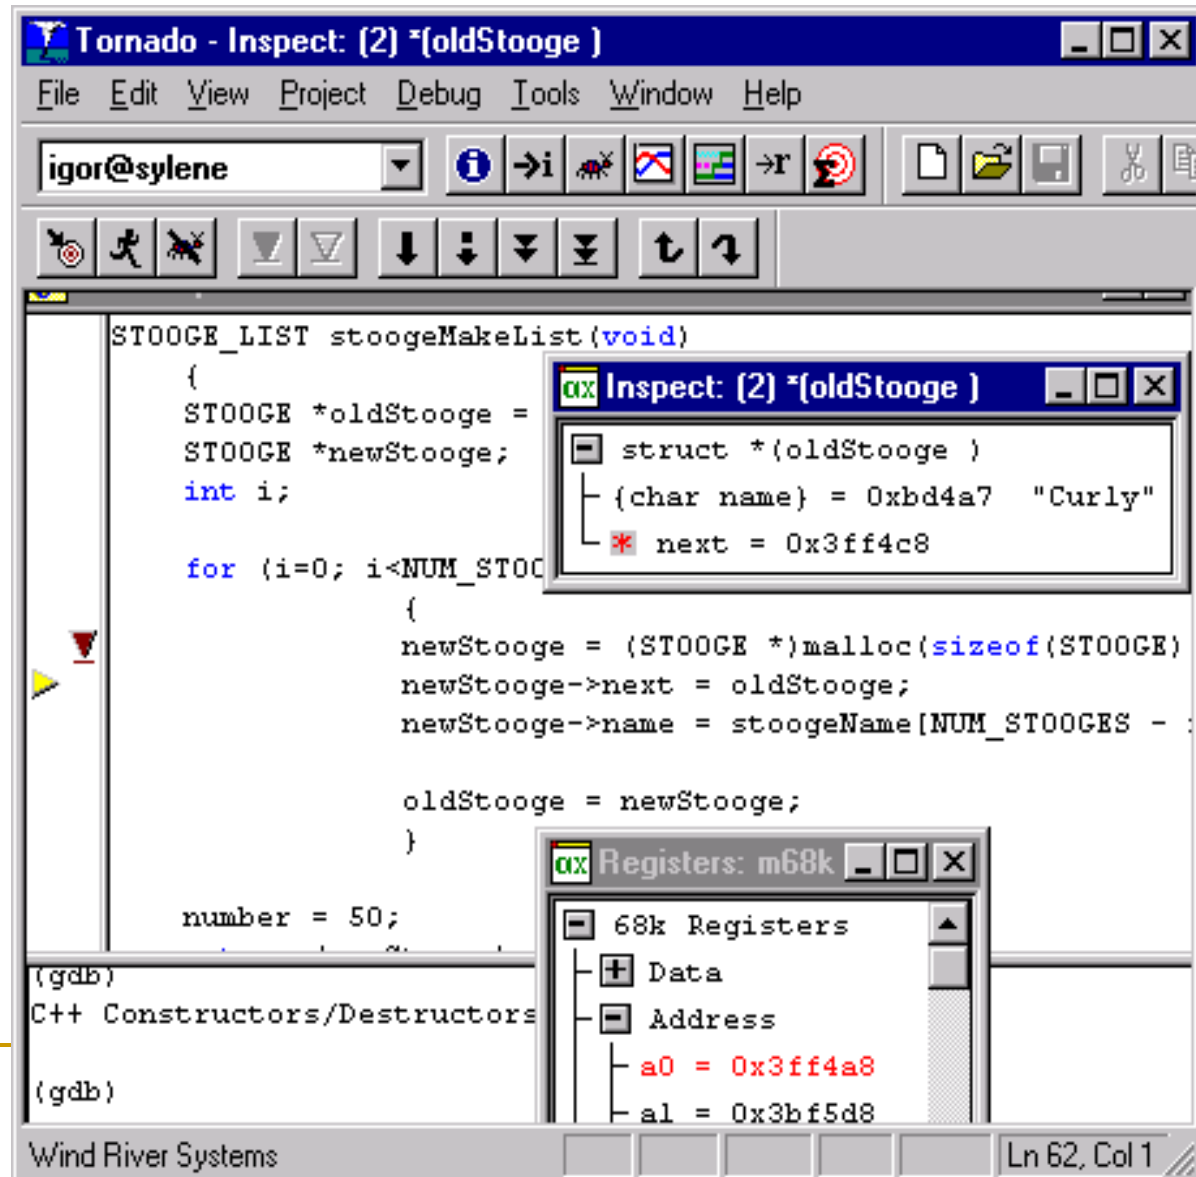


- Download object module.
- Spawn task to run function.
- Stop debugging.
- Toggle breakpoint.
- Temporary breakpoint.
- Step.
- Next (step over function call).
- Continue.
- Finish current subroutine.
- Up one stack frame.
- Down one stack frame.



CrossWind Debugger Displays

- Application expressions
- Current local variables
- CPU registers
- Memory dump
- Stack call chain
- Displays are updated when control returns to the debugger.




Browser Information

- Target Information
- Memory Usage (tools / application)
- Module Information (segments & symbols)
- Object Information (particular system objects)
- Spy Chart (CPU utilization)
- Stack Check
- Tasks



Update Now



Update Periodically




Configure Browser

WindView (Post-Processing Analysis)

WindView instruments the VxWorks kernel to record information on system (or user) events as they occur. If a high resolution timer ($\approx 1\mu\text{s}$) is available, events are assigned a timestamp.

This WindView graph illustrates a deadlock between *tTaskHi* and *tTaskLow*, and also shows task *u2* exiting (with help from the task *tExcTask*).



Summary

- Read Ch. 8.
- Read Sha's paper on priority inheritance protocols.