

# CIS 721 - Real-Time Systems

## Lecture 14: Mixing Real-Time and Non-Real-Time in Priority Driven Systems

---

Mitch Neilsen  
**neilsen@ksu.edu**

# Outline

- See Chapter 7 of Liu's text
- We discussed mixing real-time and non-real-time (aperiodic) jobs in static cyclic schedules
- We now address the same issue in priority-driven systems.
- First, we consider two straightforward scheduling algorithms for periodic and aperiodic jobs.
- Then, we look at a class of algorithms called **bandwidth-preserving algorithms** to schedule aperiodic jobs in a real-time system.

---

# Periodic and Aperiodic Tasks


- **Periodic task:**  $T_i$  is specified by  $(\phi_i, p_i, e_i, D_i)$
  - **Aperiodic tasks:** non-real-time
    - Released at arbitrary times.
    - Have no deadline and  $e_i$  may be unspecified.
  - We assume that periodic and aperiodic tasks are independent of each other.
-

# Correct and Optimal Schedules

- A **correct schedule** never results in a deadline being missed by periodic tasks.
- A **correct scheduling algorithm** only produces correct schedules.
- An **optimal aperiodic job scheduling algorithm** minimizes either
  - the response time of the aperiodic job at the head of the queue, or
  - the average response time of all aperiodic jobs.

# Scheduling Mixed Jobs

- We assume there are separate job queues for real-time (periodic) and non-real-time (aperiodic) jobs.
- How do we minimize response time for aperiodic jobs without impacting periodic jobs?



# Background Scheduling (BS)


- Periodic jobs are scheduled using any priority-driven scheduling algorithm.
- Aperiodic jobs are scheduled and executed in the background:
  - Aperiodic jobs are executed only when there is no periodic job ready to execute.
  - Simple to implement and always produces correct schedules.  
**The lowest priority task simply executes jobs from the aperiodic job queue.**
  - We can improve response times without jeopardizing deadlines by using a **slack stealing algorithm** to delay the execution of periodic jobs as long as possible.
    - This is the same thing we did with cyclic executives.
    - However, it is very expensive (in terms of overhead) to implement slack-stealing in priority-driven systems.

# Simple Periodic (Polling) Server

- Periodic jobs are scheduled using any priority-driven scheduling algorithm.
- Aperiodic jobs are executed by a special periodic server:
  - The periodic server is a periodic task  $T_p = (p_s, e_s)$ .
    - $e_s$  is called the execution budget of the server.
    - The ratio  $u_s = e_s/p_s$  is the size of the server.
  - Suspends as soon as the aperiodic queue is empty or  $T_p$  has executed for  $e_s$  time units (whichever comes first).
    - This is called exhausting its execution budget.
  - Once suspended, it cannot execute again until the start of the next period.
    - That is, the execution budget is replenished (reset to  $e_s$  time units) at the start of each period.
    - Thus, the start of each period is called the replenishment time for the simple periodic server.

# Periodic Server with RM Scheduling


- **Example Schedule:** Two tasks,  $T_1 = (3,1)$ ,  $T_2 = (10,4)$ , and a periodic server  $T_p = (2.5,0.5)$ . Assume an aperiodic job  $J_a$  arrives at  $t = 0.1$  with an execution time of  $e_a = 0.8$ .
  - The periodic server cannot execute the job that arrives at time 0.1 since it was suspended at time 0 because the aperiodic job queue was empty.






# Example (cont.)

- The Periodic Server executes its job  $J_a$  starting at 2.5, up until it exhausts its budget at time 3.



# Example (cont.)

- It finishes executing in the next period from time 5.0 to 5.3.
- So the response time of Ja is 5.2; that is, from 0.1 to 5.3.



# Improving the Periodic Server

- The problem with the periodic server is that it exhausts its execution budget whenever the aperiodic job queue is empty.
  - If an aperiodic job arrives  $\varepsilon$  time units after the start of the period, it must wait until the start of the next period ( $p_s - \varepsilon$  time units) before it can begin execution.
- We would like to preserve the execution budget of the polling server and use it later in the period to shorten the response time of aperiodic jobs:
  - **Bandwidth-Preserving Servers** do just this!

# Bandwidth-Preserving Servers

- More terminology:
  - The periodic server is **backlogged** whenever the aperiodic job queue is nonempty or the server is executing a job.
  - The server is **idle** whenever it is not backlogged.
  - The server is **eligible** for execution when it is backlogged and has an execution budget (greater than zero).
  - When the server executes, it consumes its execution budget at the rate of one time unit per unit of execution.
  - Depending on the type of periodic server, it may also consume all or a portion of its execution budget when it is idle: the simple periodic server consumed all of its execution budget when the server was idle.

# Bandwidth-Preserving Servers

- **Bandwidth-preserving servers** differ in their replenishment times and how they preserve their execution budget when idle.
- We assume the scheduler tracks the consumption of the server's execution budget and suspends the server when the budget is exhausted or the server becomes idle.
- The scheduler replenishes the servers execution budget at the appropriate replenishment times, as specified by the type of bandwidth-preserving periodic server.
- The server is only eligible for execution when it is backlogged and its execution budget is non-zero.

# Four Bandwidth-Preserving Servers

- **Deferrable Servers (1987)**
  - Oldest and simplest of the bandwidth-preserving servers.
  - Static-priority algorithms by Lehoczky, Sha, and Strosnider.
  - Deadline-driven algorithms by Ghazalie and Baker (1995).
- **Sporadic Servers (1989)**
  - Static-priority algorithms by Sprunt, Sha, and Lehoczky.
  - Deadline-driven algorithms by Ghazalie and Baker (1995).
- **Total Bandwidth Servers (1994, 1995)**
  - Deadline-driven algorithms by Spuri and Buttazzo.
- **Constant Utilization Servers (1997)**
  - Deadline-driven algorithms by Deng, Liu, and Sun.


# Deferrable Server (DS)

- Let the task  $T_{DS} = (p_s, e_s)$  denote a **deferrable server**.
- Consumption Rule:
  - The execution budget is consumed at the rate of one time unit per unit of execution.
- Replenishment Rule:
  - The execution budget is set to  $e_s$  at time instants  $k \cdot p_s$ , for  $k \geq 0$ .
  - Note: Unused execution budget cannot be carried over to the next period.
- The scheduler treats the deferrable server as a periodic task that may suspend itself during execution (i.e., when the aperiodic queue is empty).

# DS with RM Scheduling Example

**Example Schedule:** Same two tasks,  $T_1 = (3,1)$ ,  $T_2 = (10,4)$ , and deferrable server  $T_{DS} = (2.5,0.5)$ . Assume an aperiodic job  $J_a$  arrives at time  $t = 0.1$  with an execution time of  $e_a = 0.8$  (again).

The DS can execute the job that arrives at time 0.1 since it preserved its budget when the aperiodic job queue was empty.






# DS with RM Scheduling Example (cont.)

**Example Schedule:** Same two tasks,  $T_1 = (3,1)$ ,  $T_2 = (10,4)$ , and deferrable server  $T_{DS} = (2.5,0.5)$ . Assume an aperiodic job  $J_a$  arrives at time  $t = 0.1$  with an execution time of  $e_a = 0.8$  (again).


The response time of the aperiodic job  $J_a$  is 2.7  
It was 5.2 with the simple periodic server.



# Another Example

**Another Example:** Two tasks,  $T_1 = (2, 3.5, 1.5)$ ,  $T_2 = (6.5, 0.5)$ , and a deferrable server  $T_{DS} = (3, 1)$ . Assume an aperiodic job  $J_a$  arrives at time  $t = 2.8$  with and execution time of  $e_a = 1.7$ .


The response time of the aperiodic job  $J_a$  is 3.7.



# DS with EDF Scheduling Example

**Same Task Set:** Two tasks,  $T_1 = (2, 3.5, 1.5)$ ,  $T_2 = (6.5, 0.5)$ , and a deferrable server  $T_{DS} = (3, 1)$ . Assume an aperiodic job  $J_a$  arrives at time  $t = 2.8$  with an execution time of  $e_a = 1.7$ .


The response time of the aperiodic job  $J_a$  is still 3.7.



# DS with EDF vs Background Scheduling

**Same Task Set:** Two tasks,  $T_1 = (2, 3.5, 1.5)$ ,  $T_2 = (6.5, 0.5)$ , and  $T_{DS} = (3, 1)$  with background scheduling. Assume an aperiodic job  $J_a$  arrives at time  $t = 2.8$  with an execution time of  $e_a = 1.7$ .

The DS exhausts its budget at time 4.7...




# DS with EDF vs Background Scheduling

[ recall, Background Scheduling = schedule when no periodic tasks are ready ]

**Same Task Set:** Two tasks,  $T_1 = (2, 3.5, 1.5)$ ,  $T_2 = (6.5, 0.5)$ , and  $T_{DS} = (3, 1)$  with background scheduling. Assume an aperiodic job  $J_a$  arrives at time  $t = 2.8$  with an execution time of  $e_a = 1.7$ .

However, using background scheduling, the response time of the aperiodic job  $J_a$  is reduced to 2.4.




# DS with Background Scheduling

- We can also combine background scheduling of the deferrable server with RM.
  - For the deferrable server example task set, the response time doesn't change. Why?
- Why complicate things by adding background scheduling of the deferrable server?
- Why not just give the deferrable scheduler a larger execution budget? See the next slide!

# DS with RM Scheduling Revisited

**Modified Example:** Same two tasks,  $T_1 = (2, 3.5, 1.5)$ ,  $T_2 = (6.5, 0.5)$ , and deferrable server  $T_{DS} = (3, 1)$ . Assume an aperiodic job  $J_a$  arrives at time  $t_0 = 65$  with an execution time of  $e_a = 3$ .

A larger execution budget for  $T_{DS}$  would result in  $T_1$  missing a deadline.  
Time  $t_0 = 65$  is a critical instant for this task set.



# Schedulability and DS

- There are no known necessary and sufficient schedulability conditions for task sets that contain a DS with arbitrary priority. We will see why shortly.
- However, we can extend TDA and Generalized TDA to yield necessary and sufficient schedulability tests when the DS is the highest priority task in a periodic (real world sporadic) task set.
- We start with a critical instant lemma for systems with a DS.



# Critical Instants in a Fixed Priority System with a Deferrable Server (DS)

**Lemma 7-1:** [Lehoczky, Sha, and Strosnider] In a fixed-priority system in which the relative deadline of every independent, preemptable periodic task is no greater than its period and there is a deferrable server  $(p_s, e_s)$  with the highest priority among all tasks, a critical instant of every periodic task  $T_i$  occurs at time  $t_0$  when all of the following are true.

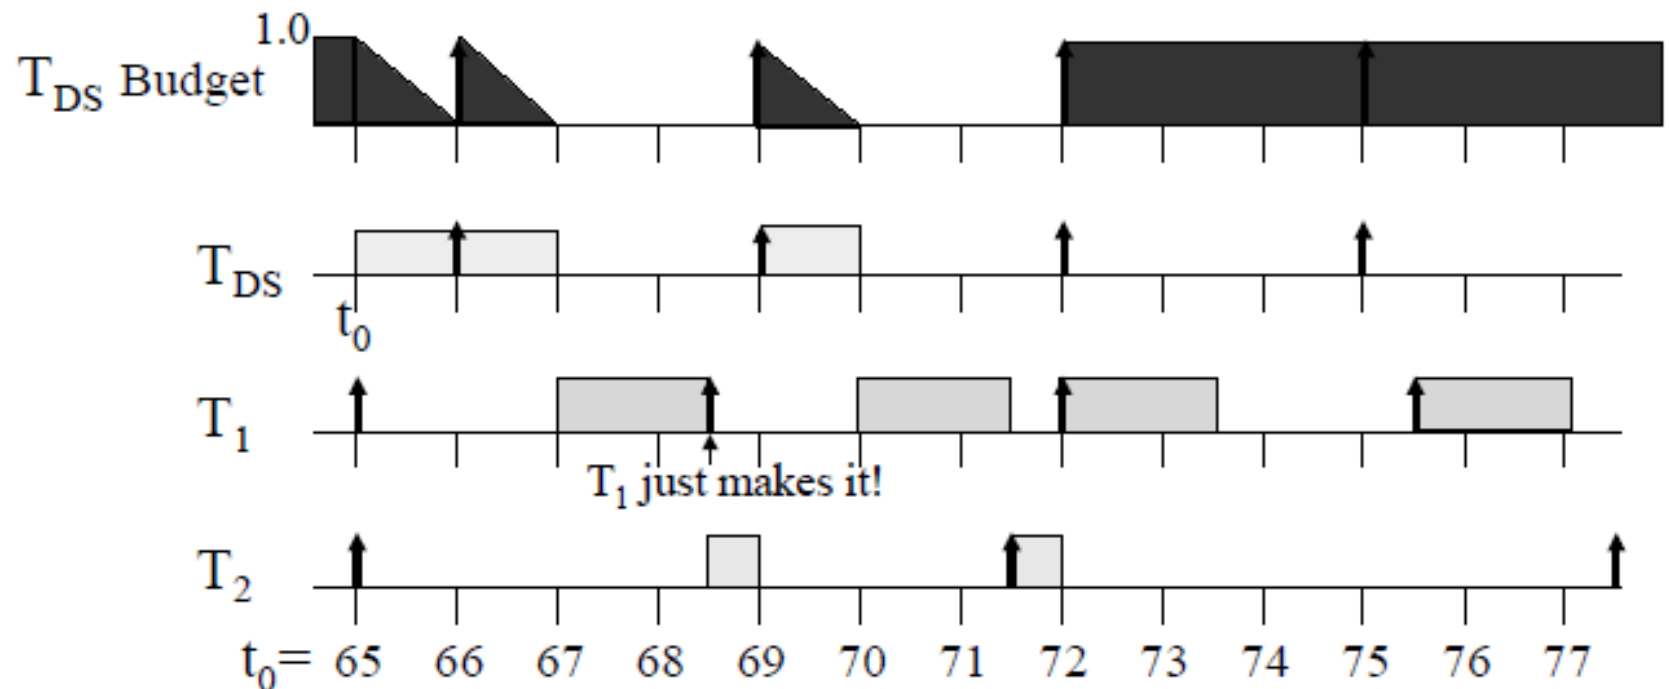
1. One of its jobs  $J_{i,c}$  is released at  $t_0$ .
2. A job in every higher-priority task is release at the same time.
3. The budget of the server is  $e_s$  at  $t_0$ , one or more aperiodic jobs are released at  $t_0$ , and they keep the server backlogged hereafter.
4. The next replenishment time of the server is  $t_0 + e_s$ .

# Observations on Lemma 7.1

- The Proof of Lemma 7.1 is a straightforward extension of the proof we gave for Theorem 6.5. Convince yourself of this!
- Note: We are not saying that  $T_{DS}, T_1, \dots, T_i$  will all necessarily release jobs at the same time, but if this does happen, we are claiming that the time of release will be a critical instant for  $T_i$ .
- We can use the critical instant  $t_0$ , defined by Lemma 7.1, to derive necessary and sufficient conditions for the schedulability of a task set when the DS has highest priority.
- First, let's take a look at a processor demand anomaly created by the bandwidth preserving DS.

# Observations on Lemma 7.1 (cont.)

All four conditions of Lemma 7.1 hold in the last example:



Notice that the processor demand created by the DS in an interval from  $[65, 68.5]$  is twice what it would be if it were an ordinary periodic task! This is because we preserve the bandwidth of the DS.

# Response Time Analysis with a DS

Observation: Cases (1) and (2) of Lemma 7.1 define a critical instant for any fixed-priority task set. When cases (3) and (4) of Lemma 7.1 are true, the processor demand created by the DS in an interval of length  $t$  can be at most

$$e_s + \left\lceil \frac{t - e_s}{p_s} \right\rceil e_s \quad (*)$$

Thus, TDA and Generalized TDA with blocking terms can be extended to systems with a DS that executes at the highest priority. The TDA function becomes:

$$w_i(t) = e_i + b_i + e_s + \left\lceil \frac{t - e_s}{p_s} \right\rceil e_s + \sum_{k=1}^{i-1} \left\lceil \frac{t}{p_k} \right\rceil \cdot e_k \quad \text{for } 0 < t \leq \min(D_i, p_i)$$

# DS with Highest Fixed Priority

- When the DS is the highest priority process in a fixed-priority system:
  - It may be able to execute an extra  $e_s$  time units more than a normal periodic task in the feasible interval of task  $T_i$ , as expressed in Equation (\*) and the modified  $w_i(t)$ .
- Thus, Response Time Analysis, using the modified  $w_i(t)$  provides a necessary and sufficient condition for fixed-priority systems with one DS executing at the **highest** priority.

# DS with Arbitrary Fixed Priority

- When the DS is not the highest priority process:
  - It may not be able to execute the extra  $e_s$  time units expressed in Equation (\*) and the modified  $w_i(t)$ .
  - However, the time-demand function of a task  $T_i$  with lower priority than an arbitrary-priority DS is bounded from above by the modified  $w_i(t)$ .
- Thus, Response Time Analysis provides a sufficient (but not necessary) condition for fixed-priority systems with one **arbitrary**-priority DS.

# Multiple Arbitrary Fixed-Priority DS

We may want to differentiate aperiodic jobs by executing them at different priorities. To do this, we use multiple DS with different priorities and task parameters  $(p_{s,k}, e_{s,k})$ .

The TDA and Generalized TDA with blocking terms can be further extended to these systems. Specifically, the time demand function  $w_i(t)$  of a periodic task  $T_i$  with a lower priority than  $m$  DS becomes:

$$w_i(t) = e_i + b_i + \sum_{k=1}^m \left( 1 + \left\lceil \frac{t - e_{s,k}}{p_{s,k}} \right\rceil \right) e_{s,k} + \sum_{k=1}^{i-1} \left\lceil \frac{t}{p_k} \right\rceil \cdot e_k \quad \text{for } 0 < t \leq \min(D_i, p_i)$$

# Schedulable Utilization with Fixed-Priority DS

- We now look at utilization-based scheduling tests for fixed-priority systems with one DS.
- There are no known necessary and sufficient schedulable utilization conditions for fixedpriority systems with a DS.
- However, there does exist a sufficient condition for RM when the DS has the shortest period plus some other conditions...



# RM Schedulable Utilization with a DS

**Theorem 7.2** Consider a system of  $n$  independent, preemptable periodic tasks whose periods satisfy the inequalities  $p_s < p_1 < p_2 < \dots < p_n < 2p_s$  and  $p_n > p_s + e_s$  and whose relative deadlines are equal to their respective periods. This system is schedulable rate monotonically with a deferrable server  $(p_s, e_s)$  if their total utilization is less than or equal to

$$U_{\text{RM/DS}}(n) = (n-1) \left[ \left( \frac{u_s + 2}{u_s + 1} \right)^{1/(n-1)} - 1 \right]$$

where  $u_s$  is the utilization  $e_s/p_s$  of the server.

**Proof:** Similar to Thm 6.11 and left as an exercise!

Note that this is only a **sufficient** schedulability test.

# RM Schedulable Utilization with a DS and Arbitrary Periods

Observe: If  $p_i < p_s$ , then task  $T_i$  is unaffected by the DS. If  $p_i > p_s$ , then it may be blocked an extra  $e_s$  time units in a feasible interval.

**Theorem**: Consider a system of  $n$  independent, preemptable periodic tasks whose relative deadlines are equal to their respective periods. Task  $T_i$  with  $p_i > p_s$  is schedulable rate monotonically with a deferrable server  $(p_s, e_s)$  if

$$U_i + u_s + \frac{e_s + b_i}{p_i} \leq U_{\text{RM}}(i+1)$$

where  $u_s$  is the utilization  $e_s/p_s$  of the server,  $U_i$  is the total utilization of the tasks  $T_1 \dots T_i$ , and  $b_i$  is the blocking time encountered by task  $T_i$  from lower priority tasks.

# Schedulability of a Deadline-Driven System with a DS

- In fixed-priority systems, the DS behaves like a periodic task  $(p_s, e_s)$  except that it could execute an extra amount of time (at most  $e_s$  time units) in the feasible interval of any lower priority job.
- In a deadline-driven system, the DS can execute at most  $e_s$  time units in the feasible interval of any job (under certain conditions).
- We present a sufficient (but not necessary) schedulability condition for the EDF algorithm.
- First, a bound on the processor demand created by a DS.

# Bounding the Demand of a DS in an EDF System

- An interval  $(a, b]$  is **post-idle** if either  $a = 0$ , or if no job with a deadline in the interval  $(a+1, b]$  executes in the interval  $(a-1, a]$ .
  - The implication of this definition is that all jobs with deadlines in  $(a+1, b]$  are “idle” during the interval  $(a-1, a]$  in the sense that all jobs released before time  $a$  with deadlines in  $(a+1, b]$  have completed execution before time  $a$  (*i.e.*, either the processor is idle in  $(a-1, a]$  or a job with deadline at time  $a$  or after time  $b$  executes in  $(a-1, a]$ ).
- The following lemma gives us a simple upper bound for the processor demand in a post-idle interval of length  $L$ .

# Maximum Demand of a DS

**Lemma:** The maximum demand  $w_{DS}(L)$  of a  $DS = (p_s, e_s)$  during a post-idle interval of length  $L$  in an EDF scheduled system of  $n$  independent, preemptable periodic tasks is bounded such that

$$w_{DS}(L) \leq u_s (L + p_s - e_s)$$


where  $u_s$  is the utilization  $e_s/p_s$  of the server.

**Proof:** Let  $(t_{-1}, t]$  be a post-idle interval. The maximum demand for DS occurs when

1. At time  $t_{-1}$ , its budget is equal to  $e_s$  and the server's deadline (and budget replenished time) is  $t_{-1} + e_s$ .
2. One or more aperiodic jobs arrive at  $t_{-1}$  and the DS is backlogged until at least time  $t$ .
3. The server's deadline  $t_{-1} + e_s$  is earlier than the deadlines of all the periodic jobs that are ready for execution in the interval  $(t_{-1}, t_{-1} + e_s]$ .

## Proof (cont.)

Maximum demand created by DS in a post-idle interval under EDF:



Observe that under these conditions, the maximum demand created by DS in the post-idle interval  $(t_{-1}, t]$  is at most

$$e_s + \left\lfloor \frac{t - (t_{-1} + e_s)}{p_s} \right\rfloor e_s = e_s + \left\lfloor \frac{t - t_{-1} - e_s}{p_s} \right\rfloor e_s$$

## Proof (cont.)

$$\begin{aligned}\text{Thus, } w_{\text{DS}}(t - t_{-1}) &\leq e_s + \left\lfloor \frac{t - t_{-1} - e_s}{p_s} \right\rfloor e_s \leq e_s + \frac{t - t_{-1} - e_s}{p_s} e_s \\ &= \frac{p_s}{p_s} e_s + \frac{t - t_{-1} - e_s}{p_s} e_s = p_s u_s + (t - t_{-1} - e_s) u_s \\ &= u_s (p_s + (t - t_{-1} - e_s)) \\ &= u_s (t - t_{-1} + p_s - e_s)\end{aligned}$$

Since  $(t_{-1}, t]$  is a post-idle interval of length  $L = t - t_{-1}$ ,

$$w_{\text{DS}}(L) \leq u_s (L + p_s - e_s).$$

# Schedulability with a DS

- Combining this result with Theorem 6.5, we obtain a Theorem 7.3 by Ghazalie and Baker:

**Theorem 7.3:** A periodic task  $T_i$  in a system of  $n$  independent, preemptable periodic tasks is schedulable with a DS=  $(p_s, e_s)$  according to the EDF algorithm if

$$\sum_{k=1}^n \frac{e_k}{\min(D_k, p_k)} + u_s \left( 1 + \frac{p_s - e_s}{D_i} \right) \leq 1 \quad (7.5)$$

where  $u_s$  is the utilization  $e_s/p_s$  of the server.



# Proof of Theorem 7.3

Suppose Equation (7.5) holds for task  $T_i$  but a deadline is missed. Let  $t_d$  be the earliest point in time at which a deadline is missed and  $t_{-1}$  be the start of last post-idle interval that includes time  $t_d$ . Thus, a deadline is missed in the post-idle interval  $(t_{-1}, t_d]$ .

From Theorem 6.2 and the previous lemma, the demand in this interval is at most

$$\sum_{k=1}^n \left\lfloor \frac{t_d - t_i}{\min(D_k, p_k)} \right\rfloor e_k + u_s (t_d - t_i + p_s - e_s)$$

Because a deadline is missed at  $t_d$ , demand over  $(t_{-1}, t_d]$  exceeds  $t_d - t_{-1}$ . Thus, we have

$$\begin{aligned} t_d - t_i &< \sum_{k=1}^n \left\lfloor \frac{t_d - t_i}{\min(D_k, p_k)} \right\rfloor e_k + u_s (t_d - t_i + p_s - e_s) \\ &\leq \sum_{k=1}^n \frac{t_d - t_i}{\min(D_k, p_k)} e_k + u_s (t_d - t_i + p_s - e_s) \end{aligned}$$

## Proof (cont.)

Dividing both sides by  $(t_d - t_{-1})$ , we get

$$\begin{aligned} 1 &< \sum_{k=1}^n \frac{e_k}{\min(D_k, p_k)} + \frac{u_s(t_d - t_i + p_s - e_s)}{t_d - t_i} \\ &= \sum_{k=1}^n \frac{e_k}{\min(D_k, p_k)} + u_s \left( 1 + \frac{p_s - e_s}{t_d - t_i} \right) \\ &\leq \sum_{k=1}^n \frac{e_k}{\min(D_k, p_k)} + u_s \left( 1 + \frac{p_s - e_s}{D_i} \right) \end{aligned}$$

Since  $D_i \leq (t_d - t_{-1})$ .

This contradicts our assumption that Equation (7.5) holds.

---

# Multiple DS

We may want to differentiate aperiodic jobs by executing them at different priorities. To do this in a deadline-driven system, we (again) use multiple DS with different priorities and task parameters  $(p_{s,k}, e_{s,k})$ .

**Corollary:** A periodic task  $T_i$  in a system of  $n$  independent, preemptable periodic tasks is schedulable with  $m$  a DS =  $(p_{s,k}, e_{s,k})$  according to the EDF algorithm if

$$\sum_{k=1}^n \frac{e_k}{\min(D_k, p_k)} + \sum_{k=1}^m u_{s,k} \left( 1 + \frac{p_{s,k} - e_{s,k}}{D_i} \right) \leq 1$$

where  $u_{s,k}$  is the utilization  $e_{s,k}/p_{s,k}$  of server  $k$ .

The proof is left as an exercise.

---

# DS Summary

- In both fixed-priority and deadline-driven systems, we see that the DS behaves like a periodic task with parameters  $(p_s, e_s)$  except it may execute an additional amount of time in the feasible interval of any lower priority job.
  - This is because, the bandwidth-preserving conditions result in a scheduling algorithm that is non-work-conserving with respect to a normal periodic task.
-

---

# Sporadic Servers

- Sporadic Servers (SS) were designed to overcome the blocking time a DS may impose on lower priority jobs.
  - All sporadic servers are bandwidth preserving, but the consumption and replenishment rules ensure that a SS, specified a  $T_s = (p_s, e_s)$  never creates more demand than a periodic (“real-world” sporadic) task with the same task parameters.
  - Thus, schedulability of a system with a SS is determined exactly as a system without a SS.
-