

CIS 721 - Real-Time Systems

Quiz #1 Review

Mitch Neilsen
neilsen@ksu.edu

Outline

- Quiz #1: Fri., Oct. 16, in class
- Topics
 - Real-Time Scheduling Theory and Algorithms
 - Ch. 1-7
- Open book, open notes

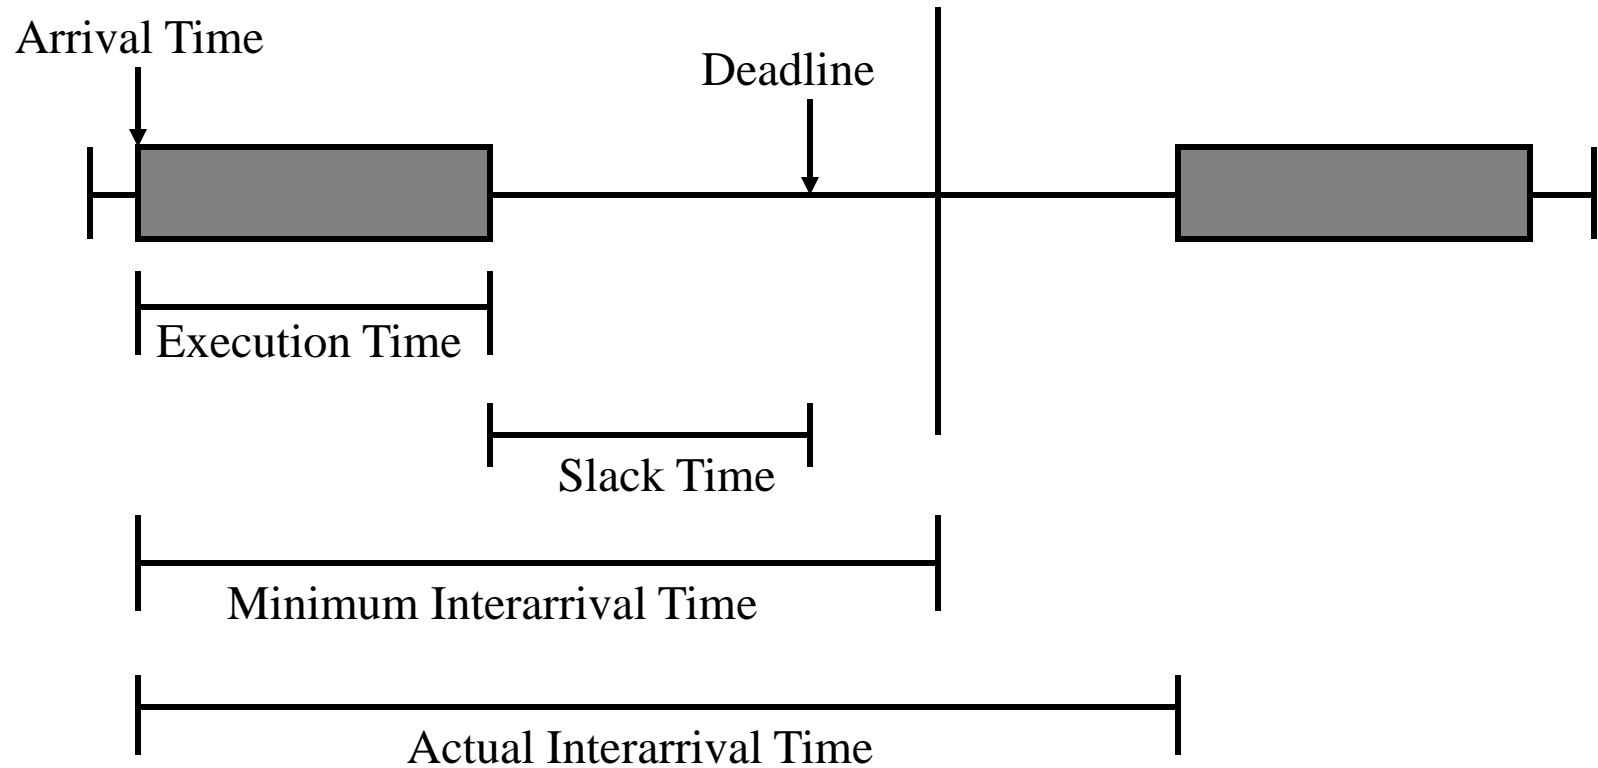
Terms and Concepts

- A **real-time system** is a system with performance deadlines on computations and actions; that is, *system **correctness** depends on the **timeliness** of the results.*
- An **embedded system** is a system that exists within a larger system.
- A **job** is a unit of work that is scheduled and executed by the system ($J_{i,k}$).
- A **task** is a set of related jobs that provide some system function $\tau_i = \{ J_{i,1}, J_{i,2}, \dots, J_{i,n} \}$; e.g., the reception of a data frame could be a job that is part of a task that provides time service.
- The **deadline** of a job is the time at which a job must be completed.

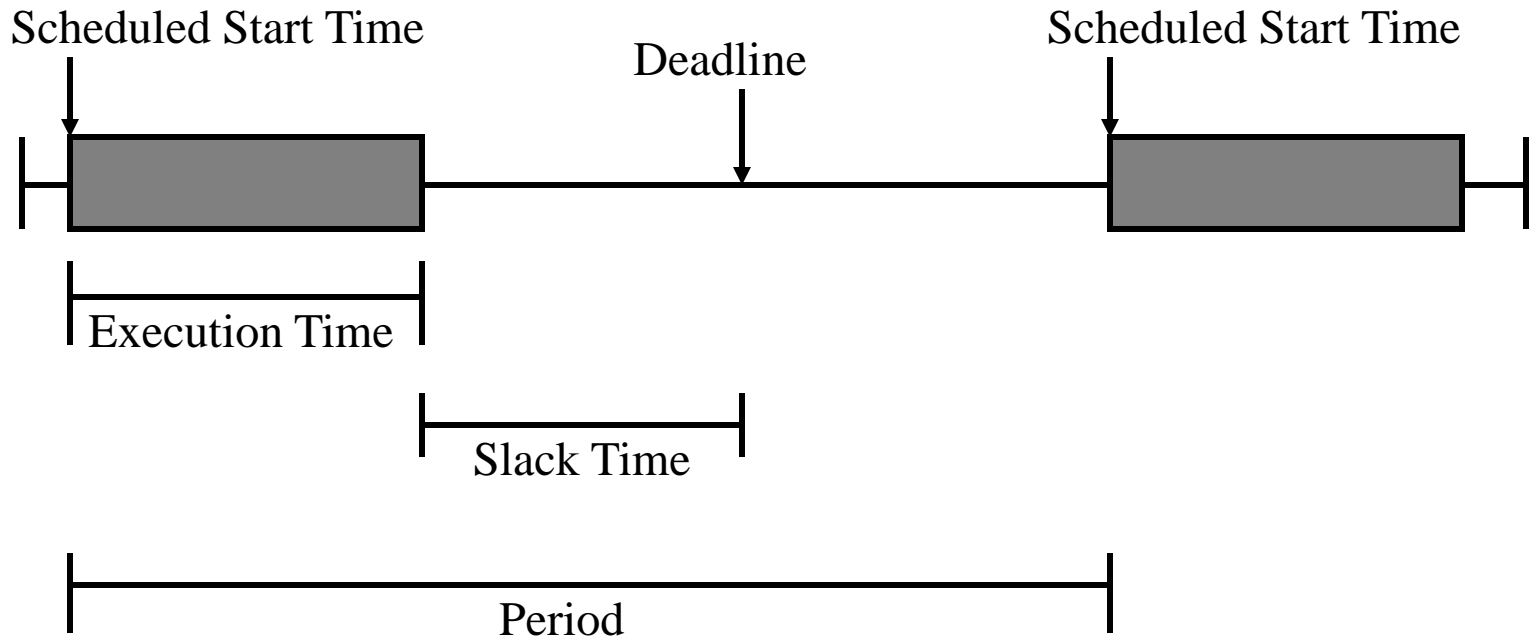
Deadlines

- The **release time** (or **arrival time**) of a job is the time at which the job becomes available for execution (r_i or R_i).
- The **response time** of a job is the length of time between the release time of the job and the time instant when it completes.
- The **relative deadline** of a job is the maximum allowable response time of a job (D_i).
- The **absolute deadline** of a job is the time at which a job must be completed ($d_i = r_i + D_i$).
- Failure to meet a **hard deadline** is considered a fatal fault, whereas a **soft deadline** can be missed as long as the average performance is optimized.

Event-Driven Task



Time-Driven Task



Scheduler

- A **scheduler** assigns jobs to processors.
- A **schedule** is an assignment of all jobs in the system on available processors (produced by scheduler).
- The **execution time** (or **run-time**) of a job is the amount of time required to complete the execution of a job once it has been scheduled (e_i or C_i).
- A constraint imposed on the timing behavior of a job is called a **timing constraint**.

Assumptions

- The scheduler works correctly; e.g., it only produces **valid schedules** that satisfy the following conditions:
 - each processor is assigned to at most one job at a time,
 - each job is assigned to at most one processor,
 - no job is scheduled before its release time, and
 - all precedence constraints and resource usage constraints are satisfied.

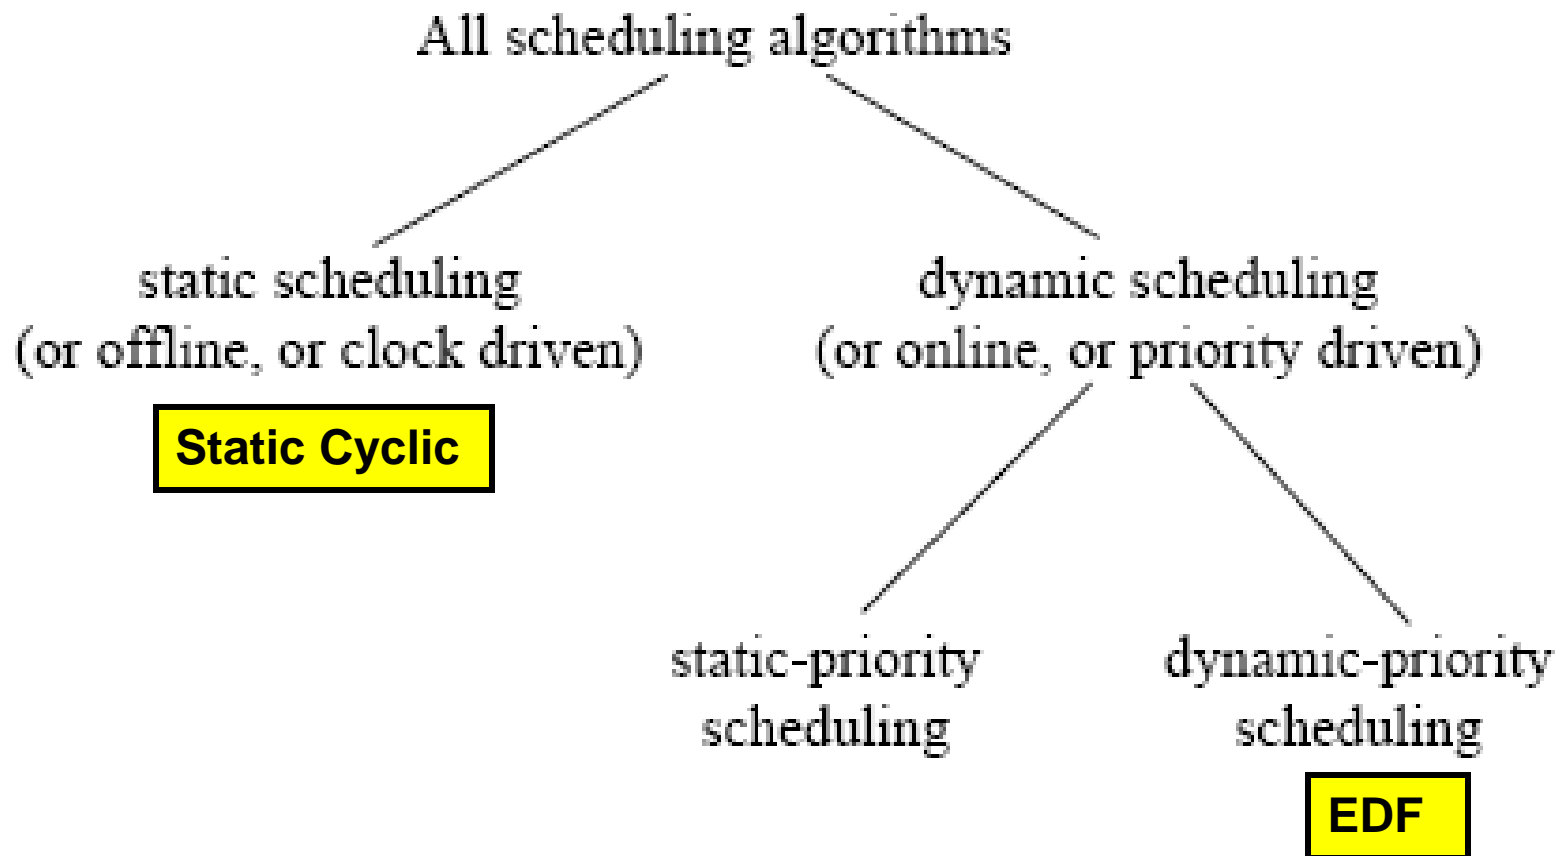
Feasible Schedule

- A valid schedule is a **feasible schedule** if every job meets its timing constraints; e.g., completes executing by its deadline.
- A set of jobs is **schedulable** according to a scheduling algorithm if (when) using the algorithm (the scheduler) always produces a feasible schedule.

Common Approaches For Real-Time Scheduling

- **Clock-Driven (Time-Driven) Approach** – scheduling decisions are made at specific time instants.
- **Weighted Round-Robin Approach** - every job joins a FIFO queue; when a job reaches the front of the queue, its weight refers to the fraction of processor time (number of time slices) allocated to the job.
- **Priority-Driven (Event-Driven) Approach** - ready jobs with highest priorities are scheduled for execution first.
 - Scheduling decisions are made when particular events occur; e.g., a job is released or a processor becomes idle. A **work-conserving** processor is busy whenever there is work to be done.

Classification of Scheduling Algorithms



EDF is optimal for scheduling preemptive tasks on a single processor.

EDF Algorithm

- **Earliest-Deadline-First (EDF) algorithm:**
 - At any time, execute the available job with the earliest deadline.
- **Theorem: (Optimality of EDF):** In a system with **one processor** and **preemption** allowed, EDF is optimal; that is, EDF can produce a feasible schedule for a given job set J with arbitrary release times and deadlines, if a feasible schedule exists.
- **Proof:** Apply schedule transformations and remove idle time.

EDF may not be optimal

- When preemption is not allowed:

$$\begin{array}{rcl} & r_i & d_i & e_i \\ J_1 & = & (0, 10, 3) \\ J_2 & = & (2, 14, 6) \\ J_3 & = & (4, 12, 4) \end{array}$$

- When more than one processor is used:

$$\begin{array}{rcl} & r_i & d_i & e_i \\ J_1 & = & (0, 4, 1) \\ J_2 & = & (0, 4, 1) \\ J_3 & = & (0, 5, 5) \end{array}$$

Periodic Task Model

- **Tasks:** T_1, \dots, T_n
- Each consists of a set of **jobs**: $T_i = \{J_{i1}, J_{i2}, \dots\}$
- ϕ_i : **phase** of task T_i = time when its first job is released
- p_i : **period** of T_i = minimum inter-release time
- H : **hyperperiod** $H = \text{lcm}(p_1, \dots, p_n)$
- e_i : **execution time** of T_i
- u_i : **utilization** of task T_i is given by $u_i = e_i / p_i$
- D_i : (relative) **deadline** of T_i , typically $D_i = p_i$

Periodic Task

- We refer to a periodic task T_i with phase ϕ_i , period p_i , execution time e_i , and relative deadline D_i by the 4-tuple (ϕ_i, p_i, e_i, D_i) .
- Example: $(1, 10, 3, 6)$
- By default, the phase of each task is 0, and its relative deadline is equal to its period.
- Example: $(0, 10, 3, 10) = (10, 3)$.

Static Cyclic Scheduling

- Periodic Task 1: (12, 3)
- Periodic Task 2: (6, 3)
- Periodic Task 3: (12, 2)

- Hyperperiod $H = \text{lcm}(12, 6, 12) = 12$
- Potential frame sizes? 3 or 6

Clock-Driven Example

\$ hi_pr < cyclic2.input > cyclic2.output

INPUT:

p max 8 12

n 1 s

n 8 t

a 1 2 3

a 1 3 3

a 1 4 3

a 1 5 2

a 2 6 6

a 2 7 6

a 3 6 6

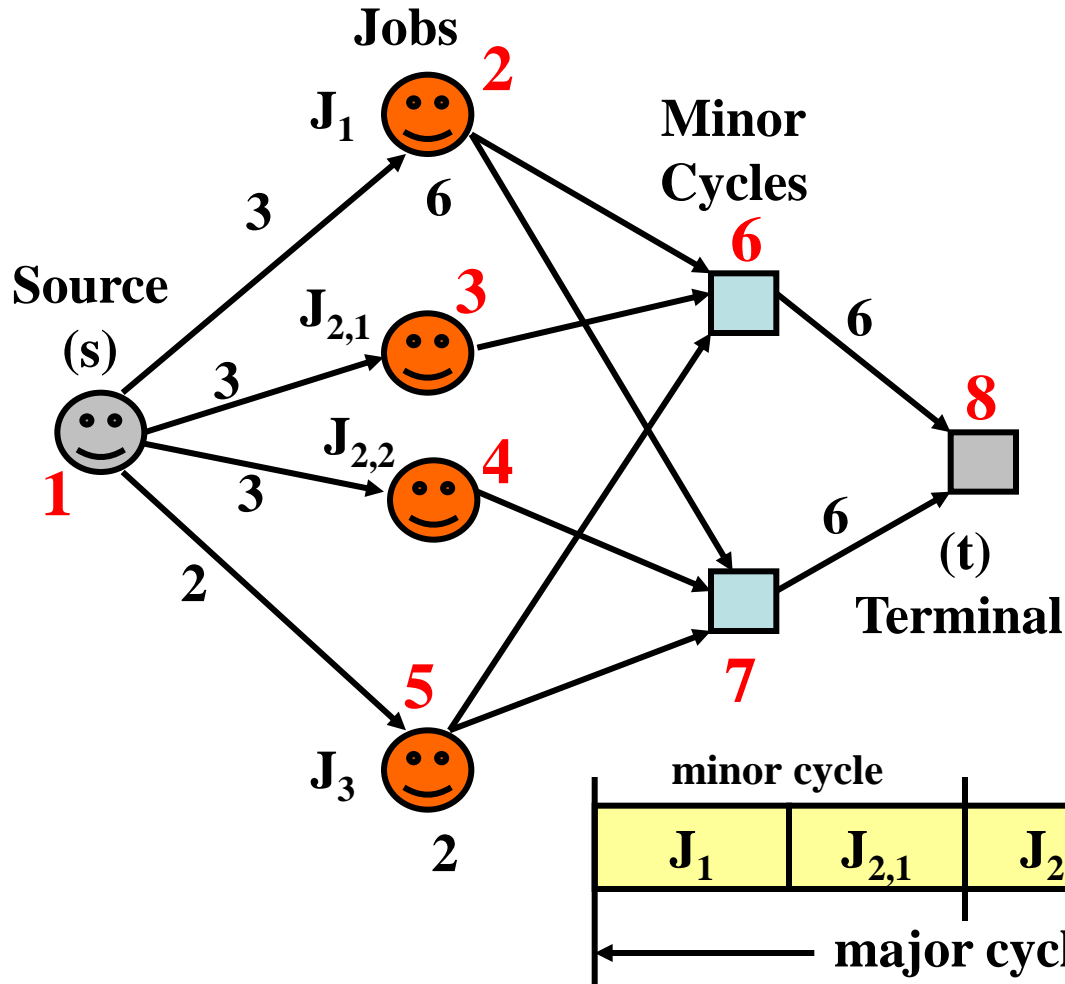
a 4 7 6

a 5 6 6

a 5 7 6

a 6 8 6

a 7 8 6



OUTPUT:

max flow:11

c flow values

f 1 2 3

f 1 4 3

f 1 3 3

f 1 5 2

f 2 6 3

f 2 7 0

f 3 6 3

f 4 7 3

f 5 7 2

f 5 6 0

f 6 8 6


f 7 8 5

c


Resulting Graphics File

```
digraph g {  
  1 -> 2 [label=3]  
  1 -> 4 [label=3]  
  1 -> 3 [label=3]  
  1 -> 5 [label=2]  
  2 -> 6 [label=3]  
  3 -> 6 [label=3]  
  4 -> 7 [label=3]  
  5 -> 7 [label=2]  
  6 -> 8 [label=6]  
  7 -> 8 [label=5]  
  { rank=same; 2; 3; 4; 5;}  
  { rank=same; 6; 7;}  
}
```

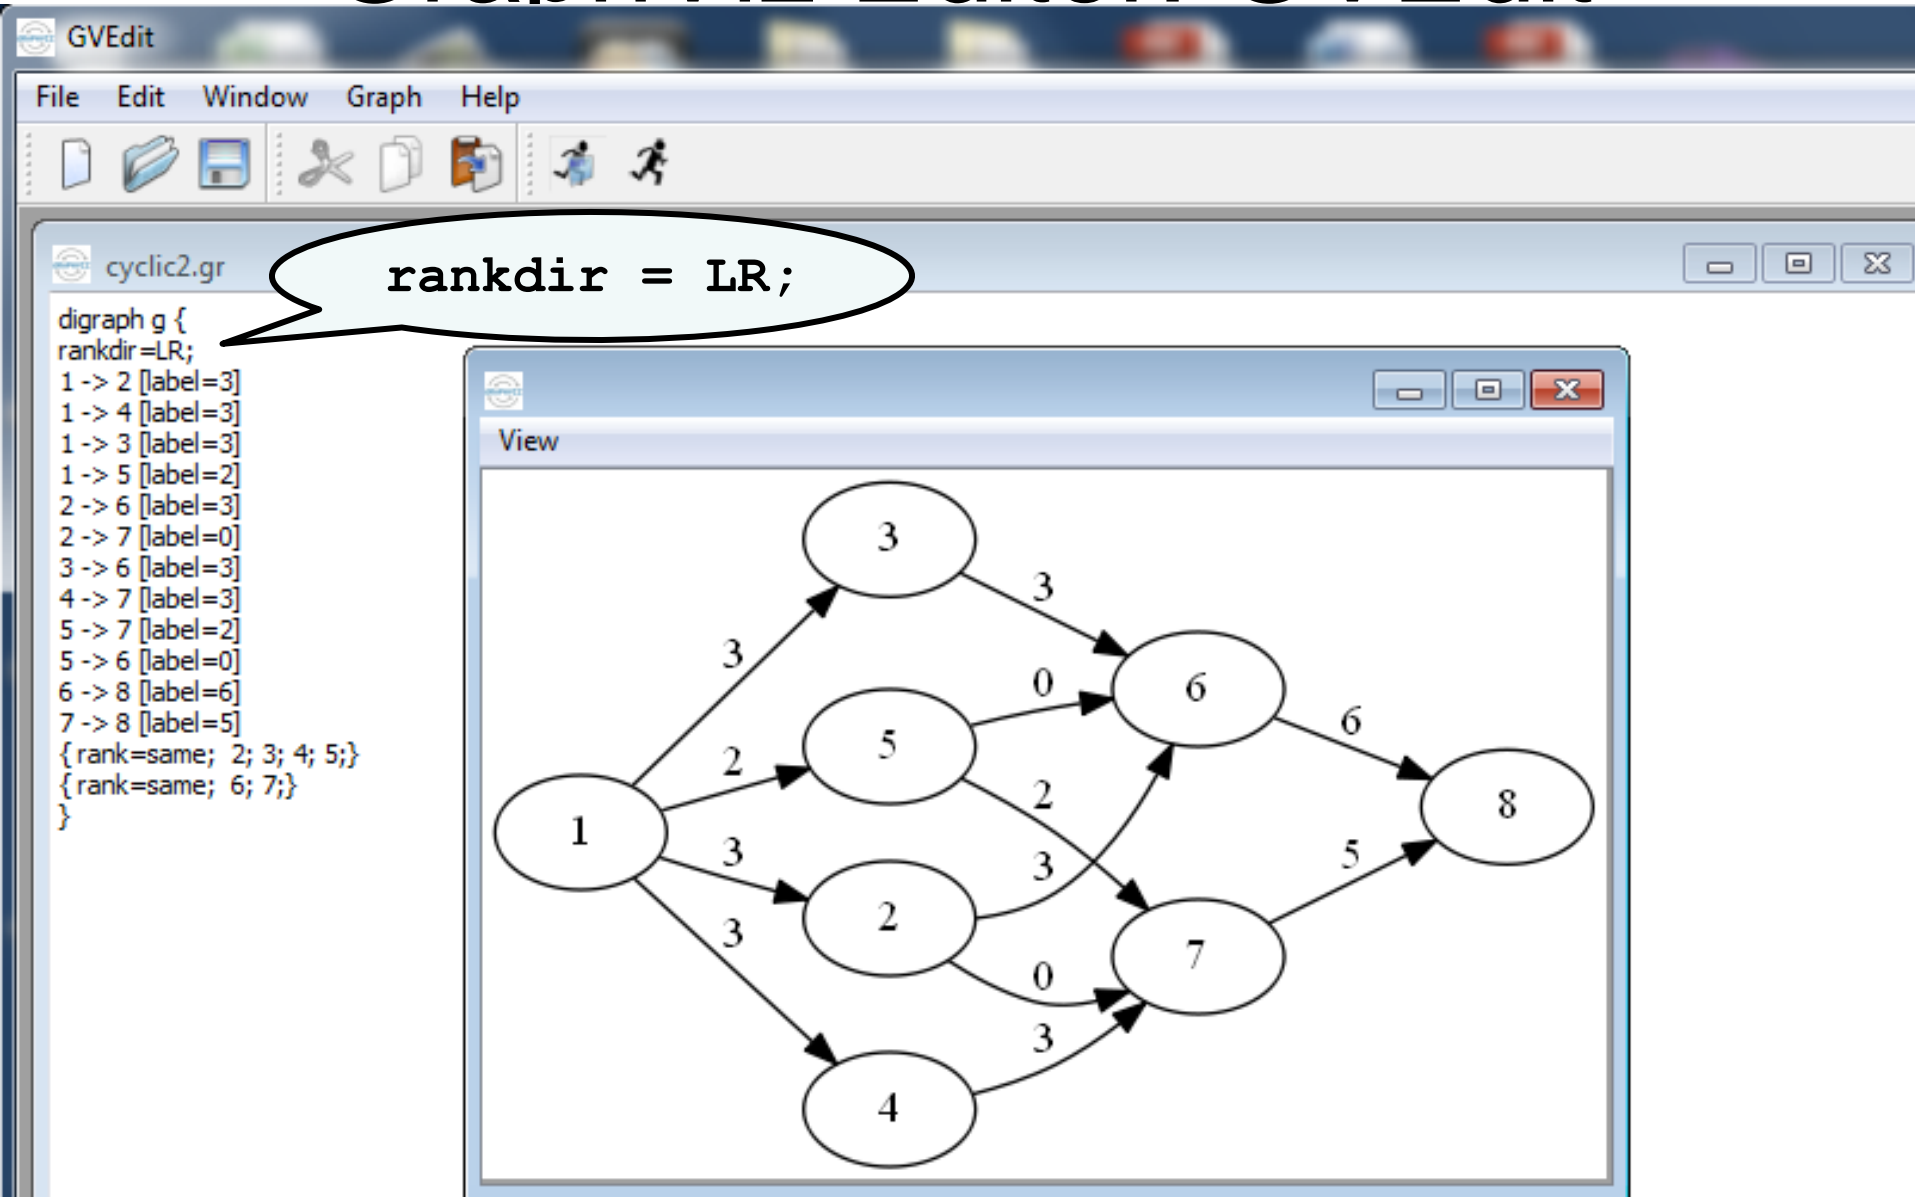
dot -Tpng fn.gr -o fn.png



Output generated




GraphViz Editor: GVEdit



Static-Priority vs. Dynamic Priority


- **Static-Priority:** All jobs in task have same priority.
- example:

Rate-Monotonic: “The shorter the period, the higher the priority.”



- **Dynamic-Priority:** May assign different priorities to individual jobs.
- example:

Earliest-Deadline-First: “The nearer the absolute deadline, the higher the priority.”



Rate-Monotonic Algorithm (RM)

- The **rate** of a task is the inverse of its period ($f_i = 1 / p_i$).
- Tasks with **higher rates (shorter periods)** are assigned **higher priorities**.
- C. L. Liu and J. W. Layland, “Scheduling Algorithms for Multiprogramming in a Hard Real-Time Environment”, JACM, Vol. 20, No. 1, pages 46-61, 1973.

Deadline-Monotonic Algorithm (DM)

- Tasks with **shorter relative deadlines** are assigned **higher priorities**.
- When tasks have relative deadlines (D_i) equal to their periods (p_i), the rate-monotonic algorithm is the same as the deadline-monotonic algorithm.

Optimal Priority Assignment

- A given priority assignment algorithm is **optimal** if whenever a task set can be scheduled by some fixed priority assignment, it can also be scheduled by the given algorithm.
- Liu and Layland show that the rate-monotonic (RM) algorithm is optimal, for preemptive, periodic task sets, with phase 0 and relative deadlines equal to their periods.


Processor Utilization

- Given a periodic task T_i , the ratio $u_i = e_i / p_i$ is called the **utilization of task T_i** .
- The **total utilization U** of all tasks in a system is the sum of the utilizations of all individual tasks:

$$U = \sum_{i=1}^n \frac{e_i}{p_i}$$


Maximum Achievable Utilization

A task set is **fully utilized** if any increase in run-time would result in a missed deadline.



Example #1

Task		Period	Deadline	Run-Time	Phase
T_i		p_i	D_i	e_i	ϕ_i
<hr/>					
A	(High Priority)	2	2	1	0
B	(Low Priority)	2	2	1	0




$U = 1/2 + 1/2 = 1.0 = 100 \%$
 The task set is fully utilized.

...

Example #2

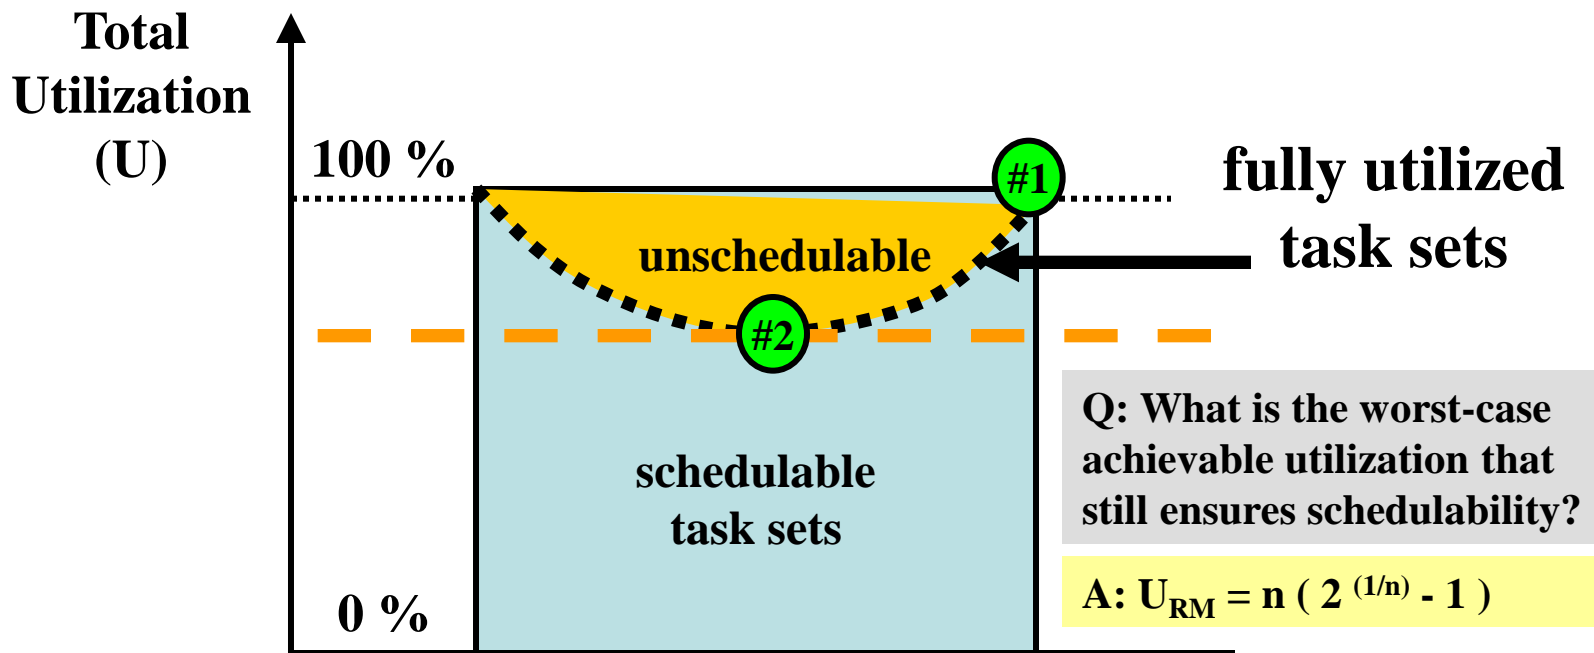
Task T_i	Period p_i	Deadline D_i	Run-Time e_i	Phase ϕ_i
<hr/>				
A (High Priority)	2	2	1	0
B (Low Priority)	3	3	1	0



$U = 1/2 + 1/3 = 0.8333$
The task set is fully utilized,
... even though $U < 1.0$.

Fully utilized task sets

A task set is **fully utilized** if any increase in run-time would result in a missed deadline.



Utilization-Based Test

- A **sufficient, but not necessary**, test that can be used to test the schedulability of a task set that is assigned priorities using the rate-monotonic algorithm, with deadlines equal to periods.
- Compute total task utilization $U(n) = U$.
- Compare with worst-case utilization bound (also called **schedulable utilization**) $U_{RM}(n) = U_{RM}$:
 - If $U > 1$, then the task set is not schedulable.
 - If $U \leq U_{RM}$, then the task set is schedulable.
 - Otherwise, no conclusion can be made.

Schedulability Analysis

- **Utilization-Based Tests**
 - Not exact (sufficient, but not necessary).
 - Not applicable to more general task models.
- **Time-Based Tests (Response Time Analysis)**
 - Use analytic approach to predict worst-case response time of each task.
 - Compare computed worst-case response times with deadlines.
 - Exact (sufficient and necessary)

Example

Task		Period	Deadline	Run-Time	Phase
τ_i		T_i	D_i	C_i	ϕ_i
<hr/>					
A	(High Priority)	7	7	3	0
B		12	12	3	0
C	(Low Priority)	20	20	5	0

- $U = 3/7 + 3/12 + 5/20 = 13/14 \approx 0.93$
- $U_{RM} = 3 (2^{1/3} - 1) \approx 0.78$
- Since $U_{RM} < U \leq 1.0$, no conclusion can be drawn using the Utilization-Based Test.

Response Time

- The **response time** (R_i) for task T_i is given by $R_i = e_i + I_i$ where:
 - e_i is the execution time of each job in T_i , and
 - I_i is the **maximum interference** caused by higher priority tasks in any interval $[t, t + R_i)$.

Maximum Interference (I_i)

- The number of releases of task T_j in $[t, t + R_i)$ is given by

$$\left\lceil \frac{R_i}{p_j} \right\rceil$$

- The interference caused by task T_j is $\left\lceil \frac{R_i}{p_j} \right\rceil * e_j$

- The maximum interference caused by all higher priority tasks is given by

$$I_i = \sum_{j \in hp(i)} \left\lceil \frac{R_i}{p_j} \right\rceil * e_j$$

where $hp(i)$ = set of all tasks with priority greater than task T_j .

Response Time Analysis

- The (**worst-case**) **response time** (w_i) for task T_i is given by the implicit equation:

$$R_i = e_i + \sum_{j \in hp(i)} \left\lceil \frac{R_i}{p_j} \right\rceil * e_j$$

- Solve by forming a recurrence relation:

$$w_i^{n+1} = e_i + \sum_{j \in hp(i)} \left\lceil \frac{w_i^n}{p_j} \right\rceil * e_j$$

$$w_i^0 = e_i$$

Solving Recurrence

- The sequence

$$w_i^0, w_i^1, w_i^2, \dots, w_i^n$$

is non-decreasing:

- If $w_i^{n+1} = w_i^n$, then a fixed point (solution) has been found.
- If $w_i^{n+1} > D_i$, then no solution exists.

Algorithm

Input: $e_1, \dots, e_m, p_1, \dots, p_m, D_1, \dots, D_m$

Output: R_1, R_2, \dots, R_m

for $i = 1$ to m

$n = 0$

$w_i^n = e_i$

 loop

$$w_i^{n+1} = e_i + \sum_{j \in hp(i)} \left\lceil \frac{w_i^n}{p_j} \right\rceil * e_j$$

 if $w_i^{n+1} = w_i^n$ then

$R_i = w_i^n$

 break out of loop { solution found }

 if $w_i^{n+1} > D_i$ then

 break out of loop { no solution }

$n = n + 1$

 end loop

end for

Response Time Analysis

- For each task, T_i , compute worst-case response time (R_i).
- If ($R_i \leq D_i$) for each task T_i , then the task set is feasible (schedulable).
- Response Time Analysis is both necessary and sufficient, and works for task sets with deadlines different than periods. If deadlines are greater than periods, then need to test jobs in task T_i over a level- i busy period and assign priorities from lowest to highest.

Preemption Thresholds Task Model

- Task Set $\Gamma = \{\tau_1, \tau_2, \tau_3, \dots, \tau_n\}$
 - Each task τ_i is characterized by (C_i, T_i, D_i) , denoted $\tau_i \sim (C_i, T_i, D_i)$.
 - Each task τ_i is assigned a priority $\pi_i \in \{1, 2, \dots, n\}$
 - and a preemption threshold $\gamma_i \in \{\pi_i, \pi_i + 1, \dots, n\}$.
- **Notes:**
 - 1 = lowest priority, n = highest priority.
 - π_i = static priority.
 - γ_i = dynamic priority.

Run-Time Model

- Modified fixed-priority, preemptive scheduling.
- When task τ_i is released, it is scheduled using its static priority π_i .
- After task τ_i starts executing, another task τ_j can preempt τ_i only if $\pi_j > \gamma_i \geq \pi_i$.


Extremes

- If $\gamma_i = \pi_i$ for each i , then the result is preemptive, priority-based scheduling.
- If $\gamma_i = n$ (max. priority) for each i , then the result is non-preemptive, priority-based scheduling.

Response Time Analysis


- The **blocking time** of task τ_i is denoted $B(\tau_i)$. Blocking occurs if a lower priority task is running and task τ_i cannot preempt it.

$$B(\tau_i) = \max_j \{ C_j / \gamma_j \geq \pi_i > \pi_j \}$$



Busy Period Analysis

- A **critical instant** occurs when all higher priority tasks arrive at the same time, and the task that contributes to the maximum blocking arrives at the critical instant - ε .



Divide Busy Period

- Divide the busy period for τ_i into two parts:
 - the length of time from the critical instant (time 0) to the point when τ_i starts executing its q^{th} job ($\mathbf{S}_i(\mathbf{q})$).
 - the length of time from the time τ_i starts executing its q^{th} job until it finishes executing its q^{th} job ($\mathbf{F}_i(\mathbf{q}) - \mathbf{S}_i(\mathbf{q})$).
- Let $q = 1, 2, \dots, m$ until we reach $q = m$ s.t. $F_i(m) \leq m T_i$ that is, the m^{th} job completes before the next job is released.
- Then,

$$R_i = \max_{q \in \{1, \dots, m\}} \{ F_i(q) - (q - 1)T_i \}$$

Worst-Case Start Time ($S_i(q)$)

$$S_i(q) = B(\tau_i) + (q-1)C_i + \sum_{\substack{j \in \{1, \dots, n\} \\ \pi_j > \pi_i}} \left(1 + \left\lfloor \frac{S_i(q)}{T_j} \right\rfloor\right) C_j$$

Worst-Case Finish Time ($F_i(q)$)

$$F_i(q) = S_i(q) + C_i + \sum_{\substack{j \in \{1, \dots, n\} \\ \pi_j > \gamma_i}} \left(\left\lceil \frac{F_i(q)}{T_j} \right\rceil - \left(1 + \left\lfloor \frac{S_i(q)}{T_j} \right\rfloor \right) \right) C_j$$

Algorithm to compute R_i

Input: $C_1, \dots, C_m, T_1, \dots, T_m, \pi_1, \dots, \pi_m, \gamma_1, \dots, \gamma_m$

Output: R_1, R_2, \dots, R_m

done = FALSE

q = 1

while (not done)

 compute $S_i(q)$ and $F_i(q)$

 if $F_i(q) \leq q T_i$ then

 done = TRUE

 m = q

 else

 q = q + 1

 end if

end while

$R_i = \max_{q \in \{1, \dots, m\}} (F_i(q) - (q - 1) T_i)$

Preemption Threshold Assignment

- Assign preemption thresholds from lowest priority to highest priority.
- Only elevate preemption threshold as much as needed to make a given task feasible.

Summary

- Quiz #1
 - Fri., Oct. 16, in class
 - Open book, open notes
 - 50 minutes