

Device Independent I/O (Block Special File)

- Read the following sections in the text: 5.6.2 (pp550-553), 5.6.3 (pp553-555), 5.6.4 (pp555-557)
- Refer to disk layout in figure 5-34 on p551
- Refer to the super block in Figure 5-35 on p552
- Refer to the i-node structure in Figure 5-36 on p556
- Refer to zone[] information in an i-n inode in Figure 5-11 on p503



Device Independent I/O (Block Special File) (cont)

- `mknod` - make block or character special files
- `mknod` [*OPTION*]... *NAME TYPE* [*MAJOR MINOR*]
 - Create the special file *NAME* of the given *TYPE*.
 - *OPTION*:
 - **-Z**, **--context=CONTEXT** set security context (quoted string)
 - **-m**, **--mode=MODE** set permission mode (as in `chmod`),
 - *TYPE* may be:
 - `b` create a block (buffered) special file
 - `c`, `u` create a character (unbuffered) special file
 - `p` create a FIFO
 - Both *MAJOR* and *MINOR* must be specified when *TYPE* is `b`, `c`, or `u`, and they must be omitted when *TYPE* is `p`.



Major Device Number and Minor Device Number

- The major device number determines which device driver process a request message is sent to
 - In Unix, the major device number determines which function is to be invoked
- The minor device number is passed in a request message to the driver process
 - In Unix, it is passed as an argument to the function to be invoked
 - It's up to the device driver process/function as to how to use the minor device number



Device Independent I/O (Block Special File) (cont)

```
struct dmap (4220 on p686)
04220 extern struct dmap {
04221     int _PROTOTYPE ((*dmap_opcl), (int, Dev_t, int, int) );
        // points to gen_opcl
04222     void _PROTOTYPE ((*dmap_io), (int, message *) );
        // request for actually I/O (points to gen_io)
04223     int dmap_driver;
        // the device driver process to send a request message to
04224     int dmap_flags;
04225 } dmap[];
```

- Initialization for a RAM driver is done at pp28235-28238
- Initialization for a FLOPPY driver is (probably) done at 28204, 28209, 28210 (process number 6 ?)
 - dm->dmap_opcl = gen_opcl;
 - dm->dmap_io = gen_io;
 - dm->dmap_driver = MEM_PROC_NR; // for the RAM driver
 - dm->dmap_flags = 0; // for the RAM driver



Open Operation on Block Special File

- `do_open` (24550)
 - calls `common_open` at 24566
- In `common_open` (24573)
 - calls `dev_open` at 24640 for `I_BLOCK_SPECIAL` (`I_CHAR_SPECIAL`, too) based on `i-node.i_mode` information
- In `dev_open` (28334)
 - calls `(*dp->dmap_opcl)(DEV_OPEN, dev, proc, flags) DEV_OPEN` at 28349
 - Recall `dp->dmap_opcl` points to "`gen_opcl`"
- (cont)



Open Operation on Block Special File (cont)

- In `gen_opcl` (28455) (note that `op` is `DEV_OPEN`)
 - identifies the `dmap` structure for the major device number in `dp` at 28466
 - calls `(*dp->dmap_io)(dp->dmap_driver, &dev_mess)` at 28474
 - `dp->dmap_io` points to "`gen_io`" (28575)
 - `dp->dmap_driver` is set to the task to send a message to (`MEM_PROC_NR` for the RAM driver)
 - `dev_mess` contains: `op` (`= DEV_OPEN`), minor device number, caller (the process that requested this operation) process number, and flag (`=0?`)
- In `gen_io` (28575),
 - sends a message to `task_nr` (`=dp->dmap_driver`) at 28593



Read/Write Operation on Block Special File

- `do_read(25030)/do_write (25627)`
 - calls `read_write (25032, 25629)`
- In `read_write (25038)`
 - calls `rw_chunk` at 25173) with i-node, in case of `I_BLOCK_SPECIAL`
- In `rw_chunk (25251)`
 - `inode->i_zone[0]` is set to dev (major and minor dev numbers) at 25278
 - then calls `get_block` with dev at 25303
- In `get_block (22426)`
 - calls `rw_block (22641)` with ptr to buf at 22511, where struct buf (21616) contains dev
- In `rw_block (22641)`
 - calls `dev_io` at 22661 with `op = DEV_READ` or `DEV_WRITE` (see 22660)
- (cont)



Read/Write Operation on Block Special File (cont)

- In `dev_io` (28406) (`op = DEV_READ` or `DEV_WRITE`, `dev` – major and minor device numbers)
 - identifies the `dmap` structure for the major device number in `dp` at 28420
 - calls `dp->dmap_io(dp->dmap_driver, &dev_mess)` at 28432, where
 - `dp->dmap_io` is set to "`gen_io`"
 - `dp->dmap_driver` is the process to send a request to (`MEM_PROC_NR` for the RAM driver)
 - `dev_mess` contains `op` (`DEV_READ` or `DEV_WRITE`), the minor device number, position, who (the requester), buffer (address), and the number of bytes
- In `gen_io` (28575)
 - sends a message to the process (`MEM_PROC_NR` for the RAM driver) at 28593



Important Data Structures (Summary)

```
struct device { // 10857
    u64_t dv_base; // address of the data area of the device
    u64_t dv_size;
};
```

```
typedef struct { // 2856 client provides an array of this structure in scattered I/O
    vir_bytes iov_addr; /* address of an I/O buffer in the client */
    vir_bytes iov_size; /* sizeof an I/O buffer */
} iovec_t;
```

```
struct partition { // in src/include/minix/partition.h
    u64_t base; /* byte offset to the partition start */
    u64_t size; /* number of bytes in the partition */
    unsigned cylinders; /* disk geometry */
    unsigned heads;
    unsigned sectors;
};
```



sys_vircopy and sys_physcopy (p23)

