

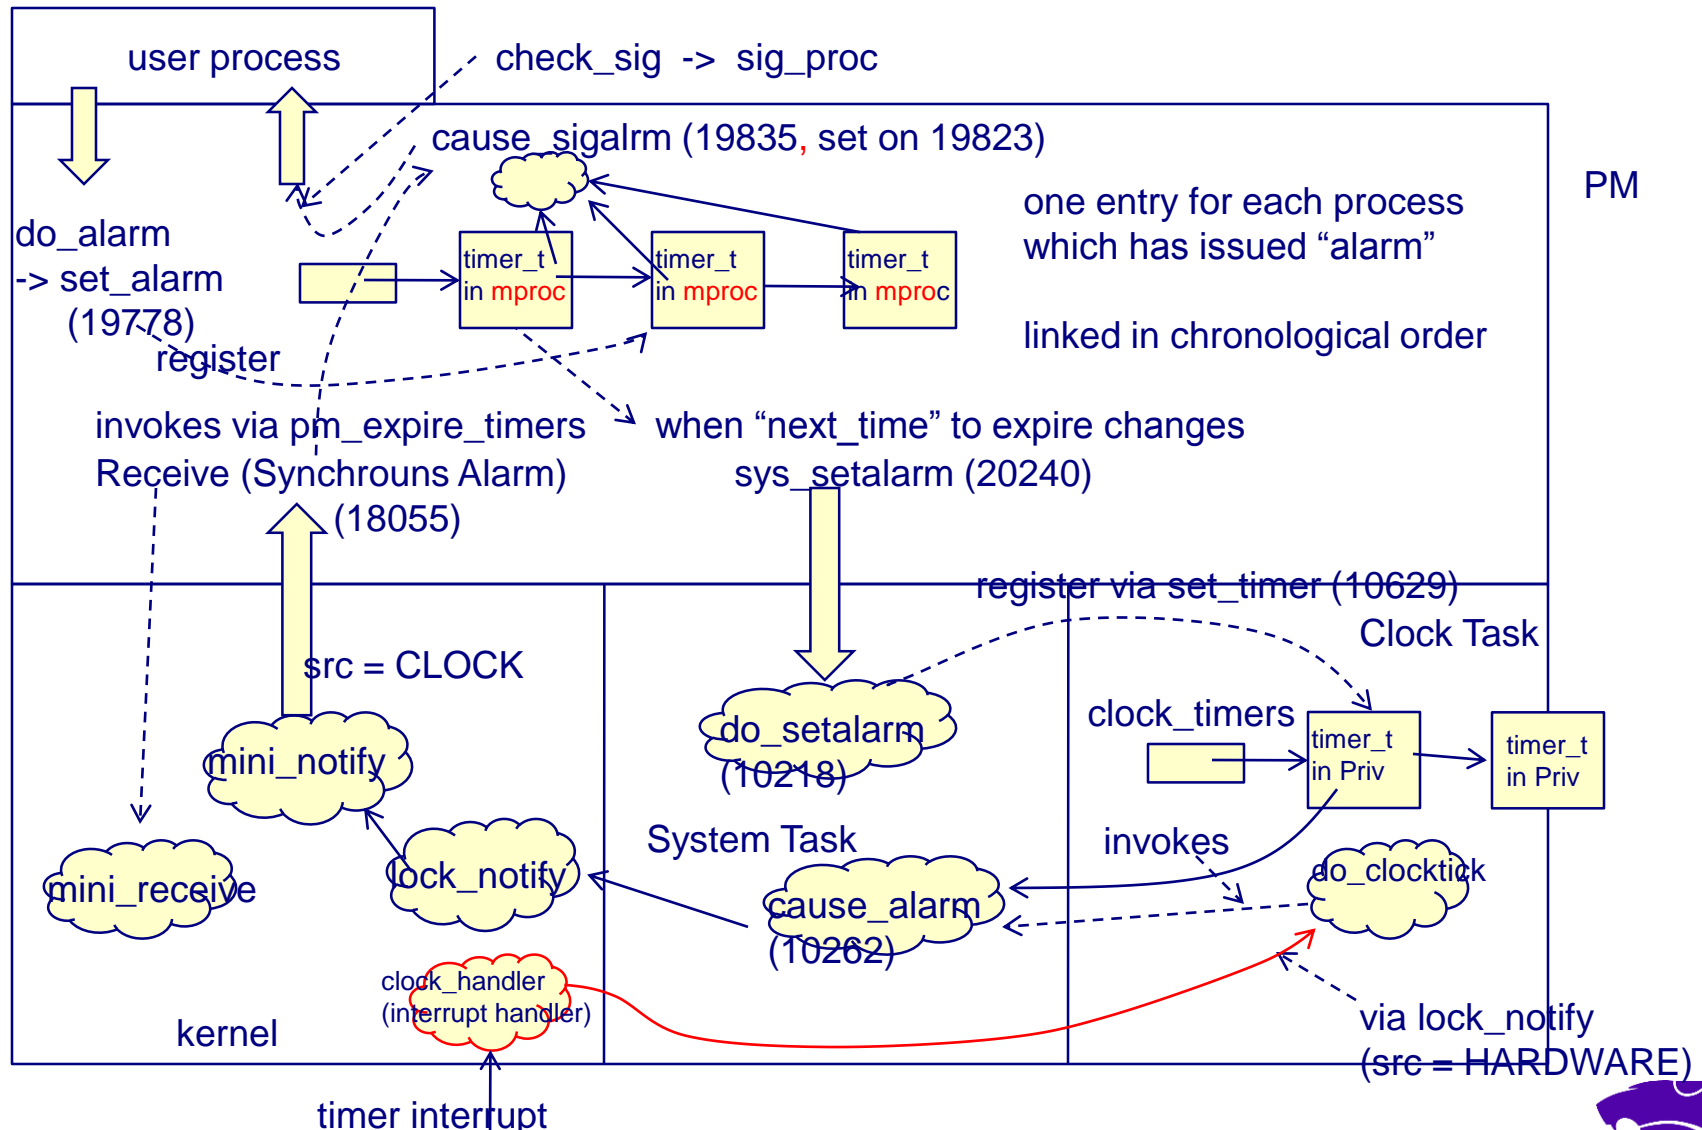
Three Ways to Deliver Signals

1. A process issues the kill system call
 - handled by `do_kill` (19689)
 2. A clock timer expires
 - A notification message is from `cause_alarm` (10970)
 - `call_nr = SYM_ALARM` is sent to PM and `pm_expire_timers` (20256) handles this case
 - for each expired alarm, it calls `cause_sigalarm` (19835)
 - `cause_sigalarm` calls `check_sig` with `SIGALRM`
 3. The kernel generates a signal (segmentation violation, `SIGINT`, `SIGQUIT`, `SIGKILL` etc) and sends a notification message to PM that contains pending signals (`SYS_SIG`)
- See the following define statements and recall notify messages (7458~7475)

```
03655 #define NOTIFY_MESSAGE      0x1000
03656 #define NOTIFY_FROM(p_nr)    (NOTIFY_MESSAGE | ((p_nr) + NR_TASKS))
03657 #define SYN_ALARM    NOTIFY_FROM(CLOCK)    /* synchronous alarm */
03658 #define SYS_SIG      NOTIFY_FROM(SYSTEM)   /* system signal */
```



Synchronous Alarm (revisited: see Chapter 6)



Handling Kernel Generated Signals

- Kernel calls `cause_sig` (9949)
 - e.g., see line 8059 in exception (8012)

Kernel (SYSTEM) side

`cause_sig (dest_proc, sig_no) (9949)` \Rightarrow add `sig_no` to `dest_proc`'s set of pending signals
↓ (9975) set `dest_proc`'s `SIGNALED`, `SIG_PENDING` and block it (STEP A)

`send_sig(PM, SIGKSIG) (9931)`
SIGKSIG is added to PM's set of pending signals in the message
↓ (9943)

`lock_notify(SYSTEM, PM) (7758)`
message with `call_nr = SYS_SIG` is sent to PM

signaled process is blocked during this period



Handling Kernel Generated Signals (cont)

In PM

PM calls `ksig_pending()` (19699) since `SIGKSIG` is set (see 18058~18060)

`ksig_pending` (19699)

`sys_getksig(&proc, &sig_map)` →

`handle_sig(proc, sigmap)` (19729)

PM becomes the sender (generator) of the signal and for each signal `s` in the pending signal set, it calls `check_sig(id, s)` for process id (id may be 0 (`SIGINT` `SIGQUIT`) or -1 (`SIGKILL`), or a single process) (`check_sig` calls `sig_proc` to deliver the signal)

`sys_endksig(proc)` →

SYSTEM

`do_getksig` (kernel/system/do_getksig)
reset `SIGNALED` of the proc (STEP B)

If a new signal arrives at the signaled process during this period, the signaled process does not become ready

`do_endksig` (kernel/system/do_endksig)
if `SIGNALED` is still reset, reset `SIG_PENDING` and unblock proc (STEP C)



Interpretation of handling kernel generated signals

- Multiple kernel generated signals may be delivered and handled by a process at once
- In `ksig_pending`, PM sends all kernel generated signals to all processes that have pending kernel signals (see the while statement at 19714)
- The system wants to deliver a kernel generated signal to a process immediately
 - the process should not proceed when the kernel generates a signal to that process
 - Therefore, the kernel blocks the process to be signaled at STEP A
 - The process is unblocked when all the kernel generated signals have been sent (its stack and PCB have been modified) at STEP C

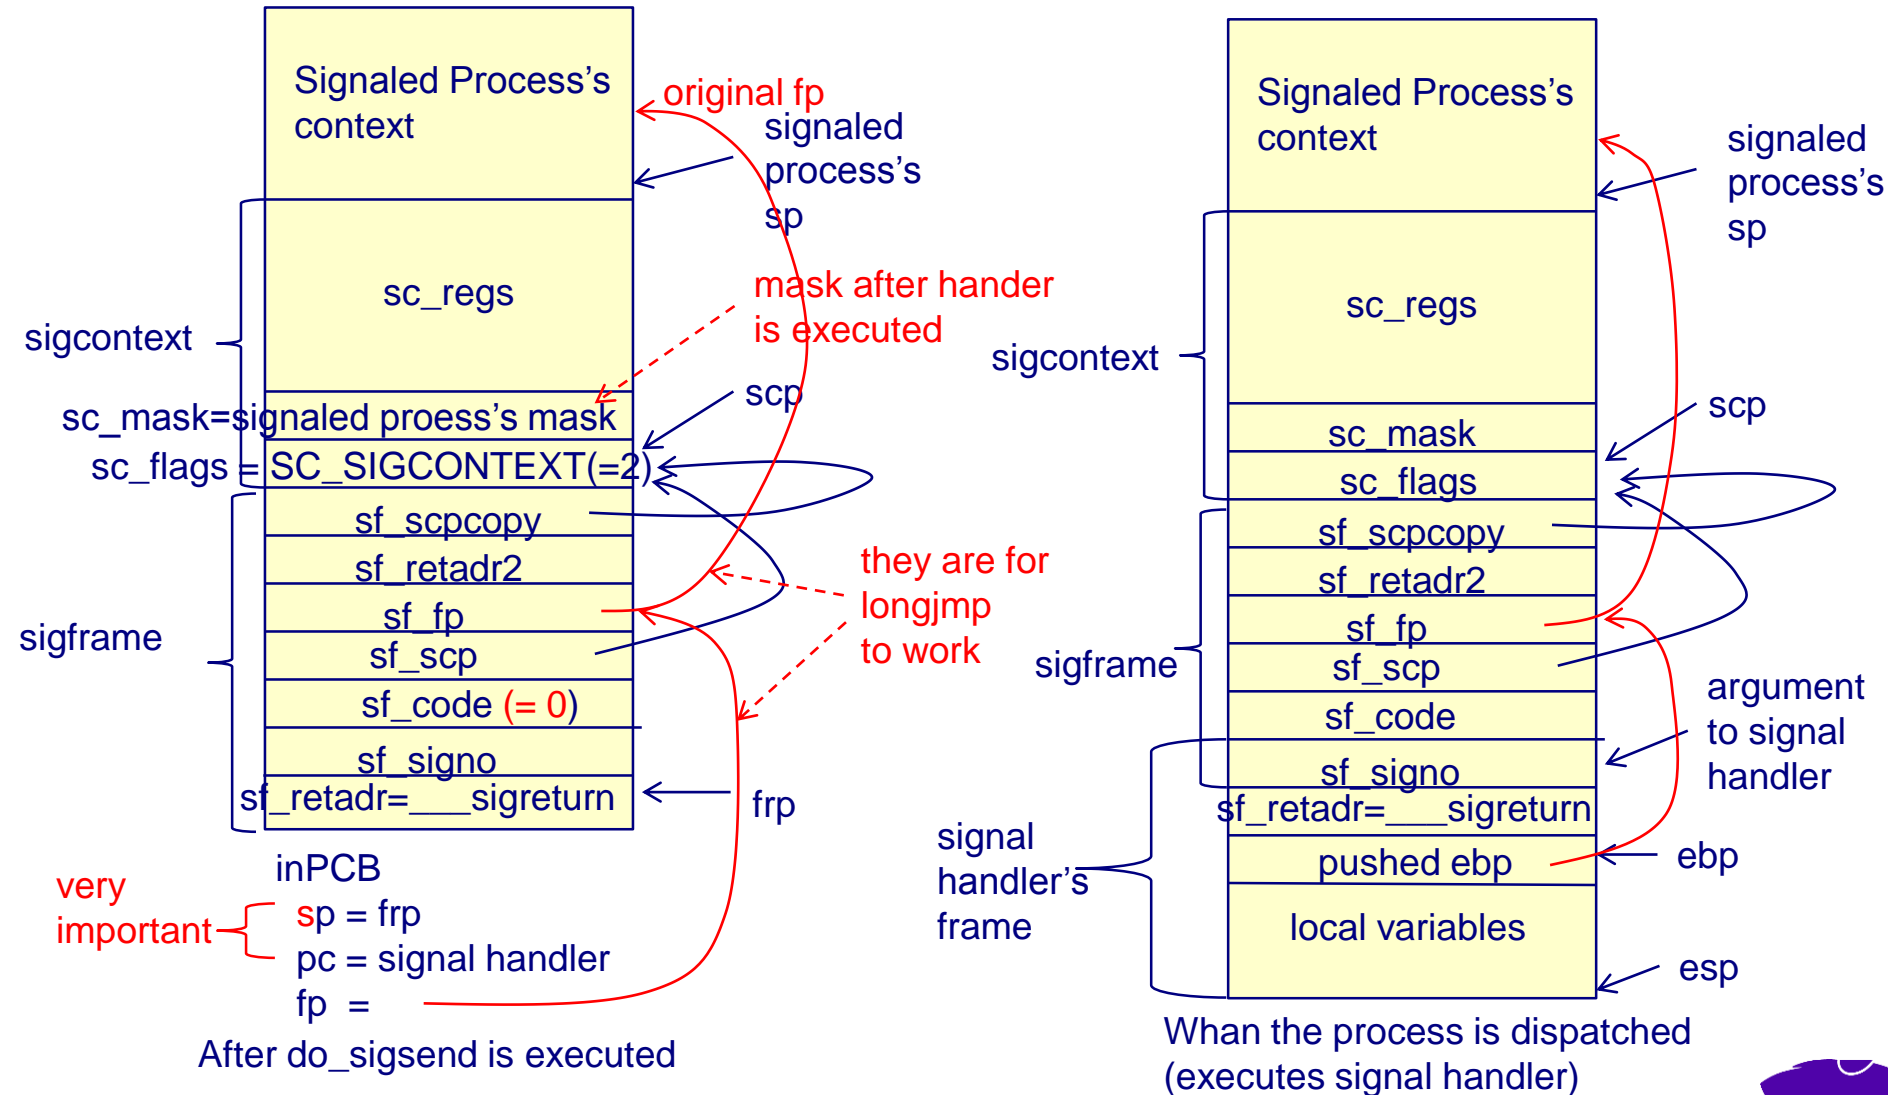


Interpretation of handling kernel generated signals (cont)

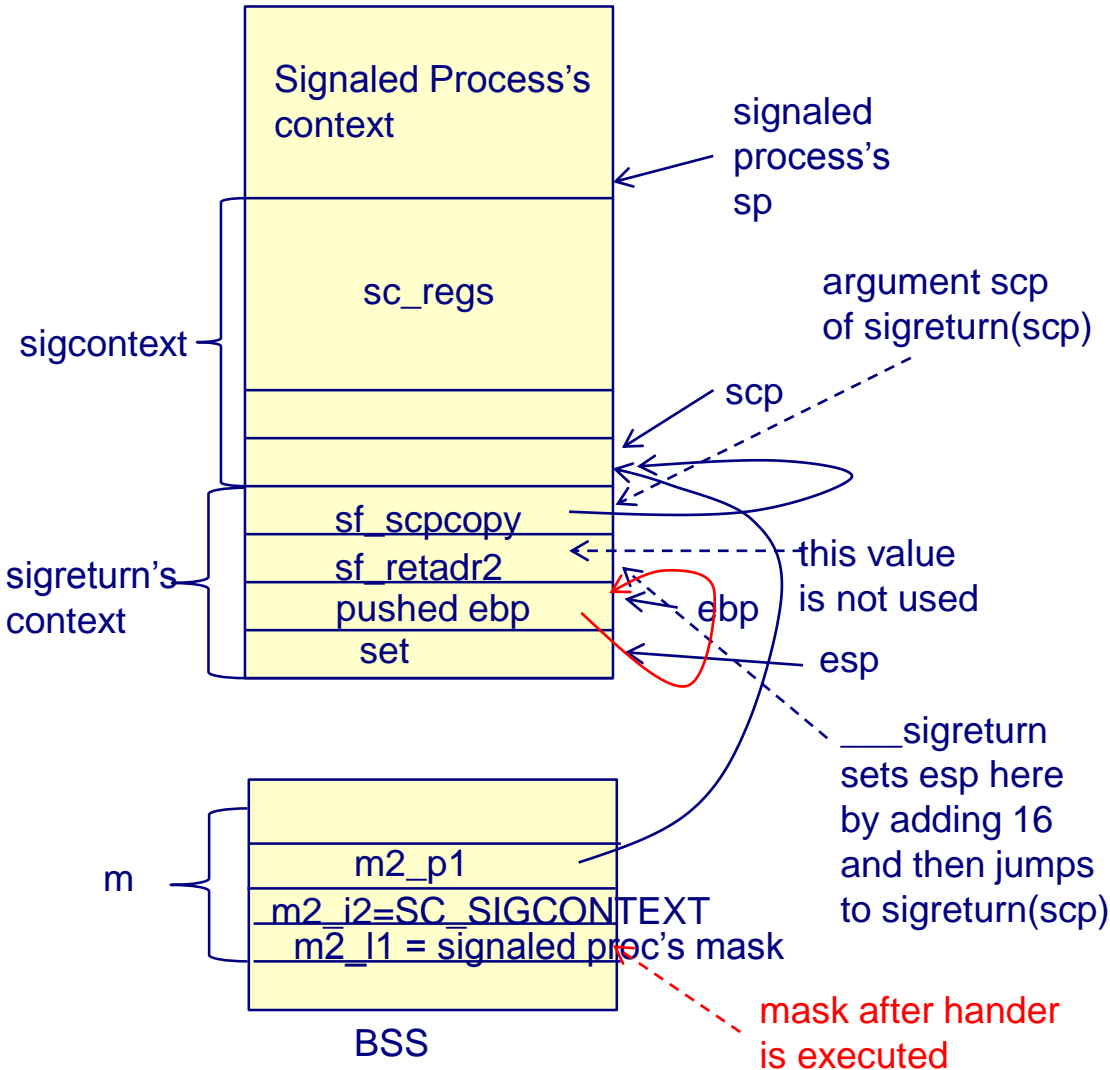
- Flag SIGNAL indicates that the kernel has generated a signal to a process (STEP A)
- When the system picks up a process to deliver pending kernel signals at STEP B, the process's SIGNAL flag is reset
- While PM is sending a signal to a process (changing its stack and PCB), the kernel may generate another signal to the process. Then, the process's SIGNAL flag is set again by STEP C
- Thus, in `do_endksig()`, the system unblocks the process only when its SIGNAL is still reset
 - If SIGNAL is set again, the process handles all (previous ones and the newly generated signals) signals (executes their signal handlers) next time when PM receives a message at 18052



Signal Processing



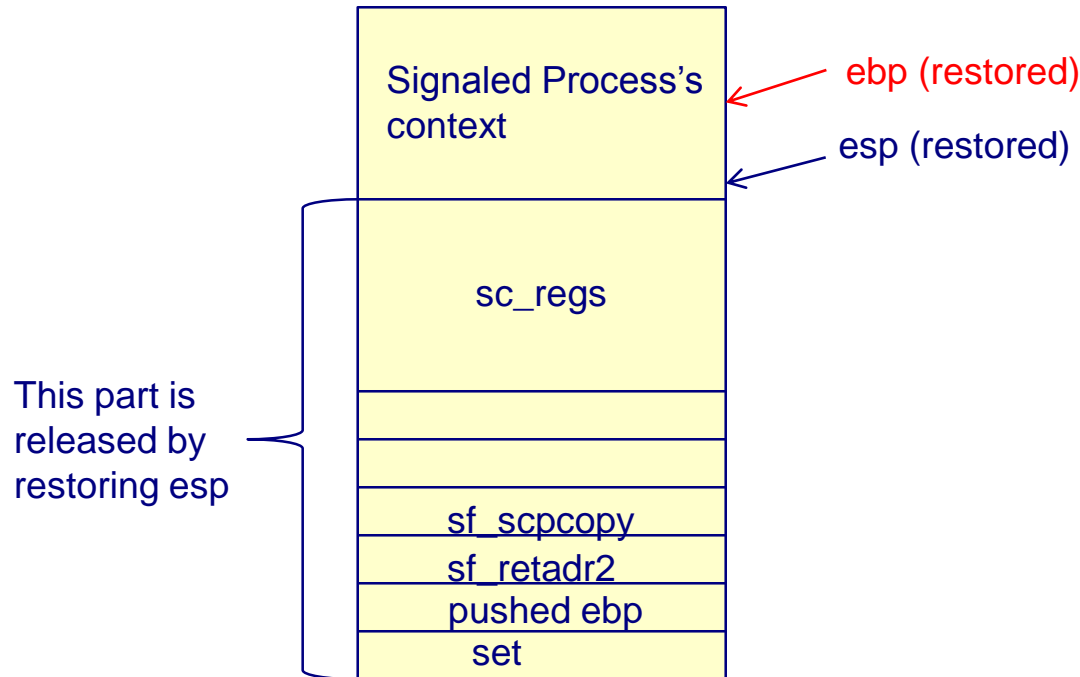
Signal Processing (cont)



- After the signal handler completes, it “returns” to `___sigreturn`
- `___sigreturn` adds 16 to `esp` and jumps to `sigreturn(csp)`
- `sigreturn(scp)` sets message `m` and passes the control to PM’s `do_sigreturn()`
- PM’s `do_sigreturn()` passes the control to Kernel’s `do_sigreturn` with pointer to `sigcontext` [by casting it to `(struct sigmsg *)`]



Signal Processing (cont)



- Kernel's `do_sigreturn` copies sigcontext to its local variable `sc`
- Then, it restores `sc's sc_regs` to the signaled process's PCB (including `eip (pc)` and `esp`)
 - their values are set in 2 slides back
- When the process is dispatched, it resumes the execution from the point of signal
 - except for `psw` (`psw` holds the value at the end of the execution of the handler)

check if this is semantically ok

