

**CIS 833 – Information Retrieval and Text Mining**

**Lecture 14**

# Language Models

October 20, 2015

Credits for slides: Hofmann, Mihalcea, Mobasher, Mooney, Schutze.

## Required Reading

Information Retrieval textbook

- Chapter 12: Language models for IR

## Query Generation Model

- How might a query look like that would ask for a specific document?
  - A common search heuristic is to use words that you expect to find in matching document(s) as your query.
  - The LM approach directly exploits that idea
    - A document is a good match to a query if the document model is likely to generate the query, which will in turn happen if the document contains the query words often.

*think of a relevant document & formulate a query by picking some of the keywords as query terms.*

environment logging ban

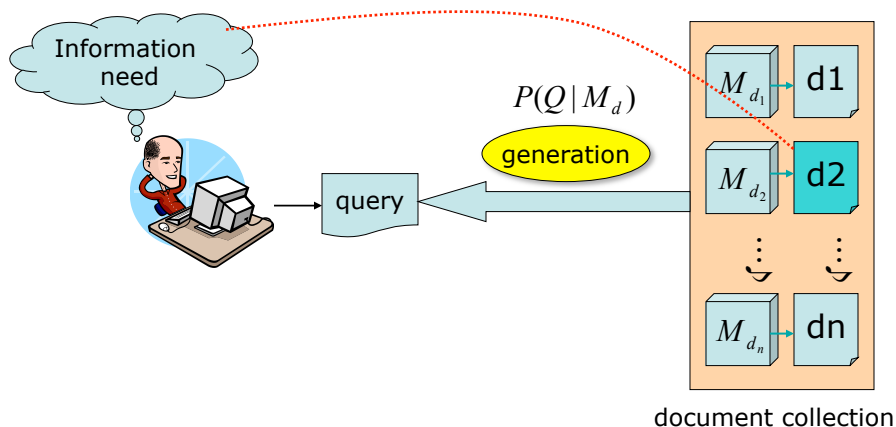
Google Search

I'm Feeling Lucky




environmentalists are blasting a Bush administration proposal to lift a ban on logging in remote areas of national forests, saying the move ignores popular support for protecting forests.

## IR based on Language Model (LM)




## Formal Language (Model)

- Traditional generative model: generates strings
  - Finite state machines or regular grammars, etc.
- Example:



## Formal Language (Model)


- Traditional generative model: generates strings
  - Finite state machines or regular grammars, etc.
- Example:



I wish  
I wish I wish  
I wish I wish I wish  
I wish I wish I wish I wish  
...

## Formal Language (Model)

- Traditional generative model: generates strings
  - Finite state machines or regular grammars, etc.
- Example:



I wish  
 I wish I wish  
 I wish I wish I wish  
 I wish I wish I wish I wish  
 ...  
 wish I wish ???

## Stochastic Language Model



Models *probability* of generating strings in the language (commonly all strings over alphabet  $\Sigma$ )

Model M

0.2	the	the	man	likes	the	woman
0.1	a					
0.01	man	0.2	0.01	0.02	0.2	0.01
0.01	woman					
0.03	said					
0.02	likes					
...						

multiply

$$P(s | M) = 0.00000008$$


## Stochastic Language Models

Model M1	Model M2					
0.2 the	0.2 the	the	class	pleaseth	yon	maiden
0.01 class	0.0001 class	—	—	—	—	—
0.0001 sayst	0.03 sayst	0.2	0.01	0.0001	0.0001	0.0005
0.0001 pleaseth	0.02 pleaseth	0.2	0.0001	0.02	0.1	0.01
0.0001 yon	0.1 yon					
0.0005 maiden	0.01 maiden					
0.01 woman	0.0001 woman					

$$P(s|M2) > P(s|M1)$$

## Stochastic Language Models

- A statistical model for generating text
  - Probability distribution over strings (words) in a given language
  - How do we calculate probabilities over sequences of strings?

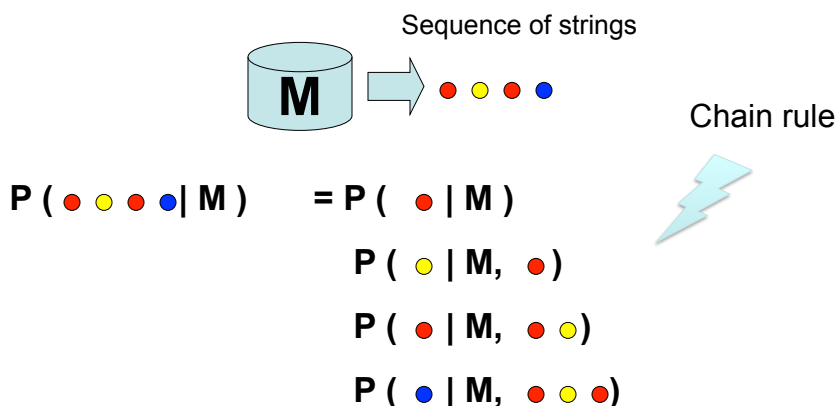


$$P(\text{red, yellow, red, blue} | M) = ???$$

Chain rule?

## Stochastic Language Models

- A statistical model for generating text
  - Probability distribution over strings (words) in a given language



## Unigram and higher-order models

$$P(\text{red yellow red blue}) \\ = P(\text{red}) P(\text{yellow} | \text{red}) P(\text{red} | \text{red yellow}) P(\text{blue} | \text{red yellow red})$$

- Unigram Language Models

$$P(\text{red}) P(\text{yellow}) P(\text{red}) P(\text{blue})$$

Easy,  
Effective!

- Bigram (generally,  $n$ -gram) Language Models

$$P(\text{red}) P(\text{yellow} | \text{red}) P(\text{red} | \text{yellow}) P(\text{blue} | \text{red})$$

- Other Language Models

- Grammar-based models (probabilistic context-free grammars or PCFGs), etc.
- Used for speech recognition, spelling correction, machine translation, etc.

## Language Models for IR


- Treat the generation of queries as a random process.
  - Users have a reasonable idea of terms that are likely to occur in documents of interest.
  - They will choose query terms that distinguish these documents from others in the collection.

## Query Likelihood Model

- Treat each document  $d$  as the basis for a model  $M_d$  (e.g., unigram language model)
  - Infer a language model for each document.
- Estimate the probability of generating the query according to each of these models
$$P(M_d | Q) = P(Q | M_d) \times P(M_d) / P(Q)$$
  - $P(Q)$  is the same for all documents, so ignore
  - $P(M_d)$  [the prior] is often treated as the same for all documents
    - But we could use criteria like authority, length, genre
  - $P(Q | M_d)$  is the probability of  $Q$  given  $d$ 's model  $M_d$
- Rank documents  $d$  based on  $P(M_d | Q)$   
(or equivalently,  $P(Q | M_d)$ )
- Very general formal approach

## Ranking

- Documents are ranked by the probability that a query would be observed as a random sample from the respective document model.



$$P(Q | M_d) = \prod_{w \in Q} P(w | M_d) = \prod_{\substack{w \text{ distinct} \\ \text{words in } Q}} P(w | M_d)^{q_w}$$

## Query generation probability

- The probability of producing the query given the language model of document  $d$  (using maximum likelihood estimate) is:

$$P(Q | M_d) = \prod_{w \in Q} P(w | M_d)$$

$$= \prod_{w \in Q} \frac{tf_{(w,d)}}{dl_d}$$

Unigram assumption:  
Given a particular language model, the query terms occur independently

$M_d$  : language model of document  $d$

$tf_{(w,d)}$  : raw tf of word  $w$  in document  $d$

$dl_d$  : total number of tokens in document  $d$



## Insufficient data

- Zero probability  $P(w | M_d) = 0$ 
  - May not wish to assign a probability of zero to a document that is missing one or more of the query terms [gives strict conjunctive semantics]
  - Need to smooth probabilities
  - There's a wide space of approaches to smoothing probability distributions to deal with this problem, such as adding 1,  $\frac{1}{2}$  or  $\epsilon$  to counts, Dirichlet priors, discounting, and interpolation
- General approach here
  - A non-occurring term is possible, but no more likely than would be expected by chance in the collection.
  - If  $tf_{(w,d)} = 0$  then  $P(w | M_d) = \frac{cf_w}{cs}$ 
    - $cs$  : raw count of word  $w$  in the collection
    - $cf_t$  : raw collection size (total number of tokens in the collection)

## Mixture model

- A simple idea that works well in practice is to use a mixture between the document multinomial and the collection multinomial distribution
$$P(w|d) = \lambda P(w|M_d) + (1 - \lambda)P(w|M_c)$$
  - *linear interpolation language model*
- Mixes the probability from the document with the general collection frequency of the word
- Correctly setting  $\lambda$  is very important
- A high value of lambda makes the search “conjunctive-like” – suitable for short queries
- A low value is more suitable for long queries
- Can tune  $\lambda$  to optimize performance
  - Perhaps make it dependent on document size (cf. Dirichlet prior or Witten-Bell smoothing)

## Basic mixture model summary

- General formulation of the LM for IR

$$p(Q | d) = \prod_{t \in Q} ((1 - \lambda)p(t | M_c) + \lambda p(t | M_d))$$

general language model

individual-document model

- The user has a document in mind, and generates the query from this document.
- The equation represents the probability that the document that the user had in mind was in fact this one.

## Collection Statistics in IR

- Collection statistics ...
  - Are integral parts of the language model.
  - Are not used heuristically as in many other approaches.
    - At least in theory...

## Example

- Document collection (2 documents)
  - $d_1$ : Xerox reports a profit but revenue is down
  - $d_2$ : Lucent narrows quarter loss but revenue decreases further
- Model: MLE unigram from documents;  $\lambda = \frac{1}{2}$
- Query: *revenue down*
  - $P(Q|d_1) = ?$
  - $P(Q|d_2) = ?$
- Ranking: ?

$$p(Q|d) = \prod_{t \in Q} ((1 - \lambda)p(t|M_c) + \lambda p(t|M_d))$$

## Language models: pros & cons

- Novel way of looking at the problem of text retrieval based on probabilistic language modeling
  - Conceptually simple and explanatory
  - Formal mathematical model
  - Natural use of collection statistics, not heuristics (almost...)
- LMs provide effective retrieval and can be improved to the extent that the following conditions can be met
  - Our language models are accurate representations of the data.
  - Users have some sense of term distribution.

## LM vs. Prob. Model for IR

- The main difference is whether “Relevance” figures explicitly in the model or not
  - LM approach attempts to do away with modeling relevance
- LM approach assumes that documents and expressions of information problems are of the same type
- Computationally tractable, intuitively appealing

## LM vs. Prob. Model for IR

- Problems of the basic LM approach
  - Assumption of equivalence between document and information problem representation is unrealistic
  - Very simple models of language
  - Relevance feedback is difficult to integrate, as are user preferences, and other general issues of relevance
  - Can't easily accommodate phrases, passages, Boolean operators
- Current extensions focus on putting relevance back into the model, etc.

## Comparison With Vector Space

- There's some relation to traditional TF-IDF models:
  - terms often used as if they were independent
  - (unscaled) term frequency is directly in model
  - the probabilities do length normalization of term frequencies
  - the effect of doing a mixture with overall collection frequencies is a little like idf: terms rare in the general collection but common in some documents will have a greater influence on the ranking

## Comparison With Vector Space

- Differences
  - Based on probability rather than similarity
    - Intuitions are probabilistic rather than geometric
  - Details of use of document length and term, document, and collection frequency differ
- Recent work has shown that LM approach performs better than TF-IDF and other retrieval models.